

Sequential Parameter Optimization Cookbook

Thomas Bartz-Beielstein

Jul 1, 2026

Table of contents

Preface	1
I. Sequential Parameter Optimization Toolbox (SPOT)	3
1. Introduction to SpotOptim	5
2. Examples	7
2.1. Complex Constrained Optimization: The Robot Arm	9
2.2. Jupyter Notebook	10
3. SpotOptim API Reference	11
4. SpotOptim.minimize	13
4.1. API	13
4.2. Parameters	13
4.3. Returns	14
4.4. See Also	14
4.5. Notes	14
4.6. Jupyter Notebook	15
5. SpotOptim Internal Methods Examples	17
5.1. Setup: Import Required Packages	17
5.2. 1. Main Optimization Method: <code>optimize()</code>	17
5.2.1. Changing the Optimizer	19
5.3. 2. Initial Design Methods	20
5.3.1. 2.1 <code>get_initial_design()</code>	20
5.3.2. 2.2 <code>curate_initial_design()</code>	21
5.3.3. 2.3 <code>rm_initial_design_NA_values()</code>	22
5.3.4. 2.4 <code>check_size_initial_design()</code>	23
5.3.5. 2.5 <code>get_best_xy_initial_design()</code>	24
5.4. 3. Surrogate Model Methods	24
5.4.1. 3.1 <code>fit_surrogate()</code>	24
5.4.2. 3.2 <code>_predict_with_uncertainty()</code>	25
5.4.3. 3.3 <code>_acquisition_function()</code>	26
5.4.4. 3.4 <code>suggest_next_infill_point()</code>	27

Table of contents

5.5.	4. OCBA Method (Optimal Computing Budget Allocation)	27
5.5.1.	4.1 <code>apply_ocba()</code>	27
5.6.	5. NaN/Inf Handling Methods	28
5.6.1.	5.1 <code>apply_penalty_NA()</code>	28
5.6.2.	5.2 <code>remove_nan()</code>	29
5.6.3.	5.3 <code>_handle_NA_new_points()</code>	30
5.7.	6. Main Loop Update Methods	31
5.7.1.	6.1 <code>_update_best_main_loop()</code>	31
5.8.	7. Termination Method	32
5.8.1.	7.1 <code>determine_termination()</code>	32
5.9.	8. Utility Methods	33
5.9.1.	8.1 <code>select_new()</code>	33
5.9.2.	8.2 <code>repair_non_numeric()</code>	34
5.9.3.	8.3 <code>map_to_factor_values()</code>	35
5.10.	9. Integration Examples	35
5.10.1.	9.1 Optimization with Custom Initial Design	35
5.10.2.	9.2 Optimization with Starting Point	36
5.10.3.	9.3 Optimization with Integer Variables	37
5.10.4.	9.4 Noisy Function Optimization with OCBA	39
5.10.5.	9.5 Optimization with NaN Handling	40
5.11.	Summary	41
5.12.	Jupyter Notebook	42
6.	Optimizing the Aircraft Wing Weight Example	43
6.1.	The AWWE Objective Function	43
6.2.	Baseline Configuration	44
6.3.	Optimization Setup	45
6.4.	Method 1: SpotOptim (Surrogate Model Based Optimization)	45
6.5.	Design Table	46
6.6.	Run optimization	46
6.7.	Result Table	48
6.8.	Progress of the Optimization	48
6.9.	Contour Plots of Most Important Hyperparameters	49
6.10.	Method 2: Nelder-Mead Simplex	52
6.11.	Method 3: BFGS (Quasi-Newton)	53
6.12.	Comparison of Results	54
6.13.	Visualization: Convergence Plots	55
6.13.1.	SpotOptim Convergence	55
6.14.	Optimal Parameter Values	56
6.15.	Analysis of Optimal Solutions	58
6.16.	Key Insights from Optimal Solutions	59
6.17.	Method Efficiency Comparison	60
6.18.	Visualization: 2D Slices of Optimal Solutions	61
6.19.	Conclusion	63
6.20.	Jupyter Notebook	64

7. Benchmarking SpotOptim with Sklearn Kriging (Matern Kernel) on 6D Rosenbrock and 10D Michalewicz Functions	65
7.1. SpotOptim with Sklearn Kriging in 6 Dimensions: Rosenbrock Function	65
7.1.1. Define the 6D Rosenbrock Function	65
7.1.2. Set up SpotOptim Parameters	66
7.1.3. Sklearn Gaussian Process Regressor as Surrogate	66
7.1.4. Visualize Optimization Progress	69
7.1.5. Evaluation of Multiple Repeats	70
7.2. SpotOptim with Sklearn Kriging in 10 Dimensions: Michalewicz Function	71
7.2.1. Define the 10D Michalewicz Function	71
7.2.2. Set up SpotOptim Parameters	72
7.2.3. Sklearn Gaussian Process Regressor as Surrogate	72
7.2.4. Visualize Optimization Progress	79
7.2.5. Evaluation of Multiple Repeats	80
7.3. Comparison: SpotOptim vs SpotPython	81
7.4. Summary	82
7.5. Jupyter Notebook	82
8. Success Rate Tracking in SpotOptim	83
8.1. What is Success Rate?	83
8.2. First Example	83
8.3. Second Example	84
8.4. Accessing Success Rate	85
8.5. Interpreting Success Rate	86
8.5.1. High Success Rate (> 0.5)	86
8.5.2. Medium Success Rate ($0.2 - 0.5$)	86
8.5.3. Low Success Rate (< 0.2)	87
8.6. TensorBoard Visualization	87
8.7. Example: Comparing Multiple Runs	87
8.8. Success Rate with Noisy Functions	89
8.9. Advanced: Custom Window Size	90
8.10. Best Practices	91
8.10.1. 1. Monitor During Long Runs	91
8.10.2. 2. Combine with TensorBoard	91
8.10.3. 3. Use as Stopping Criterion	91
8.10.4. 4. Compare Different Strategies	92
8.11. Technical Details	92
8.11.1. How Success is Counted	92
8.11.2. Rolling Window Calculation	92
8.11.3. Update Frequency	92
8.12. Summary	93
8.13. Jupyter Notebook	93
9. Parallel Evaluation	95
9.1. Sequential evaluation	95

Table of contents

9.2. Parallelism inside the objective function	95
9.3. Independent optimizations across processes	96
9.4. Jupyter Notebook	97
10. Restart Mechanism in SpotOptim	99
10.1. Key Concepts	99
10.1.1. Parameters	99
10.2. How it Works	99
10.3. Example: Triggering Restarts	100
10.3.1. Analyzing Restart Results	102
10.4. Example: Effect of <code>restart_inject_best</code>	102
10.5. When to Use Restarts?	105
II. Variables and Hyperparameters	107
11. Variable Type (<code>var_type</code>) Implementation in SpotOptim	109
11.1. Overview	109
11.2. Supported Variable Types	109
11.2.1. 1. <code>'float'</code>	109
11.2.2. 2. <code>'int'</code>	109
11.2.3. 3. <code>'factor'</code>	109
11.3. Implementation Details	110
11.3.1. Where <code>var_type</code> is Used	110
11.3.2. Core Method: <code>repair_non_numeric()</code>	110
11.4. Example Usage	111
11.4.1. Example 1: All Float Variables (Default)	111
11.4.2. Example 2: Pure Integer Optimization	111
11.4.3. Example 3: Categorical (Factor) Variables	112
11.4.4. Example 4: Mixed Variable Types	113
11.5. Key Findings	114
11.6. Recommendations	114
11.7. Future Enhancements (Optional)	114
11.8. Jupyter Notebook	115
12. Factor Variables for Categorical Hyperparameters	117
12.1. Overview	117
12.2. Quick Start	117
12.2.1. Basic Factor Variable Usage	117
12.2.2. Neural Network Activation Function Optimization	119
12.3. Mixed Variable Types	121
12.3.1. Combining Factor, Integer, and Continuous Variables	121
12.4. Multiple Factor Variables	124
12.4.1. Optimizing Both Activation and Optimizer	124

12.5. Advanced Usage	126
12.5.1. Custom Categorical Choices	126
12.5.2. Viewing All Evaluated Configurations	128
12.6. How It Works	131
12.6.1. Internal Mechanism	131
12.6.2. Variable Type Auto-Detection	132
12.7. Complete Example: Full Workflow	133
12.8. Best Practices	141
12.8.1. Do's	141
12.8.2. Don'ts	142
12.9. Troubleshooting	143
12.9.1. Common Issues	143
12.10 Summary	144
12.11 Jupyter Notebook	144
13. Variable Transformations for Search Space Scaling	145
13.1. Overview	145
13.2. Why Use Transformations?	145
13.2.1. Problem: Poorly Scaled Search Spaces	145
13.2.2. Solution: Logarithmic and Other Transformations	146
13.3. Quick Start	146
13.3.1. Basic Log-Scale Transformation	146
13.4. Supported Transformations	147
13.4.1. Transformation Guidelines	148
13.5. Detailed Examples	149
13.5.1. Example 1: Neural Network Hyperparameter Tuning	149
13.5.2. Example 2: Physics-Informed Neural Networks (PINNs)	152
13.5.3. Example 3: Mixing Transformations	153
13.6. Viewing Transformations in Tables	154
13.6.1. Design Table (Before Optimization)	154
13.6.2. Results Table (After Optimization)	155
13.7. Internal Architecture	156
13.7.1. Flow Diagram	156
13.7.2. Key Components	156
13.8. Best Practices	157
13.8.1. 1. Choose Appropriate Transformations	157
13.8.2. 2. Match Transformation to Range	157
13.8.3. 3. Validate Transformation Choice	157
13.8.4. 4. Combine with Variable Types	158
13.9. Troubleshooting	158
13.9.1. Issue: Values Out of Bounds	158
13.9.2. Issue: Poor Optimization Performance	158
13.9.3. Issue: Transformation Not Applied	159
13.10 Comparison: Manual vs Automatic Transformations	159
13.10.1 Manual Approach (Old Way)	159

Table of contents

13.10.2Automatic Approach (New Way)	160
13.11Summary	161
13.12Jupyter Notebook	162

III. Visualization 163

14. Surrogate Model Visualization 165

14.1. Overview	165
14.2. Features	165
14.3. Usage	165
14.3.1. Basic Usage	165
14.3.2. With Custom Parameters	166
14.3.3. Higher-Dimensional Problems	167
14.4. Plot Interpretation	169
14.4.1. Top Left: Prediction Surface	169
14.4.2. Top Right: Prediction Uncertainty Surface	169
14.4.3. Bottom Left: Prediction Contour with Points	170
14.4.4. Bottom Right: Uncertainty Contour with Points	170
14.5. Parameters	170
14.5.1. Dimension Selection	170
14.5.2. Appearance	170
14.5.3. Grid and Resolution	170
14.5.4. Color Scaling	170
14.5.5. Display	171
14.6. Examples	171
14.6.1. Example 1: 2D Rosenbrock Function	171
14.6.2. Example 2: Using Kriging Surrogate	172
14.6.3. Example 3: Comparing Different Dimension Pairs	173
14.7. Tips and Best Practices	176
14.8. Comparison with spotpython's plotkd()	177
14.8.1. Similarities	177
14.8.2. Differences	177
14.9. Error Handling	177
14.10Jupyter Notebook	178

15. TensorBoard Logging in SpotOptim 179

15.1. Quick Start	179
15.1.1. Enable TensorBoard Logging	179
15.1.2. View Logs in TensorBoard	180
15.1.3. Cleaning Old Logs	180
15.1.4. Use Cases	181
15.1.5. Custom Log Directory	181
15.1.6. What Gets Logged	181
15.1.7. Examples	182

15.2. TensorBoard Features	185
15.2.1. SCALARS Tab	185
15.2.2. HPARAMS Tab	185
15.2.3. Text Tab	185
15.3. TensorBoard Log Cleaning Feature in SpotOptim	186
15.3.1. Basic Usage	186
15.3.2. Use Cases	187
15.3.3. Implementation Details	187
15.3.4. Execution Flow	188
15.4. Safety Features	188
15.5. Tensorboard Demo Scripts	188
15.6. Troubleshooting	189
15.7. Related Parameters	189
15.8. Performance Notes	190
15.9. Jupyter Notebook	190

IV. Saving and Loading Models 191

16. Save and Load in SpotOptim	193
16.1. Key Concepts	193
16.1.1. Experiments vs Results	193
16.2. Quick Start	194
16.3. Distributed Workflow	194
16.3.1. Step 1: Define Experiment Locally	195
16.3.2. Step 2: Execute on Remote Machine	195
16.3.3. Step 3: Analyze Results Locally	196
16.4. Advanced Usage	198
16.4.1. Custom Filenames and Paths	198
16.4.2. Overwrite Protection	198
16.4.3. Loading and Continuing Optimization	199
16.5. Working with Noisy Functions	200
16.6. Working with Different Variable Types	201
16.7. Best Practices	203
16.7.1. 1. Re-attaching the Objective Function (Optional)	203
16.7.2. 2. Use Meaningful Prefixes	203
16.7.3. 3. Save Experiments Before Remote Execution	203
16.7.4. 4. Version Your Experiments	204
16.7.5. 5. Handle File Paths Robustly	204
16.8. Complete Example: Multi-Machine Workflow	204
16.8.1. Local Machine (Setup)	205
16.8.2. Remote Machine (Execution)	206
16.8.3. Local Machine (Analysis)	208
16.9. Technical Details	210
16.9.1. Serialization Method	210

Table of contents

16.9.2. Component Reinitialization	211
16.9.3. Excluded Components	211
16.9.4. File Format	211
16.10.Troubleshooting	211
16.10.1Issue: “AttributeError: ‘SpotOptim’ object has no attribute ‘fun’”	211
16.10.2Issue: “FileNotFoundError: Experiment file not found”	212
16.10.3Issue: “FileExistsError: File already exists”	212
16.10.4Issue: Results differ after loading	212
16.11Jupyter Notebook	213
17.Reproducibility in SpotOptim	215
17.1. Introduction	215
17.2. Basic Usage	215
17.2.1. Making Optimization Reproducible	215
17.2.2. Running Independent Experiments	216
17.3. Practical Examples	217
17.3.1. Example 1: Comparing Different Configurations	217
17.3.2. Example 2: Reproducible Research Experiment	218
17.3.3. Example 3: Multiple Independent Runs	220
17.3.4. Example 4: Reproducible Initial Design	222
17.3.5. Example 5: Custom Initial Design with Seed	223
17.4. Advanced Topics	224
17.4.1. Seed and Noisy Functions	224
17.4.2. Different Seeds for Different Exploration	225
17.5. Best Practices	226
17.5.1. 1. Always Use Seeds for Production Code	226
17.5.2. 2. Document Your Seeds	226
17.5.3. 3. Use Different Seeds for Different Experiments	226
17.5.4. 4. Test Robustness Across Multiple Seeds	227
17.6. What the Seed Controls	227
17.7. Common Questions	227
17.8. Summary	228
17.9. Jupyter Notebook	228
V. Surrogate Handling	229
18.Surrogate Model Selection in SpotOptim	231
18.1. Introduction	231
18.1.1. Acquisition Functions and Uncertainty	232
18.2. Setup and Imports	232
18.3. The AWWE Objective Function	233
18.4. 1. Default Surrogate: Gaussian Process with Matern Kernel	233
18.4.1. Visualization: Default Surrogate	234

18.5. 2. Gaussian Process with RBF (Radial Basis Function) Kernel	238
18.5.1. Visualization: RBF Kernel	240
18.6. 3. Gaussian Process with Matern nu=1.5 Kernel	243
18.6.1. Visualization: Matern nu=1.5 Kernel	245
18.7. 4. Gaussian Process with Rational Quadratic Kernel	248
18.7.1. Visualization: Rational Quadratic Kernel	250
18.8. 5. SpotOptim Kriging Model	253
18.8.1. Visualization: Kriging Model	255
18.9. 6. Random Forest Regressor	258
18.9.1. Visualization: Random Forest	260
18.10.7. XGBoost Regressor	263
18.11.8. Support Vector Regression (SVR)	265
18.11.1. Visualization: SVR	266
18.12.9. Gradient Boosting Regressor	269
18.12.1. Visualization: Gradient Boosting	270
18.13. Comprehensive Comparison	274
18.13.1. Visualization: Performance Comparison	275
18.14. Convergence Comparison	277
18.15. Key Insights and Recommendations	278
18.16. Summary Statistics	280
18.17. Conclusion	282
18.18. Jupyter Notebook	282
19. Acquisition Failure Handling in SpotOptim	283
19.1. What is Acquisition Failure?	283
19.2. Fallback Strategies	283
19.2.1. Random Space-Filling Design (Default)	283
19.3. How It Works	285
19.4. Advanced Usage: Setting Tolerance	285
19.5. Best Practices	286
19.5.1. 1. Monitor Fallback Activations	286
19.5.2. 2. Adjust Tolerance Based on Problem Scale	287
19.6. Technical Details	287
19.6.1. Random Strategy Implementation	287
19.7. Summary	288
19.8. Jupyter Notebook	288
20. Point Selection Implementation in SpotOptim	289
20.1. Overview	289
20.2. Implementation Details	289
20.2.1. Parameters	289
20.2.2. Methods	289
20.3. Key Differences from spotpython	290
20.4. Example Usage	290
20.5. Benefits	292

Table of contents

20.6. Comparison with spotpython	292
20.7. Jupyter Notebook	292

VI. Optimization on the Surrogate 293

21. Optimization on the Surrogate	295
21.1. Key Arguments	295
21.2. Examples	296
21.2.1. 1. Default Configuration (Differential Evolution)	296
21.2.2. 2. Customizing Differential Evolution	296
21.2.3. 3. Using Gradient-Based Optimization (L-BFGS-B)	297
21.2.4. 4. Using Gradient-Free Optimization (Nelder-Mead)	298
21.2.5. 5. Returning Multiple Candidates	299

VII. Kriging 301

22. Kriging Surrogate Integration	303
22.1. Overview	303
22.2. Module Structure	303
22.3. Kriging Class (src/spotoptim/surrogate/kriging.py)	303
22.4. Integration with SpotOptim	304
22.5. Documentation	304
22.6. Usage Examples	305
22.6.1. Basic Usage	305
22.6.2. Custom Parameters	305
22.6.3. Prediction with Uncertainty	306
22.7. Technical Details	307
22.7.1. Kriging vs GaussianProcessRegressor	307
22.7.2. Algorithm	307
22.7.3. Key Arguments Passed from SpotOptim	308
22.8. Benefits	308
22.9. Future Enhancements	308
22.10 Jupyter Notebook	309
23. Kriging Surrogate Models in SpotOptim	311
23.1. Introduction	311
23.1.1. What is Kriging?	311
23.1.2. Why Use Kriging in SpotOptim?	312
23.2. Basic Usage	312
23.2.1. Creating a Simple Kriging Model	312
23.2.2. Default vs Custom Surrogate	313
23.3. Understanding Kriging Parameters	314
23.3.1. Method: Interpolation vs Regression	314

23.3.2. Noise Parameter and Regularization	316
23.3.3. Correlation Length Parameters (Theta)	317
23.3.4. Isotropic vs Anisotropic Correlation	318
23.4. Advanced Features	319
23.4.1. Mixed Variable Types	319
23.4.2. Customizing the Distance Metric for Factors	322
23.4.3. Handling High-Dimensional Problems	323
23.4.4. Uncertainty Quantification	324
23.5. Practical Examples	326
23.5.1. Example 1: Optimizing the Rastrigin Function	326
23.5.2. Example 2: Robust Optimization with Noise	328
23.5.3. Example 3: Real-World Machine Learning Hyperparameter Tuning	329
23.5.4. Example 4: Comparing Kriging Methods	332
23.5.5. Example 5: Sensitivity to Theta Bounds	333
23.6. Comparing Surrogates	335
23.6.1. Kriging vs Gaussian Process vs Random Forest	335
23.7. Best Practices	336
23.7.1. 1. Choosing the Right Method	336
23.7.2. 2. Setting Model Complexity	337
23.7.3. 3. Handling Different Variable Types	337
23.7.4. 4. Reproducibility	338
23.7.5. 5. Monitoring Surrogate Quality	338
23.8. Common Issues and Solutions	339
23.8.1. Issue 1: Slow Fitting for Large Datasets	339
23.8.2. Issue 2: Poor Predictions for Categorical Variables	339
23.8.3. Issue 3: Numerical Instability	340
23.8.4. Issue 4: Overfitting to Noisy Data	340
23.9. Summary	340
23.9.1. Key Takeaways	340
23.9.2. Quick Reference	341
23.9.3. Parameter Cheat Sheet	341
23.10.Jupyter Notebook	342
24. Kriging in SpotOptim: The Forrester Implementation	343
25. Introduction	345
26.1. The Core Kriging Model (Section 2.4)	347
26.1. Correlation Function	347
26.2. Training: Concentrated Log-Likelihood	347
26.3. Prediction	348
27.2. Handling Noisy Data	349
27.1. A. Interpolation (method="interpolation")	349
27.2. B. Regression (method="regression")	349

Table of contents

27.3. C. Re-interpolation (method="reinterpolation")	350
28.3. Examples	351
28.1. Setup	351
28.2. Comparisons	351
28.3. Visualization	352
28.3.1. Interpretation	353
29.4. Infill Criteria	355
30. References	357
30.1. Jupyter Notebook	357
31. Nystroem Kernel Approximation	359
31.1. Introduction	359
31.2. The Nystroem Method	359
31.2.1. Moderate n , Large k Use Case	360
31.3. Implementation in SpotOptim	360
31.3.1. Components	360
31.3.2. Example Usage	360
31.4. Handling Mixed Variables	362
31.4.1. Using <code>SpotOptimKernel</code>	362
31.5. Jupyter Notebook	363
VIII Multiobjective Optimization	365
32. Multi-Objective Optimization Support in SpotOptim	367
32.1. Overview	368
32.2. Common Scalarization Strategies	372
32.2.1. 1. Weighted Sum	372
32.2.2. 2. Weighted Sum with Normalization	372
32.2.3. 3. Min-Max (Chebyshev)	372
32.2.4. 4. Target Achievement	373
32.2.5. 5. Product	373
32.3. Integration with Other Features	373
32.4. Implementation Details	373
32.4.1. Automatic Detection	373
32.4.2. Data Flow	374
32.4.3. Backward Compatibility	374
32.5. Limitations and Notes	374
32.5.1. What This Is	374
32.5.2. What This Is Not	374
32.5.3. For True Multi-Objective Optimization	375

32.6. Benchmark Multi-Objective Test Functions	375
32.6.1. Function Overview	375
32.6.2. Example 6: ZDT1 - Convex Pareto Front	376
32.6.3. Example 7: ZDT2 - Non-Convex Pareto Front	378
32.6.4. Example 8: ZDT3 - Disconnected Pareto Front	380
32.6.5. Example 9: ZDT4 - Multimodal with Many Local Fronts	381
32.6.6. Example 10: ZDT6 - Non-Uniform Density	383
32.6.7. Example 11: DTLZ2 - Scalable Spherical Pareto Front	385
32.6.8. Example 12: Schaffer N1 - Simple Bi-Objective	387
32.6.9. Example 13: Fonseca-Fleming - Symmetric Concave Front	389
32.7. Example 14: Kursawe - Non-Convex Disconnected Front	391
32.7.1. Comparison: Different Scalarization Strategies on ZDT1	392
32.8. A Simple Visualization Example	394
32.9. Jupyter Notebook	401

IX. Hyperparameter Tuning 403

33. The TorchObjective Class for Hyperparameter Tuning 405

33.1. Overview	405
33.2. Workflow Steps	405
33.3. Step-by-Step Example	406
33.3.1. 1. Define the Model	406
33.3.2. 2. Prepare the Data	406
33.3.3. 3. Define Hyperparameters	407
33.3.4. 4. Create Experiment Control	408
33.3.5. 5. Initialize TorchObjective	409
33.3.6. 6. Run Optimization	409
33.4. Advanced Features	410
33.4.1. Tuning Batch Size	410
33.4.2. Tuning Architecture	410
33.4.3. Log-Scale Transformations	412
33.5. Custom Models	412

34. The MLP Class 415

34.1. Overview	415
34.2. Basic Usage	415
34.2.1. Initialization	415
34.2.2. Implicit Architecture (Width & Depth)	416
34.3. Configuration Options	417
34.4. Optimizer Integration	417
34.5. Hyperparameter Tuning	418
34.5.1. Example: Tuning with SpotOptim	419

35. Physics-Informed Neural Networks (PINNs) Demo 1 421

Table of contents

36. Overview	423
36.1. The Differential Equation	423
37. Setup	425
38. The Neural Network	427
39. Generate Training Data	429
40. Collocation Points	433
41. PINN Training	435
41.1. Loss Components Explained	435
41.1.1. Neural Network Prediction on Data Points	435
41.1.2. Data Loss (Loss1)	435
41.1.3. Neural Network Prediction on Collocation Points	435
41.1.4. Computing Derivatives for the ODE	435
41.1.5. Physics Loss (Loss2)	436
41.1.6. Total Loss	436
41.2. Training Loop	436
42. Results Visualization	439
42.1. Training Progress	439
42.2. Solution Evolution During Training	440
42.3. Final Solution Comparison	442
42.4. Error Analysis	444
43. Summary	447
43.1. Key Advantages of PINNs	447
43.2. Using SpotOptim for Hyperparameter Optimization	447
43.3. Jupyter Notebook	448
44. Hyperparameter Tuning for Physics-Informed Neural Networks	449
45. Overview	451
45.1. Key Features	451
45.1.1. 1. PyTorch Dataset and DataLoader	451
45.1.2. 2. Automatic Transformation Handling	452
46. The Problem	453
47. Setup	455
48. Data Generation	457
48.1. Custom Dataset Classes	457
49. Define the PINN Training Function	461

50. Hyperparameter Optimization with SpotOptim	465
50.1. Define the Objective Function	465
50.2. Run the Optimization	467
51. Results Analysis	471
51.1. Best Configuration	471
51.1.1. Results Table with Importance Scores	472
51.2. Optimization History	472
51.3. Surrogate Visualization	473
51.4. Parameter Distribution Analysis	475
52. Train Final Model with Best Hyperparameters	477
53. Evaluate Final Model	481
54. Visualize Final Solution	483
55. Comparison with Baseline	485
56. Hyperparameter Sensitivity Analysis	489
56.1. Sensitivity Analysis (Spearman Correlation)	490
57. Summary	491
57.1. Key Findings	491
57.2. Recommendations	492
57.3. Using These Results	493
57.4. Using var_trans for Your Hyperparameter Optimization	494
57.5. Future Directions	495
57.6. Jupyter Notebook	495
58. Learning Rate Mapping for Unified Optimizer Interface	497
58.1. Overview	497
58.2. Quick Start	497
58.2.1. Basic Usage	497
58.2.2. Scaling Learning Rates	498
58.2.3. Integration with LinearRegressor	499
58.3. Function Reference	499
58.3.1. map_lr(lr_unified, optimizer_name, use_default_scale=True)	499
58.4. Supported Optimizers	500
58.5. Use Cases	501
58.5.1. Comparing Different Optimizers	501
58.5.2. Hyperparameter Optimization	502
58.5.3. Hyperparameter Optimization with SpotOptim	504
58.5.4. Log-Scale Hyperparameter Search	507
58.5.5. Custom Learning Rate Schedules	507
58.5.6. Direct Usage Without LinearRegressor	508

Table of contents

58.6. Best Practices	509
58.6.1. Choosing Unified Learning Rate	509
58.6.2. Optimizer Selection Guidelines	510
58.6.3. Common Patterns	510
58.7. Troubleshooting	511
58.7.1. Issue: Training is unstable (loss explodes)	511
58.7.2. Issue: Training is too slow (loss decreases very slowly)	511
58.7.3. Issue: Different results across optimizer runs	511
58.7.4. Issue: Want to use raw learning rate without mapping	511
58.7.5. Issue: Optimizer not supported	511
58.8. Technical Details	512
58.8.1. How It Works	512
58.8.2. Design Rationale	512
58.8.3. Default Learning Rates	513
58.9. Examples	513
58.9.1. Complete Example: Optimizer Comparison Study	513
58.10.Jupyter Notebook	517
59. SpotOptim Scaling	519
60. TorchStandardScaler in SpotOptim	521
60.1. Introduction	521
60.2. Using TorchStandardScaler Directly	521
60.3. Integrating with TorchObjective	522
60.3.1. Example: Scaling Effect	522
60.3.2. Explanation	524
X. Data Sets	527
61. Diabetes Dataset Utilities	529
61.1. Overview	529
61.2. Quick Start	529
61.2.1. Basic Usage	529
61.2.2. Training a Model	530
61.3. Function Reference	532
61.3.1. get_diabetes_data loaders()	532
61.4. DiabetesDataset Class	533
61.4.1. Manual Dataset Creation	534
61.5. Advanced Usage	535
61.5.1. Custom Transforms	535
61.5.2. Different Train/Test Splits	536
61.5.3. Without Feature Scaling	536
61.5.4. Larger Batch Sizes	537
61.5.5. GPU Training with Pin Memory	537

61.6. Complete Training Example	538
61.7. Integration with SpotOptim	541
61.8. Best Practices	543
61.8.1. 1. Always Use Feature Scaling	543
61.8.2. 2. Set Random Seeds for Reproducibility	544
61.8.3. 3. Don't Shuffle Test Data	544
61.8.4. 4. Choose Appropriate Batch Size	544
61.8.5. 5. Save the Scaler for Production	545
61.9. Troubleshooting	545
61.9.1. Issue: Out of Memory	545
61.9.2. Issue: Different Data Ranges	546
61.9.3. Issue: Non-Reproducible Results	546
61.9.4. Issue: Slow Data Loading	546
61.10 Summary	547
61.11 Jupyter Notebook	547
 XI. Details About SPOT	 549
62. SpotOptim Step-by-Step Optimization Process	551
63. Introduction	553
64. Setup and Test Functions	555
65. Standard Optimization Workflow	557
65.1. Phase: Initialization and Setup	557
65.2. Phase: Initial Design Generation	558
65.2.1. Method: <code>get_initial_design()</code>	558
65.3. Phase 3: Initial Design Curation	561
65.3.1. Method: <code>curate_initial_design()</code>	561
65.4. Phase 4: Initial Design Evaluation	561
65.4.1. Method: <code>evaluate_function()</code>	561
65.5. Phase 5: Handling Failed Evaluations	562
65.5.1. Method: <code>rm_initial_design_NA_values(X0,y0)</code>	562
65.6. Phase 6: Validation Check	563
65.6.1. Method: <code>check_size_initial_design(y0, n_evaluated)</code>	563
65.7. Phase 7: Storage Initialization	563
65.7.1. Method: <code>init_storage()</code>	563
65.8. Phase: Statistics Update	564
65.8.1. Method: <code>update_stats()</code>	564
65.9. Phase: Log initial Design to TensorBoard	564
65.9.1. Method: <code>_log_initial_design_tensorboard()</code>	564
65.10 Phase: Initial Best Point	565
65.10.1 Method: <code>get_best_xy_initial_design()</code>	565

Table of contents

65.11	Illustration of the Initial Design Phase Results	565
66.	Sequential Optimization Loop	567
66.1.	Step: Surrogate Model Fitting	567
66.1.1.	Method: <code>fit_scheduler()</code>	567
66.1.2.	Method: <code>fit_surrogate()</code>	567
66.1.3.	Method: <code>fit_selection_dispatcher()</code>	568
66.1.4.	Step: Apply OCBA	569
66.1.5.	Method: <code>apply_ocba()</code>	569
66.2.	Step: Predict with Uncertainty	570
66.2.1.	Method: <code>_predict_with_uncertainty()</code>	570
66.3.	Step: Next Point Suggestion	570
66.3.1.	Method: <code>suggest_next_infill_point()</code>	570
66.4.	Step: Acquisition Function Evaluation	571
66.4.1.	Method: <code>_acquisition_function()</code>	571
66.5.	Step: Update Repeats for Infill Points	572
66.5.1.	Method: <code>update_repeats_infill_points()</code>	572
66.6.	Append OCBA Points to Infill Points	573
66.7.	Step: Evaluation of New Points	573
66.7.1.	Method: <code>evaluate_function()</code> (again)	573
66.8.	Step: Handle Failed Evaluations (Sequential)	573
66.8.1.	Method: <code>_handle_NA_new_points()</code>	573
66.9.	Step: Update Success Rate	574
66.9.1.	Method: <code>update_success_rate(y0)</code>	574
66.10	Step: Update Storage	575
66.10.1	Internal updates	575
66.11	Update Statistics	575
66.11.1.	What happens in <code>update_stats()</code> :	575
66.12	Step: Update Best Solution	576
66.12.1	Method: <code>_update_best_main_loop()</code>	576
67.	Complete Optimization Example	577
68.	Noisy Functions with Repeats	579
68.1.	Configuration for Noisy Functions	579
68.2.	Noisy Optimization Workflow Differences	580
68.2.1.	Modified Initial Design	580
68.2.2.	Update Statistics Method	581
69.	Optimal Computing Budget Allocation (OCBA)	583
69.1.	OCBA Configuration	583
69.2.	OCBA Method	584
69.2.1.	Method: <code>apply_ocba()</code>	584

70. Handling Function Evaluation Failures	587
70.1. Example with Failures	587
70.2. Failure Handling in Initial Design	588
70.2.1. Method: <code>rm_initial_design_NA_values()</code>	588
70.3. Failure Handling in Sequential Phase	589
70.3.1. Method: <code>_handle_NA_new_points()</code>	589
70.4. Penalty Application	589
70.4.1. Method: <code>apply_penalty_NA()</code>	589
71. Complete Method Summary	591
71.1. Methods Called During <code>optimize()</code>	591
71.1.1. Preparation Phase (inside <code>optimize_sequential_run()</code> via <code>_initialize_run()</code>)	591
71.1.2. Sequential Optimization Loop (each iteration inside <code>_run_sequential_loop()</code>)	591
71.1.3. Finalization	592
71.2. Helper Methods Used	592
72. Termination Conditions	593
72.1. Method: <code>determine_termination()</code>	593
73. Performance Comparison on Noisy Functions	595
74. Summary	599
74.1. Complete Workflow Diagram	599
74.2. Key Concepts	601
74.2.1. 1. Initial Design	601
74.2.2. 2. Surrogate Model	601
74.2.3. 3. Acquisition Function	601
74.2.4. 4. Noise Handling	602
74.2.5. 5. Failure Handling	602
74.3. Best Practices	602
74.4. Conclusion	603
74.5. Jupyter Notebook	603
75. SpotOptim Execution Flow	605
76. SpotOptim Execution Flow	607
76.1. Example Without Restarts	607
76.1.1. <code>SpotOptim.__init__</code>	607
76.1.2. <code>SpotOptim.optimize</code>	608
76.1.3. <code>SpotOptim._optimize_single_run</code>	608
76.1.4. <code>SpotOptim.get_initial_design</code>	608
76.1.5. <code>SpotOptim.curate_initial_design</code>	608
76.1.6. <code>SpotOptim.evaluate_function</code> (Initial Design)	609
76.1.7. <code>SpotOptim.init_storage</code>	609

Table of contents

76.1.8. Optimization Loop Start	609
76.1.9. SpotOptim._fit_scheduler	609
76.1.10 SpotOptim.suggest_next_infill_point	609
76.1.11 SpotOptim._try_optimizer_candidates	610
76.1.12 SpotOptim.optimize_acquisition_func	610
76.1.13 SpotOptim.evaluate_function (Loop)	610
76.1.14 SpotOptim.update_storage	610
76.1.15 SpotOptim._update_best_main_loop	611
76.1.16 Termination	611
76.2. Example With Restarts	611
76.2.1. SpotOptim.optimize (Global Loop)	611
76.2.2. SpotOptim._optimize_single_run (First Run)	612
76.2.3. Restart Injection (Between Runs)	612
76.2.4. SpotOptim._optimize_single_run (Subsequent Runs)	612
76.2.5. Termination	613
76.3. Description Based on Chunks	613
76.3.1. Configuration & Helpers	613
76.3.2. Dimension Reduction	614
76.3.3. Variable Transformation	615
76.3.4. Initial Design	615
76.3.5. Surrogate & Acquisition	616
76.3.6. Optimization Loop	616
76.3.7. Results & Analysis	617
76.3.8. Plotting	618
76.3.9. TensorBoard Integration	618
76.4. Description of the Infill Point Acquisition Steps	619
76.4.1. 1. Optimize Acquisition Function	619
76.4.2. 2. Candidate Validation and Selection	619
76.4.3. 3. Failure Handling and Fallback Strategies	620
76.4.4. Key Arguments and Configuration	620
76.4.5. Example Usage	622
76.5. Using the Surrogate Model	623
76.6. Comparing the SpotOptim With the Default Surrogate Model and MLP Surrogate	624
76.7. Multi-Surrogate Switching in SpotOptim	627
76.7.1. Comparison: Default (GP) vs Multi-Surrogate (GP + MLP) on Robot Arm	627

XII.Optimization 659

77.Aircraft Wing Weight Example 661	661
77.1. AWWE Equation	661
77.2. AWWE Parameters and Equations (Part 1)	661
77.3. Goals: Understanding and Optimization	662

77.4. Properties of the Python “Solver”	664
77.5. Plot 1: Load Factor (N_z) and Aspect Ratio (A)	665
77.6. Plot 2: Taper Ratio and Fuel Weight	667
77.7. The Big Picture: Combining all Variables	668
77.8. AWE Landscape	671
77.9. Summary of the First Experiments	672
77.10. Exercise	672
77.10.1. Adding Paint Weight	672
77.11. Jupyter Notebook	673

XIII Numerical Methods 675

78. Simulation and Surrogate Modeling 677

78.1. Surrogates	677
78.1.1. Costs of Simulation	678
78.1.2. Mathematical Models and Meta-Models	678
78.1.3. Surrogates = Trained Meta-models	678
78.1.4. Computer Experiments	678
78.1.5. Limits of Mathematical Modeling	679
78.1.6. Why Computer Simulations are Necessary	679
78.1.7. Simulation Requirements	679
78.2. Applications of Surrogate Models	679
78.3. DACE and RSM	680
78.3.1. Noise Handling in RSM and DACE	681
78.4. Updating a Surrogate Model	681

79. Sampling Plans 683

79.1. Ideas and Concepts	683
79.1.1. The ‘Curse of Dimensionality’ and How to Avoid It	684
79.1.2. Physical versus Computational Experiments	685
79.1.3. Designing Preliminary Experiments (Screening)	686
79.1.4. Special Considerations When Deploying Screening Algorithms	695
79.2. Analyzing Variable Importance in Aircraft Wing Weight	695
79.3. Designing a Sampling Plan	698
79.3.1. Stratification	698
79.3.2. Latin Squares and Random Latin Hypercubes	703
79.3.3. Space-filling Designs: Maximin Plans	706
79.3.4. Memory Management	711
79.3.5. Optimizing the Morris-Mitchell Criterion Φ_q	719
79.3.6. Evolutionary Operation	721
79.3.7. Putting it all Together	722
79.4. Experimental Analysis of the Morris-Mitchell Criterion	729
79.4.1. Evaluation of Sampling Designs	729
79.4.2. Demonstrate the Impact of mmphi Parameters	732

Table of contents

79.4.3. Morris-Mitchell Criterion: Impact of Adding Points	732
79.5. A Sample-Size Invariant Version of the Morris-Mitchell Criterion . . .	737
79.5.1. Comparison of <code>mmphi()</code> and <code>mmphi_intensive()</code>	737
79.5.2. Plotting the Two Morris-Mitchell Criteria for Different Sample Sizes	738
79.6. Jupyter Notebook	741
80. Constructing a Surrogate	743
80.1. Stage One: Preparing the Data and Choosing a Modelling Approach .	743
80.2. Stage Two: Parameter Estimation and Training	745
80.3. Stage Three: Model Testing	747
80.4. Jupyter Notebook	748
81. Response Surface Methods	749
81.1. What is RSM?	749
81.1.1. Visualization: Problems in Practice	752
81.1.2. RSM: Strategies	752
81.1.3. RSM: Noise in the Empirical Model	753
81.1.4. RSM: Natural and Coded Variables	753
81.1.5. RSM Low-order Polynomials	754
81.2. First-Order Models (Main Effects Model)	754
81.2.1. First-Order Model Properties	755
81.2.2. First-order Model with Interactions in <code>python</code>	755
81.2.3. Observations: First-Order Model with Interactions	757
81.3. Second-Order Models	757
81.3.1. Second-Order Models: Properties	758
81.3.2. Example: Stationary Ridge	758
81.3.3. Observations: Second-Order Model (Ridge)	759
81.3.4. Example: Rising Ridge	760
81.3.5. Summary: Rising Ridge	761
81.3.6. Falling Ridge	761
81.3.7. Saddle Point	762
81.3.8. Interpretation: Saddle Points	763
81.3.9. Summary: Ridge Analysis	763
81.4. General RSM Models	763
81.4.1. Ordinary Least Squares	763
81.5. General Linear Regression	764
81.6. Designs	765
81.6.1. Different Designs	766
81.7. RSM Experimentation	766
81.7.1. First Step	766
81.7.2. Second Step	766
81.7.3. Third Step	766
81.8. RSM: Review and General Considerations	767
81.8.1. Historical Considerations about RSM	768

81.8.2. Status Quo	768
81.8.3. The Role of Statistics	768
81.8.4. New RSM is needed: DACE	768
81.9. Exercises	769
81.10.Jupyter Notebook	769
82. Polynomial Models	771
82.1. Fitting a Polynomial	771
82.2. Polynomial Fitting in Python	772
82.2.1. Fitting the Polynomial	772
82.2.2. Explaining the k -fold Cross-Validation	773
82.2.3. Making Predictions	774
82.2.4. Plotting the Results	775
82.3. Example One: Aerofoil Drag	776
82.4. Example Two: A Multimodal Test Case	778
82.5. Extending to Multivariable Polynomial Models	779
82.6. Jupyter Notebook	781
83. Radial Basis Function Models	783
83.1. Radial Basis Function Models	783
83.1.1. Fitting Noise-Free Data	784
83.1.2. Numerical Stability Through Positive Definite Matrices	788
83.1.3. Ill-Conditioning	789
83.1.4. Parameter Optimization: A Two-Level Approach	790
83.2. Python Implementation of the RBF Model	793
83.2.1. The Rbf Class	794
83.3. RBF Example: The One-Dimensional $\sin$ Function	800
83.4. RBF Example: The Two-Dimensional dome Function	802
83.4.1. The Connection Between RBF Models and Neural Networks	805
83.5. Radial Basis Function Models for Noisy Data	806
83.5.1. Ridge Regularization Approach	806
83.5.2. Reduced Basis Approach	806
83.6. Jupyter Notebook	807
84. Kriging (Gaussian Process Regression)	809
84.1. From Gaussian RBF to Kriging Basis Functions	809
84.2. Building the Kriging Model	810
84.3. MLE to estimate θ and p	817
84.3.1. The Log-Likelihood	817
84.3.2. Differentiation with Respect to μ	819
84.3.3. Differentiation with Respect to σ	820
84.3.4. Results of the Optimizations	821
84.3.5. The Concentrated Log-Likelihood Function	821
84.3.6. Optimizing the Parameters $\vec{\theta}$ and $\vec{p}$	822
84.3.7. Correlation and Covariance Matrices Revisited	822

Table of contents

84.4. Implementing an MLE of the Model Parameters	823
84.5. Kriging Prediction	823
84.6. Kriging Example: Sinusoid Function	825
84.6.1. Calculating the Correlation Matrix Ψ	825
84.6.2. Computing the Ψ Matrix	827
84.6.3. Selecting the New Locations	828
84.6.4. Computing the ψ Vector	828
84.6.5. Predicting at New Locations	829
84.6.6. Visualization	829
84.6.7. The Complete Python Code for the Example	830
84.7. Jupyter Notebook	836
85. Matrices	837
85.1. Derivatives of Quadratic Forms	837
85.2. The Condition Number	838
85.3. The Moore-Penrose Pseudoinverse	839
85.3.1. Definitions	839
85.3.2. Implementation in Python	839
85.4. Strictly Positive Definite Kernels	840
85.4.1. Definition	840
85.4.2. Connection to Positive Definite Matrices	840
85.4.3. Connection to RBF Models	841
85.5. Cholesky Decomposition and Positive Definite Matrices	841
85.5.1. Example of Cholesky Decomposition	844
85.5.2. Inverse Matrix Using Cholesky Decomposition	845
85.6. Nyström Approximation	846
85.6.1. What's the Big Idea?	846
85.6.2. How Does It Work?	846
85.6.3. Example	847
85.6.4. Why Is This Useful?	847
85.6.5. Applying the Nyström Approximation: How Nyström Approximation Helps Kriging	848
85.6.6. Example: Predicting Temperature with Nyström-Kriging	848
85.6.7. Details: Woodbury Matrix Identity for Avoiding the Big Inversion	849
85.6.8. The Example: Step-by-Step	850
85.7. Extending spotpython's Kriging Surrogate with Nyström Approximation for Enhanced Scalability	852
85.7.1. Introduction: Overcoming the Scalability Challenge in Kriging for Sequential Optimization	852
85.7.2. Report Objectives and Structure	853
85.8. Theoretical Foundations: The Nyström-Kriging Framework	853
85.8.1. A Primer on Kriging (Gaussian Process Regression)	853
85.8.2. The Nyström Method for Low-Rank Kernel Approximation	854
85.8.3. Justification for Landmark Selection	854

85.9. Implementation: A Scalable Kriging Class for spotpython	854
85.9.1. Updated kriging.py with Nyström Approximation (excerpt) . .	854
85.10 Implementation Details	857
85.10.1 Architectural Enhancements to <code>init</code>	857
85.10.2 The <code>fit()</code> Method: A Dual-Pathway Approach	857
85.10.3 The <code>predict()</code> Method: Conditional Prediction Logic	857
85.10.4 Critical Detail: Preserving Mixed Variable Type Functionality .	858
85.10.5 Seamless Integration into the Nyström Workflow	858
85.11 Jupyter Notebook	858
86. Infill Criteria	859
86.1. Balancing Exploitation and Exploration	860
86.2. Expected Improvement (EI)	860
86.3. Expected Improvement in the Hyperparameter Tuning Cookbook (Python Implementation)	861
86.4. Jupyter Notebook	862
87. Remote Optimization Challenge: A Benchmarking Study	863
87.1. Benchmarking Best Practices	863
87.2. Experimental Setup	863
87.2.1. The Remote Objective Function	863
87.2.2. Algorithms	864
87.2.3. Budget and Settings	864
87.3. Comparison of the Difficulty	864
87.4. Statistical Power Analysis	867
87.4.1. 1. Estimating Noise (σ)	867
87.4.2. 2. Determining Sample Size (n)	867
87.5. Experimental Comparison	868
87.6. Convergence Analysis (Noisy)	871
87.7. Comparative Analysis	873
87.8. Conclusion	875
87.9. Landscape Analysis of <code>robot_arm_hard</code>	875
87.10 References	877
87.11 Jupyter Notebook	878
XIV Data-Driven Modeling and Optimization	879
88. Basic Statistics and Data Analysis	881
88.1. Exploratory Data Analysis	881
88.1.1. Histograms	881
88.1.2. Boxplots	884
88.2. Probability Distributions	884
88.2.1. Sampling from a Distribution	885

Table of contents

88.3. Discrete Distributions	885
88.3.1. Bernoulli Distribution	885
88.3.2. Binomial Distribution	886
88.4. Continuous Distributions	887
88.4.1. Distribution functions: PDFs and CDFs	887
88.5. Background: Expectation, Mean, Standard Deviation	888
88.6. Distributions and Random Numbers in Python	890
88.6.1. Calculation of the Standard Deviation with Python	890
88.6.2. The Argument “axis”	891
88.7. Expectation (Continuous)	892
88.8. Variance and Standard Deviation (Continuous)	892
88.8.1. The Uniform Distribution	893
88.9. The Uniform Distribution	894
88.9.1. The Normal Distribution	895
88.9.2. Visualization of the Standard Deviation	896
88.9.3. Realizations of a Normal Distribution	897
88.10. The Normal Distribution	898
88.11. The Exponential Distribution	903
88.11.1. Standardization of Random Variables	903
88.12. Covariance and Correlation	903
88.12.1. The Multivariate Normal Distribution	903
88.13. Covariance	906
88.14. Correlation	908
88.14.1. Definitions	908
88.14.2. Computations	910
88.14.3. The Outer-product and the <code>np.outer</code> Function	911
88.14.4. Correlation and Independence	914
88.14.5. Pearson’s Correlation	916
88.14.6. Interpreting the Correlation: Correlation Squared	916
88.14.7. Partial Correlation	917
88.15. Hypothesis Testing and the Null-Hypothesis	920
88.15.1. Alternative Hypotheses, Main Ideas	920
88.15.2. p-values: What they are and how to interpret them	920
88.16. Statistical Power	920
88.17. The Central Limit Theorem	920
88.18. Maximum Likelihood	920
88.19. Maximum Likelihood Estimation: Multivariate Normal Distribution	922
88.19.1. The Joint Probability Density Function of the Multivariate Normal Distribution	922
88.19.2. The Log-Likelihood Function	923
88.20. Cross-Validation	923
88.21. Mutual Information	924
88.22. t-SNE	925

89. Hypothesis Testing	927
89.1. Paired-Sample Analysis	927
89.1.1. Directional Hypotheses for a Paired Experiment	928
89.1.2. The Paired Samples t-Test: A Technical Examination	930
89.1.3. A Priori Sample Size Determination for Desired Power	932
89.1.4. Effect Size (Δ)	933
89.2. Non-inferiority Test	938
90. Addressing Multicollinearity: Principle Component Analysis (PCA) and Factor Analysis (FA)	941
90.1. Introduction	941
90.2. The Car-Sales Data Set	942
90.2.1. The Target Variable	943
90.2.2. The Features	943
90.2.3. Combining Non-categorical and Categorical (encoded) Data	945
90.3. Fit the Linear Regression Model	945
90.3.1. Model Summary and Interpretation	946
90.4. Collinearity Diagnostics	947
90.4.1. The Coefficient Table	947
90.4.2. Eigenvalues and Condition Indices	948
90.4.3. Kayser-Meyer-Olkin (KMO) Measure	949
90.5. Addressing Multicollinearity with Principal Component Analysis (PCA)	950
90.5.1. Introduction to PCA	950
90.5.2. Application of PCA in Regression Problems:	950
90.5.3. Scree Plot	951
90.5.4. Loading Scores (for PCA)	951
90.5.5. PCA for Car Sales Example	952
90.5.6. Creating the Regression Model with Principal Components	956
90.5.7. Collinearity Diagnostics for PCA Regression Model	958
90.5.8. PCA: Creating the Regression Model with three Principle Components only	959
90.6. Addressing Multicollinearity and Latent Structure with Factor Analysis (FA)	960
90.6.1. Introduction to Factor Analysis	960
90.6.2. Determining the Number of Factors for Factor Analysis	961
90.6.3. Scree Plot for Factor Analysis	961
90.6.4. Factor Loadings	962
90.6.5. Factor Scores	964
90.6.6. Creating the Regression Model with Extracted Factors (from FA)	965
90.6.7. Factor Analysis: Creating the Regression Model with three Extracted Factors only	968
90.7. Summary: Comparing OLS, PCA, and Factor Analysis Models	970
90.7.1. Interpretation of the Regression Models	970
90.7.2. Differences Compared to the Standard OLS Model	971

Table of contents

90.7.3. Key Differences Between Loading Scores (PCA) and Factor Loadings (FA)	972
90.7.4. Advantages of Using PCA and FA	972
90.7.5. Disadvantages of Using PCA and FA	973
90.7.6. When to Use Which Method	973
90.8. Using Principal Components / Factors in Other Models	974
90.8.1. Random Forest Regressor with the Full Dataset	974
90.8.2. Random Forest Regressor with PCA Components	975
90.8.3. Random Forest Regressor with Extracted Factors (from FA)	976
90.8.4. Comparison of the Random Forest Models	976
90.9. Videos: Principal Component Analysis (PCA)	978
90.10Jupyter Notebook	978
91. Regression	979
91.1. Supervised and Unsupervised Learning	979
91.1.1. Starting point	979
91.1.2. Philosophy	979
91.1.3. Supervised Learning	979
91.1.4. Unsupervised Learning	980
91.2. Linear Regression	981
91.2.1. The main ideas of fitting a line to data (The main ideas of least squares and linear regression.)	981
91.2.2. Optimal Predictor	982
91.2.3. Curse of Dimensionality and Parametric Models	982
91.2.4. Assessing Model Accuracy and Bias-Variance Trade-off	984
91.3. Multiple Regression	985
91.4. R-squared	985
91.4.1. R-Squared in Simple Linear Regression	985
91.5. Assessing Confounding Effects in Multiple Regression	986
91.6. Cross-Validation	989
92. Classification	991
92.1. Classification: Some Details	991
92.2. k-Nearest Neighbor Classification	992
92.2.1. Minkowski Distance	992
92.3. Decision and Classification Trees	993
92.3.1. Decision Trees	993
92.3.2. Regression Trees	993
92.4. The Confusion Matrix	993
92.4.1. Sensitivity and Specificity	993
92.5. Naive Bayes	993
92.6. Gaussian Naive Bayes	993
93. Clustering	995
93.1. DBSCAN	995

93.2. k-Means Clustering	995
93.3. DDMO-Additional Videos	995
93.4. DDMO-Exercises	996
XV. Appendix	1003
94. Gaussian Processes—Some Background Information	1005
94.1. Gaussian Process Prior	1005
94.2. Covariance Function	1005
94.2.1. Positive Definiteness	1007
94.3. Construction of the Covariance Matrix	1007
94.4. Generation of Random Samples and Plotting the Realizations of the Random Function	1010
94.5. Properties of the 1d Example	1012
94.5.1. Several Bumps:	1012
94.5.2. Smoothness:	1013
94.5.3. Scale of Two:	1013
94.6. Jupyter Notebook	1015
References	1017

Preface

This document provides a comprehensive guide to optimization and hyperparameter tuning using the Sequential Parameter Optimization Toolbox (SPOT) for Python.

Part I.

**Sequential Parameter
Optimization Toolbox (SPOT)**

1. Introduction to SpotOptim

2. Examples

Let us consider the problem of minimizing the Rosenbrock function. This function (and its respective derivatives) is implemented in `rosen` (and `rosen_der`, `rosen_hess`) in the `scipy.optimize` module.

The function is usually evaluated on the hypercube $x_i \in [-5, 10]$, for all $i = 1, \dots, d$, although it may be restricted to the hypercube $x_i \in [-2.048, 2.048]$, for all $i = 1, \dots, d$, see <https://www.sfu.ca/~ssurjano/rosen.html>.

A simple application of the Nelder-Mead method is:

```
from scipy.optimize import minimize, rosen, rosen_der
x0 = [1.3, 0.7, 0.8, 1.9, 1.2]
res = minimize(rosen, x0, method='Nelder-Mead', tol=1e-6,
               bounds = [(-2.048, 2.048)] * 5)
res.x
```

```
array([1.00000002, 1.00000002, 1.00000007, 1.00000015, 1.00000028])
```

Now using the BFGS algorithm, using the first derivative and a few options:

```
res = minimize(rosen, x0, method='BFGS', jac=rosen_der,
               options={'gtol': 1e-6, 'disp': True})
res.x
print(res.message)
```

```
Optimization terminated successfully.
      Current function value: 0.000000
      Iterations: 26
      Function evaluations: 31
      Gradient evaluations: 31
Optimization terminated successfully.
```

Now, let's see how to solve the same problem using **SpotOptim**, which uses a Surrogate-Model Based Optimization (SMBO) approach. Unlike `minimize`, `SpotOptim` requires bounds as it samples the search space globally.

2. Examples

```
import numpy as np
from spotoptim import SpotOptim
from scipy.optimize import rosen

# SpotOptim expects function input as a 2D array (batch of points)
# So we wrap the rosen function to handle (n_samples, n_dim) input
def rosen_batch(X):
    return np.array([rosen(x) for x in X])

# Define the optimizer
# We use a 5-dimensional problem similar to the example above
# Bounds are required for SpotOptim
optimizer = SpotOptim(
    fun=rosen_batch,
    bounds = [(-2.048, 2.048)] * 5,
    max_iter=50, # Total budget of function evaluations
    n_initial=10, # Initial random samples
    seed=42
)

# Run optimization
res = optimizer.optimize()
print(res.message)
```

```
Optimization terminated: maximum evaluations (50) reached
      Current function value: 3.593232
      Iterations: 40
      Function evaluations: 50
```

```
print(f"Best solution found: {res.x}")
print(f"Best objective value: {res.fun}")
print(f"Total evaluations: {res.nfev}")
```

```
Best solution found: [ 0.21649427  0.03937146  0.00069121 -0.00588232  0.01926206]
Best objective value: 3.5932321765816138
Total evaluations: 50
```

SpotOptim is particularly useful when the objective function is expensive to evaluate (e.g., simulations, hyperparameter tuning), as it builds a model to intelligently select the next point to evaluate, often finding good solutions with fewer function calls than gradient-free local search methods.

2.1. Complex Constrained Optimization: The Robot Arm

For a more challenging example, let's consider the `robot_arm_hard` problem. This is a 10-dimensional problem where a simulated robot arm must reach a target while avoiding obstacles in a maze-like configuration. It features “hard” constraints implemented as severe penalties, creating a rugged landscape that is difficult for local optimizers.

First, let's try solving it with **Scipy's** default local optimizer. Note that we need to wrap the function to return a scalar float, as `robot_arm_hard` returns an array.

```
import numpy as np
from scipy.optimize import minimize
from spotoptim.function.so import robot_arm_hard

# Wrapper for Scipy (expects scalar return)
def objective_scipy(x):
    # robot_arm_hard handles constraints internally via penalties
    return float(robot_arm_hard(x)[0])

# Starting point
np.random.seed(42)
x0 = np.random.rand(10)
print(f"Initial cost: {objective_scipy(x0):.4f}")

# Run generic minimization (Nelder-Mead is default for gradient-free)
res_scipy = minimize(objective_scipy, x0, method='Nelder-Mead', tol=1e-4,
                    bounds = [(0.0, 1.0)] * 10)

print(f"Scipy Best cost: {res_scipy.fun:.4f}")
print(f"Scipy Success: {res_scipy.success}")
```

```
Initial cost: 54.3293
Scipy Best cost: 11.0496
Scipy Success: False
```

The local optimizer often gets stuck in local optima created by the obstacles (penalties), failing to find a path to the target.

Now let's use **SpotOptim**. We don't need a wrapper because `SpotOptim` natively supports the batch-vectorized output of `robot_arm_hard`. We simply define the bounds for the 10 joint angles.

2. Examples

```
from spotoptim import SpotOptim

# Define bounds: 10 angles normalized to [0, 1]
bounds = [(0.0, 1.0)] * 10

optimizer = SpotOptim(
    fun=robot_arm_hard, # Native array-based function
    bounds=bounds,
    max_iter=50,         # Budget
    n_initial=20,        # More exploration for complex space
    seed=42,
    max_surrogate_points=30
)

res_spot = optimizer.optimize()

print(f"SpotOptim Best cost: {res_spot.fun:.4f}")
```

SpotOptim Best cost: 23.3587

SpotOptim's use of a surrogate model and global acquisition function allows it to better explore the landscape and often jump out of the local traps that catch the local optimizer, finding a significantly better configuration for the robot arm.

2.2. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

3. SpotOptim API Reference

4. SpotOptim.minimize

Minimization of scalar function of one or more variables using Surrogate-Model Based Optimization (SMBO).

4.1. API

```
spotoptim.SpotOptim(fun, bounds, max_iter=20, n_initial=10, ...).optimize()
```

4.2. Parameters

fun : callable The objective function to be minimized. **fun(x)** -> **float** where **x** is an array with shape (n,) and **fun** returns a scalar float.

bounds : sequence or **Bounds** Bounds for variables. There are two ways to specify the bounds:

1. Instance of **scipy.optimize.Bounds** class.
2. Sequence of (**min**, **max**) pairs for each element in **x**. **None** is not supported.

max_iter : int, optional Maximum number of iterations (function evaluations). Default is 20. Note: This includes the initial design points.

n_initial : int, optional Number of initial design points (Latin Hypercube Sampling). Default is 10. Must be less than **max_iter**.

seed : int, optional Random seed for reproducibility. Default is None.

surrogate : estimator object, optional Surrogate model to use. If None, a default Gaussian Process Regressor (Matern kernel) is used.

acquisition : str, optional Acquisition function to use. Default is “y” (best value). Options: “EI” (Expected Improvement), “LCB” (Lower Confidence Bound), etc.

var_type : list of str, optional List specifying the type of each variable:

- “float” (continuous, default)
- “int” (integer)
- “factor” (categorical)

4. *SpotOptim.minimize*

var__name : list of str, optional List of names for each variable. Default is ["x0", "x1", ...].

4.3. Returns

res : OptimizeResult The optimization result represented as a `scipy.optimize.OptimizeResult` object. Important attributes are:

- **x**: the solution array
- **fun**: the value of the objective function at the solution
- **success**: a boolean flag indicating if the optimizer exited successfully
- **message**: a description of the cause of the termination
- **nfev**: number of evaluations of the objective function
- **nit**: number of iterations

4.4. See Also

`scipy.optimize.minimize` : Standard Scipy minimization. `sklearn.gaussian_process.GaussianProcess` : The default surrogate model.

4.5. Notes

This section describes the available solvers that can be selected by the ‘method’ parameter of the minimize function. However, **SpotOptim** uses a fixed SMBO method (Surrogate-Model Based Optimization) which iteratively: 1. Builds a surrogate model (e.g., Kriging/Gaussian Process) on available data. 2. Optimizes an acquisition function (e.g., Expected Improvement) to select the next point. 3. Evaluates the objective function at the new point. 4. Updates the model and repeats.

This method is particularly effective for expensive-to-evaluate black-box functions.

4.6. Jupyter Notebook

i Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

5. SpotOptim Internal Methods Examples

5.1. Setup: Import Required Packages

First, we import all necessary packages for the examples.

```
# Core imports
import numpy as np
import time
from scipy.optimize import OptimizeResult

# SpotOptim
from spotoptim import SpotOptim

# Surrogate models
from sklearn.gaussian_process import GaussianProcessRegressor
from sklearn.gaussian_process.kernels import RBF, ConstantKernel

# Set random seed for reproducibility
np.random.seed(42)

print("All packages imported successfully!")
```

All packages imported successfully!

5.2. 1. Main Optimization Method: `optimize()`

The `optimize()` method is the main entry point for running the optimization process. It coordinates all other methods in the optimization workflow:

1. **Initial Design Phase:** `get_initial_design()`, `curate_initial_design()`, `rm_initial_design_NA_values()`, `check_size_initial_design()`, `get_best_xy_initial_design()`
2. **Main Loop:** Surrogate fitting, OCBA application, point suggestion, evaluation
3. **Termination:** `determine_termination()`

5. SpotOptim Internal Methods Examples

Let's see a complete optimization example:

```
# Define a simple quadratic function
def sphere(X):
    """Sphere function:  $f(x) = \sum(x^2)$ """
    X = np.atleast_2d(X)
    return np.sum(X**2, axis=1)

# Create optimizer
opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    n_initial=5,
    max_iter=20,
    verbose=True
)

# Run optimization
result = opt.optimize()

print(f"\nBest point found: {result.x}")
print(f"Best value: {result.fun:.6f}")
print(f"Total evaluations: {result.nfev}")
print(f"Sequential iterations: {result.nit}")
print(f"Success: {result.success}")
print(f"Message: {result.message}")
```

```
TensorBoard logging disabled
Initial best:  $f(x) = 0.799873$ 
Iter 1 | Best: 0.799873 | Curr: 1.034112 | Rate: 0.00 | Evals: 30.0%
Iter 2 | Best: 0.714746 | Rate: 0.50 | Evals: 35.0%
Iter 3 | Best: 0.078402 | Rate: 0.67 | Evals: 40.0%
Iter 4 | Best: 0.078402 | Curr: 0.100492 | Rate: 0.50 | Evals: 45.0%
Iter 5 | Best: 0.018846 | Rate: 0.60 | Evals: 50.0%
Iter 6 | Best: 0.000245 | Rate: 0.67 | Evals: 55.0%
Iter 7 | Best: 0.000014 | Rate: 0.71 | Evals: 60.0%
Iter 8 | Best: 0.000012 | Rate: 0.75 | Evals: 65.0%
Iter 9 | Best: 0.000011 | Rate: 0.78 | Evals: 70.0%
Iter 10 | Best: 0.000011 | Curr: 0.000011 | Rate: 0.70 | Evals: 75.0%
Iter 11 | Best: 0.000010 | Rate: 0.73 | Evals: 80.0%
Iter 12 | Best: 0.000010 | Curr: 0.000010 | Rate: 0.67 | Evals: 85.0%
Iter 13 | Best: 0.000009 | Rate: 0.69 | Evals: 90.0%
Iter 14 | Best: 0.000009 | Rate: 0.71 | Evals: 95.0%
Iter 15 | Best: 0.000008 | Rate: 0.73 | Evals: 100.0%
```

5.2. 1. Main Optimization Method: `optimize()`

```
Best point found: [-0.00097017 -0.00272042]
Best value: 0.000008
Total evaluations: 20
Sequential iterations: 15
Success: True
Message: Optimization terminated: maximum evaluations (20) reached
        Current function value: 0.000008
        Iterations: 15
        Function evaluations: 20
```

5.2.1. Changing the Optimizer

```
# Create optimizer
opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    n_initial=5,
    max_iter=20,
    verbose=True,
    acquisition_optimizer='tricands',
)

# Run optimization
result = opt.optimize()

print(f"\nBest point found: {result.x}")
print(f"Best value: {result.fun:.6f}")
print(f"Total evaluations: {result.nfev}")
print(f"Sequential iterations: {result.nit}")
print(f"Success: {result.success}")
print(f"Message: {result.message}")
```

```
TensorBoard logging disabled
Initial best: f(x) = 6.700842
Iter 1 | Best: 3.568333 | Rate: 1.00 | Evals: 30.0%
Iter 2 | Best: 3.568333 | Curr: 3.857658 | Rate: 0.50 | Evals: 35.0%
Iter 3 | Best: 3.103296 | Rate: 0.67 | Evals: 40.0%
Iter 4 | Best: 0.553412 | Rate: 0.75 | Evals: 45.0%
Iter 5 | Best: 0.553412 | Curr: 1.113457 | Rate: 0.60 | Evals: 50.0%
Iter 6 | Best: 0.553412 | Curr: 0.621437 | Rate: 0.50 | Evals: 55.0%
Iter 7 | Best: 0.209945 | Rate: 0.57 | Evals: 60.0%
```

5. SpotOptim Internal Methods Examples

```
Iter 8 | Best: 0.068916 | Rate: 0.62 | Evals: 65.0%
Iter 9 | Best: 0.039700 | Rate: 0.67 | Evals: 70.0%
Iter 10 | Best: 0.039700 | Curr: 0.082609 | Rate: 0.60 | Evals: 75.0%
Iter 11 | Best: 0.039700 | Curr: 0.047742 | Rate: 0.55 | Evals: 80.0%
Iter 12 | Best: 0.039700 | Curr: 0.046135 | Rate: 0.50 | Evals: 85.0%
Iter 13 | Best: 0.039700 | Curr: 0.042580 | Rate: 0.46 | Evals: 90.0%
Iter 14 | Best: 0.039700 | Curr: 0.040970 | Rate: 0.43 | Evals: 95.0%
Iter 15 | Best: 0.039700 | Curr: 0.040370 | Rate: 0.40 | Evals: 100.0%
```

```
Best point found: [ 0.19777919 -0.02414536]
Best value: 0.039700
Total evaluations: 20
Sequential iterations: 15
Success: True
Message: Optimization terminated: maximum evaluations (20) reached
        Current function value: 0.039700
        Iterations: 15
        Function evaluations: 20
```

5.3. 2. Initial Design Methods

These methods handle the creation and validation of the initial design of experiments.

5.3.1. 2.1 get_initial_design()

Purpose: Generate or process initial design points

Used by: optimize() method at the start

Calls: generate_initial_design() if X0 is None

This method handles three scenarios: 1. Generate LHS design when X0=None 2. Include starting point x0 if provided 3. Transform user-provided X0

```
# Example 1: Generate default LHS design
opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    n_initial=10,
    seed=42
)
X0 = opt.get_initial_design()
print(f"Generated LHS design shape: {X0.shape}")
print(f"First 3 points:\n{X0[:3]}")
```



```
# Example 2: With starting point x0
opt_x0 = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    n_initial=10,
    x0=[0.0, 0.0],
    seed=42
)
X0_with_x0 = opt_x0.get_initial_design()
print(f"\nDesign with x0, first point: {X0_with_x0[0]}")

# Example 3: Provide custom initial design
X0_custom = np.array([[0, 0], [1, 1], [2, 2]])
X0_processed = opt.get_initial_design(X0_custom)
print(f"\nCustom design shape: {X0_processed.shape}")
```

Generated LHS design shape: (10, 2)

First 3 points:

```
[[ 2.08325584  1.08901333]
 [ 0.1234075  -4.30931841]
 [ 3.04534393  0.82490542]]
```

Design with x0, first point: [0. 0.]

Custom design shape: (3, 2)

5.3.2. 2.2 curate_initial_design()

Purpose: Remove duplicates and ensure sufficient unique points

Used by: optimize() after get_initial_design()

Handles: Duplicate removal, point generation, repetition for noisy functions

```
# Example 1: Remove duplicates
opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    n_initial=10,
    seed=42
)
X0_with_dups = np.array([[1, 2], [1, 2], [3, 4], [3, 4], [5, 6]])
X0_curated = opt.curate_initial_design(X0_with_dups)
print(f"Original points: {len(X0_with_dups)}")
```

5. SpotOptim Internal Methods Examples

```
print(f"After removing duplicates: {len(X0_curated)}")

# Example 2: With repeats for noisy functions
opt_repeat = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    n_initial=5,
    repeats_initial=3, # Repeat each point 3 times
    seed=42
)
X0 = np.array([[1, 2], [3, 4], [5, 6], [7, 8], [9, 10]])
X0_repeated = opt_repeat.curate_initial_design(X0)
print(f"\nOriginal points: {len(X0)}")
print(f"After repeating (3x): {len(X0_repeated)}")
```

Original points: 5
After removing duplicates: 10

Original points: 5
After repeating (3x): 15

5.3.3. 2.3 rm_initial_design_NA_values()

Purpose: Remove NaN/inf values from initial design evaluations

Used by: optimize() after evaluating initial design

Returns: Cleaned arrays and original count

```
opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    n_initial=10,
    seed=42
)

# Simulate initial design with some NaN values
X0 = np.array([[1, 2], [3, 4], [5, 6]])
y0 = np.array([5.0, np.nan, np.inf])

X0_clean, y0_clean, n_eval = opt.rm_initial_design_NA_values(X0, y0)

print(f"Original evaluations: {n_eval}")
print(f"Valid points remaining: {X0_clean.shape[0]}")
```

```
print(f"Clean X0:\n{X0_clean}")
print(f"Clean y0: {y0_clean}")
```

```
Original evaluations: 3
Valid points remaining: 1
Clean X0:
[[1 2]]
Clean y0: [5.]
```

5.3.4. 2.4 check_size_initial_design()

Purpose: Validate sufficient points for surrogate fitting

Used by: optimize() after handling NaN values

Raises: ValueError if insufficient points

```
opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    n_initial=10,
    seed=42
)

# Example 1: Sufficient points - no error
y0_sufficient = np.array([1.0, 2.0, 3.0, 4.0, 5.0])
try:
    opt.check_size_initial_design(y0_sufficient, n_evaluated=10)
    print("x Sufficient points - validation passed")
except ValueError as e:
    print(f" Error: {e}")

# Example 2: Insufficient points - raises error
y0_insufficient = np.array([1.0]) # Only 1 point, need at least 3 for 2D
try:
    opt.check_size_initial_design(y0_insufficient, n_evaluated=10)
    print("x Validation passed")
except ValueError as e:
    print(f" Expected error: {e}")
```

```
x Sufficient points - validation passed
```

```
Expected error: Insufficient valid initial design points: only 1 finite value(s) out of 10 evaluated
```

5. SpotOptim Internal Methods Examples

5.3.5. 2.5 get_best_xy_initial_design()

Purpose: Determine and store the best point from initial design

Used by: optimize() after initial design evaluation

Updates: self.best_x_ and self.best_y_ attributes

```
opt = SpotOptim(  
    fun=sphere,  
    bounds=[(-5, 5), (-5, 5)],  
    n_initial=5,  
    verbose=True,  
    seed=42  
)  
  
# Simulate initial design (normally done in optimize())  
opt.X_ = np.array([[1, 2], [0, 0], [2, 1]])  
opt.y_ = np.array([5.0, 0.0, 5.0])  
  
opt.get_best_xy_initial_design()  
  
print(f"\nBest x from initial design: {opt.best_x_}")  
print(f"Best y from initial design: {opt.best_y_}")
```

```
TensorBoard logging disabled  
Initial best: f(x) = 0.000000
```

```
Best x from initial design: [0 0]  
Best y from initial design: 0.0
```

5.4. 3. Surrogate Model Methods

These methods handle surrogate model operations during the optimization loop.

5.4.1. 3.1 fit_surrogate()

Purpose: Fit surrogate model to data

Used by: optimize() in main loop

Calls: fit_selection_dispatcher() if max_surrogate_points exceeded

```

opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    max_surrogate_points=10,
    seed=42
)

# Generate some training data
X = np.random.rand(50, 2) * 10 - 5 # 50 points in [-5, 5]
y = np.sum(X**2, axis=1)

# Fit surrogate (will select 10 best points)
opt.fit_surrogate(X, y)

print("Surrogate fitted successfully!")
print("Surrogate model: {type(opt.surrogate).__name__}")

```

Surrogate fitted successfully!
 Surrogate model: GaussianProcessRegressor

5.4.2. 3.2 `_predict_with_uncertainty()`

Purpose: Predict with uncertainty estimates, handling surrogates without `return_std`

Used by: `_acquisition_function()` and `plot_surrogate()`

Returns: Predictions and standard deviations

```

opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    surrogate=GaussianProcessRegressor(),
    seed=42
)

# Train surrogate
X_train = np.array([[0, 0], [1, 1], [2, 2]])
y_train = np.array([0, 2, 8])
opt.fit_surrogate(X_train, y_train)

# Predict on new points
X_test = np.array([[1.5, 1.5], [3.0, 3.0]])
preds, stds = opt._predict_with_uncertainty(X_test)

```

5. SpotOptim Internal Methods Examples

```
print(f"Test points:\n{X_test}")
print(f"Predictions: {preds}")
print(f"Uncertainties: {stds}")
```

```
Test points:
[[1.5 1.5]
 [3.  3. ]]
Predictions: [0. 0.]
Uncertainties: [4.76095228 4.76095228]
```

5.4.3. 3.3 _acquisition_function()

Purpose: Compute acquisition function value

Used by: `suggest_next_infill_point()` for optimization

Calls: `_predict_with_uncertainty()`

Supports: Expected Improvement (EI), Probability of Improvement (PI), Mean prediction

```
opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    surrogate=GaussianProcessRegressor(),
    acquisition='ei',
    seed=42
)

# Setup
X_train = np.array([[0, 0], [1, 1], [2, 2]])
y_train = np.array([0, 2, 8])
opt.fit_surrogate(X_train, y_train)
opt.y_ = y_train # Needed for acquisition function

# Evaluate acquisition function
x_eval = np.array([1.5, 1.5])
acq_value = opt._acquisition_function(x_eval)

print(f"Point to evaluate: {x_eval}")
print(f"Acquisition function value (EI): {acq_value:.6f}")
print(f"(Lower is better for minimization)")
```

```
Point to evaluate: [1.5 1.5]
Acquisition function value (EI): -1.899345
(Lower is better for minimization)
```

5.4.4. 3.4 suggest_next_infill_point()

Purpose: Suggest next point to evaluate using acquisition function optimization

Used by: optimize() in main loop

Calls: _acquisition_function(), _handle_acquisition_failure() if needed

Handles: Integer/factor rounding, duplicate avoidance

```
opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    surrogate=GaussianProcessRegressor(),
    acquisition='ei',
    seed=42
)

# Setup
X_train = np.array([[0, 0], [1, 1], [2, 2]])
y_train = np.array([0, 2, 8])
opt.fit_surrogate(X_train, y_train)
opt.X_ = X_train
opt.y_ = y_train

# Suggest next point
x_next = opt.suggest_next_infill_point()

print(f"Next point to evaluate: {x_next}")
print(f"Expected to be between known points or in unexplored regions")
```

```
Next point to evaluate: [[ 5.84038953e-04 -1.62016871e+00]]
Expected to be between known points or in unexplored regions
```

5.5. 4. OCBA Method (Optimal Computing Budget Allocation)

5.5.1. 4.1 apply_ocba()

Purpose: Apply OCBA for noisy functions to determine which points to re-evaluate

Used by: optimize() in main loop when noise=True and ocba_delta > 0

Returns: Points to re-evaluate or None

5. SpotOptim Internal Methods Examples

```
# Example with OCBA enabled
opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    repeats_initial=2, # Enable noise handling
    ocba_delta=5,
    verbose=True,
    seed=42
)

# Simulate optimization state
opt.mean_X = np.array([[1, 2], [0, 0], [2, 1], [1, 1]])
opt.mean_y = np.array([5.0, 0.1, 5.0, 2.0])
opt.var_y = np.array([0.1, 0.05, 0.15, 0.08])

X_ocba = opt.apply_ocba()

if X_ocba is not None:
    print(f"\nOCBA returned {X_ocba.shape[0]} points for re-evaluation")
else:
    print("\nOCBA: No re-evaluation needed")
```

TensorBoard logging disabled

```
In get_ocba():
means: [5.  0.1 5.  2. ]
vars:  [0.1  0.05 0.15 0.08]
delta: 5
n_designs: 4
Ratios: [0.18794252 0.81801772 0.28191379 1.          ]
Best: 1, Second best: 3
      OCBA: Adding 5 re-evaluation(s)
```

OCBA returned 5 points for re-evaluation

5.6. 5. NaN/Inf Handling Methods

5.6.1. 5.1 apply_penalty_NA()

Purpose: Replace NaN and infinite values with penalty plus random noise

Used by: `_handle_NA_new_points()` and indirectly by `optimize()`

Algorithm: $\text{penalty} = \max(\text{finite_y}) + 3 * \text{std}(\text{finite_y}) + \text{noise}$


```

opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5)],
    seed=42
)

# Historical values (used for penalty computation)
y_hist = np.array([1.0, 2.0, 3.0, 5.0])

# New evaluations with NaN/inf
y_new = np.array([4.0, np.nan, np.inf])

y_clean = opt.apply_penalty_NA(y_new, y_history=y_hist)

print(f"Original: {y_new}")
print(f"After penalty: {y_clean}")
print(f"All finite: {np.all(np.isfinite(y_clean))}")
print(f"Penalty values > max(hist): {y_clean[1] > np.max(y_hist)}")

```

```

Original: [ 4. nan inf]
After penalty: [ 4.          10.10964895 10.18824424]
All finite: True
Penalty values > max(hist): True

```

5.6.2. 5.2 remove_nan()

Purpose: Remove rows where y contains NaN or inf values

Used by: optimize() after function evaluations

Returns: Cleaned arrays

```

opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5)],
    seed=42
)

X = np.array([[1, 2], [3, 4], [5, 6]])
y = np.array([1.0, np.nan, np.inf])

X_clean, y_clean = opt.remove_nan(X, y, stop_on_zero_return=False)

print(f"Original X shape: {X.shape}")

```

5. SpotOptim Internal Methods Examples

```
print(f"Clean X shape: {X_clean.shape}")
print(f"Clean X:\n{X_clean}")
print(f"Clean y: {y_clean}")
```

```
Original X shape: (3, 2)
Clean X shape: (1, 2)
Clean X:
[[1 2]]
Clean y: [1.]
```

5.6.3. 5.3 _handle_NA_new_points()

Purpose: Handle NaN/inf values in new evaluation points during main loop

Used by: optimize() after evaluating new points

Calls: apply_penalty_NA() and remove_nan()

Returns: None, None if all evaluations invalid (skip iteration)

```
opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    n_initial=5,
    verbose=True,
    seed=42
)

# Simulate optimization state
opt.y_ = np.array([1.0, 2.0, 3.0]) # Historical values
opt.n_iter_ = 1

# Case 1: Some valid values
print("Case 1: Some valid evaluations")
x_next = np.array([[1, 2], [3, 4], [5, 6]])
y_next = np.array([5.0, np.nan, 10.0])
x_clean, y_clean = opt._handle_NA_new_points(x_next, y_next)
print(f"Valid points remaining: {x_clean.shape[0] if x_clean is not None else 0}")

# Case 2: All invalid - should return None
print("\nCase 2: All invalid evaluations")
x_all_bad = np.array([[1, 2], [3, 4]])
y_all_bad = np.array([np.nan, np.inf])
x_clean, y_clean = opt._handle_NA_new_points(x_all_bad, y_all_bad)
print(f"Result: {'None (skip iteration)' if x_clean is None else 'Valid points'}")
```

TensorBoard logging disabled

Case 1: Some valid evaluations

Warning: Removed 1 sample(s) with NaN/inf values

Valid points remaining: 2

Case 2: All invalid evaluations

Warning: All objective function values are NaN or inf. Returning empty arrays.

Warning: All new evaluations were NaN/inf, skipping iteration 1

Result: None (skip iteration)

5.7. 6. Main Loop Update Methods

5.7.1. 6.1 _update_best_main_loop()

Purpose: Update best solution found during main optimization loop

Used by: optimize() after each iteration

Updates: self.best_x_ and self.best_y_ if improvement found

```
opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    n_initial=5,
    verbose=True,
    seed=42
)

# Simulate optimization state
opt.n_iter_ = 1
opt.best_x_ = np.array([1.0, 1.0])
opt.best_y_ = 2.0

# Case 1: New best found
print("Case 1: New best found")
x_new = np.array([[0.1, 0.1], [0.5, 0.5]])
y_new = np.array([0.02, 0.5])
opt._update_best_main_loop(x_new, y_new)
print(f"Updated best_y: {opt.best_y_}\n")

# Case 2: No improvement
print("Case 2: No improvement")
opt.n_iter_ = 2
x_no_improve = np.array([[1.5, 1.5]])
y_no_improve = np.array([4.5])
```

5. SpotOptim Internal Methods Examples

```
opt._update_best_main_loop(x_no_improve, y_no_improve)
print(f"Best_y unchanged: {opt.best_y_}")
```

TensorBoard logging disabled

Case 1: New best found

Iter 1 | Best: 0.020000 | Rate: 0.00 | Evals: 0.0%

Updated best_y: 0.02

Case 2: No improvement

Iter 2 | Best: 0.020000 | Curr: 4.500000 | Rate: 0.00 | Evals: 0.0%

Best_y unchanged: 0.02

5.8. 7. Termination Method

5.8.1. 7.1 determine_termination()

Purpose: Determine termination reason for optimization

Used by: optimize() at the end

Checks: Max iterations, time limit, or successful completion

```
opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=20,
    max_time=10.0,
    seed=42
)

# Case 1: Maximum evaluations reached
print("Case 1: Maximum evaluations reached")
opt.y_ = np.zeros(20)
start_time = time.time()
msg = opt.determine_termination(start_time)
print(f"Message: {msg}\n")

# Case 2: Time limit exceeded (simulated)
print("Case 2: Time limit exceeded")
opt.y_ = np.zeros(10)
start_time = time.time() - 700 # 11.67 minutes ago
msg = opt.determine_termination(start_time)
print(f"Message: {msg}\n")
```

```
# Case 3: Successful completion
print("Case 3: Successful completion")
opt.y_ = np.zeros(10)
start_time = time.time()
msg = opt.determine_termination(start_time)
print(f"Message: {msg}")
```

Case 1: Maximum evaluations reached

Message: Optimization terminated: maximum evaluations (20) reached

Case 2: Time limit exceeded

Message: Optimization terminated: time limit (10.00 min) reached

Case 3: Successful completion

Message: Optimization finished successfully

5.9. 8. Utility Methods

5.9.1. 8.1 select_new()

Purpose: Select rows from A that are not in X (avoid duplicate evaluations)

Used by: suggest_next_infill_point() to ensure new points are different from evaluated points

```
opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5)],
    seed=42
)

A = np.array([[1, 2], [3, 4], [5, 6]])
X = np.array([[3, 4], [7, 8]])

new_A, is_new = opt.select_new(A, X)

print(f"Candidate points A:\n{A}")
print(f"\nKnown points X:\n{X}")
print(f"\nNew points from A:\n{new_A}")
print(f"\nIs new mask: {is_new}")
```

Candidate points A:

5. SpotOptim Internal Methods Examples

```
[[1 2]
 [3 4]
 [5 6]]
```

Known points X:

```
[[3 4]
 [7 8]]
```

New points from A:

```
[[1 2]
 [5 6]]
```

Is new mask: [True False True]

5.9.2. 8.2 repair_non_numeric()

Purpose: Round non-numeric values (int, factor) to integers

Used by: Various methods to ensure integer/factor variables have valid values

```
opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    var_type=['int', 'float'],
    seed=42
)

X = np.array([[1.2, 2.5], [3.7, 4.1], [5.9, 6.8]])
X_repaired = opt.repair_non_numeric(X.copy(), opt.var_type)

print(f"Original X:\n{X}")
print(f"\nRepaired X (first column rounded to int):\n{X_repaired}")
```

Original X:

```
[[1.2 2.5]
 [3.7 4.1]
 [5.9 6.8]]
```

Repaired X (first column rounded to int):

```
[[1. 2.5]
 [4. 4.1]
 [6. 6.8]]
```

5.9.3. 8.3 map_to_factor_values()**Purpose:** Map internal integer values to original factor strings**Used by:** optimize() when preparing results for user**Handles:** Factor (categorical) variables

```

opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), ('red', 'green', 'blue')],
    var_type=['float', 'factor'],
    seed=42
)

# Internal representation (integers for factors)
X_internal = np.array([[1.0, 0], [2.0, 1], [3.0, 2]])
X_mapped = opt.map_to_factor_values(X_internal)

print(f"Internal representation:\n{X_internal}")
print(f"\nMapped to factor values:\n{X_mapped}")

```

Internal representation:

```

[[1. 0.]
 [2. 1.]
 [3. 2.]]

```

Mapped to factor values:

```

[[1.0 'red']
 [2.0 'green']
 [3.0 'blue']]

```

```

/var/folders/dw/pvtj6mt91znd0hftcztb0k00000gn/T/ipykernel_31961/1870456492.py:1: UserWarning: Factor
agnostic) factor metrics. Factor integer codes will be treated as ordinal numbers, which may mislead
aware Kriging surrogate: Kriging(metric_factorial='hamming').
    opt = SpotOptim(

```

5.10. 9. Integration Examples**5.10.1. 9.1 Optimization with Custom Initial Design**

5. SpotOptim Internal Methods Examples

```
opt = SpotOptim(  
    fun=sphere,  
    bounds=[(-5, 5), (-5, 5)],  
    max_iter=15,  
    seed=42,  
    verbose=True  
)  
  
# Provide custom initial design  
X0 = np.array([[0, 0], [1, 1], [2, 2], [-1, -1], [-2, -2]])  
result = opt.optimize(X0=X0)  
  
print(f"\nBest point: {result.x}")  
print(f"Best value: {result.fun:.6f}")
```

```
TensorBoard logging disabled  
Note: Initial design size (5) is smaller than requested (10) due to NaN/inf values  
Initial best: f(x) = 0.000000  
Iter 1 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.00 | Evals: 40.0%  
Iter 2 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.00 | Evals: 46.7%  
Optimizer candidate 1/3 was duplicate/invalid.  
Iter 3 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.00 | Evals: 53.3%  
Optimizer candidate 1/3 was duplicate/invalid.  
Iter 4 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.00 | Evals: 60.0%  
Iter 5 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.00 | Evals: 66.7%  
Iter 6 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.00 | Evals: 73.3%  
Iter 7 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.00 | Evals: 80.0%  
Iter 8 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.00 | Evals: 86.7%  
Iter 9 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.00 | Evals: 93.3%  
Optimizer candidate 1/3 was duplicate/invalid.  
Optimizer candidate 2/3 was duplicate/invalid.  
Iter 10 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.00 | Evals: 100.0%  
  
Best point: [0. 0.]  
Best value: 0.000000
```

5.10.2. 9.2 Optimization with Starting Point

```
opt = SpotOptim(  
    fun=sphere,  
    bounds=[(-5, 5), (-5, 5)],
```



```

max_iter=15,
n_initial=5,
x0=[0.5, 0.5], # Starting point near optimum
seed=42,
verbose=True
)

result = opt.optimize()

print(f"\nBest point: {result.x}")
print(f"Best value: {result.fun:.6f}")

```

Starting point x0 validated and processed successfully.

Original scale: [0.5 0.5]

Internal scale: [0.5 0.5]

TensorBoard logging disabled

Including 1 starting points from x0 in initial design.

Initial best: f(x) = 0.500000

Iter 1	Best: 0.440515	Rate: 1.00	Evals: 40.0%
Iter 2	Best: 0.314161	Rate: 1.00	Evals: 46.7%
Iter 3	Best: 0.058028	Rate: 1.00	Evals: 53.3%
Iter 4	Best: 0.037546	Rate: 1.00	Evals: 60.0%
Iter 5	Best: 0.022557	Rate: 1.00	Evals: 66.7%
Iter 6	Best: 0.000061	Rate: 1.00	Evals: 73.3%
Iter 7	Best: 0.000006	Rate: 1.00	Evals: 80.0%
Iter 8	Best: 0.000003	Rate: 1.00	Evals: 86.7%
Iter 9	Best: 0.000002	Rate: 1.00	Evals: 93.3%
Iter 10	Best: 0.000002	Rate: 1.00	Evals: 100.0%

Best point: [0.0006169 0.00117845]

Best value: 0.000002

5.10.3. 9.3 Optimization with Integer Variables

```

opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    var_type=['int', 'int'],
    max_iter=20,
    n_initial=10,
    seed=42,
    verbose=True
)

```

5. SpotOptim Internal Methods Examples

```
)  
  
result = opt.optimize()  
  
print(f"\nBest point: {result.x}")  
print(f"Best value: {result.fun:.6f}")  
print(f"Point has integers: {np.all(result.x == np.round(result.x))}")
```

```
TensorBoard logging disabled  
Initial best: f(x) = 5.000000  
Iter 1 | Best: 0.000000 | Rate: 1.00 | Evals: 55.0%  
Optimizer candidate 1/3 was duplicate/invalid.  
Optimizer candidate 2/3 was duplicate/invalid.  
Optimizer candidate 3/3 was duplicate/invalid.  
Fallback attempt 1/10: Using fallback strategy  
Acquisition failure: Using random space-filling design as fallback.  
Iter 2 | Best: 0.000000 | Curr: 17.000000 | Rate: 0.50 | Evals: 60.0%  
Optimizer candidate 1/3 was duplicate/invalid.  
Optimizer candidate 2/3 was duplicate/invalid.  
Optimizer candidate 3/3 was duplicate/invalid.  
Fallback attempt 1/10: Using fallback strategy  
Acquisition failure: Using random space-filling design as fallback.  
Iter 3 | Best: 0.000000 | Curr: 13.000000 | Rate: 0.33 | Evals: 65.0%  
Optimizer candidate 1/3 was duplicate/invalid.  
Optimizer candidate 2/3 was duplicate/invalid.  
Optimizer candidate 3/3 was duplicate/invalid.  
Fallback attempt 1/10: Using fallback strategy  
Acquisition failure: Using random space-filling design as fallback.  
Iter 4 | Best: 0.000000 | Curr: 41.000000 | Rate: 0.25 | Evals: 70.0%  
Optimizer candidate 1/3 was duplicate/invalid.  
Optimizer candidate 2/3 was duplicate/invalid.  
Optimizer candidate 3/3 was duplicate/invalid.  
Fallback attempt 1/10: Using fallback strategy  
Acquisition failure: Using random space-filling design as fallback.  
Iter 5 | Best: 0.000000 | Curr: 20.000000 | Rate: 0.20 | Evals: 75.0%  
Optimizer candidate 1/3 was duplicate/invalid.  
Optimizer candidate 2/3 was duplicate/invalid.  
Optimizer candidate 3/3 was duplicate/invalid.  
Fallback attempt 1/10: Using fallback strategy  
Acquisition failure: Using random space-filling design as fallback.  
Iter 6 | Best: 0.000000 | Curr: 25.000000 | Rate: 0.17 | Evals: 80.0%  
Optimizer candidate 1/3 was duplicate/invalid.  
Optimizer candidate 2/3 was duplicate/invalid.  
Optimizer candidate 3/3 was duplicate/invalid.
```

```

Fallback attempt 1/10: Using fallback strategy
Acquisition failure: Using random space-filling design as fallback.
Iter 7 | Best: 0.000000 | Curr: 17.000000 | Rate: 0.14 | Evals: 85.0%
Optimizer candidate 1/3 was duplicate/invalid.
Optimizer candidate 2/3 was duplicate/invalid.
Optimizer candidate 3/3 was duplicate/invalid.
Fallback attempt 1/10: Using fallback strategy
Acquisition failure: Using random space-filling design as fallback.
Iter 8 | Best: 0.000000 | Curr: 13.000000 | Rate: 0.12 | Evals: 90.0%
Optimizer candidate 1/3 was duplicate/invalid.
Optimizer candidate 2/3 was duplicate/invalid.
Optimizer candidate 3/3 was duplicate/invalid.
Fallback attempt 1/10: Using fallback strategy
Acquisition failure: Using random space-filling design as fallback.
Iter 9 | Best: 0.000000 | Curr: 29.000000 | Rate: 0.11 | Evals: 95.0%
Optimizer candidate 1/3 was duplicate/invalid.
Optimizer candidate 2/3 was duplicate/invalid.
Optimizer candidate 3/3 was duplicate/invalid.
Fallback attempt 1/10: Using fallback strategy
Acquisition failure: Using random space-filling design as fallback.
Iter 10 | Best: 0.000000 | Curr: 16.000000 | Rate: 0.10 | Evals: 100.0%

Best point: [-0.  0.]
Best value: 0.000000
Point has integers: True

```

5.10.4. 9.4 Noisy Function Optimization with OCBA

```

# Noisy sphere function
def noisy_sphere(X):
    X = np.atleast_2d(X)
    return np.sum(X**2, axis=1) + np.random.normal(0, 0.1, X.shape[0])

opt = SpotOptim(
    fun=noisy_sphere,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=25,
    n_initial=10,
    repeats_initial=2, # Enable noise handling
    ocba_delta=5, # Enable OCBA with 5 re-evaluations
    seed=42,
    verbose=True
)

```

5. SpotOptim Internal Methods Examples

```
result = opt.optimize()

print(f"\nBest point: {result.x}")
print(f"Best value: {result.fun:.6f}")
print(f"Note: OCBA re-evaluated promising points to reduce uncertainty")
```

```
TensorBoard logging disabled
Initial best: f(x) = 3.734393, mean best: f(x) = 3.842155

In get_ocba():
means: [49.11419203 26.21685016 7.98802093 6.31788609 6.16395351 18.69399049
 3.84215472 5.54382744 9.93117406 36.60267168]
vars: [4.40284305e-03 6.35744853e-04 1.33640624e-08 3.37927306e-03
 1.64745915e-03 1.91555606e-03 1.16126758e-02 1.00799409e-03
 6.73858270e-13 2.56053422e-03]
delta: 5
n_designs: 10
Ratios: [6.17116575e-03 3.64806588e-03 2.23358257e-06 1.58383644e+00
 8.77929372e-01 2.49475944e-02 5.05750153e+00 1.00000000e+00
 5.22117127e-11 6.85366978e-03]
Best: 6, Second best: 7
OCBA: Adding 5 re-evaluation(s)
Iter 1 | Best: 0.027038 | Rate: 0.17 | Evals: 104.0% | Mean Best: 0.027038

Best point: [-0.11672289 0.04817731]
Best value: 0.027038
Note: OCBA re-evaluated promising points to reduce uncertainty
```

5.10.5. 9.5 Optimization with NaN Handling

```
# Function that sometimes returns NaN
def sometimes_nan(X):
    X = np.atleast_2d(X)
    y = np.sum(X**2, axis=1)
    # Return NaN for large values (penalty will be applied)
    y[y > 100] = np.nan
    return y

opt = SpotOptim(
    fun=sometimes_nan,
    bounds=[(-10, 10), (-10, 10)],
```

```

max_iter=20,
n_initial=10,
seed=42,
verbose=True
)

result = opt.optimize()

print(f"\nBest point: {result.x}")
print(f"Best value: {result.fun:.6f}")
print(f"Optimization succeeded despite NaN values: {result.success}")

```

TensorBoard logging disabled

Warning: 3 initial design point(s) returned NaN/inf and will be ignored (reduced from 10 to 7 points)

Note: Initial design size (7) is smaller than requested (10) due to NaN/inf values

Initial best: $f(x) = 15.702882$

Iter	Best	Curr	Rate	Evals
1	15.702882	16.692397	0.00	40.0%
2	9.868398		0.50	45.0%
3	1.155915		0.67	50.0%
4	0.113986		0.75	55.0%
5	0.003493		0.80	60.0%
6	0.000979		0.83	65.0%
7	0.000948		0.86	70.0%
8	0.000778		0.88	75.0%
9	0.000009		0.89	80.0%
10	0.000000		0.90	85.0%
11	0.000000	0.000000	0.82	90.0%
12	0.000000	0.000000	0.75	95.0%
13	0.000000	0.000000	0.69	100.0%

Best point: [2.13856406e-04 6.37169523e-05]

Best value: 0.000000

Optimization succeeded despite NaN values: True

5.11. Summary

This notebook demonstrated all major methods in SpotOptim with executable examples:

Core Flow:

1. **Initial Design:** Generate/process → Curate → Evaluate → Handle NaN → Validate → Get best

5. *SpotOptim Internal Methods Examples*

2. **Main Loop:** Fit surrogate → Apply OCBA → Suggest point → Evaluate → Handle NaN → Update best
3. **Termination:** Determine reason → Prepare results

Key Features:

- Automatic handling of NaN/inf values with penalties
- Support for noisy functions with OCBA re-evaluation
- Integer and factor variable support
- Flexible initial design (LHS, custom, or with starting point)
- Multiple acquisition functions (EI, PI, mean)
- Termination by max iterations or time limit

All examples can be run independently and demonstrate the modular design of SpotOptim!

5.12. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

6. Optimizing the Aircraft Wing Weight Example

i Note

- This section demonstrates optimization of the Aircraft Wing Weight Example (AWWE) function
- We compare three optimization methods:
 - **SpotOptim**: Bayesian optimization with Gaussian Process surrogate
 - **Nelder-Mead**: Derivative-free simplex method from `scipy.optimize`
 - **BFGS**: Quasi-Newton method from `scipy.optimize`
- The following Python packages are imported:

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from scipy.optimize import minimize
from spotoptim import SpotOptim
import time
import pprint
```

6.1. The AWWE Objective Function

We use the same AWWE function from Chapter 77, which models the weight of an unpainted light aircraft wing. The function accepts inputs in the unit cube $[0, 1]^9$ and returns the wing weight.

```
def wingwt(x):
    """
    Aircraft Wing Weight function.

    Args:
        x: array-like of 9 values in [0,1]
```

6. Optimizing the Aircraft Wing Weight Example

```
[Sw, Wfw, A, L, q, l, Rtc, Nz, Wdg]

Returns:
    Wing weight (scalar)
    """
    # Ensure x is a 2D array for batch evaluation
    x = np.atleast_2d(x)

    # Transform from unit cube to natural scales
    Sw = x[:, 0] * (200 - 150) + 150
    Wfw = x[:, 1] * (300 - 220) + 220
    A = x[:, 2] * (10 - 6) + 6
    L = (x[:, 3] * (10 - (-10)) - 10) * np.pi/180
    q = x[:, 4] * (45 - 16) + 16
    l = x[:, 5] * (1 - 0.5) + 0.5
    Rtc = x[:, 6] * (0.18 - 0.08) + 0.08
    Nz = x[:, 7] * (6 - 2.5) + 2.5
    Wdg = x[:, 8] * (2500 - 1700) + 1700

    # Calculate weight on natural scale
    W = 0.036 * Sw**0.758 * Wfw**0.0035 * (A/np.cos(L)**2)**0.6 * q**0.006
    W = W * l**0.04 * (100*Rtc/np.cos(L))**(-0.3) * (Nz*Wdg)**(0.49)

    return W.ravel()

# Wrapper for scipy.optimize (expects 1D input, returns scalar)
def wingwt_scipy(x):
    return float(wingwt(x.reshape(1, -1))[0])
```

6.2. Baseline Configuration

The baseline Cessna C172 Skyhawk configuration (coded in unit cube):

```
baseline_coded = np.array([0.48, 0.4, 0.38, 0.5, 0.62, 0.344, 0.4, 0.37, 0.38])
baseline_weight = wingwt(baseline_coded)[0]
print(f"Baseline wing weight: {baseline_weight:.2f} lb")
```

Baseline wing weight: 233.91 lb

6.3. Optimization Setup

We'll optimize the AWE function starting from the baseline configuration using three different methods:

1. **SpotOptim**: Bayesian optimization (good for expensive black-box functions)
2. **Nelder-Mead**: Derivative-free simplex method (robust but can be slow)
3. **BFGS**: Quasi-Newton method (fast but requires smooth functions)

```
# Starting point (baseline configuration)
x0 = baseline_coded.copy()

# Bounds for all methods (unit cube)
bounds = [(0, 1)] * 9

# Number of function evaluations budget
max_evals = 30

print(f"Starting point: {x0}")
print(f"Starting weight: {baseline_weight:.2f} lb")
```

```
Starting point: [0.48  0.4   0.38  0.5   0.62  0.344 0.4   0.37  0.38 ]
Starting weight: 233.91 lb
```

6.4. Method 1: SpotOptim (Surrogate Model Based Optimization)

```
# Start timing
start_time = time.time()

# Configure SpotOptim
optimizer_spot = SpotOptim(
    fun=wingwt,
    bounds=bounds,
    x0=None,
    max_iter=max_evals,
    n_initial=10, # Initial design points
    var_name=['Sw', 'Wfw', 'A', 'L', 'q', 'l', 'Rtc', 'Nz', 'Wdg'],
    acquisition='y', # ei: Expected Improvement
    max_surrogate_points=100,
    seed=42,
```

6. Optimizing the Aircraft Wing Weight Example

```
verbose=True,  
tensorboard_log=True,  
tensorboard_clean=True  
)
```

```
Removed old TensorBoard logs: runs/spotoptim_20260701_201813  
Cleaned 1 old TensorBoard log directory  
TensorBoard logging enabled: runs/spotoptim_20260701_201824
```

6.5. Design Table

```
pprint.pprint(optimizer_spot.get_design_table())
```

	name	type	lower	upper	default	transform
	Sw	float	0	1	0.5	-
	Wfw	float	0	1	0.5	-
	A	float	0	1	0.5	-
	L	float	0	1	0.5	-
	q	float	0	1	0.5	-
	l	float	0	1	0.5	-
	Rtc	float	0	1	0.5	-
	Nz	float	0	1	0.5	-
	Wdg	float	0	1	0.5	-

6.6. Run optimization

```
result_spot = optimizer_spot.optimize()
```

```
Initial best: f(x) = 205.911302  
Iter 1 | Best: 205.911302 | Curr: 207.213032 | Rate: 0.00 | Evals: 36.7%  
Iter 2 | Best: 123.245058 | Rate: 0.50 | Evals: 40.0%  
Iter 3 | Best: 123.050044 | Rate: 0.67 | Evals: 43.3%  
Iter 4 | Best: 122.581938 | Rate: 0.75 | Evals: 46.7%  
Iter 5 | Best: 121.850820 | Rate: 0.80 | Evals: 50.0%  
Iter 6 | Best: 121.742699 | Rate: 0.83 | Evals: 53.3%  
Iter 7 | Best: 121.577343 | Rate: 0.86 | Evals: 56.7%
```

6.6. Run optimization

```
Iter 8 | Best: 120.643326 | Rate: 0.88 | Evals: 60.0%
Iter 9 | Best: 120.013506 | Rate: 0.89 | Evals: 63.3%
Iter 10 | Best: 119.701481 | Rate: 0.90 | Evals: 66.7%
Iter 11 | Best: 119.506129 | Rate: 0.91 | Evals: 70.0%
Iter 12 | Best: 119.503886 | Rate: 0.92 | Evals: 73.3%
Iter 13 | Best: 119.503673 | Rate: 0.92 | Evals: 76.7%
Iter 14 | Best: 119.503673 | Curr: 119.503677 | Rate: 0.86 | Evals: 80.0%
Iter 15 | Best: 119.503672 | Rate: 0.87 | Evals: 83.3%
Iter 16 | Best: 119.503672 | Curr: 119.503777 | Rate: 0.81 | Evals: 86.7%
Iter 17 | Best: 119.503672 | Curr: 119.503814 | Rate: 0.76 | Evals: 90.0%
Iter 18 | Best: 119.503672 | Curr: 119.504008 | Rate: 0.72 | Evals: 93.3%
Iter 19 | Best: 119.503672 | Curr: 119.503679 | Rate: 0.68 | Evals: 96.7%
Iter 20 | Best: 119.503672 | Curr: 119.503672 | Rate: 0.65 | Evals: 100.0%
TensorBoard writer closed. View logs with: tensorboard --logdir=runs/spotoptim_20260701_201824
```

```
# End timing
spot_time = time.time() - start_time

print(f"\nSpotOptim Results:")
print(f"  Best weight: {result_spot.fun:.4f} lb")
print(f"  Function evaluations: {result_spot.nfev}")
print(f"  Time elapsed: {spot_time:.2f} seconds")
print(f"  Success: {result_spot.success}")
```

```
SpotOptim Results:
  Best weight: 119.5037 lb
  Function evaluations: 30
  Time elapsed: 6.83 seconds
  Success: True
```

```
optimizer_spot.print_best()
```

Best Solution Found:

```
-----
Sw: 0.0000
Wfw: 0.0000
A: 0.0000
L: 0.4999
q: 0.0000
l: 0.0000
Rtc: 1.0000
```

6. Optimizing the Aircraft Wing Weight Example

```
Nz: 0.0000
Wdg: 0.0000
Objective Value: 119.5037
Total Evaluations: 30
```

6.7. Result Table


```
pprint.pprint(optimizer_spot.get_results_table(show_importance=True))
```

```
('|  name |  type |  default |  lower |  upper |  tuned |  transform |
|  importance |  stars |\n'
'|-----|-----|-----|-----|-----|-----|-----|
|-----|-----|\n'
'|      Sw | float |      0.5 |      0 |      1 |      0 |      - |
'|      12.79 |      . |\n'
'|      Wfw | float |      0.5 |      0 |      1 |      0 |      - |
'|      14.41 |      . |\n'
'|      A | float |      0.5 |      0 |      1 |      0 |      - |
'|      14.97 |      . |\n'
'|      L | float |      0.5 |      0 |      1 | 0.4999 |      - |
'|      0.82 |      |\n'
'|      q | float |      0.5 |      0 |      1 |      0 |      - |
'|      8.29 |      |\n'
'|      l | float |      0.5 |      0 |      1 |      0 |      - |
'|      6.77 |      |\n'
'|      Rtc | float |      0.5 |      0 |      1 |      1 |      - |
'|      13.06 |      . |\n'
'|      Nz | float |      0.5 |      0 |      1 |      0 |      - |
'|      15.11 |      . |\n'
'|      Wdg | float |      0.5 |      0 |      1 |      0 |      - |
'|      13.78 |      . |\n'
'\n'
'Interpretation: ***: >99%, **: >75%, *: >50%, .: >10%')
```

6.8. Progress of the Optimization

```
optimizer_spot.plot_progress(log_y=False)
```

6.9. Contour Plots of Most Important Hyperparameters



6.9. Contour Plots of Most Important Hyperparameters

```
optimizer_spot.plot_important_hyperparameter_contour(max_imp=3)
```


Plotting surrogate contours for top 3 most important parameters:

Nz: importance = 15.11% (type: float)
A: importance = 14.97% (type: float)
Wfw: importance = 14.41% (type: float)

Generating 3 surrogate plots...


Plotting Nz vs A

6. Optimizing the Aircraft Wing Weight Example




Plotting N_z vs W_{fw}

6.9. Contour Plots of Most Important Hyperparameters



Plotting A vs Wfw

6. Optimizing the Aircraft Wing Weight Example



6.10. Method 2: Nelder-Mead Simplex

```
print("\n" + "=" * 60)
print("Running Nelder-Mead Simplex...")
print("=" * 60)

# Start timing
start_time = time.time()

# Run optimization
result_nm = minimize(
    wingwt_scipy,
    x0=x0,
    method='Nelder-Mead',
```


6.11. Method 3: BFGS (Quasi-Newton)

```
        bounds=bounds,
        options={'maxfev': max_evals, 'disp': True}
    )

    # End timing
    nm_time = time.time() - start_time

    print(f"\nNelder-Mead Results:")
    print(f"   Best weight: {result_nm.fun:.4f} lb")
    print(f"   Function evaluations: {result_nm.nfev}")
    print(f"   Time elapsed: {nm_time:.2f} seconds")
    print(f"   Success: {result_nm.success}")
```

```
=====
Running Nelder-Mead Simplex...
=====
```

```
Nelder-Mead Results:
  Best weight: 220.5449 lb
  Function evaluations: 30
  Time elapsed: 0.00 seconds
  Success: False
```

6.11. Method 3: BFGS (Quasi-Newton)

```
print("\n" + "=" * 60)
print("Running BFGS (Quasi-Newton)...")
print("=" * 60)

# Start timing
start_time = time.time()

# Run optimization
result_bfgs = minimize(
    wingwt_scipy,
    x0=x0,
    method='L-BFGS-B', # Bounded BFGS
    bounds=bounds,
    options={'maxfun': max_evals, 'disp': True})
```

6. Optimizing the Aircraft Wing Weight Example

```
)

# End timing
bfgs_time = time.time() - start_time

print(f"\nBFGS Results:")
print(f"  Best weight: {result_bfgs.fun:.4f} lb")
print(f"  Function evaluations: {result_bfgs.nfev}")
print(f"  Time elapsed: {bfgs_time:.2f} seconds")
print(f"  Success: {result_bfgs.success}")
```

```
=====
Running BFGS (Quasi-Newton)...
=====
```

```
BFGS Results:
  Best weight: 119.5037 lb
  Function evaluations: 60
  Time elapsed: 0.00 seconds
  Success: False
```

6.12. Comparison of Results

```
# Create comparison DataFrame
comparison = pd.DataFrame({
    'Method': ['Baseline', 'SpotOptim', 'Nelder-Mead', 'BFGS'],
    'Best Weight (lb)': [
        baseline_weight,
        result_spot.fun,
        result_nm.fun,
        result_bfgs.fun
    ],
    'Improvement (%)': [
        0.0,
        (baseline_weight - result_spot.fun) / baseline_weight * 100,
        (baseline_weight - result_nm.fun) / baseline_weight * 100,
        (baseline_weight - result_bfgs.fun) / baseline_weight * 100
    ],
    'Function Evals': [
        1,
```

```

        result_spot.nfev,
        result_nm.nfev,
        result_bfgs.nfev
    ],
    'Time (s)': [
        0.0,
        spot_time,
        nm_time,
        bfgs_time
    ],
    'Success': [
        True,
        result_spot.success,
        result_nm.success,
        result_bfgs.success
    ]
})

print("\n" + "=" * 80)
print("OPTIMIZATION COMPARISON")
print("=" * 80)
print(comparison.to_string(index=False))
print("=" * 80)

```

```

=====
OPTIMIZATION COMPARISON
=====

```

Method	Best Weight (lb)	Improvement (%)	Function Evals	Time (s)	Success
Baseline	233.908405	0.000000	1	0.000000	True
SpotOptim	119.503672	48.910057	30	6.826029	True
Nelder-Mead	220.544928	5.713124	30	0.000737	False
BFGS	119.503672	48.910057	60	0.001178	False

```

=====

```

6.13. Visualization: Convergence Plots

6.13.1. SpotOptim Convergence

```
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 5))
```


6. Optimizing the Aircraft Wing Weight Example

```
# Plot 1: Best value over iterations
y_history = optimizer_spot.y_
best_so_far = np.minimum.accumulate(y_history)

ax1.plot(range(len(y_history)), y_history, 'o', alpha=0.5, label='Evaluated points')
ax1.plot(range(len(best_so_far)), best_so_far, 'r-', linewidth=2, label='Best so far')
ax1.axhline(y=baseline_weight, color='g', linestyle='--', label='Baseline')
ax1.set_xlabel('Function Evaluation')
ax1.set_ylabel('Wing Weight (lb)')
ax1.set_title('SpotOptim Convergence')
ax1.legend()
ax1.grid(True, alpha=0.3)

# Plot 2: Improvement over baseline
improvement = (baseline_weight - best_so_far) / baseline_weight * 100
ax2.plot(range(len(improvement)), improvement, 'b-', linewidth=2)
ax2.set_xlabel('Function Evaluation')
ax2.set_ylabel('Improvement over Baseline (%)')
ax2.set_title('SpotOptim Improvement Progress')
ax2.grid(True, alpha=0.3)
ax2.axhline(y=0, color='k', linestyle='-', alpha=0.3)

plt.tight_layout()
plt.show()
```



6.14. Optimal Parameter Values

Let's examine the optimal parameter values found by each method:

6.14. Optimal Parameter Values

```
# Parameter names
param_names = ['Sw', 'Wfw', 'A', 'L', 'q', 'l', 'Rtc', 'Nz', 'Wdg']

# Transform from unit cube to natural scales
def decode_params(x):
    scales = [
        (150, 200),      # Sw
        (220, 300),      # Wfw
        (6, 10),          # A
        (-10, 10),        # L (degrees)
        (16, 45),         # q
        (0.5, 1),         # l
        (0.08, 0.18),     # Rtc
        (2.5, 6),         # Nz
        (1700, 2500)      # Wdg
    ]
    decoded = []
    for i, (low, high) in enumerate(scales):
        decoded.append(x[i] * (high - low) + low)
    return decoded

# Create comparison table
baseline_decoded = decode_params(baseline_coded)
spot_decoded = decode_params(result_spot.x)
nm_decoded = decode_params(result_nm.x)
bfgs_decoded = decode_params(result_bfgs.x)

param_comparison = pd.DataFrame({
    'Parameter': param_names,
    'Baseline': baseline_decoded,
    'SpotOptim': spot_decoded,
    'Nelder-Mead': nm_decoded,
    'BFGS': bfgs_decoded
})

print("\n" + "=" * 100)
print("OPTIMAL PARAMETER VALUES (Natural Scale)")
print("=" * 100)
print(param_comparison.to_string(index=False))
print("=" * 100)
```

```
=====
OPTIMAL PARAMETER VALUES (Natural Scale)
```

6. Optimizing the Aircraft Wing Weight Example

Parameter	Baseline	SpotOptim	Nelder-Mead	BFGS
Sw	174.000	150.00000	173.923947	150.00
Wfw	252.000	220.00000	254.534988	220.00
A	7.520	6.00000	7.484111	6.00
L	0.000	-0.00116	-0.126137	0.00
q	33.980	16.00000	34.156818	16.00
l	0.672	0.50000	0.674613	0.50
Rtc	0.120	0.18000	0.124872	0.18
Nz	3.795	2.50000	3.427058	2.50
Wdg	2004.000	1700.00000	2028.938270	1700.00

6.15. Analysis of Optimal Solutions


```
# Calculate percentage changes from baseline
changes_spot = [(spot_decoded[i] - baseline_decoded[i]) / baseline_decoded[i] * 100
                 for i in range(len(param_names))]
changes_nm = [(nm_decoded[i] - baseline_decoded[i]) / baseline_decoded[i] * 100
              for i in range(len(param_names))]
changes_bfgs = [(bfgs_decoded[i] - baseline_decoded[i]) / baseline_decoded[i] * 100
                for i in range(len(param_names))]

# Visualization
fig, axes = plt.subplots(1, 3, figsize=(18, 5))

for ax, changes, method in zip(axes,
                                [changes_spot, changes_nm, changes_bfgs],
                                ['SpotOptim', 'Nelder-Mead', 'BFGS']):
    colors = ['red' if c < 0 else 'green' for c in changes]
    ax.barh(param_names, changes, color=colors, alpha=0.6)
    ax.axvline(x=0, color='black', linestyle='-', linewidth=0.5)
    ax.set_xlabel('Change from Baseline (%)')
    ax.set_title(f'{method}: Parameter Changes')
    ax.grid(True, alpha=0.3, axis='x')

plt.tight_layout()
plt.show()
```

6.16. Key Insights from Optimal Solutions



6.16. Key Insights from Optimal Solutions

```
# Find parameters with largest changes for each method
def analyze_changes(decoded, baseline_decoded, method_name):
    changes = {param_names[i]: decoded[i] - baseline_decoded[i]
               for i in range(len(param_names))}
    sorted_changes = sorted(changes.items(), key=lambda x: abs(x[1]), reverse=True)

    print(f"\n{n{method_name}} - Top 5 Parameter Changes:")
    print("-" * 50)
    for param, change in sorted_changes[:5]:
        idx = param_names.index(param)
        pct = change / baseline_decoded[idx] * 100
        print(f" {param:5s}: {change:+8.2f} ({pct:+6.1f}%)")

analyze_changes(spot_decoded, baseline_decoded, "SpotOptim")
analyze_changes(nm_decoded, baseline_decoded, "Nelder-Mead")
analyze_changes(bfgs_decoded, baseline_decoded, "BFGS")
```

SpotOptim - Top 5 Parameter Changes:

```
-----
Wdg  :  -304.00 ( -15.2%)
Wfw  :  -32.00 ( -12.7%)
Sw   :  -24.00 ( -13.8%)
q    :  -17.98 ( -52.9%)
A    :   -1.52 ( -20.2%)
```

Nelder-Mead - Top 5 Parameter Changes:

```
-----
Wdg  :   +24.94 ( +1.2%)
Wfw  :    +2.53 ( +1.0%)
```

6. Optimizing the Aircraft Wing Weight Example

```
Nz   :   -0.37 (  -9.7%)
q     :   +0.18 ( +0.5%)
L     :   -0.13 ( -inf%)
```

BFGS - Top 5 Parameter Changes:

```
-----
Wdg   :  -304.00 ( -15.2%)
Wfw   :  -32.00 ( -12.7%)
Sw    :  -24.00 ( -13.8%)
q      :  -17.98 ( -52.9%)
A      :   -1.52 ( -20.2%)
```

6.17. Method Efficiency Comparison

```
# Calculate efficiency metrics
efficiency = pd.DataFrame({
    'Method': ['SpotOptim', 'Nelder-Mead', 'BFGS'],
    'Weight Reduction (lb)': [
        baseline_weight - result_spot.fun,
        baseline_weight - result_nm.fun,
        baseline_weight - result_bfgs.fun
    ],
    'Evals to Best': [
        np.argmin(optimizer_spot.y_) + 1,
        result_nm.nfev,
        result_bfgs.nfev
    ],
    'Time per Eval (ms)': [
        spot_time / result_spot.nfev * 1000,
        nm_time / result_nm.nfev * 1000,
        bfgs_time / result_bfgs.nfev * 1000
    ]
})

print("\n" + "=" * 80)
print("METHOD EFFICIENCY METRICS")
print("=" * 80)
print(efficiency.to_string(index=False))
print("=" * 80)
```


METHOD EFFICIENCY METRICS

Method	Weight Reduction (lb)	Evals to Best	Time per Eval (ms)
SpotOptim	114.404734	25	227.534302
Nelder-Mead	13.363478	30	0.024565
BFGS	114.404734	60	0.019638

6.18. Visualization: 2D Slices of Optimal Solutions

Let's visualize how the optimal solutions compare in the most important 2D subspaces:

```
# Create 2D slices showing optimal points
fig, axes = plt.subplots(2, 2, figsize=(14, 12))

# Important parameter pairs based on sensitivity analysis
pairs = [
    (7, 2), # Nz vs A (load factor vs aspect ratio)
    (0, 8), # Sw vs Wdg (wing area vs gross weight)
    (7, 0), # Nz vs Sw (load factor vs wing area)
    (2, 6)  # A vs Rtc (aspect ratio vs thickness ratio)
]

for ax, (i, j) in zip(axes.flat, pairs):
    # Create meshgrid for contour plot
    x_range = np.linspace(0, 1, 50)
    y_range = np.linspace(0, 1, 50)
    X, Y = np.meshgrid(x_range, y_range)

    # Evaluate function on grid (fixing other parameters at baseline)
    Z = np.zeros_like(X)
    for ii in range(X.shape[0]):
        for jj in range(X.shape[1]):
            point = baseline_coded.copy()
            point[i] = X[ii, jj]
            point[j] = Y[ii, jj]
            Z[ii, jj] = wingwt(point)[0]

    # Plot contours
    contour = ax.contourf(X, Y, Z, levels=20, cmap='viridis', alpha=0.6)
    ax.contour(X, Y, Z, levels=10, colors='black', alpha=0.3, linewidths=0.5)

# Plot optimal points
```

6. Optimizing the Aircraft Wing Weight Example


```
ax.plot(baseline_coded[i], baseline_coded[j], 'go', markersize=12,
        label='Baseline', markeredgewidth=1.5)
ax.plot(result_spot.x[i], result_spot.x[j], 'r*', markersize=15,
        label='SpotOptim', markeredgewidth=1)
ax.plot(result_nm.x[i], result_nm.x[j], 'bs', markersize=10,
        label='Nelder-Mead', markeredgewidth=1)
ax.plot(result_bfgs.x[i], result_bfgs.x[j], 'c^', markersize=10,
        label='BFGS', markeredgewidth=1)

ax.set_xlabel(param_names[i])
ax.set_ylabel(param_names[j])
ax.set_title(f'{param_names[j]} vs {param_names[i]}')
ax.legend(loc='best', fontsize=8)
ax.grid(True, alpha=0.3)

# Add colorbar
plt.colorbar(contour, ax=ax, label='Wing Weight (lb)')

plt.tight_layout()
plt.show()
```

6.19. Conclusion



6.19. Conclusion

This analysis demonstrates the application of different optimization methods to the Aircraft Wing Weight Example. Key takeaways:

1. **SpotOptim** provides efficient global optimization with good exploration of the design space
2. **Nelder-Mead** offers robust derivative-free optimization but may require more evaluations
3. **BFGS** converges quickly for smooth problems but can get trapped in local minima

For aircraft design problems with expensive simulations, Bayesian optimization (SpotOptim) offers the best balance of efficiency and solution quality, making it particularly suitable for real-world engineering applications.

6.20. Jupyter Notebook

i Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

7. Benchmarking SpotOptim with Sklearn Kriging (Matern Kernel) on 6D Rosenbrock and 10D Michalewicz Functions

i Note

These test functions were used during the Dagstuhl Seminar 25451 Bayesian Optimisation (Nov 02 – Nov 07, 2025), see [here](#).

This notebook demonstrates the use of `SpotOptim` with sklearn's Gaussian Process Regressor as a surrogate model.

7.1. SpotOptim with Sklearn Kriging in 6 Dimensions: Rosenbrock Function

This section demonstrates how to use the `SpotOptim` class with sklearn's Gaussian Process Regressor (using Matern kernel) as a surrogate on the 6-dimensional Rosenbrock function. We use a maximum of 100 function evaluations.

```
import warnings
warnings.filterwarnings("ignore")
import json
import numpy as np
from spotoptim import SpotOptim
from spotoptim.function import rosenbrock
```

7.1.1. Define the 6D Rosenbrock Function

7. Benchmarking SpotOptim with Sklearn Kriging (Matern Kernel) on 6D Rosenbrock and 10D Michalev

```
dim = 6
lower = np.full(dim, -2.0)
upper = np.full(dim, 2.0)
bounds = list(zip(lower, upper))
fun = rosenbrock
max_iter = 100
```

7.1.2. Set up SpotOptim Parameters

```
n_initial = dim
seed = 321
```

7.1.3. Sklearn Gaussian Process Regressor as Surrogate

```
from sklearn.gaussian_process import GaussianProcessRegressor
from sklearn.gaussian_process.kernels import Matern, ConstantKernel

# Use a Matern kernel instead of the standard RBF kernel
kernel = ConstantKernel(1.0, (1e-2, 1e12)) * Matern(
    length_scale=1.0,
    length_scale_bounds=(1e-4, 1e2),
    nu=2.5
)
surrogate = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=100)

# Create SpotOptim instance with sklearn surrogate
opt_rosen = SpotOptim(
    fun=fun,
    bounds=bounds,
    n_initial=n_initial,
    max_iter=max_iter,
    surrogate=surrogate,
    seed=seed,
    verbose=1
)

# Run optimization
result_rosen = opt_rosen.optimize()
```

TensorBoard logging disabled

7.1. SpotOptim with Sklearn Kriging in 6 Dimensions: Rosenbrock Function

Initial best: $f(x) = 783.489838$

Iter 1	Best: 783.489838	Curr: 2941.930042	Rate: 0.00	Evals: 7.0%
Iter 2	Best: 783.489838	Curr: 8421.000000	Rate: 0.00	Evals: 8.0%
Iter 3	Best: 783.489838	Curr: 1944.404445	Rate: 0.00	Evals: 9.0%
Iter 4	Best: 783.489838	Curr: 1654.277451	Rate: 0.00	Evals: 10.0%
Iter 5	Best: 117.869196	Rate: 0.20	Evals: 11.0%	
Iter 6	Best: 85.222268	Rate: 0.33	Evals: 12.0%	
Iter 7	Best: 61.531770	Rate: 0.43	Evals: 13.0%	
Iter 8	Best: 61.531770	Curr: 63.503227	Rate: 0.38	Evals: 14.0%
Iter 9	Best: 52.554093	Rate: 0.44	Evals: 15.0%	
Iter 10	Best: 46.042631	Rate: 0.50	Evals: 16.0%	
Iter 11	Best: 42.435012	Rate: 0.55	Evals: 17.0%	
Iter 12	Best: 42.435012	Curr: 42.493861	Rate: 0.50	Evals: 18.0%
Iter 13	Best: 40.690530	Rate: 0.54	Evals: 19.0%	
Iter 14	Best: 40.690530	Curr: 40.721119	Rate: 0.50	Evals: 20.0%
Iter 15	Best: 38.248374	Rate: 0.53	Evals: 21.0%	
Iter 16	Best: 38.248374	Curr: 83.545318	Rate: 0.50	Evals: 22.0%
Iter 17	Best: 25.633597	Rate: 0.53	Evals: 23.0%	
Iter 18	Best: 22.829288	Rate: 0.56	Evals: 24.0%	
Iter 19	Best: 21.173103	Rate: 0.58	Evals: 25.0%	
Iter 20	Best: 19.159174	Rate: 0.60	Evals: 26.0%	
Iter 21	Best: 16.453329	Rate: 0.62	Evals: 27.0%	
Iter 22	Best: 16.453329	Curr: 16.486129	Rate: 0.59	Evals: 28.0%
Iter 23	Best: 16.069485	Rate: 0.61	Evals: 29.0%	
Iter 24	Best: 13.416452	Rate: 0.62	Evals: 30.0%	
Iter 25	Best: 13.197358	Rate: 0.64	Evals: 31.0%	
Iter 26	Best: 10.990453	Rate: 0.65	Evals: 32.0%	
Iter 27	Best: 9.395567	Rate: 0.67	Evals: 33.0%	
Iter 28	Best: 6.807592	Rate: 0.68	Evals: 34.0%	
Iter 29	Best: 6.807592	Curr: 7.174969	Rate: 0.66	Evals: 35.0%
Iter 30	Best: 6.056502	Rate: 0.67	Evals: 36.0%	
Iter 31	Best: 5.923850	Rate: 0.68	Evals: 37.0%	
Iter 32	Best: 5.847398	Rate: 0.69	Evals: 38.0%	
Iter 33	Best: 5.847398	Curr: 5.881544	Rate: 0.67	Evals: 39.0%
Iter 34	Best: 5.657375	Rate: 0.68	Evals: 40.0%	
Iter 35	Best: 5.309238	Rate: 0.69	Evals: 41.0%	
Iter 36	Best: 5.268698	Rate: 0.69	Evals: 42.0%	
Iter 37	Best: 5.140954	Rate: 0.70	Evals: 43.0%	
Iter 38	Best: 4.958588	Rate: 0.71	Evals: 44.0%	
Iter 39	Best: 4.926015	Rate: 0.72	Evals: 45.0%	
Iter 40	Best: 4.926015	Curr: 4.943337	Rate: 0.70	Evals: 46.0%
Iter 41	Best: 4.885253	Rate: 0.71	Evals: 47.0%	
Iter 42	Best: 4.810985	Rate: 0.71	Evals: 48.0%	
Iter 43	Best: 4.607756	Rate: 0.72	Evals: 49.0%	
Iter 44	Best: 4.529472	Rate: 0.73	Evals: 50.0%	

7. Benchmarking SpotOptim with Sklearn Kriging (Matern Kernel) on 6D Rosenbrock and 10D Michalew

```

Iter 45 | Best: 4.439773 | Rate: 0.73 | Evals: 51.0%
Iter 46 | Best: 4.383023 | Rate: 0.74 | Evals: 52.0%
Iter 47 | Best: 4.383023 | Curr: 4.413104 | Rate: 0.72 | Evals: 53.0%
Iter 48 | Best: 4.352151 | Rate: 0.73 | Evals: 54.0%
Iter 49 | Best: 4.352151 | Curr: 4.361515 | Rate: 0.71 | Evals: 55.0%
Iter 50 | Best: 4.274633 | Rate: 0.72 | Evals: 56.0%
Iter 51 | Best: 4.066356 | Rate: 0.73 | Evals: 57.0%
Iter 52 | Best: 4.066356 | Curr: 4.070280 | Rate: 0.71 | Evals: 58.0%
Iter 53 | Best: 4.042185 | Rate: 0.72 | Evals: 59.0%
Iter 54 | Best: 4.042185 | Curr: 4.066555 | Rate: 0.70 | Evals: 60.0%
Iter 55 | Best: 4.035292 | Rate: 0.71 | Evals: 61.0%
Iter 56 | Best: 4.020523 | Rate: 0.71 | Evals: 62.0%
Iter 57 | Best: 4.020523 | Curr: 4.022247 | Rate: 0.70 | Evals: 63.0%
Iter 58 | Best: 4.020523 | Curr: 4.032639 | Rate: 0.69 | Evals: 64.0%
Iter 59 | Best: 3.998297 | Rate: 0.69 | Evals: 65.0%
Iter 60 | Best: 3.955409 | Rate: 0.70 | Evals: 66.0%
Iter 61 | Best: 3.955409 | Curr: 3.985831 | Rate: 0.69 | Evals: 67.0%
Optimizer candidate 1/3 was duplicate/invalid.
Iter 62 | Best: 3.936906 | Rate: 0.69 | Evals: 68.0%
Iter 63 | Best: 3.917257 | Rate: 0.70 | Evals: 69.0%
Iter 64 | Best: 3.832945 | Rate: 0.70 | Evals: 70.0%
Iter 65 | Best: 3.737995 | Rate: 0.71 | Evals: 71.0%
Iter 66 | Best: 3.732943 | Rate: 0.71 | Evals: 72.0%
Iter 67 | Best: 3.698151 | Rate: 0.72 | Evals: 73.0%
Iter 68 | Best: 3.542894 | Rate: 0.72 | Evals: 74.0%
Iter 69 | Best: 3.466792 | Rate: 0.72 | Evals: 75.0%
Iter 70 | Best: 3.389356 | Rate: 0.73 | Evals: 76.0%
Iter 71 | Best: 3.389356 | Curr: 3.398457 | Rate: 0.72 | Evals: 77.0%
Iter 72 | Best: 3.389356 | Curr: 3.485776 | Rate: 0.71 | Evals: 78.0%
Iter 73 | Best: 3.193561 | Rate: 0.71 | Evals: 79.0%
Iter 74 | Best: 3.193561 | Curr: 3.197722 | Rate: 0.70 | Evals: 80.0%
Iter 75 | Best: 3.193561 | Curr: 3.232813 | Rate: 0.69 | Evals: 81.0%
Iter 76 | Best: 3.193561 | Curr: 3.193589 | Rate: 0.68 | Evals: 82.0%
Iter 77 | Best: 3.193561 | Curr: 3.207034 | Rate: 0.68 | Evals: 83.0%
Iter 78 | Best: 3.193561 | Curr: 3.205703 | Rate: 0.67 | Evals: 84.0%
Iter 79 | Best: 3.120052 | Rate: 0.67 | Evals: 85.0%
Iter 80 | Best: 3.003385 | Rate: 0.68 | Evals: 86.0%
Iter 81 | Best: 2.939601 | Rate: 0.68 | Evals: 87.0%
Iter 82 | Best: 2.939601 | Curr: 2.980987 | Rate: 0.67 | Evals: 88.0%
Iter 83 | Best: 2.939283 | Rate: 0.67 | Evals: 89.0%
Iter 84 | Best: 2.901138 | Rate: 0.68 | Evals: 90.0%
Iter 85 | Best: 2.846729 | Rate: 0.68 | Evals: 91.0%
Iter 86 | Best: 2.836273 | Rate: 0.69 | Evals: 92.0%
Iter 87 | Best: 2.810150 | Rate: 0.69 | Evals: 93.0%
Iter 88 | Best: 2.810150 | Curr: 2.843165 | Rate: 0.68 | Evals: 94.0%

```


7.1. SpotOptim with Sklearn Kriging in 6 Dimensions: Rosenbrock Function

```
Iter 89 | Best: 2.757874 | Rate: 0.69 | Evals: 95.0%
Iter 90 | Best: 2.757874 | Curr: 2.759910 | Rate: 0.68 | Evals: 96.0%
Iter 91 | Best: 2.757874 | Curr: 2.766545 | Rate: 0.67 | Evals: 97.0%
Iter 92 | Best: 2.757874 | Curr: 2.760556 | Rate: 0.66 | Evals: 98.0%
Iter 93 | Best: 2.757874 | Curr: 2.800058 | Rate: 0.66 | Evals: 99.0%
Iter 94 | Best: 2.562003 | Rate: 0.66 | Evals: 100.0%
```

```
print(f"[6D] Sklearn Kriging: min y = {result_rosen.fun:.4f} at x = {result_rosen.x}")
print(f"Number of function evaluations: {result_rosen.nfev}")
print(f"Number of iterations: {result_rosen.nit}")
```


```
[6D] Sklearn Kriging: min y = 2.5620 at x = [ 7.37596458e-01  5.44350495e-
01  2.99589530e-01  9.62930430e-02
 1.58104768e-02 -7.82045238e-05]
Number of function evaluations: 100
Number of iterations: 94
```

7.1.4. Visualize Optimization Progress

```
import matplotlib.pyplot as plt

# Plot the optimization progress
plt.figure(figsize=(10, 6))
plt.semilogy(np.minimum.accumulate(opt_rosen.y_), 'b-', linewidth=2)
plt.xlabel('Function Evaluations', fontsize=12)
plt.ylabel('Best Objective Value (log scale)', fontsize=12)
plt.title('6D Rosenbrock: Sklearn Kriging Progress', fontsize=14)
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```

7. Benchmarking SpotOptim with Sklearn Kriging (Matern Kernel) on 6D Rosenbrock and 10D Michalev



7.1.5. Evaluation of Multiple Repeats

To perform 30 repeats and collect statistics:

```
# Perform 30 independent runs
n_repeats = 30
results = []

print(f"Running {n_repeats} independent optimizations...")
for i in range(n_repeats):
    kernel_i = ConstantKernel(1.0, (1e-2, 1e12)) * Matern(
        length_scale=1.0,
        length_scale_bounds=(1e-4, 1e2),
        nu=2.5
    )
    surrogate_i = GaussianProcessRegressor(kernel=kernel_i, n_restarts_optimizer=100)

    opt_i = SpotOptim(
        fun=fun,
        bounds=bounds,
        n_initial=n_initial,
        max_iter=max_iter,
        surrogate=surrogate_i,
        seed=seed + i, # Different seed for each run
```

7.2. SpotOptim with Sklearn Kriging in 10 Dimensions: Michalewicz Function

```
        verbose=0
    )

    result_i = opt_i.optimize()
    results.append(result_i.fun)

    if (i + 1) % 10 == 0:
        print(f" Completed {i + 1}/{n_repeats} runs")

# Compute statistics
mean_result = np.mean(results)
std_result = np.std(results)
min_result = np.min(results)
max_result = np.max(results)

print(f"\nResults over {n_repeats} runs:")
print(f" Mean of best values: {mean_result:.6f}")
print(f" Std of best values: {std_result:.6f}")
print(f" Min of best values: {min_result:.6f}")
print(f" Max of best values: {max_result:.6f}")
```

7.2. SpotOptim with Sklearn Kriging in 10 Dimensions: Michalewicz Function

This section demonstrates how to use the `SpotOptim` class with sklearn's Gaussian Process Regressor (using Matern kernel) as a surrogate on the 10-dimensional Michalewicz function. We use a maximum of 300 function evaluations.

7.2.1. Define the 10D Michalewicz Function

```
from spotoptim.function import michalewicz

dim = 10
lower = np.full(dim, 0.0)
upper = np.full(dim, np.pi)
bounds = list(zip(lower, upper))
fun = michalewicz
max_iter = 300
```

7.2.2. Set up SpotOptim Parameters

```
n_initial = dim
seed = 321
```

7.2.3. Sklearn Gaussian Process Regressor as Surrogate

```
from sklearn.gaussian_process import GaussianProcessRegressor
from sklearn.gaussian_process.kernels import Matern, ConstantKernel

# Use a Matern kernel instead of the standard RBF kernel
kernel = ConstantKernel(1.0, (1e-2, 1e12)) * Matern(
    length_scale=1.0,
    length_scale_bounds=(1e-4, 1e2),
    nu=2.5
)
surrogate = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=100)

# Create SpotOptim instance with sklearn surrogate
opt_micha = SpotOptim(
    fun=fun,
    bounds=bounds,
    n_initial=n_initial,
    max_iter=max_iter,
    surrogate=surrogate,
    seed=seed,
    verbose=1
)

# Run optimization
result_micha = opt_micha.optimize()
```

```
TensorBoard logging disabled
Initial best: f(x) = -3.138441
Iter 1 | Best: -3.251830 | Rate: 1.00 | Evals: 3.7%
Iter 2 | Best: -3.760942 | Rate: 1.00 | Evals: 4.0%
Iter 3 | Best: -3.760942 | Curr: -3.323915 | Rate: 0.67 | Evals: 4.3%
Iter 4 | Best: -4.055466 | Rate: 0.75 | Evals: 4.7%
Iter 5 | Best: -4.055466 | Curr: -2.924504 | Rate: 0.60 | Evals: 5.0%
Iter 6 | Best: -4.055466 | Curr: -3.905036 | Rate: 0.50 | Evals: 5.3%
Iter 7 | Best: -4.055466 | Curr: -4.029844 | Rate: 0.43 | Evals: 5.7%
```

7.2. SpotOptim with Sklearn Kriging in 10 Dimensions: Michalewicz Function

```
Iter 8 | Best: -4.055466 | Curr: -2.764374 | Rate: 0.38 | Evals: 6.0%
Iter 9 | Best: -4.055466 | Curr: -3.750592 | Rate: 0.33 | Evals: 6.3%
Iter 10 | Best: -4.094143 | Rate: 0.40 | Evals: 6.7%
Iter 11 | Best: -4.198315 | Rate: 0.45 | Evals: 7.0%
Iter 12 | Best: -4.198315 | Curr: -2.054963 | Rate: 0.42 | Evals: 7.3%
Iter 13 | Best: -4.198315 | Curr: -4.132779 | Rate: 0.38 | Evals: 7.7%
Iter 14 | Best: -4.198315 | Curr: -4.174200 | Rate: 0.36 | Evals: 8.0%
Iter 15 | Best: -4.221097 | Rate: 0.40 | Evals: 8.3%
Iter 16 | Best: -4.251756 | Rate: 0.44 | Evals: 8.7%
Iter 17 | Best: -4.289774 | Rate: 0.47 | Evals: 9.0%
Iter 18 | Best: -4.289774 | Curr: -3.972457 | Rate: 0.44 | Evals: 9.3%
Iter 19 | Best: -4.289774 | Curr: -2.281106 | Rate: 0.42 | Evals: 9.7%
Iter 20 | Best: -5.032943 | Rate: 0.45 | Evals: 10.0%
Iter 21 | Best: -5.032943 | Curr: -3.794046 | Rate: 0.43 | Evals: 10.3%
Iter 22 | Best: -5.032943 | Curr: -2.281897 | Rate: 0.41 | Evals: 10.7%
Iter 23 | Best: -5.032943 | Curr: -4.584760 | Rate: 0.39 | Evals: 11.0%
Optimizer candidate 1/3 was duplicate/invalid.
Iter 24 | Best: -5.032943 | Curr: -1.947168 | Rate: 0.38 | Evals: 11.3%
Iter 25 | Best: -5.032943 | Curr: -4.847314 | Rate: 0.36 | Evals: 11.7%
Optimizer candidate 1/3 was duplicate/invalid.
Iter 26 | Best: -5.032943 | Curr: -0.261467 | Rate: 0.35 | Evals: 12.0%
Iter 27 | Best: -5.032943 | Curr: -2.281779 | Rate: 0.33 | Evals: 12.3%
Iter 28 | Best: -5.032943 | Curr: -2.282029 | Rate: 0.32 | Evals: 12.7%
Iter 29 | Best: -5.045479 | Rate: 0.34 | Evals: 13.0%
Iter 30 | Best: -5.062816 | Rate: 0.37 | Evals: 13.3%
Iter 31 | Best: -5.062816 | Curr: -2.093269 | Rate: 0.35 | Evals: 13.7%
Iter 32 | Best: -5.062816 | Curr: -0.628600 | Rate: 0.34 | Evals: 14.0%
Iter 33 | Best: -5.087797 | Rate: 0.36 | Evals: 14.3%
Iter 34 | Best: -5.087797 | Curr: -2.247537 | Rate: 0.35 | Evals: 14.7%
Iter 35 | Best: -5.087797 | Curr: -5.081054 | Rate: 0.34 | Evals: 15.0%
Iter 36 | Best: -5.155832 | Rate: 0.36 | Evals: 15.3%
Iter 37 | Best: -5.155832 | Curr: -5.148322 | Rate: 0.35 | Evals: 15.7%
Iter 38 | Best: -5.155832 | Curr: -2.450559 | Rate: 0.34 | Evals: 16.0%
Iter 39 | Best: -5.155832 | Curr: -4.586218 | Rate: 0.33 | Evals: 16.3%
Optimizer candidate 1/3 was duplicate/invalid.
Iter 40 | Best: -5.155832 | Curr: -0.184026 | Rate: 0.33 | Evals: 16.7%
Iter 41 | Best: -5.156226 | Rate: 0.34 | Evals: 17.0%
Iter 42 | Best: -5.170901 | Rate: 0.36 | Evals: 17.3%
Iter 43 | Best: -5.171604 | Rate: 0.37 | Evals: 17.7%
Iter 44 | Best: -5.171604 | Curr: -5.170755 | Rate: 0.36 | Evals: 18.0%
Iter 45 | Best: -5.172016 | Rate: 0.38 | Evals: 18.3%
Iter 46 | Best: -5.173850 | Rate: 0.39 | Evals: 18.7%
Iter 47 | Best: -5.174029 | Rate: 0.40 | Evals: 19.0%
Iter 48 | Best: -5.174029 | Curr: -2.605367 | Rate: 0.40 | Evals: 19.3%
Iter 49 | Best: -5.174172 | Rate: 0.41 | Evals: 19.7%
```

7. Benchmarking SpotOptim with Sklearn Kriging (Matern Kernel) on 6D Rosenbrock and 10D Michalew

Optimizer candidate 1/3 was duplicate/invalid.

Iter 50	Best: -5.174172	Curr: -1.490619	Rate: 0.40	Evals: 20.0%
Iter 51	Best: -5.174172	Curr: -2.282019	Rate: 0.39	Evals: 20.3%
Iter 52	Best: -5.176808	Rate: 0.40	Evals: 20.7%	
Iter 53	Best: -5.176808	Curr: -2.288770	Rate: 0.40	Evals: 21.0%
Iter 54	Best: -5.176808	Curr: -0.795063	Rate: 0.39	Evals: 21.3%
Iter 55	Best: -5.176808	Curr: -2.627661	Rate: 0.38	Evals: 21.7%
Iter 56	Best: -5.184239	Rate: 0.39	Evals: 22.0%	
Iter 57	Best: -5.184650	Rate: 0.40	Evals: 22.3%	
Iter 58	Best: -5.184671	Rate: 0.41	Evals: 22.7%	
Iter 59	Best: -5.184671	Curr: -0.476006	Rate: 0.41	Evals: 23.0%
Iter 60	Best: -5.189888	Rate: 0.42	Evals: 23.3%	
Iter 61	Best: -5.216026	Rate: 0.43	Evals: 23.7%	
Iter 62	Best: -5.218685	Rate: 0.44	Evals: 24.0%	
Iter 63	Best: -5.228056	Rate: 0.44	Evals: 24.3%	
Iter 64	Best: -5.375679	Rate: 0.45	Evals: 24.7%	
Iter 65	Best: -5.470443	Rate: 0.46	Evals: 25.0%	
Iter 66	Best: -5.698678	Rate: 0.47	Evals: 25.3%	
Iter 67	Best: -5.698678	Curr: -0.629343	Rate: 0.46	Evals: 25.7%
Iter 68	Best: -5.698678	Curr: -2.295021	Rate: 0.46	Evals: 26.0%
Iter 69	Best: -5.698678	Curr: -5.531263	Rate: 0.45	Evals: 26.3%
Iter 70	Best: -5.705483	Rate: 0.46	Evals: 26.7%	
Iter 71	Best: -5.706956	Rate: 0.46	Evals: 27.0%	
Iter 72	Best: -5.773735	Rate: 0.47	Evals: 27.3%	
Iter 73	Best: -5.773735	Curr: -2.587389	Rate: 0.47	Evals: 27.7%
Iter 74	Best: -5.792007	Rate: 0.47	Evals: 28.0%	
Iter 75	Best: -5.792007	Curr: -2.628561	Rate: 0.47	Evals: 28.3%
Iter 76	Best: -5.798373	Rate: 0.47	Evals: 28.7%	
Iter 77	Best: -5.799245	Rate: 0.48	Evals: 29.0%	
Iter 78	Best: -5.799245	Curr: -5.791845	Rate: 0.47	Evals: 29.3%
Iter 79	Best: -5.821975	Rate: 0.48	Evals: 29.7%	
Iter 80	Best: -5.821975	Curr: -2.628587	Rate: 0.47	Evals: 30.0%
Iter 81	Best: -5.821975	Curr: -5.818729	Rate: 0.47	Evals: 30.3%
Iter 82	Best: -5.826330	Rate: 0.48	Evals: 30.7%	
Iter 83	Best: -5.826330	Curr: -2.643862	Rate: 0.47	Evals: 31.0%
Iter 84	Best: -5.826330	Curr: -3.011341	Rate: 0.46	Evals: 31.3%
Iter 85	Best: -5.826330	Curr: -2.661583	Rate: 0.46	Evals: 31.7%
Iter 86	Best: -5.826330	Curr: -2.663722	Rate: 0.45	Evals: 32.0%
Iter 87	Best: -5.831187	Rate: 0.46	Evals: 32.3%	
Iter 88	Best: -5.831187	Curr: -2.945846	Rate: 0.45	Evals: 32.7%
Iter 89	Best: -5.831187	Curr: -0.802679	Rate: 0.45	Evals: 33.0%
Iter 90	Best: -5.837935	Rate: 0.46	Evals: 33.3%	
Iter 91	Best: -5.839815	Rate: 0.46	Evals: 33.7%	
Iter 92	Best: -5.839815	Curr: -5.833454	Rate: 0.46	Evals: 34.0%
Iter 93	Best: -5.839815	Curr: -5.839101	Rate: 0.45	Evals: 34.3%

7.2. SpotOptim with Sklearn Kriging in 10 Dimensions: Michalewicz Function

```
Iter 94 | Best: -5.861861 | Rate: 0.46 | Evals: 34.7%
Iter 95 | Best: -5.861861 | Curr: -5.860331 | Rate: 0.45 | Evals: 35.0%
Iter 96 | Best: -5.861861 | Curr: -0.803144 | Rate: 0.45 | Evals: 35.3%
Iter 97 | Best: -5.861861 | Curr: -5.844693 | Rate: 0.44 | Evals: 35.7%
Iter 98 | Best: -5.870135 | Rate: 0.45 | Evals: 36.0%
Iter 99 | Best: -5.870135 | Curr: -5.869759 | Rate: 0.44 | Evals: 36.3%
Iter 100 | Best: -5.870135 | Curr: -3.098980 | Rate: 0.44 | Evals: 36.7%
Iter 101 | Best: -5.870135 | Curr: -3.109893 | Rate: 0.43 | Evals: 37.0%
Iter 102 | Best: -5.870135 | Curr: -5.868349 | Rate: 0.42 | Evals: 37.3%
Iter 103 | Best: -5.870135 | Curr: -0.713638 | Rate: 0.42 | Evals: 37.7%
Iter 104 | Best: -5.870135 | Curr: -3.102646 | Rate: 0.41 | Evals: 38.0%
Iter 105 | Best: -5.870135 | Curr: -5.868130 | Rate: 0.41 | Evals: 38.3%
Iter 106 | Best: -5.870135 | Curr: -0.637541 | Rate: 0.41 | Evals: 38.7%
Iter 107 | Best: -5.870568 | Rate: 0.42 | Evals: 39.0%
Iter 108 | Best: -5.870568 | Curr: -3.126847 | Rate: 0.42 | Evals: 39.3%
Iter 109 | Best: -5.870568 | Curr: -5.870526 | Rate: 0.42 | Evals: 39.7%
Iter 110 | Best: -5.870568 | Curr: -5.870041 | Rate: 0.41 | Evals: 40.0%
Optimizer candidate 1/3 was duplicate/invalid.
Iter 111 | Best: -5.870568 | Curr: -1.994681 | Rate: 0.40 | Evals: 40.3%
Iter 112 | Best: -5.871604 | Rate: 0.41 | Evals: 40.7%
Iter 113 | Best: -5.872649 | Rate: 0.42 | Evals: 41.0%
Iter 114 | Best: -5.873951 | Rate: 0.43 | Evals: 41.3%
Iter 115 | Best: -5.874765 | Rate: 0.43 | Evals: 41.7%
Iter 116 | Best: -5.874765 | Curr: -5.874627 | Rate: 0.42 | Evals: 42.0%
Iter 117 | Best: -5.874844 | Rate: 0.42 | Evals: 42.3%
Iter 118 | Best: -5.875052 | Rate: 0.43 | Evals: 42.7%
Iter 119 | Best: -5.875052 | Curr: -5.872627 | Rate: 0.43 | Evals: 43.0%
Iter 120 | Best: -5.875260 | Rate: 0.43 | Evals: 43.3%
Iter 121 | Best: -5.876218 | Rate: 0.44 | Evals: 43.7%
Iter 122 | Best: -5.876218 | Curr: -3.134808 | Rate: 0.44 | Evals: 44.0%
Iter 123 | Best: -5.876218 | Curr: -1.493284 | Rate: 0.44 | Evals: 44.3%
Iter 124 | Best: -5.876919 | Rate: 0.45 | Evals: 44.7%
Iter 125 | Best: -5.876919 | Curr: -5.876869 | Rate: 0.45 | Evals: 45.0%
Iter 126 | Best: -5.876919 | Curr: -5.876741 | Rate: 0.45 | Evals: 45.3%
Iter 127 | Best: -5.876919 | Curr: -5.876361 | Rate: 0.45 | Evals: 45.7%
Iter 128 | Best: -5.876919 | Curr: -5.876290 | Rate: 0.45 | Evals: 46.0%
Iter 129 | Best: -5.876919 | Curr: -1.493281 | Rate: 0.44 | Evals: 46.3%
Iter 130 | Best: -5.876919 | Curr: -2.663940 | Rate: 0.43 | Evals: 46.7%
Iter 131 | Best: -5.877020 | Rate: 0.44 | Evals: 47.0%
Iter 132 | Best: -5.877433 | Rate: 0.45 | Evals: 47.3%
Iter 133 | Best: -5.877720 | Rate: 0.45 | Evals: 47.7%
Iter 134 | Best: -5.877772 | Rate: 0.46 | Evals: 48.0%
Iter 135 | Best: -5.877772 | Curr: -0.835812 | Rate: 0.46 | Evals: 48.3%
Iter 136 | Best: -5.878110 | Rate: 0.46 | Evals: 48.7%
Iter 137 | Best: -5.878110 | Curr: -2.891125 | Rate: 0.46 | Evals: 49.0%
```

7. Benchmarking SpotOptim with Sklearn Kriging (Matern Kernel) on 6D Rosenbrock and 10D Michalew

Iter 138	Best: -5.879056	Rate: 0.47	Evals: 49.3%
Iter 139	Best: -5.881098	Rate: 0.48	Evals: 49.7%
Iter 140	Best: -5.881098	Curr: -3.132602	Rate: 0.48 Evals: 50.0%
Iter 141	Best: -5.881214	Rate: 0.48	Evals: 50.3%
Iter 142	Best: -5.881214	Curr: -0.803143	Rate: 0.47 Evals: 50.7%
Iter 143	Best: -5.881214	Curr: -3.097363	Rate: 0.46 Evals: 51.0%
Iter 144	Best: -5.881420	Rate: 0.47	Evals: 51.3%
Iter 145	Best: -5.881420	Curr: -5.881370	Rate: 0.46 Evals: 51.7%
Iter 146	Best: -5.882108	Rate: 0.46	Evals: 52.0%
Iter 147	Best: -5.882108	Curr: -5.851519	Rate: 0.45 Evals: 52.3%
Iter 148	Best: -5.882108	Curr: -5.881889	Rate: 0.45 Evals: 52.7%
Iter 149	Best: -5.882688	Rate: 0.45	Evals: 53.0%
Iter 150	Best: -5.882695	Rate: 0.46	Evals: 53.3%
Iter 151	Best: -5.882695	Curr: -5.882631	Rate: 0.46 Evals: 53.7%
Iter 152	Best: -5.882695	Curr: -5.882514	Rate: 0.45 Evals: 54.0%
Iter 153	Best: -5.882695	Curr: -3.317742	Rate: 0.45 Evals: 54.3%
Iter 154	Best: -5.882695	Curr: -3.113612	Rate: 0.45 Evals: 54.7%
Iter 155	Best: -5.882708	Rate: 0.46	Evals: 55.0%
Iter 156	Best: -5.882708	Curr: -5.882615	Rate: 0.45 Evals: 55.3%
Iter 157	Best: -5.882708	Curr: -5.882643	Rate: 0.44 Evals: 55.7%
Iter 158	Best: -5.882708	Curr: -5.882686	Rate: 0.43 Evals: 56.0%
Iter 159	Best: -5.882801	Rate: 0.44	Evals: 56.3%
Iter 160	Best: -5.882805	Rate: 0.44	Evals: 56.7%
Iter 161	Best: -5.882826	Rate: 0.44	Evals: 57.0%
Iter 162	Best: -5.882855	Rate: 0.44	Evals: 57.3%
Iter 163	Best: -5.882855	Curr: -5.882850	Rate: 0.43 Evals: 57.7%
Iter 164	Best: -5.882864	Rate: 0.43	Evals: 58.0%
Iter 165	Best: -5.882864	Curr: -3.155551	Rate: 0.42 Evals: 58.3%
Iter 166	Best: -5.882873	Rate: 0.42	Evals: 58.7%
Iter 167	Best: -5.882887	Rate: 0.43	Evals: 59.0%
Iter 168	Best: -5.882887	Curr: -5.882882	Rate: 0.43 Evals: 59.3%
Iter 169	Best: -5.882887	Curr: -3.152164	Rate: 0.43 Evals: 59.7%
Iter 170	Best: -5.882887	Curr: -5.882870	Rate: 0.42 Evals: 60.0%
Iter 171	Best: -5.882887	Curr: -3.130334	Rate: 0.41 Evals: 60.3%
Iter 172	Best: -5.882887	Curr: -1.463942	Rate: 0.40 Evals: 60.7%
Iter 173	Best: -5.882887	Curr: -5.882886	Rate: 0.40 Evals: 61.0%
Iter 174	Best: -5.882887	Curr: -5.882887	Rate: 0.39 Evals: 61.3%
Iter 175	Best: -5.882902	Rate: 0.40	Evals: 61.7%
Iter 176	Best: -5.882902	Curr: -3.111237	Rate: 0.39 Evals: 62.0%
Iter 177	Best: -5.882924	Rate: 0.39	Evals: 62.3%
Iter 178	Best: -5.883006	Rate: 0.40	Evals: 62.7%
Iter 179	Best: -5.883053	Rate: 0.40	Evals: 63.0%
Iter 180	Best: -5.883073	Rate: 0.41	Evals: 63.3%
Iter 181	Best: -5.883073	Curr: -3.193552	Rate: 0.41 Evals: 63.7%
Iter 182	Best: -5.883073	Curr: -1.491380	Rate: 0.40 Evals: 64.0%

7.2. SpotOptim with Sklearn Kriging in 10 Dimensions: Michalewicz Function

Iter 183	Best: -5.883096	Rate: 0.41	Evals: 64.3%	
Iter 184	Best: -5.883096	Curr: -5.883091	Rate: 0.41	Evals: 64.7%
Iter 185	Best: -5.883122	Rate: 0.42	Evals: 65.0%	
Iter 186	Best: -5.883123	Rate: 0.43	Evals: 65.3%	
Iter 187	Best: -5.883123	Curr: -1.687059	Rate: 0.42	Evals: 65.7%
Iter 188	Best: -5.883123	Curr: -3.172346	Rate: 0.42	Evals: 66.0%
Iter 189	Best: -5.883126	Rate: 0.43	Evals: 66.3%	
Iter 190	Best: -5.883157	Rate: 0.43	Evals: 66.7%	
Iter 191	Best: -5.883187	Rate: 0.43	Evals: 67.0%	
Iter 192	Best: -5.883187	Curr: -5.883182	Rate: 0.43	Evals: 67.3%
Iter 193	Best: -5.883196	Rate: 0.44	Evals: 67.7%	
Iter 194	Best: -5.883196	Curr: -3.185474	Rate: 0.43	Evals: 68.0%
Iter 195	Best: -5.883196	Curr: -5.883193	Rate: 0.43	Evals: 68.3%
Iter 196	Best: -5.883196	Curr: -5.883185	Rate: 0.43	Evals: 68.7%
Iter 197	Best: -5.883196	Curr: -3.177133	Rate: 0.43	Evals: 69.0%
Iter 198	Best: -5.883196	Curr: -5.883182	Rate: 0.42	Evals: 69.3%
Iter 199	Best: -5.883196	Curr: -5.883183	Rate: 0.42	Evals: 69.7%
Iter 200	Best: -5.883196	Curr: -3.299827	Rate: 0.42	Evals: 70.0%
Iter 201	Best: -5.883196	Curr: -3.308466	Rate: 0.42	Evals: 70.3%
Iter 202	Best: -5.883199	Rate: 0.43	Evals: 70.7%	
Iter 203	Best: -5.883201	Rate: 0.44	Evals: 71.0%	
Iter 204	Best: -5.883201	Curr: -5.883195	Rate: 0.44	Evals: 71.3%
Iter 205	Best: -5.883202	Rate: 0.45	Evals: 71.7%	
Iter 206	Best: -5.883202	Curr: -3.218223	Rate: 0.45	Evals: 72.0%
Iter 207	Best: -5.883202	Curr: -3.830932	Rate: 0.44	Evals: 72.3%
Iter 208	Best: -5.883202	Curr: -3.307359	Rate: 0.44	Evals: 72.7%
Iter 209	Best: -5.883202	Curr: -5.883200	Rate: 0.44	Evals: 73.0%
Iter 210	Best: -5.883205	Rate: 0.45	Evals: 73.3%	
Iter 211	Best: -5.883205	Curr: -5.883198	Rate: 0.45	Evals: 73.7%
Iter 212	Best: -5.883205	Curr: -3.332788	Rate: 0.44	Evals: 74.0%
Iter 213	Best: -5.883209	Rate: 0.44	Evals: 74.3%	
Iter 214	Best: -5.883215	Rate: 0.44	Evals: 74.7%	
Iter 215	Best: -5.883215	Curr: -5.883198	Rate: 0.43	Evals: 75.0%
Iter 216	Best: -5.883215	Curr: -3.279944	Rate: 0.43	Evals: 75.3%
Iter 217	Best: -5.883215	Rate: 0.43	Evals: 75.7%	
Iter 218	Best: -5.883215	Curr: -5.883211	Rate: 0.42	Evals: 76.0%
Iter 219	Best: -5.883215	Curr: -5.883210	Rate: 0.42	Evals: 76.3%
Iter 220	Best: -5.883215	Curr: -5.883208	Rate: 0.41	Evals: 76.7%
Iter 221	Best: -5.883215	Curr: -3.338531	Rate: 0.40	Evals: 77.0%
Iter 222	Best: -5.883215	Curr: -5.883211	Rate: 0.40	Evals: 77.3%
Iter 223	Best: -5.883215	Rate: 0.41	Evals: 77.7%	
Iter 224	Best: -5.883215	Curr: -5.883213	Rate: 0.40	Evals: 78.0%
Iter 225	Best: -5.883215	Curr: -5.883202	Rate: 0.40	Evals: 78.3%
Iter 226	Best: -5.883215	Curr: -5.883213	Rate: 0.40	Evals: 78.7%
Iter 227	Best: -5.883215	Curr: -3.392006	Rate: 0.40	Evals: 79.0%

7. Benchmarking SpotOptim with Sklearn Kriging (Matern Kernel) on 6D Rosenbrock and 10D Michalew

```

Iter 228 | Best: -5.883215 | Curr: -5.883208 | Rate: 0.40 | Evals: 79.3%
Iter 229 | Best: -5.883215 | Curr: -3.227816 | Rate: 0.40 | Evals: 79.7%
Iter 230 | Best: -5.883215 | Rate: 0.41 | Evals: 80.0%
Iter 231 | Best: -5.883215 | Curr: -5.875494 | Rate: 0.40 | Evals: 80.3%
Iter 232 | Best: -5.883215 | Curr: -5.883203 | Rate: 0.39 | Evals: 80.7%
Optimizer candidate 1/3 was duplicate/invalid.
Iter 233 | Best: -5.883215 | Curr: -5.871366 | Rate: 0.38 | Evals: 81.0%
Iter 234 | Best: -5.883215 | Curr: -3.314476 | Rate: 0.37 | Evals: 81.3%
Iter 235 | Best: -5.883215 | Curr: -1.997759 | Rate: 0.37 | Evals: 81.7%
Iter 236 | Best: -5.883215 | Curr: -5.883213 | Rate: 0.36 | Evals: 82.0%
Iter 237 | Best: -5.883216 | Rate: 0.37 | Evals: 82.3%
Iter 238 | Best: -5.883216 | Curr: -5.883212 | Rate: 0.36 | Evals: 82.7%
Iter 239 | Best: -5.883216 | Curr: -5.847127 | Rate: 0.35 | Evals: 83.0%
Iter 240 | Best: -5.883216 | Curr: -1.493860 | Rate: 0.35 | Evals: 83.3%
Iter 241 | Best: -5.883216 | Curr: -5.883214 | Rate: 0.34 | Evals: 83.7%
Iter 242 | Best: -5.883216 | Curr: -5.883210 | Rate: 0.34 | Evals: 84.0%
Iter 243 | Best: -5.883216 | Curr: -5.883215 | Rate: 0.34 | Evals: 84.3%
Iter 244 | Best: -5.883216 | Rate: 0.34 | Evals: 84.7%
Iter 245 | Best: -5.883220 | Rate: 0.35 | Evals: 85.0%
Iter 246 | Best: -5.883220 | Curr: -1.688846 | Rate: 0.34 | Evals: 85.3%
Iter 247 | Best: -5.883220 | Curr: -3.410620 | Rate: 0.34 | Evals: 85.7%
Iter 248 | Best: -5.883220 | Curr: -3.372583 | Rate: 0.34 | Evals: 86.0%
Iter 249 | Best: -5.883223 | Rate: 0.34 | Evals: 86.3%
Iter 250 | Best: -5.883223 | Curr: -3.406556 | Rate: 0.33 | Evals: 86.7%
Iter 251 | Best: -5.883223 | Curr: -5.883220 | Rate: 0.33 | Evals: 87.0%
Iter 252 | Best: -5.883223 | Curr: -3.491518 | Rate: 0.33 | Evals: 87.3%
Iter 253 | Best: -5.883224 | Rate: 0.34 | Evals: 87.7%
Iter 254 | Best: -5.883224 | Curr: -5.883209 | Rate: 0.34 | Evals: 88.0%
Iter 255 | Best: -5.883224 | Curr: -1.998252 | Rate: 0.33 | Evals: 88.3%
Iter 256 | Best: -5.883224 | Curr: -1.708372 | Rate: 0.33 | Evals: 88.7%
Iter 257 | Best: -5.883224 | Curr: -3.484460 | Rate: 0.33 | Evals: 89.0%
Iter 258 | Best: -5.883224 | Curr: -1.997583 | Rate: 0.33 | Evals: 89.3%
Iter 259 | Best: -5.883224 | Curr: -5.883219 | Rate: 0.32 | Evals: 89.7%
Iter 260 | Best: -5.883224 | Curr: -5.883218 | Rate: 0.31 | Evals: 90.0%
Iter 261 | Best: -5.883224 | Curr: -3.381047 | Rate: 0.30 | Evals: 90.3%
Iter 262 | Best: -5.883224 | Curr: -3.511563 | Rate: 0.29 | Evals: 90.7%
Iter 263 | Best: -5.883224 | Curr: -5.883221 | Rate: 0.29 | Evals: 91.0%
Iter 264 | Best: -5.883227 | Rate: 0.29 | Evals: 91.3%
Optimizer candidate 1/3 was duplicate/invalid.
Iter 265 | Best: -5.883227 | Curr: -5.867939 | Rate: 0.29 | Evals: 91.7%
Iter 266 | Best: -5.883231 | Rate: 0.29 | Evals: 92.0%
Iter 267 | Best: -5.883231 | Curr: -5.883111 | Rate: 0.28 | Evals: 92.3%
Iter 268 | Best: -5.883232 | Rate: 0.29 | Evals: 92.7%
Iter 269 | Best: -5.883232 | Curr: -5.883230 | Rate: 0.29 | Evals: 93.0%
Iter 270 | Best: -5.883232 | Curr: -2.024735 | Rate: 0.29 | Evals: 93.3%

```

7.2. SpotOptim with Sklearn Kriging in 10 Dimensions: Michalewicz Function

Iter 271	Best: -5.883232	Curr: -3.509925	Rate: 0.29	Evals: 93.7%
Iter 272	Best: -5.883232	Curr: -5.883227	Rate: 0.29	Evals: 94.0%
Iter 273	Best: -5.883232	Curr: -3.542415	Rate: 0.29	Evals: 94.3%
Iter 274	Best: -5.883232	Curr: -5.883231	Rate: 0.29	Evals: 94.7%
Iter 275	Best: -5.883232	Curr: -5.883218	Rate: 0.28	Evals: 95.0%
Iter 276	Best: -5.883232	Curr: -2.155440	Rate: 0.28	Evals: 95.3%
Iter 277	Best: -5.883232	Curr: -5.883215	Rate: 0.27	Evals: 95.7%
Iter 278	Best: -5.883232	Curr: -5.883224	Rate: 0.26	Evals: 96.0%
Iter 279	Best: -5.883232	Curr: -0.321598	Rate: 0.25	Evals: 96.3%
Iter 280	Best: -5.883232	Curr: -5.883211	Rate: 0.24	Evals: 96.7%
Iter 281	Best: -5.883232	Curr: -5.883218	Rate: 0.24	Evals: 97.0%
Iter 282	Best: -5.883232	Curr: -5.883221	Rate: 0.24	Evals: 97.3%
Iter 283	Best: -5.883232	Curr: -5.883215	Rate: 0.23	Evals: 97.7%
Iter 284	Best: -5.883232	Curr: -5.883220	Rate: 0.23	Evals: 98.0%
Iter 285	Best: -5.883232	Curr: -2.169018	Rate: 0.22	Evals: 98.3%
Iter 286	Best: -5.883237	Rate: 0.22	Evals: 98.7%	
Iter 287	Best: -5.883237	Curr: -5.883161	Rate: 0.22	Evals: 99.0%
Iter 288	Best: -5.883237	Curr: -3.526599	Rate: 0.22	Evals: 99.3%
Iter 289	Best: -5.883237	Curr: -3.394948	Rate: 0.21	Evals: 99.7%
Iter 290	Best: -5.883237	Curr: -5.883157	Rate: 0.20	Evals: 100.0%

```
print(f"[10D] Sklearn Kriging: min y = {result_micha.fun:.4f} at x = {result_micha.x}")
print(f"Number of function evaluations: {result_micha.nfev}")
print(f"Number of iterations: {result_micha.nit}")
```

```
[10D] Sklearn Kriging: min y = -5.8832 at x = [2.20300664 1.57050975 2.21940341 1.92314804 2.7656385
0.84268532 0.11299598 1.95893977 1.70658005]
Number of function evaluations: 300
Number of iterations: 290
```


7.2.4. Visualize Optimization Progress

```
import matplotlib.pyplot as plt

# Plot the optimization progress
plt.figure(figsize=(10, 6))
plt.plot(np.minimum.accumulate(opt_micha.y_), 'b-', linewidth=2)
plt.xlabel('Function Evaluations', fontsize=12)
plt.ylabel('Best Objective Value', fontsize=12)
plt.title('10D Michalewicz: Sklearn Kriging Progress', fontsize=14)
plt.grid(True, alpha=0.3)
```

7. Benchmarking SpotOptim with Sklearn Kriging (Matern Kernel) on 6D Rosenbrock and 10D Michalewicz

```
plt.tight_layout()
plt.show()
```



7.2.5. Evaluation of Multiple Repeats

To perform 30 repeats and collect statistics:

```
# Perform 30 independent runs
n_repeats = 30
results = []

print(f"Running {n_repeats} independent optimizations...")
for i in range(n_repeats):
    kernel_i = ConstantKernel(1.0, (1e-2, 1e12)) * Matern(
        length_scale=1.0,
        length_scale_bounds=(1e-4, 1e2),
        nu=2.5
    )
    surrogate_i = GaussianProcessRegressor(kernel=kernel_i, n_restarts_optimizer=100)

    opt_i = SpotOptim(
        fun=fun,
        bounds=bounds,
```

```

        n_initial=n_initial,
        max_iter=max_iter,
        surrogate=surrogate_i,
        seed=seed + i, # Different seed for each run
        verbose=0
    )

    result_i = opt_i.optimize()
    results.append(result_i.fun)

    if (i + 1) % 10 == 0:
        print(f" Completed {i + 1}/{n_repeats} runs")

# Compute statistics
mean_result = np.mean(results)
std_result = np.std(results)
min_result = np.min(results)
max_result = np.max(results)

print(f"\nResults over {n_repeats} runs:")
print(f" Mean of best values: {mean_result:.6f}")
print(f" Std of best values: {std_result:.6f}")
print(f" Min of best values: {min_result:.6f}")
print(f" Max of best values: {max_result:.6f}")

```

7.3. Comparison: SpotOptim vs SpotPython

The SpotOptim package provides a scipy-compatible interface for Bayesian optimization with the following key features:

1. **Scipy-compatible API:** Returns `OptimizeResult` objects that work seamlessly with scipy's optimization ecosystem
2. **Custom Surrogates:** Supports any sklearn-compatible surrogate model (as demonstrated with `GaussianProcessRegressor`)
3. **Flexible Interface:** Simplified parameter specification with bounds, `n_initial`, and `max_iter`
4. **Analytical Test Functions:** Built-in test functions (rosenbrock, ackley, michalewicz) for benchmarking

The main differences from spotpython are:

- **SpotOptim:** Uses bounds, `n_initial`, `max_iter` parameters with scipy-style interface

7. Benchmarking *SpotOptim* with *Sklearn Kriging (Matern Kernel)* on 6D Rosenbrock and 10D Michalewicz

- **SpotPython:** Uses `fun_control`, `design_control`, `surrogate_control` with more complex configuration

Both packages support custom surrogates and provide powerful Bayesian optimization capabilities.

7.4. Summary

This notebook demonstrated how to:

1. Use **SpotOptim** with sklearn's Gaussian Process Regressor (Matern kernel) as a surrogate
2. Optimize 6D Rosenbrock function with 100 evaluations
3. Optimize 10D Michalewicz function with 300 evaluations
4. Visualize optimization progress
5. Perform multiple independent runs for statistical analysis

The results show that **SpotOptim** with sklearn surrogates provides effective Bayesian optimization for challenging benchmark functions.

7.5. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

8. Success Rate Tracking in SpotOptim

SpotOptim tracks the **success rate** of the optimization process, which measures how often the optimizer finds improvements over recent evaluations. This metric helps you understand whether the optimization is making progress or has stalled.

8.1. What is Success Rate?

The success rate is a **rolling metric** that tracks the percentage of recent evaluations that improved upon the best value found so far. It's calculated over a sliding window of the last 100 evaluations.

Key Points: - A “success” occurs when a new evaluation finds a value **better (smaller)** than the best found so far - The rate is computed over the last **100 evaluations** (window size) - Values range from **0.0** (no recent improvements) to **1.0** (all recent evaluations improved) - Helps identify when optimization is stalling and may need adjustment

8.2. First Example

- Start TensorBoard to visualize success rate in real-time:

```
tensorboard --logdir=runs
```

The execute the following code:

```
import numpy as np
from spotoptim import SpotOptim

def rosenbrock(X):
    """Rosenbrock function - challenging optimization problem"""
    x = X[:, 0]
    y = X[:, 1]
    return (1 - x)**2 + 100 * (y - x**2)**2
```

8. Success Rate Tracking in SpotOptim

```
# Run optimization with periodic success rate checks
optimizer = SpotOptim(
    fun=rosenbrock,
    bounds=[(-2, 2), (-2, 2)],
    max_iter=20,
    n_initial=10,
    tensorboard_log=True,
    tensorboard_clean=True,
    seed=42
)

result = optimizer.optimize()

# Analyze final success rate
print(f"\nOptimization Results:")
print(f"Best value: {result.fun:.6f}")
print(f"Total evaluations: {optimizer.counter}")
print(f"Final success rate: {optimizer.success_rate:.2%}")

# Interpret the result
if optimizer.success_rate > 0.5:
    print("→ High success rate: Optimization is still making good progress")
elif optimizer.success_rate > 0.2:
    print("→ Medium success rate: Approaching convergence")
else:
    print("→ Low success rate: Optimization has likely converged")
```

```
Optimization Results:
Best value: 0.092839
Total evaluations: 20
Final success rate: 50.00%
→ Medium success rate: Approaching convergence
```

8.3. Second Example

```
from spotoptim import SpotOptim
import numpy as np

def sphere(X):
    """Simple sphere function:  $f(x) = \sum(x^2)$ """
```



```

    return np.sum(X**2, axis=1)

# Create optimizer
optimizer = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5), (-5, 5)],
    max_iter=20,
    n_initial=10,
    verbose=True
)

# Run optimization
result = optimizer.optimize()

# Check success rate
print(f"Final success rate: {optimizer.success_rate:.2%}")
print(f"Total evaluations: {optimizer.counter}")

```

```

TensorBoard logging disabled
Initial best: f(x) = 3.602357
Iter 1 | Best: 3.602357 | Curr: 14.274920 | Rate: 0.00 | Evals: 55.0%
Iter 2 | Best: 3.596919 | Rate: 0.50 | Evals: 60.0%
Iter 3 | Best: 2.616611 | Rate: 0.67 | Evals: 65.0%
Iter 4 | Best: 0.799488 | Rate: 0.75 | Evals: 70.0%
Iter 5 | Best: 0.085721 | Rate: 0.80 | Evals: 75.0%
Iter 6 | Best: 0.018392 | Rate: 0.83 | Evals: 80.0%
Iter 7 | Best: 0.013000 | Rate: 0.86 | Evals: 85.0%
Iter 8 | Best: 0.001551 | Rate: 0.88 | Evals: 90.0%
Iter 9 | Best: 0.000065 | Rate: 0.89 | Evals: 95.0%
Iter 10 | Best: 0.000006 | Rate: 0.90 | Evals: 100.0%
Final success rate: 90.00%
Total evaluations: 20

```

8.4. Accessing Success Rate

The success rate is stored in the `success_rate` attribute:

```

import numpy as np
from spotoptim import SpotOptim

def sphere(X):
    return np.sum(X**2, axis=1)

```

8. Success Rate Tracking in SpotOptim

```
optimizer = SpotOptim(fun=sphere, bounds=[(-5, 5), (-5, 5)], max_iter=20)
result = optimizer.optimize()

# Access success rate
current_rate = optimizer.success_rate
print(f"Success rate: {current_rate:.2%}")

# Also available via getter method
rate = optimizer.get_success_rate()
print(f"Via getter method: {rate:.2%}")
```

Success rate: 90.00%
Via getter method: 90.00%

8.5. Interpreting Success Rate

8.5.1. High Success Rate (> 0.5)

Success Rate: 75%

Interpretation: The optimizer is finding improvements frequently. This typically indicates: - The optimization is in an exploratory phase - The surrogate model is effectively guiding the search - There's still room for improvement in the search space

Action: Continue optimization - progress is good!

8.5.2. Medium Success Rate (0.2 - 0.5)

Success Rate: 35%

Interpretation: The optimizer occasionally finds improvements. This suggests: - The search is becoming more refined - The optimizer is balancing exploration and exploitation - Approaching a local or global optimum

Action: Monitor progress and consider stopping criteria.

8.5.3. Low Success Rate (< 0.2)

Success Rate: 8%

Interpretation: Few recent evaluations improve the best value. This may indicate: - The optimization has converged to a (local) optimum - The search is stuck in a plateau region - The budget may be exhausted in terms of meaningful progress

Action: Consider stopping optimization or adjusting parameters.

8.6. TensorBoard Visualization

When TensorBoard logging is enabled, success rate is automatically logged and can be visualized in real-time:

```
optimizer = SpotOptim(
    fun=objective,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=20,
    n_initial=10,
    tensorboard_log=True, # Enable logging
    verbose=True
)

result = optimizer.optimize()
```

View in TensorBoard:

```
tensorboard --logdir=runs
```

In the TensorBoard interface, look for: - **SCALARS** tab → **success_rate**: Rolling success rate over iterations - Compare multiple runs side-by-side - Identify when optimization stalls

8.7. Example: Comparing Multiple Runs

8. Success Rate Tracking in SpotOptim

```
import numpy as np
from spotoptim import SpotOptim

def ackley(X):
    """Ackley function - multimodal test function"""
    a = 20
    b = 0.2
    c = 2 * np.pi
    n = X.shape[1]

    sum_sq = np.sum(X**2, axis=1)
    sum_cos = np.sum(np.cos(c * X), axis=1)

    return -a * np.exp(-b * np.sqrt(sum_sq / n)) - np.exp(sum_cos / n) + a + np.e

# Run with different configurations
configs = [
    {"n_initial": 10, "max_iter": 30, "name": "Small initial"},
    {"n_initial": 20, "max_iter": 30, "name": "Large initial"},
]

results = []
for config in configs:
    optimizer = SpotOptim(
        fun=ackley,
        bounds=[(-5, 5), (-5, 5)],
        n_initial=config["n_initial"],
        max_iter=config["max_iter"],
        seed=42,
        verbose=False
    )
    result = optimizer.optimize()

    results.append({
        "name": config["name"],
        "best_value": result.fun,
        "success_rate": optimizer.success_rate,
        "n_evals": optimizer.counter
    })

print(f"\n{config['name']}:")
print(f"  Best value: {result.fun:.6f}")
print(f"  Success rate: {optimizer.success_rate:.2%}")
print(f"  Evaluations: {optimizer.counter}")
```

```
# Find best configuration
best = min(results, key=lambda x: x["best_value"])
print(f"\nBest configuration: {best['name']}")
print(f"  Achieved: f(x) = {best['best_value']:.6f}")
print(f"  Final success rate: {best['success_rate']:.2%}")
```

```
Small initial:
  Best value: 0.000738
  Success rate: 40.00%
  Evaluations: 30
```

```
Large initial:
  Best value: 0.035877
  Success rate: 60.00%
  Evaluations: 30
```

```
Best configuration: Small initial
  Achieved: f(x) = 0.000738
  Final success rate: 40.00%
```

8.8. Success Rate with Noisy Functions

For noisy functions (when `repeats_initial > 1` or `repeats_surrogate > 1`), the success rate tracks improvements in the **raw** y values, not the aggregated means:

```
import numpy as np
from spotoptim import SpotOptim

def noisy_sphere(X):
    """Sphere function with Gaussian noise"""
    base = np.sum(X**2, axis=1)
    noise = np.random.normal(0, 0.5, size=base.shape)
    return base + noise

optimizer = SpotOptim(
    fun=noisy_sphere,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=30,
    n_initial=10,
    repeats_initial=3,    # 3 evaluations per initial point
```

8. Success Rate Tracking in SpotOptim

```
repeats_surrogate=2, # 2 evaluations per new point
seed=42,
verbose=True
)

result = optimizer.optimize()

print(f"\nNoisy Optimization Results:")
print(f"Best raw value: {optimizer.min_y:.6f}")
print(f"Best mean value: {optimizer.min_mean_y:.6f}")
print(f"Success rate: {optimizer.success_rate:.2%}")
print(f"Total evaluations: {optimizer.counter}")
print(f"Unique design points: {optimizer.mean_X.shape[0]}")
```

TensorBoard logging disabled
Initial best: $f(x) = 3.219569$, mean best: $f(x) = 3.783274$

Noisy Optimization Results:
Best raw value: 3.219569
Best mean value: 3.783274
Success rate: 0.00%
Total evaluations: 30
Unique design points: 10

Note: With noisy functions, the success rate may be lower because: - Noise can mask true improvements - Multiple evaluations of the same point contribute to the window - Focus on the mean values (`min_mean_y`) for better assessment

8.9. Advanced: Custom Window Size

The success rate is calculated over a window of 100 evaluations by default. This is controlled by the `window_size` attribute:

```
import numpy as np
from spotoptim import SpotOptim

def sphere(X):
    return np.sum(X**2, axis=1)

optimizer = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
```

```

    max_iter=30,
    n_initial=10
)

# Check default window size
print(f"Window size: {optimizer.window_size}") # 100

# The window size is set during initialization
# To use a different window, you would need to modify it
# before running optimization (not typically recommended)

```

Window size: 100

8.10. Best Practices

8.10.1. 1. Monitor During Long Runs

For expensive optimization runs, periodically check success rate:

```

# Could be implemented with callbacks in future versions
# For now, success rate is updated automatically and logged to TensorBoard

```

8.10.2. 2. Combine with TensorBoard

Always enable TensorBoard logging for visual monitoring:

```

optimizer = SpotOptim(
    fun=expensive_function,
    bounds=bounds,
    max_iter=25,
    tensorboard_log=True, # Track success_rate visually
    tensorboard_path="runs/long_optimization"
)

```

8.10.3. 3. Use as Stopping Criterion

Consider stopping when success rate drops very low:

8. Success Rate Tracking in SpotOptim

```
# Manual stopping check (conceptual)
if optimizer.success_rate < 0.05 and optimizer.counter > 50:
    print("Success rate very low - optimization has likely converged")
```

8.10.4. 4. Compare Different Strategies

Use success rate to compare optimization strategies:

```
strategies = ["ei", "pi", "y"] # Different acquisition functions
for acq in strategies:
    opt = SpotOptim(fun=obj, bounds=bnds, acquisition=acq, max_iter=25)
    result = opt.optimize()
    print(f"{acq}: success_rate={opt.success_rate:.2%}, best={result.fun:.6f}")
```

8.11. Technical Details

8.11.1. How Success is Counted

A new evaluation `y_new` is considered a success if:

```
y_new < best_y_so_far
```

where `best_y_so_far` is the minimum value found in all previous evaluations.

8.11.2. Rolling Window Calculation

The success rate is computed as:

```
success_rate = (number of successes in last 100 evals) / (window size)
```

- Window size defaults to 100
- If fewer than 100 evaluations have been performed, the window size is the number of evaluations
- The window slides forward with each new evaluation

8.11.3. Update Frequency

The success rate is updated after: 1. Initial design evaluation 2. Each iteration's new point evaluation(s) 3. OCBA re-evaluations (if applicable)

8.12. Summary

- **Success rate** measures the percentage of recent evaluations that improve the best value
- Calculated over a rolling window of the last **100 evaluations**
- Values range from **0.0** to **1.0**
- High rates (>0.5) indicate active progress
- Low rates (<0.2) suggest convergence
- Automatically logged to **TensorBoard** when logging is enabled
- Available via `optimizer.success_rate` attribute after optimization

Use success rate to: - x Monitor optimization progress in real-time - x Identify when to stop optimization - x Compare different optimization strategies - x Assess optimization difficulty for different problems

8.13. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

9. Parallel Evaluation

This chapter explains how `SpotOptim` evaluates the objective function and how a user can still exploit several processor cores even though the optimizer itself runs sequentially.

9.1. Sequential evaluation

As of version 1.0.0, `SpotOptim` evaluates the objective function sequentially. The earlier steady-state asynchronous scheme, controlled by an `n_jobs` argument, was removed when the core was reduced to a lean sequential implementation. The constructor therefore no longer accepts `n_jobs` or `eval_batch_size`, and the public entry point is the single `optimize()` call shown throughout this book.

The motivation for the sequential design is reproducibility and a smaller dependency surface. With a fixed `seed`, a sequential run reproduces the same sequence of infill points on every machine, which matters for the benchmark and leaderboard results that depend on this library. Removing the parallel executor also dropped the heavy process-management dependencies from the core install.

Sequential evaluation does not prevent the use of multiple cores. Two complementary strategies remain available, and they cover the cases that the old `n_jobs` option addressed.

9.2. Parallelism inside the objective function

`SpotOptim` passes the candidate points to the objective function as a two-dimensional array whose rows are individual points. A vectorized objective that operates on the whole array uses the optimized linear-algebra backend, which is already multi-threaded, so a single call evaluates the batch efficiently without any explicit parallel code.

```
import warnings
import numpy as np
from spotoptim import SpotOptim

warnings.filterwarnings("ignore")
```

9. Parallel Evaluation

```
def sphere_batch(X):
    # X has shape (n_points, n_dim); evaluate all rows at once.
    return np.sum(X**2, axis=1)

optimizer = SpotOptim(
    fun=sphere_batch,
    bounds=[(-5, 5)] * 2,
    n_initial=5,
    max_iter=15,
    seed=42,
    verbose=False,
)
result = optimizer.optimize()

print(f"best x    : {np.round(result.x, 4)}")
print(f"best f(x): {result.fun:.6f}")
print(f"n eval   : {result.nfev}")
```

```
best x    : [0.0002 0.0004]
best f(x): 0.000000
n eval   : 15
```

When the per-point cost is dominated by a routine that releases the global interpreter lock, such as a NumPy or PyTorch operation, the objective can also spread the rows across a thread pool inside the function body. The optimizer sees a single call and remains deterministic, while the heavy work runs concurrently.

9.3. Independent optimizations across processes

The second strategy parallelizes at the level of whole optimization runs rather than single evaluations. Independent runs, for example repeated restarts from different seeds or the same tuning problem applied to several data splits, are embarrassingly parallel and can be dispatched to separate processes with the standard library. Each process executes one self-contained sequential `optimize()` call.

The following pattern illustrates the approach. It is shown for reference and is not executed during the rendering of this page, because process pools require a script entry point rather than a notebook cell.

```

from concurrent.futures import ProcessPoolExecutor
import numpy as np
from spotoptim import SpotOptim

def sphere_batch(X):
    return np.sum(X**2, axis=1)

def run_once(seed):
    optimizer = SpotOptim(
        fun=sphere_batch,
        bounds=[(-5, 5)] * 2,
        n_initial=5,
        max_iter=15,
        seed=seed,
        verbose=False,
    )
    result = optimizer.optimize()
    return seed, float(result.fun)

if __name__ == "__main__":
    seeds = [42, 43, 44, 45]
    with ProcessPoolExecutor() as pool:
        for seed, best in pool.map(run_once, seeds):
            print(f"seed {seed}: best f(x) = {best:.6f}")

```

This keeps each optimization reproducible, since the seed is fixed per run, while the wall-clock time scales with the number of available cores rather than with the number of runs.

9.4. Jupyter Notebook

i Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

10. Restart Mechanism in SpotOptim

SpotOptim includes a built-in restart mechanism designed to help the optimizer escape local optima or recover from stagnation. This feature is particularly useful for difficult landscapes where the optimizer might get stuck in a suboptimal region.

10.1. Key Concepts

The restart mechanism monitors the optimization progress and triggers a complete reset (restart) of the optimization run if no improvement is observed for a specified number of iterations.

10.1.1. Parameters

Two key parameters control this behavior:

- **restart_after_n** (int, default=100): The number of consecutive iterations with a success rate of 0.0 (no improvement) required to trigger a restart.
- **restart_inject_best** (bool, default=True): If **True**, the best solution found in all previous runs is injected into the initial design of the new restart run. This ensures that the global search does not lose the best-known solution while exploring new regions.

10.2. How it Works

1. **Monitoring:** During optimization, SpotOptim tracks the **success_rate** (percentage of valid and improved points in the current window).
2. **Triggering:** If the success rate drops to 0.0 and stays there for **restart_after_n** consecutive iterations, the current run is terminated.
3. **Restarting:** A new optimization run is initialized.
 - A new random seed is generated (if running sequentially) to ensure a different random start.
 - A new initial design (LHS) is created.

10. Restart Mechanism in SpotOptim

4. **Injection:** If `restart_inject_best=True`, the overall best point found so far is added to this new initial design.
5. **Aggregation:** When the global `max_iter` or `max_time` is reached, results from all runs are aggregated. The final returned result corresponds to the best run found.

10.3. Example: Triggering Restarts

In this example, we set `restart_after_n` to a very small value (5) to intentionally force restarts and demonstrate the mechanism. We use a multimodal function where getting stuck is possible.

```
import numpy as np
from spotoptim import SpotOptim

def multimodal_function(X):
    """A simple 2D multimodal function (Ackley-like structure)"""
    X = np.atleast_2d(X)
    return -20 * np.exp(-0.2 * np.sqrt(0.5 * np.sum(X**2, axis=1))) - \
        np.exp(0.5 * np.sum(np.cos(2 * np.pi * X), axis=1)) + \
        20 + np.exp(1)

# Configure optimizer with aggressive restart strategy
optimizer = SpotOptim(
    fun=multimodal_function,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=50,          # Total global budget
    n_initial=5,
    restart_after_n=5,    # Restart after only 5 iterations of no improvement
    restart_inject_best=True,
    seed=42,
    verbose=True          # Verbose output shows restart messages
)

result = optimizer.optimize()
```

```
TensorBoard logging disabled
Initial best: f(x) = 6.220297
Iter 1 | Best: 6.219898 | Rate: 1.00 | Evals: 12.0%
Iter 2 | Best: 6.073743 | Rate: 1.00 | Evals: 14.0%
Iter 3 | Best: 4.573169 | Rate: 1.00 | Evals: 16.0%
Iter 4 | Best: 3.122935 | Rate: 1.00 | Evals: 18.0%
```


10.3. Example: Triggering Restarts

```
Iter 5 | Best: 2.198210 | Rate: 1.00 | Evals: 20.0%
Iter 6 | Best: 2.198210 | Curr: 4.141608 | Rate: 0.80 | Evals: 22.0%
Iter 7 | Best: 1.314662 | Rate: 0.80 | Evals: 24.0%
Iter 8 | Best: 1.314662 | Curr: 3.943707 | Rate: 0.60 | Evals: 26.0%
Iter 9 | Best: 1.314662 | Curr: 1.323912 | Rate: 0.40 | Evals: 28.0%
Iter 10 | Best: 1.314662 | Curr: 1.458468 | Rate: 0.20 | Evals: 30.0%
Iter 11 | Best: 0.051957 | Rate: 0.40 | Evals: 32.0%
Iter 12 | Best: 0.051957 | Curr: 0.730561 | Rate: 0.20 | Evals: 34.0%
Iter 13 | Best: 0.049611 | Rate: 0.40 | Evals: 36.0%
Iter 14 | Best: 0.012923 | Rate: 0.60 | Evals: 38.0%
Iter 15 | Best: 0.012923 | Curr: 0.108862 | Rate: 0.60 | Evals: 40.0%
Iter 16 | Best: 0.009616 | Rate: 0.60 | Evals: 42.0%
Iter 17 | Best: 0.003379 | Rate: 0.80 | Evals: 44.0%
Iter 18 | Best: 0.003379 | Curr: 0.435782 | Rate: 0.60 | Evals: 46.0%
Iter 19 | Best: 0.003379 | Curr: 0.446607 | Rate: 0.40 | Evals: 48.0%
Iter 20 | Best: 0.003379 | Curr: 13.146239 | Rate: 0.40 | Evals: 50.0%
Iter 21 | Best: 0.003379 | Curr: 0.122371 | Rate: 0.20 | Evals: 52.0%
Iter 22 | Best: 0.003379 | Curr: 0.028373 | Rate: 0.00 | Evals: 54.0%
Iter 23 | Best: 0.002457 | Rate: 0.20 | Evals: 56.0%
Iter 24 | Best: 0.002457 | Curr: 0.321400 | Rate: 0.20 | Evals: 58.0%
Iter 25 | Best: 0.000968 | Rate: 0.40 | Evals: 60.0%
Iter 26 | Best: 0.000968 | Curr: 10.483186 | Rate: 0.40 | Evals: 62.0%
Iter 27 | Best: 0.000968 | Curr: 0.224885 | Rate: 0.40 | Evals: 64.0%
Iter 28 | Best: 0.000968 | Curr: 0.007594 | Rate: 0.20 | Evals: 66.0%
Iter 29 | Best: 0.000968 | Curr: 0.244888 | Rate: 0.20 | Evals: 68.0%
Iter 30 | Best: 0.000968 | Curr: 0.165361 | Rate: 0.00 | Evals: 70.0%
Iter 31 | Best: 0.000968 | Curr: 12.590502 | Rate: 0.00 | Evals: 72.0%
Iter 32 | Best: 0.000968 | Curr: 9.466552 | Rate: 0.00 | Evals: 74.0%
Iter 33 | Best: 0.000968 | Curr: 0.107702 | Rate: 0.00 | Evals: 76.0%
Restarting optimization: success_rate 0 for 5 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.00016887 0.00029626]
  Internal scale: [0.00016887 0.00029626]
Restart injection: Using best found point so far as starting point (f(x)=0.000968).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 0.000968
Iter 1 | Best: 0.000968 | Curr: 0.012614 | Rate: 0.00 | Evals: 88.0%
Iter 2 | Best: 0.000968 | Curr: 3.327522 | Rate: 0.00 | Evals: 90.0%
Iter 3 | Best: 0.000968 | Curr: 0.236847 | Rate: 0.00 | Evals: 92.0%
Iter 4 | Best: 0.000968 | Curr: 0.944229 | Rate: 0.00 | Evals: 94.0%
Restarting optimization: success_rate 0 for 5 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.00016887 0.00029626]
  Internal scale: [0.00016887 0.00029626]
```

10. Restart Mechanism in SpotOptim

Restart injection: Using best found point so far as starting point ($f(x)=0.000968$).
Global budget exhausted. Stopping restarts.

10.3.1. Analyzing Restart Results

You can access the results of each individual restart run via the `restarts_results_` attribute.

```
print(f"Total global evaluations: {result.nfev}")
print(f"Number of restarts performed: {len(optimizer.restarts_results_) - 1}")
print(f"Best value found globally: {result.fun:.6f}")

print("\nBreakdown by run:")
for i, res in enumerate(optimizer.restarts_results_):
    print(f"  Run {i+1}: {res.nfev} evals, Best: {res.fun:.6f}, Status: {res.message}")
```

```
Total global evaluations: 38
Number of restarts performed: 1
Best value found globally: 0.000968
```

Breakdown by run:

```
Run 1: 38 evals, Best: 0.000968, Status: Restart triggered due to lack of improvement
Run 2: 9 evals, Best: 0.000968, Status: Restart triggered due to lack of improvement
```

10.4. Example: Effect of `restart_inject_best`

The `restart_inject_best` parameter is crucial for efficiency. It ensures that “knowledge” is transferred between restarts.

- **True:** The new run starts with the best point found so far included in its initial set. This allows the surrogate model to immediately be aware of the high-quality region, potentially refining it further or using it as a baseline to explore elsewhere.
- **False:** Each restart is completely independent. This is equivalent to running the optimizer multiple times in parallel with different seeds and taking the best result.

Run with injection DISABLED:

```
opt_no_inject = SpotOptim(
    fun=multimodal_function,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=40,
```

10.4. Example: Effect of `restart_inject_best`

```
n_initial=5,  
restart_after_n=3,  
restart_inject_best=False, # Disable injection  
seed=42,  
verbose=True  
)  
res_no_inject = opt_no_inject.optimize()
```

TensorBoard logging disabled

Initial best: $f(x) = 6.220297$

Iter 1	Best: 6.219898	Rate: 1.00	Evals: 15.0%
Iter 2	Best: 6.073743	Rate: 1.00	Evals: 17.5%
Iter 3	Best: 4.573169	Rate: 1.00	Evals: 20.0%
Iter 4	Best: 3.122935	Rate: 1.00	Evals: 22.5%
Iter 5	Best: 2.198210	Rate: 1.00	Evals: 25.0%
Iter 6	Best: 2.198210	Curr: 4.141608	Rate: 0.67 Evals: 27.5%
Iter 7	Best: 1.314662	Rate: 0.67	Evals: 30.0%
Iter 8	Best: 1.314662	Curr: 3.943707	Rate: 0.33 Evals: 32.5%
Iter 9	Best: 1.314662	Curr: 1.323912	Rate: 0.33 Evals: 35.0%
Iter 10	Best: 1.314662	Curr: 1.458468	Rate: 0.00 Evals: 37.5%
Iter 11	Best: 0.051957	Rate: 0.33	Evals: 40.0%
Iter 12	Best: 0.051957	Curr: 0.730561	Rate: 0.33 Evals: 42.5%
Iter 13	Best: 0.049611	Rate: 0.67	Evals: 45.0%
Iter 14	Best: 0.012923	Rate: 0.67	Evals: 47.5%
Iter 15	Best: 0.012923	Curr: 0.108862	Rate: 0.67 Evals: 50.0%
Iter 16	Best: 0.009616	Rate: 0.67	Evals: 52.5%
Iter 17	Best: 0.003379	Rate: 0.67	Evals: 55.0%
Iter 18	Best: 0.003379	Curr: 0.435782	Rate: 0.67 Evals: 57.5%
Iter 19	Best: 0.003379	Curr: 0.446607	Rate: 0.33 Evals: 60.0%
Iter 20	Best: 0.003379	Curr: 13.146239	Rate: 0.00 Evals: 62.5%
Iter 21	Best: 0.003379	Curr: 0.122371	Rate: 0.00 Evals: 65.0%

Restarting optimization: success_rate 0 for 3 iterations.

Initial best: $f(x) = 6.528733$

Optimizer candidate 1/3 was duplicate/invalid.

Optimizer candidate 2/3 was duplicate/invalid.

Optimizer candidate 3/3 was duplicate/invalid.

Fallback attempt 1/10: Using fallback strategy

Acquisition failure: Using random space-filling design as fallback.

Iter 1	Best: 6.528733	Curr: 8.764308	Rate: 0.00 Evals: 80.0%
Iter 2	Best: 6.106511	Rate: 0.50	Evals: 82.5%
Iter 3	Best: 6.106511	Curr: 6.141971	Rate: 0.33 Evals: 85.0%
Iter 4	Best: 5.963654	Rate: 0.67	Evals: 87.5%
Iter 5	Best: 5.929160	Rate: 0.67	Evals: 90.0%
Iter 6	Best: 5.929160	Curr: 6.386849	Rate: 0.67 Evals: 92.5%

10. Restart Mechanism in SpotOptim

```
Iter 7 | Best: 5.457406 | Rate: 0.67 | Evals: 95.0%
Iter 8 | Best: 5.404737 | Rate: 0.67 | Evals: 97.5%
Iter 9 | Best: 5.386152 | Rate: 1.00 | Evals: 100.0%
```

Run with injection ENABLED:

```
opt_inject = SpotOptim(
    fun=multimodal_function,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=40,
    n_initial=5,
    restart_after_n=3,
    restart_inject_best=True, # Enable injection
    seed=42,
    verbose=True
)
res_inject = opt_inject.optimize()
```

TensorBoard logging disabled

Initial best: $f(x) = 6.220297$

```
Iter 1 | Best: 6.219898 | Rate: 1.00 | Evals: 15.0%
Iter 2 | Best: 6.073743 | Rate: 1.00 | Evals: 17.5%
Iter 3 | Best: 4.573169 | Rate: 1.00 | Evals: 20.0%
Iter 4 | Best: 3.122935 | Rate: 1.00 | Evals: 22.5%
Iter 5 | Best: 2.198210 | Rate: 1.00 | Evals: 25.0%
Iter 6 | Best: 2.198210 | Curr: 4.141608 | Rate: 0.67 | Evals: 27.5%
Iter 7 | Best: 1.314662 | Rate: 0.67 | Evals: 30.0%
Iter 8 | Best: 1.314662 | Curr: 3.943707 | Rate: 0.33 | Evals: 32.5%
Iter 9 | Best: 1.314662 | Curr: 1.323912 | Rate: 0.33 | Evals: 35.0%
Iter 10 | Best: 1.314662 | Curr: 1.458468 | Rate: 0.00 | Evals: 37.5%
Iter 11 | Best: 0.051957 | Rate: 0.33 | Evals: 40.0%
Iter 12 | Best: 0.051957 | Curr: 0.730561 | Rate: 0.33 | Evals: 42.5%
Iter 13 | Best: 0.049611 | Rate: 0.67 | Evals: 45.0%
Iter 14 | Best: 0.012923 | Rate: 0.67 | Evals: 47.5%
Iter 15 | Best: 0.012923 | Curr: 0.108862 | Rate: 0.67 | Evals: 50.0%
Iter 16 | Best: 0.009616 | Rate: 0.67 | Evals: 52.5%
Iter 17 | Best: 0.003379 | Rate: 0.67 | Evals: 55.0%
Iter 18 | Best: 0.003379 | Curr: 0.435782 | Rate: 0.67 | Evals: 57.5%
Iter 19 | Best: 0.003379 | Curr: 0.446607 | Rate: 0.33 | Evals: 60.0%
Iter 20 | Best: 0.003379 | Curr: 13.146239 | Rate: 0.00 | Evals: 62.5%
Iter 21 | Best: 0.003379 | Curr: 0.122371 | Rate: 0.00 | Evals: 65.0%
Restarting optimization: success_rate 0 for 3 iterations.
Starting point x0 validated and processed successfully.
Original scale: [-0.00105728  0.00052747]
```

10.5. When to Use Restarts?

```
Internal scale: [-0.00105728  0.00052747]
Restart injection: Using best found point so far as starting point (f(x)=0.003379).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 0.003379
Iter 1 | Best: 0.003379 | Curr: 0.010678 | Rate: 0.00 | Evals: 80.0%
Iter 2 | Best: 0.003379 | Curr: 2.766309 | Rate: 0.00 | Evals: 82.5%
Restarting optimization: success_rate 0 for 3 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [-0.00105728  0.00052747]
  Internal scale: [-0.00105728  0.00052747]
Restart injection: Using best found point so far as starting point (f(x)=0.003379).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 0.003379
Iter 1 | Best: 0.003379 | Curr: 0.006894 | Rate: 0.00 | Evals: 97.5%
Iter 2 | Best: 0.003379 | Curr: 2.245994 | Rate: 0.00 | Evals: 100.0%
```

```
print(f"Best without injection: {res_no_inject.fun:.6f}")
print(f"Best with injection:    {res_inject.fun:.6f}")
```

```
Best without injection: 0.003379
Best with injection:    0.003379
```

10.5. When to Use Restarts?

- **Complex Landscapes:** When the objective function has many local optima.
- **Stagnation:** When the optimizer tends to “flatline” early but `max_iter` is large.
- **exploration vs Exploitation:** Restarts favor exploration (by jumping to a new random initial design) when exploitation (local improvement) has seemingly exhausted the current basin of attraction.

Setting `restart_after_n` depends on your problem:

- **Low values (e.g., 10-20):** Aggressive restarts. Good if function evaluation is cheap and you want to explore many basins.
- **High values (e.g., 100+):** Conservative. Gives the optimizer plenty of time to refine the solution in the current basin before giving up.

Part II.

Variables and Hyperparameters

11. Variable Type (`var_type`) Implementation in SpotOptim

11.1. Overview

This document describes the `var_type` implementation in SpotOptim, which allows users to specify different data types for optimization variables.

11.2. Supported Variable Types

SpotOptim supports three main data types:

11.2.1. 1. 'float'

- **Purpose:** Continuous optimization with Python floats
- **Behavior:** No rounding applied, values remain continuous
- **Use case:** Standard continuous optimization variables
- **Example:** Temperature (23.5°C), Distance (1.234m)

11.2.2. 2. 'int'

- **Purpose:** Discrete integer optimization
- **Behavior:** Float values are automatically rounded to integers
- **Use case:** Count variables, discrete parameters
- **Example:** Number of layers (5), Population size (100)

11.2.3. 3. 'factor'

- **Purpose:** Unordered categorical data
- **Behavior:** Internally mapped to integer values (0, 1, 2, ...)
- **Use case:** Categorical choices like colors, algorithms, modes
- **Example:** Color ("red"→0, "green"→1, "blue"→2)

11. Variable Type (*var_type*) Implementation in *SpotOptim*

- **Note:** The actual string-to-int mapping is external to *SpotOptim*; the optimizer works with the integer representation

11.3. Implementation Details

11.3.1. Where *var_type* is Used

The *var_type* parameter is properly propagated throughout the optimization process:

1. **Initialization** (*__init__*):
 - Stored as `self.var_type`
 - Default: `["float"] * n_dim` if not specified
2. **Initial Design Generation** (*generate_initial_design*):
 - Applies type constraints via `repair_non_numeric()`
 - Ensures initial points respect variable types
3. **New Point Suggestion** (*suggest_next_infill_point*):
 - Applies type constraints to acquisition function optimization results
 - Ensures suggested points respect variable types
4. **User-Provided Initial Design** (*optimize*):
 - Applies type constraints to *X0* if provided
 - Ensures consistency regardless of input source
5. **Mesh Grid Generation** (*_generate_mesh_grid*):
 - Used for plotting, respects variable types
 - Ensures visualization shows correct discrete/continuous behavior

11.3.2. Core Method: `repair_non_numeric()`

This method enforces variable type constraints:

```
def repair_non_numeric(self, X: np.ndarray, var_type: List[str]) -> np.ndarray:
    """Round non-continuous values to integers."""
    mask = np.isin(var_type, ["float"], invert=True)
    X[:, mask] = np.around(X[:, mask])
    return X
```

Logic:

- Variables with type 'float': No change (continuous)
- Variables with type 'int' or 'factor': Rounded to integers

11.4. Example Usage

11.4.1. Example 1: All Float Variables (Default)

```
import numpy as np
from spotoptim import SpotOptim

# Example 1: All float variables (default)
opt1 = SpotOptim(
    fun=lambda X: np.sum(X**2, axis=1),
    bounds=[(0, 10), (0, 10), (0, 10)],
    max_iter=20,
    n_initial=10,
    seed=42
    # var_type defaults to ["float", "float", "float"]
)
result1 = opt1.optimize()
print(f"Best value: {result1.fun:.6f}")
print(f"Best point (floats): {result1.x}")
```

```
Best value: 0.000000
Best point (floats): [0. 0. 0.]
```

11.4.2. Example 2: Pure Integer Optimization

```
import numpy as np
from spotoptim import SpotOptim

def discrete_func(X):
    return np.sum(X**2, axis=1)

bounds = [(-5, 5), (-5, 5)]
var_type = ["int", "int"]

opt = SpotOptim(
    fun=discrete_func,
```

11. Variable Type (*var_type*) Implementation in *SpotOptim*

```
        bounds=bounds,
        var_type=var_type,
        max_iter=20,
        n_initial=10,
        seed=42
    )

    result = opt.optimize()
    print(f"Best value: {result.fun:.6f}")
    print(f"Best point (integers): {result.x}")
    print(f>Note: Values are rounded to integers")
```

```
Best value: 0.000000
Best point (integers): [-0.  0.]
Note: Values are rounded to integers
```

11.4.3. Example 3: Categorical (Factor) Variables

```
import numpy as np
from spotoptim import SpotOptim

def categorical_func(X):
    # Assume X[:, 0] represents 3 categories: 0, 1, 2
    # Category 0 is best
    return (X[:, 0]**2) + (X[:, 1]**2)

bounds = [(0, 2), (0, 3)] # 3 and 4 categories respectively
var_type = ["factor", "factor"]

opt = SpotOptim(
    fun=categorical_func,
    bounds=bounds,
    var_type=var_type,
    max_iter=20,
    n_initial=10,
    seed=42
)

result = opt.optimize()
print(f"Best value: {result.fun:.6f}")
print(f"Best point (categories): {result.x}")
print(f>Note: Values are integers representing categories")
```

Best value: 0.000000
 Best point (categories): [0. 0.]
 Note: Values are integers representing categories

11.4.4. Example 4: Mixed Variable Types

```
import numpy as np
from spotoptim import SpotOptim

def mixed_func(X):
    # X[:, 0]: continuous temperature
    # X[:, 1]: discrete number of iterations
    # X[:, 2]: categorical algorithm choice (0, 1, 2)
    return X[:, 0]**2 + X[:, 1]**2 + X[:, 2]**2

bounds = [(-5, 5), (1, 100), (0, 2)]
var_type = ["float", "int", "factor"]
var_name = ["temperature", "iterations", "algorithm"]

opt = SpotOptim(
    fun=mixed_func,
    bounds=bounds,
    var_type=var_type,
    var_name=var_name,
    max_iter=20,
    n_initial=10,
    seed=42
)

result = opt.optimize()
print(f"Best value: {result.fun:.6f}")
print(f"Best point: {result.x}")
print(f"  {var_name[0]} (float): {result.x[0]:.6f}")
print(f"  {var_name[1]} (int): {int(result.x[1])}")
print(f"  {var_name[2]} (factor): {int(result.x[2])}")
```

Best value: 1.000010
 Best point: [0.00314215 1. 0.]
 temperature (float): 0.003142
 iterations (int): 1
 algorithm (factor): 0

11.5. Key Findings

1. **Type Persistence:** Variable types are correctly maintained throughout the entire optimization process, from initial design through all iterations.
2. **Automatic Enforcement:** The `repair_non_numeric()` method is called at all critical points, ensuring type constraints are never violated.
3. **Three Explicit Types:** Only 'float', 'int', and 'factor' are supported. The legacy 'num' type has been removed for clarity.
4. **User-Provided Data:** Type constraints are applied even to user-provided initial designs, ensuring consistency.
5. **Plotting Compatibility:** The plotting functionality respects variable types, ensuring correct visualization of discrete vs. continuous variables.

11.6. Recommendations

1. **Always specify `var_type` explicitly** for clarity, especially in mixed-type problems
2. **Use appropriate bounds** for factor variables (e.g., `(0, n_categories-1)`)
3. **External mapping** for string categories: Maintain your own mapping dictionary outside *SpotOptim* (e.g., `{"red": 0, "green": 1, "blue": 2}`)
4. **Validation:** The current implementation doesn't validate `var_type` length matches bounds length - users should ensure this manually

11.7. Future Enhancements (Optional)

Potential improvements that could be added:

1. **Validation:** Add validation in `__init__` to check `len(var_type) == len(bounds)`
2. **String Categories:** Add built-in support for automatic string-to-int mapping
3. **Ordered Categories:** Support ordered categorical variables (ordinal data)
4. **Type Checking:** Validate that `var_type` values are one of the allowed strings
5. **Bounds Checking:** Warn if factor bounds are not integer ranges

11.8. Jupyter Notebook

i Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

12. Factor Variables for Categorical Hyperparameters

SpotOptim supports factor variables for optimizing categorical hyperparameters, such as activation functions, optimizers, or any discrete string-based choices. Factor variables are automatically converted between string values (external interface) and integers (internal optimization), making categorical optimization seamless.

12.1. Overview

What are Factor Variables?

Factor variables allow you to specify categorical choices as tuples of strings in the bounds. SpotOptim handles the conversion:

1. **String tuples in bounds** → Internal integer mapping (0, 1, 2, ...)
2. **Optimization uses integers** internally for surrogate modeling
3. **Objective function receives strings** after automatic conversion
4. **Results return strings** (not integers)

Module: `spotoptim.SpotOptim`

Key Features:

- Define categorical choices as string tuples: ("ReLU", "Sigmoid", "Tanh")
- Automatic integer↔string conversion
- Seamless integration with neural network hyperparameters
- Mix factor variables with numeric/integer variables

12.2. Quick Start

12.2.1. Basic Factor Variable Usage

12. Factor Variables for Categorical Hyperparameters

```
from spotoptim import SpotOptim
import numpy as np

def objective_function(X):
    """Objective function receives string values."""
    results = []
    for params in X:
        activation = params[0] # This is a string!
        print(f"Testing activation: {activation}")

        # Simple scoring based on activation choice (for demonstration)
        # In real use, you would train a model and return actual performance
        scores = {
            "ReLU": 3500.0,
            "Sigmoid": 4200.0,
            "Tanh": 3800.0,
            "LeakyReLU": 3600.0
        }
        score = scores.get(activation, 5000.0) + np.random.normal(0, 100)
        results.append(score)
    return np.array(results) # Return numpy array

# Define bounds with factor variable
optimizer = SpotOptim(
    fun=objective_function,
    bounds=[("ReLU", "Sigmoid", "Tanh", "LeakyReLU")],
    var_type=["factor"],
    max_iter=20,
    seed=42
)

result = optimizer.optimize()
print(f"\nBest activation: {result.x[0]}") # Returns string, e.g., "ReLU"
print(f"Best score: {result.fun:.4f}")
```

/var/folders/dw/pvtj6mt91znd0hftcztqb0k00000gn/T/ipykernel_37163/3329959676.py:24: Use
agnostic) factor metrics. Factor integer codes will be treated as ordinal numbers, wh
aware Kriging surrogate: Kriging(metric_factorial='hamming').
optimizer = SpotOptim(

Testing activation: ReLU
Testing activation: Sigmoid
Testing activation: Tanh

```

Testing activation: LeakyReLU
Testing activation: Tanh
Testing activation: Tanh
Testing activation: Sigmoid
Testing activation: Tanh
Testing activation: Tanh
Testing activation: Tanh
Testing activation: Sigmoid
Testing activation: Sigmoid
Testing activation: ReLU
Testing activation: Tanh
Testing activation: LeakyReLU
Testing activation: LeakyReLU
Testing activation: LeakyReLU
Testing activation: Tanh
Testing activation: Sigmoid
Testing activation: LeakyReLU

```

```

Best activation: LeakyReLU
Best score: 3427.5082

```

12.2.2. Neural Network Activation Function Optimization

```

import torch
import torch.nn as nn
from spotoptim import SpotOptim
from spotoptim.data import get_diabetes_dataloaders
from spotoptim.nn.linear_regressor import LinearRegressor
import numpy as np

def train_and_evaluate(X):
    """Train models with different activation functions."""
    results = []

    for params in X:
        activation = params[0] # String: "ReLU", "Sigmoid", etc.

        # Load data
        train_loader, test_loader, _ = get_diabetes_dataloaders()

        # Create model with the activation function
        model = LinearRegressor(
            input_dim=10,

```

12. Factor Variables for Categorical Hyperparameters

```
        output_dim=1,
        l1=64,
        num_hidden_layers=2,
        activation=activation # Pass string directly!
    )

    # Train model
    optimizer = model.get_optimizer("Adam", lr=0.01)
    criterion = nn.MSELoss()

    for epoch in range(50):
        model.train()
        for batch_X, batch_y in train_loader:
            predictions = model(batch_X)
            loss = criterion(predictions, batch_y)
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

    # Evaluate
    model.eval()
    test_loss = 0.0
    with torch.no_grad():
        for batch_X, batch_y in test_loader:
            predictions = model(batch_X)
            test_loss += criterion(predictions, batch_y).item()

    avg_loss = test_loss / len(test_loader)
    results.append(avg_loss)

    return np.array(results) # Return numpy array

# Optimize activation function choice
optimizer = SpotOptim(
    fun=train_and_evaluate,
    bounds=[("ReLU", "Sigmoid", "Tanh", "LeakyReLU", "ELU")],
    var_type=["factor"],
    max_iter=30
)

result = optimizer.optimize()
print(f"Best activation function: {result.x[0]}")
print(f"Best test MSE: {result.fun:.4f}")
```

```

/var/folders/dw/pvtj6mt91znd0hftcztb0k00000gn/T/ipykernel_37163/1439913637.py:54: UserWarning: Factor
agnostic) factor metrics. Factor integer codes will be treated as ordinal numbers, which may mislead
aware Kriging surrogate: Kriging(metric_factorial='hamming').
optimizer = SpotOptim(

```

```

Best activation function: Sigmoid
Best test MSE: 26374.1868

```

12.3. Mixed Variable Types

12.3.1. Combining Factor, Integer, and Continuous Variables

```

import numpy as np
import torch
import torch.nn as nn
from spotoptim import SpotOptim
from spotoptim.data import get_diabetes_dataloaders
from spotoptim.nn.linear_regressor import LinearRegressor

def comprehensive_optimization(X):
    """Optimize learning rate, layer size, depth, and activation."""
    results = []

    for params in X:
        log_lr = params[0]      # Continuous (log scale)
        l1 = int(params[1])     # Integer
        n_layers = int(params[2]) # Integer
        activation = params[3]   # Factor (string)

        lr = 10 ** log_lr      # Convert from log scale

        print(f"lr={lr:.6f}, l1={l1}, layers={n_layers}, activation={activation}")

        # Load data
        train_loader, test_loader, _ = get_diabetes_dataloaders(
            batch_size=32,
            random_state=42
        )

        # Create model
        model = LinearRegressor(

```

12. Factor Variables for Categorical Hyperparameters

```
        input_dim=10,
        output_dim=1,
        l1=l1,
        num_hidden_layers=n_layers,
        activation=activation
    )

    # Train
    optimizer = model.get_optimizer("Adam", lr=lr)
    criterion = nn.MSELoss()

    for epoch in range(30):
        model.train()
        for batch_X, batch_y in train_loader:
            predictions = model(batch_X)
            loss = criterion(predictions, batch_y)
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

    # Evaluate
    model.eval()
    test_loss = 0.0
    with torch.no_grad():
        for batch_X, batch_y in test_loader:
            predictions = model(batch_X)
            test_loss += criterion(predictions, batch_y).item()

    results.append(test_loss / len(test_loader))

    return np.array(results)

# Optimize all four hyperparameters simultaneously
optimizer = SpotOptim(
    fun=comprehensive_optimization,
    bounds=[
        (-4, -2), # log10(learning_rate)
        (16, 128), # l1 (neurons per layer)
        (0, 4), # num_hidden_layers
        ("ReLU", "Sigmoid", "Tanh", "LeakyReLU") # activation function
    ],
    var_type=["float", "int", "int", "factor"],
    max_iter=50
)
```

```

result = optimizer.optimize()

# Results contain original string values
print("\nOptimization Results:")
print(f"Best learning rate: {10**result.x[0]:.6f}")
print(f"Best layer size: {int(result.x[1])}")
print(f"Best num layers: {int(result.x[2])}")
print(f"Best activation: {result.x[3]}") # String value!
print(f"Best test MSE: {result.fun:.4f}")

```

```

lr=0.000467, l1=17, layers=1, activation=Tanh
lr=0.000220, l1=94, layers=4, activation=ReLU
lr=0.000690, l1=111, layers=2, activation=Sigmoid

```

```

/var/folders/dw/pvtj6mt9l2nd0hftczqb0k00000gn/T/ipykernel_37163/2535631531.py:63: UserWarning: Factor
agnostic) factor metrics. Factor integer codes will be treated as ordinal numbers, which may mislead
aware Kriging surrogate: Kriging(metric_factorial='hamming').

```

```

optimizer = SpotOptim(

```

```

lr=0.006473, l1=49, layers=0, activation=Sigmoid
lr=0.000123, l1=125, layers=1, activation=LeakyReLU
lr=0.001590, l1=98, layers=4, activation=ReLU
lr=0.001419, l1=65, layers=3, activation=LeakyReLU
lr=0.004632, l1=36, layers=2, activation=Tanh
lr=0.003901, l1=53, layers=3, activation=Tanh
lr=0.000334, l1=77, layers=1, activation=Sigmoid
lr=0.001596, l1=77, layers=1, activation=LeakyReLU
lr=0.009535, l1=88, layers=2, activation=Sigmoid
lr=0.006323, l1=115, layers=2, activation=Sigmoid
lr=0.000309, l1=74, layers=1, activation=ReLU
lr=0.000590, l1=42, layers=2, activation=Sigmoid
lr=0.000570, l1=43, layers=0, activation=LeakyReLU
lr=0.000493, l1=114, layers=3, activation=Sigmoid
lr=0.001018, l1=33, layers=2, activation=Sigmoid
lr=0.000179, l1=78, layers=3, activation=Sigmoid
lr=0.000315, l1=124, layers=1, activation=Tanh
lr=0.009564, l1=88, layers=2, activation=Sigmoid
lr=0.006655, l1=73, layers=2, activation=Tanh
lr=0.005195, l1=103, layers=4, activation=ReLU
lr=0.001384, l1=110, layers=4, activation=Tanh
lr=0.000421, l1=107, layers=2, activation=LeakyReLU
lr=0.000708, l1=28, layers=3, activation=Sigmoid
lr=0.005019, l1=66, layers=1, activation=LeakyReLU

```

12. Factor Variables for Categorical Hyperparameters

```
lr=0.001990, l1=85, layers=3, activation=Tanh
lr=0.001617, l1=18, layers=0, activation=LeakyReLU
lr=0.000643, l1=44, layers=2, activation=LeakyReLU
lr=0.000149, l1=49, layers=3, activation=Sigmoid
lr=0.009535, l1=88, layers=2, activation=Sigmoid
lr=0.000776, l1=63, layers=4, activation=Sigmoid
lr=0.002738, l1=112, layers=3, activation=Tanh
lr=0.000102, l1=52, layers=2, activation=Sigmoid
lr=0.001336, l1=117, layers=4, activation=Tanh
lr=0.003807, l1=72, layers=0, activation=ReLU
lr=0.003646, l1=88, layers=1, activation=Tanh
lr=0.000738, l1=50, layers=1, activation=ReLU
lr=0.009536, l1=88, layers=2, activation=Sigmoid
lr=0.000317, l1=26, layers=1, activation=Tanh
lr=0.004723, l1=36, layers=2, activation=Tanh
lr=0.008777, l1=49, layers=4, activation=Tanh
lr=0.001696, l1=53, layers=1, activation=Tanh
lr=0.008277, l1=57, layers=2, activation=ReLU
lr=0.004042, l1=18, layers=0, activation=Tanh
lr=0.001416, l1=22, layers=3, activation=Sigmoid
lr=0.000119, l1=118, layers=3, activation=Sigmoid
lr=0.003565, l1=116, layers=3, activation=ReLU
lr=0.000190, l1=81, layers=4, activation=Tanh
```

Optimization Results:

Best learning rate: 0.009535

Best layer size: 88

Best num layers: 2

Best activation: Sigmoid

Best test MSE: 26403.0527

12.4. Multiple Factor Variables

12.4.1. Optimizing Both Activation and Optimizer

```
from spotoptim import SpotOptim
from spotoptim.data import get_diabetes_data loaders
from spotoptim.nn.linear_regressor import LinearRegressor
import torch.nn as nn
import numpy as np

def optimize_activation_and_optimizer(X):
```



```

"""Optimize both activation function and optimizer choice."""
results = []

for params in X:
    activation = params[0]      # Factor variable 1
    optimizer_name = params[1] # Factor variable 2
    lr = 10 ** params[2]       # Continuous variable

    train_loader, test_loader, _ = get_diabetes_data loaders()

    model = LinearRegressor(
        input_dim=10,
        output_dim=1,
        l1=64,
        num_hidden_layers=2,
        activation=activation
    )

    # Use the optimizer string
    optimizer = model.get_optimizer(optimizer_name, lr=lr)
    criterion = nn.MSELoss()

    # Train
    for epoch in range(30):
        model.train()
        for batch_X, batch_y in train_loader:
            predictions = model(batch_X)
            loss = criterion(predictions, batch_y)
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

    # Evaluate
    model.eval()
    test_loss = 0.0
    with torch.no_grad():
        for batch_X, batch_y in test_loader:
            predictions = model(batch_X)
            test_loss += criterion(predictions, batch_y).item()

    results.append(test_loss / len(test_loader))

return np.array(results) # Return numpy array

```

12. Factor Variables for Categorical Hyperparameters

```
# Two factor variables + one continuous
opt = SpotOptim(
    fun=optimize_activation_and_optimizer,
    bounds=[
        ("ReLU", "Tanh", "Sigmoid", "LeakyReLU"),    # Activation
        ("Adam", "SGD", "RMSprop", "AdamW"),          # Optimizer
        (-4, -2)                                       # log10(lr)
    ],
    var_type=["factor", "factor", "float"],
    max_iter=40
)

result = opt.optimize()
print(f"Best activation: {result.x[0]}")
print(f"Best optimizer: {result.x[1]}")
print(f"Best learning rate: {10**result.x[2]:.6f}")
```

```
/var/folders/dw/pvtj6mt91znd0hftcztqb0k00000gn/T/ipykernel_37163/2475134545.py:53: UserWarning: Using non-agnostic factor metrics. Factor integer codes will be treated as ordinal numbers, which may not be appropriate for all metrics. Consider using the Kriging surrogate: Kriging(metric_factorial='hamming').
```

```
    opt = SpotOptim(

Best activation: ReLU
Best optimizer: SGD
Best learning rate: 0.007895
```

12.5. Advanced Usage

12.5.1. Custom Categorical Choices

Factor variables work with any string values, not just activation functions:

```
from spotoptim import SpotOptim
import numpy as np

def train_model_with_config(dropout_policy, batch_norm, weight_init):
    """Simulate model training with different configurations."""
    # In real use, this would train an actual model
    # Here we return synthetic scores for demonstration
    base_score = 3000.0
```

```

# Dropout impact
dropout_scores = {"none": 200, "light": 0, "heavy": 100}
# Batch norm impact
bn_scores = {"before": -50, "after": 0, "none": 150}
# Weight init impact
init_scores = {"xavier": 0, "kaiming": -30, "normal": 100}

score = (base_score +
         dropout_scores.get(dropout_policy, 0) +
         bn_scores.get(batch_norm, 0) +
         init_scores.get(weight_init, 0) +
         np.random.normal(0, 50))

return score

def train_with_config(X):
    """Objective function with various categorical choices."""
    results = []

    for params in X:
        dropout_policy = params[0] # "none", "light", "heavy"
        batch_norm = params[1]      # "before", "after", "none"
        weight_init = params[2]     # "xavier", "kaiming", "normal"

        # Use these strings to configure your model
        score = train_model_with_config(
            dropout_policy=dropout_policy,
            batch_norm=batch_norm,
            weight_init=weight_init
        )
        results.append(score)

    return np.array(results) # Return numpy array

optimizer = SpotOptim(
    fun=train_with_config,
    bounds=[
        ("none", "light", "heavy"),          # Dropout policy
        ("before", "after", "none"),          # Batch norm position
        ("xavier", "kaiming", "normal")       # Weight initialization
    ],
    var_type=["factor", "factor", "factor"],
    max_iter=25,
    seed=42

```

12. Factor Variables for Categorical Hyperparameters

```
)

result = optimizer.optimize()
print("Best configuration:")
print(f"  Dropout: {result.x[0]}")
print(f"  Batch norm: {result.x[1]}")
print(f"  Weight init: {result.x[2]}")
print(f"  Score: {result.fun:.4f}")
```

```
/var/folders/dw/pvtj6mt91znd0hftcztqb0k00000gn/T/ipykernel_37163/2135006710.py:44: Use
agnostic) factor metrics. Factor integer codes will be treated as ordinal numbers, wh
aware Kriging surrogate: Kriging(metric_factorial='hamming').
    optimizer = SpotOptim(
```

```
Best configuration:
  Dropout: light
  Batch norm: before
  Weight init: xavier
  Score: 2854.3360
```

12.5.2. Viewing All Evaluated Configurations

```
import torch
import torch.nn as nn
from spotoptim import SpotOptim
from spotoptim.data import get_diabetes_dataloaders
from spotoptim.nn.linear_regressor import LinearRegressor
import numpy as np

def train_and_evaluate(X):
    """Train models with different activation functions."""
    results = []

    for params in X:
        l1 = int(params[0])          # Integer: layer size
        activation = params[1]       # String: activation function

        # Load data
        train_loader, test_loader, _ = get_diabetes_dataloaders()

        # Create model with the activation function
```

```

model = LinearRegressor(
    input_dim=10,
    output_dim=1,
    l1=l1,
    num_hidden_layers=2,
    activation=activation # Pass string directly!
)

# Train model
optimizer = model.get_optimizer("Adam", lr=0.01)
criterion = nn.MSELoss()

for epoch in range(50):
    model.train()
    for batch_X, batch_y in train_loader:
        predictions = model(batch_X)
        loss = criterion(predictions, batch_y)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

# Evaluate
model.eval()
test_loss = 0.0
with torch.no_grad():
    for batch_X, batch_y in test_loader:
        predictions = model(batch_X)
        test_loss += criterion(predictions, batch_y).item()

avg_loss = test_loss / len(test_loader)
results.append(avg_loss)

return np.array(results)

optimizer = SpotOptim(
    fun=train_and_evaluate,
    bounds=[
        (16, 128), # Layer size
        ("ReLU", "Sigmoid", "Tanh", "LeakyReLU") # Activation
    ],
    var_type=["int", "factor"], # IMPORTANT: Specify variable types!
    max_iter=30,
    seed=42
)

```

12. Factor Variables for Categorical Hyperparameters

```
result = optimizer.optimize()

# Access all evaluated configurations
print("\nAll evaluated configurations:")
print("Layer Size | Activation | Test MSE")
print("-" * 42)
for i in range(min(10, len(result.X))): # Show first 10
    l1 = int(result.X[i, 0])
    activation = result.X[i, 1] # String value!
    loss = result.y[i]
    print(f"{l1:10d} | {activation:10s} | {loss:.4f}")

# Find top 5 configurations
sorted_indices = result.y.argsort()[:5]
print("\nTop 5 configurations:")
for idx in sorted_indices:
    print(f"l1={int(result.X[idx, 0]):3d}, "
          f"activation={result.X[idx, 1]:10s}, "
          f"MSE={result.y[idx]:.4f}")
```

```
/var/folders/dw/pvtj6mt91znd0hftcztqb0k00000gn/T/ipykernel_37163/2283186313.py:54: UserWarning: Using non-integer codes for categorical factor metrics. Factor integer codes will be treated as ordinal numbers, which may lead to incorrect Kriging surrogate: Kriging(metric_factorial='hamming').
optimizer = SpotOptim(
```

```
All evaluated configurations:
Layer Size | Activation | Test MSE
```

```
-----
      95 | Tanh      | 26554.0299
      73 | ReLU      | 26522.1647
     106 | Tanh      | 26513.4258
      61 | Tanh      | 26607.4414
     126 | LeakyReLU | 26389.1042
      43 | Sigmoid   | 26641.9941
      92 | Sigmoid   | 26559.3444
      56 | Sigmoid   | 26421.8913
      17 | LeakyReLU | 26653.0911
      30 | ReLU      | 26574.7884
```

```
Top 5 configurations:
l1=119, activation=Sigmoid    , MSE=26214.9948
l1=125, activation=Sigmoid    , MSE=26293.1582
```

```
l1=126, activation=Sigmoid    , MSE=26314.9245
l1=126, activation=LeakyReLU  , MSE=26389.1042
l1=118, activation=ReLU       , MSE=26409.5560
```

12.6. How It Works

12.6.1. Internal Mechanism

SpotOptim handles factor variables through automatic conversion:

1. **Initialization:** String tuples in bounds are detected

```
bounds = [("ReLU", "Sigmoid", "Tanh")]
# Internally mapped to: {0: "ReLU", 1: "Sigmoid", 2: "Tanh"}
# Bounds become: [(0, 2)]
```

2. **Sampling:** Initial design samples from $[0, n_levels-1]$ and rounds to integers

```
# Samples might be: [0.3, 1.8, 2.1]
# After rounding: [0, 2, 2]
```

3. **Evaluation:** Before calling objective function, integers $\rightarrow$ strings

```
# [0, 2, 2]  $\rightarrow$  ["ReLU", "Tanh", "Tanh"]
# Objective function receives strings
```

4. **Optimization:** Surrogate model works with integers $[0, n_levels-1]$

5. **Results:** Final results mapped back to strings

```
result.x[0] # Returns "ReLU", not 0
result.X    # All rows contain strings for factor variables
```

```
array([[95.0, 'Tanh'],
       [73.0, 'ReLU'],
       [106.0, 'Tanh'],
       [61.0, 'Tanh'],
       [126.0, 'LeakyReLU'],
       [43.0, 'Sigmoid'],
       [92.0, 'Sigmoid'],
       [56.0, 'Sigmoid'],
       [17.0, 'LeakyReLU'],
       [30.0, 'ReLU'],
       [124.0, 'Tanh'],
       [128.0, 'LeakyReLU'],
       [125.0, 'LeakyReLU'],
```

12. Factor Variables for Categorical Hyperparameters

```
[125.0, 'Sigmoid'],
[123.0, 'Tanh'],
[127.0, 'ReLU'],
[55.0, 'Sigmoid'],
[127.0, 'Tanh'],
[119.0, 'Sigmoid'],
[126.0, 'Sigmoid'],
[126.0, 'ReLU'],
[118.0, 'Sigmoid'],
[118.0, 'ReLU'],
[125.0, 'Tanh'],
[120.0, 'Tanh'],
[124.0, 'ReLU'],
[127.0, 'Sigmoid'],
[56.0, 'ReLU'],
[52.0, 'LeakyReLU'],
[104.0, 'ReLU']], dtype=object)
```

12.6.2. Variable Type Auto-Detection

If you don't specify `var_type`, SpotOptim automatically detects factor variables:

```
# Example 1: Explicit var_type (recommended)
# This shows the syntax - replace my_function with your actual function

# optimizer = SpotOptim(
#     fun=my_function,
#     bounds=[(-4, -2), ("ReLU", "Tanh")],
#     var_type=["float", "factor"] # Explicit
# )

# Example 2: Auto-detection (works but less explicit)
# optimizer = SpotOptim(
#     fun=my_function,
#     bounds=[(-4, -2), ("ReLU", "Tanh")]
#     # var_type automatically set to ["float", "factor"]
# )

# Here's a working example:
from spotoptim import SpotOptim
import numpy as np

def demo_function(X):
    results = []
```


12.7. Complete Example: Full Workflow

```
for params in X:
    lr = 10 ** params[0] # Continuous parameter
    activation = params[1] # Factor parameter
    score = 3000 + lr * 100 + {"ReLU": 0, "Tanh": 50}.get(activation, 100)
    results.append(score + np.random.normal(0, 10))
return np.array(results)

# With explicit var_type (recommended)
optimizer = SpotOptim(
    fun=demo_function,
    bounds=[(-4, -2), ("ReLU", "Tanh")],
    var_type=["float", "factor"], # Explicit is clearer
    max_iter=10,
    seed=42
)

result = optimizer.optimize()
print(f"Best lr: {10**result.x[0]:.6f}, Best activation: {result.x[1]}")
```

Best lr: 0.000305, Best activation: ReLU

/var/folders/dw/pvtj6mt9lzd0hftcztqb0k00000gn/T/ipykernel_37163/3134039290.py:31: UserWarning: Factor metrics. Factor integer codes will be treated as ordinal numbers, which may mislead aware Kriging surrogate: Kriging(metric_factorial='hamming').

```
optimizer = SpotOptim(
```

12.7. Complete Example: Full Workflow

```
"""
Complete example: Neural network hyperparameter optimization with factor variables.
"""

import numpy as np
import torch
import torch.nn as nn
from spotoptim import SpotOptim
from spotoptim.data import get_diabetes_dataloaders
from spotoptim.nn.linear_regressor import LinearRegressor

def objective_function(X):
    """Train and evaluate models with given hyperparameters."""
```

12. Factor Variables for Categorical Hyperparameters

```
results = []

for params in X:
    # Extract hyperparameters
    log_lr = params[0]
    l1 = int(params[1])
    num_layers = int(params[2])
    activation = params[3] # String!

    lr = 10 ** log_lr

    print(f"Testing: lr={lr:.6f}, l1={l1}, layers={num_layers}, "
          f"activation={activation}")

    # Load data
    train_loader, test_loader, _ = get_diabetes_dataloaders(
        test_size=0.2,
        batch_size=32,
        random_state=42
    )

    # Create and train model
    model = LinearRegressor(
        input_dim=10,
        output_dim=1,
        l1=l1,
        num_hidden_layers=num_layers,
        activation=activation
    )

    optimizer = model.get_optimizer("Adam", lr=lr)
    criterion = nn.MSELoss()

    # Training loop
    num_epochs = 30
    for epoch in range(num_epochs):
        model.train()
        for batch_X, batch_y in train_loader:
            predictions = model(batch_X)
            loss = criterion(predictions, batch_y)
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()
```

```

# Evaluation
model.eval()
test_loss = 0.0
with torch.no_grad():
    for batch_X, batch_y in test_loader:
        predictions = model(batch_X)
        loss = criterion(predictions, batch_y)
        test_loss += loss.item()

    avg_test_loss = test_loss / len(test_loader)
    results.append(avg_test_loss)
    print(f" → Test MSE: {avg_test_loss:.4f}")

return np.array(results)

def main():
    print("=" * 80)
    print("Neural Network Hyperparameter Optimization with Factor Variables")
    print("=" * 80)

    # Define optimization problem
    optimizer = SpotOptim(
        fun=objective_function,
        bounds=[
            (-4, -2), # log10(learning_rate)
            (16, 128), # l1 (neurons)
            (0, 4), # num_hidden_layers
            ("ReLU", "Sigmoid", "Tanh", "LeakyReLU") # activation (factor!)
        ],
        var_type=["float", "int", "int", "factor"],
        max_iter=50,
        seed=42
    )

    # Run optimization
    print("\nStarting optimization...")
    result = optimizer.optimize()

    # Display results
    print("\n" + "=" * 80)
    print("OPTIMIZATION RESULTS")
    print("=" * 80)
    print(f"Best learning rate: {10**result.x[0]:.6f}")

```

12. Factor Variables for Categorical Hyperparameters

```
print(f"Best layer size (l1): {int(result.x[1])}")
print(f"Best num hidden layers: {int(result.x[2])}")
print(f"Best activation function: {result.x[3]}") # String value!
print(f"Best test MSE: {result.fun:.4f}")

# Show top 5 configurations
print("\n" + "=" * 80)
print("TOP 5 CONFIGURATIONS")
print("=" * 80)
sorted_indices = result.y.argsort()[:5]
print(f"{'Rank':<6} {'LR':<12} {'L1':<6} {'Layers':<8} "
      f"{'Activation':<12} {'MSE':<10}")
print("-" * 80)
for rank, idx in enumerate(sorted_indices, 1):
    lr = 10 ** result.X[idx, 0]
    l1 = int(result.X[idx, 1])
    layers = int(result.X[idx, 2])
    activation = result.X[idx, 3]
    mse = result.y[idx]
    print(f"{'rank':<6} {'lr':<12.6f} {'l1':<6} {'layers':<8} "
          f"{'activation':<12} {'mse':<10.4f}")

# Train final model with best configuration
print("\n" + "=" * 80)
print("TRAINING FINAL MODEL")
print("=" * 80)

best_lr = 10 ** result.x[0]
best_l1 = int(result.x[1])
best_layers = int(result.x[2])
best_activation = result.x[3]

print(f"Configuration: lr={best_lr:.6f}, l1={best_l1}, "
      f"layers={best_layers}, activation={best_activation}")

train_loader, test_loader, _ = get_diabetes_data loaders(
    test_size=0.2,
    batch_size=32,
    random_state=42
)

final_model = LinearRegressor(
    input_dim=10,
    output_dim=1,
```

```

        l1=best_l1,
        num_hidden_layers=best_layers,
        activation=best_activation
    )

    optimizer_final = final_model.get_optimizer("Adam", lr=best_lr)
    criterion = nn.MSELoss()

    # Extended training
    num_epochs = 100
    print(f"\nTraining for {num_epochs} epochs...")
    for epoch in range(num_epochs):
        final_model.train()
        train_loss = 0.0
        for batch_X, batch_y in train_loader:
            predictions = final_model(batch_X)
            loss = criterion(predictions, batch_y)
            optimizer_final.zero_grad()
            loss.backward()
            optimizer_final.step()
            train_loss += loss.item()

        if (epoch + 1) % 20 == 0:
            avg_train_loss = train_loss / len(train_loader)
            print(f"Epoch {epoch+1}/{num_epochs}: Train MSE = {avg_train_loss:.4f}")

    # Final evaluation
    final_model.eval()
    final_test_loss = 0.0
    with torch.no_grad():
        for batch_X, batch_y in test_loader:
            predictions = final_model(batch_X)
            final_test_loss += criterion(predictions, batch_y).item()

    final_avg_loss = final_test_loss / len(test_loader)
    print(f"\nFinal Test MSE: {final_avg_loss:.4f}")
    print("=" * 80)

if __name__ == "__main__":
    main()

```

=====

Neural Network Hyperparameter Optimization with Factor Variables

12. Factor Variables for Categorical Hyperparameters

```
=====

Starting optimization...
Testing: lr=0.000414, l1=62, layers=1, activation=Tanh
  → Test MSE: 26660.6276
Testing: lr=0.000256, l1=81, layers=1, activation=Sigmoid
  → Test MSE: 26721.7930
Testing: lr=0.009331, l1=90, layers=4, activation=Tanh

/var/folders/dw/pvtj6mt91znd0hftcztqb0k00000gn/T/ipykernel_37163/156832938.py:80: UserWarning: Using a
agnostic) factor metrics. Factor integer codes will be treated as ordinal numbers, which may be
aware Kriging surrogate: Kriging(metric_factorial='hamming').
  optimizer = SpotOptim(

  → Test MSE: 26526.8822
Testing: lr=0.001413, l1=106, layers=2, activation=Sigmoid
  → Test MSE: 26668.0417
Testing: lr=0.000652, l1=27, layers=3, activation=Tanh
  → Test MSE: 26651.3522
Testing: lr=0.003696, l1=45, layers=3, activation=Tanh
  → Test MSE: 26621.2611
Testing: lr=0.002433, l1=32, layers=2, activation=ReLU
  → Test MSE: 26621.0866
Testing: lr=0.000114, l1=52, layers=3, activation=ReLU
  → Test MSE: 26598.5514
Testing: lr=0.006093, l1=99, layers=0, activation=Sigmoid
  → Test MSE: 26724.3027
Testing: lr=0.000249, l1=118, layers=1, activation=LeakyReLU
  → Test MSE: 26637.8073
Testing: lr=0.009335, l1=90, layers=4, activation=Tanh
  → Test MSE: 26554.9277
Testing: lr=0.000265, l1=35, layers=3, activation=Sigmoid
  → Test MSE: 26611.9941
Testing: lr=0.000159, l1=79, layers=3, activation=LeakyReLU
  → Test MSE: 26629.7676
Testing: lr=0.005810, l1=77, layers=3, activation=ReLU
  → Test MSE: 26565.9492
Testing: lr=0.009330, l1=90, layers=4, activation=Tanh
  → Test MSE: 26538.5618
Testing: lr=0.002670, l1=62, layers=3, activation=Sigmoid
  → Test MSE: 26507.8125
Testing: lr=0.001138, l1=78, layers=3, activation=Sigmoid
  → Test MSE: 26576.3301
Testing: lr=0.000917, l1=48, layers=2, activation=Sigmoid
```

12.7. Complete Example: Full Workflow

```
→ Test MSE: 26571.5521
Testing: lr=0.000317, l1=19, layers=2, activation=Sigmoid
→ Test MSE: 26660.5202
Testing: lr=0.007670, l1=32, layers=2, activation=Sigmoid
→ Test MSE: 26596.0104
Testing: lr=0.001611, l1=112, layers=3, activation=ReLU
→ Test MSE: 26635.1263
Testing: lr=0.000209, l1=69, layers=1, activation=LeakyReLU
→ Test MSE: 26624.6940
Testing: lr=0.001362, l1=67, layers=2, activation=ReLU
→ Test MSE: 26623.6823
Testing: lr=0.000628, l1=51, layers=2, activation=LeakyReLU
→ Test MSE: 26596.9701
Testing: lr=0.000465, l1=30, layers=3, activation=Sigmoid
→ Test MSE: 26607.1113
Testing: lr=0.001686, l1=66, layers=2, activation=Tanh
→ Test MSE: 26610.1615
Testing: lr=0.000454, l1=55, layers=3, activation=Tanh
→ Test MSE: 26595.5046
Testing: lr=0.000603, l1=115, layers=3, activation=Sigmoid
→ Test MSE: 26403.9635
Testing: lr=0.009157, l1=103, layers=3, activation=Tanh
→ Test MSE: 26545.4720
Testing: lr=0.001266, l1=41, layers=3, activation=Sigmoid
→ Test MSE: 26594.5371
Testing: lr=0.008903, l1=43, layers=1, activation=Sigmoid
→ Test MSE: 26720.3561
Testing: lr=0.002573, l1=117, layers=3, activation=LeakyReLU
→ Test MSE: 26608.9674
Testing: lr=0.003325, l1=20, layers=0, activation=Tanh
→ Test MSE: 26618.4727
Testing: lr=0.000176, l1=76, layers=4, activation=LeakyReLU
→ Test MSE: 26604.9805
Testing: lr=0.002366, l1=52, layers=2, activation=Tanh
→ Test MSE: 26640.5104
Testing: lr=0.001022, l1=60, layers=2, activation=LeakyReLU
→ Test MSE: 26619.7891
Testing: lr=0.000805, l1=50, layers=1, activation=LeakyReLU
→ Test MSE: 26666.4902
Testing: lr=0.001737, l1=38, layers=4, activation=Sigmoid
→ Test MSE: 26717.5605
Testing: lr=0.000209, l1=88, layers=2, activation=ReLU
→ Test MSE: 26608.3913
Testing: lr=0.000115, l1=79, layers=3, activation=Tanh
→ Test MSE: 26611.1159
```

12. Factor Variables for Categorical Hyperparameters

```
Testing: lr=0.001429, l1=74, layers=0, activation=Tanh
  → Test MSE: 26493.1680
Testing: lr=0.000480, l1=42, layers=1, activation=Tanh
  → Test MSE: 26672.6230
Testing: lr=0.005419, l1=29, layers=3, activation=Sigmoid
  → Test MSE: 26596.9805
Testing: lr=0.001159, l1=62, layers=1, activation=Sigmoid
  → Test MSE: 26695.5312
Testing: lr=0.002008, l1=102, layers=3, activation=Sigmoid
  → Test MSE: 26715.1042
Testing: lr=0.001383, l1=80, layers=3, activation=Tanh
  → Test MSE: 26622.8477
Testing: lr=0.000977, l1=83, layers=4, activation=Sigmoid
  → Test MSE: 26617.5130
Testing: lr=0.005980, l1=82, layers=2, activation=LeakyReLU
  → Test MSE: 26594.1654
Testing: lr=0.004110, l1=110, layers=2, activation=Tanh
  → Test MSE: 26629.9134
Testing: lr=0.002603, l1=57, layers=2, activation=ReLU
  → Test MSE: 26650.2799
```

OPTIMIZATION RESULTS

```
Best learning rate: 0.000603
Best layer size (l1): 115
Best num hidden layers: 3
Best activation function: Sigmoid
Best test MSE: 26403.9635
```

TOP 5 CONFIGURATIONS

Rank	LR	L1	Layers	Activation	MSE
1	0.000603	115	3	Sigmoid	26403.9635
2	0.001429	74	0	Tanh	26493.1680
3	0.002670	62	3	Sigmoid	26507.8125
4	0.009331	90	4	Tanh	26526.8822
5	0.009330	90	4	Tanh	26538.5618

TRAINING FINAL MODEL

```
Configuration: lr=0.000603, l1=115, layers=3, activation=Sigmoid
```



```

Training for 100 epochs...
Epoch 20/100: Train MSE = 27529.2773
Epoch 40/100: Train MSE = 27848.2657
Epoch 60/100: Train MSE = 27897.5425
Epoch 80/100: Train MSE = 28427.8789
Epoch 100/100: Train MSE = 28886.2425

```

```

Final Test MSE: 26578.0579
=====

```

12.8. Best Practices

12.8.1. Do's

☒ Use descriptive string values

```

bounds=[("xavier_uniform", "kaiming_normal", "orthogonal")]

```

☒ Explicitly specify `var_type` for clarity

```

var_type=["float", "int", "factor"]

```

☒ Access results as strings

```

# Example: Accessing factor variable results as strings
# (This assumes you've run an optimization with activation as a factor variable)

# If you have a result from the previous examples:
# best_activation = result.x[3] # For 4-parameter optimization
# Or for simpler cases:
# best_activation = result.x[0] # For single-parameter optimization

# Example with inline optimization:
from spotoptim import SpotOptim
import numpy as np

def quick_test(X):
    results = []
    for params in X:
        activation = params[0]
        score = {"ReLU": 3500, "Tanh": 3600}.get(activation, 4000)

```

12. Factor Variables for Categorical Hyperparameters

```
        results.append(score + np.random.normal(0, 50))
    return np.array(results)

opt = SpotOptim(
    fun=quick_test,
    bounds=[("ReLU", "Tanh")],
    var_type=["factor"],
    max_iter=10,
    seed=42
)
result = opt.optimize()

# Access as string - this is the correct way
best_activation = result.x[0] # String value like "ReLU"
print(f"Best activation: {best_activation} (type: {type(best_activation).__name__})")

# You can use it directly in your model
# model = LinearRegressor(activation=best_activation)
```

```
/var/folders/dw/pvtj6mt91znd0hftcztqb0k00000gn/T/ipykernel_37163/48594712.py:21: UserWarning: Ignoring (non-agnostic) factor metrics. Factor integer codes will be treated as ordinal numbers, which may not be appropriate for Kriging surrogate: Kriging(metric_factorial='hamming').
  opt = SpotOptim(
```

Best activation: ReLU (type: str)

⚠ Mix factor variables with numeric/integer variables

```
bounds=[(-4, -2), (16, 128), ("ReLU", "Tanh")]
var_type=["float", "int", "factor"]
```

12.8.2. Don'ts

⚠ Don't use integers in factor bounds

```
# Wrong: Use strings, not integers
bounds=[(0, 1, 2)] # Wrong!
bounds=[("ReLU", "Sigmoid", "Tanh")] # Correct!
```

⚠ Don't expect integers in objective function

```
def objective(X):
    activation = X[0][2]
    # activation is a string, not an integer!
    # Don't do: if activation == 0: # Wrong!
    # Do: if activation == "ReLU": # Correct!
```

❌ Don't manually convert factor variables

```
# SpotOptim handles conversion automatically
# Don't do manual mapping in your objective function
```

❌ Don't use empty tuples

```
# Wrong: Empty tuple
bounds=[()]

# Correct: At least one string
bounds=[("ReLU",)] # Single choice (will be treated as fixed)
```

12.9. Troubleshooting

12.9.1. Common Issues

Issue: Objective function receives integers instead of strings

Solution: Ensure you're using the latest version of SpotOptim with factor variable support. Factor variables are automatically converted before calling the objective function.

Issue: ValueError: could not convert string to float

Solution: This occurs if there's a version mismatch. Update SpotOptim to ensure the object array conversion is implemented correctly.

Issue: Results show integers instead of strings

Solution: Check that you're accessing `result.x` (mapped values) instead of internal arrays. The result object automatically maps factor variables to their original strings.

Issue: Single-level factor variables cause dimension reduction

Behavior: If a factor variable has only one choice, e.g., ("ReLU",), SpotOptim treats it as a fixed dimension and may reduce the dimensionality. This is expected behavior.

Solution: Use at least two choices for optimization, or remove single-choice dimensions from bounds.

12.10. Summary

Factor variables in SpotOptim enable:

- **Categorical optimization:** Optimize over discrete string choices
- **Automatic conversion:** Seamless integer↔string mapping
- **Neural network hyperparameters:** Optimize activation functions, optimizers, etc.
- **Mixed variable types:** Combine with continuous and integer variables
- **Clean interface:** Objective functions work with strings directly
- **String results:** Final results contain original string values

Factor variables make categorical hyperparameter optimization as easy as continuous optimization!

12.11. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

13. Variable Transformations for Search Space Scaling

SpotOptim supports automatic variable transformations to improve optimization in scaled search spaces. Instead of manually handling transformations (e.g., log-scale for learning rates), you can specify transformations via the `var_trans` parameter, and SpotOptim handles everything internally.

13.1. Overview

What are Variable Transformations?

Variable transformations allow you to specify how search space dimensions should be scaled during optimization:

1. **Original scale** (user interface): Input bounds, output results, plots
2. **Transformed scale** (internal): Surrogate modeling, acquisition optimization
3. **Automatic conversion**: SpotOptim handles all transformations transparently

Module: `spotoptim.SpotOptim`

Key Features:

- Define transformations via `var_trans` parameter: `["log10", "sqrt", None, ...]`
- Optimization occurs in transformed space (better for surrogate models)
- All external interfaces use original scale (intuitive for users)
- Supported transformations: `log10`, `log/ln`, `sqrt`, `exp`, `square`, `cube`, `inv/reciprocal`
- Mix transformed and non-transformed variables

13.2. Why Use Transformations?

13.2.1. Problem: Poorly Scaled Search Spaces

Some hyperparameters span multiple orders of magnitude:

13. Variable Transformations for Search Space Scaling

- **Learning rates:** 0.0001 to 1.0 (4 orders of magnitude)
- **Regularization:** 0.001 to 100 (5 orders of magnitude)
- **Network sizes:** 10 to 1000 neurons

Direct optimization in these spaces is inefficient because:

1. Surrogate models struggle with extreme scales
2. Uniform sampling wastes evaluations in unimportant regions
3. Acquisition functions behave poorly with skewed distributions

13.2.2. Solution: Logarithmic and Other Transformations

Transform the space for optimization while maintaining user-friendly interfaces:

```
# Without transformations (manual approach)
bounds = [(-4, 0)] # log10(lr): awkward for users
lr = 10 ** params[0] # Manual transformation in objective

# With transformations (automatic)
bounds = [(0.0001, 1.0)] # lr in natural scale
var_trans = ["log10"] # SpotOptim handles transformation
lr = params[0] # Already in original scale!
```

13.3. Quick Start

13.3.1. Basic Log-Scale Transformation

```
from spotoptim import SpotOptim
import numpy as np

def objective_function(X):
    """Objective receives parameters in ORIGINAL scale."""
    results = []
    for params in X:
        lr = params[0] # Already in [0.001, 0.1] - original scale!
        alpha = params[1] # Already in [0.01, 1.0] - original scale!

        # Simulate model training
        score = (lr - 0.01)**2 + (alpha - 0.1)**2 + np.random.normal(0, 0.01)
        results.append(score)
```

```

    return np.array(results)

# Create optimizer with transformations
optimizer = SpotOptim(
    fun=objective_function,
    bounds=[
        (0.001, 0.1),    # learning rate (original scale)
        (0.01, 1.0)      # alpha (original scale)
    ],
    var_trans=["log10", "log10"], # Both use log10 transformation
    var_name=["lr", "alpha"],
    max_iter=20,
    seed=42
)

# Run optimization
result = optimizer.optimize()

print(f"Best lr: {result.x[0]:.6f}")    # In original scale
print(f"Best alpha: {result.x[1]:.6f}") # In original scale
print(f"Best score: {result.fun:.6f}")

```

```

Best lr: 0.006994
Best alpha: 0.125349
Best score: -0.018481

```

13.4. Supported Transformations

SpotOptim supports the following transformations:

Transformation	Forward ($x \rightarrow t$)	Inverse ($t \rightarrow x$)	Use Case
"log10"	$t = \log_{10}(x)$	$x = 10^t$	Learning rates, regularization
"log" or "ln"	$t = \ln(x)$	$x = e^t$	Natural exponential scales
"sqrt"	$t = \sqrt{x}$	$x = t^2$	Moderate scaling
"exp"	$t = e^x$	$x = \ln(t)$	Inverse of natural log

13. Variable Transformations for Search Space Scaling

Transformation	Forward ($x \rightarrow t$)	Inverse ($t \rightarrow x$)	Use Case
"square"	$t = x^2$	$x = \sqrt{t}$	Inverse of sqrt
"cube"	$t = x^3$	$x = \sqrt[3]{t}$	Strong scaling
"inv" or "reciprocal"	$t = 1/x$	$x = 1/t$	Reciprocal relationships
"pow(base, x)"	$t = \text{base}^x$	$x = \log_{\text{base}}(t)$	Power scaling (e.g. <code>pow(2, x)</code>)
"log(x, base)"	$t = \log_{\text{base}}(x)$	$x = \text{base}^t$	Log scaling with custom base
None or "id"	$t = x$	$x = t$	No transformation

13.4.1. Transformation Guidelines

When to use "log10" or "log":

- Parameters spanning multiple orders of magnitude
- Learning rates: (`1e-5`, `1e-1`) $\rightarrow$ uniform sampling in log space
- Regularization parameters: (`1e-6`, `1e2`)
- Batch sizes, hidden units when range is large

When to use "sqrt":

- Moderate scaling (1-2 orders of magnitude)
- Batch sizes: (`16`, `512`)
- Number of neurons: (`32`, `256`)

When to use "inv" (reciprocal):

- Inverse relationships (e.g., $1/\text{temperature}$)
- When smaller values are more important

Dynamic Transformations:

- "pow(base, x)": Exponential scaling. Example: `transform="pow(2, x)"` with bounds `[2, 5]` leads to internal search on `[4, 32]`.
- "log(x, base)": Logarithmic scaling with custom base. Example: `transform="log(x, 2)"`.

When to use None:

- Parameters with narrow ranges
- Already well-scaled parameters
- Categorical indices (use with `var_type=["factor"]`)

13.5. Detailed Examples

13.5.1. Example 1: Neural Network Hyperparameter Tuning

```
import torch
import torch.nn as nn
from spotoptim import SpotOptim
from spotoptim.data import get_diabetes_dataloaders
from spotoptim.nn.linear_regressor import LinearRegressor
import numpy as np

def train_neural_network(X):
    """Train neural network with hyperparameters in original scale."""
    results = []

    for params in X:
        # All parameters in original scale
        hidden_size = int(params[0]) # [16, 256]
        num_layers = int(params[1]) # [1, 4]
        lr = params[2] # [0.0001, 0.1]
        weight_decay = params[3] # [1e-6, 0.01]

        print(f"Training: hidden={hidden_size}, layers={num_layers}, "
              f"lr={lr:.6f}, wd={weight_decay:.6f}")

        # Load data
        train_loader, test_loader, _ = get_diabetes_dataloaders(batch_size=32)

        # Create model
        model = LinearRegressor(
            input_dim=10,
            output_dim=1,
            l1=hidden_size,
            num_hidden_layers=num_layers,
            activation="ReLU",
            lr=lr
        )
```

13. Variable Transformations for Search Space Scaling

```
# Get optimizer with weight decay
optimizer = torch.optim.Adam(model.parameters(),
                              lr=lr,
                              weight_decay=weight_decay)

# Train
model.train()
for epoch in range(50):
    for batch_x, batch_y in train_loader:
        optimizer.zero_grad()
        outputs = model(batch_x)
        loss = nn.MSELoss()(outputs, batch_y)
        loss.backward()
        optimizer.step()

# Evaluate
model.eval()
total_loss = 0
with torch.no_grad():
    for batch_x, batch_y in test_loader:
        outputs = model(batch_x)
        loss = nn.MSELoss()(outputs, batch_y)
        total_loss += loss.item()

avg_loss = total_loss / len(test_loader)
results.append(avg_loss)

return np.array(results)

# Create optimizer with appropriate transformations
optimizer = SpotOptim(
    fun=train_neural_network,
    bounds=[
        (16, 256),          # hidden_size: moderate range
        (1, 4),              # num_layers: small range
        (0.0001, 0.1),       # lr: 3 orders of magnitude
        (1e-6, 0.01)         # weight_decay: 4 orders of magnitude
    ],
    var_trans=[
        "sqrt",              # sqrt for hidden_size
        None,                 # no transformation for num_layers
        "log10",              # log10 for learning rate
        "log10"               # log10 for weight_decay
    ],
```

```

var_type=["int", "int", "float", "float"],
var_name=["hidden_size", "num_layers", "lr", "weight_decay"],
max_iter=30,
n_initial=10,
seed=42
)

result = optimizer.optimize()

print("\nBest Configuration:")
print(f"  Hidden Size: {int(result.x[0])}")
print(f"  Num Layers: {int(result.x[1])}")
print(f"  Learning Rate: {result.x[2]:.6f}")
print(f"  Weight Decay: {result.x[3]:.8f}")
print(f"  Best Loss: {result.fun:.6f}")

```

```

Training: hidden=59, layers=2, lr=0.000865, wd=0.000189
Training: hidden=42, layers=3, lr=0.000399, wd=0.000008
Training: hidden=250, layers=3, lr=0.067021, wd=0.001348
Training: hidden=119, layers=3, lr=0.004622, wd=0.000093
Training: hidden=79, layers=1, lr=0.015013, wd=0.000395
Training: hidden=180, layers=2, lr=0.028644, wd=0.001817
Training: hidden=152, layers=1, lr=0.002881, wd=0.000002
Training: hidden=19, layers=2, lr=0.010144, wd=0.000004
Training: hidden=216, layers=3, lr=0.000118, wd=0.000022
Training: hidden=41, layers=4, lr=0.000249, wd=0.009479
Training: hidden=28, layers=1, lr=0.000291, wd=0.006011
Training: hidden=79, layers=1, lr=0.013124, wd=0.000362
Training: hidden=67, layers=3, lr=0.002156, wd=0.000117
Training: hidden=21, layers=2, lr=0.018989, wd=0.000001
Training: hidden=88, layers=1, lr=0.100000, wd=0.000842
Training: hidden=77, layers=2, lr=0.044713, wd=0.000630
Training: hidden=59, layers=4, lr=0.000287, wd=0.001440
Training: hidden=152, layers=2, lr=0.005498, wd=0.000228
Training: hidden=16, layers=3, lr=0.011726, wd=0.000003
Training: hidden=61, layers=3, lr=0.001824, wd=0.000030
Training: hidden=141, layers=3, lr=0.003106, wd=0.000137
Training: hidden=77, layers=1, lr=0.055908, wd=0.000566
Training: hidden=86, layers=1, lr=0.024038, wd=0.001470
Training: hidden=159, layers=2, lr=0.019535, wd=0.001514
Training: hidden=182, layers=1, lr=0.009794, wd=0.000211
Training: hidden=83, layers=1, lr=0.029090, wd=0.000823
Training: hidden=101, layers=1, lr=0.030585, wd=0.009004
Training: hidden=162, layers=2, lr=0.015001, wd=0.000030

```

13. Variable Transformations for Search Space Scaling

Training: hidden=174, layers=1, lr=0.000295, wd=0.001219

Training: hidden=188, layers=1, lr=0.100000, wd=0.000239

Best Configuration:

Hidden Size: 86

Num Layers: 1

Learning Rate: 0.024038

Weight Decay: 0.00146977

Best Loss: 2571.753906

13.5.2. Example 2: Physics-Informed Neural Networks (PINNs)

```
from spotoptim import SpotOptim
import numpy as np

def train_pinn(X):
    """Train PINN with hyperparameters in original scale."""
    results = []

    for params in X:
        neurons = int(params[0])      # [16, 128]
        layers = int(params[1])       # [1, 4]
        lr = params[2]                 # [0.1, 10.0]
        alpha = params[3]              # [0.01, 1.0]

        # Simulate PINN training
        # In practice, this would train an actual PINN model
        val_error = 0.1 * (1/lr) + 0.05 * alpha + np.random.normal(0, 0.01)
        results.append(val_error)

    return np.array(results)

# Define optimization with transformations
optimizer = SpotOptim(
    fun=train_pinn,
    bounds=[
        (16, 128),      # neurons
        (1, 4),          # layers
        (0.1, 10.0),     # lr: covers 2 orders of magnitude
        (0.01, 1.0)      # alpha: covers 2 orders of magnitude
    ],
    var_trans=[None, None, "log10", "log10"],
    var_type=["int", "int", "float", "float"],
```

```

var_name=["neurons", "layers", "lr", "alpha"],
max_iter=20,
seed=42
)

result = optimizer.optimize()
optimizer.print_best(result)

```

Best Solution Found:

```

-----
neurons: 38
layers: 1
lr: 3.6651
alpha: 0.1074
Objective Value: 0.0154
Total Evaluations: 20

```

13.5.3. Example 3: Mixing Transformations

```

from spotoptim import SpotOptim
import numpy as np

def complex_objective(X):
    """Objective with multiple transformation types."""
    results = []

    for params in X:
        # All in original scale
        x1 = params[0] # sqrt-transformed: [10, 1000]
        x2 = params[1] # log10-transformed: [0.001, 1.0]
        x3 = params[2] # no transformation: [-5, 5]
        x4 = params[3] # reciprocal: [0.1, 10]

        # Complex function
        result = (
            (x1/100 - 5)**2 +
            (np.log10(x2) + 1.5)**2 +
            x3**2 +
            (1/x4 - 0.2)**2
        )
        results.append(result)

```

13. Variable Transformations for Search Space Scaling

```
        return np.array(results)

optimizer = SpotOptim(
    fun=complex_objective,
    bounds=[
        (10, 1000),      # x1: moderate range
        (0.001, 1.0),    # x2: log scale
        (-5, 5),          # x3: symmetric range
        (0.1, 10)         # x4: for reciprocal
    ],
    var_trans=["sqrt", "log10", None, "inv"],
    var_name=["x1_sqrt", "x2_log", "x3_linear", "x4_inv"],
    max_iter=30,
    seed=42
)

result = optimizer.optimize()
```

13.6. Viewing Transformations in Tables

The transformation type is displayed in the “trans” column of both design and results tables:

13.6.1. Design Table (Before Optimization)

```
from spotoptim import SpotOptim
import numpy as np

optimizer = SpotOptim(
    fun=lambda X: np.sum(X**2, axis=1),
    bounds=[
        (0.001, 1.0),
        (0.01, 10.0),
        (10, 1000),
        (-5, 5)
    ],
    var_trans=["log10", "log", "sqrt", None],
    var_name=["lr", "alpha", "neurons", "bias"],
    max_iter=10
)
```

13.6. Viewing Transformations in Tables

```
# Display design table
print(optimizer.get_design_table())
```

name	type	lower	upper	default	transform
lr	float	0.001	1	0.5005	log10
alpha	float	0.01	10	5.005	log
neurons	float	10	1000	505	sqrt
bias	float	-5	5	0	-

Output:

name	type	lower	upper	default	trans
lr	num	0.0010	1.0000	0.5005	log10
alpha	num	0.0100	10.0000	5.0050	log
neurons	num	10.0000	1000.0000	505.0000	sqrt
bias	num	-5.0000	5.0000	0.0000	-

13.6.2. Results Table (After Optimization)

```
result = optimizer.optimize()

# Display results with transformations
print(optimizer.get_results_table())
```


name	type	default	lower	upper	tuned	transform
lr	float	0.5005	0.001	1	0.0016	log10
alpha	float	5.005	0.01	10	2.0835	log
neurons	float	505	10	1000	21.4937	sqrt
bias	float	0	-5	5	-1.3697	-

Output shows the “trans” column with transformation types, helping you understand which parameters were optimized in which scale.

13.7. Internal Architecture

Understanding how transformations work internally can help debug issues and understand behavior:

13.7.1. Flow Diagram



13.7.2. Key Components

1. **Bounds Transformation** (`_transform_bounds()`):
 - Called during initialization
 - Transforms `_original_lower` and `_original_upper` → `lower` and `upper`
 - Updates `self.bounds` for internal use
2. **Forward Transformation** (`transform_X()`):
 - Converts from original scale to transformed scale
 - Used before surrogate fitting
 - Used when comparing distances
3. **Inverse Transformation** (`inverse_transform_X()`):
 - Converts from transformed scale to original scale
 - Used before function evaluation
 - Used when storing results
4. **Storage**:
 - `self.X_` stores in **original scale**

- `self.best_x_` stores in **original scale**
- All external-facing data in original scale

13.8. Best Practices

13.8.1. 1. Choose Appropriate Transformations

```
# Good: Log scale for learning rate
bounds = [(1e-5, 1e-1)]
var_trans = ["log10"]

# Bad: No transformation for wide range
bounds = [(1e-5, 1e-1)]
var_trans = [None] # Poor sampling distribution
```

13.8.2. 2. Match Transformation to Range

```
# Wide range (>3 orders of magnitude): use log
bounds = [(1e-6, 1e-2)]
var_trans = ["log10"]

# Moderate range (1-2 orders): use sqrt
bounds = [(10, 500)]
var_trans = ["sqrt"]

# Narrow range (<1 order): no transformation
bounds = [(-1, 1)]
var_trans = [None]
```

13.8.3. 3. Validate Transformation Choice

```
# Check if transformation makes sense
import numpy as np

# Original space
x_orig = np.linspace(0.001, 1.0, 10)
print("Original:", x_orig)
```

13. Variable Transformations for Search Space Scaling

```
# Log10 transformed space
x_trans = np.log10(x_orig)
print("Transformed:", x_trans)
print("Range ratio:", np.ptp(x_trans) / np.ptp(x_orig))
# Should be much more uniform distribution
```

13.8.4. 4. Combine with Variable Types

```
# Mix transformations with variable types
optimizer = SpotOptim(
    fun=objective,
    bounds=[
        (10, 200),                # int with sqrt
        ("ReLU", "Tanh", "Sigmoid"), # factor (no transform)
        (0.0001, 0.1),            # num with log10
        (0.01, 1.0)               # num with log10
    ],
    var_type=["int", "factor", "float", "float"],
    var_trans=["sqrt", None, "log10", "log10"],
    var_name=["neurons", "activation", "lr", "dropout"]
)
```

13.9. Troubleshooting

13.9.1. Issue: Values Out of Bounds

Problem: Objective function receives values outside specified bounds.

Solution: This should not happen with transformations. If it does:

```
# Check transformation is applied correctly
print(f"Original bounds: {optimizer._original_lower} to {optimizer._original_upper}")
print(f"Transformed bounds: {optimizer.lower} to {optimizer.upper}")
print(f"Transformations: {optimizer.var_trans}")
```

13.9.2. Issue: Poor Optimization Performance

Problem: Optimization doesn't find good solutions.

Possible causes:

13.10. Comparison: Manual vs Automatic Transformations

1. Wrong transformation type for the parameter scale
2. Transformation not needed (adding unnecessary complexity)
3. Bounds too wide or too narrow

Solution:

```
# Try different transformations
for trans in [None, "log10", "sqrt"]:
    optimizer = SpotOptim(
        fun=objective,
        bounds=[(0.001, 1.0)],
        var_trans=[trans],
        max_iter=20,
        seed=42
    )
    result = optimizer.optimize()
    print(f"Transformation: {trans}, Best: {result.fun:.6f}")
```

13.9.3. Issue: Transformation Not Applied

Problem: Transformation doesn't seem to affect optimization.

Check:

```
# Verify var_trans length matches dimensions
print(f"Number of dimensions: {len(optimizer.bounds)}")
print(f"Number of transformations: {len(optimizer.var_trans)}")
# These must match!

# Check transformation is not None/"id"
print(f"Transformations: {optimizer.var_trans}")
```

13.10. Comparison: Manual vs Automatic Transformations

13.10.1. Manual Approach (Old Way)

```
def objective_manual(X):
    """Manual transformation - error-prone!"""
    results = []
```

13. Variable Transformations for Search Space Scaling

```
for params in X:
    # Must remember to transform
    lr = 10 ** params[0] # Was in log scale
    alpha = 10 ** params[1] # Was in log scale

    # Use parameters
    score = compute_score(lr, alpha)
    results.append(score)
return np.array(results)

# Bounds in log scale - confusing!
optimizer = SpotOptim(
    fun=objective_manual,
    bounds=[(-4, -1), (-2, 0)], # log10 scale
    var_name=["log10_lr", "log10_alpha"] # Confusing names
)

result = optimizer.optimize()
# Must transform back for interpretation
best_lr = 10 ** result.x[0]
best_alpha = 10 ** result.x[1]
```

13.10.2. Automatic Approach (New Way)

```
def objective_auto(X):
    """Automatic transformation - clean!"""
    results = []
    for params in X:
        # Already in original scale
        lr = params[0]
        alpha = params[1]

        # Use parameters directly
        score = compute_score(lr, alpha)
        results.append(score)
    return np.array(results)

# Bounds in natural scale - intuitive!
optimizer = SpotOptim(
    fun=objective_auto,
    bounds=[(0.0001, 0.1), (0.01, 1.0)],
    var_trans=["log10", "log10"], # Specify transformation
```

```

    var_name=["lr", "alpha"] # Natural names
)

result = optimizer.optimize()
# Results already in original scale
best_lr = result.x[0]
best_alpha = result.x[1]

```

13.11. Summary

Key Takeaways:

1. ☒ Use `var_trans` to specify transformations for each dimension
2. ☒ Transformations improve optimization for poorly scaled spaces
3. ☒ All user interfaces (bounds, results, plots) use original scale
4. ☒ Optimization happens internally in transformed space
5. ☒ Common transformations: "log10" for learning rates, "sqrt" for moderate scaling
6. ☒ View transformations in tables with “trans” column

When to Use:

- Parameters spanning multiple orders of magnitude → "log10" or "log"
- Moderate scaling (1-2 orders) → "sqrt"
- Reciprocal relationships → "inv"
- Well-scaled parameters → None

Benefits:

- Better surrogate model performance
- More efficient sampling
- Improved optimization convergence
- User-friendly interface (no manual transformations in objective function)

13.12. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

Part III.

Visualization

14. Surrogate Model Visualization

This document describes the `plot_surrogate()` method added to the `SpotOptim` class, which provides visualization capabilities similar to the `plotkd()` function in the `spotpython` package.

14.1. Overview

The `plot_surrogate()` method creates a comprehensive 4-panel visualization of the fitted surrogate model, showing both predictions and uncertainty estimates across two selected dimensions.

14.2. Features

- **3D Surface Plots:** Visualize the surrogate's predictions and uncertainty as 3D surfaces
- **Contour Plots:** View 2D contours with overlaid evaluation points
- **Multi-dimensional Support:** Visualize any two dimensions of higher-dimensional problems
- **Customizable Appearance:** Control colors, resolution, transparency, and more

14.3. Usage

14.3.1. Basic Usage


```
import numpy as np
from spotoptim import SpotOptim

# Define objective function
def sphere(X):
    return np.sum(X**2, axis=1)
```

14. Surrogate Model Visualization

```
# Run optimization
optimizer = SpotOptim(fun=sphere, bounds=[(-5, 5), (-5, 5)], max_iter=20)
result = optimizer.optimize()

# Visualize the surrogate model
optimizer.plot_surrogate(i=0, j=1, show=True)
```




14.3.2. With Custom Parameters

```
optimizer.plot_surrogate(  
    i=0, # First dimension to plot  
    j=1, # Second dimension to plot  
    var_name=['x1', 'x2'], # Variable names for axes  
    add_points=True, # Show evaluated points
```

```

cmap='viridis',          # Colormap
alpha=0.7,               # Surface transparency
num=100,                 # Grid resolution
contour_levels=25,       # Number of contour levels
grid_visible=True,       # Show grid on contours
figsize=(12, 10),        # Figure size
show=True                # Display immediately
)

```



14.3.3. Higher-Dimensional Problems

For problems with more than 2 dimensions, `plot_surrogate()` creates a 2D slice by fixing all other dimensions at their mean values:


14. Surrogate Model Visualization

```
# 4D optimization problem
def sphere_4d(X):
    return np.sum(X**2, axis=1)


bounds = [(-3, 3)] * 4
optimizer = SpotOptim(fun=sphere_4d, bounds=bounds, max_iter=20)
result = optimizer.optimize()

# Visualize dimensions 0 and 2 (dimensions 1 and 3 fixed at mean)
optimizer.plot_surrogate(
    i=0, j=2,
    var_name=['x0', 'x1', 'x2', 'x3']
)

# Visualize different dimension pair
optimizer.plot_surrogate(i=1, j=3, var_name=['x0', 'x1', 'x2', 'x3'])
```



14.4. Plot Interpretation



14.4. Plot Interpretation

The visualization consists of 4 panels:

14.4.1. Top Left: Prediction Surface

- Shows the surrogate model's predicted function values as a 3D surface
- Helps understand the model's belief about the objective function landscape
- Lower values (blue in default colormap) indicate predicted minima

14.4.2. Top Right: Prediction Uncertainty Surface

- Shows the standard deviation of predictions as a 3D surface
- Indicates where the model is uncertain and might benefit from more samples

14. Surrogate Model Visualization

- Lower values (blue) indicate high confidence, higher values (red) indicate uncertainty

14.4.3. Bottom Left: Prediction Contour with Points

- 2D contour plot of predictions
- Red dots show the actual points evaluated during optimization
- Useful for understanding the exploration-exploitation trade-off

14.4.4. Bottom Right: Uncertainty Contour with Points

- 2D contour plot of prediction uncertainty
- Shows how uncertainty decreases around evaluated points
- Helps identify unexplored regions

14.5. Parameters

14.5.1. Dimension Selection

- `i` (int, default=0): Index of first dimension to plot
- `j` (int, default=1): Index of second dimension to plot

14.5.2. Appearance

- `var_name` (list of str, optional): Names for each dimension
- `cmap` (str, default='jet'): Matplotlib colormap name
- `alpha` (float, default=0.8): Surface transparency (0=transparent, 1=opaque)
- `figsize` (tuple, default=(12, 10)): Figure size in inches (width, height)

14.5.3. Grid and Resolution

- `num` (int, default=100): Number of grid points per dimension
- `contour_levels` (int, default=30): Number of contour levels
- `grid_visible` (bool, default=True): Show grid lines on contour plots

14.5.4. Color Scaling

- `vmin` (float, optional): Minimum value for color scale
- `vmax` (float, optional): Maximum value for color scale

14.5.5. Display

- `show` (bool, default=True): Display plot immediately
- `add_points` (bool, default=True): Overlay evaluated points on contours

14.6. Examples

14.6.1. Example 1: 2D Rosenbrock Function


```
import numpy as np
from spotoptim import SpotOptim

def rosenbrock(X):
    X = np.atleast_2d(X)
    x, y = X[:, 0], X[:, 1]
    return (1 - x)**2 + 100 * (y - x**2)**2

optimizer = SpotOptim(
    fun=rosenbrock,
    bounds=[(-2, 2), (-2, 2)],
    max_iter=30,
    seed=42
)
result = optimizer.optimize()

# Visualize with custom colormap
optimizer.plot_surrogate(
    var_name=['x', 'y'],
    cmap='coolwarm',
    add_points=True
)
```

14. Surrogate Model Visualization



14.6.2. Example 2: Using Kriging Surrogate


```
import numpy as np
from spotoptim import SpotOptim, Kriging

def sphere(X):
    return np.sum(X**2, axis=1)

optimizer = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    surrogate=Kriging(seed=42), # Use Kriging instead of GP
    max_iter=20
)
result = optimizer.optimize()
```



```
# The plotting works the same with any surrogate
optimizer.plot_surrogate(var_name=['x1', 'x2'])
```



14.6.3. Example 3: Comparing Different Dimension Pairs

```
# 3D problem - visualize all dimension pairs
def sphere_3d(X):
    return np.sum(X**2, axis=1)

optimizer = SpotOptim(
    fun=sphere_3d,
    bounds=[(-5, 5)] * 3,
    max_iter=25
)
```


14. Surrogate Model Visualization

```
result = optimizer.optimize()


# Dimensions 0 vs 1
optimizer.plot_surrogate(i=0, j=1, var_name=['x0', 'x1', 'x2'])

# Dimensions 0 vs 2
optimizer.plot_surrogate(i=0, j=2, var_name=['x0', 'x1', 'x2'])


# Dimensions 1 vs 2
optimizer.plot_surrogate(i=1, j=2, var_name=['x0', 'x1', 'x2'])
```



14.6. Examples



14. Surrogate Model Visualization



14.7. Tips and Best Practices

1. **Run Optimization First:** Always call `optimize()` before `plot_surrogate()`
2. **Choose Dimensions Wisely:** For high-dimensional problems, plot dimensions that you suspect are most important or interactive
3. **Adjust Resolution:** Use lower `num` values (e.g., 50) for faster plotting, higher values (e.g., 200) for smoother surfaces
4. **Color Scales:** Set `vmin` and `vmax` explicitly when comparing multiple plots to ensure consistent color scales
5. **Uncertainty Analysis:** High uncertainty areas (bright colors in uncertainty plots) are good candidates for additional sampling
6. **Exploration vs Exploitation:** Red dots clustered in low-prediction areas show exploitation; spread-out dots show exploration

14.8. Comparison with spotpython's plotkd()

The `plot_surrogate()` method is inspired by spotpython's `plotkd()` function but adapted for SpotOptim's simplified interface:

14.8.1. Similarities

- Same 4-panel layout (2 surfaces + 2 contours)
- Visualizes predictions and uncertainty
- Supports dimension selection and customization

14.8.2. Differences

- **Integration:** Method of SpotOptim class (no separate function needed)
- **Simpler:** Fewer parameters, more sensible defaults
- **Automatic:** Uses optimizer's bounds and data automatically
- **Type Handling:** Automatically applies variable type constraints (int/float/factor)

14.9. Error Handling

The method validates inputs and provides clear error messages:

```
# Before optimization runs
optimizer.plot_surrogate() # ValueError: No optimization data available

# Invalid dimension indices
optimizer.plot_surrogate(i=5, j=1) # ValueError: i must be less than n_dim

# Same dimension twice
optimizer.plot_surrogate(i=0, j=0) # ValueError: i and j must be different
```

14.10. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

15. TensorBoard Logging in SpotOptim

SpotOptim supports TensorBoard logging for monitoring optimization progress in real-time.

15.1. Quick Start

15.1.1. Enable TensorBoard Logging

```
from spotoptim import SpotOptim
import numpy as np

def sphere(X):
    return np.sum(X**2, axis=1)

optimizer = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=20,
    n_initial=10,
    tensorboard_log=True, # Enable logging
    tensorboard_clean=True,
    verbose=True,
    seed=42
)

result = optimizer.optimize()
print(f"Best value: {result.fun:.6f}")
print(f"Logs saved to: runs/{optimizer.tensorboard_path}")
```

```
Removed old TensorBoard logs: runs/spotoptim_20260701_204408
Cleaned 1 old TensorBoard log directory
TensorBoard logging enabled: runs/spotoptim_20260701_204415
Initial best: f(x) = 3.925721
Iter 1 | Best: 0.020505 | Rate: 1.00 | Evals: 55.0%
```

15. TensorBoard Logging in SpotOptim

```
Iter 2 | Best: 0.000751 | Rate: 1.00 | Evals: 60.0%
Iter 3 | Best: 0.000024 | Rate: 1.00 | Evals: 65.0%
Iter 4 | Best: 0.000000 | Rate: 1.00 | Evals: 70.0%
Iter 5 | Best: 0.000000 | Rate: 1.00 | Evals: 75.0%
Iter 6 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.83 | Evals: 80.0%
Optimizer candidate 1/3 was duplicate/invalid.
Optimizer candidate 2/3 was duplicate/invalid.
Optimizer candidate 3/3 was duplicate/invalid.
Fallback attempt 1/10: Using fallback strategy
Acquisition failure: Using random space-filling design as fallback.
Iter 7 | Best: 0.000000 | Curr: 19.241544 | Rate: 0.71 | Evals: 85.0%
Iter 8 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.62 | Evals: 90.0%
Iter 9 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.56 | Evals: 95.0%
Iter 10 | Best: 0.000000 | Curr: 0.000001 | Rate: 0.50 | Evals: 100.0%
TensorBoard writer closed. View logs with: tensorboard --logdir=runs/spotoptim_20260701_204415
Best value: 0.000000
Logs saved to: runs/runs/spotoptim_20260701_204415
```

15.1.2. View Logs in TensorBoard

In a separate terminal, run:

```
tensorboard --logdir=runs
```

Then open your browser to <http://localhost:6006>

15.1.3. Cleaning Old Logs

You can automatically remove old TensorBoard logs before starting a new optimization:

```
optimizer = SpotOptim(
    fun=objective,
    bounds=[(-5, 5), (-5, 5)],
    tensorboard_log=True,
    tensorboard_clean=True, # Remove old logs from 'runs' directory
    verbose=True
)
```


i Warning

This permanently deletes all subdirectories in the `runs` folder. Make sure to save important logs elsewhere before enabling this feature.

15.1.4. Use Cases

1. **Clean Start** - Remove old logs and create new one:

```
tensorboard_log=True, tensorboard_clean=True
```

2. **Preserve History** - Keep old logs and add new one (default):

```
tensorboard_log=True, tensorboard_clean=False
```

3. **Just Clean** - Remove old logs without new logging:

```
tensorboard_log=False, tensorboard_clean=True
```

15.1.5. Custom Log Directory

Specify a custom path for TensorBoard logs:

```
optimizer = SpotOptim(
    fun=objective,
    bounds=[(-5, 5), (-5, 5)],
    tensorboard_log=True,
    tensorboard_path="my_experiments/run_001",
    ...
)
```

15.1.6. What Gets Logged**15.1.6.1. Scalar Metrics**

For Deterministic Functions:

- `y_values/min`: Best (minimum) `y` value found so far
- `y_values/last`: Most recently evaluated `y` value
- `X_best/x0`, `X_best/x1`, ...: Coordinates of the best point

For Noisy Functions (repeats > 1):

- `y_values/min`: Best single evaluation

15. TensorBoard Logging in SpotOptim

- `y_values/mean_best`: Best mean y value
- `y_values/last`: Most recent evaluation
- `y_variance_at_best`: Variance at the best mean point
- `X_mean_best/x0`, `X_mean_best/x1`, ...: Coordinates of best mean point

15.1.6.2. Hyperparameters

Each function evaluation is logged with:

- Input coordinates (`x0`, `x1`, `x2`, ...)
- Function value (`hp_metric`)

This allows you to explore the relationship between hyperparameters and objective values in the HPARAMS tab.

15.1.7. Examples

Example 15.1 (Basic Tensorboard Usage).

```
import numpy as np
from spotoptim import SpotOptim

optimizer = SpotOptim(
    fun=lambda X: np.sum(X**2, axis=1),
    bounds=[(-5, 5), (-5, 5)],
    max_iter=20,
    n_initial=10,
    tensorboard_log=True,
    tensorboard_clean=True,
    verbose=True,
    seed=42
)
result = optimizer.optimize()
print(f"Best value: {result.fun:.6f}")
```

```
Removed old TensorBoard logs: runs/spotoptim_20260701_204415
Cleaned 1 old TensorBoard log directory
TensorBoard logging enabled: runs/spotoptim_20260701_204419
Initial best: f(x) = 3.925721
Iter 1 | Best: 0.020505 | Rate: 1.00 | Evals: 55.0%
Iter 2 | Best: 0.000751 | Rate: 1.00 | Evals: 60.0%
Iter 3 | Best: 0.000024 | Rate: 1.00 | Evals: 65.0%
Iter 4 | Best: 0.000000 | Rate: 1.00 | Evals: 70.0%
```

```

Iter 5 | Best: 0.000000 | Rate: 1.00 | Evals: 75.0%
Iter 6 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.83 | Evals: 80.0%
Optimizer candidate 1/3 was duplicate/invalid.
Optimizer candidate 2/3 was duplicate/invalid.
Optimizer candidate 3/3 was duplicate/invalid.
Fallback attempt 1/10: Using fallback strategy
Acquisition failure: Using random space-filling design as fallback.
Iter 7 | Best: 0.000000 | Curr: 19.241544 | Rate: 0.71 | Evals: 85.0%
Iter 8 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.62 | Evals: 90.0%
Iter 9 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.56 | Evals: 95.0%
Iter 10 | Best: 0.000000 | Curr: 0.000001 | Rate: 0.50 | Evals: 100.0%
TensorBoard writer closed. View logs with: tensorboard --logdir=runs/spotoptim_20260701_204419
Best value: 0.000000

```

Example 15.2 (Noisy Optimization).

```

import numpy as np
from spotoptim import SpotOptim

def noisy_objective(X):
    base = np.sum(X**2, axis=1)
    noise = np.random.normal(0, 0.1, size=base.shape)
    return base + noise

optimizer = SpotOptim(
    fun=noisy_objective,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=20,
    n_initial=10,
    repeats_initial=3,
    repeats_surrogate=2,
    tensorboard_log=True,
    tensorboard_clean=True,
    tensorboard_path="runs/noisy_exp",
    seed=42
)
result = optimizer.optimize()
print(f"Best value: {result.fun:.6f}")

```

Best value: 3.784490

Example 15.3 (With OCBA).

15. TensorBoard Logging in SpotOptim

```
import numpy as np
from spotoptim import SpotOptim

def noisy_objective(X):
    base = np.sum(X**2, axis=1)
    noise = np.random.normal(0, 0.1, size=base.shape)
    return base + noise

optimizer = SpotOptim(
    fun=noisy_objective,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=20,
    n_initial=10,
    repeats_initial=2,
    ocba_delta=3, # Re-evaluate 3 promising points per iteration
    tensorboard_log=True,
    tensorboard_clean=True,
    seed=42
)
result = optimizer.optimize()
print(f"Best value: {result.fun:.6f}")
```

Best value: 3.734393

Example 15.4 (Comparing Multiple Runs). Run multiple optimizations with different settings:

```
# Run 1: Standard
opt1 = SpotOptim(..., tensorboard_path="runs/standard")
opt1.optimize()

# Run 2: With OCBA
opt2 = SpotOptim(..., ocba_delta=3, tensorboard_path="runs/with_ocba")
opt2.optimize()

# Run 3: More initial points
opt3 = SpotOptim(..., n_initial=20, tensorboard_path="runs/more_initial")
opt3.optimize()
```

Then view all runs together:

```
tensorboard --logdir=runs
```

15.2. TensorBoard Features

15.2.1. SCALARS Tab

- View convergence curves
- Compare optimization progress across runs
- Track how metrics change over iterations

15.2.2. HPARAMS Tab

- Explore hyperparameter space
- See which parameter combinations work best
- Identify patterns in successful configurations

15.2.3. Text Tab

- View configuration details
- Check run metadata

Tips

1. **Organize Experiments:** Use descriptive tensorboard_path names:

```
tensorboard_path=f"runs/exp_{date}_{config_name}"
```

2. **Compare Algorithms:** Run multiple optimization strategies and compare:

```
# Different acquisition functions
for acq in ['ei', 'pi', 'y']:
    opt = SpotOptim(..., acquisition=acq, tensorboard_path=f"runs/acq_{acq}")
    opt.optimize()
```

3. **Clean Up Old Runs:** Use tensorboard_clean=True for automatic cleanup, or manually:

```
rm -rf runs/old_experiment
```

4. **Port Conflicts:** If port 6006 is busy, use a different port:

```
tensorboard --logdir=runs --port=6007
```

15.3. TensorBoard Log Cleaning Feature in SpotOptim

Automatic cleaning of old TensorBoard log directories with the `tensorboard_clean` parameter.

15.3.1. Basic Usage

```
import numpy as np
from spotoptim import SpotOptim

def sphere(X):
    """Simple sphere function"""
    return np.sum(X**2, axis=1)

# Remove old logs and create new log directory
optimizer = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=20,
    n_initial=10,
    tensorboard_log=True,
    tensorboard_clean=True, # Removes all subdirectories in 'runs'
    verbose=True,
    seed=42
)

result = optimizer.optimize()
print(f"Best value: {result.fun:.6f}")
print(f"Logs saved to: runs/{optimizer.tensorboard_path}")
```

```
Removed old TensorBoard logs: runs/spotoptim_20260701_204423
Cleaned 1 old TensorBoard log directory
TensorBoard logging enabled: runs/spotoptim_20260701_204423
Initial best: f(x) = 3.925721
Iter 1 | Best: 0.020505 | Rate: 1.00 | Evals: 55.0%
Iter 2 | Best: 0.000751 | Rate: 1.00 | Evals: 60.0%
Iter 3 | Best: 0.000024 | Rate: 1.00 | Evals: 65.0%
```

15.3. TensorBoard Log Cleaning Feature in SpotOptim

```
Iter 4 | Best: 0.000000 | Rate: 1.00 | Evals: 70.0%
Iter 5 | Best: 0.000000 | Rate: 1.00 | Evals: 75.0%
Iter 6 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.83 | Evals: 80.0%
Optimizer candidate 1/3 was duplicate/invalid.
Optimizer candidate 2/3 was duplicate/invalid.
Optimizer candidate 3/3 was duplicate/invalid.
Fallback attempt 1/10: Using fallback strategy
Acquisition failure: Using random space-filling design as fallback.
Iter 7 | Best: 0.000000 | Curr: 19.241544 | Rate: 0.71 | Evals: 85.0%
Iter 8 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.62 | Evals: 90.0%
Iter 9 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.56 | Evals: 95.0%
Iter 10 | Best: 0.000000 | Curr: 0.000001 | Rate: 0.50 | Evals: 100.0%
TensorBoard writer closed. View logs with: tensorboard --logdir=runs/spotoptim_20260701_204423
Best value: 0.000000
Logs saved to: runs/runs/spotoptim_20260701_204423
```

15.3.2. Use Cases

tensorboard_log	tensorboard_clean	Behavior
True	True	Clean old logs, create new log directory
True	False	Preserve old logs, create new log directory
False	True	Clean old logs, no new logging
False	False	No logging, no cleaning (default)

15.3.3. Implementation Details

15.3.3.1. Cleaning Method

```
def _clean_tensorboard_logs(self) -> None:
    """Clean old TensorBoard log directories from the runs folder."""
    if self.tensorboard_clean:
        runs_dir = "runs"
        if os.path.exists(runs_dir) and os.path.isdir(runs_dir):
            # Get all subdirectories in runs
            subdirs = [
                os.path.join(runs_dir, d)
                for d in os.listdir(runs_dir)
                if os.path.isdir(os.path.join(runs_dir, d))
            ]
```

15. TensorBoard Logging in SpotOptim

```
# Remove each subdirectory
for subdir in subdirs:
    try:
        shutil.rmtree(subdir)
        if self.verbose:
            print(f"Removed old TensorBoard logs: {subdir}")
    except Exception as e:
        if self.verbose:
            print(f"Warning: Could not remove {subdir}: {e}")
```

15.3.4. Execution Flow

1. User creates `SpotOptim` instance with `tensorboard_clean=True`
2. During initialization, `_clean_tensorboard_logs()` is called
3. Method checks if 'runs' directory exists
4. Removes all subdirectories (but preserves files)
5. If `tensorboard_log=True`, a new log directory is created
6. Optimization proceeds normally

15.4. Safety Features

- Only removes **directories**, not files in 'runs' folder
- Handles missing 'runs' directory gracefully
- Error handling for permission issues
- Verbose output shows what's being removed
- Default is `False` to prevent accidental deletion

Warning

- Setting `tensorboard_clean=True` permanently deletes all subdirectories in the 'runs' folder. Make sure to save important logs elsewhere before enabling this feature.

15.5. Tensorboard Demo Scripts

Run the comprehensive TensorBoard demo:

```
python demo_tensorboard.py
```


This demonstrates:

- Deterministic optimization (Rosenbrock function)
- Noisy optimization with repeated evaluations
- OCBA for intelligent re-evaluation

Run the log cleaning demo:

```
python demo_tensorboard_clean.py
```

This demonstrates:

- Creating multiple log directories
- Preserving old logs (default behavior)
- Cleaning old logs automatically
- Cleaning without creating new logs

15.6. Troubleshooting

Q: TensorBoard shows “No dashboards are active” A: Make sure you’ve run an optimization with `tensorboard_log=True` first.

Q: Can’t see my latest run A: Refresh TensorBoard (click the reload button in the upper right).

Q: How do I stop TensorBoard? A: Press Ctrl+C in the terminal where TensorBoard is running.

Q: Logs taking up too much space? A: Use `tensorboard_clean=True` to automatically remove old logs, or manually delete old run directories.

Q: How do I remove all old logs at once? A: Set `tensorboard_clean=True` when creating your optimizer. This will remove all subdirectories in the `runs` folder.

15.7. Related Parameters

- `tensorboard_log` (bool): Enable/disable logging (default: False)
- `tensorboard_path` (str): Custom log directory (default: auto-generated with timestamp)
- `tensorboard_clean` (bool): Remove old logs from ‘runs’ directory before starting (default: False)
- `verbose` (bool): Print progress to console (default: False)
- `var_name` (list): Custom names for variables (used in TensorBoard labels)

15.8. Performance Notes

TensorBoard logging has minimal overhead:

- Event files are efficiently buffered and written
- Writer is properly closed after optimization completes

15.9. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

Part IV.

Saving and Loading Models

16. Save and Load in SpotOptim

SpotOptim provides comprehensive save and load functionality for serializing optimization configurations and results. This enables distributed workflows where experiments are defined locally, executed remotely, and analyzed back on the local machine.

16.1. Key Concepts

16.1.1. Experiments vs Results

SpotOptim distinguishes between two types of saved data:

- **Experiment** (*_exp.pkl): Configuration only, excluding the objective function and results. Used to transfer optimization setup to remote machines.
- **Result** (*_res.pkl): Complete optimization state including configuration, all evaluations, and results. Used to save and analyze completed optimizations.

Table 16.1 shows what gets saved in each file type.

Table 16.1.: Robust Function Saving

Component	Experiment	Result
Configuration (bounds, parameters)	x	x
Objective function	x	x
Evaluations (X, y)		x
Best solution		x
Surrogate model	Excluded*	x
TensorBoard writer		

SpotOptim uses `dill` for serialization, which allows robust saving of objective functions, including lambda functions and locally defined functions.

16.2. Quick Start

```
import numpy as np
from spotoptim import SpotOptim

def sphere(X):
    """Simple sphere function"""
    return np.sum(X**2, axis=1)

# Create and configure optimizer
optimizer = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=20,
    n_initial=10,
    seed=42
)

# Run optimization
result = optimizer.optimize()
print(f"Best value: {result.fun:.6f}")

# Save complete results
optimizer.save_result(prefix="sphere_opt")
# Creates: sphere_opt_res.pkl

# Later: load and analyze results
loaded_opt = SpotOptim.load_result("sphere_opt_res.pkl")
print(f"Loaded best value: {loaded_opt.best_y_:.6f}")
print(f"Total evaluations: {loaded_opt.counter}")
```

```
Best value: 0.000000
Experiment saved to sphere_opt_res.pkl
Result saved to sphere_opt_res.pkl
Loaded result from sphere_opt_res.pkl
Loaded best value: 0.000000
Total evaluations: 20
```

16.3. Distributed Workflow

The save/load functionality enables a powerful workflow for distributed optimization:

16.3.1. Step 1: Define Experiment Locally

Generate the `remote_job_001_exp.pkl` file by running the following code:

```
import numpy as np
from spotoptim import SpotOptim
from spotoptim.function import rosenbrock

optimizer = SpotOptim(
    fun=rosenbrock,
    bounds=[(-2, 2), (-2, 2)],
    max_iter=50,
    n_initial=10,
    seed=42,
    verbose=True)

optimizer.save_experiment(prefix="remote_job_001")
```

TensorBoard logging disabled
Experiment saved to `remote_job_001_exp.pkl`

16.3.2. Step 2: Execute on Remote Machine

Generate the `remote_job_001_res.pkl` file by running the following code on the remote machine:

```
from spotoptim import SpotOptim
optimizer = SpotOptim.load_experiment("remote_job_001_exp.pkl")
result = optimizer.optimize()
optimizer.save_result(prefix="remote_job_001")
```

```
Loaded experiment from remote_job_001_exp.pkl
Initial best: f(x) = 6.724883
Iter 1 | Best: 5.101214 | Rate: 1.00 | Evals: 22.0%
Iter 2 | Best: 5.101214 | Curr: 5.993415 | Rate: 0.50 | Evals: 24.0%
Iter 3 | Best: 5.101214 | Curr: 27.962855 | Rate: 0.33 | Evals: 26.0%
Iter 4 | Best: 3.502923 | Rate: 0.50 | Evals: 28.0%
Iter 5 | Best: 0.109107 | Rate: 0.60 | Evals: 30.0%
Iter 6 | Best: 0.109107 | Curr: 1.546493 | Rate: 0.50 | Evals: 32.0%
Iter 7 | Best: 0.109107 | Curr: 0.209959 | Rate: 0.43 | Evals: 34.0%
Iter 8 | Best: 0.094949 | Rate: 0.50 | Evals: 36.0%
Iter 9 | Best: 0.094949 | Curr: 0.106009 | Rate: 0.44 | Evals: 38.0%
Iter 10 | Best: 0.092839 | Rate: 0.50 | Evals: 40.0%
```

16. Save and Load in SpotOptim

```
Iter 11 | Best: 0.090991 | Rate: 0.55 | Evals: 42.0%
Iter 12 | Best: 0.090991 | Curr: 0.141817 | Rate: 0.50 | Evals: 44.0%
Iter 13 | Best: 0.088978 | Rate: 0.54 | Evals: 46.0%
Iter 14 | Best: 0.087946 | Rate: 0.57 | Evals: 48.0%
Iter 15 | Best: 0.086988 | Rate: 0.60 | Evals: 50.0%
Iter 16 | Best: 0.086155 | Rate: 0.62 | Evals: 52.0%
Iter 17 | Best: 0.085517 | Rate: 0.65 | Evals: 54.0%
Iter 18 | Best: 0.084230 | Rate: 0.67 | Evals: 56.0%
Iter 19 | Best: 0.082585 | Rate: 0.68 | Evals: 58.0%
Iter 20 | Best: 0.080877 | Rate: 0.70 | Evals: 60.0%
Iter 21 | Best: 0.078872 | Rate: 0.71 | Evals: 62.0%
Iter 22 | Best: 0.076996 | Rate: 0.73 | Evals: 64.0%
Iter 23 | Best: 0.074527 | Rate: 0.74 | Evals: 66.0%
Iter 24 | Best: 0.072343 | Rate: 0.75 | Evals: 68.0%
Iter 25 | Best: 0.069881 | Rate: 0.76 | Evals: 70.0%
Iter 26 | Best: 0.067964 | Rate: 0.77 | Evals: 72.0%
Iter 27 | Best: 0.066997 | Rate: 0.78 | Evals: 74.0%
Iter 28 | Best: 0.065905 | Rate: 0.79 | Evals: 76.0%
Iter 29 | Best: 0.064781 | Rate: 0.79 | Evals: 78.0%
Iter 30 | Best: 0.064713 | Rate: 0.80 | Evals: 80.0%
Iter 31 | Best: 0.064065 | Rate: 0.81 | Evals: 82.0%
Iter 32 | Best: 0.063214 | Rate: 0.81 | Evals: 84.0%
Iter 33 | Best: 0.061810 | Rate: 0.82 | Evals: 86.0%
Iter 34 | Best: 0.060027 | Rate: 0.82 | Evals: 88.0%
Iter 35 | Best: 0.059298 | Rate: 0.83 | Evals: 90.0%
Iter 36 | Best: 0.059298 | Curr: 0.059606 | Rate: 0.81 | Evals: 92.0%
Iter 37 | Best: 0.058625 | Rate: 0.81 | Evals: 94.0%
Iter 38 | Best: 0.057821 | Rate: 0.82 | Evals: 96.0%
Iter 39 | Best: 0.056506 | Rate: 0.82 | Evals: 98.0%
Iter 40 | Best: 0.054879 | Rate: 0.82 | Evals: 100.0%
Experiment saved to remote_job_001_res.pkl
Result saved to remote_job_001_res.pkl
```

16.3.3. Step 3: Analyze Results Locally

```
from spotoptim import SpotOptim
optimizer = SpotOptim.load_result("remote_job_001_res.pkl")
print(f"Best value found: {optimizer.best_y:.6f}")
print(f"Best point: {optimizer.best_x}")
print(f"Total evaluations: {optimizer.counter}")
print(f"Number of iterations: {optimizer.n_iter}")
```

Loaded result from remote_job_001_res.pkl

Best value found: 0.054879
 Best point: [0.76577181 0.58600184]
 Total evaluations: 50
 Number of iterations: 40

```
optimizer.plot_progress(log_y=True)
```




Figure 16.1.: Optimization Progress

Access all evaluated points as follows:

```
print(f"\nAll evaluated points shape: {optimizer.X_.shape}")
print(f"All objective values shape: {optimizer.y_.shape}")
```

All evaluated points shape: (50, 2)
 All objective values shape: (50,)

16.4. Advanced Usage

16.4.1. Custom Filenames and Paths

```
import os
from spotoptim import SpotOptim
import numpy as np

def objective(X):
    return np.sum(X**2, axis=1)

optimizer = SpotOptim(
    fun=objective,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=30,
    seed=42
)

# Save with custom filename
optimizer.save_experiment(
    filename="custom_name.pkl",
    verbosity=1
)

# Save to specific directory
os.makedirs("experiments/batch_001", exist_ok=True)
optimizer.save_experiment(
    prefix="exp_001",
    path="experiments/batch_001",
    verbosity=1
)
# Creates: experiments/batch_001/exp_001_exp.pkl
```

Experiment saved to custom_name.pkl

Experiment saved to experiments/batch_001/exp_001_exp.pkl

16.4.2. Overwrite Protection

```
from spotoptim import SpotOptim
import numpy as np
```

```

def sphere(X):
    return np.sum(X**2, axis=1)

optimizer = SpotOptim(fun=sphere, bounds=[(-5, 5), (-5, 5)], max_iter=20)
result = optimizer.optimize()

# First save
optimizer.save_result(prefix="my_result")

# Try to save again - raises FileExistsError by default
try:
    optimizer.save_result(prefix="my_result")
except FileExistsError as e:
    print(f"File already exists: {e}")

# Explicitly allow overwriting
optimizer.save_result(prefix="my_result", overwrite=True)
print("File overwritten successfully")

```

```

Experiment saved to my_result_res.pkl
Result saved to my_result_res.pkl
Experiment saved to my_result_res.pkl
Result saved to my_result_res.pkl
Experiment saved to my_result_res.pkl
Result saved to my_result_res.pkl
File overwritten successfully

```

16.4.3. Loading and Continuing Optimization

```

from spotoptim import SpotOptim
import numpy as np

def objective(X):
    return np.sum(X**2, axis=1)

# Initial optimization
opt1 = SpotOptim(
    fun=objective,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=20,
    seed=42

```

16. Save and Load in SpotOptim

```
)
result1 = opt1.optimize()
opt1.save_result(prefix="checkpoint")

print(f"Initial optimization: {result1.nfev} evaluations, best={result1.fun:.6f}")

# Load and continue
opt2 = SpotOptim.load_result("checkpoint_res.pkl")
# No need to re-attach function as it is loaded with dill
opt2.max_iter = 25 # Increase budget

# Continue optimization
result2 = opt2.optimize()
print(f"After continuation: {result2.nfev} evaluations, best={result2.fun:.6f}")
```

```
Experiment saved to checkpoint_res.pkl
Result saved to checkpoint_res.pkl
Initial optimization: 20 evaluations, best=0.000000
Loaded result from checkpoint_res.pkl
After continuation: 25 evaluations, best=0.000000
```

16.5. Working with Noisy Functions

Save and load preserves noise statistics for reproducible analysis:

```
import numpy as np
from spotoptim import SpotOptim

def noisy_objective(X):
    """Objective with measurement noise"""
    true_value = np.sum(X**2, axis=1)
    noise = np.random.normal(0, 0.1, X.shape[0])
    return true_value + noise

# Optimize noisy function with repeated evaluations
optimizer = SpotOptim(
    fun=noisy_objective,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=20,
    n_initial=10,
    repeats_initial=3, # Repeat initial points
    repeats_surrogate=2, # Repeat surrogate points
```

```

    seed=42,
    verbose=True
)

result = optimizer.optimize()

# Save results (includes noise statistics)
optimizer.save_result(prefix="noisy_opt")

# Load and analyze noise statistics
loaded_opt = SpotOptim.load_result("noisy_opt_res.pkl")

print(f"Repeats initial: {loaded_opt.repeats_initial}")
print(f"Repeats surrogate: {loaded_opt.repeats_surrogate}")
print(f"Best mean value: {loaded_opt.best_y_:.6f}")

if loaded_opt.mean_y is not None:
    print(f"Mean values available: {len(loaded_opt.mean_y)}")
    print(f"Variance values available: {len(loaded_opt.var_y)}")

```

```

TensorBoard logging disabled
Initial best: f(x) = 3.784490, mean best: f(x) = 3.897231
Experiment saved to noisy_opt_res.pkl
Result saved to noisy_opt_res.pkl
Loaded result from noisy_opt_res.pkl
Repeats initial: 3
Repeats surrogate: 2
Best mean value: 3.784490
Mean values available: 10
Variance values available: 10

```

16.6. Working with Different Variable Types

Save and load preserves variable type information:

```

import numpy as np
from spotoptim import SpotOptim

def mixed_objective(X):
    """Objective with mixed variable types"""
    return np.sum(X**2, axis=1)

```

16. Save and Load in SpotOptim

```
# Create optimizer with mixed variable types
optimizer = SpotOptim(
    fun=mixed_objective,
    bounds=[(-5, 5), (-5, 5), (-5, 5), (-5, 5)],
    var_type=["float", "int", "factor", "float"],
    var_name=["continuous", "integer", "categorical", "another_cont"],
    max_iter=20,
    n_initial=10,
    seed=42
)

result = optimizer.optimize()

# Save results
optimizer.save_result(prefix="mixed_vars")

# Load results
loaded_opt = SpotOptim.load_result("mixed_vars_res.pkl")

print("Variable types preserved:")
print(f"  var_type: {loaded_opt.var_type}")
print(f"  var_name: {loaded_opt.var_name}")

# Verify integer variables are still integers
print(f"\nInteger variable (dim 1) values:")
print(loaded_opt.X[:, 5, 1]) # Should be integers
```

```
Experiment saved to mixed_vars_res.pkl
Result saved to mixed_vars_res.pkl
Loaded result from mixed_vars_res.pkl
Variable types preserved:
  var_type: ['float', 'int', 'factor', 'float']
  var_name: ['continuous', 'integer', 'categorical', 'another_cont']

Integer variable (dim 1) values:
[-1.  1.  2.  3. -4.]
```

16.7. Best Practices

16.7.1. 1. Re-attaching the Objective Function (Optional)

Since `dill` is used for serialization, the objective function is automatically loaded. You only need to re-attach the function if:

1. You want to change the objective function (e.g., to a different implementation).
2. The function relies on external resources that cannot be serialized.

```
# Load experiment
optimizer = SpotOptim.load_experiment("experiment_exp.pkl")

# OPTIONAL: Replace the function if needed
# optimizer.fun = new_objective_function

# Run optimization
result = optimizer.optimize()
```

16.7.2. 2. Use Meaningful Prefixes

Organize your experiments with descriptive prefixes:

```
# Good practice: descriptive prefixes
optimizer.save_experiment(prefix="sphere_d10_seed42")
optimizer.save_experiment(prefix="rosenbrock_n100_lhs")
optimizer.save_result(prefix="final_run_2024_11_15")

# Avoid: generic names
optimizer.save_experiment(prefix="exp1") # Not descriptive
optimizer.save_result(prefix="result")  # Hard to track
```

16.7.3. 3. Save Experiments Before Remote Execution

```
# Define locally
optimizer = SpotOptim(bounds=bounds, max_iter=20, seed=42)
optimizer.save_experiment(prefix="remote_job")

# Transfer file to remote machine
# Execute remotely
```

16. Save and Load in SpotOptim

```
# Transfer results back
# Analyze locally
```

16.7.4. 4. Version Your Experiments

```
import datetime

# Add timestamp to prefix
timestamp = datetime.datetime.now().strftime("%Y%m%d_%H%M%S")
prefix = f"experiment_{timestamp}"

optimizer.save_experiment(prefix=prefix)
# Creates: experiment_20241115_143022_exp.pkl
```

16.7.5. 5. Handle File Paths Robustly

```
import os

# Create directory structure
exp_dir = "experiments/batch_001"
os.makedirs(exp_dir, exist_ok=True)

# Save with full path
optimizer.save_experiment(
    prefix="exp_001",
    path=exp_dir
)

# Load with full path
exp_file = os.path.join(exp_dir, "exp_001_exp.pkl")
loaded_opt = SpotOptim.load_experiment(exp_file)
```

16.8. Complete Example: Multi-Machine Workflow

Here's a complete example demonstrating the entire workflow:

16.8.1. Local Machine (Setup)

```

# setup_experiment.py
import numpy as np
from spotoptim import SpotOptim
import os

DIRNAME = "experiments_compare_de_tricands_bfgs"

# Placeholder function for experiment setup
def placeholder_func(X):
    return np.sum(X**2, axis=1)

# Create experiments directory
os.makedirs(DIRNAME, exist_ok=True)
# guarantee that the directory is empty
for file in os.listdir(DIRNAME):
    os.remove(os.path.join(DIRNAME, file))

# Define multiple experiments
experiments = [
    {"acquisition_optimizer": "differential_evolution", "max_iter": 20, "prefix": "exp_de"},
    {"acquisition_optimizer": "tricands", "max_iter": 20, "prefix": "exp_tricands"},
    {"acquisition_optimizer": "L-BFGS-B", "max_iter": 20, "prefix": "exp_bfgs"},
]

for exp_config in experiments:
    optimizer = SpotOptim(
        fun=placeholder_func, # Placeholder - will be replaced remotely
        bounds=[(-10, 10), (-10, 10), (-10, 10)],
        max_iter=exp_config["max_iter"],
        n_initial=10,
        seed=42,
        acquisition_optimizer=exp_config["acquisition_optimizer"],
        verbose=True
    )

    optimizer.save_experiment(
        prefix=exp_config["prefix"],
        path=DIRNAME
    )

```

16. Save and Load in SpotOptim

```
print(f"Created: {DIRNAME}/{exp_config['prefix']}_exp.pkl")  
print("\nAll experiments created. Transfer '{DIRNAME}' folder to remote machine.")
```

```
TensorBoard logging disabled  
Experiment saved to experiments_compare_de_tricands_bfgs/exp_de_exp.pkl  
Created: experiments_compare_de_tricands_bfgs/exp_de_exp.pkl  
TensorBoard logging disabled  
Experiment saved to experiments_compare_de_tricands_bfgs/exp_tricands_exp.pkl  
Created: experiments_compare_de_tricands_bfgs/exp_tricands_exp.pkl  
TensorBoard logging disabled  
Experiment saved to experiments_compare_de_tricands_bfgs/exp_bfgs_exp.pkl  
Created: experiments_compare_de_tricands_bfgs/exp_bfgs_exp.pkl
```

All experiments created. Transfer '{DIRNAME}' folder to remote machine.

16.8.2. Remote Machine (Execution)

```
# run_experiments.py  
import numpy as np  
from spotoptim import SpotOptim  
import os  
import glob  
  
def complex_objective(X):  
    """Complex multimodal objective function"""  
    term1 = np.sum(X**2, axis=1)  
    term2 = 10 * np.sum(np.cos(2 * np.pi * X), axis=1)  
    term3 = 0.1 * np.sum(np.sin(5 * np.pi * X), axis=1)  
    return term1 - term2 + term3  
  
# Find all experiment files  
exp_files = glob.glob(f"{DIRNAME}/*_exp.pkl")  
print(f"Found {len(exp_files)} experiments to run")  
  
# Run each experiment  
for exp_file in exp_files:  
    print(f"\nProcessing: {exp_file}")  
  
    # Load experiment  
    optimizer = SpotOptim.load_experiment(exp_file)
```

16.8. Complete Example: Multi-Machine Workflow

```
# Attach objective
optimizer.fun = complex_objective

# Run optimization
result = optimizer.optimize()
print(f" Best value: {result.fun:.6f}")

# Save result (same prefix, different suffix)
prefix = os.path.basename(exp_file).replace("_exp.pkl", "")
optimizer.save_result(
    prefix=prefix,
    path=DIRNAME
)
print(f" Saved: {DIRNAME}/{prefix}_res.pkl")

print("\nAll experiments completed. Transfer results back to local machine.")
```

Found 3 experiments to run

Processing: experiments_compare_de_tricands_bfgs/exp_de_exp.pkl
Loaded experiment from experiments_compare_de_tricands_bfgs/exp_de_exp.pkl
Initial best: $f(x) = 24.447486$
Iter 1 | Best: 24.447486 | Curr: 46.009409 | Rate: 0.00 | Evals: 55.0%
Optimizer candidate 1/3 was duplicate/invalid.
Iter 2 | Best: 24.447486 | Curr: 86.823725 | Rate: 0.00 | Evals: 60.0%
Iter 3 | Best: 24.447486 | Curr: 32.303496 | Rate: 0.00 | Evals: 65.0%
Iter 4 | Best: 24.447486 | Curr: 32.699441 | Rate: 0.00 | Evals: 70.0%
Iter 5 | Best: 24.447486 | Curr: 25.589825 | Rate: 0.00 | Evals: 75.0%
Iter 6 | Best: 9.027734 | Rate: 0.17 | Evals: 80.0%
Iter 7 | Best: 9.027734 | Curr: 25.531014 | Rate: 0.14 | Evals: 85.0%
Iter 8 | Best: 9.027734 | Curr: 11.243210 | Rate: 0.12 | Evals: 90.0%
Iter 9 | Best: 6.618538 | Rate: 0.22 | Evals: 95.0%
Iter 10 | Best: 6.618538 | Curr: 16.111614 | Rate: 0.20 | Evals: 100.0%
Best value: 6.618538
Experiment saved to experiments_compare_de_tricands_bfgs/exp_de_res.pkl
Result saved to experiments_compare_de_tricands_bfgs/exp_de_res.pkl
Saved: experiments_compare_de_tricands_bfgs/exp_de_res.pkl

Processing: experiments_compare_de_tricands_bfgs/exp_bfgs_exp.pkl
Loaded experiment from experiments_compare_de_tricands_bfgs/exp_bfgs_exp.pkl
Initial best: $f(x) = 24.447486$
Iter 1 | Best: 24.447486 | Curr: 115.305737 | Rate: 0.00 | Evals: 55.0%
Iter 2 | Best: 24.447486 | Curr: 73.901317 | Rate: 0.00 | Evals: 60.0%
Iter 3 | Best: 24.447486 | Curr: 25.511124 | Rate: 0.00 | Evals: 65.0%

16. Save and Load in SpotOptim

```
Iter 4 | Best: 24.447486 | Curr: 80.640861 | Rate: 0.00 | Evals: 70.0%
Iter 5 | Best: 24.447486 | Curr: 42.258181 | Rate: 0.00 | Evals: 75.0%
Iter 6 | Best: 24.447486 | Curr: 32.949358 | Rate: 0.00 | Evals: 80.0%
Iter 7 | Best: 15.174307 | Rate: 0.14 | Evals: 85.0%
Iter 8 | Best: 14.707912 | Rate: 0.25 | Evals: 90.0%
Iter 9 | Best: 14.707912 | Curr: 65.102882 | Rate: 0.22 | Evals: 95.0%
Iter 10 | Best: 14.707912 | Curr: 53.449070 | Rate: 0.20 | Evals: 100.0%
```

Best value: 14.707912

Experiment saved to experiments_compare_de_tricands_bfgs/exp_bfgs_res.pkl

Result saved to experiments_compare_de_tricands_bfgs/exp_bfgs_res.pkl

Saved: experiments_compare_de_tricands_bfgs/exp_bfgs_res.pkl

Processing: experiments_compare_de_tricands_bfgs/exp_tricands_exp.pkl

Loaded experiment from experiments_compare_de_tricands_bfgs/exp_tricands_exp.pkl

Initial best: $f(x) = 24.447486$

```
Iter 1 | Best: 24.447486 | Curr: 60.586280 | Rate: 0.00 | Evals: 55.0%
Iter 2 | Best: 20.457316 | Rate: 0.50 | Evals: 60.0%
Iter 3 | Best: 20.457316 | Curr: 44.287891 | Rate: 0.33 | Evals: 65.0%
Iter 4 | Best: 20.457316 | Curr: 54.350759 | Rate: 0.25 | Evals: 70.0%
Iter 5 | Best: 20.457316 | Curr: 29.059194 | Rate: 0.20 | Evals: 75.0%
Iter 6 | Best: 20.457316 | Curr: 38.314818 | Rate: 0.17 | Evals: 80.0%
Iter 7 | Best: 20.457316 | Curr: 22.880271 | Rate: 0.14 | Evals: 85.0%
Iter 8 | Best: 20.457316 | Curr: 47.526104 | Rate: 0.12 | Evals: 90.0%
Iter 9 | Best: 10.765448 | Rate: 0.22 | Evals: 95.0%
Iter 10 | Best: 10.765448 | Curr: 16.195813 | Rate: 0.20 | Evals: 100.0%
```

Best value: 10.765448

Experiment saved to experiments_compare_de_tricands_bfgs/exp_tricands_res.pkl

Result saved to experiments_compare_de_tricands_bfgs/exp_tricands_res.pkl

Saved: experiments_compare_de_tricands_bfgs/exp_tricands_res.pkl

All experiments completed. Transfer results back to local machine.

16.8.3. Local Machine (Analysis)

```
# analyze_results.py
import numpy as np
from spotoptim import SpotOptim
import glob
import matplotlib.pyplot as plt

# Find all result files
result_files = glob.glob(f"{DIRNAME}/*_res.pkl")
print(f"Found {len(result_files)} results to analyze")
```

```

# Load and compare results
results = []
for res_file in result_files:
    opt = SpotOptim.load_result(res_file)
    results.append({
        "file": res_file,
        "best_value": opt.best_y_,
        "best_point": opt.best_x_,
        "n_evals": opt.counter,
        "seed": opt.seed,
        "acquisition_optimizer": opt.acquisition_optimizer,
    })
    print(f"{res_file}: best={opt.best_y_:.6f}, evals={opt.counter}")

# Find best overall result
best = min(results, key=lambda x: x["best_value"])
print(f"\nBest result:")
print(f"  File: {best['file']}")
print(f"  Value: {best['best_value']:.6f}")
print(f"  Point: {best['best_point']}")
print(f"  Seed: {best['seed']}")
print(f"  Acquisition optimizer: {best['acquisition_optimizer']}")

# Plot convergence comparison
plt.figure(figsize=(12, 6))

for res_file in result_files:
    opt = SpotOptim.load_result(res_file)
    seed = opt.seed
    acquisition_optimizer = opt.acquisition_optimizer
    cummin = [opt.y_[i+1].min() for i in range(len(opt.y_))]
    plt.plot(cummin, label=f"{acquisition_optimizer}", linewidth=2, alpha=0.7)

plt.xlabel("Iteration", fontsize=12)
plt.ylabel("Best Value Found", fontsize=12)
plt.title("Optimization Progress Comparison", fontsize=14)
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.savefig(f"{DIRNAME}/convergence_comparison.png", dpi=150)
print("\nConvergence plot saved to: {DIRNAME}/convergence_comparison.png")

```

Found 3 results to analyze

16. Save and Load in SpotOptim

```
Loaded result from experiments_compare_de_tricands_bfgs/exp_bfgs_res.pkl
experiments_compare_de_tricands_bfgs/exp_bfgs_res.pkl: best=14.707912, evals=20
Loaded result from experiments_compare_de_tricands_bfgs/exp_de_res.pkl
experiments_compare_de_tricands_bfgs/exp_de_res.pkl: best=6.618538, evals=20
Loaded result from experiments_compare_de_tricands_bfgs/exp_tricands_res.pkl
experiments_compare_de_tricands_bfgs/exp_tricands_res.pkl: best=10.765448, evals=20
```

Best result:

File: experiments_compare_de_tricands_bfgs/exp_de_res.pkl

Value: 6.618538

Point: [-3.72183897 -1.94716806 1.86706201]

Seed: 42


Acquisition optimizer: differential_evolution

Loaded result from experiments_compare_de_tricands_bfgs/exp_bfgs_res.pkl

Loaded result from experiments_compare_de_tricands_bfgs/exp_de_res.pkl

Loaded result from experiments_compare_de_tricands_bfgs/exp_tricands_res.pkl

Convergence plot saved to: {DIRNAME}/convergence_comparison.png



16.9. Technical Details

16.9.1. Serialization Method

SpotOptim uses Python's built-in `pickle` module for serialization. This provides:

- **Standard library:** No additional dependencies required

- **Compatibility:** Works with numpy arrays, sklearn models, scipy functions
- **Performance:** Efficient serialization of large datasets

16.9.2. Component Reinitialization

When loading experiments, certain components are automatically recreated:

- **Surrogate model:** Gaussian Process with default kernel
- **LHS sampler:** Latin Hypercube Sampler with original seed

This ensures loaded experiments can continue optimization without manual configuration.

16.9.3. Excluded Components

Some components cannot be pickled and are automatically excluded:

- **Objective function (fun):** Included (via dill) in both experiment and result files.
- **TensorBoard writer (tb_writer):** File handles cannot be serialized
- **Surrogate model (experiments only):** Recreated on load for experiments

16.9.4. File Format

Files are saved using pickle's highest protocol:

```
with open(filename, "wb") as handle:
    pickle.dump(optimizer_state, handle, protocol=pickle.HIGHEST_PROTOCOL)
```

16.10. Troubleshooting

16.10.1. Issue: “AttributeError: ‘SpotOptim’ object has no attribute ‘fun’”

This should not happen if dill is working correctly. If it does, your function might not be picklable. Try defining the function at the top level of your script or module. If all else fails, you can manually re-attach it:

16. Save and Load in SpotOptim

```
#| eval: false
opt = SpotOptim.load_experiment("exp.pkl")
opt.fun = my_objective_function # Manually re-attach
result = opt.optimize()
```

16.10.2. Issue: “FileNotFoundError: Experiment file not found”

Cause: Incorrect file path or file doesn't exist.

Solution: Check file path and ensure file exists:

```
import os

filename = "experiment_exp.pkl"
if os.path.exists(filename):
    opt = SpotOptim.load_experiment(filename)
else:
    print(f"File not found: {filename}")
```

16.10.3. Issue: “FileExistsError: File already exists”

Cause: Attempting to save over an existing file without `overwrite=True`.

Solution: Either use a different prefix or enable overwriting:

```
# Option 1: Use different prefix
optimizer.save_result(prefix="my_result_v2")

# Option 2: Enable overwriting
optimizer.save_result(prefix="my_result", overwrite=True)
```

16.10.4. Issue: Results differ after loading

Cause: Random state not preserved or function behavior changed.

Solution: Ensure you're using the same seed and function definition:

```
# When saving
optimizer = SpotOptim(..., seed=42) # Use fixed seed

# When loading and continuing
```



```
loaded_opt = SpotOptim.load_result("result_res.pkl")  
# loaded_opt.fun is already attached
```

16.11. Jupyter Notebook

i Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

17. Reproducibility in SpotOptim

17.1. Introduction

SpotOptim provides full support for reproducible optimization runs through the `seed` parameter. This is essential for:

- **Scientific research:** Ensuring experiments can be replicated
- **Debugging:** Reproducing specific optimization behaviors
- **Benchmarking:** Fair comparison between different configurations
- **Production:** Consistent results in deployed applications

When you specify a seed, SpotOptim guarantees that running the same optimization multiple times will produce identical results. Without a seed, each run explores the search space differently, which can be useful for robustness testing.

17.2. Basic Usage

17.2.1. Making Optimization Reproducible

To ensure reproducible results, simply specify the `seed` parameter when creating the optimizer:

```
import numpy as np
from spotoptim import SpotOptim

def sphere(X):
    """Simple sphere function:  $f(x) = \sum(x^2)$ """
    return np.sum(X**2, axis=1)

# Reproducible optimization
optimizer = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=30,
    n_initial=15,
```

17. Reproducibility in SpotOptim

```
seed=42, # This ensures reproducibility
verbose=True
)

result = optimizer.optimize()
print(f"Best solution: {result.x}")
print(f"Best value: {result.fun}")
```

```
TensorBoard logging disabled
Initial best: f(x) = 0.049607
Iter 1 | Best: 0.000052 | Rate: 1.00 | Evals: 53.3%
Iter 2 | Best: 0.000016 | Rate: 1.00 | Evals: 56.7%
Iter 3 | Best: 0.000003 | Rate: 1.00 | Evals: 60.0%
Iter 4 | Best: 0.000000 | Rate: 1.00 | Evals: 63.3%
Iter 5 | Best: 0.000000 | Rate: 1.00 | Evals: 66.7%
Iter 6 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.83 | Evals: 70.0%
Iter 7 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.71 | Evals: 73.3%
Iter 8 | Best: 0.000000 | Rate: 0.75 | Evals: 76.7%
Iter 9 | Best: 0.000000 | Rate: 0.78 | Evals: 80.0%
Iter 10 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.70 | Evals: 83.3%
Iter 11 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.64 | Evals: 86.7%
Iter 12 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.58 | Evals: 90.0%
Iter 13 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.54 | Evals: 93.3%
Iter 14 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.50 | Evals: 96.7%
Iter 15 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.47 | Evals: 100.0%
Best solution: [-4.39940747e-05 -2.08039138e-04]
Best value: 4.521576166022012e-08
```

Key Point: Running this code multiple times (even on different days or machines) will always produce the same result.

17.2.2. Running Independent Experiments

If you don't specify a seed, each optimization run will explore the search space differently:

```
# Non-reproducible: different results each time
optimizer = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=30,
    n_initial=15
```

```

    # No seed specified
)

result = optimizer.optimize()
# Results will vary between runs

```

This is useful when you want to: - Explore different regions of the search space - Test the robustness of your results - Run multiple independent optimization attempts

17.3. Practical Examples

17.3.1. Example 1: Comparing Different Configurations

When comparing different optimizer settings, use the same seed for fair comparison:

```

import numpy as np
from spotoptim import SpotOptim

def rosenbrock(X):
    """Rosenbrock function"""
    x = X[:, 0]
    y = X[:, 1]
    return (1 - x)**2 + 100 * (y - x**2)**2

# Configuration 1: More initial points
opt1 = SpotOptim(
    fun=rosenbrock,
    bounds=[(-2, 2), (-2, 2)],
    max_iter=50,
    n_initial=20,
    seed=42 # Same seed for fair comparison
)
result1 = opt1.optimize()

# Configuration 2: Fewer initial points, more iterations
opt2 = SpotOptim(
    fun=rosenbrock,
    bounds=[(-2, 2), (-2, 2)],
    max_iter=50,
    n_initial=10,
    seed=42 # Same seed
)

```

17. Reproducibility in SpotOptim

```
result2 = opt2.optimize()

print(f"Config 1 (more initial): {result1.fun:.6f}")
print(f"Config 2 (fewer initial): {result2.fun:.6f}")
```

```
Config 1 (more initial): 0.000220
Config 2 (fewer initial): 0.054879
```

17.3.2. Example 2: Reproducible Research Experiment

For scientific papers or reports, always use a fixed seed and document it:

```
import numpy as np
from spotoptim import SpotOptim

def rastrigin(X):
    """Rastrigin function (multimodal)"""
    A = 10
    n = X.shape[1]
    return A * n + np.sum(X**2 - A * np.cos(2 * np.pi * X), axis=1)

# Documented seed for reproducibility
RANDOM_SEED = 12345

optimizer = SpotOptim(
    fun=rastrigin,
    bounds=[(-5.12, 5.12), (-5.12, 5.12), (-5.12, 5.12)],
    max_iter=100,
    n_initial=30,
    seed=RANDOM_SEED,
    verbose=True
)

result = optimizer.optimize()

print(f"\nExperiment Results (seed={RANDOM_SEED}):")
print(f"Best solution: {result.x}")
print(f"Best value: {result.fun}")
print(f"Iterations: {result.nit}")
print(f"Function evaluations: {result.nfev}")

# These results can now be cited in a paper
```

17.3. Practical Examples

TensorBoard logging disabled

Initial best: $f(x) = 5.256150$

Iter 1	Best: 5.256150	Curr: 8.980822	Rate: 0.00	Evals: 31.0%
Iter 2	Best: 5.256150	Curr: 16.506647	Rate: 0.00	Evals: 32.0%
Iter 3	Best: 5.256150	Curr: 26.407455	Rate: 0.00	Evals: 33.0%
Iter 4	Best: 5.256150	Curr: 5.573478	Rate: 0.00	Evals: 34.0%
Iter 5	Best: 4.391963	Rate: 0.20	Evals: 35.0%	
Iter 6	Best: 4.365740	Rate: 0.33	Evals: 36.0%	
Iter 7	Best: 4.365740	Curr: 23.387719	Rate: 0.29	Evals: 37.0%
Iter 8	Best: 4.365740	Curr: 26.955948	Rate: 0.25	Evals: 38.0%
Iter 9	Best: 4.141973	Rate: 0.33	Evals: 39.0%	
Iter 10	Best: 3.393072	Rate: 0.40	Evals: 40.0%	
Iter 11	Best: 3.073942	Rate: 0.45	Evals: 41.0%	
Iter 12	Best: 3.073942	Curr: 27.948277	Rate: 0.42	Evals: 42.0%
Iter 13	Best: 3.068052	Rate: 0.46	Evals: 43.0%	
Iter 14	Best: 3.013474	Rate: 0.50	Evals: 44.0%	
Iter 15	Best: 2.992072	Rate: 0.53	Evals: 45.0%	
Iter 16	Best: 2.985140	Rate: 0.56	Evals: 46.0%	
Iter 17	Best: 2.985140	Curr: 19.619123	Rate: 0.53	Evals: 47.0%
Iter 18	Best: 2.984970	Rate: 0.56	Evals: 48.0%	
Iter 19	Best: 2.984970	Curr: 2.984973	Rate: 0.53	Evals: 49.0%
Iter 20	Best: 2.984964	Rate: 0.55	Evals: 50.0%	
Iter 21	Best: 2.984963	Rate: 0.57	Evals: 51.0%	
Iter 22	Best: 2.984963	Curr: 28.577831	Rate: 0.55	Evals: 52.0%
Iter 23	Best: 2.984963	Curr: 2.984967	Rate: 0.52	Evals: 53.0%
Iter 24	Best: 2.984963	Curr: 2.984966	Rate: 0.50	Evals: 54.0%
Iter 25	Best: 2.984963	Curr: 2.984968	Rate: 0.48	Evals: 55.0%
Iter 26	Best: 2.984963	Curr: 28.565596	Rate: 0.46	Evals: 56.0%
Iter 27	Best: 2.984963	Curr: 31.747890	Rate: 0.44	Evals: 57.0%
Iter 28	Best: 2.984960	Rate: 0.46	Evals: 58.0%	
Iter 29	Best: 2.984958	Rate: 0.48	Evals: 59.0%	
Iter 30	Best: 2.984958	Curr: 86.774141	Rate: 0.47	Evals: 60.0%
Iter 31	Best: 2.984957	Rate: 0.48	Evals: 61.0%	
Iter 32	Best: 2.984957	Curr: 86.774141	Rate: 0.47	Evals: 62.0%
Iter 33	Best: 2.984945	Rate: 0.48	Evals: 63.0%	
Iter 34	Best: 2.984945	Curr: 67.151312	Rate: 0.47	Evals: 64.0%
Iter 35	Best: 2.984932	Rate: 0.49	Evals: 65.0%	
Iter 36	Best: 2.984913	Rate: 0.50	Evals: 66.0%	
Iter 37	Best: 2.984913	Curr: 2.986002	Rate: 0.49	Evals: 67.0%
Iter 38	Best: 2.984878	Rate: 0.50	Evals: 68.0%	
Iter 39	Best: 2.984878	Curr: 68.868074	Rate: 0.49	Evals: 69.0%
Iter 40	Best: 2.984878	Rate: 0.50	Evals: 70.0%	
Iter 41	Best: 2.984878	Curr: 2.987673	Rate: 0.49	Evals: 71.0%
Iter 42	Best: 2.984878	Curr: 66.972659	Rate: 0.48	Evals: 72.0%
Iter 43	Best: 2.984878	Curr: 2.984879	Rate: 0.47	Evals: 73.0%

17. Reproducibility in SpotOptim

Iter 44	Best: 2.984878	Curr: 2.984879	Rate: 0.45	Evals: 74.0%
Iter 45	Best: 2.984878	Curr: 2.984879	Rate: 0.44	Evals: 75.0%
Iter 46	Best: 2.984878	Curr: 68.668796	Rate: 0.43	Evals: 76.0%
Iter 47	Best: 2.984878	Curr: 19.507980	Rate: 0.43	Evals: 77.0%
Iter 48	Best: 2.984878	Curr: 2.984879	Rate: 0.42	Evals: 78.0%
Iter 49	Best: 2.984878	Curr: 2.984879	Rate: 0.41	Evals: 79.0%
Iter 50	Best: 2.984878	Curr: 2.984880	Rate: 0.40	Evals: 80.0%
Iter 51	Best: 2.984878	Curr: 22.068452	Rate: 0.39	Evals: 81.0%
Iter 52	Best: 2.984878	Curr: 2.984879	Rate: 0.38	Evals: 82.0%
Iter 53	Best: 2.984878	Curr: 2.984879	Rate: 0.38	Evals: 83.0%
Iter 54	Best: 2.984878	Curr: 2.984879	Rate: 0.37	Evals: 84.0%
Iter 55	Best: 2.984878	Curr: 29.565839	Rate: 0.36	Evals: 85.0%
Iter 56	Best: 2.984878	Curr: 44.114074	Rate: 0.36	Evals: 86.0%
Iter 57	Best: 2.984878	Curr: 2.984879	Rate: 0.35	Evals: 87.0%
Iter 58	Best: 2.984878	Curr: 2.984879	Rate: 0.34	Evals: 88.0%
Iter 59	Best: 2.984878	Curr: 73.017405	Rate: 0.34	Evals: 89.0%
Iter 60	Best: 2.984878	Curr: 2.984881	Rate: 0.33	Evals: 90.0%
Iter 61	Best: 2.984878	Curr: 2.984879	Rate: 0.33	Evals: 91.0%
Iter 62	Best: 2.984878	Curr: 2.984879	Rate: 0.32	Evals: 92.0%
Iter 63	Best: 2.984878	Curr: 2.984879	Rate: 0.32	Evals: 93.0%
Iter 64	Best: 2.984878	Curr: 46.307440	Rate: 0.31	Evals: 94.0%
Iter 65	Best: 2.984878	Curr: 2.984878	Rate: 0.31	Evals: 95.0%
Iter 66	Best: 2.984878	Curr: 2.984879	Rate: 0.30	Evals: 96.0%
Iter 67	Best: 2.984878	Curr: 2.984879	Rate: 0.30	Evals: 97.0%
Iter 68	Best: 2.984878	Curr: 2.984880	Rate: 0.29	Evals: 98.0%
Iter 69	Best: 2.984878	Curr: 2.984879	Rate: 0.29	Evals: 99.0%
Iter 70	Best: 2.984878	Curr: 86.774141	Rate: 0.29	Evals: 100.0%

Experiment Results (seed=12345):

Best solution: [0.99495067 -0.99492308 -0.9949319]

Best value: 2.9848775762946076

Iterations: 70

Function evaluations: 100

17.3.3. Example 3: Multiple Independent Runs

To test robustness, run the same optimization with different seeds:

```
import numpy as np
from spotoptim import SpotOptim

def ackley(X):
    """Ackley function"""
    a = 20
```



```

b = 0.2
c = 2 * np.pi
n = X.shape[1]

sum_sq = np.sum(X**2, axis=1)
sum_cos = np.sum(np.cos(c * X), axis=1)

return -a * np.exp(-b * np.sqrt(sum_sq / n)) - np.exp(sum_cos / n) + a + np.e

# Run 5 independent optimizations
results = []
seeds = [42, 123, 456, 789, 1011]

for seed in seeds:
    optimizer = SpotOptim(
        fun=ackley,
        bounds=[(-5, 5), (-5, 5)],
        max_iter=40,
        n_initial=20,
        seed=seed,
        verbose=False
    )
    result = optimizer.optimize()
    results.append(result.fun)
    print(f"Run with seed {seed:4d}: f(x) = {result.fun:.6f}")

# Analyze robustness
print(f"\nBest result: {min(results):.6f}")
print(f"Worst result: {max(results):.6f}")
print(f"Mean: {np.mean(results):.6f}")
print(f"Std dev: {np.std(results):.6f}")

```

```

Run with seed 42: f(x) = 0.000371
Run with seed 123: f(x) = 0.000152
Run with seed 456: f(x) = 0.000095
Run with seed 789: f(x) = 0.002216
Run with seed 1011: f(x) = 0.000406

```

```

Best result: 0.000095
Worst result: 0.002216
Mean: 0.000648
Std dev: 0.000793

```

17.3.4. Example 4: Reproducible Initial Design

The seed ensures that even the initial design points are reproducible:

```
import numpy as np
from spotoptim import SpotOptim

def simple_quadratic(X):
    return np.sum((X - 1)**2, axis=1)

# Create two optimizers with same seed
opt1 = SpotOptim(
    fun=simple_quadratic,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=25,
    n_initial=10,
    seed=999
)

opt2 = SpotOptim(
    fun=simple_quadratic,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=25,
    n_initial=10,
    seed=999 # Same seed
)

# Run both optimizations
result1 = opt1.optimize()
result2 = opt2.optimize()

# Verify identical results
print("Initial design points are identical:",
      np.allclose(opt1.X_[:10], opt2.X_[:10]))
print("All evaluated points are identical:",
      np.allclose(opt1.X_, opt2.X_))
print("All function values are identical:",
      np.allclose(opt1.y_, opt2.y_))
print("Best solutions are identical:",
      np.allclose(result1.x, result2.x))
```

```
Initial design points are identical: True
All evaluated points are identical: True
All function values are identical: True
Best solutions are identical: True
```

17.3.5. Example 5: Custom Initial Design with Seed

Even when providing a custom initial design, the seed ensures reproducible subsequent iterations:

```
import numpy as np
from spotoptim import SpotOptim

def beale(X):
    """Beale function"""
    x = X[:, 0]
    y = X[:, 1]
    term1 = (1.5 - x + x * y)**2
    term2 = (2.25 - x + x * y**2)**2
    term3 = (2.625 - x + x * y**3)**2
    return term1 + term2 + term3

# Custom initial design (e.g., from previous knowledge)
X_start = np.array([
    [0.0, 0.0],
    [1.0, 1.0],
    [2.0, 2.0],
    [-1.0, -1.0]
])

# Run twice with same seed and initial design
opt1 = SpotOptim(
    fun=beale,
    bounds=[(-4.5, 4.5), (-4.5, 4.5)],
    max_iter=30,
    n_initial=10,
    seed=777
)
result1 = opt1.optimize(X0=X_start)

opt2 = SpotOptim(
    fun=beale,
    bounds=[(-4.5, 4.5), (-4.5, 4.5)],
    max_iter=30,
    n_initial=10,
    seed=777 # Same seed
)
result2 = opt2.optimize(X0=X_start)
```

17. Reproducibility in SpotOptim

```
print("Results are identical:", np.allclose(result1.x, result2.x))
print(f"Best value: {result1.fun:.6f}")
```

```
Results are identical: True
Best value: 1.347608
```

17.4. Advanced Topics

17.4.1. Seed and Noisy Functions

When optimizing noisy functions with repeated evaluations, the seed ensures reproducible noise:

```
import numpy as np
from spotoptim import SpotOptim

def noisy_sphere(X):
    """Sphere function with Gaussian noise"""
    base = np.sum(X**2, axis=1)
    noise = np.random.normal(0, 0.1, size=base.shape)
    return base + noise

optimizer = SpotOptim(
    fun=noisy_sphere,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=40,
    n_initial=20,
    repeats_initial=3, # 3 evaluations per point
    repeats_surrogate=2,
    seed=42 # Ensures same noise pattern
)

result = optimizer.optimize()
print(f"Best mean value: {optimizer.min_mean_y:.6f}")
print(f"Variance at best: {optimizer.min_var_y:.6f}")
```

```
Best mean value: 0.093359
Variance at best: 0.010943
```

Important: With the same seed, even the noise will be identical across runs!

17.4.2. Different Seeds for Different Exploration

Use different seeds to explore different regions systematically:

```
import numpy as np
from spotoptim import SpotOptim

def griewank(X):
    """Griewank function"""
    sum_sq = np.sum(X**2 / 4000, axis=1)
    prod_cos = np.prod(np.cos(X / np.sqrt(np.arange(1, X.shape[1] + 1))), axis=1)
    return sum_sq - prod_cos + 1

# Systematic exploration with different seeds
best_overall = float('inf')
best_seed = None

for seed in range(10, 20): # Seeds 10-19
    optimizer = SpotOptim(
        fun=griewank,
        bounds=[(-600, 600), (-600, 600)],
        max_iter=50,
        n_initial=25,
        seed=seed
    )
    result = optimizer.optimize()

    if result.fun < best_overall:
        best_overall = result.fun
        best_seed = seed

    print(f"Seed {seed}: f(x) = {result.fun:.6f}")

print(f"\nBest result with seed {best_seed}: {best_overall:.6f}")
```

```
Seed 10: f(x) = 1.360689
Seed 11: f(x) = 0.331196
Seed 12: f(x) = 1.615590
Seed 13: f(x) = 3.986211
Seed 14: f(x) = 2.311520
Seed 15: f(x) = 1.674165
Seed 16: f(x) = 1.410071
Seed 17: f(x) = 0.535659
Seed 18: f(x) = 1.321237
```

17. Reproducibility in SpotOptim

Seed 19: $f(x) = 4.052091$

Best result with seed 11: 0.331196

17.5. Best Practices

17.5.1. 1. Always Use Seeds for Production Code

```
# Good: Reproducible
optimizer = SpotOptim(fun=objective, bounds=bounds, seed=42)

# Risky: Non-reproducible
optimizer = SpotOptim(fun=objective, bounds=bounds)
```

17.5.2. 2. Document Your Seeds

```
# Configuration for experiment reported in Section 4.2
EXPERIMENT_SEED = 2024
MAX_ITERATIONS = 100

optimizer = SpotOptim(
    fun=my_objective,
    bounds=my_bounds,
    max_iter=MAX_ITERATIONS,
    seed=EXPERIMENT_SEED
)
```

17.5.3. 3. Use Different Seeds for Different Experiments

```
# Different experiments should use different seeds
BASELINE_SEED = 100
EXPERIMENT_A_SEED = 200
EXPERIMENT_B_SEED = 300
```

17.5.4. 4. Test Robustness Across Multiple Seeds

```
# Run same optimization with multiple seeds
for seed in [42, 123, 456, 789, 1011]:
    optimizer = SpotOptim(fun=objective, bounds=bounds, seed=seed)
    result = optimizer.optimize()
# Analyze results
```

17.6. What the Seed Controls

The `seed` parameter ensures reproducibility by controlling:

1. **Initial Design Generation:** Latin Hypercube Sampling produces the same initial points
2. **Surrogate Model:** Gaussian Process random initialization is identical
3. **Acquisition Optimization:** Differential evolution explores the same candidates
4. **Random Sampling:** Any random exploration uses the same random numbers

This guarantees that the entire optimization pipeline is deterministic and reproducible.

17.7. Common Questions

Q: Can I use `seed=0`?

A: Yes, any integer (including 0) is a valid seed.

Q: Will different Python versions give the same results?

A: Generally yes, but minor numerical differences may occur due to underlying library changes. Use the same environment for exact reproducibility.

Q: Does the seed affect the objective function?

A: No, the seed only affects SpotOptim's internal random processes. If your objective function has its own randomness, you'll need to control that separately.

Q: How do I choose a good seed value?

A: Any integer works. Common choices are 42, 123, or dates (e.g., 20241112). What matters is consistency, not the specific value.

17.8. Summary

- Use `seed` parameter for reproducible optimization
- Same seed → identical results (every time)
- No seed → different results (random exploration)
- Essential for research, debugging, and production
- Document your seeds for transparency
- Test robustness with multiple different seeds

17.9. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

Part V.

Surrogate Handling

18. Surrogate Model Selection in SpotOptim

This section demonstrates how to select and configure different surrogate models in SpotOptim. We compare various surrogate options:

- **Gaussian Process** with different kernels (Matern, RBF, Rational Quadratic)
- **SpotOptim Kriging** model
- **Random Forest** regressor
- **XGBoost** regressor
- **Support Vector Regression (SVR)**
- **Gradient Boosting** regressor

All methods are evaluated on the Aircraft Wing Weight Example (AWWE) function. We visualize the fitted surrogates and compare optimization performance.

18.1. Introduction

SpotOptim supports any scikit-learn compatible regressor, allowing you to choose the surrogate that best fits your problem characteristics:

- **Gaussian Processes:** Provide uncertainty estimates, smooth interpolation, good for continuous functions. Support Expected Improvement (EI) acquisition.
- **Kriging:** Similar to GP but with customizable correlation functions. Supports EI acquisition.
- **Random Forests:** Robust to noise, handle discontinuities. Don't provide uncertainty, so use `acquisition='y'` (greedy).
- **XGBoost:** Excellent for high-dimensional problems, fast training and prediction. Use `acquisition='y'`.
- **SVR:** Good for high-dimensional problems with smooth structure. Use `acquisition='y'`.
- **Gradient Boosting:** Strong performance on structured problems. Use `acquisition='y'`.

18.1.1. Acquisition Functions and Uncertainty

Models that provide uncertainty estimates (Gaussian Process, Kriging) work with all acquisition functions: 'ei' (Expected Improvement), 'pi' (Probability of Improvement), and 'y' (greedy).

Tree-based and other models (Random Forest, XGBoost, SVR, Gradient Boosting) don't provide uncertainty estimates by default, so they should use `acquisition='y'` for greedy optimization. SpotOptim automatically handles this gracefully.

18.2. Setup and Imports

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from spotoptim import SpotOptim
from spotoptim.surrogate import Kriging
import time
import warnings
warnings.filterwarnings('ignore')

# Scikit-learn models
from sklearn.gaussian_process import GaussianProcessRegressor
from sklearn.gaussian_process.kernels import (
    Matern, RBF, ConstantKernel, WhiteKernel, RationalQuadratic
)
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.svm import SVR
```

- XGBoost (if available)

```
import os
os.environ["OMP_NUM_THREADS"] = "1" # Prevent OpenMP conflict on macOS

try:
    import xgboost as xgb
    XGBOOST_AVAILABLE = True
except ImportError:
    XGBOOST_AVAILABLE = False
    print("XGBoost not available. Install with: pip install xgboost")

print("Libraries imported successfully!")
```

XGBoost not available. Install with: `pip install xgboost`
 Libraries imported successfully!

18.3. The AWE Objective Function

We use the Aircraft Wing Weight Example function, which models the weight of an unpainted light aircraft wing. The function accepts inputs in the unit cube $[0, 1]^9$ and returns the wing weight in pounds.

```
from spotoptim.function.so import wingwt

# Problem setup
bounds = [(0, 1)] * 9
param_names = ['Sw', 'Wfw', 'A', 'L', 'q', 'l', 'Rtc', 'Nz', 'Wdg']
max_iter = 20
n_initial = 15
seed = 42

print(f"Problem dimension: {len(bounds)}")
print(f"Optimization budget: {max_iter} evaluations")
```

Problem dimension: 9
 Optimization budget: 20 evaluations

18.4. 1. Default Surrogate: Gaussian Process with Matern Kernel

SpotOptim's default surrogate is a Gaussian Process with a Matern kernel ($\nu=2.5$), which provides twice-differentiable sample paths and good performance for most optimization problems.

```
start_time = time.time()

# Default GP (no surrogate specified)
optimizer_default = SpotOptim(
    fun=wingwt,
    bounds=bounds,
    max_iter=max_iter,
    n_initial=n_initial,
    var_name=param_names,
```

18. Surrogate Model Selection in SpotOptim

```
        acquisition='ei',
        seed=seed,
        verbose=False
    )

    result_default = optimizer_default.optimize()
    time_default = time.time() - start_time

    print(f"\nResults:")
    print(f"    Best weight: {result_default.fun:.4f} lb")
    print(f"    Function evaluations: {result_default.nfev}")
    print(f"    Time: {time_default:.2f}s")
    print(f"    Success: {result_default.success}")


    # Store for comparison
    results_comparison = [{
        'Surrogate': 'GP Matern nu=2.5 (Default)',
        'Best Weight': result_default.fun,
        'Evaluations': result_default.nfev,
        'Time (s)': time_default,
        'Success': result_default.success
    }]
```

```
Results:
    Best weight: 119.8666 lb
    Function evaluations: 20
    Time: 1.04s
    Success: True
```

18.4.1. Visualization: Default Surrogate

```
# Plot convergence
optimizer_default.plot_progress(log_y=False, figsize=(10, 5))
```

18.4. 1. Default Surrogate: Gaussian Process with Matern Kernel



```
# Plot most important hyperparameters
optimizer_default.plot_important_hyperparameter_contour(max_imp=3)
plt.suptitle('Default GP Matern nu=2.5: Most Important Parameters', y=1.02)
plt.show()
```


Plotting surrogate contours for top 3 most important parameters:

Nz: importance = 22.46% (type: float)
A: importance = 20.27% (type: float)
Rtc: importance = 17.83% (type: float)

Generating 3 surrogate plots...


Plotting Nz vs A

18. Surrogate Model Selection in SpotOptim




Plotting N_z vs R_{tc}

18.4. 1. Default Surrogate: Gaussian Process with Matern Kernel



Plotting A vs R_{tc}

18. Surrogate Model Selection in SpotOptim



<Figure size 1650x1050 with 0 Axes>

18.5. 2. Gaussian Process with RBF (Radial Basis Function) Kernel

The RBF kernel (also called squared exponential) produces infinitely differentiable sample paths, resulting in very smooth predictions.

```
start_time = time.time()

# Configure GP with RBF kernel
kernel_rbf = ConstantKernel(1.0, (1e-3, 1e3)) * RBF(
    length_scale=1.0,
    length_scale_bounds=(1e-2, 1e2)
)
```


18.5. 2. Gaussian Process with RBF (Radial Basis Function) Kernel

```
gp_rbf = GaussianProcessRegressor(  
    kernel=kernel_rbf,  
    n_restarts_optimizer=10,  
    normalize_y=True,  
    random_state=seed  
)  
  
optimizer_rbf = SpotOptim(  
    fun=wingwt,  
    bounds=bounds,  
    surrogate=gp_rbf,  
    max_iter=max_iter,  
    n_initial=n_initial,  
    var_name=param_names,  
    acquisition='ei',  
    seed=seed,  
    verbose=False  
)  
  
result_rbf = optimizer_rbf.optimize()  
time_rbf = time.time() - start_time  
  
print(f"\nResults:")  
print(f"  Best weight: {result_rbf.fun:.4f} lb")  
print(f"  Function evaluations: {result_rbf.nfev}")  
print(f"  Time: {time_rbf:.2f}s")  
print(f"  Success: {result_rbf.success}")  
  
results_comparison.append({  
    'Surrogate': 'GP RBF',  
    'Best Weight': result_rbf.fun,  
    'Evaluations': result_rbf.nfev,  
    'Time (s)': time_rbf,  
    'Success': result_rbf.success  
})
```

Results:
Best weight: 120.1093 lb
Function evaluations: 20
Time: 0.36s
Success: True

18.5.1. Visualization: RBF Kernel

```
optimizer_rbf.plot_progress(log_y=False, figsize=(10, 5))
```



```
optimizer_rbf.plot_important_hyperparameter_contour(max_imp=3)
plt.suptitle('GP RBF Kernel: Most Important Parameters', y=1.02)
plt.show()
```


Plotting surrogate contours for top 3 most important parameters:

Nz: importance = 22.94% (type: float)
A: importance = 20.71% (type: float)
Rtc: importance = 18.21% (type: float)

Generating 3 surrogate plots...


Plotting Nz vs A

18.5. 2. Gaussian Process with RBF (Radial Basis Function) Kernel




Plotting N_z vs R_{tc}

18. Surrogate Model Selection in SpotOptim



Plotting A vs R_{tc}

18.6. 3. Gaussian Process with Matern nu=1.5 Kernel



<Figure size 1650x1050 with 0 Axes>

18.6. 3. Gaussian Process with Matern nu=1.5 Kernel

The Matern nu=1.5 kernel produces once-differentiable sample paths, allowing for more flexible (less smooth) fits than nu=2.5.

```
start_time = time.time()

# Configure GP with Matern nu=1.5
kernel_matern15 = ConstantKernel(1.0, (1e-3, 1e3)) * Matern(
    length_scale=1.0,
    length_scale_bounds=(1e-2, 1e2),
    nu=1.5
)
```


18. Surrogate Model Selection in SpotOptim

```
gp_matern15 = GaussianProcessRegressor(  
    kernel=kernel_matern15,  
    n_restarts_optimizer=10,  
    normalize_y=True,  
    random_state=seed  
)  
  
optimizer_matern15 = SpotOptim(  
    fun=wingwt,  
    bounds=bounds,  
    surrogate=gp_matern15,  
    max_iter=max_iter,  
    n_initial=n_initial,  
    var_name=param_names,  
    acquisition='ei',  
    seed=seed,  
    verbose=False  
)  
  
result_matern15 = optimizer_matern15.optimize()  
time_matern15 = time.time() - start_time  
  
print(f"\nResults:")  
print(f"   Best weight: {result_matern15.fun:.4f} lb")  
print(f"   Function evaluations: {result_matern15.nfev}")  
print(f"   Time: {time_matern15:.2f}s")  
print(f"   Success: {result_matern15.success}")  
  
results_comparison.append({  
    'Surrogate': 'GP Matern nu=1.5',  
    'Best Weight': result_matern15.fun,  
    'Evaluations': result_matern15.nfev,  
    'Time (s)': time_matern15,  
    'Success': result_matern15.success  
})
```

```
Results:  
   Best weight: 119.5466 lb  
   Function evaluations: 20  
   Time: 0.35s  
   Success: True
```


18.6.1. Visualization: Matern $\nu=1.5$ Kernel

```
optimizer_matern15.plot_progress(log_y=False, figsize=(10, 5))
```



```
optimizer_matern15.plot_important_hyperparameter_contour(max_imp=3)
plt.suptitle('GP Matern nu=1.5 Kernel: Most Important Parameters', y=1.02)
plt.show()
```


Plotting surrogate contours for top 3 most important parameters:

Nz: importance = 22.42% (type: float)
A: importance = 20.23% (type: float)
Rtc: importance = 17.79% (type: float)

Generating 3 surrogate plots...


Plotting Nz vs A

18. Surrogate Model Selection in SpotOptim




Plotting N_z vs R_{tc}

18.6. 3. Gaussian Process with Matern $\nu=1.5$ Kernel



Plotting A vs R_{tc}

18. Surrogate Model Selection in SpotOptim



<Figure size 1650x1050 with 0 Axes>

18.7. 4. Gaussian Process with Rational Quadratic Kernel

The Rational Quadratic kernel is a scale mixture of RBF kernels with different length scales, providing more flexibility than a single RBF.

```
start_time = time.time()

# Configure GP with Rational Quadratic kernel
kernel_rq = ConstantKernel(1.0, (1e-3, 1e3)) * RationalQuadratic(
    length_scale=1.0,
    alpha=1.0,
    length_scale_bounds=(1e-2, 1e2),
```

```

    alpha_bounds=(1e-2, 1e2)
)

gp_rq = GaussianProcessRegressor(
    kernel=kernel_rq,
    n_restarts_optimizer=10,
    normalize_y=True,
    random_state=seed
)

optimizer_rq = SpotOptim(
    fun=wingwt,
    bounds=bounds,
    surrogate=gp_rq,
    max_iter=max_iter,
    n_initial=n_initial,
    var_name=param_names,
    acquisition='ei',
    seed=seed,
    verbose=False
)

result_rq = optimizer_rq.optimize()
time_rq = time.time() - start_time

print(f"\nResults:")
print(f"  Best weight: {result_rq.fun:.4f} lb")
print(f"  Function evaluations: {result_rq.nfev}")
print(f"  Time: {time_rq:.2f}s")
print(f"  Success: {result_rq.success}")

results_comparison.append({
    'Surrogate': 'GP Rational Quadratic',
    'Best Weight': result_rq.fun,
    'Evaluations': result_rq.nfev,
    'Time (s)': time_rq,
    'Success': result_rq.success
})

```

Results:

```

Best weight: 120.1102 lb
Function evaluations: 20
Time: 0.48s


```

18. Surrogate Model Selection in SpotOptim

Success: True

18.7.1. Visualization: Rational Quadratic Kernel

```
optimizer_rq.plot_progress(log_y=False, figsize=(10, 5))
```



```
optimizer_rq.plot_important_hyperparameter_contour(max_imp=3)
plt.suptitle('GP Rational Quadratic Kernel: Most Important Parameters', y=1.02)
plt.show()
```

Plotting surrogate contours for top 3 most important parameters:

Nz: importance = 22.91% (type: float)


A: importance = 20.68% (type: float)

Rtc: importance = 18.19% (type: float)

Generating 3 surrogate plots...


Plotting Nz vs A

18.7. 4. Gaussian Process with Rational Quadratic Kernel




Plotting Nz vs Rtc

18. Surrogate Model Selection in SpotOptim



Plotting A vs R_{tc}

18.8. 5. SpotOptim Kriging Model



<Figure size 1650x1050 with 0 Axes>

18.8. 5. SpotOptim Kriging Model

SpotOptim includes its own Kriging implementation optimized for sequential design. It uses Gaussian correlation function and optimizes hyperparameters via differential evolution.

```
start_time = time.time()

# Configure Kriging model
kriging_model = Kriging(
    noise=1e-10,      # Regularization parameter
    kernel='gauss',   # Gaussian/RBF kernel
    n_theta=None,     # Auto: use number of dimensions
```

18. Surrogate Model Selection in SpotOptim

```
    min_theta=-3.0,          # Min log10(theta) bound
    max_theta=2.0,           # Max log10(theta) bound
    seed=seed
)

optimizer_kriging = SpotOptim(
    fun=wingwt,
    bounds=bounds,
    surrogate=kriging_model,
    max_iter=max_iter,
    n_initial=n_initial,
    var_name=param_names,
    acquisition='ei',
    seed=seed,
    verbose=False
)

result_kriging = optimizer_kriging.optimize()
time_kriging = time.time() - start_time


print(f"\nResults:")
print(f"  Best weight: {result_kriging.fun:.4f} lb")
print(f"  Function evaluations: {result_kriging.nfev}")
print(f"  Time: {time_kriging:.2f}s")
print(f"  Success: {result_kriging.success}")

results_comparison.append({
    'Surrogate': 'SpotOptim Kriging',
    'Best Weight': result_kriging.fun,
    'Evaluations': result_kriging.nfev,
    'Time (s)': time_kriging,
    'Success': result_kriging.success
})
```

```
Results:
  Best weight: 121.2932 lb
  Function evaluations: 20
  Time: 6.25s
  Success: True
```

18.8.1. Visualization: Kriging Model

```
optimizer_kriging.plot_progress(log_y=False, figsize=(10, 5))
```



```
optimizer_kriging.plot_important_hyperparameter_contour(max_imp=3)
plt.suptitle('SpotOptim Kriging: Most Important Parameters', y=1.02)
plt.show()
```


Plotting surrogate contours for top 3 most important parameters:

Nz: importance = 20.85% (type: float)
 A: importance = 18.80% (type: float)
 Rtc: importance = 16.51% (type: float)

Generating 3 surrogate plots...


Plotting Nz vs A

18. Surrogate Model Selection in SpotOptim




Plotting N_z vs R_{tc}

18.8. 5. SpotOptim Kriging Model



Plotting A vs R_{tc}

18. Surrogate Model Selection in SpotOptim



<Figure size 1650x1050 with 0 Axes>

18.9. 6. Random Forest Regressor

Random Forests are ensemble methods that handle noise well and can model discontinuities. They don't naturally provide uncertainty estimates, so the acquisition function uses predictions only.

```
start_time = time.time()

# Configure Random Forest
rf_model = RandomForestRegressor(
    n_estimators=100,
    max_depth=15,
    min_samples_split=2,
```

```

    min_samples_leaf=1,
    random_state=seed,
)

optimizer_rf = SpotOptim(
    fun=wingwt,
    bounds=bounds,
    surrogate=rf_model,
    max_iter=max_iter,
    n_initial=n_initial,
    var_name=param_names,
    acquisition='y', # Use 'y' (greedy) since RF doesn't provide std
    seed=seed,
    verbose=False
)

result_rf = optimizer_rf.optimize()
time_rf = time.time() - start_time

print(f"\nResults:")
print(f"  Best weight: {result_rf.fun:.4f} lb")
print(f"  Function evaluations: {result_rf.nfev}")
print(f"  Time: {time_rf:.2f}s")
print(f"  Success: {result_rf.success}")
print(f"  Note: Using acquisition='y' (greedy) since RF doesn't provide uncertainty")

results_comparison.append({
    'Surrogate': 'Random Forest',
    'Best Weight': result_rf.fun,
    'Evaluations': result_rf.nfev,
    'Time (s)': time_rf,
    'Success': result_rf.success
})

```

Results:


```

Best weight: 145.9792 lb
Function evaluations: 20
Time: 0.48s
Success: True
Note: Using acquisition='y' (greedy) since RF doesn't provide uncertainty

```

18.9.1. Visualization: Random Forest

```
optimizer_rf.plot_progress(log_y=False, figsize=(10, 5))
```



```
optimizer_rf.plot_important_hyperparameter_contour(max_imp=3)
plt.suptitle('Random Forest: Most Important Parameters', y=1.02)
plt.show()
```


Plotting surrogate contours for top 3 most important parameters:

Nz: importance = 21.66% (type: float)
A: importance = 19.61% (type: float)
Rtc: importance = 15.41% (type: float)

Generating 3 surrogate plots...


Plotting Nz vs A

18.9. 6. Random Forest Regressor




Plotting N_z vs R_{tc}

18. Surrogate Model Selection in SpotOptim



Plotting A vs R_{tc}

18.10. 7. XGBoost Regressor



<Figure size 1650x1050 with 0 Axes>

18.10. 7. XGBoost Regressor

XGBoost is a gradient boosting implementation known for excellent performance on structured data and fast training/prediction times.

```
if XGBOOST_AVAILABLE:
    start_time = time.time()

    # Configure XGBoost
    xgb_model = xgb.XGBRegressor(
        n_estimators=100,
        max_depth=6,
        learning_rate=0.1,
```

18. Surrogate Model Selection in SpotOptim

```
        subsample=0.8,
        colsample_bytree=0.8,
        random_state=seed,
        nthread=1,
        verbosity=0,
    )

    optimizer_xgb = SpotOptim(
        fun=wingwt,
        bounds=bounds,
        surrogate=xgb_model,
        max_iter=max_iter,
        n_initial=n_initial,
        var_name=param_names,
        acquisition='y', # Use 'y' (greedy) since XGBoost doesn't provide std
        seed=seed,
        verbose=False
    )

    result_xgb = optimizer_xgb.optimize()
    time_xgb = time.time() - start_time

    print(f"\nResults:")
    print(f"  Best weight: {result_xgb.fun:.4f} lb")
    print(f"  Function evaluations: {result_xgb.nfev}")
    print(f"  Time: {time_xgb:.2f}s")
    print(f"  Success: {result_xgb.success}")
    print(f"  Note: Using acquisition='y' (greedy) since XGBoost doesn't provide uncer

    results_comparison.append({
        'Surrogate': 'XGBoost',
        'Best Weight': result_xgb.fun,
        'Evaluations': result_xgb.nfev,
        'Time (s)': time_xgb,
        'Success': result_xgb.success
    })

    # Visualization
    optimizer_xgb.plot_progress(log_y=False, figsize=(10, 5))
    plt.title('XGBoost: Convergence')
    plt.show()

    optimizer_xgb.plot_important_hyperparameter_contour(max_imp=3)
    plt.suptitle('XGBoost: Most Important Parameters', y=1.02)
```

```
plt.show()
else:
    print("=" * 80)
    print("7. XGBoost Regressor - SKIPPED (not installed)")
    print("=" * 80)
    print("Install XGBoost with: pip install xgboost")
```

```
=====
7. XGBoost Regressor - SKIPPED (not installed)
=====
```

```
Install XGBoost with: pip install xgboost
```

18.11. 8. Support Vector Regression (SVR)

SVR with RBF kernel can model complex non-linear relationships. It's particularly good for high-dimensional problems with smooth structure.

```
start_time = time.time()

# Configure SVR
svr_model = SVR(
    kernel='rbf',
    C=100.0,
    epsilon=0.1,
    gamma='scale'
)

optimizer_svr = SpotOptim(
    fun=wingwt,
    bounds=bounds,
    surrogate=svr_model,
    max_iter=max_iter,
    n_initial=n_initial,
    var_name=param_names,
    acquisition='y', # Use 'y' (greedy) since SVR doesn't provide std by default
    seed=seed,
    verbose=False
)

result_svr = optimizer_svr.optimize()
time_svr = time.time() - start_time
```

18. Surrogate Model Selection in SpotOptim

```
print(f"\nResults:")
print(f"  Best weight: {result_svr.fun:.4f} lb")
print(f"  Function evaluations: {result_svr.nfev}")
print(f"  Time: {time_svr:.2f}s")
print(f"  Success: {result_svr.success}")


results_comparison.append({
    'Surrogate': 'SVR (RBF)',
    'Best Weight': result_svr.fun,
    'Evaluations': result_svr.nfev,
    'Time (s)': time_svr,
    'Success': result_svr.success
})
```

Results:

Best weight: 156.7071 lb
Function evaluations: 20
Time: 0.06s
Success: True

18.11.1. Visualization: SVR

```
optimizer_svr.plot_progress(log_y=False, figsize=(10, 5))
```



18.11. 8. Support Vector Regression (SVR)


```
optimizer_svr.plot_important_hyperparameter_contour(max_imp=3)
plt.suptitle('Support Vector Regression: Most Important Parameters', y=1.02)
plt.show()
```

Plotting surrogate contours for top 3 most important parameters:

Nz: importance = 25.67% (type: float)
A: importance = 22.43% (type: float)
Rtc: importance = 16.38% (type: float)


Generating 3 surrogate plots...

Plotting Nz vs A




Plotting Nz vs Rtc

18. Surrogate Model Selection in SpotOptim



Plotting A vs R_{tc}

18.12. 9. Gradient Boosting Regressor



<Figure size 1650x1050 with 0 Axes>

18.12. 9. Gradient Boosting Regressor

Gradient Boosting from scikit-learn is another ensemble method that builds trees sequentially, with each tree correcting errors of the previous ones.

```
start_time = time.time()

# Configure Gradient Boosting
gb_model = GradientBoostingRegressor(
    n_estimators=100,
    max_depth=5,
    learning_rate=0.1,
    subsample=0.8,
```

18. Surrogate Model Selection in SpotOptim

```
        random_state=seed
    )

    optimizer_gb = SpotOptim(
        fun=wingwt,
        bounds=bounds,
        surrogate=gb_model,
        max_iter=max_iter,
        n_initial=n_initial,
        var_name=param_names,
        acquisition='y', # Use 'y' (greedy) since GB doesn't provide std
        seed=seed,
        verbose=False
    )

    result_gb = optimizer_gb.optimize()
    time_gb = time.time() - start_time

    print(f"\nResults:")
    print(f"  Best weight: {result_gb.fun:.4f} lb")
    print(f"  Function evaluations: {result_gb.nfev}")
    print(f"  Time: {time_gb:.2f}s")
    print(f"  Success: {result_gb.success}")


    results_comparison.append({
        'Surrogate': 'Gradient Boosting',
        'Best Weight': result_gb.fun,
        'Evaluations': result_gb.nfev,
        'Time (s)': time_gb,
        'Success': result_gb.success
    })
```

Results:

```
Best weight: 135.4721 lb
Function evaluations: 20
Time: 0.14s
Success: True
```

18.12.1. Visualization: Gradient Boosting

```
optimizer_gb.plot_progress(log_y=False, figsize=(10, 5))
```



```
optimizer_gb.plot_important_hyperparameter_contour(max_imp=3)
plt.suptitle('Gradient Boosting: Most Important Parameters', y=1.02)
plt.show()
```


Plotting surrogate contours for top 3 most important parameters:

Nz: importance = 18.61% (type: float)
 A: importance = 16.46% (type: float)
 Rtc: importance = 14.06% (type: float)

Generating 3 surrogate plots...


Plotting Nz vs A

18. Surrogate Model Selection in SpotOptim




Plotting N_z vs R_{tc}

18.12. 9. Gradient Boosting Regressor



Plotting A vs R_{tc}

18. Surrogate Model Selection in SpotOptim



<Figure size 1650x1050 with 0 Axes>

18.13. Comprehensive Comparison

Now let's compare all surrogate models side-by-side.

```
# Create comparison DataFrame
df_comparison = pd.DataFrame(results_comparison)

# Calculate improvement from best
best_weight = df_comparison['Best Weight'].min()
df_comparison['Gap to Best (%)'] = (
    (df_comparison['Best Weight'] - best_weight) / best_weight * 100
)
```

```
# Sort by best weight
df_comparison = df_comparison.sort_values('Best Weight')

print("\n" + "=" * 100)
print("SURROGATE MODEL COMPARISON")
print("=" * 100)
print(df_comparison.to_string(index=False))
print("=" * 100)
```

SURROGATE MODEL COMPARISON

	Surrogate	Best Weight	Evaluations	Time (s)	Success	Gap to Best (%)
	GP Matern nu=1.5	119.546577	20	0.347332	True	0.000000
GP Matern nu=2.5 (Default)		119.866592	20	1.043850	True	0.267690
	GP RBF	120.109278	20	0.358928	True	0.470696
	GP Rational Quadratic	120.110159	20	0.483553	True	0.471433
	SpotOptim Kriging	121.293180	20	6.252107	True	1.461022
	Gradient Boosting	135.472131	20	0.135436	True	13.321631
	Random Forest	145.979231	20	0.477507	True	22.110757
	SVR (RBF)	156.707097	20	0.058865	True	31.084554

18.13.1. Visualization: Performance Comparison

```
fig, axes = plt.subplots(2, 2, figsize=(16, 12))

# Plot 1: Best weight comparison
ax1 = axes[0, 0]
colors = ['green' if i == 0 else 'steelblue' for i in range(len(df_comparison))]
ax1.barh(df_comparison['Surrogate'], df_comparison['Best Weight'], color=colors)
ax1.set_xlabel('Best Weight (lb)')
ax1.set_title('Best Weight Found by Each Surrogate')
ax1.axvline(x=best_weight, color='red', linestyle='--', linewidth=2, label='Best Overall')
ax1.legend()
ax1.grid(True, alpha=0.3, axis='x')

# Plot 2: Computational time
ax2 = axes[0, 1]
ax2.barh(df_comparison['Surrogate'], df_comparison['Time (s)'], color='coral')
ax2.set_xlabel('Time (seconds)')
```

18. Surrogate Model Selection in SpotOptim


```
ax2.set_title('Computational Time')
ax2.grid(True, alpha=0.3, axis='x')

# Plot 3: Gap to best
ax3 = axes[1, 0]
colors_gap = ['green' if gap < 0.1 else 'orange' if gap < 1.0 else 'red'
              for gap in df_comparison['Gap to Best (%)']]
ax3.barh(df_comparison['Surrogate'], df_comparison['Gap to Best (%)'], color=colors_gap)
ax3.set_xlabel('Gap to Best Solution (%)')
ax3.set_title('Solution Quality (Lower is Better)')
ax3.axvline(x=1.0, color='black', linestyle='--', alpha=0.5, linewidth=1)
ax3.grid(True, alpha=0.3, axis='x')

# Plot 4: Efficiency (weight reduction per second)
ax4 = axes[1, 1]
baseline_weight = wingwt(np.array([0.48, 0.4, 0.38, 0.5, 0.62, 0.344, 0.4, 0.37, 0.38]))
df_comparison['Efficiency'] = (baseline_weight - df_comparison['Best Weight']) / df_comparison['Time']
ax4.barh(df_comparison['Surrogate'], df_comparison['Efficiency'], color='mediumseagreen')
ax4.set_xlabel('Weight Reduction per Second (lb/s)')
ax4.set_title('Optimization Efficiency')
ax4.grid(True, alpha=0.3, axis='x')

plt.tight_layout()
plt.show()
```


18.14. Convergence Comparison



18.14. Convergence Comparison

Let's compare how different surrogates converge over iterations.

```
fig, ax = plt.subplots(figsize=(14, 8))

# Collect convergence data from all optimizers
optimizers = [
    (optimizer_default, 'GP Matern nu=2.5 (Default)', 'blue'),
    (optimizer_rbf, 'GP RBF', 'green'),
    (optimizer_matern15, 'GP Matern nu=1.5', 'red'),
    (optimizer_rq, 'GP Rational Quadratic', 'purple'),
    (optimizer_kriging, 'SpotOptim Kriging', 'orange'),
    (optimizer_rf, 'Random Forest', 'brown'),
    (optimizer_svr, 'SVR', 'pink'),
    (optimizer_gb, 'Gradient Boosting', 'gray')
]
```


18. Surrogate Model Selection in SpotOptim

```
if XGBOOST_AVAILABLE:
    optimizers.append((optimizer_xgb, 'XGBoost', 'cyan'))

for opt, label, color in optimizers:
    y_history = opt.y_
    best_so_far = np.minimum.accumulate(y_history)
    ax.plot(range(len(best_so_far)), best_so_far, linewidth=2, label=label, color=col

ax.set_xlabel('Function Evaluation', fontsize=12)
ax.set_ylabel('Best Wing Weight Found (lb)', fontsize=12)
ax.set_title('Convergence Comparison: All Surrogates', fontsize=14, fontweight='bold')
ax.legend(loc='upper right', fontsize=10)
ax.grid(True, alpha=0.3)
ax.set_xlim(0, max_iter)

plt.tight_layout()
plt.show()
```



18.15. Key Insights and Recommendations

```
print("\n" + "=" * 100)
print("KEY INSIGHTS AND RECOMMENDATIONS")
```

```

print("=" * 100)

# Find best surrogate
best_surrogate = df_comparison.iloc[0]['Surrogate']
best_value = df_comparison.iloc[0]['Best Weight']
best_time = df_comparison.iloc[0]['Time (s)']

print(f"\n1. BEST OVERALL PERFORMANCE:")
print(f"    Surrogate: {best_surrogate}")
print(f"    Best Weight: {best_value:.4f} lb")
print(f"    Computation Time: {best_time:.2f}s")

# Find fastest
fastest_idx = df_comparison['Time (s)'].idxmin()
fastest_surrogate = df_comparison.loc[fastest_idx, 'Surrogate']
fastest_time = df_comparison.loc[fastest_idx, 'Time (s)']

print(f"\n2. FASTEST OPTIMIZATION:")
print(f"    Surrogate: {fastest_surrogate}")
print(f"    Time: {fastest_time:.2f}s")
print(f"    Best Weight: {df_comparison.loc[fastest_idx, 'Best Weight']:.4f} lb")

# Find most efficient
most_efficient_idx = df_comparison['Efficiency'].idxmax()
most_efficient = df_comparison.loc[most_efficient_idx, 'Surrogate']

print(f"\n3. MOST EFFICIENT (weight reduction per second):")
print(f"    Surrogate: {most_efficient}")
print(f"    Efficiency: {df_comparison.loc[most_efficient_idx, 'Efficiency']:.4f} lb/s")

print(f"\n4. RECOMMENDATIONS BY PROBLEM TYPE:")
print(f"    - Smooth, continuous functions: Gaussian Process with RBF or Matern nu=2.5")
print(f"    - Functions with noise: Random Forest or Gradient Boosting")
print(f"    - High-dimensional problems (>20D): XGBoost or Random Forest")
print(f"    - Limited budget (<50 evals): Gaussian Process with Expected Improvement")
print(f"    - Fast evaluation needed: XGBoost or Random Forest")
print(f"    - Need uncertainty estimates: Gaussian Process or Kriging")
print(f"    - Non-smooth/discontinuous: Random Forest or Gradient Boosting")

print(f"\n5. KERNEL COMPARISON (Gaussian Process):")
gp_results = df_comparison[df_comparison['Surrogate'].str.contains('GP')]
print(gp_results[['Surrogate', 'Best Weight', 'Time (s)']].to_string(index=False))

print("\n" + "=" * 100)

```

18. Surrogate Model Selection in SpotOptim

=====

KEY INSIGHTS AND RECOMMENDATIONS

=====

1. BEST OVERALL PERFORMANCE:
Surrogate: GP Matern nu=1.5
Best Weight: 119.5466 lb
Computation Time: 0.35s
2. FASTEST OPTIMIZATION:
Surrogate: SVR (RBF)
Time: 0.06s
Best Weight: 156.7071 lb
3. MOST EFFICIENT (weight reduction per second):
Surrogate: SVR (RBF)
Efficiency: 1311.4961 lb/s
4. RECOMMENDATIONS BY PROBLEM TYPE:
 - Smooth, continuous functions: Gaussian Process with RBF or Matern nu=2.5
 - Functions with noise: Random Forest or Gradient Boosting
 - High-dimensional problems (>20D): XGBoost or Random Forest
 - Limited budget (<50 evals): Gaussian Process with Expected Improvement
 - Fast evaluation needed: XGBoost or Random Forest
 - Need uncertainty estimates: Gaussian Process or Kriging
 - Non-smooth/discontinuous: Random Forest or Gradient Boosting
5. KERNEL COMPARISON (Gaussian Process):

	Surrogate	Best Weight	Time (s)
	GP Matern nu=1.5	119.546577	0.347332
GP Matern nu=2.5 (Default)		119.866592	1.043850
	GP RBF	120.109278	0.358928
	GP Rational Quadratic	120.110159	0.483553

=====

18.16. Summary Statistics

```
# Summary statistics
print("\n" + "=" * 100)
print("SUMMARY STATISTICS")
```

```

print("=" * 100)

summary_stats = pd.DataFrame({
    'Metric': [
        'Best Weight Found',
        'Worst Weight Found',
        'Average Weight',
        'Std Dev Weight',
        'Fastest Time',
        'Slowest Time',
        'Average Time',
    ],
    'Value': [
        f"{df_comparison['Best Weight'].min():.4f} lb",
        f"{df_comparison['Best Weight'].max():.4f} lb",
        f"{df_comparison['Best Weight'].mean():.4f} lb",
        f"{df_comparison['Best Weight'].std():.4f} lb",
        f"{df_comparison['Time (s)'].min():.2f} s",
        f"{df_comparison['Time (s)'].max():.2f} s",
        f"{df_comparison['Time (s)'].mean():.2f} s",
    ]
})

print(summary_stats.to_string(index=False))
print("=" * 100)

```

```

=====
SUMMARY STATISTICS
=====

```

```

      Metric      Value
Best Weight Found 119.5466 lb
Worst Weight Found 156.7071 lb
Average Weight 129.8855 lb
Std Dev Weight 14.5497 lb
Fastest Time 0.06 s
Slowest Time 6.25 s
Average Time 1.14 s
=====

```

18.17. Conclusion

This comprehensive comparison demonstrates that:

1. **Gaussian Processes** with appropriate kernels (Matern, RBF) provide excellent performance for smooth optimization problems and naturally support Expected Improvement acquisition
2. **SpotOptim Kriging** offers a lightweight alternative to sklearn's GP with comparable performance
3. **Random Forest** and **XGBoost** are robust alternatives that handle noise and discontinuities well, though they require greedy acquisition
4. **SVR** and **Gradient Boosting** offer middle-ground solutions with good scalability
5. The choice of surrogate should be based on:
 - Function smoothness
 - Computational budget
 - Need for uncertainty quantification
 - Problem dimensionality
 - Noise characteristics

For the AWWE problem, Gaussian Process surrogates generally performed best due to the function's smooth structure, but tree-based methods (RF, XGBoost, GB) can be preferable for more complex, noisy, or high-dimensional problems.

18.18. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

19. Acquisition Failure Handling in SpotOptim

SpotOptim provides sophisticated fallback strategies for handling acquisition function failures during optimization. This ensures robust optimization even when the surrogate model struggles to suggest new points.

19.1. What is Acquisition Failure?

During surrogate-based optimization, the acquisition function suggests new points to evaluate. However, sometimes the suggested point is **too close** to existing points (within `tolerance_x` distance), which would provide little new information. When this happens, SpotOptim uses a **fallback strategy** to propose an alternative point.

19.2. Fallback Strategies

SpotOptim uses a **fallback strategy** to propose an alternative point. The `acquisition_failure_strategy` parameter controls this behavior, defaulting to "random".

19.2.1. Random Space-Filling Design (Default)

Strategy name: "random"

This strategy uses Latin Hypercube Sampling (LHS) to generate a new space-filling point. LHS ensures good coverage of the search space by dividing each dimension into equal-probability intervals.

When to use:

- General-purpose optimization
- When you want simplicity and good space-filling properties
- Default choice for most problems

19. Acquisition Failure Handling in SpotOptim

Example:

```
from spotoptim import SpotOptim
import numpy as np

def sphere(X):
    return np.sum(X**2, axis=1)

optimizer = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=50,
    n_initial=10,
    acquisition_failure_strategy="random", # Default
    verbose=True
)

result = optimizer.optimize()
```

```
TensorBoard logging disabled
Initial best: f(x) = 4.615750
Iter 1 | Best: 4.615750 | Curr: 13.598656 | Rate: 0.00 | Evals: 22.0%
Iter 2 | Best: 0.020058 | Rate: 0.50 | Evals: 24.0%
Iter 3 | Best: 0.000021 | Rate: 0.67 | Evals: 26.0%
Iter 4 | Best: 0.000002 | Rate: 0.75 | Evals: 28.0%
Iter 5 | Best: 0.000001 | Rate: 0.80 | Evals: 30.0%
Iter 6 | Best: 0.000001 | Curr: 0.000002 | Rate: 0.67 | Evals: 32.0%
Iter 7 | Best: 0.000001 | Curr: 0.000002 | Rate: 0.57 | Evals: 34.0%
Iter 8 | Best: 0.000001 | Rate: 0.62 | Evals: 36.0%
Iter 9 | Best: 0.000001 | Rate: 0.67 | Evals: 38.0%
Iter 10 | Best: 0.000001 | Curr: 0.000002 | Rate: 0.60 | Evals: 40.0%
Iter 11 | Best: 0.000001 | Curr: 0.000002 | Rate: 0.55 | Evals: 42.0%
Iter 12 | Best: 0.000001 | Curr: 0.000001 | Rate: 0.50 | Evals: 44.0%
Iter 13 | Best: 0.000001 | Curr: 0.000002 | Rate: 0.46 | Evals: 46.0%
Iter 14 | Best: 0.000001 | Curr: 0.000002 | Rate: 0.43 | Evals: 48.0%
Iter 15 | Best: 0.000001 | Rate: 0.47 | Evals: 50.0%
Iter 16 | Best: 0.000001 | Curr: 0.000001 | Rate: 0.44 | Evals: 52.0%
Iter 17 | Best: 0.000001 | Curr: 0.000001 | Rate: 0.41 | Evals: 54.0%
Iter 18 | Best: 0.000000 | Rate: 0.44 | Evals: 56.0%
Iter 19 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.42 | Evals: 58.0%
Iter 20 | Best: 0.000000 | Curr: 0.000000 | Rate: 0.40 | Evals: 60.0%
Iter 21 | Best: 0.000000 | Curr: 0.000001 | Rate: 0.38 | Evals: 62.0%
Iter 22 | Best: 0.000000 | Rate: 0.41 | Evals: 64.0%
Iter 23 | Best: 0.000000 | Rate: 0.43 | Evals: 66.0%
```


Iter 24	Best: 0.000000	Curr: 0.000000	Rate: 0.42	Evals: 68.0%
Iter 25	Best: 0.000000	Curr: 0.000000	Rate: 0.40	Evals: 70.0%
Iter 26	Best: 0.000000	Curr: 0.000000	Rate: 0.38	Evals: 72.0%
Iter 27	Best: 0.000000	Curr: 0.000000	Rate: 0.37	Evals: 74.0%
Iter 28	Best: 0.000000	Curr: 0.000000	Rate: 0.36	Evals: 76.0%
Iter 29	Best: 0.000000	Curr: 0.000000	Rate: 0.34	Evals: 78.0%
Iter 30	Best: 0.000000	Curr: 0.000000	Rate: 0.33	Evals: 80.0%
Iter 31	Best: 0.000000	Curr: 0.000000	Rate: 0.32	Evals: 82.0%
Iter 32	Best: 0.000000	Curr: 0.000003	Rate: 0.31	Evals: 84.0%
Iter 33	Best: 0.000000	Curr: 0.000000	Rate: 0.30	Evals: 86.0%
Iter 34	Best: 0.000000	Curr: 0.000010	Rate: 0.29	Evals: 88.0%
Iter 35	Best: 0.000000	Curr: 0.000000	Rate: 0.29	Evals: 90.0%
Iter 36	Best: 0.000000	Curr: 0.000000	Rate: 0.28	Evals: 92.0%
Iter 37	Best: 0.000000	Curr: 0.000000	Rate: 0.27	Evals: 94.0%
Iter 38	Best: 0.000000	Curr: 0.000182	Rate: 0.26	Evals: 96.0%
Iter 39	Best: 0.000000	Curr: 0.000000	Rate: 0.26	Evals: 98.0%
Iter 40	Best: 0.000000	Curr: 0.000001	Rate: 0.25	Evals: 100.0%

19.3. How It Works

The acquisition failure handling is integrated into the optimization process:

1. **Acquisition optimization:** SpotOptim uses differential evolution to optimize the acquisition function
2. **Distance check:** The proposed point is checked against existing points using `tolerance_x`
3. **Fallback activation:** If the point is too close, `_handle_acquisition_failure()` is called
4. **Strategy execution:** The configured fallback strategy generates a new point
5. **Evaluation:** The fallback point is evaluated and added to the dataset

19.4. Advanced Usage: Setting Tolerance

The `tolerance_x` parameter controls when the fallback strategy is triggered. A larger tolerance means points need to be farther apart, triggering the fallback more often:

```
def simple_objective(X):
    """Simple quadratic function for demonstration"""
    return np.sum(X**2, axis=1)

bounds_demo = [(-5, 5), (-5, 5)]
```

19. Acquisition Failure Handling in SpotOptim

```
# Strict tolerance (smaller value) - fewer fallbacks
optimizer_strict = SpotOptim(
    fun=simple_objective,
    bounds=bounds_demo,
    tolerance_x=1e-6, # Very small - almost never triggers fallback
    max_iter=20,
    seed=42
)

# Relaxed tolerance (larger value) - more fallbacks
optimizer_relaxed = SpotOptim(
    fun=simple_objective,
    bounds=bounds_demo,
    tolerance_x=0.5, # Larger - triggers fallback more often
    max_iter=20,
    seed=42
)

print(f"Strict tolerance setup complete")
print(f"Relaxed tolerance setup complete")
```

```
Strict tolerance setup complete
Relaxed tolerance setup complete
```

19.5. Best Practices

19.5.1. 1. Monitor Fallback Activations

Enable verbose mode to see when fallbacks are triggered:

```
def test_objective(X):
    return np.sum(X**2, axis=1)

optimizer = SpotOptim(
    fun=test_objective,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=20,
    verbose=True, # Shows fallback messages
    seed=42
)
print("Optimizer with verbose mode created")
```

TensorBoard logging disabled
 Optimizer with verbose mode created

19.5.2. 2. Adjust Tolerance Based on Problem Scale

For problems with small search spaces, use smaller tolerance:

```
def scale_objective(X):
    return np.sum(X**2, axis=1)

# Small search space
optimizer_small = SpotOptim(
    fun=scale_objective,
    bounds=[(-1, 1), (-1, 1)],
    tolerance_x=0.01, # Small tolerance for small space
    max_iter=20,
    seed=42
)

# Large search space
optimizer_large = SpotOptim(
    fun=scale_objective,
    bounds=[(-100, 100), (-100, 100)],
    tolerance_x=1.0, # Larger tolerance for large space
    max_iter=20,
    seed=42
)

print(f"Small space optimizer created (bounds: [-1, 1])")
print(f"Large space optimizer created (bounds: [-100, 100])")
```

Small space optimizer created (bounds: [-1, 1])
 Large space optimizer created (bounds: [-100, 100])

19.6. Technical Details

19.6.1. Random Strategy Implementation

The fallback strategy:

1. Generates a single point using Latin Hypercube Sampling
2. Ensures the point is within bounds

19. Acquisition Failure Handling in SpotOptim

3. Applies variable type repairs (rounding for int/factor variables)

This is computationally efficient while maintaining good space-filling properties.

19.7. Summary

- **Strategy:** SpotOptim uses a random space-filling strategy (LHS) when acquisition fails.
- **Trigger:** Activated when acquisition-proposed point is too close to existing points (within `tolerance_x`)
- **Monitoring:** Enable `verbose=True` to see when fallbacks occur

19.8. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

20. Point Selection Implementation in SpotOptim

20.1. Overview

This feature automatically selects a subset of evaluated points for surrogate model training when the total number of points exceeds a specified threshold.

It is implemented as a point selection mechanism for SpotOptim that mirrors the functionality in spotpython's `Spot` class.

20.2. Implementation Details

20.2.1. Parameters

Added to `SpotOptim.__init__`:

- `max_surrogate_points` (int, optional): Maximum number of points to use for surrogate fitting
- `selection_method` (str, default='distant'): Method for selecting points ('distant' or 'best')

20.2.2. Methods

1. `select_distant_points(X, y, k)`

- Uses K-means clustering to find k clusters
- Selects the point closest to each cluster center
- Ensures space-filling properties for surrogate training
- Mimics `spotpython.utils.aggregate.select_distant_points`

2. `select_best_cluster(X, y, k)`

- Uses K-means clustering to find k clusters
- Computes mean objective value for each cluster
- Selects all points from the cluster with the best (lowest) mean value

20. Point Selection Implementation in SpotOptim

- Mimics `spotpython.utils.aggregate.select_best_cluster`

3. `_selection_dispatcher(X, y)`

- Dispatcher method that routes to the appropriate selection function
- Returns all points if `max_surrogate_points` is None
- Mimics `spotpython.spot.spot.Spot.selection_dispatcher`

The method `_fit_surrogate(X, y)` checks if `X.shape[0] > self.max_surrogate_points`. If true, it calls `_selection_dispatcher` to get a subset. Then, it fits the surrogate only on the selected points. This implementation matches the logic in `spotpython.spot.spot.Spot.fit_surrogate`

20.3. Key Differences from spotpython

While the implementation follows `spotpython`'s design, there is a difference: `spotoptim` uses a simplified clustering, it uses `sklearn`'s `KMeans` directly instead of a custom implementation.

20.4. Example Usage

This example demonstrates the point selection feature with a limited number of surrogate points. Increase `MAX_ITER`, `N_INITIAL`, and `MAX_SURROGATE_POINTS` to see more pronounced effects.

```
from spotoptim import SpotOptim
import numpy as np

MAX_ITER = 20
N_INITIAL = 5
MAX_SURROGATE_POINTS = 10

# Define an example objective function
def sphere(X):
    """Simple sphere function for demonstration"""
    return np.sum(X**2, axis=1)

bounds = [(-5, 5), (-5, 5), (-5, 5)]

# Without point selection (default behavior)
optimizer1 = SpotOptim(
    fun=sphere,
```

```

    bounds=bounds,
    max_iter=MAX_ITER,
    n_initial=N_INITIAL,
    seed=42
)
result1 = optimizer1.optimize()
print(f"Without selection - Best value: {result1.fun:.6f}")
print(f"Total points evaluated: {result1.nfev}")

# With point selection using distant method
optimizer2 = SpotOptim(
    fun=sphere,
    bounds=bounds,
    max_iter=MAX_ITER,
    n_initial=N_INITIAL,
    max_surrogate_points=MAX_SURROGATE_POINTS,
    selection_method='distant',
    seed=42
)
result2 = optimizer2.optimize()
print(f"\nWith 'distant' selection - Best value: {result2.fun:.6f}")
print(f"Total points evaluated: {result2.nfev}")
print(f"Max surrogate points: {optimizer2.max_surrogate_points}")

# With point selection using best cluster method
optimizer3 = SpotOptim(
    fun=sphere,
    bounds=bounds,
    max_iter=MAX_ITER,
    n_initial=N_INITIAL,
    max_surrogate_points=MAX_SURROGATE_POINTS,
    selection_method='best',
    seed=42
)
result3 = optimizer3.optimize()
print(f"\nWith 'best' selection - Best value: {result3.fun:.6f}")
print(f"Total points evaluated: {result3.nfev}")
print(f"Max surrogate points: {optimizer3.max_surrogate_points}")

```

Without selection - Best value: 0.000000
Total points evaluated: 20

With 'distant' selection - Best value: 1.034865
Total points evaluated: 20

20. Point Selection Implementation in SpotOptim

Max surrogate points: 10

With 'best' selection - Best value: 3.365132

Total points evaluated: 20

Max surrogate points: 10

20.5. Benefits

1. **Scalability:** Enables efficient optimization with many function evaluations
2. **Computational efficiency:** Reduces surrogate training time for large datasets
3. **Maintained accuracy:** Careful point selection preserves model quality
4. **Flexibility:** Two selection methods for different optimization scenarios

20.6. Comparison with spotpython

Feature	spotpython	SpotOptim
Point selection via clustering	x	x
'distant' method	x	x
'best' method	x	x
Selection dispatcher	x	x
Nyström approximation	x	
Modular design	x (utils.aggregate)	x (class methods)

20.7. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

Part VI.

Optimization on the Surrogate

21. Optimization on the Surrogate

In Surrogate-Model-Based Optimization (SMBO), determining the next candidate point to evaluate involves solving an internal optimization problem. This problem aims to maximize an **acquisition function** (like Expected Improvement) based on the current surrogate model. This step is also known as “optimization on the surrogate” or “infill criterion optimization”.

`SpotOptim` provides extensive control over this process, allowing you to choose the acquisition function, the optimization algorithm used to maximize it, and fine-tune that optimizer’s parameters.

21.1. Key Arguments

The following arguments in `SpotOptim` control the optimization on the surrogate:

- **acquisition** (str): The acquisition function to use. Common choices are "y" (Surrogate Prediction, typically for minimization), "EI" (Expected Improvement), "LCB" (Lower Confidence Bound).
- **acquisition_optimizer** (str or callable): The optimization algorithm used to maximize the acquisition function. Defaults to "differential_evolution". You can also specify any method supported by `scipy.optimize.minimize` (e.g., "L-BFGS-B", "Nelder-Mead") or provide a custom callable.
- **acquisition_optimizer_kwargs** (dict): A dictionary of keyword arguments passed to the **acquisition_optimizer**. This allows you to tune parameters like population size, max iterations, or tolerance. Defaults to `{'maxiter': 10000, 'gtol': 1e-9}` if not provided.
 - **Note:** These kwargs are *also* passed to the Surrogate Model’s internal optimizer (e.g., for hyperparameter tuning of the Gaussian Process) if applicable.
- **acquisition_fun_return_size** (int): The number of candidate points to return from the acquisition optimization. Defaults to 3. This is useful for batch evaluation or providing multiple starting points.
- **acquisition_failure_strategy** (str): Strategy to use if the acquisition optimization fails or yields poor results. Options include "random" (sample random points) or "best" (use best found so far).

21.2. Examples

The following examples demonstrate how to configure these parameters.

21.2.1. 1. Default Configuration (Differential Evolution)

By default, `SpotOptim` uses Differential Evolution (`scipy.optimize.differential_evolution`).

```
import numpy as np
from spotoptim import SpotOptim

def obj_fun(X):
    return np.sum(X**2, axis=1)

# Default behavior
spot = SpotOptim(
    fun=obj_fun,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=10,
    n_initial=2,
    acquisition="EI",
    # Default: acquisition_optimizer="differential_evolution"
)
spot.optimize()
print("Best y:", spot.best_y_)
```

Best y: 0.18696636523305848

21.2.2. 2. Customizing Differential Evolution

You can use `acquisition_optimizer_kwargs` to adjust Differential Evolution parameters, such as increasing `maxiter` or changing the `popsiz`.

```
import numpy as np
from spotoptim import SpotOptim

def obj_fun(X):
    return np.sum(X**2, axis=1)

# Configure DE parameters
de_kwargs = {
    "maxiter": 200,    # Increase max iterations
```

```

    "popsize": 30,      # Increase population size
    "mutation": (0.6, 1.1)
}

spot = SpotOptim(
    fun=obj_fun,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=10,
    n_initial=2,
    acquisition="EI",
    acquisition_optimizer="differential_evolution",
    acquisition_optimizer_kwargs=de_kwargs
)
spot.optimize()
print("Best y with Custom DE:", spot.best_y_)

```

Best y with Custom DE: 0.6513632578887001

21.2.3. 3. Using Gradient-Based Optimization (L-BFGS-B)

You can switch to a gradient-based optimizer like L-BFGS-B by specifying it in `acquisition_optimizer`. Note that for `minimize`-based methods, parameters are usually passed via an `options` dictionary within `acquisition_optimizer_kwargs`.

```

import numpy as np
from spotoptim import SpotOptim

def obj_fun(X):
    return np.sum(X**2, axis=1)

# Configure L-BFGS-B parameters
lbfgs_kwargs = {
    "method": "L-BFGS-B", # Explicitly state method (good practice)
    "options": {
        "maxiter": 100,
        "ftol": 1e-9
    }
}

spot = SpotOptim(
    fun=obj_fun,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=10,

```

21. Optimization on the Surrogate

```
n_initial=2,
acquisition="EI",
acquisition_optimizer="L-BFGS-B",
acquisition_optimizer_kwargs=lbfgs_kwargs
)
spot.optimize()
print("Best y with L-BFGS-B:", spot.best_y_)
```

Best y with L-BFGS-B: 0.6306878460280021

21.2.4. 4. Using Gradient-Free Optimization (Nelder-Mead)

For non-smooth acquisition landscapes or when robustness is needed without gradients, Nelder-Mead is a good choice. SpotOptim automatically handles the interface to ensure compatibility.

```
import numpy as np
from spotoptim import SpotOptim

def obj_fun(X):
    return np.sum(X**2, axis=1)

# Configure Nelder-Mead
nm_kwargs = {
    "method": "Nelder-Mead",
    "options": {
        "maxiter": 500,
        "adaptive": True
    }
}

spot = SpotOptim(
    fun=obj_fun,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=10,
    n_initial=2,
    acquisition="EI",
    acquisition_optimizer="Nelder-Mead",
    acquisition_optimizer_kwargs=nm_kwargs
)
spot.optimize()
print("Best y with Nelder-Mead:", spot.best_y_)
```

Best y with Nelder-Mead: 0.2585730176337664

21.2.5. 5. Returning Multiple Candidates

Setting `acquisition_fun_return_size > 1` forces the optimizer to return multiple candidate points (e.g., the top N from the final population).

```
import numpy as np
from spotoptim import SpotOptim

def obj_fun(X):
    return np.sum(X**2, axis=1)

spot = SpotOptim(
    fun=obj_fun,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=5, # Short run just to demo config
    n_initial=2,
    acquisition="EI",
    acquisition_fun_return_size=5 # Return top 5 candidates
)
# The internal optimization loop handles these candidates automatically
spot.optimize()
```

```
message: Optimization terminated: maximum evaluations (5) reached
        Current function value: 12.337862
        Iterations: 3
        Function evaluations: 5

success: True
  fun: 12.337862160925518
   x: [ 3.461e+00  5.988e-01]
  X: [[-4.607e+00  3.318e+00]
      [ 3.723e+00 -8.692e-01]
      [ 3.167e+00  2.435e+00]
      [ 5.000e+00 -5.000e+00]
      [ 3.461e+00  5.988e-01]]
  nit: 3
 nfev: 5
   y: [ 3.223e+01  1.462e+01  1.596e+01  5.000e+01  1.234e+01]
```


Part VII.

Kriging

22. Kriging Surrogate Integration

22.1. Overview

Implementation of a Kriging (Gaussian Process) surrogate model to SpotOptim, providing an alternative to scikit-learn's GaussianProcessRegressor.

22.2. Module Structure

```
src/spotoptim/surrogate/  
███ __init__.py          # Module exports  
███ kriging.py           # Kriging implementation (~350 lines)  
███ README.md            # Module documentation
```

22.3. Kriging Class (src/spotoptim/surrogate/kriging.py)

Key Features:

- Scikit-learn compatible interface (`fit()`, `predict()`)
- Gaussian (RBF) kernel: $R = \exp(-D)$
- Automatic hyperparameter optimization via maximum likelihood
- Cholesky decomposition for efficient linear algebra
- Prediction with uncertainty (`return_std=True`)
- Reproducible results via seed parameter

Implementation Details:

- lean, well-documented code
- No external dependencies beyond NumPy, SciPy
- Simplified from spotpython.surrogate.kriging
- Focused on core functionality needed for SpotOptim

Parameters:

- `noise`: Regularization (nugget effect)

22. Kriging Surrogate Integration

- `kernel`: Currently 'gauss' (Gaussian/RBF)
- `n_theta`: Number of length scale parameters
- `min_theta`, `max_theta`: Bounds for hyperparameter optimization
- `seed`: Random seed for reproducibility

22.4. Integration with SpotOptim

No Changes Required to SpotOptim Core!

The existing `surrogate` parameter already supports any scikit-learn compatible model:

```
from spotoptim import SpotOptim, Kriging
import numpy as np

def rosenbrock(X):
    """Rosenbrock function"""
    x = X[:, 0]
    y = X[:, 1]
    return (1 - x)**2 + 100 * (y - x**2)**2

kriging = Kriging(seed=42)
optimizer = SpotOptim(
    fun=rosenbrock,
    bounds=[(-2, 2), (-2, 2)],
    surrogate=kriging, # Just pass the Kriging instance
    max_iter=30,
    seed=42
)
result = optimizer.optimize()
print(f"Best value: {result.fun:.6f}")
print(f"Best point: {result.x}")
```

Best value: 0.057427

Best point: [0.76200682 0.57784969]

22.5. Documentation

Added Example to `notebooks/demos.ipynb`

- Demonstrates Kriging vs GP comparison
- Shows custom parameter usage

22.6. Usage Examples

22.6.1. Basic Usage

```
from spotoptim import SpotOptim, Kriging
import numpy as np

def sphere(X):
    """Sphere function:  $f(x) = \sum(x^2)$ """
    return np.sum(X**2, axis=1)

kriging = Kriging(noise=1e-6, seed=42)
optimizer = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    surrogate=kriging,
    max_iter=20,
    seed=42
)
result = optimizer.optimize()
print(f"Best value: {result.fun:.6f}")
print(f"Best point: {result.x}")
```

Best value: 0.000017

Best point: [-0.00259053 0.0031664]

22.6.2. Custom Parameters

```
import numpy as np

def ackley(X):
    """Ackley function - multimodal test function"""
    a = 20
    b = 0.2
    c = 2 * np.pi
    n = X.shape[1]

    sum_sq = np.sum(X**2, axis=1)
    sum_cos = np.sum(np.cos(c * X), axis=1)
```

22. Kriging Surrogate Integration

```
        return -a * np.exp(-b * np.sqrt(sum_sq / n)) - np.exp(sum_cos / n) + a + np.e

kriging = Kriging(
    noise=1e-4,
    min_theta=-2.0,
    max_theta=3.0,
    seed=123
)

optimizer = SpotOptim(
    fun=ackley,
    bounds=[(-5, 5), (-5, 5)],
    surrogate=kriging,
    max_iter=40,
    seed=123
)
result = optimizer.optimize()
print(f"Best value: {result.fun:.6f}")
print(f"Best point: {result.x}")
```

Best value: 0.004575

Best point: [-0.00112193 -0.00113152]

22.6.3. Prediction with Uncertainty

```
from spotoptim import Kriging
import numpy as np

# Generate training data
np.random.seed(42)
X_train = np.random.uniform(-5, 5, (20, 2))
y_train = np.sum(X_train**2, axis=1)

# Generate test data
X_test = np.random.uniform(-5, 5, (10, 2))

# Fit model and predict with uncertainty
model = Kriging(seed=42)
model.fit(X_train, y_train)
y_pred, y_std = model.predict(X_test, return_std=True)

print(f"Predictions shape: {y_pred.shape}")
```

```
print(f"Uncertainties shape: {y_std.shape}")
print(f"Mean prediction: {y_pred.mean():.4f}")
print(f"Mean uncertainty: {y_std.mean():.4f}")
```

```
Predictions shape: (10,)
Uncertainties shape: (10,)
Mean prediction: 20.8065
Mean uncertainty: 0.0302
```

22.7. Technical Details

22.7.1. Kriging vs GaussianProcessRegressor

Aspect	Kriging	GaussianProcessRegressor
Lines of code	~350	Complex internal implementation
Dependencies	NumPy, SciPy	scikit-learn + dependencies
Kernel	Gaussian (RBF)	Multiple types (Matern, RQ, etc.)
Hyperparameter opt	Differential Evolution	L-BFGS-B with restarts
Use case	Simplified, explicit	Production, flexible

22.7.2. Algorithm

1. Correlation Matrix:

- Compute squared distances: $D_{ij} = \sum_k \mathbb{I}_k (x_{ik} - x_{jk})^2$
- Apply kernel: $R_{ij} = \exp(-D_{ij})$
- Add nugget: $R_{ii} += \text{noise}$

2. Maximum Likelihood:

- Optimize $\mathbb{I}$ via differential evolution
- Minimize: $(n/2)\log(\mathbb{I}^2) + (1/2)\log|R|$
- Concentrated likelihood (μ profiled out)

3. Prediction:

- Mean: $\hat{f}(x) = \hat{\mu} + \mathbb{I}(x)\mathbb{R}\mathbb{I}^{-1}r$
- Variance: $s^2(x) = \mathbb{I}^2[1 + \mathbb{I} - \mathbb{I}(x)\mathbb{R}\mathbb{I}^{-1}\mathbb{I}(x)]$
- Uses Cholesky decomposition for efficiency

22.7.3. Key Arguments Passed from SpotOptim

SpotOptim passes these to the surrogate via the standard interface:

During fit:

```
# Example of how SpotOptim uses the surrogate internally
surrogate.fit(X, y)
```

- **X**: Training points (`n_initial` or accumulated evaluations)
- **y**: Function values

During predict:

```
# Example of internal usage
mu = surrogate.predict(x)[0] # For acquisition='y'
mu, sigma = surrogate.predict(x, return_std=True) # For acquisition='ei', 'pi'
```

Implicit parameters via seed:

- `random_state=seed` (for `GaussianProcessRegressor`)
- `seed=seed` (for Kriging)

22.8. Benefits

1. **Self-contained**: No heavy scikit-learn dependency for surrogate
2. **Explicit**: Clear hyperparameter bounds and optimization
3. **Educational**: Readable implementation of Kriging/GP
4. **Flexible**: Easy to extend with new kernels or features
5. **Compatible**: Works seamlessly with existing SpotOptim API

22.9. Future Enhancements

Potential additions:

- ☐ Additional kernels (Matern, Exponential, Cubic)
- ☐ Anisotropic hyperparameters (separate θ per dimension)
- ☐ Gradient-enhanced predictions
- ☐ Batch predictions for efficiency
- ☐ Parallel hyperparameter optimization
- ☐ ARD (Automatic Relevance Determination)

22.10. Jupyter Notebook

i Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

23. Kriging Surrogate Models in SpotOptim

23.1. Introduction

SpotOptim provides a powerful Kriging (Gaussian Process) surrogate model that can significantly enhance optimization performance. This tutorial introduces the **Kriging** class, explains its theory and parameters, and demonstrates how to use it effectively with SpotOptim.

23.1.1. What is Kriging?

Kriging, also known as Gaussian Process regression, is a sophisticated interpolation method that:

- **Predicts function values** at untested points based on observed data
- **Provides uncertainty estimates** indicating confidence in predictions
- **Models smooth functions** common in engineering and scientific optimization
- **Handles noisy observations** through regression approaches

The mathematical foundation of Kriging is based on the assumption that the objective function $f(\mathbf{x})$ can be modeled as:

$$f(\mathbf{x}) = \mu + Z(\mathbf{x})$$

where μ is a constant mean and $Z(\mathbf{x})$ is a Gaussian process with correlation structure. The correlation between two points $\mathbf{x}_i$ and $\mathbf{x}_j$ is typically modeled using a Gaussian (RBF) correlation function:

$$R(\mathbf{x}_i, \mathbf{x}_j) = \exp \left(- \sum_{k=1}^d \theta_k |x_{i,k} - x_{j,k}|^2 \right)$$

where $\theta_k > 0$ are the correlation length parameters that control smoothness in each dimension.

23.1.2. Why Use Kriging in SpotOptim?

1. **Better Surrogate Models:** More accurate predictions than simple interpolation
2. **Uncertainty Quantification:** Know where the model is uncertain
3. **Mixed Variable Types:** Handle continuous, integer, and categorical variables
4. **Multiple Methods:** Choose between interpolation, regression, and reinterpolation
5. **Customizable:** Control regularization, correlation parameters, and more

23.2. Basic Usage

23.2.1. Creating a Simple Kriging Model

Let's start with the most basic usage - creating a Kriging model with default settings:

```
import numpy as np
from spotoptim import SpotOptim
from spotoptim.surrogate import Kriging

# Simple objective function
def sphere(X):
    """Sphere function:  $f(x) = \sum(x^2)$ """
    return np.sum(X**2, axis=1)

# Create Kriging surrogate with defaults
kriging = Kriging(seed=42)

# Use with SpotOptim
optimizer = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    surrogate=kriging, # Use Kriging instead of default GP
    max_iter=20,
    n_initial=10,
    seed=42,
    verbose=True
)

result = optimizer.optimize()
print(f"\nBest solution: {result.x}")
print(f"Best value: {result.fun:.6f}")
```

```

TensorBoard logging disabled
Initial best: f(x) = 3.925721
Iter 1 | Best: 0.000460 | Rate: 1.00 | Evals: 55.0%
Iter 2 | Best: 0.000155 | Rate: 1.00 | Evals: 60.0%
Iter 3 | Best: 0.000017 | Rate: 1.00 | Evals: 65.0%
Iter 4 | Best: 0.000017 | Curr: 0.000572 | Rate: 0.75 | Evals: 70.0%
Iter 5 | Best: 0.000017 | Curr: 0.000244 | Rate: 0.60 | Evals: 75.0%
Iter 6 | Best: 0.000017 | Curr: 0.000178 | Rate: 0.50 | Evals: 80.0%
Iter 7 | Best: 0.000017 | Curr: 0.000139 | Rate: 0.43 | Evals: 85.0%
Iter 8 | Best: 0.000017 | Curr: 0.000099 | Rate: 0.38 | Evals: 90.0%
Iter 9 | Best: 0.000017 | Curr: 0.000081 | Rate: 0.33 | Evals: 95.0%
Iter 10 | Best: 0.000017 | Curr: 0.000039 | Rate: 0.30 | Evals: 100.0%

Best solution: [-0.00259053  0.0031664 ]
Best value: 0.000017

```

23.2.2. Default vs Custom Surrogate

SpotOptim uses a Gaussian Process Regressor by default. Here's how Kriging compares:

```

import numpy as np
from spotoptim import SpotOptim
from spotoptim.surrogate import Kriging

def rosenbrock(X):
    """Rosenbrock function: (1-x)^2 + 100(y-x^2)^2"""
    x = X[:, 0]
    y = X[:, 1]
    return (1 - x)**2 + 100 * (y - x**2)**2

# With default Gaussian Process
opt_default = SpotOptim(
    fun=rosenbrock,
    bounds=[(-2, 2), (-2, 2)],
    max_iter=30,
    n_initial=15,
    seed=42,
    verbose=False
)
result_default = opt_default.optimize()

# With Kriging surrogate
kriging = Kriging(method='regression', seed=42)

```

23. Kriging Surrogate Models in SpotOptim

```
opt_kriging = SpotOptim(  
    fun=rosenbrock,  
    bounds=[(-2, 2), (-2, 2)],  
    surrogate=kriging,  
    max_iter=30,  
    n_initial=15,  
    seed=42,  
    verbose=False  
)  
result_kriging = opt_kriging.optimize()  
  
print("Comparison:")  
print(f"Default GP:   f(x) = {result_default.fun:.6f}")  
print(f"Kriging:      f(x) = {result_kriging.fun:.6f}")
```

```
Comparison:  
Default GP:   f(x) = 0.126472  
Kriging:      f(x) = 0.091470
```

Both approaches work well, but Kriging offers more control over the surrogate behavior.

23.3. Understanding Kriging Parameters

23.3.1. Method: Interpolation vs Regression

The `method` parameter controls how Kriging handles data:

- **“interpolation”**: Exact interpolation, passes through all training points
- **“regression”**: Allows smoothing, better for noisy data
- **“reinterpolation”**: Hybrid approach

```
import numpy as np  
from spotoptim.surrogate import Kriging  
import matplotlib.pyplot as plt  
  
# Create noisy training data  
np.random.seed(42)  
X_train = np.linspace(0, 2*np.pi, 10).reshape(-1, 1)  
y_train = np.sin(X_train.ravel()) + 0.1 * np.random.randn(10)  
  
# Test data for smooth predictions
```

23.3. Understanding Kriging Parameters

```
X_test = np.linspace(0, 2*np.pi, 100).reshape(-1, 1)
y_true = np.sin(X_test.ravel())


# Compare methods
methods = ['interpolation', 'regression', 'reinterpolation']
fig, axes = plt.subplots(1, 3, figsize=(15, 4))

for ax, method in zip(axes, methods):
    model = Kriging(method=method, seed=42, model_fun_evals=50)
    model.fit(X_train, y_train)
    y_pred, y_std = model.predict(X_test, return_std=True)

    # Plot
    ax.plot(X_test, y_true, 'k--', label='True function', linewidth=2)
    ax.plot(X_test, y_pred, 'b-', label='Prediction', linewidth=2)
    ax.fill_between(X_test.ravel(),
                   y_pred - 2*y_std,
                   y_pred + 2*y_std,
                   alpha=0.3, label='95% CI')
    ax.scatter(X_train, y_train, c='r', s=50, zorder=5, label='Training data')
    ax.set_title(f'Method: {method}')
    ax.set_xlabel('x')
    ax.set_ylabel('y')
    ax.legend()
    ax.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

print("Key differences:")
print("- Interpolation: Passes exactly through training points")
print("- Regression: Smooths over noisy data (recommended)")
print("- Reinterpolation: Balances between the two")
```



Key differences:

23. Kriging Surrogate Models in SpotOptim

- Interpolation: Passes exactly through training points
- Regression: Smooths over noisy data (recommended)
- Reinterpolation: Balances between the two

Recommendation: Use `method='regression'` for most optimization problems, especially with noisy objectives.

23.3.2. Noise Parameter and Regularization

The noise parameter adds a small nugget effect for numerical stability:

```
import numpy as np
from spotoptim import SpotOptim
from spotoptim.surrogate import Kriging

def noisy_ackley(X):
    """Ackley function with observation noise"""
    a = 20
    b = 0.2
    c = 2 * np.pi
    n = X.shape[1]

    sum_sq = np.sum(X**2, axis=1)
    sum_cos = np.sum(np.cos(c * X), axis=1)

    base = -a * np.exp(-b * np.sqrt(sum_sq / n)) - np.exp(sum_cos / n) + a + np.e
    noise = np.random.normal(0, 0.5, size=base.shape) # Add noise
    return base + noise

# Very small noise (near interpolation)
kriging_small = Kriging(noise=1e-10, method='interpolation', seed=42)
opt_small = SpotOptim(
    fun=noisy_ackley,
    bounds=[(-5, 5), (-5, 5)],
    surrogate=kriging_small,
    max_iter=25,
    n_initial=12,
    seed=42,
    verbose=False
)
result_small = opt_small.optimize()

# Larger noise (more robust to noise)
```



```

kriging_large = Kriging(noise=0.01, method='interpolation', seed=42)
opt_large = SpotOptim(
    fun=noisy_ackley,
    bounds=[(-5, 5), (-5, 5)],
    surrogate=kriging_large,
    max_iter=25,
    n_initial=12,
    seed=42,
    verbose=False
)
result_large = opt_large.optimize()

print("Impact of noise parameter:")
print(f"Small noise (1e-10): f(x) = {result_small.fun:.6f}")
print(f"Large noise (0.01): f(x) = {result_large.fun:.6f}")
print("\nNote: For noisy functions, use 'regression' method instead,")
print("      which automatically optimizes the Lambda (nugget) parameter.")

```

```

Impact of noise parameter:
Small noise (1e-10): f(x) = -0.196059
Large noise (0.01): f(x) = 2.312981

```

Note: For noisy functions, use 'regression' method instead,
which automatically optimizes the Lambda (nugget) parameter.

Best Practice: For noisy functions, use method='regression' which automatically optimizes the regularization parameter instead of fixing it.

23.3.3. Correlation Length Parameters (Theta)

The `theta` parameters control smoothness. SpotOptim optimizes these automatically:

- **min_theta:** Minimum $\log_{10}(\theta)$ value (default: -3.0 $\rightarrow \theta = 0.001$)
- **max_theta:** Maximum $\log_{10}(\theta)$ value (default: 2.0 $\rightarrow \theta = 100$)

Smaller $\theta \rightarrow$ smoother function, larger $\theta \rightarrow$ more wiggly function.

```

import numpy as np
from spotoptim.surrogate import Kriging

# Sample data
X_train = np.array([[0.0], [0.5], [1.0], [1.5], [2.0]])
y_train = np.sin(X_train.ravel())

```

23. Kriging Surrogate Models in SpotOptim

```
# Tight bounds (smoother)
model_smooth = Kriging(min_theta=-1.0, max_theta=0.0, seed=42, model_fun_evals=30)
model_smooth.fit(X_train, y_train)

# Wide bounds (more flexible)
model_flexible = Kriging(min_theta=-3.0, max_theta=2.0, seed=42, model_fun_evals=30)
model_flexible.fit(X_train, y_train)

print("Theta bounds and optimal values:")
print(f"Smooth model:   bounds=[-1.0, 0.0], theta={model_smooth.theta}")
print(f"Flexible model: bounds=[-3.0, 2.0], theta={model_flexible.theta}")
print("\nThe optimizer automatically finds the best theta within the bounds.")
```

Theta bounds and optimal values:

Smooth model: bounds=[-1.0, 0.0], theta=[-0.95216449]

Flexible model: bounds=[-3.0, 2.0], theta=[-0.95309049]

The optimizer automatically finds the best theta within the bounds.

Default Values: The defaults `min_theta=-3.0` and `max_theta=2.0` work well for most problems.

23.3.4. Isotropic vs Anisotropic Correlation

- **Isotropic** (`isotropic=True`): Single θ for all dimensions (faster, fewer parameters)
- **Anisotropic** (`isotropic=False`): Different θ per dimension (more flexible, default)

```
import numpy as np
from spotoptim.surrogate import Kriging

# 3D problem
np.random.seed(42)
X_train = np.random.rand(15, 3) * 10 - 5
y_train = X_train[:, 0]**2 + 0.1*X_train[:, 1]**2 + 2*X_train[:, 2]**2

# Isotropic: one theta for all dimensions
model_iso = Kriging(isotropic=True, seed=42, model_fun_evals=40)
model_iso.fit(X_train, y_train)

# Anisotropic: different theta per dimension
```

```

model_aniso = Kriging(isotropic=False, seed=42, model_fun_evals=40)
model_aniso.fit(X_train, y_train)

print("Correlation structure:")
print(f"Isotropic:   n_theta={model_iso.n_theta}, theta={model_iso.theta}")
print(f"Anisotropic: n_theta={model_aniso.n_theta}, theta={model_aniso.theta}")
print("\nAnisotropic can capture different smoothness in each dimension.")

# Test predictions
X_test = np.array([[0.0, 0.0, 0.0]])
y_iso = model_iso.predict(X_test)
y_aniso = model_aniso.predict(X_test)
print(f"\nPrediction at origin:")
print(f"Isotropic:   {y_iso[0]:.4f}")
print(f"Anisotropic: {y_aniso[0]:.4f}")

```

```

Correlation structure:
Isotropic:   n_theta=1, theta=[-3.]
Anisotropic: n_theta=3, theta=[-2.23292358 -3.          -1.97228776]

Anisotropic can capture different smoothness in each dimension.

Prediction at origin:
Isotropic:   -0.0666
Anisotropic: -1.0267

```

When to Use:

- Use **isotropic** for faster fitting when dimensions have similar characteristics
- Use **anisotropic** (default) when dimensions behave differently

23.4. Advanced Features

23.4.1. Mixed Variable Types

Kriging supports mixed variable types: continuous (`float`), integer (`int`), and categorical (`factor`):

```

import numpy as np
from spotoptim import SpotOptim
from spotoptim.surrogate import Kriging

```

23. Kriging Surrogate Models in SpotOptim

```
def mixed_objective(X):
    """
    Optimize a system with:
    - x0: continuous learning rate (0.001 to 0.1)
    - x1: integer number of layers (1 to 5)
    - x2: categorical activation (0=ReLU, 1=Tanh, 2=Sigmoid)
    """
    X = np.atleast_2d(X)
    learning_rate = X[:, 0]
    n_layers = X[:, 1]
    activation = X[:, 2]

    # Simplified model: prefer lr~0.01, 3 layers, ReLU
    loss = (learning_rate - 0.01)**2 + (n_layers - 3)**2
    loss += np.where(activation == 0, 0.0, # ReLU is best
                    np.where(activation == 1, 0.5, # Tanh is ok
                             1.0)) # Sigmoid is worst

    return loss

# Define bounds and types
bounds = [
    (0.001, 0.1), # learning_rate (float)
    (1, 5),      # n_layers (int)
    (0, 2)       # activation (factor: 0, 1, or 2)
]

var_type = ['float', 'int', 'factor']

# Create Kriging with mixed types
kriging_mixed = Kriging(
    method='regression',
    var_type=var_type,
    seed=42,
    model_fun_evals=50
)

# Optimize
optimizer = SpotOptim(
    fun=mixed_objective,
    bounds=bounds,
    var_type=var_type,
    surrogate=kriging_mixed,
    max_iter=30,
    n_initial=15,
```

```

    seed=42,
    verbose=True
)

result = optimizer.optimize()

print(f"\nOptimal configuration:")
print(f"Learning rate: {result.x[0]:.4f}")
print(f"Num layers:      {int(result.x[1])}")
print(f"Activation:        {int(result.x[2])} (0=ReLU, 1=Tanh, 2=Sigmoid)")
print(f"Loss:              {result.fun:.6f}")

```

```

TensorBoard logging disabled
Initial best: f(x) = 0.500035
Iter 1 | Best: 0.500035 | Curr: 0.500081 | Rate: 0.00 | Evals: 53.3%
Iter 2 | Best: 0.000081 | Rate: 0.50 | Evals: 56.7%
Optimizer candidate 1/3 was duplicate/invalid.
Iter 3 | Best: 0.000081 | Curr: 0.500036 | Rate: 0.33 | Evals: 60.0%
Optimizer candidate 1/3 was duplicate/invalid.
Iter 4 | Best: 0.000081 | Curr: 0.500000 | Rate: 0.25 | Evals: 63.3%
Optimizer candidate 1/3 was duplicate/invalid.
Iter 5 | Best: 0.000081 | Curr: 1.000014 | Rate: 0.20 | Evals: 66.7%
Iter 6 | Best: 0.000021 | Rate: 0.33 | Evals: 70.0%
Optimizer candidate 1/3 was duplicate/invalid.
Iter 7 | Best: 0.000021 | Curr: 1.000009 | Rate: 0.29 | Evals: 73.3%
Iter 8 | Best: 0.000015 | Rate: 0.38 | Evals: 76.7%
Iter 9 | Best: 0.000015 | Curr: 0.500001 | Rate: 0.33 | Evals: 80.0%
Iter 10 | Best: 0.000015 | Curr: 0.500000 | Rate: 0.30 | Evals: 83.3%
Iter 11 | Best: 0.000001 | Rate: 0.36 | Evals: 86.7%
Iter 12 | Best: 0.000001 | Curr: 0.500001 | Rate: 0.33 | Evals: 90.0%
Iter 13 | Best: 0.000001 | Curr: 1.000001 | Rate: 0.31 | Evals: 93.3%
Iter 14 | Best: 0.000001 | Curr: 0.500000 | Rate: 0.29 | Evals: 96.7%
Iter 15 | Best: 0.000001 | Curr: 0.000001 | Rate: 0.27 | Evals: 100.0%

```

```

Optimal configuration:
Learning rate: 0.0093
Num layers:    3
Activation:    0 (0=ReLU, 1=Tanh, 2=Sigmoid)
Loss:          0.000001

```

Key Points:

- Set `var_type` in both Kriging and SpotOptim
- Factor variables use specialized distance metrics (default: Canberra)

23. Kriging Surrogate Models in SpotOptim

- Integer variables are treated as ordered but discrete

23.4.2. Customizing the Distance Metric for Factors

For categorical variables, you can choose different distance metrics.

```
import numpy as np
from spotoptim.surrogate import Kriging

# Categorical data: 2 factor variables
X_train = np.array([
    [0, 0], # Category A, Color Red
    [1, 0], # Category B, Color Red
    [0, 1], # Category A, Color Blue
    [1, 1], # Category B, Color Blue
    [2, 2], # Category C, Color Green
])
y_train = np.array([1.0, 1.5, 1.2, 2.0, 3.5])

# Different distance metrics
metrics = ['canberra', 'hamming']

for metric in metrics:
    model = Kriging(
        method='regression',
        var_type=['factor', 'factor'],
        metric_factorial=metric,
        seed=42,
        model_fun_evals=40
    )
    model.fit(X_train, y_train)

    # Test prediction
    X_test = np.array([[1, 1]])
    y_pred = model.predict(X_test)

    print(f"Metric: {metric:10s} | Prediction at [1,1]: {y_pred[0]:.4f}")

print("\nAvailable metrics: 'canberra' (default), 'hamming', 'jaccard', etc.")
```

```
Metric: canberra    | Prediction at [1,1]: 1.9463
Metric: hamming     | Prediction at [1,1]: 2.0000
```

```
Available metrics: 'canberra' (default), 'hamming', 'jaccard', etc.
```

23.4.3. Handling High-Dimensional Problems

For high-dimensional problems, Kriging can become computationally expensive. Here are strategies:

```
import numpy as np
from spotoptim import SpotOptim
from spotoptim.surrogate import Kriging

def high_dim_sphere(X):
    """10-dimensional sphere function"""
    return np.sum(X**2, axis=1)

# Strategy 1: Use isotropic correlation (fewer parameters)
kriging_iso = Kriging(
    method='regression',
    isotropic=True, # Single theta for all 10 dimensions
    seed=42,
    model_fun_evals=50 # Faster optimization
)

optimizer_iso = SpotOptim(
    fun=high_dim_sphere,
    bounds=[(-5, 5)] * 10, # 10 dimensions
    surrogate=kriging_iso,
    max_iter=50,
    n_initial=30,
    seed=42,
    verbose=False
)

result_iso = optimizer_iso.optimize()

# Strategy 2: Use max_surrogate_points to limit training set size
kriging_limited = Kriging(method='regression', seed=42)

optimizer_limited = SpotOptim(
    fun=high_dim_sphere,
    bounds=[(-5, 5)] * 10,
    surrogate=kriging_limited,
    max_iter=50,
    n_initial=30,
    max_surrogate_points=100, # Limit surrogate training set
    seed=42,
    verbose=False
)
```

23. Kriging Surrogate Models in SpotOptim

```
)

result_limited = optimizer_limited.optimize()

print("High-dimensional optimization strategies:")
print(f"Isotropic Kriging:      f(x) = {result_iso.fun:.6f}")
print(f"Limited training set:  f(x) = {result_limited.fun:.6f}")
print("\nFor >5 dimensions, consider isotropic=True or limit training set.")
```

```
High-dimensional optimization strategies:
Isotropic Kriging:      f(x) = 0.002313
Limited training set:  f(x) = 0.083729
```

For >5 dimensions, consider isotropic=True or limit training set.

23.4.4. Uncertainty Quantification

Kriging provides uncertainty estimates, useful for exploration vs exploitation:

```
import numpy as np
from spotoptim.surrogate import Kriging
import matplotlib.pyplot as plt

# 1D example for visualization
np.random.seed(42)
X_train = np.array([[0.0], [1.0], [3.0], [5.0], [6.0]])
y_train = np.sin(X_train.ravel()) + 0.05 * np.random.randn(5)

# Fit Kriging model
model = Kriging(method='regression', seed=42, model_fun_evals=50)
model.fit(X_train, y_train)

# Dense test points
X_test = np.linspace(-0.5, 6.5, 200).reshape(-1, 1)
y_pred, y_std = model.predict(X_test, return_std=True)

# True function
y_true = np.sin(X_test.ravel())

# Plot
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
```



```

plt.plot(X_test, y_true, 'k--', label='True function', linewidth=2)
plt.plot(X_test, y_pred, 'b-', label='Kriging prediction', linewidth=2)
plt.fill_between(X_test.ravel(),
                 y_pred - 2*y_std,
                 y_pred + 2*y_std,
                 alpha=0.3, color='blue', label='95% confidence')
plt.scatter(X_train, y_train, c='red', s=100, zorder=5,
            edgecolors='black', label='Training points')
plt.xlabel('x', fontsize=12)
plt.ylabel('y', fontsize=12)
plt.title('Kriging Predictions with Uncertainty', fontsize=14)
plt.legend()
plt.grid(True, alpha=0.3)


plt.subplot(1, 2, 2)
plt.plot(X_test, y_std, 'r-', linewidth=2)
plt.scatter(X_train, np.zeros_like(X_train), c='blue', s=100,
            zorder=5, edgecolors='black', label='Training points')
plt.xlabel('x', fontsize=12)
plt.ylabel('Standard Deviation', fontsize=12)
plt.title('Uncertainty Estimates', fontsize=14)
plt.legend()
plt.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

print("Key observations:")
print("- Uncertainty is low near training points")
print("- Uncertainty is high far from training data")
print("- This guides acquisition functions (e.g., EI exploits low uncertainty)")

```

23. Kriging Surrogate Models in SpotOptim



Key observations:

- Uncertainty is low near training points
- Uncertainty is high far from training data
- This guides acquisition functions (e.g., EI exploits low uncertainty)

23.5. Practical Examples

23.5.1. Example 1: Optimizing the Rastrigin Function

The Rastrigin function is highly multimodal - a challenging test case:

```
import numpy as np
from spotoptim import SpotOptim
from spotoptim.surrogate import Kriging

def rastrigin(X):
    """
    Rastrigin function: highly multimodal
    Global minimum:  $f(0, \dots, 0) = 0$ 
    """
    A = 10
    n = X.shape[1]
    return A * n + np.sum(X**2 - A * np.cos(2 * np.pi * X), axis=1)

# Configure Kriging for multimodal function
kriging = Kriging(
    method='regression',      # Smooth over local variations
    min_theta=-3.0,          # Allow flexible correlation
```

```

max_theta=2.0,
seed=42,
model_fun_evals=100      # More budget for better surrogate
)

optimizer = SpotOptim(
    fun=rastrigin,
    bounds=[(-5.12, 5.12), (-5.12, 5.12)],
    surrogate=kriging,
    acquisition='ei',      # Expected Improvement for exploration
    max_iter=60,
    n_initial=30,
    seed=42,
    verbose=True
)

result = optimizer.optimize()

print(f"\nRastrigin optimization:")
print(f"Best solution: {result.x}")
print(f"Best value:      {result.fun:.6f} (global optimum: 0.0)")
print(f"Distance to optimum: {np.linalg.norm(result.x):.6f}")

```

TensorBoard logging disabled

Initial best: $f(x) = 16.581088$

Iter 1	Best: 16.581088	Curr: 32.343746	Rate: 0.00	Evals: 51.7%
Iter 2	Best: 16.581088	Curr: 29.443523	Rate: 0.00	Evals: 53.3%
Iter 3	Best: 16.581088	Curr: 38.997047	Rate: 0.00	Evals: 55.0%
Iter 4	Best: 16.581088	Curr: 40.561338	Rate: 0.00	Evals: 56.7%
Iter 5	Best: 16.581088	Curr: 39.590541	Rate: 0.00	Evals: 58.3%
Iter 6	Best: 16.581088	Curr: 29.177251	Rate: 0.00	Evals: 60.0%
Iter 7	Best: 16.537179	Rate: 0.14	Evals: 61.7%	
Iter 8	Best: 11.879924	Rate: 0.25	Evals: 63.3%	
Iter 9	Best: 11.879924	Curr: 15.985887	Rate: 0.22	Evals: 65.0%
Iter 10	Best: 8.399681	Rate: 0.30	Evals: 66.7%	
Iter 11	Best: 8.399681	Curr: 29.969447	Rate: 0.27	Evals: 68.3%
Iter 12	Best: 6.176571	Rate: 0.33	Evals: 70.0%	
Iter 13	Best: 6.176571	Curr: 12.034154	Rate: 0.31	Evals: 71.7%
Iter 14	Best: 6.176571	Curr: 6.866442	Rate: 0.29	Evals: 73.3%
Iter 15	Best: 6.176571	Curr: 17.848506	Rate: 0.27	Evals: 75.0%
Iter 16	Best: 6.176571	Curr: 22.932958	Rate: 0.25	Evals: 76.7%
Iter 17	Best: 6.176571	Curr: 10.109510	Rate: 0.24	Evals: 78.3%
Iter 18	Best: 6.176571	Curr: 27.148627	Rate: 0.22	Evals: 80.0%
Iter 19	Best: 2.554349	Rate: 0.26	Evals: 81.7%	

23. Kriging Surrogate Models in SpotOptim

```
Iter 20 | Best: 1.223834 | Rate: 0.30 | Evals: 83.3%
Iter 21 | Best: 1.016496 | Rate: 0.33 | Evals: 85.0%
Iter 22 | Best: 1.016496 | Curr: 1.068918 | Rate: 0.32 | Evals: 86.7%
Iter 23 | Best: 1.016496 | Curr: 10.193493 | Rate: 0.30 | Evals: 88.3%
Iter 24 | Best: 1.016496 | Curr: 18.559615 | Rate: 0.29 | Evals: 90.0%
Iter 25 | Best: 1.016496 | Curr: 22.110057 | Rate: 0.28 | Evals: 91.7%
Iter 26 | Best: 1.016496 | Curr: 16.563339 | Rate: 0.27 | Evals: 93.3%
Iter 27 | Best: 1.016496 | Curr: 2.346268 | Rate: 0.26 | Evals: 95.0%
Iter 28 | Best: 1.016496 | Curr: 21.868452 | Rate: 0.25 | Evals: 96.7%
Iter 29 | Best: 1.016496 | Curr: 9.862757 | Rate: 0.24 | Evals: 98.3%
Iter 30 | Best: 1.016496 | Curr: 3.221800 | Rate: 0.23 | Evals: 100.0%
```

Rastrigin optimization:

Best solution: [-0.00313486 -1.00489573]

Best value: 1.016496 (global optimum: 0.0)

Distance to optimum: 1.004901

23.5.2. Example 2: Robust Optimization with Noise

When optimizing noisy functions, Kriging's regression mode helps:

```
import numpy as np
from spotoptim import SpotOptim
from spotoptim.surrogate import Kriging

def noisy_beale(X):
    """
    Beale function with Gaussian noise
    Global minimum: f(3, 0.5) = 0
    """
    x = X[:, 0]
    y = X[:, 1]

    term1 = (1.5 - x + x * y)**2
    term2 = (2.25 - x + x * y**2)**2
    term3 = (2.625 - x + x * y**3)**2

    base = term1 + term2 + term3
    noise = np.random.normal(0, 0.5, size=base.shape)

    return base + noise

# Use regression method for noisy objectives
kriging_robust = Kriging(
```

```

    method='regression',          # Automatically optimizes Lambda (nugget)
    min_Lambda=-9.0,             # Allow small regularization
    max_Lambda=-2.0,             # But not too large
    seed=42,
    model_fun_evals=80
)

# Use repeated evaluations for noise handling
optimizer = SpotOptim(
    fun=noisy_beale,
    bounds=[(-4.5, 4.5), (-4.5, 4.5)],
    surrogate=kriging_robust,
    max_iter=50,
    n_initial=20,
    repeats_initial=3,           # Evaluate each initial point 3 times
    repeats_surrogate=2,        # Evaluate each new point 2 times
    seed=42,
    verbose=True
)

result = optimizer.optimize()

print(f"\nNoisy Beale optimization:")
print(f"Best solution: {result.x}")
print(f"Best mean value: {optimizer.min_mean_y:.6f}")
print(f"Variance at best: {optimizer.min_var_y:.6f}")
print(f"True optimum: [3.0, 0.5]")
print(f"Distance: {np.linalg.norm(result.x - np.array([3.0, 0.5])):.4f}")

```

TensorBoard logging disabled

Initial best: $f(x) = 8.638354$, mean best: $f(x) = 8.806788$

Noisy Beale optimization:

Best solution: [2.25151734 -0.78841507]

Best mean value: 8.806788

Variance at best: 0.056351

True optimum: [3.0, 0.5]

Distance: 1.4900

23.5.3. Example 3: Real-World Machine Learning Hyperparameter Tuning

Optimize hyperparameters for a neural network:

23. Kriging Surrogate Models in SpotOptim

```
import numpy as np
from spotoptim import SpotOptim
from spotoptim.surrogate import Kriging

def ml_objective(X):
    """
    Simulate training a neural network
    Variables:
    - x0: learning_rate (log scale: 1e-4 to 1e-1)
    - x1: num_layers (integer: 1 to 5)
    - x2: hidden_size (integer: 16, 32, 64, 128, 256)
    - x3: dropout_rate (float: 0.0 to 0.5)

    Returns validation error
    """
    X = np.atleast_2d(X)
    lr = X[:, 0]
    n_layers = X[:, 1]
    hidden_size = X[:, 2]
    dropout = X[:, 3]

    # Simulate validation error (lower is better)
    # Optimal around: lr=0.001, 3 layers, 64 hidden, 0.2 dropout
    error = (np.log10(lr) + 3)**2 # Prefer lr ~ 0.001
    error += (n_layers - 3)**2 # Prefer 3 layers
    error += ((hidden_size - 64) / 32)**2 # Prefer 64 hidden units
    error += (dropout - 0.2)**2 * 10 # Prefer 0.2 dropout

    # Add some noise to simulate training variance
    error += np.random.normal(0, 0.1, size=error.shape)

    return error

# Setup optimization
bounds = [
    (1e-4, 1e-1), # learning_rate
    (1, 5), # num_layers
    (16, 256), # hidden_size
    (0.0, 0.5) # dropout_rate
]

var_type = ['float', 'int', 'int', 'float']

# Use Kriging with mixed types
```

```

kriging_ml = Kriging(
    method='regression',
    var_type=var_type,
    seed=42,
    model_fun_evals=60
)

optimizer = SpotOptim(
    fun=ml_objective,
    bounds=bounds,
    var_type=var_type,
    var_trans=['log10', None, None, None], # Log scale for learning rate
    surrogate=kriging_ml,
    max_iter=40,
    n_initial=20,
    seed=42,
    verbose=True
)

result = optimizer.optimize()

print(f"\nOptimal hyperparameters:")
print(f"Learning rate: {result.x[0]:.6f}")
print(f"Num layers:    {int(result.x[1])}")
print(f"Hidden size:     {int(result.x[2])}")
print(f"Dropout rate:    {result.x[3]:.3f}")
print(f"Validation error: {result.fun:.6f}")

```

```

TensorBoard logging disabled
Initial best: f(x) = 1.937667
Iter 1 | Best: 1.937667 | Curr: 7.989495 | Rate: 0.00 | Evals: 52.5%
Iter 2 | Best: 1.937667 | Curr: 5.729961 | Rate: 0.00 | Evals: 55.0%
Iter 3 | Best: 1.937667 | Curr: 5.701930 | Rate: 0.00 | Evals: 57.5%
Iter 4 | Best: 1.398150 | Rate: 0.25 | Evals: 60.0%
Iter 5 | Best: 1.357760 | Rate: 0.40 | Evals: 62.5%
Iter 6 | Best: 1.357760 | Curr: 1.441755 | Rate: 0.33 | Evals: 65.0%
Iter 7 | Best: 1.357760 | Curr: 1.382557 | Rate: 0.29 | Evals: 67.5%
Iter 8 | Best: 1.357760 | Curr: 1.437570 | Rate: 0.25 | Evals: 70.0%
Iter 9 | Best: 1.357760 | Curr: 2.161225 | Rate: 0.22 | Evals: 72.5%
Iter 10 | Best: 1.357760 | Curr: 1.488995 | Rate: 0.20 | Evals: 75.0%
Iter 11 | Best: 1.218966 | Rate: 0.27 | Evals: 77.5%
Iter 12 | Best: 1.218966 | Curr: 2.182884 | Rate: 0.25 | Evals: 80.0%
Iter 13 | Best: 1.164571 | Rate: 0.31 | Evals: 82.5%
Iter 14 | Best: 0.955484 | Rate: 0.36 | Evals: 85.0%

```

23. Kriging Surrogate Models in SpotOptim

```
Iter 15 | Best: 0.955484 | Curr: 1.084840 | Rate: 0.33 | Evals: 87.5%
Iter 16 | Best: 0.955484 | Curr: 1.241790 | Rate: 0.31 | Evals: 90.0%
Iter 17 | Best: 0.955484 | Curr: 1.101066 | Rate: 0.29 | Evals: 92.5%
Iter 18 | Best: 0.866854 | Rate: 0.33 | Evals: 95.0%
Iter 19 | Best: 0.866854 | Curr: 0.930139 | Rate: 0.32 | Evals: 97.5%
Iter 20 | Best: 0.866854 | Curr: 1.082757 | Rate: 0.30 | Evals: 100.0%
```

```
Optimal hyperparameters:
Learning rate: 0.000100
Num layers:    3
Hidden size:   56
Dropout rate:  0.206
Validation error: 0.866854
```

23.5.4. Example 4: Comparing Kriging Methods

Let's compare all three Kriging methods on the same problem:

```
import numpy as np
from spotoptim import SpotOptim
from spotoptim.surrogate import Kriging

def levy(X):
    """
    Levy function N. 13
    Global minimum: f(1, 1) = 0
    """
    x = X[:, 0]
    y = X[:, 1]

    w1 = 1 + (x - 1) / 4
    w2 = 1 + (y - 1) / 4

    term1 = np.sin(np.pi * w1)**2
    term2 = (w1 - 1)**2 * (1 + 10 * np.sin(np.pi * w1 + 1)**2)
    term3 = (w2 - 1)**2 * (1 + np.sin(2 * np.pi * w2)**2)

    return term1 + term2 + term3

methods = ['interpolation', 'regression', 'reinterpolation']
results_methods = []

for method in methods:
    kriging = Kriging(
```



```

        method=method,
        seed=42,
        model_fun_evals=60
    )

    optimizer = SpotOptim(
        fun=levy,
        bounds=[(-10, 10), (-10, 10)],
        surrogate=kriging,
        max_iter=40,
        n_initial=20,
        seed=42,
        verbose=False
    )

    result = optimizer.optimize()
    results_methods.append((method, result.fun, result.x))

    print(f"{method:15s} | f(x) = {result.fun:8.6f} | x = {result.x}")

print(f"\nGlobal optimum: f(1, 1) = 0.0")
print("\nAll methods find good solutions, but 'regression' is most robust.")

```

```

interpolation    | f(x) = 0.000481 | x = [0.97975967 0.98011761]
regression       | f(x) = 0.000082 | x = [0.99254945 0.98221295]
reinterpolation | f(x) = 0.000082 | x = [0.99254945 0.98221295]

```

Global optimum: f(1, 1) = 0.0

All methods find good solutions, but 'regression' is most robust.

23.5.5. Example 5: Sensitivity to Theta Bounds

Theta bounds control the range of smoothness. Let's see their impact:

```

import numpy as np
from spotoptim import SpotOptim
from spotoptim.surrogate import Kriging

def griewank(X):
    """Griewank function"""
    sum_sq = np.sum(X**2 / 4000, axis=1)
    prod_cos = np.prod(np.cos(X / np.sqrt(np.arange(1, X.shape[1] + 1))), axis=1)

```

23. Kriging Surrogate Models in SpotOptim

```
        return sum_sq - prod_cos + 1

# Different theta bounds
theta_configs = [
    (-2.0, 1.0, "Tight (smoother)"),
    (-3.0, 2.0, "Default"),
    (-4.0, 3.0, "Wide (more flexible)")
]

for min_theta, max_theta, label in theta_configs:
    kriging = Kriging(
        method='regression',
        min_theta=min_theta,
        max_theta=max_theta,
        seed=42,
        model_fun_evals=50
    )

    optimizer = SpotOptim(
        fun=griewank,
        bounds=[(-600, 600), (-600, 600)],
        surrogate=kriging,
        max_iter=35,
        n_initial=18,
        seed=42,
        verbose=False
    )

    result = optimizer.optimize()
    print(f"{label:20s} [{min_theta:5.1f}, {max_theta:4.1f}] | f(x) = {result.fun:8.6f}")

print("\nDefault bounds work well for most problems.")
```

```
Tight (smoother)      [ -2.0,  1.0] | f(x) = 6.900924
Default               [ -3.0,  2.0] | f(x) = 7.159682
Wide (more flexible) [ -4.0,  3.0] | f(x) = 6.243257
```

Default bounds work well for most problems.

23.6. Comparing Surrogates

23.6.1. Kriging vs Gaussian Process vs Random Forest

Let's compare different surrogate models:

```
import numpy as np
from spotoptim import SpotOptim
from spotoptim.surrogate import Kriging, SimpleKriging
from sklearn.gaussian_process import GaussianProcessRegressor
from sklearn.gaussian_process.kernels import Matern, ConstantKernel
from sklearn.ensemble import RandomForestRegressor

def schwefel(X):
    """Schwefel function"""
    return 418.9829 * X.shape[1] - np.sum(X * np.sin(np.sqrt(np.abs(X))), axis=1)

bounds = [(-500, 500), (-500, 500)]

# 1. Kriging (full-featured)
kriging = Kriging(method='regression', seed=42, model_fun_evals=60)
opt_kriging = SpotOptim(fun=schwefel, bounds=bounds, surrogate=kriging,
                        max_iter=35, n_initial=18, seed=42, verbose=False)
result_kriging = opt_kriging.optimize()

# 2. SimpleKriging (lightweight)
simple_kriging = SimpleKriging(noise=1e-10, seed=42)
opt_simple = SpotOptim(fun=schwefel, bounds=bounds, surrogate=simple_kriging,
                       max_iter=35, n_initial=18, seed=42, verbose=False)
result_simple = opt_simple.optimize()

# 3. Gaussian Process (sklearn default)
kernel = ConstantKernel(1.0) * Matern(length_scale=1.0, nu=2.5)
gp = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=10,
                              normalize_y=True, random_state=42)
opt_gp = SpotOptim(fun=schwefel, bounds=bounds, surrogate=gp,
                   max_iter=35, n_initial=18, seed=42, verbose=False)
result_gp = opt_gp.optimize()

# 4. Random Forest
rf = RandomForestRegressor(n_estimators=100, random_state=42)
opt_rf = SpotOptim(fun=schwefel, bounds=bounds, surrogate=rf,
                   max_iter=35, n_initial=18, seed=42, verbose=False)
result_rf = opt_rf.optimize()
```

23. Kriging Surrogate Models in SpotOptim

```
print("Surrogate Model Comparison:")
print(f"Kriging (full):      f(x) = {result_kriging.fun:8.2f}")
print(f"SimpleKriging:      f(x) = {result_simple.fun:8.2f}")
print(f"Gaussian Process:   f(x) = {result_gp.fun:8.2f}")
print(f"Random Forest:      f(x) = {result_rf.fun:8.2f}")
print(f"\nGlobal optimum: f(420.9687, 420.9687) = 0.0")
print("\nKriging offers:")
print("- Mixed variable types (continuous, integer, categorical)")
print("- Multiple methods (interpolation, regression, reinterpolation)")
print("- Explicit control over regularization and correlation")
```

Surrogate Model Comparison:

Kriging (full):	f(x) =	163.48
SimpleKriging:	f(x) =	118.55
Gaussian Process:	f(x) =	118.45
Random Forest:	f(x) =	183.47

Global optimum: f(420.9687, 420.9687) = 0.0

Kriging offers:

- Mixed variable types (continuous, integer, categorical)
- Multiple methods (interpolation, regression, reinterpolation)
- Explicit control over regularization and correlation

23.7. Best Practices

23.7.1. 1. Choosing the Right Method

```
# For smooth, deterministic functions
kriging = Kriging(method='interpolation', noise=1e-10, seed=42)

# For general optimization (RECOMMENDED)
kriging = Kriging(method='regression', seed=42)

# For noisy functions with repeated evaluations
kriging = Kriging(method='regression', seed=42)
# Use with repeats_initial and repeats_surrogate in SpotOptim
```

23.7.2. 2. Setting Model Complexity

```
# For low-dimensional problems (<5D)
kriging = Kriging(
    method='regression',
    isotropic=False,          # Anisotropic (default)
    model_fun_evals=100,      # More budget for better fit
    seed=42
)

# For high-dimensional problems (>5D)
kriging = Kriging(
    method='regression',
    isotropic=True,           # Fewer parameters
    model_fun_evals=50,       # Faster fitting
    seed=42
)
```

23.7.3. 3. Handling Different Variable Types

```
# Mixed types example
bounds = [
    (0.0, 10.0),    # continuous
    (1, 10),        # integer
    (0, 3)          # categorical (4 categories)
]

var_type = ['float', 'int', 'factor']

kriging = Kriging(
    method='regression',
    var_type=var_type,
    metric_factorial='canberra', # For factor variables
    seed=42
)

optimizer = SpotOptim(
    fun=objective,
    bounds=bounds,
    var_type=var_type,
    surrogate=kriging,
```

23. Kriging Surrogate Models in SpotOptim

```
seed=42  
)
```

23.7.4. 4. Reproducibility

Always set the seed for reproducible results:

```
# Both Kriging and SpotOptim should have seeds  
kriging = Kriging(method='regression', seed=42)  
  
optimizer = SpotOptim(  
    fun=objective,  
    bounds=bounds,  
    surrogate=kriging,  
    seed=42 # Same seed or different seed  
)
```

23.7.5. 5. Monitoring Surrogate Quality

Check the negative log-likelihood after fitting:

```
import numpy as np  
from spotoptim.surrogate import Kriging  
  
# Sample data  
X = np.random.rand(20, 2) * 10 - 5  
y = np.sum(X**2, axis=1)  
  
model = Kriging(method='regression', seed=42)  
model.fit(X, y)  
  
print(f"Negative log-likelihood: {model.negLnLike:.4f}")  
print(f"Optimized theta: {model.theta_}")  
print(f"Optimized Lambda: {model.Lambda_:.4f}")  
print(f"Condition number of Psi: {model.cnd_Psi:.2e}")  
  
print("\nLower negLnLike = better fit")  
print("Check condition number for numerical stability")
```

```
Negative log-likelihood: -15.4700  
Optimized theta: [-2.99943411 -2.99998454]
```

Optimized Lambda: -8.9984
 Condition number of Psi: 3.12e+14

Lower negLnLike = better fit
 Check condition number for numerical stability

23.8. Common Issues and Solutions

23.8.1. Issue 1: Slow Fitting for Large Datasets

Problem: Kriging becomes slow with many training points.

Solution: Limit surrogate training set size:

```
# Use max_surrogate_points in SpotOptim
optimizer = SpotOptim(
    fun=objective,
    bounds=bounds,
    surrogate=kriging,
    max_surrogate_points=200, # Limit to 200 points
    selection_method='distant', # or 'best'
    seed=42
)
```

23.8.2. Issue 2: Poor Predictions for Categorical Variables

Problem: Kriging doesn't handle factors well.

Solution: Try different distance metrics:

```
# Try different metrics
for metric in ['canberra', 'hamming', 'jaccard']:
    kriging = Kriging(
        method='regression',
        var_type=['float', 'factor'],
        metric_factorial=metric,
        seed=42
    )
# Test performance
```

23.8.3. Issue 3: Numerical Instability

Problem: Correlation matrix is nearly singular.

Solution: Increase regularization:

```
# For interpolation method
kriging = Kriging(
    method='interpolation',
    noise=1e-6, # Increase from default 1e-8
    seed=42
)

# Or use regression method (recommended)
kriging = Kriging(
    method='regression',
    min_Lambda=-6.0, # Don't go too small
    seed=42
)
```

23.8.4. Issue 4: Overfitting to Noisy Data

Problem: Kriging fits noise instead of underlying function.

Solution: Use regression method with reasonable Lambda bounds:

```
kriging = Kriging(
    method='regression',
    min_Lambda=-5.0, # Not too small
    max_Lambda=-1.0, # Not too large
    seed=42
)
```

23.9. Summary

23.9.1. Key Takeaways

1. Use `method='regression'` for most optimization problems
2. Set `seed` for reproducibility
3. Use `var_type` for mixed variable types
4. Default parameters work well for most cases
5. Use `isotropic=True` for high-dimensional problems (>5D)

23.9.2. Quick Reference

```

from spotoptim import SpotOptim
from spotoptim.surrogate import Kriging

# Recommended configuration for general use
kriging = Kriging(
    method='regression',      # Smooths over noise
    seed=42                   # Reproducibility
)

optimizer = SpotOptim(
    fun=objective,
    bounds=bounds,
    surrogate=kriging,
    max_iter=50,
    n_initial=20,
    seed=42,
    verbose=True
)

result = optimizer.optimize()

```

23.9.3. Parameter Cheat Sheet

Parameter	Default	When to Change
method	'regression'	Use 'interpolation' for exact fit
noise	$\sqrt{\text{eps}}$	Rarely; use method='regression' instead
min_theta	-3.0	Adjust if default range doesn't work
max_theta	2.0	Adjust if default range doesn't work
isotropic	False	Set True for >5 dimensions
var_type	['num']	Set for mixed variable types
metric_factorial	'canberra'	Try 'hamming' for factors
model_fun_evals	100	Reduce for faster fitting
seed	124	Always set for reproducibility

23.10. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

24. Kriging in SpotOptim: The Forrester Implementation

25. Introduction

The **Kriging** surrogate model in `spotoptim` is an implementation of the methods described in the classic text “**Engineering Design via Surrogate Modelling**” by Forrester, Sóbester, and Keane (2008). This tutorial explains the mathematical foundations of the implementation and demonstrates how to use the different fitting methods available in the library. The implementation has been validated against the original Matlab code provided with the book (`likelihood.m`, `pred.m`, etc.).

26. 1. The Core Kriging Model (Section 2.4)

Kriging, or Gaussian Process Regression, treats the response $y(\mathbf{x})$ as a realization of a stochastic process:

$$Y(\mathbf{x}) = \mu + Z(\mathbf{x})$$

where μ is the mean (often modeled as a constant in Ordinary Kriging) and $Z(\mathbf{x})$ is a Gaussian process with zero mean and covariance:

$$\text{Cov}(Z(\mathbf{x}^{(i)}), Z(\mathbf{x}^{(l)})) = \sigma^2 \Psi(\mathbf{x}^{(i)}, \mathbf{x}^{(l)})$$

26.1. Correlation Function

We use the Gaussian correlation function (Eq 2.4 in the book), which is infinitely differentiable and assumes a smooth underlying response:

$$\Psi(\mathbf{x}^{(i)}, \mathbf{x}^{(l)}) = \exp \left(- \sum_{j=1}^k \theta_j |x_j^{(i)} - x_j^{(l)}|^{p_j} \right)$$

In `spotoptim`, we typically set $p_j = 2$ (Gaussian) and optimize θ_j (length-scales).

26.2. Training: Concentrated Log-Likelihood

To fit the model, we maximize the **Concentrated Log-Likelihood** (Eq 2.30/2.33) with respect to the hyperparameters θ :

$$\ln(L) \approx -\frac{n}{2} \ln(\hat{\sigma}^2) - \frac{1}{2} \ln |\Psi|$$

where $\hat{\sigma}^2$ is the estimated variance (Eq 2.28). This matches the implementation in the book's `likelihood.m`.

26.3. Prediction

The prediction at a new point $\mathbf{x}$ is given by (Eq 2.15):

$$\hat{y}(\mathbf{x}) = \hat{\mu} + \psi^T \Psi^{-1}(\mathbf{y} - \mathbf{1}\hat{\mu})$$

where ψ is the correlation vector between the new point and the training points. This matches `pred.m`.

27. 2. Handling Noisy Data

Real-world data is often noisy. Forrester et al. (Section 6) describe three approaches to handle this, all available in `spotoptim` via the `method` argument.

27.1. A. Interpolation (`method="interpolation"`)

This is the standard Kriging formulation. It assumes **no noise** in the data.

- The model passes exactly through the training points.
- A very small “nugget” (noise) is added only for numerical stability of the matrix inversion.

27.2. B. Regression (`method="regression"`)

Described in Section 6.2 (“Regression Kriging”), this method assumes the observations contain noise ϵ :

$$y = \mu + Z(\mathbf{x}) + \epsilon$$

We introduce a regularization parameter λ (Lambda) to the diagonal of the correlation matrix:

$$\mathbf{C} = \Psi + \Psi^T + (1 + \lambda)\mathbf{I}$$

Both θ and λ are optimized simultaneously. The resulting surrogate:

- **Does not pass through the training points** (it smooths them).
- Is suitable for noisy objective functions.

27.3. C. Re-interpolation (method="reinterpolation")

Described in Section 6.3, this is a hybrid approach.

1. We fit the hyperparameters (θ and λ) using the **Regression** method to learn the noise level.
2. For prediction, we use the “noise-free” correlation matrix (removing λ) to predict the underlying smooth trend.

This “re-interpolates” the regression model, providing a predictor that typically has better error properties for design optimization than pure regression.

28. 3. Examples

We will verify these methods with a simple 1D example.

28.1. Setup

```
import numpy as np
import matplotlib.pyplot as plt
from spotoptim.surrogate import Kriging

# Define a function with noise
def true_function(x):
    return (x * 6 - 2)**2 * np.sin(x * 12 - 4)

np.random.seed(42)
X = np.linspace(0, 1, 11).reshape(-1, 1)
noise = np.random.normal(0, 1.0, X.shape)
y = true_function(X).flatten() + noise.flatten() * 3.0 # Add significant noise

X_plot = np.linspace(0, 1, 100).reshape(-1, 1)
y_true = true_function(X_plot)
```

28.2. Comparisons

We will fit three models and compare them.

```
# 1. Interpolation (Forces fit through noisy points)
model_interp = Kriging(method="interpolation", seed=42)
model_interp.fit(X, y)
y_interp, s_interp = model_interp.predict(X_plot, return_std=True)

# 2. Regression (Smooths the data by learning Lambda)
model_reg = Kriging(method="regression", seed=42, minLambda=-1, maxLambda=10)
model_reg.fit(X, y)
```

28. 3. Examples

```
y_reg, s_reg = model_reg.predict(X_plot, return_std=True)

# 3. Re-interpolation (Smooth trend based on regression parameters)
model_reinterp = Kriging(method="reinterpolation", seed=42)
model_reinterp.fit(X, y)
y_reinterp, s_reinterp = model_reinterp.predict(X_plot, return_std=True)
```


28.3. Visualization

```
fig, axes = plt.subplots(1, 3, figsize=(18, 5))

def plot_model(ax, title, y_pred, std):
    ax.plot(X_plot, y_true, 'k--', label="True Function")
    ax.scatter(X, y, c='r', marker='x', s=100, label="Noisy Data")
    ax.plot(X_plot, y_pred, 'b-', label="Surrogate")
    ax.fill_between(X_plot.flatten(),
                   y_pred - 3*std,
                   y_pred + 3*std,
                   color='b', alpha=0.2, label="Uncertainty (3 $\sigma$ )")
    ax.set_title(title)
    ax.legend()
    ax.set_ylim(-15, 20)

plot_model(axes[0], "Interpolation\n(Fits noise exactly)", y_interp, s_interp)
plot_model(axes[1], f"Regression\n(Smooths data. Lambda={model_reg.Lambda_:.2e})", y_reg, s_reg)
plot_model(axes[2], "Re-interpolation\n(Smooth trend, clean signal)", y_reinterp, s_reinterp)

plt.tight_layout()
plt.show()
```



28.3.1. Interpretation

1. **Interpolation:** Wiggles aggressively to hit every noisy point. This is “overfitting” the noise.
2. **Regression:** Produces a much smoother curve that does not hit the points. It has correctly identified that the data is noisy (optimized $\lambda > 0$).
3. **Re-interpolation:** Similar to regression in shape, but mathematically grounded in predicting the *process* rather than the *data*.

29. 4. Infill Criteria

The book also describes “Infill Criteria” (Section 3) for deciding where to sample next. This is implemented in `SpotOptim` itself (not the surrogate class). When you use `SpotOptim(..., acquisition="ei")`, you are using the **Expected Improvement** criterion described in Section 3.2.1 of the book.

```
from spotoptim import SpotOptim

# Use Kriging with EI
optimizer = SpotOptim(
    fun=lambda x: np.sum(x**2),
    bounds=[(-5, 5)],
    surrogate=model_reinterp, # Pass your customized Kriging model
    acquisition="ei",         # Expected Improvement
    max_iter=20
)
```


30. References

1. Forrester, A., Sobester, A., & Keane, A. (2008). *Engineering Design via Surrogate Modelling: A Practical Guide*. Wiley.
2. **SpotOptim Source Code:** `spotoptim/surrogate/kriging.py` (Validated against book's Matlab code).

30.1. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

31. Nystroem Kernel Approximation

This chapter introduces the Nystroem method for approximating kernel maps, which enables scalable Gaussian Process regression for datasets with a large number of features (k) and a moderate number of samples (n).

31.1. Introduction

Gaussian Process (GP) models, such as Kriging, are powerful tools for regression and surrogate modeling. However, they suffer from computational scalability issues. The standard implementation scales as $\mathcal{O}(n^3)$, where n is the number of samples. Additionally, constructing the distance matrix scales as $\mathcal{O}(n^2k)$, where k is the number of features (dimensions).

When k is very large (high-dimensional data), the distance calculation itself can become a bottleneck. The Nystroem method offers a solution by approximating the feature map of the kernel using a subset of the training data.

31.2. The Nystroem Method

The Nystroem method approximates a kernel $K(x, y)$ by mapping inputs into a lower-dimensional feature space $\mathbb{R}^D$ (where D is the number of components).

$$x \mapsto \phi(x) \in \mathbb{R}^D$$

such that the inner product in this space approximates the kernel:

$$\langle \phi(x), \phi(y) \rangle \approx K(x, y)$$

31. Nystroem Kernel Approximation

31.2.1. Moderate n , Large k Use Case

While Nystroem is often cited for its ability to handle large n (by reducing the effective sample size for kernel construction), it is uniquely powerful for the “Moderate n , Large k ” scenario:

1. **Dimensionality Reduction:** It acts like a non-linear Kernel PCA. It projects high-dimensional data (k) into a feature space of dimension `n_components`.
2. **Efficiency:** If `n_components` $\ll k$, subsequent distance calculations in the GP model become significantly faster ($\mathcal{O}(n^2 \cdot \text{n_components})$) instead of $\mathcal{O}(n^2 k)$.

31.3. Implementation in SpotOptim

`spotoptim` provides a modular implementation compatible with its `Kriging` surrogate. We use a `Pipeline` to chain the Nystroem approximation with the Kriging model.

31.3.1. Components

- `spotoptim.surrogate.nystroem.Nystroem`: The feature map approximator.
- `spotoptim.surrogate.pipeline.Pipeline`: Utility to chain transformers and estimators.
- `spotoptim.surrogate.kernels`: Modular kernels (`RBF`, `WhiteKernel`).

31.3.2. Example Usage

Here is how to set up a GP model with Nystroem approximation for a high-dimensional problem.

```
import numpy as np
from spotoptim.surrogate.kriging import Kriging
from spotoptim.surrogate.nystroem import Nystroem
from spotoptim.surrogate.pipeline import Pipeline
from spotoptim.surrogate.kernels import RBF, WhiteKernel

def define_nystroem_gp_model(n_components=50, noise_level=1e-5):
    """
    Defines a Pipeline with Nystroem approximation and Kriging.

    Args:
        n_components (int): Dimensionality of the target feature space.
        noise_level (float): Noise level for the WhiteKernel.
```

```

Returns:
    Pipeline: A configured pipeline ready to be fitted.
    """
    # 1. Define the kernel for Nystroem
    # We use an RBF kernel plus a small WhiteKernel for stability
    kernel = 1.0 * RBF(length_scale=1.0) + WhiteKernel(noise_level=noise_level)

    # 2. Define Nystroem approximation
    # This will project data from k dimensions -> n_components dimensions
    nystroem = Nystroem(kernel=kernel, n_components=n_components, random_state=42)

    # 3. Define the Kriging model
    # This model will receive inputs of dimension n_components
    gp_model = Kriging(method='regression', seed=42)

    # 4. Create the Pipeline
    gp_model_pipeline = Pipeline([
        ('nystroem', nystroem),
        ('gp', gp_model)
    ])

    return gp_model_pipeline

# --- Demonstration ---
# Generate high-dimensional synthetic data (Moderate n=50, Large k=1000)
rng = np.random.RandomState(42)
n_samples = 50
n_features = 1000

X_train = rng.randn(n_samples, n_features)
# Target depends only on first few features
y_train = np.sum(X_train[:, :5]**2, axis=1) + rng.normal(0, 0.1, n_samples)

# Create and fit model
pipeline = define_nystroem_gp_model(n_components=20) # Reduce 1000 dims -> 20 dims
pipeline.fit(X_train, y_train)

# Predict
X_test = rng.randn(10, n_features)
y_pred = pipeline.predict(X_test)

print(f"Predicted shape: {y_pred.shape}")
print(f"Sample predictions: {y_pred[:3]}")

```

31. Nystroem Kernel Approximation

```
Predicted shape: (10,)  
Sample predictions: [4.6967762 4.6967762 4.6967762]
```

In this example, the Kriging model operates on 20 features instead of 1000, drastically reducing the cost of distance matrix computations during hyperparameter optimization.

31.4. Handling Mixed Variables

Standard kernel functions (like RBF) usually assume continuous input data. `spotoptim` provides a specialized `SpotOptimKernel` that can handle mixed variable types (continuous, integer, factor) correctly. This is essential if your high-dimensional dataset contains categorical variables.

31.4.1. Using SpotOptimKernel

To use mixed variables with Nystroem, you must initialize Nystroem with a `SpotOptimKernel` configured with your variable types.

```
from spotoptim.surrogate.kernels import SpotOptimKernel  
  
# Define variable types for your data features  
# Example: 2 floats, 1 integer, 1 factor  
var_types = ['float', 'float', 'int', 'factor']  
  
# Initial weights/theta for the kernel (usually ones, or optimized externally)  
theta = np.ones(len(var_types))  
  
# 1. Create the Mixed Kernel  
mixed_kernel = SpotOptimKernel(theta=theta, var_type=var_types)  
  
# 2. Use it in Nystroem  
nystroem = Nystroem(kernel=mixed_kernel, n_components=50)  
  
# The pipeline setup remains the same  
pipeline = Pipeline([  
    ('nystroem', nystroem),  
    ('gp', Kriging())  
])
```

When `fit` or `transform` is called, `SpotOptimKernel` will compute distances using the appropriate metric for each variable type (e.g., Hamming/Canberra distance for factors) before the Nystroem projection occurs.

31.5. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

Part VIII.

Multiobjective Optimization

32. Multi-Objective Optimization Support in SpotOptim

```
import copy
import numpy as np
import pandas as pd
import numpy as np
import warnings

from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import MinMaxScaler
from sklearn.gaussian_process import GaussianProcessRegressor
from sklearn.kernel_approximation import Nystroem
from sklearn.gaussian_process.kernels import RBF, WhiteKernel
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score

import matplotlib.pyplot as plt
import seaborn as sns

from scipy.stats.qmc import LatinHypercube
from scipy.optimize import dual_annealing

from spotdesirability.utils.desirability import DMin, DMax, DOverall

from spotoptim.sampling.mm import mmphi_intensive, mmphi_intensive_update, mm_improvement_contour
from spotoptim.inspection.predictions import plot_actual_vs_predicted
from spotoptim.inspection import (generate_mdi, generate_imp, plot_importances, plot_feature_importance)

from spotoptim.utils.boundaries import get_boundaries, map_to_original_scale
from spotoptim.mo.pareto import is_pareto_efficient
from spotoptim.plot.mo import plot_mo
from spotoptim.eda import plot_ip_boxplots, plot_ip_histograms
from spotoptim.utils import get_combinations
```

32. Multi-Objective Optimization Support in SpotOptim

```
from spotoptim.sampling.lhs import rlh
from spotoptim.sampling.mm import (bestlh, plot_mmphi_vs_n_lhs, mmphi, mmphi_intensive)
from spotoptim.mo import mo_mm_desirability_function, mo_mm_desirability_optimizer
from spotoptim.mo.mo_mm import mo_xy_desirability_plot
from spotoptim.function.mo import mo_conv2_max
from spotoptim.mo.pareto import (mo_xy_surface, mo_xy_contour, mo_pareto_optx_plot)
from spotoptim.utils.eval import mo_cv_models, mo_eval_models
warnings.filterwarnings("ignore")
```

```
n_estimators=100
max_depth = None
random_state=42
```

32.1. Overview

SpotOptim supports multi-objective optimization functions with automatic detection and flexible scalarization strategies. The parameter `fun_mo2so` can be used to convert multi-objective values to single-objective values. If the function `fun` returns several objective function values, say $f_1, f_2, \dots, f_p$, these values are stored in `y_mo` and can be used by `fun_mo2so`. The `fun_mo2so` parameter is optional. If it is `None`, which is the default, only the first objective f_1 is used as single-objective and the remaining objectives $f_2, \dots, f_p$ are stored in `y_mo`, but not further used in the optimization process.

`fun_mo2so`

- `fun_mo2so` (callable, optional): Function to convert multi-objective values to single-objective
 - Takes array of shape `(n_samples, n_objectives)`
 - Returns array of shape `(n_samples,)`
 - If `None`, uses first objective (default behavior)

Attributes

- `y_mo` (ndarray or `None`): Stores all multi-objective function values
 - Shape: `(n_samples, n_objectives)` for multi-objective problems
 - `None` for single-objective problems

i Methods

- `get_shape(y)`: Get shape of objective function output
- `store_mo(y_mo)`: Store multi-objective values with automatic appending
- `mo2so(y_mo)`: Convert multi-objective to single-objective values

`SpotOptim`'s method `evaluate_function(X)` automatically detects multi-objective functions. It calls `mo2so()` to convert multi-objective to single-objective. It also stores the original multi-objective values in `y_mo`. And it returns single-objective values for optimization.

Example 32.1 (Default Behavior (Use First Objective)).

Example 32.2.

```
import numpy as np
from spotoptim import SpotOptim

def bi_objective(X):
    """Two conflicting objectives."""
    obj1 = np.sum(X**2, axis=1)          # Minimize at origin
    obj2 = np.sum((X - 2)**2, axis=1)    # Minimize at (2, 2)
    return np.column_stack([obj1, obj2])

optimizer = SpotOptim(
    fun=bi_objective,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=30,
    n_initial=15,
    seed=42,
    max_surrogate_points=30
)

result = optimizer.optimize()

print(f"Best x: {result.x}")            # Near [0, 0]
print(f"Best f(x): {result.fun}")       # Minimizes obj1
print(f"M0 values stored: {optimizer.y_mo.shape}") # (30, 2)
```

Best x: [-4.39940747e-05 -2.08039138e-04]

Best f(x): 4.521576166022012e-08

M0 values stored: (30, 2)

Example 32.3 (Weighted Sum Scalarization).

32. Multi-Objective Optimization Support in SpotOptim

Example 32.4.

```
def weighted_sum(y_mo):
    """Equal weighting of objectives."""
    return 0.5 * y_mo[:, 0] + 0.5 * y_mo[:, 1]

optimizer = SpotOptim(
    fun=bi_objective,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=30,
    n_initial=15,
    fun_mo2so=weighted_sum, # Custom conversion
    seed=42,
    max_surrogate_points=30
)

result = optimizer.optimize()
print(f"Compromise solution: {result.x}") # Near [1, 1]
```

Compromise solution: [0.99819069 0.99988984]

Example 32.5 (Min-Max Scalarization).

Example 32.6.

```
def min_max(y_mo):
    """Minimize the maximum objective."""
    return np.max(y_mo, axis=1)

optimizer = SpotOptim(
    fun=bi_objective,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=30,
    n_initial=15,
    fun_mo2so=min_max,
    seed=42,
    max_surrogate_points=30
)

result = optimizer.optimize()
# Finds solution with balanced objective values
```

Example 32.7 (Three or More Objectives).

Example 32.8.

```
def three_objectives(X):
    """Three different norms."""
    obj1 = np.sum(X**2, axis=1)      # L2 norm
    obj2 = np.sum(np.abs(X), axis=1) # L1 norm
    obj3 = np.max(np.abs(X), axis=1) # L-infinity norm
    return np.column_stack([obj1, obj2, obj3])

def custom_scalarization(y_mo):
    """Weighted combination."""
    return 0.4 * y_mo[:, 0] + 0.3 * y_mo[:, 1] + 0.3 * y_mo[:, 2]

optimizer = SpotOptim(
    fun=three_objectives,
    bounds=[(-5, 5), (-5, 5), (-5, 5)],
    max_iter=35,
    n_initial=20,
    fun_mo2so=custom_scalarization,
    seed=42,
    max_surrogate_points=30
)

result = optimizer.optimize()
```

Example 32.9 (Noise Handling).**Example 32.10.**

```
def noisy_bi_objective(X):
    """Noisy multi-objective function."""
    noise1 = np.random.normal(0, 0.05, X.shape[0])
    noise2 = np.random.normal(0, 0.05, X.shape[0])

    obj1 = np.sum(X**2, axis=1) + noise1
    obj2 = np.sum((X - 1)**2, axis=1) + noise2
    return np.column_stack([obj1, obj2])

optimizer = SpotOptim(
    fun=noisy_bi_objective,
    bounds=[(-5, 5), (-5, 5)],
    max_iter=40,
    n_initial=10,
    repeats_initial=3,      # Handle noise
)
```

32. Multi-Objective Optimization Support in SpotOptim

```
    repeats_surrogate=2,  
    seed=42,  
    max_surrogate_points=30  
)  
  
result = optimizer.optimize()  
# Works seamlessly with noise handling
```

32.2. Common Scalarization Strategies

32.2.1. 1. Weighted Sum

```
def weighted_sum(y_mo, weights=[0.5, 0.5]):  
    return sum(w * y_mo[:, i] for i, w in enumerate(weights))
```

Use **when**: Objectives have similar scales and you want linear trade-offs

32.2.2. 2. Weighted Sum with Normalization

```
def normalized_weighted_sum(y_mo, weights=[0.5, 0.5]):  
    # Normalize each objective to [0, 1]  
    y_norm = (y_mo - y_mo.min(axis=0)) / (y_mo.max(axis=0) - y_mo.min(axis=0) + 1e-10)  
    return sum(w * y_norm[:, i] for i, w in enumerate(weights))
```

Use **when**: Objectives have very different scales

32.2.3. 3. Min-Max (Chebyshev)

```
def min_max(y_mo):  
    return np.max(y_mo, axis=1)
```

Use **when**: You want balanced performance across all objectives

32.2.4. 4. Target Achievement

```
def target_achievement(y_mo, targets=[0.0, 0.0]):
    # Minimize deviation from targets
    return np.sum((y_mo - targets)**2, axis=1)
```

Use when: You have specific target values for each objective

32.2.5. 5. Product

```
def product(y_mo):
    return np.prod(y_mo + 1e-10, axis=1) # Add small value to avoid zero
```

Use when: All objectives should be minimized together

32.3. Integration with Other Features

Multi-objective support works seamlessly with:

- ☒ **Noise Handling** - Use `repeats_initial` and `repeats_surrogate`
- ☒ **OCBA** - Use `ocba_delta` for intelligent re-evaluation
- ☒ **TensorBoard Logging** - Logs converted single-objective values
- ☒ **Dimension Reduction** - Fixed dimensions work normally
- ☒ **Custom Variable Names** - `var_name` parameter supported

32.4. Implementation Details

32.4.1. Automatic Detection

SpotOptim automatically detects multi-objective functions:

- If function returns 2D array (`n_samples`, `n_objectives`), it's multi-objective
- If function returns 1D array (`n_samples`), it's single-objective

32.4.2. Data Flow

User Function $\rightarrow$ y_mo (raw) $\rightarrow$ mo2so() $\rightarrow$ y_ (single-objective)
 $\downarrow$
y_mo (stored)

1. Function returns multi-objective values
2. `store_mo()` saves them in `y_mo` attribute
3. `mo2so()` converts to single-objective using `fun_mo2so` or default
4. Surrogate model optimizes the single-objective values
5. All original multi-objective values remain accessible in `y_mo`

32.4.3. Backward Compatibility

- ☒ Fully backward compatible:

- Single-objective functions work unchanged
- `fun_mo2so` defaults to `None`
- `y_mo` is `None` for single-objective problems
- No breaking changes to existing code

32.5. Limitations and Notes

32.5.1. What This Is

- ☒ Scalarization approach to multi-objective optimization
- ☒ Single solution found per optimization run
- ☒ Different scalarizations $\rightarrow$ different Pareto solutions
- ☒ Suitable for preference-based multi-objective optimization

32.5.2. What This Is Not

- ❌ Not a true multi-objective optimizer (doesn't find Pareto front)
- ❌ Doesn't generate multiple solutions in one run
- ❌ Not suitable for discovering entire Pareto front

32.5.3. For True Multi-Objective Optimization

For finding the complete Pareto front, consider specialized tools:

- **pymoo**: Comprehensive multi-objective optimization framework
- **platypus**: Multi-objective optimization library
- **NSGA-II**, **MOEA/D**: Dedicated multi-objective algorithms

32.6. Benchmark Multi-Objective Test Functions

SpotOptim includes 12 multi-objective benchmark functions in `spotoptim.function.mo`. These functions are widely used to evaluate and compare multi-objective optimization algorithms.

32.6.1. Function Overview

Function	Objectives	Variables	Pareto Front	Characteristics
ZDT1	2	30	Convex	Unimodal, minimization
ZDT2	2	30	Non-convex	Unimodal, minimization
ZDT3	2	30	Disconnected	Multimodal, minimization
ZDT4	2	10	Convex	Many local fronts
ZDT6	2	10	Non-convex	Non-uniform density
DTLZ1	Scalable	$n_{\text{obj}}+4$	Linear	Multimodal
DTLZ2	Scalable	$n_{\text{obj}}+9$	Spherical	Unimodal
Schaffer N1	2	1	Convex	Simple, minimization
Fonseca-Fleming	2	2-10	Concave	Symmetric, minimization
mo_conv2_min	2	2	Convex	Convex, minimization
mo_conv2_max	2	2	Convex	Convex, maximization
Kursawe	2	3	Disconnected	Non-convex, disconnected minimization

32. Multi-Objective Optimization Support in SpotOptim

`mo_conv2_min` is a smooth convex minimization pair with objectives $f_1(x, y) = x^2 + y^2$ and $f_2(x, y) = (x - 1)^2 + (y - 1)^2$ on $[0, 1]^2$, producing a convex quadratic Pareto front along $x = y$. `mo_conv2_max` flips both objectives for maximization via $f_1(x, y) = 2 - (x^2 + y^2)$ and $f_2(x, y) = 2 - ((1 - x)^2 + (1 - y)^2)$, keeping objective values bounded in $[0, 2]$.

32.6.2. Example 6: ZDT1 - Convex Pareto Front

ZDT1 is a classical bi-objective test problem with a convex Pareto front, widely used for benchmarking.

```
import numpy as np
import matplotlib.pyplot as plt
from spotoptim.function.mo import zdt1
from spotoptim import SpotOptim

# Generate data for visualization
n_points = 100
x1_range = np.linspace(0, 1, n_points)

# Pareto front: x1 in [0,1], rest = 0
X_pareto = np.column_stack([x1_range, np.zeros((n_points, 2))])
y_pareto = zdt1(X_pareto)

# Non-optimal points for comparison
X_non_optimal = np.random.rand(200, 3)
y_non_optimal = zdt1(X_non_optimal)
```

Figure 32.1 visualizes the Pareto front:

```
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))

# Plot objective space
ax1.scatter(y_non_optimal[:, 0], y_non_optimal[:, 1],
            alpha=0.3, label='Non-optimal', s=20)
ax1.plot(y_pareto[:, 0], y_pareto[:, 1],
        'r-', linewidth=2, label='Pareto Front')
ax1.set_xlabel('f1(x) = x1')
ax1.set_ylabel('f2(x) = g(x) * [1 - sqrt(x1/g(x))]\')
ax1.set_title('ZDT1: Convex Pareto Front')
ax1.legend()
ax1.grid(True, alpha=0.3)
```

32.6. Benchmark Multi-Objective Test Functions

```
# Plot decision space (first two variables)
ax2.scatter(X_non_optimal[:, 0], X_non_optimal[:, 1],
            alpha=0.3, label='Non-optimal', s=20)
ax2.scatter(X_pareto[:, 0], X_pareto[:, 1],
            c='red', label='Pareto Optimal', s=30)
ax2.set_xlabel('x1')
ax2.set_ylabel('x2')
ax2.set_title('Decision Space (first 2 dimensions)')
ax2.legend()
ax2.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```




Figure 32.1.: ZDT1 Convex Pareto Front

```
print("ZDT1 Characteristics:")
print("- Convex Pareto front:  $f_2 = 1 - \sqrt{f_1}$ ")
print("- Optimal when  $x_2, x_3, \dots, x_n = 0$ ")
print("- Easy to solve, good for testing basic functionality")
```

ZDT1 Characteristics:

- Convex Pareto front: $f_2 = 1 - \sqrt{f_1}$
- Optimal when $x_2, x_3, \dots, x_n = 0$
- Easy to solve, good for testing basic functionality

Now optimize using SpotOptim with weighted sum:

32. Multi-Objective Optimization Support in SpotOptim

```
def weighted_sum_zdt1(y_mo):
    """Equal weighting for ZDT1."""
    return 0.5 * y_mo[:, 0] + 0.5 * y_mo[:, 1]

optimizer = SpotOptim(
    fun=zdt1,
    bounds=[(0, 1)] * 30, # 30 variables in [0, 1]
    max_iter=50,
    n_initial=20,
    fun_mo2so=weighted_sum_zdt1,
    seed=42,
    max_surrogate_points=30
)

result = optimizer.optimize()

print(f"\nOptimization Result:")
print(f"Best x (first 5 dims): {result.x[:5]}")
print(f"Best scalarized value: {result.fun:.6f}")

# Evaluate multi-objective values
y_best = zdt1(result.x.reshape(1, -1))
print(f"f1 = {y_best[0, 0]:.6f}, f2 = {y_best[0, 1]:.6f}")
```

```
Optimization Result:
Best x (first 5 dims): [1. 0. 0. 0. 0.]
Best scalarized value: 0.500000
f1 = 1.000000, f2 = 0.000000
```

32.6.3. Example 7: ZDT2 - Non-Convex Pareto Front

ZDT2 has a non-convex Pareto front, making it more challenging than ZDT1.

```
from spotoptim.function.mo import zdt2

# Generate Pareto front
X_pareto = np.column_stack([x1_range, np.zeros((n_points, 2))])
y_pareto = zdt2(X_pareto)

# Non-optimal points
y_non_optimal = zdt2(X_non_optimal)
```

Figure 32.2 visualizes the Pareto front:

```
fig, ax = plt.subplots(figsize=(8, 6))
ax.scatter(y_non_optimal[:, 0], y_non_optimal[:, 1],
          alpha=0.3, label='Non-optimal', s=20)
ax.plot(y_pareto[:, 0], y_pareto[:, 1],
       'r-', linewidth=2, label='Pareto Front')
ax.set_xlabel('f1(x) = x1')
ax.set_ylabel('f2(x) = g(x) * [1 - (x1/g(x))2]')
ax.set_title('ZDT2: Non-Convex Pareto Front')
ax.legend()
ax.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```




Figure 32.2.: ZDT2 Non-Convex Pareto Front

```
print("ZDT2 Characteristics:")
print("- Non-convex Pareto front: f2 = 1 - f12")
```

32. Multi-Objective Optimization Support in SpotOptim

```
print("- Tests algorithm's ability to handle non-convexity")
print("- Optimal when x2, x3, ..., xn = 0")
```

ZDT2 Characteristics:

- Non-convex Pareto front: $f_2 = 1 - f_1^2$
- Tests algorithm's ability to handle non-convexity
- Optimal when $x_2, x_3, \dots, x_n = 0$

32.6.4. Example 8: ZDT3 - Disconnected Pareto Front

ZDT3 has a discontinuous Pareto front consisting of 5 separate regions.

```
from spotoptim.function.mo import zdt3

# Generate Pareto front
X_pareto = np.column_stack([x1_range, np.zeros((n_points, 2))])
y_pareto = zdt3(X_pareto)

# Non-optimal points
y_non_optimal = zdt3(X_non_optimal)
```

Figure 32.3 visualizes the Pareto front:

```
fig, ax = plt.subplots(figsize=(8, 6))
ax.scatter(y_non_optimal[:, 0], y_non_optimal[:, 1],
           alpha=0.3, label='Non-optimal', s=20)
ax.plot(y_pareto[:, 0], y_pareto[:, 1],
        'r-', linewidth=2, label='Pareto Front')
ax.set_xlabel('f1(x) = x1')
ax.set_ylabel('f2(x) with sin term')
ax.set_title('ZDT3: Disconnected Pareto Front')
ax.legend()
ax.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```





Figure 32.3.: ZDT3 Disconnected Pareto Front

```
print("ZDT3 Characteristics:")
print("- Disconnected Pareto front with 5 separate regions")
print("- Tests diversity maintenance in algorithms")
print("- sin(10*x1) term creates discontinuities")
```

```
ZDT3 Characteristics:
- Disconnected Pareto front with 5 separate regions
- Tests diversity maintenance in algorithms
- sin(10*x1) term creates discontinuities
```

32.6.5. Example 9: ZDT4 - Multimodal with Many Local Fronts

ZDT4 has 21^9 local Pareto fronts, testing the algorithm's ability to escape local optima.

32. Multi-Objective Optimization Support in SpotOptim

```
from spotoptim.function.mo import zdt4

# Generate Pareto front (same as ZDT1 when optimal)
X_pareto = np.column_stack([x1_range, np.zeros((n_points, 9))])
y_pareto = zdt4(X_pareto)

# Non-optimal points with multimodal g-function
X_non_optimal_zdt4 = np.random.rand(200, 10)
X_non_optimal_zdt4[:, 0] = np.clip(X_non_optimal_zdt4[:, 0], 0, 1) # x1 in [0,1]
X_non_optimal_zdt4[:, 1:] = X_non_optimal_zdt4[:, 1:] * 10 - 5 # others in [-5,5]
y_non_optimal = zdt4(X_non_optimal_zdt4)
```

Figure 32.4 visualizes the Pareto front:

```
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))

# Objective space
ax1.scatter(y_non_optimal[:, 0], y_non_optimal[:, 1],
            alpha=0.3, label='Non-optimal', s=20)
ax1.plot(y_pareto[:, 0], y_pareto[:, 1],
        'r-', linewidth=2, label='Pareto Front')
ax1.set_xlabel('f1(x) = x1')
ax1.set_ylabel('f2(x)')
ax1.set_title('ZDT4: Multimodal - Objective Space')
ax1.legend()
ax1.grid(True, alpha=0.3)

# Show multimodal landscape for x2
x2_range = np.linspace(-5, 5, 100)
g_vals = 1 + 10 * 9 + (x2_range**2 - 10 * np.cos(4 * np.pi * x2_range)) * 9
ax2.plot(x2_range, g_vals)
ax2.set_xlabel('x2 (or x3, ..., x10)')
ax2.set_ylabel('Contribution to g(x)')
ax2.set_title('Multimodal Landscape')
ax2.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```

32.6. Benchmark Multi-Objective Test Functions




Figure 32.4.: ZDT4 Multimodal Objective Space

```
print("ZDT4 Characteristics:")
print("- 21^9 local Pareto fronts")
print("- Tests global search capability")
print("- x1 in [0,1], x_i in [-5,5] for i=2,...,10")
```

ZDT4 Characteristics:

- 21^9 local Pareto fronts
- Tests global search capability
- x_1 in $[0,1]$, x_i in $[-5,5]$ for $i=2,\dots,10$

32.6.6. Example 10: ZDT6 - Non-Uniform Density

ZDT6 has a non-uniform search space with low density of solutions near the Pareto front.

```
from spotoptim.function.mo import zdt6

# Generate Pareto front
X_pareto = np.column_stack([x1_range, np.zeros((n_points, 2))])
y_pareto = zdt6(X_pareto)

# Non-optimal points
y_non_optimal = zdt6(X_non_optimal)
```

Figure 32.5 visualizes the Pareto front:

32. Multi-Objective Optimization Support in SpotOptim

```
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))

# Objective space
ax1.scatter(y_non_optimal[:, 0], y_non_optimal[:, 1],
            alpha=0.3, label='Non-optimal', s=20)
ax1.plot(y_pareto[:, 0], y_pareto[:, 1],
        'r-', linewidth=2, label='Pareto Front')
ax1.set_xlabel('f1(x)')
ax1.set_ylabel('f2(x)')
ax1.set_title('ZDT6: Non-Uniform Density')
ax1.legend()
ax1.grid(True, alpha=0.3)

# Show f1 non-uniformity
ax2.plot(x1_range, 1 - np.exp(-4 * x1_range) * (np.sin(6 * np.pi * x1_range)**6))
ax2.set_xlabel('x1')
ax2.set_ylabel('f1(x1)')
ax2.set_title('Non-uniform f1 mapping')
ax2.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```




Figure 32.5.: ZDT6 Non-Uniform Density

```
print("ZDT6 Characteristics:")
print("- Non-uniform density of Pareto optimal solutions")
print("- Low density near the Pareto front")
print("- Tests diversity maintenance")
```

ZDT6 Characteristics:

- Non-uniform density of Pareto optimal solutions
- Low density near the Pareto front
- Tests diversity maintenance

32.6.7. Example 11: DTLZ2 - Scalable Spherical Pareto Front

DTLZ2 is a scalable test problem with a concave spherical Pareto front.

```
from spotoptim.function.mo import dtlz2

# 3D visualization for 3 objectives
n_samples = 500
X_samples = np.random.rand(n_samples, 12)
y_samples = dtlz2(X_samples, n_obj=3)

# Generate Pareto optimal points (on unit sphere)
theta = np.linspace(0, np.pi/2, 30)
phi = np.linspace(0, np.pi/2, 30)
theta_grid, phi_grid = np.meshgrid(theta, phi)

f1 = np.cos(theta_grid) * np.cos(phi_grid)
f2 = np.cos(theta_grid) * np.sin(phi_grid)
f3 = np.sin(theta_grid)
```

Figure 32.6 visualizes the Pareto front:

```
fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111, projection='3d')

# Pareto front surface
ax.plot_surface(f1, f2, f3, alpha=0.3, color='red', label='Pareto Front')

# Sample points
colors = np.linalg.norm(y_samples, axis=1)
scatter = ax.scatter(y_samples[:, 0], y_samples[:, 1], y_samples[:, 2],
                    c=colors, cmap='viridis', alpha=0.5, s=10)

ax.set_xlabel('f1(x)')
ax.set_ylabel('f2(x)')
ax.set_zlabel('f3(x)')
ax.set_title('DTLZ2: Spherical Pareto Front (3 objectives)')
plt.colorbar(scatter, label='Distance from origin')
```

32. Multi-Objective Optimization Support in SpotOptim

```
plt.tight_layout()
plt.show()
```




Figure 32.6.: DTLZ2 Spherical Pareto Front (3 objectives)

```
print("DTLZ2 Characteristics:")
print("- Scalable number of objectives")
print("- Pareto front is unit sphere:  $\sum f_i^2 = 1$ ")
print("- Concave, unimodal")
print(f"- Sample: sum of squares = {np.sum(y_samples[0]**2):.4f}")
```

```
DTLZ2 Characteristics:
- Scalable number of objectives
- Pareto front is unit sphere:  $\sum f_i^2 = 1$ 
```

32.6. Benchmark Multi-Objective Test Functions

- Concave, unimodal
- Sample: sum of squares = 4.0966

Optimize DTLZ2 with 3 objectives:

```
def weighted_sum_3obj(y_mo):
    """Equal weighting for 3 objectives."""
    return np.mean(y_mo, axis=1)

optimizer = SpotOptim(
    fun=lambda X: dtlz2(X, n_obj=3),
    bounds=[(0, 1)] * 12, # 12 variables for 3 objectives
    max_iter=50,
    n_initial=25,
    fun_mo2so=weighted_sum_3obj,
    seed=42,
    max_surrogate_points=30
)

result = optimizer.optimize()

# Evaluate
y_best = dtlz2(result.x.reshape(1, -1), n_obj=3)
print(f"\nOptimization Result:")
print(f"f1 = {y_best[0, 0]:.4f}, f2 = {y_best[0, 1]:.4f}, f3 = {y_best[0, 2]:.4f}")
print(f"Sum of squares: {np.sum(y_best[0]**2):.4f} (should be ~1.0)")
```

Optimization Result:
f1 = 0.0000, f2 = 0.0000, f3 = 1.2504
Sum of squares: 1.5635 (should be ~1.0)

32.6.8. Example 12: Schaffer N1 - Simple Bi-Objective

Schaffer N1 is one of the earliest and simplest multi-objective test functions.

```
from spotoptim.function.mo import schaffer_n1

# Generate data
x_range = np.linspace(-10, 10, 200).reshape(-1, 1)
y_schaffer = schaffer_n1(x_range)

# Pareto optimal range [0, 2]
```

32. Multi-Objective Optimization Support in SpotOptim

```
x_pareto = np.linspace(0, 2, 100).reshape(-1, 1)
y_pareto = schaffer_n1(x_pareto)
```

Figure 32.7 visualizes the Pareto front:

```
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))

# Objective space
ax1.plot(y_schaffer[:, 0], y_schaffer[:, 1], 'b-', alpha=0.3, label='All points')
ax1.plot(y_pareto[:, 0], y_pareto[:, 1], 'r-', linewidth=2, label='Pareto Front')
ax1.set_xlabel('f1(x) = x2')
ax1.set_ylabel('f2(x) = (x-2)2')
ax1.set_title('Schaffer N1: Objective Space')
ax1.legend()
ax1.grid(True, alpha=0.3)

# Decision space
ax2.plot(x_range, y_schaffer[:, 0], label='f1(x) = x2')
ax2.plot(x_range, y_schaffer[:, 1], label='f2(x) = (x-2)2')
ax2.axvspan(0, 2, alpha=0.2, color='red', label='Pareto Optimal')
ax2.set_xlabel('x')
ax2.set_ylabel('Objective values')
ax2.set_title('Decision Space')
ax2.legend()
ax2.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```




Figure 32.7.: Schaffer N1 Objective Space


```
print("Schaffer N1 Characteristics:")
print("- Simplest multi-objective function")
print("- 1D decision variable")
print("- Pareto optimal: x ∈ [0, 2]")
print("- Convex Pareto front")
```

```
Schaffer N1 Characteristics:
- Simplest multi-objective function
- 1D decision variable
- Pareto optimal: x ∈ [0, 2]
- Convex Pareto front
```

32.6.9. Example 13: Fonseca-Fleming - Symmetric Concave Front

Fonseca-Fleming is a classical bi-objective problem with a concave Pareto front. It is defined for 2 to 10 decision variables as follows:

$$f_1(x) = 1 - \exp\left(-\sum_{i=1}^n \left(x_i - \frac{1}{\sqrt{n}}\right)^2\right) \quad f_2(x) = 1 - \exp\left(-\sum_{i=1}^n \left(x_i + \frac{1}{\sqrt{n}}\right)^2\right),$$

where $x_i \in [-2, 2]$. The Pareto optimal solutions satisfy:

$$\sum_{i=1}^n x_i^2 = \frac{1}{n}.$$

```
from spotoptim.function.mo import fonseca_fleming

# Generate data for 2D
n_grid = 50
x1_ff = np.linspace(-2, 2, n_grid)
x2_ff = np.linspace(-2, 2, n_grid)
X1_grid, X2_grid = np.meshgrid(x1_ff, x2_ff)

# Evaluate on grid
X_grid = np.column_stack([X1_grid.ravel(), X2_grid.ravel()])
y_grid = fonseca_fleming(X_grid)

# Pareto optimal points (approximately)
X_pareto_ff = np.linspace(-1/np.sqrt(2), 1/np.sqrt(2), 100)
X_pareto_2d = np.column_stack([X_pareto_ff, X_pareto_ff])
y_pareto_ff = fonseca_fleming(X_pareto_2d)
```

Figure 32.8 visualizes the Pareto front:

32. Multi-Objective Optimization Support in SpotOptim

```
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))

# Objective space
scatter = ax1.scatter(y_grid[:, 0], y_grid[:, 1],
                      c=np.linalg.norm(y_grid, axis=1),
                      cmap='viridis', alpha=0.5, s=10)
ax1.plot(y_pareto_ff[:, 0], y_pareto_ff[:, 1],
        'r-', linewidth=2, label='Approx. Pareto Front')
ax1.set_xlabel('f1(x)')
ax1.set_ylabel('f2(x)')
ax1.set_title('Fonseca-Fleming: Concave Front')
ax1.legend()
ax1.grid(True, alpha=0.3)
plt.colorbar(scatter, ax=ax1, label='Distance from origin')

# Decision space
ax2.contourf(X1_grid, X2_grid, y_grid[:, 0].reshape(n_grid, n_grid),
             levels=20, cmap='viridis', alpha=0.6)
ax2.plot(X_pareto_2d[:, 0], X_pareto_2d[:, 1],
        'r-', linewidth=2, label='Pareto Optimal Region')
ax2.set_xlabel('x1')
ax2.set_ylabel('x2')
ax2.set_title('Decision Space (f1)')
ax2.legend()
plt.colorbar(ax2.contourf(X1_grid, X2_grid, y_grid[:, 0].reshape(n_grid, n_grid),
                        levels=20, cmap='viridis', alpha=0.6), ax=ax2)

plt.tight_layout()
plt.show()
```

32.7. Example 14: Kursawe - Non-Convex Disconnected Front




Figure 32.8.: Fonseca-Fleming Concave Front

```
print("Fonseca-Fleming Characteristics:")
print("- Concave Pareto front (minimization)")
print("- Symmetric problem")
print("- Scalable to higher dimensions")
```

Fonseca-Fleming Characteristics:

- Concave Pareto front (minimization)
- Symmetric problem
- Scalable to higher dimensions

32.7. Example 14: Kursawe - Non-Convex Disconnected Front

Optimize Kursawe with min-max scalarization:

```
from spotoptim.function.mo import kursawe
def min_max_kursawe(y_mo):
    """Minimize maximum objective (Chebyshev)."""
    return np.max(y_mo, axis=1)

optimizer = SpotOptim(
    fun=kursawe,
    bounds=[(-5, 5)] * 3,
    max_iter=50,
    n_initial=20,
```

```

        fun_mo2so=min_max_kursawe,
        seed=42,
        max_surrogate_points=30
    )

result = optimizer.optimize()

y_best = kursawe(result.x.reshape(1, -1))
print(f"\nOptimization Result:")
print(f"Best x: {result.x}")
print(f"f1 = {y_best[0, 0]:.4f}, f2 = {y_best[0, 1]:.4f}")
print(f"Max objective: {max(y_best[0]):.4f}")

```

```

Optimization Result:
Best x: [-3.59284321  4.89413033 -0.90817413]
f1 = -6.6645, f2 = -3.7006
Max objective: -3.7006

```

32.7.1. Comparison: Different Scalarization Strategies on ZDT1

Let's compare how different scalarization strategies find different Pareto solutions:

```

scalarizations = {
    'First Objective': None, # Default
    'Equal Weight': lambda y: 0.5 * y[:, 0] + 0.5 * y[:, 1],
    'Favor f1': lambda y: 0.8 * y[:, 0] + 0.2 * y[:, 1],
    'Favor f2': lambda y: 0.2 * y[:, 0] + 0.8 * y[:, 1],
    'Min-Max': lambda y: np.max(y, axis=1),
}

results = {}
for name, scalarization in scalarizations.items():
    opt = SpotOptim(
        fun=zdt1,
        bounds=[(0, 1)] * 30,
        max_iter=40,
        n_initial=15,
        fun_mo2so=scalarization,
        seed=42,
        verbose=0,
        max_surrogate_points=30
    )

```

32.7. Example 14: Kursawe - Non-Convex Disconnected Front

```
res = opt.optimize()
y_res = zdt1(res.x.reshape(1, -1))
results[name] = (res.x, y_res[0])
```

Figure 32.9 visualizes the results:

```
fig, ax = plt.subplots(figsize=(10, 6))

# Plot Pareto front
X_pareto = np.column_stack([x1_range, np.zeros((n_points, 29))])
y_pareto = zdt1(X_pareto)
ax.plot(y_pareto[:, 0], y_pareto[:, 1], 'k-', linewidth=2,
        label='True Pareto Front', alpha=0.3)

# Plot optimization results
colors = ['blue', 'green', 'orange', 'purple', 'red']
for (name, (x, y)), color in zip(results.items(), colors):
    ax.scatter(y[0], y[1], s=100, c=color, marker='*',
              label=f'{name}: ({y[0]:.3f}, {y[1]:.3f})', zorder=5)

ax.set_xlabel('f1(x)')
ax.set_ylabel('f2(x)')
ax.set_title('Different Scalarization Strategies Find Different Pareto Solutions')
ax.legend()
ax.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

print("\nComparison Results:")
for name, (x, y) in results.items():
    print(f"{name:20s}: f1={y[0]:.4f}, f2={y[1]:.4f}")
```

32. Multi-Objective Optimization Support in SpotOptim




Figure 32.9.

Comparison Results:

First Objective	: f1=0.0000, f2=5.5965
Equal Weight	: f1=1.0000, f2=0.1656
Favor f1	: f1=0.0000, f2=1.7932
Favor f2	: f1=1.0000, f2=0.3476
Min-Max	: f1=1.0000, f2=0.9458

32.8. A Simple Visualization Example

The `mo_mm_desirability_function` computes the negative combined desirability for a candidate point $\mathbf{x}$. It predicts the objective values using the provided surrogate models, computes the Morris-Mitchell improvement if requested, and then calculates the overall desirability using the provided `DOverall` object. The function returns the negative desirability (for minimization) and the list of individual objective values. Consider the following example usage:

Generate `X_base` in the range $[0,1]$ as input data and evaluate the two-objective maximization function `mo_conv2_max`.

A smooth, convex two-objective maximization problem on $[0, 1]^2$ using flipped versions of the minimization objectives:

32.8. A Simple Visualization Example

$$f1(x, y) = 2 - (x^2 + y^2) \quad (32.1)$$

$$f2(x, y) = 2 - ((1 - x)^2 + (1 - y)^2) \quad (32.2)$$

It has the following properties:

- Domain: $[0, 1]^2$
- Objectives: maximize both $f1$ and $f2$
- Ideal points: $(0, 0)$ for $f1$ (gives $f1 = 2$); $(1, 1)$ for $f2$ (gives $f2 = 2$)
- Pareto set: line $x = y$ in $[0, 1]$
- Pareto front: convex quadratic trade-off $f1 = 2 - 2t^2$, $f2 = 2 - 2(1 - t)^2$, $t \in [0, 1]$

```
X_base = np.random.rand(500, 2)
y = mo_conv2_max(X_base)
```

The RandomForestRegressor models are fitted for each objective.

```
models = []
for i in range(y.shape[1]):
    model = RandomForestRegressor(n_estimators=n_estimators, random_state=random_state, max_depth=max_depth)
    model.fit(X_base, y[:, i])
    models.append(model)
```

```
# calculate base Morris-Mitchell stats
phi_base, J_base, d_base = mmphi_intensive(X_base, q=2, p=2)
print(f"phi_base: {phi_base}, J_base: {J_base}, d_base: {d_base}")
```

```
phi_base: 6.212013258255464, J_base: [1 1 1 ... 1 1 1], d_base: [0.00133909 0.00224301 0.00243956 ...]
```

Figure 32.10 shows the surfaces of the two objectives using the fitted models.

```
x_min = 0
x_max = 1
combinations = [(0, 1)]
target_names = ["y1", "y2"]

mo_xy_surface(models, bounds=[(x_min, x_max), (x_min, x_max)], target_names=target_names)
```

32. Multi-Objective Optimization Support in SpotOptim




Figure 32.10.: Surfaces of the two objectives of the mo_conv2_max function

Figure 32.11 shows the contours of the two objectives using the fitted models.

```
x_min = 0
x_max = 1
mo_xy_contour(models, bounds=[(x_min, x_max), (x_min, x_max)], target_names=target_names)
```




Figure 32.11.: Contours of the two objectives of the mo_conv2_max function

Figure 32.12 shows the Pareto front of the original data.

32.8. A Simple Visualization Example

```
plot_mo(y_orig=y, target_names=target_names, combinations=combinations, pareto="max", pareto_front_c
```




Figure 32.12.: Pareto front of the original data

Figure 32.13 visualizes the Pareto-optimal points, i.e., the points in the input space that correspond to the Pareto front:

```
mo_pareto_optx_plot(X=X_base, Y=y, minimize=False)
```




Figure 32.13.: Pareto-optimal points in the input space

32. Multi-Objective Optimization Support in SpotOptim


Compute desirability functions for each objective:

```
d_funcs = []
for i in range(y.shape[1]):
    d_func = DMax(low=np.min(y[:, i]), high=np.max(y[:, i]))
    d_funcs.append(d_func)
```


The desirability functions can be visualized as follows:

```
d_funcs[0].plot(xlabel=target_names[0], figsize=(6,4))
d_funcs[1].plot(xlabel=target_names[1], figsize=(6,4))
```

32.8. A Simple Visualization Example



(a) Desirability functions of the `mo_conv2_max` function



(b)

Figure 32.14.

32. Multi-Objective Optimization Support in SpotOptim

The minimal expected improvement of the Morris-Mitchell criterion is defined as 1 percent of the base value, and the maximal expected improvement as 10 percent of the base value. Then, the desirability function for Morris-Mitchell objective can be defined as follows:

```
# imp_min = phi_base * 0.01
# imp_max = phi_base * 0.1

# d_mm = DMax(low=imp_min, high=imp_max)
# d_mm.plot()
```

Combine into overall desirability

```
# D_overall = DOverall(*d_funcs, d_mm)
D_overall = DOverall(*d_funcs)
```

Now we can call the `mo_mm_desirability_function` with a test point:

```
x_test = np.array([0.5, 0.5]) # Example test point
neg_D, objectives = mo_mm_desirability_function(x_test, models, X_base, J_base, d_base)
print(f"Negative Desirability: {neg_D}")
print(f"Objectives: {objectives}")
```

```
Negative Desirability: -0.7374994411771443
Objectives: [np.float64(1.4906734308429819), np.float64(1.5123762382740051)]
```

```
mo_xy_desirability_plot(models, bounds=[(x_min, x_max), (x_min, x_max)], X_base=X_base)
```




Figure 32.15.: Desirability function for multi-objective optimization

32.9. Jupyter Notebook

i Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

Part IX.

Hyperparameter Tuning

33. The TorchObjective Class for Hyperparameter Tuning

SpotOptim provides a `TorchObjective` class that simplifies the process of hyperparameter tuning for PyTorch models. It acts as a bridge between your PyTorch code (models, data, training logic) and the `SpotOptim` optimizer.

This tutorial guides you through the entire workflow, from defining a model to running a comprehensive hyperparameter optimization experiment.

33.1. Overview

The `TorchObjective` class automates several tedious tasks:

1. **Translating Hyperparameters:** Converts the vector of values provided by `SpotOptim` into a dictionary of named parameters with correct types (int, float, categorical).
2. **Model Instantiation:** Creates a new instance of your model for each evaluation with the specific hyperparameters.
3. **Data Loading:** Handles creating PyTorch `DataLoaders`, including dynamic batch sizes.
4. **Training & Evaluation:** Runs the training loop and returns validation metrics (e.g., validation loss, MSE).

33.2. Workflow Steps

The typical workflow involves 5 steps:

1. **Define Model:** Create a standard PyTorch `nn.Module`.
2. **Prepare Data:** Wrap your data in a `SpotOptim` data container.
3. **Define Hyperparameters:** Specify what to tune using `ParameterSet`.
4. **Create Experiment:** Bundle everything into an `ExperimentControl` object.
5. **Optimize:** Use `SpotOptim` to find the best configuration.

33.3. Step-by-Step Example

Let's walk through a complete example tuning a Multi-Layer Perceptron (MLP) on the California Housing dataset.

33.3.1. 1. Define the Model

You can use any standard PyTorch model. For this example, we'll use the built-in MLP class from `spotoptim.nn.mlp`, but you could define your own `nn.Module`.

```
import torch
import torch.nn as nn
from spotoptim.nn.mlp import MLP
```

33.3.2. 2. Prepare the Data

Load your data (numpy arrays or tensors) and wrap it in `SpotDataFromArray`.

```
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from spotoptim.core.data import SpotDataFromArray

# Load data
data = fetch_california_housing()
X, y = data.data, data.target.reshape(-1, 1)

# Split and scale
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)
scaler_X = StandardScaler()
scaler_y = StandardScaler()

X_train = scaler_X.fit_transform(X_train)
X_val = scaler_X.transform(X_val)
y_train = scaler_y.fit_transform(y_train)
y_val = scaler_y.transform(y_val)

# Create SpotData container
spot_data = SpotDataFromArray(
    x_train=X_train, y_train=y_train,
    x_val=X_val, y_val=y_val
)
```

33.3.3. 3. Define Hyperparameters

Use `ParameterSet` to define the search space. You can tune floats, integers, and categorical factors.

```
from spotoptim.hyperparameters import ParameterSet

params = ParameterSet()

# Tuning Learning Rate (log scale)
params.add_float("lr", 1e-4, 1e-1, default=1e-3, transform="log10")

# Tuning Model Architecture
params.add_int("num_hidden_layers", 1, 3, default=2)
params.add_int("l1", 16, 128, default=32)
params.add_float("dropout", 0.0, 0.5, default=0.0)

# Tuning Optimization Parameters
params.add_int("epochs", 10, 50, default=20)
params.add_int("batch_size", 32, 256, default=64) # Batch size can now be tuned!
```

```
ParameterSet(
    lr=Parameter(
        name='lr',
        var_name='lr',
        bounds=Bounds(low=0.0001, high=0.1),
        default=0.001,
        transform='log10',
        type='float'
    ),
    num_hidden_layers=Parameter(
        name='num_hidden_layers',
        var_name='num_hidden_layers',
        bounds=Bounds(low=1, high=3),
        default=2,
        transform=None,
        type='int'
    ),
    l1=Parameter(
        name='l1',
        var_name='l1',
        bounds=Bounds(low=16, high=128),
        default=32,
        transform=None,
```

33. The TorchObjective Class for Hyperparameter Tuning

```
        type='int'
    ),
    dropout=Parameter(
        name='dropout',
        var_name='dropout',
        bounds=Bounds(low=0.0, high=0.5),
        default=0.0,
        transform=None,
        type='float'
    ),
    epochs=Parameter(
        name='epochs',
        var_name='epochs',
        bounds=Bounds(low=10, high=50),
        default=20,
        transform=None,
        type='int'
    ),
    batch_size=Parameter(
        name='batch_size',
        var_name='batch_size',
        bounds=Bounds(low=32, high=256),
        default=64,
        transform=None,
        type='int'
    ),
)
```

33.3.4. 4. Create Experiment Control

Bundles the configuration together.

```
import torch
from spotoptim.core.experiment import ExperimentControl

experiment = ExperimentControl(
    experiment_name="california_housing_tuning",
    model_class=MLP,
    dataset=spot_data,
    hyperparameters=params,
    metrics=["val_loss"],
    device="cpu", # or "cuda" if available
    loss_function=nn.MSELoss(),
    num_workers=0,
```

```
    batch_size=64 # Default fallback if not tuned
)
```

33.3.5. 5. Initialize TorchObjective

This wraps the experiment into a callable function that `SpotOptim` can use.

```
from spotoptim.function.torch_objective import TorchObjective

objective = TorchObjective(experiment)
```

33.3.6. 6. Run Optimization

Now use `SpotOptim` to find the best hyperparameters.

```
from spotoptim import SpotOptim
import numpy as np

optimizer = SpotOptim(
    fun=objective,
    bounds=objective.bounds,
    var_type=objective.var_type,
    var_name=objective.var_name,
    var_trans=objective.var_trans, # Use defined transformations (e.g., log10 for lr)
    n_initial=5, # Number of initial random points
    max_iter=10, # Total evaluations (initial + sequential)
    seed=42, # For reproducibility
    verbose=True
)

result = optimizer.optimize()

print("\nBest Configuration Found:")
best_params = objective._get_hyperparameters(result.x)
for k, v in best_params.items():
    print(f"{k}: {v}")

print(f"\nBest Validation Loss: {result.fun:.5f}")
```

```
TensorBoard logging disabled
Initial best: f(x) = 0.451601
```

33. The TorchObjective Class for Hyperparameter Tuning

```
Iter 1 | Best: 0.451601 | Curr: 0.566953 | Rate: 0.00 | Evals: 60.0%
Optimizer candidate 1/3 was duplicate/invalid.
Iter 2 | Best: 0.446584 | Rate: 0.50 | Evals: 70.0%
Iter 3 | Best: 0.446584 | Curr: 1.018047 | Rate: 0.33 | Evals: 80.0%
Iter 4 | Best: 0.446584 | Curr: 0.504958 | Rate: 0.25 | Evals: 90.0%
Iter 5 | Best: 0.446584 | Curr: 0.504960 | Rate: 0.20 | Evals: 100.0%
```

```
Best Configuration Found:
lr: 0.028975719181608448
num_hidden_layers: 2
l1: 35
dropout: 0.07075743906046855
epochs: 20
batch_size: 128
```

Best Validation Loss: 0.44658

33.4. Advanced Features

33.4.1. Tuning Batch Size

TorchObjective supports dynamic batch size tuning. By adding `batch_size` to your `ParameterSet` (as shown above), the objective function will automatically recreate `DataLoaders` with the specific batch size for each evaluation.

This allows you to optimize the trade-off between training speed (larger batches) and convergence quality (often better with smaller batches).

33.4.2. Tuning Architecture

You can tune structural parameters like:

- **num_hidden_layers**: Depth of the network.
- **l1**: Width of the layers (first hidden layer size, propagated if others are not specified).
- **dropout**: Regularization strength.
- **Optimization method**: You can even tune the optimizer itself (e.g., “Adam” vs “SGD”) if you add it as a factor hyperparameter and your model class supports a `get_optimizer` method (like MLP does).

Example adding optimizer tuning:

```
# Add optimizer choice
params.add_factor("optimizer", ["Adam", "SGD", "RMSprop"], default="Adam")

ParameterSet(
    lr=Parameter(
        name='lr',
        var_name='lr',
        bounds=Bounds(low=0.0001, high=0.1),
        default=0.001,
        transform='log10',
        type='float'
    ),
    num_hidden_layers=Parameter(
        name='num_hidden_layers',
        var_name='num_hidden_layers',
        bounds=Bounds(low=1, high=3),
        default=2,
        transform=None,
        type='int'
    ),
    l1=Parameter(
        name='l1',
        var_name='l1',
        bounds=Bounds(low=16, high=128),
        default=32,
        transform=None,
        type='int'
    ),
    dropout=Parameter(
        name='dropout',
        var_name='dropout',
        bounds=Bounds(low=0.0, high=0.5),
        default=0.0,
        transform=None,
        type='float'
    ),
    epochs=Parameter(
        name='epochs',
        var_name='epochs',
        bounds=Bounds(low=10, high=50),
        default=20,
        transform=None,
        type='int'
    ),
)
```

33. The TorchObjective Class for Hyperparameter Tuning

```
batch_size=Parameter(
    name='batch_size',
    var_name='batch_size',
    bounds=Bounds(low=32, high=256),
    default=64,
    transform=None,
    type='int'
),
optimizer=Parameter(
    name='optimizer',
    var_name='optimizer',
    bounds=['Adam', 'SGD', 'RMSprop'],
    default='Adam',
    transform=None,
    type='factor'
),
)
```

33.4.3. Log-Scale Transformations

For parameters that span several orders of magnitude (like Learning Rate), it is highly recommended to use a log transformation. `ParameterSet` handles this easily:

```
# Search lr in [0.0001, 0.1] but optimizer sees [-4, -1]
params.add_float("lr", 1e-4, 1e-1, transform="log10")
```

`SpotOptim` will propose values in the transformed space (e.g., -3.0), and `TorchObjective` will automatically untransform them ($10^{-3} = 0.001$) before passing them to the model.

33.5. Custom Models

To use your own model, simply define a class inheriting from `nn.Module` whose `__init__` accepts the hyperparameters you defined.

```
class MyCustomModel(nn.Module):
    def __init__(self, input_dim, output_dim, hidden_size, activation):
        super().__init__()
        # ... verify args match your ParameterSet names ...
        self.layer1 = nn.Linear(input_dim, hidden_size)
        self.act = getattr(nn, activation)()
```



```
self.layer2 = nn.Linear(hidden_size, output_dim)

def forward(self, x):
    return self.layer2(self.act(self.layer1(x)))
```

Ensure your `ParameterSet` includes `hidden_size` and `activation` (as a factor). `input_dim` and `output_dim` are automatically passed from the dataset info.

34. The MLP Class

SpotOptim provides a MLP class that implements a flexible Multi-Layer Perceptron (MLP) architecture using PyTorch. It is designed to be easily plug-and-play for both standalone usage and hyperparameter optimization with SpotOptim.

34.1. Overview

The MLP class extends `torch.nn.Sequential` and offers:

1. **Flexible Architecture:** Define layers explicitly via `hidden_channels` or using compact hyperparameters like `l1` (width) and `num_hidden_layers` (depth).
2. **Integrated Components:** Built-in support for normalization layers, activation functions, and dropout.
3. **Optimization Helpers:** Includes a `get_optimizer` method to easily instantiate optimizers with unified learning rates.
4. **Tuning Ready:** A `get_default_parameters` static method returns a `ParameterSet` ready for SpotOptim.

34.2. Basic Usage

34.2.1. Initialization

You can initialize an MLP by describing its architecture explicitly.

```
import torch
from spotoptim.nn.mlp import MLP

# Input: 10 features
# Hidden: 32 neurons, then 16 neurons
# Output: 1 feature (regression)
model = MLP(in_channels=10, hidden_channels=[32, 16, 1])
print(model)

# Forward pass
```

34. The MLP Class

```
x = torch.randn(5, 10)
output = model(x)
print("Output shape:", output.shape)
```

```
MLP(
  (0): Linear(in_features=10, out_features=32, bias=True)
  (1): ReLU()
  (2): Dropout(p=0.0, inplace=False)
  (3): Linear(in_features=32, out_features=16, bias=True)
  (4): ReLU()
  (5): Dropout(p=0.0, inplace=False)
  (6): Linear(in_features=16, out_features=1, bias=True)
  (7): Dropout(p=0.0, inplace=False)
)
Output shape: torch.Size([5, 1])
```

34.2.2. Implicit Architecture (Width & Depth)

For hyperparameter tuning, it is often easier to control the network's size with just two numbers: width and depth.

```
# Create a network with 3 hidden layers, each having 64 neurons
model_compact = MLP(
    in_channels=10,
    l1=64,          # Width (neurons per hidden layer)
    num_hidden_layers=3, # Depth (number of hidden layers)
    output_dim=1
)
print(model_compact)
```

```
MLP(
  (0): Linear(in_features=10, out_features=64, bias=True)
  (1): ReLU()
  (2): Dropout(p=0.0, inplace=False)
  (3): Linear(in_features=64, out_features=64, bias=True)
  (4): ReLU()
  (5): Dropout(p=0.0, inplace=False)
  (6): Linear(in_features=64, out_features=64, bias=True)
  (7): ReLU()
  (8): Dropout(p=0.0, inplace=False)
  (9): Linear(in_features=64, out_features=1, bias=True)
  (10): Dropout(p=0.0, inplace=False)
)
```

i Note

If `hidden_channels` is provided, `l1` and `num_hidden_layers` are ignored.

34.3. Configuration Options

The MLP constructor supports several customization options:

- **activation_layer**: The activation function class (default: `torch.nn.ReLU`).
- **norm_layer**: Optional normalization layer (e.g., `torch.nn.BatchNorm1d`).
- **dropout**: Dropout probability applied after each layer (default: 0.0).
- **bias**: Whether to use bias in linear layers (default: `True`).

```
model_custom = MLP(
    in_channels=10,
    hidden_channels=[32, 1],
    activation_layer=torch.nn.Tanh,
    dropout=0.2
)
print(model_custom)
```

```
MLP(
  (0): Linear(in_features=10, out_features=32, bias=True)
  (1): Tanh()
  (2): Dropout(p=0.2, inplace=False)
  (3): Linear(in_features=32, out_features=1, bias=True)
  (4): Dropout(p=0.2, inplace=False)
)
```

34.4. Optimizer Integration

The MLP class simplifies optimizer creation, specifically handling the “unified learning rate” concept used in SpotOptim (where different optimizers have their default learning rates mapped to a common scale).

```
# Create model with a unified learning rate of 1.0 (default)
model = MLP(in_channels=10, hidden_channels=[32, 1], lr=1.0)

# Get Adam optimizer (lr=1.0 maps to 0.001)
opt_adam = model.get_optimizer("Adam")
```

34. The MLP Class

```
print(f"Adam lr: {opt_adam.param_groups[0]['lr']}")

# Get SGD optimizer (lr=1.0 maps to 0.01)
opt_sgd = model.get_optimizer("SGD")
print(f"SGD lr: {opt_sgd.param_groups[0]['lr']}")
```

```
Adam lr: 0.001
SGD lr: 0.01
```

You can also pass extra arguments to the optimizer:

```
# SGD with momentum
opt_sgd_mom = model.get_optimizer("SGD", momentum=0.9)
print(opt_sgd_mom)
```

```
SGD (
Parameter Group 0
  dampening: 0
  differentiable: False
  foreach: None
  fused: None
  lr: 0.01
  maximize: False
  momentum: 0.9
  nesterov: False
  weight_decay: 0
)
```

34.5. Hyperparameter Tuning

One of the key features of the MLP class is its ability to suggest a default `ParameterSet` for tuning. This provides a great starting point for finding the best architecture.

```
from spotoptim.hyperparameters import ParameterSet

# Get default search space
params = MLP.get_default_parameters()
print("Default tunable parameters:", params.names())
```

```
Default tunable parameters: ['l1', 'num_hidden_layers', 'activation', 'lr', 'optimizer']
```

34.5.1. Example: Tuning with SpotOptim

Here is how you can use the MLP class in a full SpotOptim tuning loop using TorchObjective.

```
from spotoptim import SpotOptim
from spotoptim.core.experiment import ExperimentControl
from spotoptim.core.data import SpotDataFromArray
from spotoptim.function.torch_objective import TorchObjective
import numpy as np

# 1. Dummy Data
X = np.random.rand(100, 10)
y = np.random.rand(100, 1)
data = SpotDataFromArray(X, y)

# 2. Get Default Parameters & Add Custom Ones
params = MLP.get_default_parameters()

# Customize: fix the optimizer to Adam, but tune epochs
params.add_int("epochs", 5, 20, default=10)

# 3. Setup Experiment
experiment = ExperimentControl(
    experiment_name="mlp_tuning_demo",
    model_class=MLP,
    dataset=data,
    hyperparameters=params,
    metrics=["val_loss"],
    device="cpu",
    batch_size=16
)

# 4. Create Objective
objective = TorchObjective(experiment)

# 5. Optimize
optimizer = SpotOptim(
    fun=objective,
    bounds=objective.bounds,
    var_type=objective.var_type,
    var_name=objective.var_name,
    var_trans=objective.var_trans,
    n_initial=3,
    max_iter=5,
```

34. The MLP Class

```
seed=42,  
verbose=False  
)  
  
res = optimizer.optimize()  
  
print("Best Parameters:")  
print(objective._get_hyperparameters(res.x))
```

```
/var/folders/dw/pvtj6mt91znd0hftcztqb0k00000gn/T/ipykernel_40142/2724594128.py:33: Use  
agnostic) factor metrics. Factor integer codes will be treated as ordinal numbers, wh  
aware Kriging surrogate: Kriging(metric_factorial='hamming').  
optimizer = SpotOptim(
```

Best Parameters:

```
{'l1': 68, 'num_hidden_layers': 2, 'activation': 'LeakyReLU', 'lr': 0.240637430047912
```

This setup automatically tunes the architecture (`l1`, `num_hidden_layers`), usage of activation functions (`activation`), learning rate (`lr`), and optimization method (`optimizer`) if left in the parameter set.

35. Physics-Informed Neural Networks (PINNs) Demo 1

Solving ODEs with SpotOptim's LinearRegressor

36. Overview

This tutorial demonstrates how to use Physics-Informed Neural Networks (PINNs) to solve ordinary differential equations (ODEs) using SpotOptim's **LinearRegressor** class. We'll solve a first-order ODE with an initial condition, combining data-driven learning with physics constraints.

36.1. The Differential Equation

We want to solve the following ODE:

$$\frac{dy}{dt} + 0.1y - \sin\left(\frac{\pi t}{2}\right) = 0$$

with initial condition:

$$y(0) = 0$$

37. Setup

First, let's import the necessary libraries:

```
import matplotlib.pyplot as plt
import numpy as np
import torch
import torch.nn as nn
from typing import Tuple
from spoptim.nn.linear_regressor import LinearRegressor

# Set random seed for reproducibility
torch.manual_seed(123)
```

```
<torch._C.Generator at 0x117daf730>
```


38. The Neural Network

We'll use SpotOptim's `LinearRegressor` class to create a neural network with:

- 1 input (time t)
- 1 output (solution y)
- 32 neurons per hidden layer
- 3 hidden layers
- Tanh activation function

```
model = LinearRegressor(  
    input_dim=1,  
    output_dim=1,  
    l1=32,  
    num_hidden_layers=3,  
    activation="Tanh",  
    lr=1.0  
)  
  
# Get optimizer with custom learning rate  
optimizer = model.get_optimizer("Adam", lr=3.0) # 3.0 * 0.001 = 0.003  
  
print(f"Model architecture:")  
print(model)
```

Model architecture:

```
LinearRegressor(  
  (activation): Tanh()  
  (network): Sequential(  
    (0): Linear(in_features=1, out_features=32, bias=True)  
    (1): Tanh()  
    (2): Linear(in_features=32, out_features=32, bias=True)  
    (3): Tanh()  
    (4): Linear(in_features=32, out_features=32, bias=True)  
    (5): Tanh()  
    (6): Linear(in_features=32, out_features=1, bias=True)  
  )  
)
```


39. Generate Training Data

We'll generate exact solution data using a numerical solver (RK2 method):

```
def oscillator(
    n_steps: int = 3000,
    t_min: float = 0.0,
    t_max: float = 30.0,
    y0: float = 0.0,
    alpha: float = 0.1,
    omega: float = np.pi / 2
) -> Tuple[torch.Tensor, torch.Tensor]:
    """
    Solve ODE:  $dy/dt + \alpha y - \sin(\omega t) = 0$ 
    using RK2 (midpoint method).

    Args:
        n_steps: Number of time steps
        t_min: Start time
        t_max: End time
        y0: Initial condition  $y(t_{\min})$ 
        alpha: Damping coefficient
        omega: Forcing frequency

    Returns:
        Tuple of (time points, solution values) as torch tensors
    """
    t_step = (t_max - t_min) / n_steps
    t_points = np.arange(t_min, t_min + n_steps * t_step, t_step)[:n_steps]

    y = [y0]

    for t_current_step_end in t_points[1:]:
        t_midpoint = t_current_step_end - t_step / 2.0
        y_prev = y[-1]

        # Stage 1: intermediate value
        slope_at_t_mid = -alpha * y_prev + np.sin(omega * t_midpoint)
```

39. Generate Training Data

```
y_intermediate = y_prev + (t_step / 2.0) * slope_at_t_mid

# Stage 2: final value
slope_at_t_end = -alpha * y_intermediate + np.sin(omega * t_current_step_end)
y_next = y_prev + t_step * slope_at_t_end
y.append(y_next)

t_tensor = torch.tensor(t_points, dtype=torch.float32).view(-1, 1)
y_tensor = torch.tensor(y, dtype=torch.float32).view(-1, 1)

return t_tensor, y_tensor
```

Generate the exact solution and sample training data points:

```
# Generate exact solution (3000 points)
x, y = oscillator()

# Sample training data (every 119th point, giving ~25 points)
x_data = x[0:3000:119]
y_data = y[0:3000:119]

print(f"Exact solution points: {len(x)}")
print(f"Training data points: {len(x_data)}")
```

Exact solution points: 3000
Training data points: 26

Visualize the exact solution and training data:

```
plt.figure(figsize=(10, 6))
plt.plot(x.numpy(), y.numpy(), 'b-', linewidth=2, label="Exact solution y(t)")
plt.scatter(x_data.numpy(), y_data.numpy(), color="tab:orange",
            s=50, label="Training data", zorder=5)
plt.xlabel("Time t", fontsize=12)
plt.ylabel("Solution y(t)", fontsize=12)
plt.title("ODE Solution and Training Data", fontsize=14)
plt.legend(fontsize=11)
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```



40. Collocation Points

Create collocation points where we'll enforce the physics (ODE) constraints:

```
# 50 evenly spaced points in [0, 30]
x_physics = torch.linspace(0, 30, 50).view(-1, 1).requires_grad_(True)

print(f"Collocation points: {len(x_physics)}")
print(f"Range: [{x_physics.min().item():.2f}, {x_physics.max().item():.2f}]" )
```

```
Collocation points: 50
Range: [0.00, 30.00]
```


41. PINN Training

Now we'll train the neural network using two loss components:

1. **Data Loss (Loss1)**: Ensures the network fits the training data
2. **Physics Loss (Loss2)**: Ensures the network satisfies the ODE

41.1. Loss Components Explained

41.1.1. Neural Network Prediction on Data Points

The network predicts values at the training data points:

```
yh = model(x_data)
```

41.1.2. Data Loss (Loss1)

Mean squared error between predictions and actual data:

```
loss1 = torch.mean((yh - y_data)**2)
```

41.1.3. Neural Network Prediction on Collocation Points

The network predicts values at the physics enforcement points:

```
yhp = model(x_physics)
```

41.1.4. Computing Derivatives for the ODE

We compute dy/dt using automatic differentiation:

41. PINN Training

```
dyhp_dxphysics = torch.autograd.grad(
    yhp, x_physics,
    torch.ones_like(yhp),
    create_graph=True
)[0]
```

41.1.5. Physics Loss (Loss2)

The ODE residual at collocation points:

$$\text{residual} = \frac{dy}{dt} + 0.1y - \sin\left(\frac{\pi t}{2}\right)$$

```
physics = dyhp_dxphysics + 0.1 * yhp - torch.sin(np.pi * x_physics / 2)
loss2 = torch.mean(physics**2)
```

41.1.6. Total Loss

Combined loss with weighting factor α :

```
loss = loss1 + alpha * loss2
```

41.2. Training Loop

```
loss_history_pinn = []
loss2_history_pinn = []
plot_data_points_pinn = []

alpha = 6e-2 # Weight for physics loss
n_epochs = 48000

print("Training Physics-Informed Neural Network...")
print(f"Total epochs: {n_epochs}")
print(f"Physics loss weight (alpha): {alpha}")

for i in range(n_epochs):
    optimizer.zero_grad()

    # Data Loss: Fit the training data
```



```

yh = model(x_data)
loss1 = torch.mean((yh - y_data)**2)

# Physics Loss: Satisfy the ODE at collocation points
yhp = model(x_physics)
dyhp_dxphysics = torch.autograd.grad(
    yhp, x_physics,
    torch.ones_like(yhp),
    create_graph=True
)[0]
physics = dyhp_dxphysics + 0.1 * yhp - torch.sin(np.pi * x_physics / 2)
loss2 = torch.mean(physics**2)

# Total Loss
loss = loss1 + alpha * loss2
loss.backward()
optimizer.step()

# Store history every 100 steps
if (i + 1) % 100 == 0:
    loss_history_pinn.append(loss.detach())
    loss2_history_pinn.append(loss2.detach())

# Store snapshots for visualization every 10000 steps
if (i + 1) % 10000 == 0:
    current_yh_full = model(x).detach()
    plot_data_points_pinn.append({
        'yh': current_yh_full,
        'step': i + 1
    })
    print(f"Epoch {i+1}/{n_epochs}: Total Loss = {loss.item():.6f}, "
          f"Physics Loss = {loss2.item():.6f}")

print("Training completed!")

```

```

Training Physics-Informed Neural Network...
Total epochs: 48000
Physics loss weight (alpha): 0.06
Epoch 10000/48000: Total Loss = 0.000321, Physics Loss = 0.002949
Epoch 20000/48000: Total Loss = 0.000020, Physics Loss = 0.000309
Epoch 30000/48000: Total Loss = 0.000095, Physics Loss = 0.000285
Epoch 40000/48000: Total Loss = 0.000823, Physics Loss = 0.001919
Training completed!

```


42. Results Visualization


42.1. Training Progress

```
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 5))

# Plot total loss
ax1.plot(loss_history_pinn, 'b-', linewidth=1.5)
ax1.set_xlabel('Iteration (×100)', fontsize=11)
ax1.set_ylabel('Total Loss', fontsize=11)
ax1.set_title('Training Progress: Total Loss', fontsize=12)
ax1.grid(True, alpha=0.3)
ax1.set_yscale('log')

# Plot physics loss
ax2.plot(loss2_history_pinn, 'r-', linewidth=1.5)
ax2.set_xlabel('Iteration (×100)', fontsize=11)
ax2.set_ylabel('Physics Loss', fontsize=11)
ax2.set_title('Training Progress: Physics Loss (ODE Residual)', fontsize=12)
ax2.grid(True, alpha=0.3)
ax2.set_yscale('log')

plt.tight_layout()
plt.show()
```



42.2. Solution Evolution During Training

Visualize how the neural network solution evolves during training:

```

xp_plot = x_physics.detach()

for plot_info in plot_data_points_pinn:
    plt.figure(figsize=(10, 6))

    # Plot exact solution
    plt.plot(x.numpy(), y.numpy(), 'b-', linewidth=2,
             label='Exact solution', alpha=0.7)

    # Plot training data
    plt.scatter(x_data.numpy(), y_data.numpy(), color='tab:orange',
               s=50, label='Training data', zorder=5)



    # Plot neural network prediction
    plt.plot(x.numpy(), plot_info['yh'].numpy(), 'r--', linewidth=2,
             label='PINN prediction', alpha=0.8)

    # Plot collocation points
    plt.scatter(xp_plot.numpy(),
               model(xp_plot).detach().numpy(),
               color='green', marker='x', s=30,
               label='Collocation points', alpha=0.6, zorder=4)



    plt.xlabel('Time t', fontsize=12)
    plt.ylabel('Solution y(t)', fontsize=12)
    plt.title(f'PINN Solution at Epoch {plot_info["step"]}', fontsize=14)
    plt.legend(fontsize=11)
    plt.grid(True, alpha=0.3)
    plt.tight_layout()
    plt.show()

```

42.2. Solution Evolution During Training



42. Results Visualization



42.3. Final Solution Comparison

```

# Final prediction
model.eval()
with torch.no_grad():
    y_final = model(x)

plt.figure(figsize=(12, 7))

# Plot exact solution
plt.plot(x.numpy(), y.numpy(), 'b-', linewidth=2.5,
         label='Exact solution', alpha=0.8)


# Plot PINN prediction
plt.plot(x.numpy(), y_final.numpy(), 'r--', linewidth=2,
         label='PINN prediction', alpha=0.8)

# Plot training data
plt.scatter(x_data.numpy(), y_data.numpy(), color='tab:orange',
           s=80, label='Training data', zorder=5, edgecolors='black', linewidth=0.5)

# Plot collocation points
plt.scatter(x_physics.detach().numpy(),
           model(x_physics).detach().numpy(),
           color='green', marker='x', s=50,
           label='Collocation points', alpha=0.7, zorder=4)

plt.xlabel('Time t', fontsize=13)
plt.ylabel('Solution y(t)', fontsize=13)
plt.title('Final PINN Solution vs Exact Solution', fontsize=15, fontweight='bold')
plt.legend(fontsize=12, loc='best')
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

```




42.4. Error Analysis

```
# Compute absolute error
error = torch.abs(y_final - y)

plt.figure(figsize=(10, 6))
plt.plot(x.numpy(), error.numpy(), 'r-', linewidth=2)
plt.xlabel('Time t', fontsize=12)
plt.ylabel('Absolute Error |y_exact - y_PINN|', fontsize=12)
plt.title('PINN Approximation Error', fontsize=14)
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

print("\nError Statistics:")
print(f"Maximum absolute error: {error.max().item():.6f}")
print(f"Mean absolute error: {error.mean().item():.6f}")
print(f"Root mean squared error: {torch.sqrt(torch.mean(error**2)).item():.6f}")
```


42.4. Error Analysis



Error Statistics:

Maximum absolute error: 0.125975

Mean absolute error: 0.010835

Root mean squared error: 0.021054

43. Summary

This tutorial demonstrated how to use SpotOptim's `LinearRegressor` class to implement a Physics-Informed Neural Network (PINN) for solving ordinary differential equations. Key takeaways:

1. **Network Architecture:** Used a 3-layer neural network with 32 neurons per layer and Tanh activation
2. **Dual Loss Function:** Combined data fitting loss with physics constraint loss
3. **Automatic Differentiation:** Leveraged PyTorch's autograd to compute derivatives for the ODE
4. **Collocation Method:** Enforced the ODE at specific points in the domain
5. **Training Strategy:** Balanced data-driven and physics-informed learning with weight λ

The PINN successfully learned to approximate the ODE solution using only sparse training data (~25 points) by incorporating the underlying physics through the differential equation constraint.

43.1. Key Advantages of PINNs

- **Data Efficiency:** Can learn with very few data points
- **Physics Consistency:** Solutions automatically satisfy the governing equations
- **Generalization:** Better extrapolation beyond training data
- **Flexibility:** Can handle complex geometries and boundary conditions

43.2. Using SpotOptim for Hyperparameter Optimization

The `LinearRegressor` class integrates seamlessly with SpotOptim for hyperparameter tuning:

```
from spotoptim import SpotOptim

def train_pinn(X):
    """Objective function for hyperparameter optimization."""
    results = []
```

43. Summary

```
for params in X:
    l1 = int(params[0])          # Hidden layer size
    num_layers = int(params[1])  # Number of layers
    lr_unified = 10 ** params[2] # Learning rate (log scale)

    # Create model
    model = LinearRegressor(
        input_dim=1, output_dim=1,
        l1=l1, num_hidden_layers=num_layers,
        activation="Tanh", lr=lr_unified
    )

    # Train PINN and compute validation error
    # ... training code ...

    results.append(validation_error)
return np.array(results)

# Optimize hyperparameters
optimizer = SpotOptim(
    fun=train_pinn,
    bounds=[(16, 128), (1, 5), (-4, 0)],
    var_type=["int", "int", "float"],
    var_name=["layer_size", "num_layers", "log_lr"],
    max_iter=50
)
result = optimizer.optimize()
```

This approach allows you to systematically find the best network architecture and learning rate for your specific PINN problem.

43.3. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

44. Hyperparameter Tuning for Physics-Informed Neural Networks

Using SpotOptim to Optimize PINN Architecture and Training

45. Overview

This tutorial demonstrates how to use SpotOptim for hyperparameter optimization of Physics-Informed Neural Networks (PINNs). We'll optimize the network architecture and training parameters to find the best configuration for solving an ordinary differential equation.

Building on the basic PINN demo, we'll now systematically search for optimal:

- Number of neurons per hidden layer
- Number of hidden layers
- Activation function (categorical)
- Optimizer algorithm (categorical)
- Learning rate (log-scale)
- Physics loss weight (log-scale)

45.1. Key Features

45.1.1. 1. PyTorch Dataset and DataLoader

Following PyTorch best practices from the official tutorial, this tutorial implements:

- **Custom Dataset Classes:** Separate classes for supervised data (`PINNDataset`) and collocation points (`CollocationDataset`)
- **DataLoader Integration:** Efficient batch processing with configurable batch size, shuffling, and parallel loading
- **Proper Data Separation:** Clean separation of training, validation, and collocation data
- **Gradient Tracking:** Automatic gradient handling for collocation points needed in physics loss

Benefits:

- **Modularity:** Clean separation between data and model code
- **Efficiency:** Batch processing and optional parallel data loading
- **Scalability:** Easy to extend to larger datasets
- **Best Practices:** Follows PyTorch conventions used across the ecosystem

45.1.2. 2. Automatic Transformation Handling

This tutorial also showcases SpotOptim's `var_trans` feature for automatic variable transformations. Learning rates and regularization parameters are often best explored on a log scale, but manually transforming values is tedious and error-prone. With `var_trans`, you simply specify:

```
var_trans = [None, None, "log10", "log10"]
```

SpotOptim then:

- Optimizes internally in log-transformed space (efficient exploration)
- Passes original-scale values to your objective function (no manual conversion needed)
- Displays all results in original scale (easy interpretation)

This eliminates the need for manual `10**x` conversions throughout your code!

46. The Problem

We're solving the same ODE as in the basic PINN demo:

$$\frac{dy}{dt} + 0.1y - \sin\left(\frac{\pi t}{2}\right) = 0$$

with initial condition $y(0) = 0$.

47. Setup

```
import matplotlib.pyplot as plt
import numpy as np
import torch
import torch.nn as nn
from typing import Tuple
from spotoptim import SpotOptim
from spotoptim.nn.linear_regressor import LinearRegressor
```

```
# Set random seed for reproducibility
torch.manual_seed(42)
np.random.seed(42)
```

Set number of epochs for training, maximum number of iterations for hyperparameter optimization, and initial design size. For this example we use a small number of epochs and iterations to keep the runtime short.

```
N_EPOCHS = 500
MAX_ITER = 15
N_INITIAL = 10
```


48. Data Generation

Following PyTorch best practices, we'll create custom Dataset classes for our PINN data.

48.1. Custom Dataset Classes

We'll create two dataset types:

1. PINNDataset for supervised data (training and validation)
2. CollocationDataset for physics-informed collocation points

```
from torch.utils.data import Dataset, DataLoader

def oscillator(
    n_steps: int = 3000,
    t_min: float = 0.0,
    t_max: float = 30.0,
    y0: float = 0.0,
    alpha: float = 0.1,
    omega: float = np.pi / 2
) -> Tuple[torch.Tensor, torch.Tensor]:
    """
    Solve ODE:  $dy/dt + \alpha y - \sin(\omega t) = 0$ 
    using RK2 (midpoint method).

    Returns:
        t_tensor: Time points, shape (n_steps, 1)
        y_tensor: Solution values, shape (n_steps, 1)
    """
    t_step = (t_max - t_min) / n_steps
    t_points = np.arange(t_min, t_min + n_steps * t_step, t_step)[:n_steps]

    y = [y0]

    for t_current_step_end in t_points[1:]:
        t_midpoint = t_current_step_end - t_step / 2.0
```

```

        y_prev = y[-1]

        slope_at_t_mid = -alpha * y_prev + np.sin(omega * t_midpoint)
        y_intermediate = y_prev + (t_step / 2.0) * slope_at_t_mid

        slope_at_t_end = -alpha * y_intermediate + np.sin(omega * t_current_step_end)
        y_next = y_prev + t_step * slope_at_t_end
        y.append(y_next)

    t_tensor = torch.tensor(t_points, dtype=torch.float32).view(-1, 1)
    y_tensor = torch.tensor(y, dtype=torch.float32).view(-1, 1)

    return t_tensor, y_tensor

```

```

class PINNDataset(Dataset):
    """PyTorch Dataset for PINN supervised data (training/validation).

    This dataset stores time-solution pairs (t, y) for supervised learning.

    Args:
        t (torch.Tensor): Time points, shape (n_samples, 1)
        y (torch.Tensor): Solution values, shape (n_samples, 1)
    """

    def __init__(self, t: torch.Tensor, y: torch.Tensor):
        self.t = t
        self.y = y

    def __len__(self) -> int:
        return len(self.t)

    def __getitem__(self, idx: int) -> Tuple[torch.Tensor, torch.Tensor]:
        return self.t[idx], self.y[idx]

```

```

class CollocationDataset(Dataset):
    """PyTorch Dataset for PINN collocation points.

    This dataset stores time points where physics loss is evaluated.
    Gradients are required for computing derivatives in the PDE.

    Args:
        t (torch.Tensor): Collocation time points, shape (n_points, 1)

```

```

"""

def __init__(self, t: torch.Tensor):
    # Store collocation points with gradient tracking
    self.t = t.requires_grad_(True)

def __len__(self) -> int:
    return len(self.t)

def __getitem__(self, idx: int) -> torch.Tensor:
    # Return single collocation point (still requires_grad)
    return self.t[idx].unsqueeze(0)

```

Generate exact solution using Runge-Kutta 2 (RK2) method:

```

x_exact, y_exact = oscillator()

# Create training data (sparse sampling)
t_train = x_exact[0:3000:119]
y_train = y_exact[0:3000:119]

# Create validation data (different sampling for unbiased evaluation)
t_val = x_exact[50:3000:120]
y_val = y_exact[50:3000:120]

# Create collocation points for physics loss
t_physics = torch.linspace(0, 30, 50).view(-1, 1)

# Create Dataset objects
train_dataset = PINNDataset(t_train, y_train)
val_dataset = PINNDataset(t_val, y_val)
collocation_dataset = CollocationDataset(t_physics)

print(f"Training dataset size: {len(train_dataset)}")
print(f"Validation dataset size: {len(val_dataset)}")
print(f"Collocation dataset size: {len(collocation_dataset)}")
print(f"\nSample from training dataset:")
t_sample, y_sample = train_dataset[0]
print(f"  t: {t_sample.item():.4f}, y: {y_sample.item():.4f}")

```

```

Training dataset size: 26
Validation dataset size: 25
Collocation dataset size: 50

```

48. Data Generation

```
Sample from training dataset:  
  t: 0.0000, y: 0.0000
```


49. Define the PINN Training Function

This function creates DataLoaders and trains a PINN with given hyperparameters. Following PyTorch best practices, we use DataLoader for efficient batch processing.

```
def train_pinn(
    l1: int,
    num_layers: int,
    activation: str,
    optimizer_name: str,
    lr_unified: float,
    alpha: float,
    n_epochs: int = N_EPOCHS,
    batch_size: int = 16,
    verbose: bool = False
) -> float:
    """
    Train a PINN with specified hyperparameters using DataLoaders.

    Args:
        l1: Number of neurons per hidden layer
        num_layers: Number of hidden layers
        activation: Activation function ("Tanh", "ReLU", "Sigmoid", "GELU")
        optimizer_name: Optimizer algorithm ("Adam", "SGD", "RMSprop", "AdamW")
        lr_unified: Unified learning rate multiplier
        alpha: Weight for physics loss
        n_epochs: Number of training epochs
        batch_size: Batch size for DataLoader
        verbose: Whether to print progress

    Returns:
        Validation mean squared error
    """
    # Set seed for reproducibility
    torch.manual_seed(42)

    # Create DataLoaders
    train_loader = DataLoader(
```

49. Define the PINN Training Function

```
        train_dataset,
        batch_size=batch_size,
        shuffle=True,
        num_workers=0,
        pin_memory=False
    )

    val_loader = DataLoader(
        val_dataset,
        batch_size=batch_size,
        shuffle=False,
        num_workers=0,
        pin_memory=False
    )

    # For collocation points, we can use full batch since it's small
    collocation_loader = DataLoader(
        collocation_dataset,
        batch_size=len(collocation_dataset),
        shuffle=False,
        num_workers=0
    )

    # Create model
    model = LinearRegressor(
        input_dim=1,
        output_dim=1,
        l1=l1,
        num_hidden_layers=num_layers,
        activation=activation,
        lr=lr_unified
    )

    # Get optimizer
    optimizer = model.get_optimizer(optimizer_name)

    # Training loop
    for epoch in range(n_epochs):
        model.train()
        epoch_loss = 0.0

        # Get collocation points (full batch)
        t_physics_batch = next(iter(collocation_loader))
        # Ensure gradients are enabled
```

```

t_physics_batch = t_physics_batch.requires_grad_(True)

# Iterate over training batches
for batch_t, batch_y in train_loader:
    optimizer.zero_grad()

    # Data Loss
    y_pred = model(batch_t)
    loss_data = torch.mean((y_pred - batch_y)**2)

    # Physics Loss (computed on full collocation set)
    y_physics = model(t_physics_batch)
    dy_dt = torch.autograd.grad(
        y_physics,
        t_physics_batch,
        torch.ones_like(y_physics),
        create_graph=True,
        retain_graph=True
    )[0]

    # PDE residual: dy/dt + 0.1*y - sin(pi*t/2) = 0
    physics_residual = dy_dt + 0.1 * y_physics - torch.sin(np.pi * t_physics_batch / 2)
    loss_physics = torch.mean(physics_residual**2)

    # Total Loss
    loss = loss_data + alpha * loss_physics
    loss.backward()
    optimizer.step()

    epoch_loss += loss.item()

if verbose and (epoch + 1) % 2000 == 0:
    avg_loss = epoch_loss / len(train_loader)
    print(f" Epoch {epoch+1}/{n_epochs}: Avg Loss = {avg_loss:.6f}")

# Evaluate on validation set
model.eval()
val_mse = 0.0
with torch.no_grad():
    for batch_t, batch_y in val_loader:
        y_pred = model(batch_t)
        val_mse += torch.mean((batch_y - y_pred)**2).item()

val_mse /= len(val_loader)

```

49. Define the PINN Training Function

```
        return val_mse

# Test the function with default parameters
print("Testing PINN training function with DataLoaders...")
test_error = train_pinn(
    l1=32,
    num_layers=2,
    activation="Tanh",
    optimizer_name="Adam",
    lr_unified=3.0,
    alpha=0.06,
    batch_size=16,
    n_epochs=N_EPOCHS,
    verbose=True
)
print(f"\nTest validation MSE: {test_error:.6f}")
```

Testing PINN training function with DataLoaders...

Test validation MSE: 0.188752

50. Hyperparameter Optimization with SpotOptim

Now we'll use SpotOptim to find the best hyperparameters:

50.1. Define the Objective Function

```
def objective_pinn(X):
    """
    Objective function for SpotOptim.

    Args:
        X: Array of hyperparameter configurations, shape (n_configs, 6)
        Each row: [l1, num_layers, activation, optimizer, lr_unified, alpha]
        Note: SpotOptim handles log transformations and factor mapping automatically

    Returns:
        Array of validation errors
    """
    results = []

    for i, params in enumerate(X):
        # Extract parameters (already in original scale thanks to var_trans)
        # Factor variables (activation, optimizer) are returned as strings
        l1 = int(params[0])                # Number of neurons
        num_layers = int(params[1])        # Number of hidden layers
        activation = params[2]              # Activation function
        optimizer_name = params[3]         # Optimizer algorithm
        lr_unified = params[4]              # Learning rate
        alpha = params[5]                   # Physics weight

        print(f"\nConfiguration {i+1}/{len(X)}:")
        print(f"  l1={l1}, num_layers={num_layers}, activation={activation}, ")
        print(f"  optimizer={optimizer_name}, lr_unified={lr_unified:.4f}, alpha={alpha:.4f}")
```

50. Hyperparameter Optimization with SpotOptim

```
# Train PINN with these hyperparameters
val_error = train_pinn(
    l1=l1,
    num_layers=num_layers,
    activation=activation,
    optimizer_name=optimizer_name,
    lr_unified=lr_unified,
    alpha=alpha,
    n_epochs=N_EPOCHS,
    verbose=False
)

print(f" Validation MSE: {val_error:.6f}")
results.append(val_error)

return np.array(results)
```

We test the objective function with two configurations:

```
X_test = np.array([
    [32, 2, "Tanh", "Adam", 3.0, 0.06], # Baseline config
    [64, 3, "ReLU", "AdamW", 2.0, 0.04] # Alternative config
], dtype=object)
test_results = objective_pinn(X_test)
print(f"\nTest results: {test_results}")
```

Configuration 1/2:

```
l1=32, num_layers=2, activation=Tanh,
optimizer=Adam, lr_unified=3.0000, alpha=0.0600
Validation MSE: 0.188752
```

Configuration 2/2:

```
l1=64, num_layers=3, activation=ReLU,
optimizer=AdamW, lr_unified=2.0000, alpha=0.0400
Validation MSE: 0.168650
```

Test results: [0.18875157 0.16864995]

50.2. Run the Optimization

Use `tensorboard --logdir=runs` from a shell in the current directory (where this notebook is located) to visualize the optimization process.

Setting the bounds for the search space:

```
# Define search space with var_trans for automatic log-scale handling
bounds = [
    (16, 128),           # l1: neurons per layer (16 to 128)
    (1, 4),             # num_layers: 1 to 4 hidden layers
    ("Tanh", "ReLU", "Sigmoid", "GELU"), # activation: activation function
    ("Adam", "SGD", "RMSprop", "AdamW"), # optimizer: optimizer algorithm
    (0.1, 10.0),        # lr_unified: learning rate (0.1 to 10)
    (0.01, 1.0)         # alpha: physics weight (0.01 to 1.0)
]
```

Specify the variable types and transformations. Use `var_trans` to handle log-scale transformations automatically, factor variables don't need transformations (None):

```
var_type = ["int", "int", "factor", "factor", "float", "float"]
var_name = ["l1", "num_layers", "activation", "optimizer", "lr_unified", "alpha"]
var_trans = [None, None, None, None, "log10", "log10"]
```

```
# Create optimizer
optimizer = SpotOptim(
    fun=objective_pinn,
    bounds=bounds,
    var_type=var_type,
    var_name=var_name,
    var_trans=var_trans, # Automatic log-scale handling!
    max_iter=MAX_ITER,
    n_initial=N_INITIAL,
    seed=42,
    verbose=True,
    tensorboard_clean=True,
    tensorboard_log=True
)
```

Factor variable at dimension 2:

Levels: ['Tanh', 'ReLU', 'Sigmoid', 'GELU']

Mapped to integers: 0 to 3

Factor variable at dimension 3:

50. Hyperparameter Optimization with SpotOptim

```
Levels: ['Adam', 'SGD', 'RMSprop', 'AdamW']
Mapped to integers: 0 to 3
Removed old TensorBoard logs: runs/spotoptim_20260701_211045
Cleaned 1 old TensorBoard log directory
TensorBoard logging enabled: runs/spotoptim_20260701_211149
```

```
/var/folders/dw/pvtj6mt91znd0hftcztqb0k00000gn/T/ipykernel_40499/2553769193.py:2: UserWarning: Using a
agnostic) factor metrics. Factor integer codes will be treated as ordinal numbers, which may be
aware Kriging surrogate: Kriging(metric_factorial='hamming').
optimizer = SpotOptim(
```

Display search space configuration. The transcolumn shows applied transformations. `lr_unified` and `alpha` use `log10` transformation internally. This enables efficient exploration of log-scale parameters. All values shown are in original scale (not transformed).

```
design_table = optimizer.get_design_table(tablefmt="github")
print(design_table)
```

name	type	lower	upper	default	transform
l1	int	16.0000	128.0000	72	-
num_layers	int	1.0000	4.0000	2	-
activation	factor	-	-	Sigmoid	-
optimizer	factor	-	-	RMSprop	-
lr_unified	float	0.1000	10.0000	5.0500	log10
alpha	float	0.0100	1.0000	0.5050	log10

Run optimization

```
result = optimizer.optimize()
```

Configuration 1/10:

```
l1=28, num_layers=2, activation=Sigmoid,
optimizer=RMSprop, lr_unified=1.0211, alpha=0.0923
Validation MSE: 0.192541
```

Configuration 2/10:

```
l1=117, num_layers=3, activation=ReLU,
optimizer=SGD, lr_unified=0.4832, alpha=0.0146
Validation MSE: 0.205611
```


50.2. Run the Optimization

Configuration 3/10:

l1=92, num_layers=1, activation=Tanh,
optimizer=SGD, lr_unified=6.5194, alpha=0.6267
Validation MSE: nan

Configuration 4/10:

l1=41, num_layers=4, activation=ReLU,
optimizer=Adam, lr_unified=0.2742, alpha=0.1697
Validation MSE: 0.185978

Configuration 5/10:

l1=105, num_layers=2, activation=ReLU,
optimizer=AdamW, lr_unified=0.1136, alpha=0.0176
Validation MSE: 0.202467

Configuration 6/10:

l1=24, num_layers=1, activation=Sigmoid,
optimizer=RMSprop, lr_unified=4.4408, alpha=0.0293
Validation MSE: 0.189705

Configuration 7/10:

l1=83, num_layers=3, activation=GELU,
optimizer=AdamW, lr_unified=0.2277, alpha=0.2652
Validation MSE: 0.181881

Configuration 8/10:

l1=65, num_layers=2, activation=Tanh,
optimizer=SGD, lr_unified=3.4389, alpha=0.9830
Validation MSE: 0.195487

Configuration 9/10:

l1=116, num_layers=3, activation=Sigmoid,
optimizer=RMSprop, lr_unified=0.8521, alpha=0.1020
Validation MSE: 0.180443

Configuration 10/10:

l1=51, num_layers=4, activation=Sigmoid,
optimizer=Adam, lr_unified=1.6729, alpha=0.0458
Validation MSE: 0.205772

Warning: 1 initial design point(s) returned NaN/inf and will be ignored (reduced from 10 to 9 points)

Note: Initial design size (9) is smaller than requested (10) due to NaN/inf values

Initial best: $f(x) = 0.180443$

Configuration 1/1:

50. Hyperparameter Optimization with SpotOptim

```
l1=100, num_layers=2, activation=Sigmoid,  
optimizer=RMSprop, lr_unified=0.1469, alpha=0.0131  
Validation MSE: 0.196912  
Iter 1 | Best: 0.180443 | Curr: 0.196912 | Rate: 0.00 | Evals: 66.7%
```

```
Configuration 1/1:  
l1=94, num_layers=1, activation=Sigmoid,  
optimizer=SGD, lr_unified=3.6048, alpha=0.0374  
Validation MSE: 0.196974  
Iter 2 | Best: 0.180443 | Curr: 0.196974 | Rate: 0.00 | Evals: 73.3%
```

```
Configuration 1/1:  
l1=96, num_layers=3, activation=GELU,  
optimizer=RMSprop, lr_unified=0.2996, alpha=0.1593  
Validation MSE: 0.169332  
Iter 3 | Best: 0.169332 | Rate: 0.33 | Evals: 80.0%
```

```
Configuration 1/1:  
l1=127, num_layers=1, activation=Sigmoid,  
optimizer=SGD, lr_unified=1.9791, alpha=0.2477  
Validation MSE: 0.208434  
Iter 4 | Best: 0.169332 | Curr: 0.208434 | Rate: 0.25 | Evals: 86.7%
```

```
Configuration 1/1:  
l1=78, num_layers=3, activation=Sigmoid,  
optimizer=SGD, lr_unified=0.5950, alpha=0.4743  
Validation MSE: 0.228979  
Iter 5 | Best: 0.169332 | Curr: 0.228979 | Rate: 0.20 | Evals: 93.3%
```

```
Configuration 1/1:  
l1=81, num_layers=2, activation=Sigmoid,  
optimizer=Adam, lr_unified=0.2781, alpha=0.0168  
Validation MSE: 0.202393  
Iter 6 | Best: 0.169332 | Curr: 0.202393 | Rate: 0.17 | Evals: 100.0%
```

TensorBoard writer closed. View logs with: `tensorboard --logdir=runs/spotoptim_202607`

51. Results Analysis

51.1. Best Configuration

Display best hyperparameters using `print_best()` method. With `var_trans`, results are already in original scale!

```
optimizer.print_best(result)
```

Best Solution Found:

```
-----  
l1: 96  
num_layers: 3  
activation: GELU  
optimizer: RMSprop  
lr_unified: 0.2996  
alpha: 0.1593  
Objective Value: 0.1693  
Total Evaluations: 15
```

Store values for later use in visualizations. Values are already in original scale thanks to `var_trans`. Factor variables are returned as strings.

```
best_l1 = int(result.x[0])  
best_num_layers = int(result.x[1])  
best_activation = result.x[2]  
best_optimizer = result.x[3]  
best_lr_unified = result.x[4]  
best_alpha = result.x[5]  
best_val_error = result.fun  
  
print(f"Best activation: {best_activation}")  
print(f"Best optimizer: {best_optimizer}")
```

```
Best activation: GELU  
Best optimizer: RMSprop
```

51. Results Analysis

51.1.1. Results Table with Importance Scores

Display comprehensive results table with importance scores


```
table = optimizer.get_results_table(show_importance=True, tablefmt="github")
print(table)
```

name	type	default	lower	upper	tuned	transform	imp
11	int	72	16.0000	128.0000	96	-	
num_layers	int	2	1.0000	4.0000	3	-	
activation	factor	Sigmoid	-	-	GELU	-	
optimizer	factor	RMSprop	-	-	RMSprop	-	
lr_unified	float	5.0500	0.1000	10.0000	0.2996	log10	
alpha	float	0.5050	0.0100	1.0000	0.1593	log10	

Interpretation: ***: >99%, **: >75%, *: >50%, .: >10%

51.2. Optimization History

```
optimizer.plot_progress(log_y=True, ylabel="Validation MSE")
```



51.3. Surrogate Visualization

Visualize the surrogate model's learned response surface for the most important hyperparameter combinations:


```
# Plot top 3 most important hyperparameter combinations
optimizer.plot_important_hyperparameter_contour(max_imp=3)
```

Plotting surrogate contours for top 3 most important parameters:

optimizer: importance = 35.84% (type: factor)
 activation: importance = 28.00% (type: factor)
 alpha: importance = 12.62% (type: float)


Generating 3 surrogate plots...

Plotting optimizer vs activation




51. Results Analysis

Plotting optimizer vs alpha



Plotting activation vs alpha


51.4. Parameter Distribution Analysis



51.4. Parameter Distribution Analysis

```
optimizer.plot_parameter_scatter(
    result,
    ylabel="Validation MSE",
    cmap="plasma",
    figsize=(14, 12)
)
```

51. Results Analysis



52. Train Final Model with Best Hyperparameters

Now let's train a final model with the optimized hyperparameters using DataLoaders:

```
print("Training final model with best hyperparameters using DataLoaders...")
print(f"Training for 30,000 epochs...")

# Set seed for reproducibility
torch.manual_seed(42)

# Create DataLoaders for final training
final_batch_size = 16
train_loader_final = DataLoader(
    train_dataset,
    batch_size=final_batch_size,
    shuffle=True,
    num_workers=0
)

collocation_loader_final = DataLoader(
    collocation_dataset,
    batch_size=len(collocation_dataset),
    shuffle=False,
    num_workers=0
)

# Create model with best hyperparameters
final_model = LinearRegressor(
    input_dim=1,
    output_dim=1,
    l1=best_l1,
    num_hidden_layers=best_num_layers,
    activation=best_activation,
    lr=best_lr_unified
)
```

52. Train Final Model with Best Hyperparameters

```
optimizer_final = final_model.get_optimizer(best_optimizer)

# Training with history tracking
loss_history = []
n_epochs_final = 30000

for epoch in range(n_epochs_final):
    final_model.train()
    epoch_loss = 0.0

    # Get collocation points
    t_physics_batch = next(iter(collocation_loader_final))
    t_physics_batch = t_physics_batch.requires_grad_(True)

    # Iterate over training batches
    for batch_t, batch_y in train_loader_final:
        optimizer_final.zero_grad()

        # Data Loss
        y_pred = final_model(batch_t)
        loss_data = torch.mean((y_pred - batch_y)**2)

        # Physics Loss
        y_physics = final_model(t_physics_batch)
        dy_dt = torch.autograd.grad(
            y_physics,
            t_physics_batch,
            torch.ones_like(y_physics),
            create_graph=True,
            retain_graph=True
        )[0]

        physics_residual = dy_dt + 0.1 * y_physics - torch.sin(np.pi * t_physics_batch)
        loss_physics = torch.mean(physics_residual**2)

        # Total Loss
        loss = loss_data + best_alpha * loss_physics
        loss.backward()
        optimizer_final.step()

        epoch_loss += loss.item()

    # Record average loss every 100 epochs
    if (epoch + 1) % 100 == 0:
```

```

        avg_loss = epoch_loss / len(train_loader_final)
        loss_history.append(avg_loss)

    if (epoch + 1) % 5000 == 0:
        avg_loss = epoch_loss / len(train_loader_final)
        print(f"   Epoch {epoch+1}/{n_epochs_final}: Avg Loss = {avg_loss:.6f}")

print("Training completed!")

# Plot training history
plt.figure(figsize=(10, 5))
plt.plot(loss_history, linewidth=1.5)
plt.xlabel('Iteration (×100)', fontsize=11)
plt.ylabel('Average Total Loss per Batch', fontsize=11)
plt.title('Final Model Training History', fontsize=12)
plt.grid(True, alpha=0.3)
plt.yscale('log')
plt.tight_layout()
plt.show()

```

Training final model with best hyperparameters using DataLoaders...

Training for 30,000 epochs...


```

Epoch 5000/30000: Avg Loss = 0.182355
Epoch 10000/30000: Avg Loss = 0.143730
Epoch 15000/30000: Avg Loss = 0.133864
Epoch 20000/30000: Avg Loss = 0.126419
Epoch 25000/30000: Avg Loss = 0.084283
Epoch 30000/30000: Avg Loss = 0.002897

```

Training completed!

52. Train Final Model with Best Hyperparameters



53. Evaluate Final Model

```
# Create validation DataLoader for evaluation
val_loader_final = DataLoader(
    val_dataset,
    batch_size=len(val_dataset),
    shuffle=False,
    num_workers=0
)

# Evaluate on validation set
final_model.eval()
with torch.no_grad():
    # Validation MSE using DataLoader
    val_mse_total = 0.0
    for batch_t, batch_y in val_loader_final:
        y_pred = final_model(batch_t)
        val_mse_total += torch.mean((y_pred - batch_y)**2).item()
    final_val_mse = val_mse_total / len(val_loader_final)

    # Predict on full domain for visualization
    y_pred_full = final_model(x_exact)
    full_mse = torch.mean((y_pred_full - y_exact)**2).item()

    # Compute maximum absolute error
    max_error = torch.max(torch.abs(y_pred_full - y_exact)).item()

print("\nFinal Model Performance:")
print("-" * 50)
print(f"  Validation MSE: {final_val_mse:.6f}")
print(f"  Full domain MSE: {full_mse:.6f}")
print(f"  Maximum absolute error: {max_error:.6f}")
```

Final Model Performance:

Validation MSE: 0.002740

53. *Evaluate Final Model*

Full domain MSE: 0.002078
Maximum absolute error: 0.119569

54. Visualize Final Solution

```
# Generate predictions
final_model.eval()
with torch.no_grad():
    y_pred = final_model(x_exact)

fig, axes = plt.subplots(2, 1, figsize=(12, 10))

# Plot 1: Solution comparison
ax1 = axes[0]
ax1.plot(x_exact.numpy(), y_exact.numpy(), 'b-', linewidth=2.5,
        label='Exact solution', alpha=0.8)
ax1.plot(x_exact.numpy(), y_pred.numpy(), 'r--', linewidth=2,
        label='PINN prediction (optimized)', alpha=0.8)

# Plot training data from dataset
ax1.scatter(train_dataset.t.numpy(), train_dataset.y.numpy(),
        color='tab:orange', s=80, label='Training data',
        zorder=5, edgecolors='black', linewidth=0.5)

# Plot collocation points
t_collocation = collocation_dataset.t.detach()
ax1.scatter(t_collocation.numpy(),
        final_model(t_collocation).detach().numpy(),
        color='green', marker='x', s=50,
        label='Collocation points', alpha=0.7, zorder=4)


ax1.set_xlabel('Time t', fontsize=12)
ax1.set_ylabel('Solution y(t)', fontsize=12)
ax1.set_title('Optimized PINN Solution vs Exact Solution', fontsize=13, fontweight='bold')
ax1.legend(fontsize=11, loc='best')
ax1.grid(True, alpha=0.3)

# Plot 2: Error
ax2 = axes[1]
error = torch.abs(y_pred - y_exact)
```

54. Visualize Final Solution

```
ax2.plot(x_exact.numpy(), error.numpy(), 'r-', linewidth=2)
ax2.axhline(y=max_error, color='gray', linestyle='--', linewidth=1,
            label=f'Max error = {max_error:.6f}')
ax2.set_xlabel('Time t', fontsize=12)
ax2.set_ylabel('Absolute Error |y_exact - y_PINN|', fontsize=12)
ax2.set_title('Approximation Error (Optimized Model)', fontsize=13, fontweight='bold')
ax2.legend(fontsize=11)
ax2.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```



55. Comparison with Baseline

Let's compare the optimized configuration with a baseline:

```
print("Training baseline model for comparison...")

# Baseline configuration (from basic PINN demo)
baseline_config = {
    'l1': 32,
    'num_layers': 3,
    'activation': 'Tanh',
    'optimizer': 'Adam',
    'lr_unified': 3.0,
    'alpha': 0.06
}

print(f"\nBaseline Configuration:")
for key, val in baseline_config.items():
    print(f"  {key}: {val}")

# Train baseline
baseline_error = train_pinn(
    l1=baseline_config['l1'],
    num_layers=baseline_config['num_layers'],
    activation=baseline_config['activation'],
    optimizer_name=baseline_config['optimizer'],
    lr_unified=baseline_config['lr_unified'],
    alpha=baseline_config['alpha'],
    n_epochs=N_EPOCHS,
    verbose=False
)

print(f"\nValidation MSE Comparison:")
print("-" * 50)
print(f"  Baseline: {baseline_error:.6f}")
print(f"  Optimized: {best_val_error:.6f}")
print(f"  Improvement: {(1 - best_val_error/baseline_error)*100:.1f}%")
```

55. Comparison with Baseline

```
# Bar plot comparison
fig, ax = plt.subplots(figsize=(8, 6))
configs = ['Baseline', 'Optimized']
errors = [baseline_error, best_val_error]
colors = ['tab:blue', 'tab:green']

bars = ax.bar(configs, errors, color=colors, alpha=0.7, edgecolor='black', linewidth=1)

# Add value labels on bars
for bar, error in zip(bars, errors):
    height = bar.get_height()
    ax.text(bar.get_x() + bar.get_width()/2., height,
            f'{error:.6f}',
            ha='center', va='bottom', fontsize=11, fontweight='bold')

ax.set_ylabel('Validation MSE', fontsize=12)
ax.set_title('Model Comparison: Baseline vs Optimized', fontsize=13, fontweight='bold')
ax.grid(True, alpha=0.3, axis='y')
plt.tight_layout()
plt.show()
```


Training baseline model for comparison...

Baseline Configuration:

```
l1: 32
num_layers: 3
activation: Tanh
optimizer: Adam
lr_unified: 3.0
alpha: 0.06
```

Validation MSE Comparison:


```
-----
Baseline: 0.185787
Optimized: 0.169332
Improvement: 8.9%
```



56. Hyperparameter Sensitivity Analysis

Let's analyze how sensitive the model is to each hyperparameter using the enhanced `plot_parameter_scatter()` method with Spearman correlation:

```
# Use the enhanced plot_parameter_scatter method with correlation display
optimizer.plot_parameter_scatter(
    result,
    ylabel="Validation MSE",
    cmap="plasma",
    figsize=(15, 10),
    show_correlation=True,
    log_y=True
)
```



56.1. Sensitivity Analysis (Spearman Correlation)

```
# Use the new sensitivity_spearman() method for tabular output
optimizer.sensitivity_spearman()
```

Sensitivity Analysis (Spearman Correlation):

```
-----
l1                : +0.196 (p=0.483)
num_layers        : -0.158 (p=0.573)
activation         : (categorical variable, use visual inspection)
optimizer         : (categorical variable, use visual inspection)
lr_unified         : +0.118 (p=0.676)
alpha             : -0.150 (p=0.594)
```

57. Summary

57.1. Key Findings

```
# Get optimization history for statistics
history = optimizer.y_

print("\n" + "="*70)
print("HYPERPARAMETER OPTIMIZATION SUMMARY")
print("="*70)

print("\n1. BEST CONFIGURATION FOUND:")
print(f"  - Neurons per layer (l1): {best_l1}")
print(f"  - Number of hidden layers: {best_num_layers}")
print(f"  - Activation function: {best_activation}")
print(f"  - Optimizer: {best_optimizer}")
print(f"  - Learning rate: {best_lr_unified:.4f}")
print(f"  - Physics weight (alpha): {best_alpha:.4f}")

print("\n2. PERFORMANCE:")
print(f"  - Validation MSE: {best_val_error:.6f}")
print(f"  - Full domain MSE: {full_mse:.6f}")
print(f"  - Maximum absolute error: {max_error:.6f}")

print("\n3. OPTIMIZATION STATISTICS:")
print(f"  - Total evaluations: {result.nfev}")
print(f"  - Initial best: {history[0]:.6f}")
print(f"  - Final best: {best_val_error:.6f}")
print(f"  - Improvement: {(1 - best_val_error/history[0])*100:.1f}%")

print("\n4. COMPARISON TO BASELINE:")
print(f"  - Baseline MSE: {baseline_error:.6f}")
print(f"  - Optimized MSE: {best_val_error:.6f}")
print(f"  - Improvement: {(1 - best_val_error/baseline_error)*100:.1f}%")

print("\n" + "="*70)
```

57. Summary

=====

HYPERPARAMETER OPTIMIZATION SUMMARY

=====

1. BEST CONFIGURATION FOUND:
 - Neurons per layer (l1): 96
 - Number of hidden layers: 3
 - Activation function: GELU
 - Optimizer: RMSprop
 - Learning rate: 0.2996
 - Physics weight (alpha): 0.1593
2. PERFORMANCE:
 - Validation MSE: 0.169332
 - Full domain MSE: 0.002078
 - Maximum absolute error: 0.119569
3. OPTIMIZATION STATISTICS:
 - Total evaluations: 15
 - Initial best: 0.192541
 - Final best: 0.169332
 - Improvement: 12.1%
4. COMPARISON TO BASELINE:
 - Baseline MSE: 0.185787
 - Optimized MSE: 0.169332
 - Improvement: 8.9%

=====

57.2. Recommendations

Based on the hyperparameter optimization results:

1. Network Architecture:

- The optimal architecture was found with {best_l1} neurons and {best_num_layers} hidden layers
- Best activation function: {best_activation}
- This balances model capacity with training efficiency

2. Optimizer Selection:

- Best optimizer: {best_optimizer}
- Different optimizers have different convergence characteristics for PINNs

3. Learning Rate:

- Optimal unified learning rate: {best_lr_unified:.4f}
- This translates to an actual Adam learning rate of {best_lr_unified * 0.001:.6f}

4. Physics Loss Weight:

- Optimal alpha: {best_alpha:.4f}
- This balances data fitting with physics constraint satisfaction

5. Training Strategy:

- Start with a broad search space to explore different architectures
- Use `var_trans` with “log10” for learning rate and physics weight parameters
- This enables efficient exploration of log-scale parameters without manual transformations
- Validate on held-out data to prevent overfitting to training points

6. Benefits of `var_trans` and Factor Variables:

- **Factor variables:** Categorical choices (activation, optimizer) handled automatically
- `SpotOptim` maps strings to integers internally and back to strings in results
- **Cleaner code:** No manual `10**x` conversions in objective function
- **Fewer errors:** Eliminates confusion about which scale values are in
- **Better optimization:** Searches efficiently in transformed space
- **Easier interpretation:** All results displayed in original scale

57.3. Using These Results

To use the optimized configuration in your own PINN problems:

```
# Create optimized PINN
model = LinearRegressor(
    input_dim=1,
    output_dim=1,
    l1={best_l1},
    num_hidden_layers={best_num_layers},
    activation="{best_activation}",
    lr={best_lr_unified:.4f}
)
```

```
optimizer = model.get_optimizer("{best_optimizer}")

# Use alpha={best_alpha:.4f} for physics loss weight
loss = data_loss + {best_alpha:.4f} * physics_loss
```

57.4. Using var_trans for Your Hyperparameter Optimization

When setting up optimization for your own PINN problems:

```
from spotoptim import SpotOptim

# Define search space with factor variables and log-scale parameters
bounds = [
    (16, 128),                # neurons (integer)
    (1, 4),                  # layers (integer)
    ("Tanh", "ReLU", "Sigmoid", "GELU"), # activation (factor)
    ("Adam", "SGD", "RMSprop", "AdamW"),  # optimizer (factor)
    (0.1, 10.0),             # learning rate (log-scale)
    (0.01, 1.0)              # physics weight (log-scale)
]

var_type = ["int", "int", "factor", "factor", "float", "float"]
var_trans = [None, None, None, None, "log10", "log10"]

opt = SpotOptim(
    fun=your_objective_function,
    bounds=bounds,
    var_type=var_type,
    var_trans=var_trans, # Automatic log-scale and factor handling!
    max_iter=MAX_ITER,
    n_initial=N_INITIAL
)

result = opt.optimize()
```

Your objective function receives parameters in **original scale** - no manual transformations needed!

57.5. Future Directions

Consider exploring:

1. **Adaptive physics weights** that change during training
2. **Architecture search** including skip connections or residual blocks
3. **Batch size optimization** as an additional hyperparameter
4. **Multi-objective optimization** balancing accuracy and computational cost
5. **Transfer learning** from pre-optimized configurations
6. **Learning rate schedules** with different decay strategies

Note: The specific optimal values depend on the problem, data distribution, and computational budget. Always validate results on held-out test data.

57.6. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

58. Learning Rate Mapping for Unified Optimizer Interface

SpotOptim provides a sophisticated learning rate mapping system through the `map_lr()` function, enabling a unified interface for learning rates across different PyTorch optimizers. This solves the challenge that different optimizers operate on vastly different learning rate scales.

58.1. Overview

Different PyTorch optimizers use different default learning rates and optimal ranges:

- **Adam**: default 0.001, typical range 0.0001-0.01
- **SGD**: default 0.01, typical range 0.001-0.1
- **RMSprop**: default 0.01, typical range 0.001-0.1

This makes it difficult to compare optimizer performance fairly or optimize learning rates across different optimizers. The `map_lr()` function provides a unified scale where **`lr_unified=1.0` corresponds to each optimizer's PyTorch default.**

Module: `spotoptim.utils.mapping`

Key Features:

- Unified learning rate scale across all optimizers
- Fair comparison when evaluating different optimizers
- Simplified hyperparameter optimization
- Based on official PyTorch default learning rates
- Supports 13 major PyTorch optimizers

58.2. Quick Start

58.2.1. Basic Usage

58. Learning Rate Mapping for Unified Optimizer Interface

```
from spotoptim.utils.mapping import map_lr

# Get optimizer-specific learning rate from unified scale
lr_adam = map_lr(1.0, "Adam")      # Returns 0.001 (Adam's default)
lr_sgd = map_lr(1.0, "SGD")        # Returns 0.01 (SGD's default)
lr_rmsprop = map_lr(1.0, "RMSprop") # Returns 0.01 (RMSprop's default)

print(f"Unified lr=1.0:")
print(f"  Adam:    {lr_adam}")
print(f"  SGD:      {lr_sgd}")
print(f"  RMSprop: {lr_rmsprop}")
```

```
Unified lr=1.0:
  Adam:    0.001
  SGD:      0.01
  RMSprop: 0.01
```

58.2.2. Scaling Learning Rates

```
from spotoptim.utils.mapping import map_lr

# Scale all learning rates by the same factor
unified_lr = 0.5

lr_adam = map_lr(unified_lr, "Adam")      # 0.5 * 0.001 = 0.0005
lr_sgd = map_lr(unified_lr, "SGD")        # 0.5 * 0.01 = 0.005
lr_rmsprop = map_lr(unified_lr, "RMSprop") # 0.5 * 0.01 = 0.005

print(f"Unified lr={unified_lr}:")
print(f"  Adam:    {lr_adam}")
print(f"  SGD:      {lr_sgd}")
print(f"  RMSprop: {lr_rmsprop}")
```

```
Unified lr=0.5:
  Adam:    0.0005
  SGD:      0.005
  RMSprop: 0.005
```

58.2.3. Integration with LinearRegressor

```
from spotoptim.nn.linear_regressor import LinearRegressor

# Create model with unified learning rate
model = LinearRegressor(
    input_dim=10,
    output_dim=1,
    l1=32,
    num_hidden_layers=2,
    lr=1.0 # Unified learning rate
)

# Get optimizer - automatically uses mapped learning rate
optimizer_adam = model.get_optimizer("Adam")      # Gets 1.0 * 0.001 = 0.001
optimizer_sgd = model.get_optimizer("SGD")        # Gets 1.0 * 0.01 = 0.01

# Verify the actual learning rates
print(f"Adam actual lr: {optimizer_adam.param_groups[0]['lr']}")
print(f"SGD actual lr: {optimizer_sgd.param_groups[0]['lr']}")
```

Adam actual lr: 0.001

SGD actual lr: 0.01

58.3. Function Reference

58.3.1. `map_lr(lr_unified, optimizer_name, use_default_scale=True)`

Maps a unified learning rate to an optimizer-specific learning rate.

Parameters:

- `lr_unified` (float): Unified learning rate multiplier. A value of 1.0 corresponds to the optimizer's default learning rate. Typical range: [0.001, 100.0].
- `optimizer_name` (str): Name of the PyTorch optimizer. Must be one of: "Adadelata", "Adagrad", "Adam", "AdamW", "SparseAdam", "Adamax", "ASGD", "LBFGS", "NAdam", "RAdam", "RMSprop", "Rprop", "SGD".
- `use_default_scale` (bool, optional): Whether to scale by the optimizer's default learning rate. If `True` (default), `lr_unified` is multiplied by the default lr. If `False`, returns `lr_unified` directly.

Returns:

- **float:** The optimizer-specific learning rate.

Raises:

- **ValueError:** If `optimizer_name` is not supported.
- **ValueError:** If `lr_unified` is not positive.

Example:

```
from spotoptim.utils.mapping import map_lr

# Get default learning rates (unified lr = 1.0)
lr = map_lr(1.0, "Adam")      # 0.001
lr = map_lr(1.0, "SGD")      # 0.01
lr = map_lr(1.0, "RMSprop")  # 0.01

# Scale learning rates
lr = map_lr(0.5, "Adam")     # 0.0005
lr = map_lr(2.0, "SGD")     # 0.02

# Without default scaling
lr = map_lr(0.01, "Adam", use_default_scale=False) # 0.01 (direct)
```

58.4. Supported Optimizers

All major PyTorch optimizers are supported with their default learning rates:

Optimizer	Default LR	Typical Range	Notes
Adam	0.001	0.0001-0.01	Most popular, good default
AdamW	0.001	0.0001-0.01	Adam with weight decay
Adamax	0.002	0.0001-0.01	Adam variant with infinity norm
NAdam	0.002	0.0001-0.01	Adam with Nesterov momentum
RAdam	0.001	0.0001-0.01	Rectified Adam
SparseAdam	0.001	0.0001-0.01	For sparse gradients
SGD	0.01	0.001-0.1	Classic, needs momentum
RMSprop	0.01	0.001-0.1	Good for RNNs
Adagrad	0.01	0.001-0.1	Adaptive learning rate
Adadelat	1.0	0.1-10.0	Extension of Adagrad
ASGD	0.01	0.001-0.1	Averaged SGD
LBFGS	1.0	0.1-10.0	Second-order optimizer
Rprop	0.01	0.001-0.1	Resilient backpropagation

58.5. Use Cases

58.5.1. Comparing Different Optimizers

```
import torch
import torch.nn as nn
from spotoptim.nn.linear_regressor import LinearRegressor
from spotoptim.data import get_diabetes_dataloaders

# Load data
train_loader, test_loader, _ = get_diabetes_dataloaders(batch_size=32, random_state=42)

# Test different optimizers with unified learning rate
unified_lr = 1.0
optimizers_to_test = ["Adam", "SGD", "RMSprop", "AdamW"]
results = {}

for opt_name in optimizers_to_test:
    # Reset for fair comparison
    torch.manual_seed(42)
    model = LinearRegressor(input_dim=10, output_dim=1, l1=32,
                           num_hidden_layers=2, lr=unified_lr)

    # Create optimizer with mapped learning rate
    if opt_name == "SGD":
        optimizer = model.get_optimizer(opt_name, momentum=0.9)
    else:
        optimizer = model.get_optimizer(opt_name)

    criterion = nn.MSELoss()

    # Train
    model.train()
    for epoch in range(50):
        for batch_X, batch_y in train_loader:
            optimizer.zero_grad()
            predictions = model(batch_X)
            loss = criterion(predictions, batch_y)
            loss.backward()
            optimizer.step()

    # Evaluate
    model.eval()
```

```

test_loss = 0.0
with torch.no_grad():
    for batch_X, batch_y in test_loader:
        predictions = model(batch_X)
        test_loss += criterion(predictions, batch_y).item()

avg_test_loss = test_loss / len(test_loader)
results[opt_name] = avg_test_loss

print(f"{opt_name:10s}: Test MSE = {avg_test_loss:.4f} "
      f"(actual lr = {optimizer.param_groups[0]['lr']:.6f})")

# Find best optimizer
best_opt = min(results, key=results.get)
print(f"\nBest optimizer: {best_opt} with MSE = {results[best_opt]:.4f}")

```

```

Adam      : Test MSE = 3859.4478 (actual lr = 0.001000)
SGD       : Test MSE = 5271.6840 (actual lr = 0.010000)
RMSprop   : Test MSE = 2769.5163 (actual lr = 0.010000)
AdamW     : Test MSE = 3866.1199 (actual lr = 0.001000)

```

```
Best optimizer: RMSprop with MSE = 2769.5163
```

58.5.2. Hyperparameter Optimization

```

from spotoptim import SpotOptim
from spotoptim.nn.linear_regressor import LinearRegressor
from spotoptim.data import get_diabetes_data loaders
import torch
import torch.nn as nn
import numpy as np

def train_model(X):
    """Objective function for hyperparameter optimization."""
    results = []

    # Load data once
    train_loader, test_loader, _ = get_diabetes_data loaders(batch_size=32, random_state=42)

    for params in X:
        lr_unified = 10 ** params[0] # Log scale: [-4, 0]
        optimizer_name = params[1]   # Factor: "Adam", "SGD", "RMSprop"

```

```

# Create model with unified lr - automatically scaled per optimizer
torch.manual_seed(42)
model = LinearRegressor(input_dim=10, output_dim=1, l1=32, num_hidden_layers=2, lr=lr_unified)
optimizer = model.get_optimizer(optimizer_name)

criterion = nn.MSELoss()

# Train
model.train()
for epoch in range(30):
    for batch_X, batch_y in train_loader:
        optimizer.zero_grad()
        predictions = model(batch_X)
        loss = criterion(predictions, batch_y)
        loss.backward()
        optimizer.step()

# Evaluate
model.eval()
test_loss = 0.0
with torch.no_grad():
    for batch_X, batch_y in test_loader:
        predictions = model(batch_X)
        test_loss += criterion(predictions, batch_y).item()

avg_test_loss = test_loss / len(test_loader)
results.append(avg_test_loss)

return np.array(results)

# Optimize unified lr across different optimizers
spot_optimizer = SpotOptim(
    fun=train_model,
    bounds=[(-4, 0), ("Adam", "SGD", "RMSprop")],
    var_type=["float", "factor"],
    max_iter=10, # Small for demo
    n_initial=5,
    seed=42
)
result = spot_optimizer.optimize()

print(f"\nBest unified lr: {10**result.x[0]:.6f}")
print(f"Best optimizer: {result.x[1]}")
print(f"Best test MSE: {result.fun:.4f}")

```

58. Learning Rate Mapping for Unified Optimizer Interface

```
# Show actual learning rate used
from spotoptim.utils.mapping import map_lr
actual_lr = map_lr(10**result.x[0], result.x[1])
print(f"Actual {result.x[1]} learning rate: {actual_lr:.6f}")
```

```
/var/folders/dw/pvtj6mt91znd0hftcztqb0k00000gn/T/ipykernel_40600/3866398708.py:50: Use
agnostic) factor metrics. Factor integer codes will be treated as ordinal numbers, wh
aware Kriging surrogate: Kriging(metric_factorial='hamming').
    spot_optimizer = SpotOptim(
```

```
Best unified lr: 0.024883
Best optimizer: SGD
Best test MSE: 12830.5124
Actual SGD learning rate: 0.000249
```

58.5.3. Hyperparameter Optimization with SpotOptim

Note, N_INITIAL and MAX_ITER are kept small for demonstration; increase for real use.

```
from spotoptim import SpotOptim
from spotoptim.nn.linear_regressor import LinearRegressor
from spotoptim.data import get_diabetes_data loaders
import torch.nn as nn
import torch
import numpy as np

MAX_ITER = 10
N_INITIAL = 5

def train_and_evaluate(X):
    """Objective function for hyperparameter optimization."""
    results = []

    # Load data once
    train_loader, test_loader, _ = get_diabetes_data loaders(
        batch_size=32, random_state=42
    )

    for params in X:
        # Extract hyperparameters
```

```

lr_unified = 10 ** params[0] # Log scale
optimizer_name = params[1]    # Factor variable
l1 = int(params[2])           # Integer
num_layers = int(params[3])   # Integer

# Create model with unified learning rate
model = LinearRegressor(
    input_dim=10,
    output_dim=1,
    l1=l1,
    num_hidden_layers=num_layers,
    lr=lr_unified # Automatically mapped per optimizer
)

# Get optimizer (lr already mapped internally)
if optimizer_name == "SGD":
    optimizer = model.get_optimizer(optimizer_name, momentum=0.9)
else:
    optimizer = model.get_optimizer(optimizer_name)

criterion = nn.MSELoss()

# Train
model.train()
for epoch in range(30):
    for batch_X, batch_y in train_loader:
        optimizer.zero_grad()
        predictions = model(batch_X)
        loss = criterion(predictions, batch_y)
        loss.backward()
        optimizer.step()

# Evaluate
model.eval()
test_loss = 0.0
with torch.no_grad():
    for batch_X, batch_y in test_loader:
        predictions = model(batch_X)
        test_loss += criterion(predictions, batch_y).item()

avg_test_loss = test_loss / len(test_loader)
results.append(avg_test_loss)

return np.array(results)

```

58. Learning Rate Mapping for Unified Optimizer Interface

```
# Optimize learning rate, optimizer choice, and architecture
optimizer = SpotOptim(
    fun=train_and_evaluate,
    bounds=[
        (-4, 0), # log10(lr_unified): [0.0001, 1.0]
        ("Adam", "SGD", "RMSprop", "AdamW"), # Optimizer choice
        (16, 128), # Layer size
        (1, 3) # Number of hidden layers
    ],
    var_type=["float", "factor", "int", "int"],
    max_iter=MAX_ITER,
    n_initial=N_INITIAL,
    seed=42
)

result = optimizer.optimize()

# Display results
print("\nOptimization Results:")
print(f"Best unified lr: {10**result.x[0]:.6f}")
print(f"Best optimizer: {result.x[1]}")
print(f"Best layer size: {int(result.x[2])}")
print(f"Best num layers: {int(result.x[3])}")
print(f"Best test MSE: {result.fun:.4f}")

# Show actual learning rate used
from spotoptim.utils.mapping import map_lr
actual_lr = map_lr(10**result.x[0], result.x[1])
print(f"Actual {result.x[1]} learning rate: {actual_lr:.6f}")
```

```
/var/folders/dw/pvtj6mt91znd0hftcztqb0k00000gn/T/ipykernel_40600/2835539742.py:68: Use
agnostic) factor metrics. Factor integer codes will be treated as ordinal numbers, wh
aware Kriging surrogate: Kriging(metric_factorial='hamming').
optimizer = SpotOptim(
```

```
Optimization Results:
Best unified lr: 1.000000
Best optimizer: RMSprop
Best layer size: 64
Best num layers: 2
Best test MSE: 2659.2904
Actual RMSprop learning rate: 0.010000
```

58.5.4. Log-Scale Hyperparameter Search

```

from spotoptim.utils.mapping import map_lr
import numpy as np

# Common pattern: sample unified lr from log scale
log_lr_range = np.linspace(-4, 0, 10) # [-4, -3.56, ..., 0]
optimizers = ["Adam", "SGD", "RMSprop"]

print("Log-scale learning rate search:")
print()
print(f"{'log_lr':<10} {'unified_lr':<12} {'Adam':<12} {'SGD':<12} {'RMSprop':<12}")
print("-" * 60)

for log_lr in log_lr_range:
    lr_unified = 10 ** log_lr
    lr_adam = map_lr(lr_unified, "Adam")
    lr_sgd = map_lr(lr_unified, "SGD")
    lr_rmsprop = map_lr(lr_unified, "RMSprop")

    print(f"{'log_lr':<10.2f} {'lr_unified':<12.6f} {'lr_adam':<12.8f} "
          f"{'lr_sgd':<12.8f} {'lr_rmsprop':<12.8f}")

```

Log-scale learning rate search:

log_lr	unified_lr	Adam	SGD	RMSprop
-4.00	0.000100	0.00000010	0.00000100	0.00000100
-3.56	0.000278	0.00000028	0.00000278	0.00000278
-3.11	0.000774	0.00000077	0.00000774	0.00000774
-2.67	0.002154	0.00000215	0.00002154	0.00002154
-2.22	0.005995	0.00000599	0.00005995	0.00005995
-1.78	0.016681	0.00001668	0.00016681	0.00016681
-1.33	0.046416	0.00004642	0.00046416	0.00046416
-0.89	0.129155	0.00012915	0.00129155	0.00129155
-0.44	0.359381	0.00035938	0.00359381	0.00359381
0.00	1.000000	0.00100000	0.01000000	0.01000000

58.5.5. Custom Learning Rate Schedules

58. Learning Rate Mapping for Unified Optimizer Interface

```
import torch
import torch.nn as nn
from spotoptim.nn.linear_regressor import LinearRegressor
from spotoptim.utils.mapping import map_lr

# Create model with unified lr
model = LinearRegressor(input_dim=10, output_dim=1, lr=1.0)

# Get initial optimizer
optimizer = model.get_optimizer("Adam")
initial_lr = optimizer.param_groups[0]['lr']
print(f"Initial learning rate: {initial_lr}")

# Use PyTorch learning rate scheduler
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=30, gamma=0.1)

# Training with scheduler
for epoch in range(100):
    # ... training code ...
    scheduler.step()

    if (epoch + 1) % 30 == 0:
        current_lr = optimizer.param_groups[0]['lr']
        print(f"Epoch {epoch+1}: lr = {current_lr:.8f}")
```

Initial learning rate: 0.001
Epoch 30: lr = 0.00010000
Epoch 60: lr = 0.00001000
Epoch 90: lr = 0.00000100

58.5.6. Direct Usage Without LinearRegressor

```
import torch
import torch.nn as nn
from spotoptim.utils.mapping import map_lr

# Define your own model
class MyModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(10, 32)
```



```

        self.fc2 = nn.Linear(32, 1)

    def forward(self, x):
        x = torch.relu(self.fc1(x))
        return self.fc2(x)

model = MyModel()

# Use map_lr to get optimizer-specific learning rate
unified_lr = 2.0
optimizer_name = "Adam"

actual_lr = map_lr(unified_lr, optimizer_name)
optimizer = torch.optim.Adam(model.parameters(), lr=actual_lr)

print(f"Unified lr: {unified_lr}")
print(f"Actual {optimizer_name} lr: {actual_lr}")

```

```

Unified lr: 2.0
Actual Adam lr: 0.002

```

58.6. Best Practices

58.6.1. Choosing Unified Learning Rate

For initial experiments:

- Start with `lr=1.0` (gives defaults for all optimizers)
- Test with `lr=0.1`, `lr=1.0`, `lr=10.0` to get a sense of scale

For hyperparameter optimization:

- Use log scale: sample from `[-4, 0]` or `[-3, 1]`
- Convert with `lr_unified = 10 ** log_lr`
- This gives reasonable ranges for all optimizers

For fine-tuning:

- If training is unstable: try smaller `lr` (e.g., 0.1 or 0.5)
- If training is too slow: try larger `lr` (e.g., 2.0 or 5.0)
- Monitor loss curves to adjust

58.6.2. Optimizer Selection Guidelines

Adam family (Adam, AdamW, NAdam, RAdam):

- ☐ Good default choice for most tasks
- ☐ Adaptive learning rates per parameter
- ☐ Works well out of the box
- Use `lr=1.0` as starting point

SGD:

- ☐ Good for large datasets
- ☐ Often achieves better generalization
- ☐ Requires momentum (e.g., 0.9)
- Use `lr=1.0` with `momentum=0.9`

RMSprop:

- ☐ Good for recurrent networks
- ☐ Handles non-stationary objectives
- Use `lr=1.0` as starting point

Others (Adadelata, Adagrad, etc.):

- Specialized use cases
- Start with `lr=1.0` and adjust

58.6.3. Common Patterns

```
# Pattern 1: Quick optimizer comparison
model = LinearRegressor(input_dim=10, output_dim=1, lr=1.0)
for opt in ["Adam", "SGD", "RMSprop"]:
    optimizer = model.get_optimizer(opt)
    # ... train and compare ...

# Pattern 2: Hyperparameter optimization
def objective(X):
    lr_unified = 10 ** X[:, 0] # Log scale
    optimizer_name = X[:, 1]   # Factor
    # ... use unified lr ...

# Pattern 3: Override model's lr
model = LinearRegressor(input_dim=10, output_dim=1, lr=1.0)
optimizer = model.get_optimizer("Adam", lr=2.0) # Override with 2.0
```

```
# Pattern 4: Direct mapping
from spotoptim.utils.mapping import map_lr
lr_actual = map_lr(unified_lr, optimizer_name)
optimizer = torch.optim.Adam(params, lr=lr_actual)
```

58.7. Troubleshooting

58.7.1. Issue: Training is unstable (loss explodes)

Solution: Learning rate is too high. Try:

```
model = LinearRegressor(input_dim=10, output_dim=1, lr=0.1) # Reduce from 1.0
```

58.7.2. Issue: Training is too slow (loss decreases very slowly)

Solution: Learning rate is too low. Try:

```
model = LinearRegressor(input_dim=10, output_dim=1, lr=5.0) # Increase from 1.0
```

58.7.3. Issue: Different results across optimizer runs

Solution: Set random seed for reproducibility:

```
import torch
torch.manual_seed(42)
```

58.7.4. Issue: Want to use raw learning rate without mapping

Solution: Use `use_default_scale=False`:

```
from spotoptim.utils.mapping import map_lr
lr = map_lr(0.001, "Adam", use_default_scale=False) # Returns 0.001 directly
```

58.7.5. Issue: Optimizer not supported

Solution: Check supported optimizers:

```
from spotoptim.utils.mapping import OPTIMIZER_DEFAULT_LR
print("Supported optimizers:", list(OPTIMIZER_DEFAULT_LR.keys()))
```

58.8. Technical Details

58.8.1. How It Works

The mapping is simple but effective:

$$\text{actual_lr} = \text{lr_unified} * \text{default_lr}[\text{optimizer_name}]$$

For example:

- $\text{map_lr}(1.0, \text{"Adam"}) \rightarrow 1.0 * 0.001 = 0.001$
- $\text{map_lr}(0.5, \text{"SGD"}) \rightarrow 0.5 * 0.01 = 0.005$
- $\text{map_lr}(2.0, \text{"RMSprop"}) \rightarrow 2.0 * 0.01 = 0.02$

This ensures that the same unified learning rate gives optimizer-specific learning rates in their typical working ranges.

58.8.2. Design Rationale

Why use defaults as scaling factors?

PyTorch's default learning rates are carefully chosen to work well for typical use cases. By using them as scaling factors:

1. $\text{lr}=1.0$ always gives sensible defaults
2. Scaling preserves the relative relationships between optimizers
3. Each optimizer stays in its optimal range
4. Easy to understand and explain

Comparison with spotPython's approach:

spotPython uses $\text{lr} = \text{lr_mult} * \text{default_lr}$ in `optimizer_handler()`. Our implementation:

- ☐ Separates mapping logic (testable, reusable)
- ☐ Provides standalone function (`map_lr()`)
- ☐ Comprehensive error handling and validation
- ☐ Extensive documentation and examples
- ☐ Full integration with `LinearRegressor`

58.8.3. Default Learning Rates

All values verified against PyTorch documentation:

```
OPTIMIZER_DEFAULT_LR = {
    "Adadelta": 1.0,
    "Adagrad": 0.01,
    "Adam": 0.001,
    "AdamW": 0.001,
    "SparseAdam": 0.001,
    "Adamax": 0.002,
    "ASGD": 0.01,
    "LBFGS": 1.0,
    "NAdam": 0.002,
    "RAdam": 0.001,
    "RMSprop": 0.01,
    "Rprop": 0.01,
    "SGD": 0.01,
}
```

58.9. Examples

58.9.1. Complete Example: Optimizer Comparison Study

```
"""
Complete example: Compare optimizers with unified learning rate interface.
"""

import torch
import torch.nn as nn
from spotoptim.nn.linear_regressor import LinearRegressor
from spotoptim.data import get_diabetes_dataloaders
from spotoptim.utils.mapping import map_lr
import matplotlib.pyplot as plt

# Set seed for reproducibility
torch.manual_seed(42)

# Load data
train_loader, test_loader, _ = get_diabetes_dataloaders(
    batch_size=32,
    random_state=42
```

```

)

# Test configurations
optimizers = ["Adam", "SGD", "RMSprop", "AdamW"]
unified_lrs = [0.5, 1.0, 2.0]

# Store results
results = {}

print("Training models with different optimizers and learning rates...")
print()

for unified_lr in unified_lrs:
    results[unified_lr] = {}

    for opt_name in optimizers:
        # Reset model for fair comparison
        torch.manual_seed(42)

        # Create model with unified lr
        model = LinearRegressor(
            input_dim=10,
            output_dim=1,
            l1=32,
            num_hidden_layers=2,
            lr=unified_lr
        )

        # Get optimizer
        if opt_name == "SGD":
            optimizer = model.get_optimizer(opt_name, momentum=0.9)
        else:
            optimizer = model.get_optimizer(opt_name)

        actual_lr = optimizer.param_groups[0]['lr']
        criterion = nn.MSELoss()

        # Track training loss
        train_losses = []

        # Train
        model.train()
        for epoch in range(50):
            epoch_loss = 0.0

```

```

        for batch_X, batch_y in train_loader:
            optimizer.zero_grad()
            predictions = model(batch_X)
            loss = criterion(predictions, batch_y)
            loss.backward()
            optimizer.step()
            epoch_loss += loss.item()

        avg_epoch_loss = epoch_loss / len(train_loader)
        train_losses.append(avg_epoch_loss)

    # Evaluate on test set
    model.eval()
    test_loss = 0.0
    with torch.no_grad():
        for batch_X, batch_y in test_loader:
            predictions = model(batch_X)
            test_loss += criterion(predictions, batch_y).item()

    avg_test_loss = test_loss / len(test_loader)
    results[unified_lr][opt_name] = {
        'train_losses': train_losses,
        'test_loss': avg_test_loss,
        'actual_lr': actual_lr
    }

    print(f"Unified lr={unified_lr:.1f}, {opt_name:10s}: "
          f"actual_lr={actual_lr:.6f}, test_MSE={avg_test_loss:.4f}")

# Display summary
print()
print("=" * 70)
print("Summary: Best configurations")
print("=" * 70)

for unified_lr in unified_lrs:
    best_opt = min(results[unified_lr].items(),
                   key=lambda x: x[1]['test_loss'])
    opt_name, metrics = best_opt

    print(f"Unified lr={unified_lr:.1f}: {opt_name:10s} "
          f"(test MSE={metrics['test_loss']:.4f}, "
          f"actual lr={metrics['actual_lr']:.6f})")

```

58. Learning Rate Mapping for Unified Optimizer Interface

```
# Find overall best
best_overall = None
best_overall_loss = float('inf')

for unified_lr in unified_lrs:
    for opt_name, metrics in results[unified_lr].items():
        if metrics['test_loss'] < best_overall_loss:
            best_overall_loss = metrics['test_loss']
            best_overall = (unified_lr, opt_name, metrics['actual_lr'])

print()
print(f"Overall best: unified_lr={best_overall[0]:.1f}, "
      f"optimizer={best_overall[1]}, "
      f"test_MSE={best_overall_loss:.4f}")
print(f"Actual learning rate used: {best_overall[2]:.6f}")
```

Training models with different optimizers and learning rates...

Unified lr=0.5, Adam	: actual_lr=0.000500, test_MSE=5095.8745
Unified lr=0.5, SGD	: actual_lr=0.005000, test_MSE=5291.3903
Unified lr=0.5, RMSprop	: actual_lr=0.005000, test_MSE=2828.7086
Unified lr=0.5, AdamW	: actual_lr=0.000500, test_MSE=5107.9596
Unified lr=1.0, Adam	: actual_lr=0.001000, test_MSE=3859.4478
Unified lr=1.0, SGD	: actual_lr=0.010000, test_MSE=5271.6840
Unified lr=1.0, RMSprop	: actual_lr=0.010000, test_MSE=2769.5163
Unified lr=1.0, AdamW	: actual_lr=0.001000, test_MSE=3866.1199
Unified lr=2.0, Adam	: actual_lr=0.002000, test_MSE=3160.6882
Unified lr=2.0, SGD	: actual_lr=0.020000, test_MSE=nan
Unified lr=2.0, RMSprop	: actual_lr=0.020000, test_MSE=2883.3341
Unified lr=2.0, AdamW	: actual_lr=0.002000, test_MSE=3162.8033

=====
Summary: Best configurations

=====
Unified lr=0.5: RMSprop (test MSE=2828.7086, actual lr=0.005000)
Unified lr=1.0: RMSprop (test MSE=2769.5163, actual lr=0.010000)
Unified lr=2.0: RMSprop (test MSE=2883.3341, actual lr=0.020000)

Overall best: unified_lr=1.0, optimizer=RMSprop, test_MSE=2769.5163
Actual learning rate used: 0.010000

58.10. Jupyter Notebook

i Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

59. SpotOptim Scaling

60. TorchStandardScaler in SpotOptim

This tutorial demonstrates the usage of `TorchStandardScaler` and its integration with `TorchObjective` in `SpotOptim`.

60.1. Introduction

Scaling input features is a crucial step in machine learning, especially for models trained with gradient-based optimization like Neural Networks. `SpotOptim` provides a convenient `TorchStandardScaler` (mimicking `sklearn`'s `StandardScaler`) that handles PyTorch tensors correctly.

Additionally, `TorchObjective` can automatically apply this scaler to your data when `use_scaler=True` is set during initialization. This ensures that:

1. Scaling is fit **only** on the training set (preventing data leakage).
2. The same scaling transformation is applied to validation and test data.
3. The process is seamless and integrated into the objective evaluation.

60.2. Using TorchStandardScaler Directly

You can use the scaler independently for any PyTorch data processing tasks.

```
import torch
from spotoptim.utils.scaler import TorchStandardScaler

# Create synthetic data
# Feature 0: Scale ~100
# Feature 1: Scale ~1
X = torch.tensor([
    [100.0, 1.0],
    [110.0, 1.2],
    [90.0, 0.8]
])

scaler = TorchStandardScaler()
```

```
# Fit and Transform
X_scaled = scaler.fit_transform(X)

print("Original Mean:\n", X.mean(dim=0))
print("Scaled Mean:\n", X_scaled.mean(dim=0))
print("Scaled Std:\n", X_scaled.std(dim=0, unbiased=False))
```

```
Original Mean:
  tensor([100.,   1.])
Scaled Mean:
  tensor([0.0000e+00, 1.1921e-07])
Scaled Std:
  tensor([1.0000, 1.0000])
```

60.3. Integrating with TorchObjective

The most convenient usage is within TorchObjective.

When you initialize TorchObjective with `use_scaler=True`, it automatically intercepts the data preparation phase. It scales the features provided in the `SpotData` object before they are loaded into the PyTorch `DataLoader`.

60.3.1. Example: Scaling Effect

The following fully self-contained example demonstrates the effect of scaling. We will train a simple Linear Regression model on a dataset with vastly different feature scales.

1. **Without Scaling:** The model might struggle to converge or require a very carefully tuned learning rate.
2. **With Scaling:** The model typically converges faster and more reliably.

```
import torch
import torch.nn as nn
import numpy as np
import random
from spotoptim.core.experiment import ExperimentControl
from spotoptim.function.torch_objective import TorchObjective
from spotoptim.hyperparameters import ParameterSet
from spotoptim.core.data import SpotDataFromArray
```

```

def set_seed(seed=42):
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)

# 1. Create a Dataset with Disparate Scales
# Feature 1: range [0, 1]
# Feature 2: range [0, 1000]
set_seed(42)
n_samples = 200
X1 = np.random.rand(n_samples, 1)
X2 = np.random.rand(n_samples, 1) * 1000.0
X = np.hstack([X1, X2])

# Target: y = 2*x1 + 0.005*x2 + noise
# Note: Coefficient for x2 is small because x2 is large
y = 2 * X1 + 0.005 * X2 + np.random.normal(0, 0.1, size=(n_samples, 1))

# Manual split for validation
split_idx = int(n_samples * 0.8)
X_train, X_val = X[:split_idx], X[split_idx:]
y_train, y_val = y[:split_idx], y[split_idx:]

dataset = SpotDataFromArray(x_train=X_train, y_train=y_train, x_val=X_val, y_val=y_val)

# 2. Define a Simple Model
class LinearReg(nn.Module):
    def __init__(self, input_dim, output_dim, **kwargs):
        super().__init__()
        self.fc = nn.Linear(input_dim, output_dim)

    def forward(self, x):
        return self.fc(x)

# 3. Setup Experiment
# We will use SGD which is sensitive to feature scaling
params = ParameterSet()
params.add_float("lr", 1e-5, 1e-1, default=1e-3)

exp = ExperimentControl(
    experiment_name="scaling_test",
    model_class=LinearReg,
    dataset=dataset,
    hyperparameters=params,

```

```

    metrics=["val_loss"],
    epochs=50,
    batch_size=32,
    device="cpu", # Force CPU for simplicity in example
    loss_function=nn.MSELoss()
)

# 4. Run Without Scaling
print("--- No Scaling ---")
objective_no_scale = TorchObjective(exp, seed=42, use_scaler=False)
# Evaluate with a learning rate that might be tricky for unscaled data
# lr=0.001 is often too high for feature x2 (scale 1000) -> gradients will be huge
res_no_scale = objective_no_scale(np.array([[1e-4]])) # Use small LR to avoid explosion
print(f"Val Loss (No Scale): {res_no_scale[0, 0]:.4f}")

# 5. Run With Scaling
print("\n--- With Scaling ---")
objective_scaled = TorchObjective(exp, seed=42, use_scaler=True)
# With scaling, standard LRs like 0.01 or 0.001 work fine
res_scaled = objective_scaled(np.array([[1e-2]])) # Can use larger LR
print(f"Val Loss (Scaled): {res_scaled[0, 0]:.4f}")

# Comparison with same LR (if stable)
# A very small LR is needed for unscaled data to not diverge,
# but that makes convergence slow for the small feature x1.
# Scaled data allows balanced learning.

```

```

--- No Scaling ---
Val Loss (No Scale): 127379.6953

```

```

--- With Scaling ---
Val Loss (Scaled): 2.3018

```

60.3.2. Explanation

In the example above, X2 has values up to 1000.

- **Without Scaling:** The gradients with respect to weights for X2 will be 1000x larger than for X1. To prevent divergence, the learning rate must be very small (e.g., $1e-4$ or $1e-5$). However, this small learning rate will make learning the weight for X1 extremely slow.

60.3. Integrating with TorchObjective

- **With Scaling:** Both `X1` and `X2` are transformed to have approximately zero mean and unit variance. The error surface becomes more spherical (isotropic), allowing standard learning rates (e.g., `1e-2`) to efficiently train weights for both features simultaneously.

By simply adding `use_scaler=True`, `TorchObjective` handles this best practice for you automatically.

Part X.

Data Sets

61. Diabetes Dataset Utilities

SpotOptim provides convenient utilities for working with the sklearn diabetes dataset, including PyTorch `Dataset` and `DataLoader` implementations. These utilities simplify data loading, preprocessing, and model training for regression tasks.

61.1. Overview

The diabetes dataset contains 10 baseline variables (age, sex, body mass index, average blood pressure, and six blood serum measurements) for 442 diabetes patients. The target is a quantitative measure of disease progression one year after baseline.

Module: `spotoptim.data.diabetes`

Key Components:

- `DiabetesDataset`: PyTorch Dataset class
- `get_diabetes_data loaders()`: Convenience function for complete data pipeline

61.2. Quick Start

61.2.1. Basic Usage

```
from spotoptim.data import get_diabetes_data loaders
from sklearn.datasets import load_diabetes
from spotoptim.data.diabetes import DiabetesDataset
import numpy as np

# Load data
diabetes = load_diabetes()
X = diabetes.data
y = diabetes.target.reshape(-1, 1)

# Now create the dataset
dataset = DiabetesDataset(X, y, transform=None, target_transform=None)
```

61. Diabetes Dataset Utilities

```
# Load data with default settings
train_loader, test_loader, scaler = get_diabetes_dataloaders()

# Iterate through batches
for batch_X, batch_y in train_loader:
    print(f"Batch features: {batch_X.shape}") # (32, 10)
    print(f"Batch targets: {batch_y.shape}") # (32, 1)
    break
```

```
Batch features: torch.Size([32, 10])
Batch targets: torch.Size([32, 1])
```

61.2.2. Training a Model

```
import torch
import torch.nn as nn
from spotoptim.data import get_diabetes_dataloaders
from spotoptim.nn.linear_regressor import LinearRegressor

# Load data
train_loader, test_loader, scaler = get_diabetes_dataloaders(
    test_size=0.2,
    batch_size=32,
    scale_features=True,
    random_state=42
)

# Create model
model = LinearRegressor(
    input_dim=10,
    output_dim=1,
    l1=64,
    num_hidden_layers=2,
    activation="ReLU"
)

# Setup training
criterion = nn.MSELoss()
optimizer = model.get_optimizer("Adam", lr=0.01)

# Training loop
num_epochs = 100
```

```

for epoch in range(num_epochs):
    model.train()
    train_loss = 0.0

    for batch_X, batch_y in train_loader:
        # Forward pass
        predictions = model(batch_X)
        loss = criterion(predictions, batch_y)

        # Backward pass
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        train_loss += loss.item()

    avg_train_loss = train_loss / len(train_loader)

    if (epoch + 1) % 20 == 0:
        print(f"Epoch {epoch+1}/{num_epochs}: Loss = {avg_train_loss:.4f}")

# Evaluation
model.eval()
test_loss = 0.0

with torch.no_grad():
    for batch_X, batch_y in test_loader:
        predictions = model(batch_X)
        loss = criterion(predictions, batch_y)
        test_loss += loss.item()

avg_test_loss = test_loss / len(test_loader)
print(f"Test MSE: {avg_test_loss:.4f}")

```

```

Epoch 20/100: Loss = 28036.1889
Epoch 40/100: Loss = 27898.2701
Epoch 60/100: Loss = 27592.0660
Epoch 80/100: Loss = 27952.7178
Epoch 100/100: Loss = 33014.2591
Test MSE: 26482.4193

```

61.3. Function Reference

61.3.1. `get_diabetes_dataloaders()`

Loads the sklearn diabetes dataset and returns configured PyTorch DataLoaders.

Signature:

```
get_diabetes_dataloaders(
    test_size=0.2,
    batch_size=32,
    shuffle_train=True,
    shuffle_test=False,
    random_state=42,
    scale_features=True,
    num_workers=0,
    pin_memory=False
)
```

```
(<torch.utils.data.dataloader.DataLoader at 0x124fb0050>,
 <torch.utils.data.dataloader.DataLoader at 0x124f49a30>,
 StandardScaler())
```

Parameters:

Parameter	Type	Default	Description
<code>test_size</code>	float	0.2	Proportion of dataset for testing (0.0 to 1.0)
<code>batch_size</code>	int	32	Number of samples per batch
<code>shuffle_train</code>	bool	True	Whether to shuffle training data
<code>shuffle_test</code>	bool	False	Whether to shuffle test data
<code>random_state</code>	int	42	Random seed for train/test split
<code>scale_features</code>	bool	True	Whether to standardize features
<code>num_workers</code>	int	0	Number of subprocesses for data loading
<code>pin_memory</code>	bool	False	Whether to pin memory (useful for GPU)

Returns:

- `train_loader` (`DataLoader`): Training data loader
- `test_loader` (`DataLoader`): Test data loader
- `scaler` (`StandardScaler` or `None`): Fitted scaler if `scale_features=True`, else `None`

Example:

```
from spotoptim.data import get_diabetes_dataloaders

# Custom configuration
train_loader, test_loader, scaler = get_diabetes_dataloaders(
    test_size=0.3,
    batch_size=64,
    shuffle_train=True,
    scale_features=True,
    random_state=123
)

print(f"Training batches: {len(train_loader)}")
print(f"Test batches: {len(test_loader)}")
print(f"Scaler mean: {scaler.mean_[0:3]}") # First 3 features
```

```
Training batches: 5
Test batches: 3
Scaler mean: [-0.00056537  0.00132258  0.00027836]
```

61.4. DiabetesDataset Class

PyTorch Dataset implementation for the diabetes dataset.

Signature:

```
DiabetesDataset(X, y, transform=None, target_transform=None)
```

```
<spotoptim.data.diabetes.DiabetesDataset at 0x124f7d250>
```

Parameters:

- `X` (`np.ndarray`): Feature matrix of shape `(n_samples, n_features)`
- `y` (`np.ndarray`): Target values of shape `(n_samples,)` or `(n_samples, 1)`
- `transform` (callable, optional): Transform to apply to features

61. Diabetes Dataset Utilities

- `target_transform` (callable, optional): Transform to apply to targets

Attributes:

- `X` (torch.Tensor): Feature tensor (n_samples, n_features)
- `y` (torch.Tensor): Target tensor (n_samples, 1)
- `n_features` (int): Number of features (10 for diabetes)
- `n_samples` (int): Number of samples

Methods:

- `__len__()`: Returns number of samples
- `__getitem__(idx)`: Returns tuple (features, target) for given index

61.4.1. Manual Dataset Creation

```
from spotoptim.data import DiabetesDataset
from sklearn.datasets import load_diabetes
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from torch.utils.data import DataLoader

# Load raw data
diabetes = load_diabetes()
X, y = diabetes.data, diabetes.target

# Split data
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# Scale features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Create datasets
train_dataset = DiabetesDataset(X_train, y_train)
test_dataset = DiabetesDataset(X_test, y_test)

# Create dataloaders
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)
```

```
# Inspect dataset
print(f"Dataset size: {len(train_dataset)}")
print(f"Features shape: {train_dataset.X.shape}")
print(f"Targets shape: {train_dataset.y.shape}")

# Get a sample
features, target = train_dataset[0]
print(f"Sample features: {features.shape}") # (10,)
print(f"Sample target: {target.shape}")    # (1,)
```

```
Dataset size: 353
Features shape: torch.Size([353, 10])
Targets shape: torch.Size([353, 1])
Sample features: torch.Size([10])
Sample target: torch.Size([1])
```

61.5. Advanced Usage

61.5.1. Custom Transforms

```
from spotoptim.data import DiabetesDataset
from sklearn.datasets import load_diabetes
import torch

# Define custom transforms
def add_noise(x):
    """Add Gaussian noise to features."""
    return x + torch.randn_like(x) * 0.01

def log_transform(y):
    """Apply log transform to target."""
    return torch.log1p(y)

# Load data
diabetes = load_diabetes()
X, y = diabetes.data, diabetes.target

# Create dataset with transforms
dataset = DiabetesDataset(
    X, y,
    transform=add_noise,
```

```
        target_transform=log_transform
    )

    # Transforms are applied when accessing items
    features, target = dataset[0]
```

61.5.2. Different Train/Test Splits

```
from spotoptim.data import get_diabetes_dataloaders

# 70/30 split
train_loader, test_loader, scaler = get_diabetes_dataloaders(
    test_size=0.3,
    random_state=42
)
print(f"Training samples: {len(train_loader.dataset)}") # ~310
print(f"Test samples: {len(test_loader.dataset)}")      # ~132

# 90/10 split
train_loader, test_loader, scaler = get_diabetes_dataloaders(
    test_size=0.1,
    random_state=42
)
print(f"Training samples: {len(train_loader.dataset)}") # ~398
print(f"Test samples: {len(test_loader.dataset)}")      # ~44
```

```
Training samples: 309
Test samples: 133
Training samples: 397
Test samples: 45
```

61.5.3. Without Feature Scaling

```
from spotoptim.data import get_diabetes_dataloaders

# Load without scaling (useful for tree-based models)
train_loader, test_loader, scaler = get_diabetes_dataloaders(
    scale_features=False
)
```

```
print(f"Scaler: {scaler}") # None

# Data is in original scale
for batch_X, batch_y in train_loader:
    print(f"Mean: {batch_X.mean(dim=0)[:3]}") # Non-zero values
    break
```

```
Scaler: None
Mean: tensor([ 0.0103, -0.0029, -0.0006])
```

61.5.4. Larger Batch Sizes

```
from spotoptim.data import get_diabetes_dataloaders

# Larger batches for faster training (if memory allows)
train_loader, test_loader, scaler = get_diabetes_dataloaders(
    batch_size=128
)
print(f"Batches per epoch: {len(train_loader)}") # Fewer batches

# Smaller batches for more gradient updates
train_loader, test_loader, scaler = get_diabetes_dataloaders(
    batch_size=8
)
print(f"Batches per epoch: {len(train_loader)}") # More batches
```

```
Batches per epoch: 3
Batches per epoch: 45
```

61.5.5. GPU Training with Pin Memory

```
import torch
from spotoptim.data import get_diabetes_dataloaders

# Enable pin_memory for faster GPU transfer
train_loader, test_loader, scaler = get_diabetes_dataloaders(
    batch_size=32,
    pin_memory=True # Set to True when using GPU
)
```

61. Diabetes Dataset Utilities

```
# Move model to GPU
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = model.to(device)

# Training loop with GPU
for batch_X, batch_y in train_loader:
    # Data is already pinned, faster transfer to GPU
    batch_X = batch_X.to(device, non_blocking=True)
    batch_y = batch_y.to(device, non_blocking=True)

    # ... training code ...
```

```
/Users/bartz/workspace/spotoptim-cookbook/.venv/lib/python3.14/site-
packages/torch/utils/data/dataloader.py:752: UserWarning: 'pin_memory' argument is set
super().__init__(loader)
```

61.6. Complete Training Example

Here's a complete example showing data loading, model training, and evaluation:

```
import torch
import torch.nn as nn
from spotoptim.data import get_diabetes_dataloaders
from spotoptim.nn.linear_regressor import LinearRegressor

def train_diabetes_model():
    """Train a neural network on the diabetes dataset."""

    # Load data
    train_loader, test_loader, scaler = get_diabetes_dataloaders(
        test_size=0.2,
        batch_size=32,
        scale_features=True,
        random_state=42
    )

    # Create model
    model = LinearRegressor(
        input_dim=10,
        output_dim=1,
        l1=128,
```

```

        num_hidden_layers=3,
        activation="ReLU"
    )

    # Setup training
    criterion = nn.MSELoss()
    optimizer = model.get_optimizer("Adam", lr=0.001, weight_decay=1e-5)

    # Training configuration
    num_epochs = 200
    best_test_loss = float('inf')

    print("Starting training...")
    print(f"Training samples: {len(train_loader.dataset)}")
    print(f"Test samples: {len(test_loader.dataset)}")
    print(f"Batches per epoch: {len(train_loader)}")
    print("-" * 60)

    for epoch in range(num_epochs):
        # Training phase
        model.train()
        train_loss = 0.0

        for batch_X, batch_y in train_loader:
            # Forward pass
            predictions = model(batch_X)
            loss = criterion(predictions, batch_y)

            # Backward pass
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

            train_loss += loss.item()

        avg_train_loss = train_loss / len(train_loader)

        # Evaluation phase
        model.eval()
        test_loss = 0.0

        with torch.no_grad():
            for batch_X, batch_y in test_loader:
                predictions = model(batch_X)

```

```

        loss = criterion(predictions, batch_y)
        test_loss += loss.item()

    avg_test_loss = test_loss / len(test_loader)

    # Track best model
    if avg_test_loss < best_test_loss:
        best_test_loss = avg_test_loss
        # Could save model here: torch.save(model.state_dict(), 'best_model.pt')

    # Print progress
    if (epoch + 1) % 20 == 0:
        print(f"Epoch {epoch+1:3d}/{num_epochs}: "
              f"Train Loss = {avg_train_loss:.4f}, "
              f"Test Loss = {avg_test_loss:.4f}")

    print("-" * 60)
    print(f"Training complete!")
    print(f"Best test loss: {best_test_loss:.4f}")

    return model, best_test_loss

# Run training
if __name__ == "__main__":
    model, best_loss = train_diabetes_model()

```

Starting training...

Training samples: 353

Test samples: 89

Batches per epoch: 12

```

-----
Epoch 20/200: Train Loss = 31197.9932, Test Loss = 26610.6426
Epoch 40/200: Train Loss = 29353.8112, Test Loss = 26605.9531
Epoch 60/200: Train Loss = 31738.0871, Test Loss = 26601.1530
Epoch 80/200: Train Loss = 35027.6898, Test Loss = 26596.4368
Epoch 100/200: Train Loss = 30155.0223, Test Loss = 26591.5130
Epoch 120/200: Train Loss = 36371.8045, Test Loss = 26586.4290
Epoch 140/200: Train Loss = 31162.0645, Test Loss = 26581.1094
Epoch 160/200: Train Loss = 30484.9867, Test Loss = 26575.4570
Epoch 180/200: Train Loss = 30675.4717, Test Loss = 26569.7292
Epoch 200/200: Train Loss = 27566.1583, Test Loss = 26563.7344
-----

```

Training complete!

Best test loss: 26563.7344

61.7. Integration with SpotOptim

Use the diabetes dataset for hyperparameter optimization with SpotOptim:

```
import numpy as np
import torch
import torch.nn as nn
from spotoptim import SpotOptim
from spotoptim.data import get_diabetes_dataloaders
from spotoptim.nn.linear_regressor import LinearRegressor

def evaluate_model(X):
    """Objective function for SpotOptim.

    Args:
        X: Array of hyperparameters [lr, l1, num_hidden_layers]

    Returns:
        Array of validation losses
    """
    results = []

    for params in X:
        lr, l1, num_hidden_layers = params
        lr = 10 ** lr # Log scale for learning rate
        l1 = int(l1)
        num_hidden_layers = int(num_hidden_layers)

        # Load data
        train_loader, test_loader, _ = get_diabetes_dataloaders(
            test_size=0.2,
            batch_size=32,
            random_state=42
        )

        # Create model
        model = LinearRegressor(
            input_dim=10,
            output_dim=1,
            l1=l1,
            num_hidden_layers=num_hidden_layers,
            activation="ReLU"
        )
```

```

# Train briefly
criterion = nn.MSELoss()
optimizer = model.get_optimizer("Adam", lr=lr)

num_epochs = 50
for epoch in range(num_epochs):
    model.train()
    for batch_X, batch_y in train_loader:
        predictions = model(batch_X)
        loss = criterion(predictions, batch_y)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

# Evaluate
model.eval()
test_loss = 0.0
with torch.no_grad():
    for batch_X, batch_y in test_loader:
        predictions = model(batch_X)
        loss = criterion(predictions, batch_y)
        test_loss += loss.item()

results.append(test_loss / len(test_loader))

return np.array(results)

# Optimize hyperparameters
optimizer = SpotOptim(
    fun=evaluate_model,
    bounds=[
        (-4, -2), # log10(lr): 0.0001 to 0.01
        (16, 128), # l1: number of neurons
        (0, 4) # num_hidden_layers
    ],
    var_type=["float", "int", "int"],
    max_iter=30,
    n_initial=10,
    seed=42,
    verbose=True
)

result = optimizer.optimize()
print(f"Best hyperparameters found:")

```

```
print(f"  Learning rate: {10**result.x[0]:.6f}")
print(f"  Hidden neurons (l1): {int(result.x[1])}")
print(f"  Hidden layers: {int(result.x[2])}")
print(f"  Best MSE: {result.fun:.4f}")
```

TensorBoard logging disabled

Initial best: $f(x) = 26536.645833$

Iter 1		Best: 26536.645833		Curr: 26620.356120		Rate: 0.00		Evals: 36.7%
Iter 2		Best: 26536.645833		Curr: 26568.520833		Rate: 0.00		Evals: 40.0%
Iter 3		Best: 26536.645833		Curr: 26575.061198		Rate: 0.00		Evals: 43.3%
Iter 4		Best: 26536.645833		Curr: 26654.526693		Rate: 0.00		Evals: 46.7%
Iter 5		Best: 26536.645833		Curr: 26606.641927		Rate: 0.00		Evals: 50.0%
Iter 6		Best: 26536.645833		Curr: 26630.983073		Rate: 0.00		Evals: 53.3%
Iter 7		Best: 26536.645833		Curr: 26596.570964		Rate: 0.00		Evals: 56.7%
Iter 8		Best: 26536.645833		Curr: 26655.807292		Rate: 0.00		Evals: 60.0%
Iter 9		Best: 26536.645833		Curr: 26600.944661		Rate: 0.00		Evals: 63.3%
Iter 10		Best: 26536.645833		Curr: 26568.430339		Rate: 0.00		Evals: 66.7%
Iter 11		Best: 26536.645833		Curr: 26650.149740		Rate: 0.00		Evals: 70.0%
Optimizer candidate 1/3 was duplicate/invalid.								
Iter 12		Best: 26536.645833		Curr: 26618.976562		Rate: 0.00		Evals: 73.3%
Optimizer candidate 1/3 was duplicate/invalid.								
Iter 13		Best: 26536.645833		Curr: 26615.913411		Rate: 0.00		Evals: 76.7%
Iter 14		Best: 26536.645833		Curr: 26628.293620		Rate: 0.00		Evals: 80.0%
Iter 15		Best: 26536.645833		Curr: 26614.167969		Rate: 0.00		Evals: 83.3%
Iter 16		Best: 26536.645833		Curr: 26568.373047		Rate: 0.00		Evals: 86.7%
Iter 17		Best: 26536.645833		Curr: 26594.309896		Rate: 0.00		Evals: 90.0%
Iter 18		Best: 26536.645833		Curr: 26586.642578		Rate: 0.00		Evals: 93.3%
Iter 19		Best: 26536.645833		Curr: 26610.666016		Rate: 0.00		Evals: 96.7%
Iter 20		Best: 26532.748698		Rate: 0.05		Evals: 100.0%		

Best hyperparameters found:

```
  Learning rate: 0.000240
  Hidden neurons (l1): 91
  Hidden layers: 0
  Best MSE: 26532.7487
```

61.8. Best Practices

61.8.1. 1. Always Use Feature Scaling

```
# Good: Features are standardized
train_loader, test_loader, scaler = get_diabetes_dataloaders(
```

61. Diabetes Dataset Utilities

```
    scale_features=True
)
```

Neural networks typically perform better with normalized inputs.

61.8.2. 2. Set Random Seeds for Reproducibility

```
# Reproducible train/test splits
train_loader, test_loader, scaler = get_diabetes_dataloaders(
    random_state=42
)

# Also set PyTorch seed
import torch
torch.manual_seed(42)
```

```
<torch._C.Generator at 0x11dcffa30>
```

61.8.3. 3. Don't Shuffle Test Data

```
# Good: Test data in consistent order
train_loader, test_loader, scaler = get_diabetes_dataloaders(
    shuffle_train=True,    # Shuffle training data
    shuffle_test=False     # Don't shuffle test data
)
```

This ensures consistent evaluation metrics across runs.

61.8.4. 4. Choose Appropriate Batch Size

```
# Small dataset (442 samples) - moderate batch size works well
train_loader, test_loader, scaler = get_diabetes_dataloaders(
    batch_size=32 # Good balance for this dataset
)
```

Too large: Fewer gradient updates per epoch
Too small: Noisy gradients, slower training

61.8.5. 5. Save the Scaler for Production

```
import pickle
import numpy as np
from spotoptim.data import get_diabetes_data loaders

# Train with scaling
train_loader, test_loader, scaler = get_diabetes_data loaders(
    scale_features=True
)

# Save scaler for production use
with open('scaler.pkl', 'wb') as f:
    pickle.dump(scaler, f)

# Later: Load and use on new data
with open('scaler.pkl', 'rb') as f:
    loaded_scaler = pickle.load(f)

# Create some example new data (same shape as diabetes features)
new_data = np.random.randn(5, 10) # 5 samples, 10 features
new_data_scaled = loaded_scaler.transform(new_data)

print(f"Original data shape: {new_data.shape}")
print(f"Scaled data shape: {new_data_scaled.shape}")
print(f"Scaled data mean: {new_data_scaled.mean(axis=0)[:3]}") # Should be close to 0
```

```
Original data shape: (5, 10)
Scaled data shape: (5, 10)
Scaled data mean: [7.03934676 5.01490242 3.47123618]
```

61.9. Troubleshooting

61.9.1. Issue: Out of Memory

Solution: Reduce batch size or disable pin_memory

```
train_loader, test_loader, scaler = get_diabetes_data loaders(
    batch_size=16,      # Smaller batches
    pin_memory=False    # Disable if not using GPU
)
```

61.9.2. Issue: Different Data Ranges

Symptom: Model not converging, loss is NaN

Solution: Ensure feature scaling is enabled

```
train_loader, test_loader, scaler = get_diabetes_dataloaders(
    scale_features=True # Must be True for neural networks
)
```

61.9.3. Issue: Non-Reproducible Results

Solution: Set all random seeds

```
import torch
import numpy as np

# Set all seeds
torch.manual_seed(42)
np.random.seed(42)

train_loader, test_loader, scaler = get_diabetes_dataloaders(
    random_state=42,
    shuffle_train=False # Disable shuffle for full reproducibility
)
```

61.9.4. Issue: Slow Data Loading

Solution: Use multiple workers (if not on Windows)

```
train_loader, test_loader, scaler = get_diabetes_dataloaders(
    num_workers=4, # Use 4 subprocesses
    pin_memory=True # Enable for GPU
)
```

Note: On Windows, set `num_workers=0` to avoid multiprocessing issues.

61.10. Summary

The diabetes dataset utilities in SpotOptim provide:

- **Easy data loading:** One function call gets complete data pipeline
- **PyTorch integration:** Native Dataset and DataLoader support
- **Preprocessing included:** Automatic feature scaling and train/test splitting
- **Flexible configuration:** Control batch size, splitting, scaling, and more
- **Production ready:** Save scalers and ensure reproducibility

61.11. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

Part XI.

Details About SPOT

62. SpotOptim Step-by-Step Optimization Process

63. Introduction

This document provides a comprehensive step-by-step explanation of the optimization process in the `SpotOptim` class. We'll use the 2-dimensional Rosenbrock function as our example and cover all methods involved in each stage of optimization.

Topics Covered:

- Standard optimization workflow
- Handling noisy functions with repeats
- Handling function evaluation failures (NaN/inf values)

64. Setup and Test Functions

Let's start by defining our test functions, including variants with noise and occasional failures. The two-dimensional Rosenbrock function is defined as:

$$f(x, y) = (a - x)^2 + b(y - x^2)^2$$

where typically ($a = 1$) and ($b = 100$). The generalized form for n dimensions is:

$$f(X) = \sum_{i=1}^{N-1} [100 \cdot (x_{i+1} - x_i^2)^2 + (1 - x_i)^2]$$

The documentation can be found here: [DOC](#)

```
import numpy as np
import time
from spotoptim import SpotOptim
from spotoptim.function import rosenbrock
import warnings
warnings.filterwarnings('ignore')
```

Set random seed for reproducibility:

```
np.random.seed(42)
```

Based on the standard Rosenbrock function, we define two variants:

1. the noisy Rosenbrock function:

```
def rosenbrock_noisy(X, noise_std=0.1):
    """
    Rosenbrock with Gaussian noise for testing noisy optimization.
    """
    X = np.atleast_2d(X)
    base_values = rosenbrock(X)
    noise = np.random.normal(0, noise_std, size=base_values.shape)
    return base_values + noise
```

and 2. the Rosenbrock function with occasional failures:

64. Setup and Test Functions

```
def rosenbrock_with_failures(X, failure_prob=0.15):
    """
    Rosenbrock that occasionally returns NaN to simulate evaluation failures.
    """
    X = np.atleast_2d(X)
    values = rosenbrock(X)

    # Randomly inject failures
    for i in range(len(values)):
        if np.random.random() < failure_prob:
            values[i] = np.nan

    return values
```


65. Standard Optimization Workflow

Let's trace through a complete optimization run to understand each step.

65.1. Phase: Initialization and Setup

When you create a `SpotOptim` instance, several initialization steps occur during the `__init__()` method, see DOC.

1. Objective function (`fun`) and bounds are stored.
2. Acquisition function (`acquisition`) is stored (default is `y`).
3. Determine if noise handling is active (based on `repeats_initial` and `repeats_surrogate`)
4. For dimensions with tuple bounds (factor variables), internal integer mappings are created and bounds are replaced with $(0, n_levels-1)$. This is handled by `_process_factor_bounds()`. So, e.g., `[('red', 'green', 'blue')]` is mapped to the integer interval $[(0, 2)]$ internally.
5. The dimension of the problem (`n_dim`) is inferred from the bounds.
6. If `var_type` is not provided, it defaults to all “float” (continuous) variables, except for factor variables which are set to “factor” via `_process_factor_bounds()`.
7. Default variable names (`var_name`) are set if not provided.
8. Default variable transformations (`var_transform`) are set if not provided.
9. Transformations are applied to bounds based on `var_transform` settings. Natural bounds are stored in `_original_lower` and `_original_upper`.
10. Dimension reduction by identifying fixed dimensions (if any) is performed via `setup_dimension_reduction()`.
11. The surrogate is initialized (default: Gaussian Process with Matérn kernel) as follows:

```
kernel = ConstantKernel(  
    constant_value=1.0, constant_value_bounds=(1e-3, 1e3)  
) * Matern(length_scale=1.0, length_scale_bounds=(1e-2, 1e2), nu=2.5)  
self.surrogate = GaussianProcessRegressor(  
    kernel=kernel,  
    n_restarts_optimizer=10,  
    normalize_y=True,
```

```

    random_state=self.seed,
)

```

12. The Design generator is initialized (default: Latin Hypercube Sampling).
13. The storage for results ins initialized. This includes the following attributes:

- `X_`: Evaluated design points
- `y_`: Corresponding function values
- `y_mo`: For multi-objective functions, stores all objectives
- `best_x_`: Best point found so far
- `best_y_`: Best function value found so far
- `n_iter_`: Number of iterations completed
- `mean_X`, `mean_y`, `var_y`: For noisy functions, mean and variance tracking
- `min_mean_X`: Best mean point for noisy functions
- `min_mean_y`: Best mean value for noisy functions
- `min_var_y`: Variance at best mean point
- `min_X`: Best point found (deterministic)
- `min_y`: Best function value found (deterministic)
- `counter`: Total number of function evaluations
- `success_rate`: Ratio of successful evaluations
- `success_counter`: Count of successful evaluations
- `window_size`: For moving average of success rate
- `success_history`: History of success/failure for evaluations

```

# Create optimizer
opt = SpotOptim(
    fun=rosenbrock,
    bounds=[(-2, 2), (1, 2)],
    n_initial=10,
    max_iter=30,
    verbose=True,
    seed=42,
    var_trans=["id", "log"]
)

```

TensorBoard logging disabled

65.2. Phase: Initial Design Generation

65.2.1. Method: `get_initial_design()`

This method generates or processes the initial sample points. It is manually called here for demonstration (normally done inside `optimize()`).

What happens in `get_initial_design()`:

1. If `X0=None`: Generate space-filling design using Latin Hypercube Sampling (LHS)
2. If `x0` (starting point) provided: Include it as first point and transform user's points to internal scale
3. Apply dimension reduction if configured
4. Round integer/factor variables

```
X0 = opt.get_initial_design(X0=None)

print(f"Generated {len(X0)} initial design points using Latin Hypercube Sampling")
print(f"Design shape: {X0.shape}")
print(f"\nFirst 3 points (internal scale):")
print(X0[:3])
```

```
Generated 10 initial design points using Latin Hypercube Sampling
Design shape: (10, 2)
```

```
First 3 points (internal scale):
[[0.83330234 0.42205824]
 [0.049363   0.0478744 ]
 [1.21813757 0.40375168]]
```

Note, because `bounds` were set to `[(-2, 2), (-2, 2)]` and `n_initial=10`, the generated points are within this range and the design has two-dimensional shape (10, 2).

- Print initial design in original scale and in internal scale:

```
X0_original = opt.inverse_transform_X(X0)
print(f"\nFirst 3 points (original scale):")
print(X0_original[:3])
print(f"\nFirst 3 points (internal scale):")
print(X0[:3])
```


```
First 3 points (original scale):
[[0.83330234 1.52509735]
 [0.049363   1.04903889]
 [1.21813757 1.49743205]]
```

```
First 3 points (internal scale):
[[0.83330234 0.42205824]
 [0.049363   0.0478744 ]
 [1.21813757 0.40375168]]
```

65. Standard Optimization Workflow

- Plot initial design in original scale:

```
import matplotlib.pyplot as plt
plt.figure(figsize=(6, 6))
plt.scatter(X0_original[:, 0], X0_original[:, 1], c='blue', label='Initial Design Points')
plt.xlabel("X1")
plt.ylabel("X2")
plt.title("Initial Design Points (Original Scale)")
plt.legend()
plt.grid(True)
plt.show()
```



65.3. Phase 3: Initial Design Curation

65.3.1. Method: `curate_initial_design()`

This method ensures we have sufficient unique points and handles repeats.

What happens in `curate_initial_design()`:

1. Remove duplicate points (can occur after rounding integers)
2. Generate additional points if duplicates reduced count below `n_initial`
3. Repeat each point `repeats_initial` times if > 1 (for noisy functions)

```
X0_curated = opt.curate_initial_design(X0)
print(f"Curated design shape: {X0_curated.shape}")
print(f"Unique points: {len(np.unique(X0_curated, axis=0))}")
print(f"Total points (with repeats): {len(X0_curated)}")
```

```
Curated design shape: (10, 2)
Unique points: 10
Total points (with repeats): 10
```

65.4. Phase 4: Initial Design Evaluation

65.4.1. Method: `evaluate_function()`

Evaluates the objective function at all initial design points. The points are converted back to the original scale for evaluation.

What happens in `evaluate_function()`:

1. Convert points from internal to original scale
2. Call objective function with batch of points
3. Convert multi-objective to single-objective if needed
4. Return array of function values

```
X0_original = opt.inverse_transform_X(X0_curated)
y0 = opt.evaluate_function(X0_curated)

print(f"Evaluated {len(y0)} points")
print(f"Function values shape: {y0.shape}")
print(f"First 5 points with function values:")
for i in range(min(5, len(y0))):
    print(f"  Point {i+1}: {X0_original[i]} → f(x)={y0[i]:.6f}")
```

65. Standard Optimization Workflow

```
print(f"\nBest initial value: {np.min(y0):.6f}")
print(f"Worst initial value: {np.max(y0):.6f}")
print(f"Mean initial value: {np.mean(y0):.6f}")
```

Evaluated 10 points

Function values shape: (10,)

First 5 points with function values:

```
Point 1: [0.83330234 1.52509735] → f(x)=69.034795
Point 2: [0.049363  1.04903889] → f(x)=110.441323
Point 3: [1.21813757 1.49743205] → f(x)=0.066006
Point 4: [-0.39865126 1.65266768] → f(x)=225.083594
Point 5: [1.93984753 1.81713524] → f(x)=379.525559
```

Best initial value: 0.066006

Worst initial value: 379.525559

Mean initial value: 144.606182

65.5. Phase 5: Handling Failed Evaluations

65.5.1. Method: `rm_initial_design_NA_values(X0,y0)`

Removes points that returned NaN or inf values. In contrast to later phases, no penalties are applied here; invalid points of the initial design are simply removed.

What happens in `rm_initial_design_NA_values()`:

1. Identify NaN/inf values in function evaluations
2. Remove corresponding design points
3. Return cleaned arrays and original count
4. Note: No penalties applied in initial design - invalid points simply removed

```
n_before = len(y0)
X0_clean, y0_clean, n_evaluated = opt.rm_initial_design_NA_values(X0_curated, y0)

print(f"Points before filtering: {n_before}")
print(f"Points after filtering: {len(y0_clean)}")
print(f"Removed: {n_before - len(y0_clean)} NaN/inf values")
print(f"\nAll remaining values finite: {np.all(np.isfinite(y0_clean))}")
```

Points before filtering: 10

Points after filtering: 10

Removed: 0 NaN/inf values

All remaining values finite: True

65.6. Phase 6: Validation Check

65.6.1. Method: `check_size_initial_design(y0, n_evaluated)`

Ensures we have enough valid points to continue. The minimum required is the smaller of:

- typical minimum for surrogate fitting (3 for multi-dimensional, 2 for 1D), or
- what the user requested (`n_initial`).

What happens in `check_size_initial_design()`:

1. Check if at least 1 valid point exists
2. Raise error if all initial evaluations failed
3. Print warnings if many points were invalid

```
try:
    opt.check_size_initial_design(y0_clean, n_evaluated)
    print(f"x Validation passed: {len(y0_clean)} valid points available")
    print(f"  Minimum required: 1 point")
    print(f"  Original evaluated: {n_evaluated} points")
except ValueError as e:
    print(f" Validation failed: {e}")
```

```
x Validation passed: 10 valid points available
  Minimum required: 1 point
  Original evaluated: 10 points
```

65.7. Phase 7: Storage Initialization

65.7.1. Method: `init_storage()`

Store evaluated points and initialize tracking variables.

What happens during storage initialization:

1. `init_storage()`: Initialize statistics tracking variables.
 - Initialize storage (as done in `optimize()`)

65. Standard Optimization Workflow

```
# Initialize storage and statistics using the new init_storage() method
opt.init_storage(X0_clean, y0_clean)
```

```
print(f"X_ (evaluated points): shape {opt.X_.shape}")
print(f"y_ (function values): shape {opt.y_.shape}")
print(f"n_iter_ (iterations): {opt.n_iter_}")
```

```
X_ (evaluated points): shape (10, 2)
y_ (function values): shape (10,)
n_iter_ (iterations): 0
```

65.8. Phase: Statistics Update

65.8.1. Method: `update_stats()`

Update statistics like means and variances for noisy evaluations.

*** What happens in `update_stats()`:***

1. If noise is set (i.e., `repeats_initial > 1`):
 - Compute mean and variance of function values for each unique point
 - Store in `mean_X`, `mean_y`, and `var_y`

```
print(f"\nStatistics updated:")
if opt.repeats_initial > 1 or opt.repeats_surrogate > 1:
    print(f" - mean_X: {opt.mean_X.shape}")
    print(f" - mean_y: {opt.mean_y.shape}")
    print(f" - var_y: {opt.var_y.shape}")
else:
    print(f" - No noise tracking (repeats=1)")
```

```
Statistics updated:
- No noise tracking (repeats=1)
```

65.9. Phase: Log initial Design to TensorBoard

65.9.1. Method: `_log_initial_design_tensorboard()`

This method logs the initial design points and their evaluations to TensorBoard for visualization.

65.10. Phase: Initial Best Point

65.10.1. Method: `get_best_xy_initial_design()`

Identify and report the best point from initial design.

What happens in `get_best_xy_initial_design()`:

1. Find minimum value in `y_` (or `mean_y` if noise)
2. Store as `best_y_`
3. Store corresponding point as `best_x_`
4. Print progress if verbose

```
opt.get_best_xy_initial_design()

print(f"Best point found: {opt.best_x_}")
print(f"Best value found: {opt.best_y_:.6f}")
print(f"\nOptimum location: [1, 1]")
print(f"Optimum value: 0")
print(f"Current gap: {opt.best_y_ - 0:.6f}")
```

```
Initial best: f(x) = 0.066006
Best point found: [1.21813757 1.49743205]
Best value found: 0.066006
```

```
Optimum location: [1, 1]
Optimum value: 0
Current gap: 0.066006
```

65.11. Illustration of the Initial Design Phase Results


To visualize the results, we generate a contour plot with contour lines of the objective function and mark the best point. The other evaluated points are shown as well:

```
# Create a grid of points for contour plotting
x = np.linspace(-2, 2, 400)
y = np.linspace(-2, 2, 400)
X, Y = np.meshgrid(x, y)
grid_points = np.column_stack([X.ravel(), Y.ravel()])
Z = rosenbrock(grid_points).reshape(X.shape)
plt.figure(figsize=(8, 6))
# Contour plot
```

65. Standard Optimization Workflow

```
contour = plt.contour(X, Y, Z, levels=np.logspace(-0.5, 3.5, 20), cmap='viridis')
plt.clabel(contour, inline=True, fontsize=8)

# Plot evaluated points from the initial design phase:
plt.scatter(opt.X[:, 0], opt.X[:, 1], c='blue', label='Evaluated Points', alpha=0.6)
# Mark best point
plt.scatter(opt.best_x[0], opt.best_x[1], c='red', s=100, label='Best Point', edgecolor='black')
plt.xlabel('X1')
plt.ylabel('X2')
plt.title('Objective Function Contours with Evaluated Points')
plt.legend()
plt.grid(True)
plt.show()
```



66. Sequential Optimization Loop

After initialization, the main optimization loop begins. Each iteration follows these steps:

66.1. Step: Surrogate Model Fitting

66.1.1. Method: `fit_scheduler()`

Fit surrogate model using appropriate data based on noise handling.

*** What happens in `fit_scheduler()`:***

First, the method transforms the input data `X` to the internal scale used by the optimizer. Then, it decides which data to use for fitting the surrogate model based on whether noise handling is enabled:

1. If `noise` is set (i.e., `repeats_surrogate > 1`):
 - Fit surrogate using mean points (`mean_X`, `mean_y`)
2. Else:
 - Fit surrogate using all evaluated points (`X_`, `y_`)

`fit_scheduler()` then calls `fit_surrogate()` with the selected data.

66.1.2. Method: `fit_surrogate()`

Fit a surrogate model (Gaussian Process) to current data.

What happens in `fit_surrogate()`:

1. If `max_surrogate_points` set and exceeded: Select subset of points using the `fit_selection_dispatcher()` method.
2. Fit surrogate model using `surrogate.fit(X, y)`
3. Surrogate learns the function landscape and uncertainty

```
opt.fit_scheduler()

print(f"Surrogate fitted with {len(opt.y_)} points")
print(f"Surrogate type: {type(opt.surrogate).__name__}")
print(f"Kernel: {opt.surrogate.kernel_}")
```

```
Surrogate fitted with 10 points
Surrogate type: GaussianProcessRegressor
Kernel: 1.08**2 * Matern(length_scale=0.436, nu=2.5)
```

66.1.3. Method: `fit_selection_dispatcher()`

What happens in `fit_selection_dispatcher()`:


1. If `max_surrogate_points` is set and exceeded:
 - Select a subset of points from $\mathbf{X}_$ and $\mathbf{y}_$ for fitting the surrogate model.
 - Strategies may include random sampling, clustering, or other heuristics to reduce the dataset size while preserving important information.

Surrogate Model Selection:

- Default: Gaussian Process with Matérn kernel
- Provides: Mean prediction $\mu(x)$ and uncertainty $\sigma(x)$
- Can be replaced with Random Forest, Kriging, etc.
- The fitted surrogate can be plotted with the following code:

```
opt.plot_surrogate()
```

66.1. Step: Surrogate Model Fitting



66.1.4. Step: Apply OCBA

66.1.5. Method: `apply_ocba()`

What happens in `apply_ocba()`:

1. Compute the optimality criteria for each point in the design space.
2. Select the most promising points based on the criteria.
3. Update the surrogate model with the selected points.

```
X_ocba = opt.apply_ocba()
if X_ocba is not None:
    print(f"OCBA selected {X_ocba.shape[0]} points for re-evaluation")
    print(f"OCBA points shape: {X_ocba.shape}")
else:
    print("OCBA not applied (noise=False or ocba_delta=0)")
```

OCBA not applied (noise=False or ocba_delta=0)

66.2. Step: Predict with Uncertainty

66.2.1. Method: `_predict_with_uncertainty()`

Make predictions at new locations with uncertainty estimates.

What happens in `_predict_with_uncertainty()`:

1. Call `surrogate.predict(X, return_std=True)`
2. Returns mean $\mu(x)$ and standard deviation $\sigma(x)$
3. Used by acquisition function to balance exploitation (low μ) and exploration (high σ)

Here we test prediction at a few points:

```
X_test = np.array([[0.5, 0.5], [1.0, 1.0], [-0.5, 0.5]])
mu, sigma = opt._predict_with_uncertainty(X_test)

print(f"Test points: {X_test.shape}")
print(f"\nPredictions ( $\mu \pm \sigma$ ):")
for i, (x, m, s) in enumerate(zip(X_test, mu, sigma)):
    print(f"  x={x} →  $\mu$ ={m:.4f},  $\sigma$ ={s:.4f}")
```

Test points: (3, 2)

Predictions ($\mu \pm \sigma$):

```
x=[0.5 0.5] →  $\mu$ =118.4846,  $\sigma$ =92.5824
x=[1. 1.] →  $\mu$ =123.3180,  $\sigma$ =128.5104
x=[-0.5 0.5] →  $\mu$ =200.6074,  $\sigma$ =35.0723
```

66.3. Step: Next Point Suggestion

66.3.1. Method: `suggest_next_infill_point()`

Optimize acquisition function to find next evaluation point.

What happens in `suggest_next_infill_point()`:

1. Use `differential_evolution` to optimize acquisition function
2. Find point that maximizes EI (or minimizes predicted value for 'y')

3. Apply rounding for integer/factor variables
4. Check distance to existing points (avoid duplicates)
5. If duplicate or too close: Use fallback strategy
6. Return suggested point

```
x_next = opt.suggest_next_infill_point()
print(f"Next point suggested: {x_next}")
```

Next point suggested: [[1.17717813 0.34623968]]

Predict at suggested point:

```
mu_next, sigma_next = opt._predict_with_uncertainty(x_next.reshape(1, -1))
print(f"\nPrediction at suggested point:")
print(f"   $\mu(x_{\text{next}})$  = {mu_next[0]:.4f}")
print(f"   $\Sigma(x_{\text{next}})$  = {sigma_next[0]:.4f}")
```

Prediction at suggested point:

```
 $\mu(x_{\text{next}})$  = -4.4329
 $\Sigma(x_{\text{next}})$  = 26.2575
```

Fallback Strategies (if acquisition optimization fails):

- 'best': Re-evaluate current best point
- 'mean': Select point with best predicted mean
- 'random': Random point in search space

66.4. Step: Acquisition Function Evaluation

66.4.1. Method: `_acquisition_function()`

Compute acquisition function value to guide search. The acquisition function is used as an objective function by the optimizer on the surrogate, e.g., differential evolution. It determines where to sample next.

```
# Evaluate acquisition at test points
print(f"Acquisition type: {opt.acquisition} (Expected Improvement)")
print(f"\nAcquisition values at test points:")

for x in X_test:
    acq_val = opt._acquisition_function(x)
    print(f"  x={x} → acq={acq_val:.6f}")
```

66. Sequential Optimization Loop

Acquisition type: y (Expected Improvement)

Acquisition values at test points:

x=[0.5 0.5] → acq=118.484575

x=[1. 1.] → acq=123.317975

x=[-0.5 0.5] → acq=200.607425

Acquisition function can be used to balance:

- exploitation, i.e., low predicted mean $\mu(x)$ and
- exploration, i.e., high uncertainty $\sigma(x)$

Acquisition Functions Available:

1. **Expected Improvement (EI)** - Default, best balance

$$\text{EI}(x) = (f^* - \mu(x))\Phi(Z) + \sigma(x)\phi(Z)$$

$$\text{where } Z = \frac{f^* - \mu(x)}{\sigma(x)}$$

2. **Probability of Improvement (PI)** - More exploitative

$$\text{PI}(x) = \Phi\left(\frac{f^* - \mu(x)}{\sigma(x)}\right)$$

3. **Mean ('y')** - Pure exploitation

$$\text{acq}(x) = \mu(x)$$

66.5. Step: Update Repeats for Infill Points

66.5.1. Method: `update_repeats_infill_points()`

Repeat the infill point for noisy function evaluation if `repeats_surrogate > 1`.

What happens in `update_repeats_infill_points()`:

1. Takes the suggested next point (`x_next`)
2. If `repeats_surrogate > 1`: Creates multiple copies for repeated evaluation
3. Otherwise: Returns point as 2D array (shape: $1 \times n_{\text{features}}$)
4. Returns array ready for function evaluation

```
x_next_repeated = opt.update_repeats_infill_points(x_next)
print(f"Shape before repeating: {x_next.shape}")
print(f"Shape after repeating: {x_next_repeated.shape}")
print(f"Number of evaluations planned: {x_next_repeated.shape[0]}")
```


Shape before repeating: (1, 2)
 Shape after repeating: (1, 2)
 Number of evaluations planned: 1

66.6. Append OCBA Points to Infill Points

Combines OCBA selected points with the next suggested point for evaluation.

```
if X_ocba is not None:
    x_next_repeated = np.append(X_ocba, x_next_repeated, axis=0)
```

66.7. Step: Evaluation of New Points

66.7.1. Method: `evaluate_function()` (again)

Evaluate the objective function at the suggested points.

```
x_next_2d = x_next.reshape(1, -1)
y_next = opt.evaluate_function(x_next_2d)

print(f"Evaluated point: {x_next}")
print(f"Function value: {y_next[0]:.6f}")
print(f"Current best: {opt.best_y:.6f}")
```

Evaluated point: [[1.17717813 0.34623968]]
 Function value: 0.109753
 Current best: 0.066006

66.8. Step: Handle Failed Evaluations (Sequential)

66.8.1. Method: `_handle_NA_new_points()`

Handle NaN/inf values in new evaluations with penalty approach.

What happens in `_handle_NA_new_points()`:

1. **Apply penalty** to NaN/inf values (unlike initial design)
 - $\text{Penalty} = \max(\text{history}) + 3 \times \text{std}(\text{history}) + \text{random noise}$

66. Sequential Optimization Loop

2. **Remove** remaining invalid values after penalty
3. **Return None** if all evaluations failed → skip iteration
4. **Continue** if any valid evaluations exist

```
x_clean, y_clean = opt._handle_NA_new_points(x_next_2d, y_next)

if x_clean is not None:
    print(f"x Valid evaluations: {len(y_clean)}")
    print(f" All values finite: {np.all(np.isfinite(y_clean))}")
else:
    print(f" All evaluations failed - iteration would be skipped")
```

```
x Valid evaluations: 1
All values finite: True
```

Why penalties in sequential phase?

- Preserves optimization history
- Allows surrogate to learn “bad” regions
- Random noise prevents identical penalties

66.9. Step: Update Success Rate

66.9.1. Method: `update_success_rate(y0)`

Update success rate BEFORE updating storage (initial design - all should be successes since starting from scratch)

What happens in `update_success_rate()`:

1. Calculate success rate as ratio of valid evaluations to total evaluated

```
opt.update_success_rate(y0_clean)
print(f"\nSuccess rate updated: {opt.success_rate:.2%} (valid evaluations / total evals)")
```

```
Success rate updated: 0.00% (valid evaluations / total evaluations)
```

66.10. Step: Update Storage

66.10.1. Internal updates

Add new evaluations to storage.

```
opt.update_storage(x_next_repeated, y_next)
```

66.11. Update Statistics

Update statistics after new evaluations.

66.11.1. What happens in `update_stats()`:

Always updated:

- `min_y`: Best (minimum) function value
- `min_X`: Design point with best value
- `counter`: Total number of evaluations

For noisy functions only (if `noise=True`):

- `mean_X`: Unique design points
- `mean_y`: Mean values per design point
- `var_y`: Variance per design point

```
opt.update_stats()
print(f"Basic statistics:")
print(f"  min_y: {opt.min_y:.6f}")
print(f"  min_X: {opt.min_X}")
print(f"  counter: {opt.counter}")

if opt.repeats_initial > 1 or opt.repeats_surrogate > 1:
    print(f"\nNoise statistics:")
    print(f"  mean_X shape: {opt.mean_X.shape}")
    print(f"  mean_y shape: {opt.mean_y.shape}")
    print(f"  var_y shape: {opt.var_y.shape}")
else:
    print(f"\nNoise handling: disabled (deterministic function)")
```

66. Sequential Optimization Loop

```
Basic statistics:
  min_y: 0.066006
  min_X: [1.21813757 1.49743205]
  counter: 11
```

Noise handling: disabled (deterministic function)

66.12. Step: Update Best Solution

66.12.1. Method: `_update_best_main_loop()`

Update the best solution if improvement found.

```
best_before = opt.best_y_
opt._update_best_main_loop(x_clean, y_clean, start_time=time.time())

print(f"Best before: {best_before:.6f}")
print(f"Best after: {opt.best_y_:.6f}")

if opt.best_y_ < best_before:
    print(f"\nx New best found!")
    print(f"  Location: {opt.best_x_}")
    print(f"  Value: {opt.best_y_:.6f}")
else:
    print(f"\n○ Best unchanged")
```

```
Iter 0 | Best: 0.066006 | Curr: 0.109753 | Rate: 0.00 | Evals: 36.7%
Best before: 0.066006
Best after: 0.066006
```

○ Best unchanged

Loop Termination Conditions:

The optimization continues until:

1. `len(y_) >= max_iter` (reached evaluation budget), OR
2. `elapsed_time >= max_time` (reached time limit)

67. Complete Optimization Example

Now let's run a complete optimization to see all steps in action:

```
# Create fresh optimizer
opt_complete = SpotOptim(
    fun=rosenbrock,
    bounds=[(-2, 2), (-2, 2)],
    n_initial=8,
    max_iter=25,
    acquisition='ei',
    verbose=False, # Set to False for cleaner output
    seed=42
)

# Run optimization
result = opt_complete.optimize()

print(f"\nOptimization Result:")
print(f"{'='*70}")
print(f"Best point found: {result.x}")
print(f"Best value: {result.fun:.6f}")
print(f"True optimum: [1.0, 1.0]")
print(f"True minimum: 0.0")
print(f"Gap to optimum: {result.fun:.6f}")
print(f"\nFunction evaluations: {result.nfev}")
print(f"Sequential iterations: {result.nit}")
print(f"Success: {result.success}")
print(f"Message: {result.message}")
```

Optimization Result:

```
=====
Best point found: [0.06611585 0.01145995]
Best value: 0.877165
True optimum: [1.0, 1.0]
True minimum: 0.0
Gap to optimum: 0.877165
```

67. Complete Optimization Example

```
Function evaluations: 25
Sequential iterations: 17
Success: True
Message: Optimization terminated: maximum evaluations (25) reached
        Current function value: 0.877165
        Iterations: 17
        Function evaluations: 25
```

68. Noisy Functions with Repeats

When dealing with noisy objective functions, SpotOptim can evaluate each point multiple times and track statistics.

68.1. Configuration for Noisy Functions

```
NOISE_STD = 10.0
print("\n" + "="*70)
print("NOISY FUNCTION OPTIMIZATION")
print("="*70)
print(f"\nConfiguration for noisy optimization with noise std = {NOISE_STD}:")

# Wrapper to add noise
def rosenbrock_noisy_wrapper(X):
    return rosenbrock_noisy(X, noise_std=NOISE_STD)

opt_noisy = SpotOptim(
    fun=rosenbrock_noisy_wrapper,
    bounds=[(-2, 2), (-2, 2)],
    n_initial=6,
    max_iter=20,
    repeats_initial=3,      # Evaluate each initial point 3 times
    repeats_surrogate=2,    # Evaluate each sequential point 2 times
    verbose=False,
    seed=42
)

print("Configuration for noisy optimization:")
print(f"  repeats_initial: {opt_noisy.repeats_initial}")
print(f"  repeats_surrogate: {opt_noisy.repeats_surrogate}")
```

=====

NOISY FUNCTION OPTIMIZATION

=====

Configuration for noisy optimization with noise std = 10.0:

Configuration for noisy optimization:

repeats_initial: 3

repeats_surrogate: 2

68.2. Noisy Optimization Workflow Differences

68.2.1. Modified Initial Design

With repeats_initial > 1:

```
result_noisy = opt_noisy.optimize()

print(f"\nInitial design with repeats:")
print(f"  n_initial = {opt_noisy.n_initial}")
print(f"  repeats_initial = {opt_noisy.repeats_initial}")
print(f"  Total initial evaluations: {opt_noisy.n_initial * opt_noisy.repeats_initial}")

print(f"\nStatistics tracked:")
print(f"  mean_X shape: {opt_noisy.mean_X.shape} (unique points)")
print(f"  mean_y shape: {opt_noisy.mean_y.shape} (mean values)")
print(f"  var_y shape: {opt_noisy.var_y.shape} (variances)")

print(f"\nExample statistics for first point:")
idx = 0
print(f"  Point: {opt_noisy.mean_X[idx]}")
print(f"  Mean value: {opt_noisy.mean_y[idx]:.4f}")
print(f"  Variance: {opt_noisy.var_y[idx]:.4f}")
print(f"  Std dev: {np.sqrt(opt_noisy.var_y[idx]):.4f}")
```

Initial design with repeats:

n_initial = 6

repeats_initial = 3

Total initial evaluations: 18

Statistics tracked:

mean_X shape: (7, 2) (unique points)

mean_y shape: (7,) (mean values)


```
var_y shape: (7,) (variances)
```

Example statistics for first point:

```
Point: [-1.94449611 -1.27399111]
```

```
Mean value: 2567.3832
```

```
Variance: 11.5968
```

```
Std dev: 3.4054
```

Key Differences with Noise:

1. **Repeated Evaluations:** Each point evaluated multiple times
2. **Statistics Tracking:**
 - `mean_X`: Unique evaluation locations
 - `mean_y`: Mean function values at each location
 - `var_y`: Variance of function values
 - `n_eval`: Number of evaluations per location
3. **Surrogate Fitting:** Uses `mean_y` instead of `y_`
4. **Best Selection:** Based on `mean_y` not individual `y_`

68.2.2. Update Statistics Method

```
print(f"After optimization:")
print(f"  Total evaluations: {len(opt_noisy.y_)}")
print(f"  Unique points: {len(opt_noisy.mean_X)}")
print(f"  Average repeats per point: {len(opt_noisy.y_) / len(opt_noisy.mean_X):.2f}")

print(f"\nVariance statistics:")
print(f"  Mean variance: {np.mean(opt_noisy.var_y):.6f}")
print(f"  Max variance: {np.max(opt_noisy.var_y):.6f}")
print(f"  Min variance: {np.min(opt_noisy.var_y):.6f}")
```

After optimization:

```
Total evaluations: 20
```

```
Unique points: 7
```

```
Average repeats per point: 2.86
```

Variance statistics:

```
Mean variance: 43.630970
```

```
Max variance: 94.990756
```

```
Min variance: 6.357449
```


69. Optimal Computing Budget Allocation (OCBA)

OCBA intelligently allocates additional evaluations to distinguish between competing solutions.

69.1. OCBA Configuration

```
print("\n" + "="*70)
print("OPTIMAL COMPUTING BUDGET ALLOCATION (OCBA)")
print("="*70)

opt_ocba = SpotOptim(
    fun=rosenbrock_noisy_wrapper,
    bounds=[(-2, 2), (-2, 2)],
    n_initial=8,
    max_iter=30,
    repeats_initial=2,
    repeats_surrogate=1,
    ocba_delta=3,          # Allocate 3 additional evaluations per iteration
    verbose=False,
    seed=42
)

print("OCBA Configuration:")
print(f"  ocba_delta: {opt_ocba.ocba_delta}")
print(f"  Purpose: Intelligently re-evaluate existing points")
print(f"  Benefit: Better distinguish between similar solutions")
```

```
=====
OPTIMAL COMPUTING BUDGET ALLOCATION (OCBA)
=====
OCBA Configuration:
```

69. Optimal Computing Budget Allocation (OCBA)

ocba_delta: 3
Purpose: Intelligently re-evaluate existing points
Benefit: Better distinguish between similar solutions

69.2. OCBA Method

69.2.1. Method: apply_ocba()

```
result_ocba = opt_ocba.optimize()

print(f"\nOCBA applied during optimization")
print(f"  Final total evaluations: {result_ocba.nfev}")
print(f"  Expected without OCBA: ~{opt_ocba.n_initial * opt_ocba.repeats_initial + res")
print(f"  Additional OCBA evaluations: ~{result_ocba.nit * opt_ocba.ocba_delta}")

print(f"\nOCBA intelligently allocated extra evaluations to:")
print(f"  - Current best candidate (confirm it's truly best)")
print(f"  - Close competitors (distinguish between similar solutions)")
print(f"  - High-variance points (reduce uncertainty)")
```

OCBA applied during optimization
Final total evaluations: 30
Expected without OCBA: ~27
Additional OCBA evaluations: ~33

OCBA intelligently allocated extra evaluations to:

- Current best candidate (confirm it's truly best)
- Close competitors (distinguish between similar solutions)
- High-variance points (reduce uncertainty)

OCBA Algorithm:

1. Identify best solution (lowest mean value)
2. Calculate allocation ratios based on:
 - Distance to best solution
 - Variance of each solution
3. Allocate `ocba_delta` additional evaluations
4. Returns points to re-evaluate
5. Requires: ≥ 3 points with variance > 0

OCBA Activation Conditions: - `noise = True` (repeats > 1) - `ocba_delta > 0` -
At least 3 design points exist - All points have variance > 0

70. Handling Function Evaluation Failures

SpotOptim robustly handles functions that occasionally fail (return NaN/inf).

70.1. Example with Failures

```
opt_failures = SpotOptim(  
    fun=rosenbrock_with_failures,  
    bounds=[(-2, 2), (-2, 2)],  
    n_initial=12,  
    max_iter=35,  
    verbose=False,  
    seed=42  
)  
  
result_failures = opt_failures.optimize()  
  
print(f"\nOptimization with ~15% random failure rate:")  
print(f"  Function evaluations: {result_failures.nfev}")  
print(f"  Sequential iterations: {result_failures.nit}")  
print(f"  Success: {result_failures.success}")  
  
# Count how many values are non-finite in raw evaluations  
n_total = len(opt_failures.y_)  
n_finite = np.sum(np.isfinite(opt_failures.y_))  
print(f"\nEvaluation statistics:")  
print(f"  Total evaluations: {n_total}")  
print(f"  Finite values: {n_finite}")  
print(f"  Note: Failed evaluations handled transparently")
```

```
Optimization with ~15% random failure rate:  
  Function evaluations: 35
```

70. Handling Function Evaluation Failures

Sequential iterations: 30
Success: True

Evaluation statistics:
Total evaluations: 35
Finite values: 35
Note: Failed evaluations handled transparently

70.2. Failure Handling in Initial Design

70.2.1. Method: `rm_initial_design_NA_values()`

```
print("Initial design failure handling:")
print("  1. Identify NaN/inf values")
print("  2. Remove invalid points entirely")
print("  3. Continue with valid points only")
print("  4. No penalties applied")
print("  5. Require at least 1 valid point")

print("\nRationale:")
print("  - Initial design should be clean")
print("  - Invalid regions identified naturally")
print("  - Surrogate trained on good data only")
```

Initial design failure handling:

1. Identify NaN/inf values
2. Remove invalid points entirely
3. Continue with valid points only
4. No penalties applied
5. Require at least 1 valid point

Rationale:

- Initial design should be clean
- Invalid regions identified naturally
- Surrogate trained on good data only

70.3. Failure Handling in Sequential Phase

70.3.1. Method: `_handle_NA_new_points()`

```
print("Sequential phase failure handling:")
print("  1. Apply penalty to NaN/inf values")
print("    - Penalty = max(history) + 3*std(history)")
print("    - Add random noise to avoid duplicates")
print("  2. Remove remaining invalid values")
print("  3. Skip iteration if all evaluations failed")
print("  4. Continue if any valid evaluations")

print("\nPenalty approach benefits:")
print("  x Preserves optimization history")
print("  x Surrogate learns to avoid bad regions")
print("  x Better exploration-exploitation balance")
print("  x More robust convergence")
```

Sequential phase failure handling:

1. Apply penalty to NaN/inf values
 - Penalty = max(history) + 3*std(history)
 - Add random noise to avoid duplicates
2. Remove remaining invalid values
3. Skip iteration if all evaluations failed
4. Continue if any valid evaluations

Penalty approach benefits:

- x Preserves optimization history
- x Surrogate learns to avoid bad regions
- x Better exploration-exploitation balance
- x More robust convergence

70.4. Penalty Application

70.4.1. Method: `apply_penalty_NA()`

Let's demonstrate penalty calculation:

70. Handling Function Evaluation Failures

```
# Simulate historical values
y_history_sim = np.array([10.0, 15.0, 8.0, 12.0, 20.0, 9.0])
y_new_sim = np.array([7.0, np.nan, 11.0, np.inf])

print("Historical values:", y_history_sim)
print("New evaluations:", y_new_sim)

# Apply penalty
y_repaired = opt_failures.apply_penalty_NA(
    y_new_sim,
    y_history=y_history_sim,
    penalty_value=None, # Compute adaptively
    sd=0.1
)

print(f"\nAfter penalty application:", y_repaired)
print(f"All finite: {np.all(np.isfinite(y_repaired))}")

# Show penalty calculation
max_hist = np.max(y_history_sim)
std_hist = np.std(y_history_sim, ddof=1)
penalty_base = max_hist + 3 * std_hist

print(f"\nPenalty calculation:")
print(f"  max(history) = {max_hist:.2f}")
print(f"  std(history) = {std_hist:.2f}")
print(f"  Base penalty = {max_hist:.2f} + 3×{std_hist:.2f} = {penalty_base:.2f}")
print(f"  Actual penalty = {penalty_base:.2f} + noise")
```

Historical values: [10. 15. 8. 12. 20. 9.]

New evaluations: [7. nan 11. inf]

After penalty application: [7. 33.50646145 11. 33.29891622]

All finite: True

Penalty calculation:

max(history) = 20.00

std(history) = 4.50

Base penalty = 20.00 + 3×4.50 = 33.51

Actual penalty = 33.51 + noise

71. Complete Method Summary

71.1. Methods Called During `optimize()`

The `optimize()` method orchestrates the full optimization process with restart support. It delegates to `execute_optimization_run()`, which runs the sequential optimization loop in `optimize_sequential_run()`.

71.1.1. Preparation Phase (inside `optimize_sequential_run()` via `_initialize_run()`)

1. `set_seed()` - Set random seed for reproducibility
2. `get_initial_design()` - Generate/process initial sample points
3. `curate_initial_design()` - Remove duplicates, handle repeats
4. `evaluate_function()` - Evaluate objective function
5. `rm_initial_design_NA_values()` - Remove NaN/inf from initial design
6. `check_size_initial_design()` - Validate sufficient points
7. `init_storage()` - Initialize storage (`X_`, `y_`, `n_iter_`)
8. `update_stats()` - Compute mean/variance for noisy functions
9. `_init_tensorboard()` - Log initial design to TensorBoard (if enabled)
10. `get_best_xy_initial_design()` - Identify initial best

71.1.2. Sequential Optimization Loop (each iteration inside `_run_sequential_loop()`)

11. `fit_scheduler()` - Select data and fit surrogate based on noise handling
 - Internally calls `transform_X()` - Transform to internal scale
 - Internally calls `fit_surrogate()` - Fit Gaussian Process to data
 - Internally calls `fit_selection_dispatcher()` - Subset selection for large datasets
12. `apply_ocba()` - OCBA allocation (if enabled)
13. `suggest_next_infill_point()` - Optimize acquisition function
 - Internally calls `_acquisition_function()`
 - Internally calls `_predict_with_uncertainty()`

71. Complete Method Summary

14. `update_repeats_infill_points()` - Repeat suggested point for noisy functions
15. `evaluate_function()` - Evaluate at suggested point(s)
16. `_handle_NA_new_points()` - Handle failures with penalties
 - Internally calls `apply_penalty_NA()`
 - Internally calls `remove_nan()`
17. `update_success_rate()` - Update success tracking
18. Check for restart (if `success_rate == 0` for `restart_after_n` consecutive iterations)
19. `update_storage()` - Append new evaluations to storage
 - Internally calls `inverse_transform_X()` - Convert back to original scale
20. `update_stats()` - Update mean/variance statistics
21. `_write_tensorboard_hparams()` - Log hyperparameters (if enabled)
22. `_write_tensorboard_scalars()` - Log scalar metrics (if enabled)
23. `_update_best_main_loop()` - Update best solution

71.1.3. Finalization

24. `to_all_dim()` - Expand to full dimensions (if dimensionality reduction used)
25. `determine_termination()` - Determine termination reason
26. `_close_tensorboard_writer()` - Close logging (if enabled)
27. `map_to_factor_values()` - Convert factors back to strings
28. Return `OptimizeResult` object

71.2. Helper Methods Used

- `generate_initial_design()` - LHS generation
- `repair_non_numeric()` - Round integer/factor variables
- `_select_new()` - Check for duplicate points
- `_handle_acquisition_failure()` - Fallback strategies
- `to_red_dim()` - Dimension reduction (if enabled)
- `fit_selection_dispatcher()` - Subset selection for large datasets
- `fit_select_distant_points()` - K-means based distant point selection
- `fit_select_best_cluster()` - Cluster-based best point selection

72. Termination Conditions

72.1. Method: `determine_termination()`

```
print("\n" + "="*70)
print("TERMINATION CONDITIONS")
print("="*70)
print("\nMethod: determine_termination()")
print("-" * 70)

print("Optimization terminates when:")
print("  1. len(y_) >= max_iter (evaluation budget exhausted)")
print("  2. elapsed_time >= max_time (time limit reached)")
print("  3. Whichever comes first")

print(f"\nExample from previous run:")
print(f"  Message: {result_failures.message}")
print(f"  Evaluations: {result_failures.nfev}/{opt_failures.max_iter}")
print(f"  Iterations: {result_failures.nit}")
```

```
=====
TERMINATION CONDITIONS
=====

Method: determine_termination()
-----

Optimization terminates when:
  1. len(y_) >= max_iter (evaluation budget exhausted)
  2. elapsed_time >= max_time (time limit reached)
  3. Whichever comes first

Example from previous run:
  Message: Optimization terminated: maximum evaluations (35) reached
  Current function value: 0.001758
  Iterations: 30
```

72. Termination Conditions

Function evaluations: 35
Evaluations: 35/35
Iterations: 30

73. Performance Comparison on Noisy Functions

Let's compare standard, noisy, and noisy+OCBA optimization:

```
print("\n" + "="*70)
print("PERFORMANCE COMPARISON")
print("="*70)

# Standard optimization
opt_std = SpotOptim(
    fun=rosenbrock_noisy_wrapper,
    bounds=[(-2, 2), (-2, 2)],
    n_initial=8,
    max_iter=50,
    repeats_initial=1,
    repeats_surrogate=1,
    verbose=False,
    seed=10
)
result_std = opt_std.optimize()

# With repeats (no OCBA)
opt_rep = SpotOptim(
    fun=rosenbrock_noisy_wrapper,
    bounds=[(-2, 2), (-2, 2)],
    n_initial=8,
    max_iter=50,
    repeats_initial=2,
    repeats_surrogate=1,
    ocba_delta=0,
    verbose=False,
    seed=10
)
result_rep = opt_rep.optimize()

# With repeats + OCBA
```

73. Performance Comparison on Noisy Functions

```
opt_ocba_comp = SpotOptim(
    fun=rosenbrock_noisy_wrapper,
    bounds=[(-2, 2), (-2, 2)],
    n_initial=8,
    max_iter=50,
    repeats_initial=1,
    repeats_surrogate=1,
    ocba_delta=1,
    verbose=False,
    seed=10
)
result_ocba_comp = opt_ocba_comp.optimize()

# Evaluate the found optima with the TRUE (noise-free) Rosenbrock function
# to get the actual quality of the solutions
print(f"\n" + "="*70)
print("TRUE FUNCTION VALUE EVALUATION")
print("="*70)
print("\nRe-evaluating found optima with noise-free Rosenbrock function:")
print("-" * 70)

# Evaluate each solution with the TRUE function (no noise)
true_val_std = rosenbrock(result_std.x.reshape(1, -1))[0]
true_val_rep = rosenbrock(result_rep.x.reshape(1, -1))[0]
true_val_ocba = rosenbrock(result_ocba_comp.x.reshape(1, -1))[0]

print(f"\nStandard (no repeats):")
print(f"  Found x: {result_std.x}")
print(f"  Noisy value (from optimization): {result_std.fun:.6f}")
print(f"  TRUE value (noise-free): {true_val_std:.6f}")

print(f"\nWith repeats (no OCBA):")
print(f"  Found x: {result_rep.x}")
print(f"  Noisy value (from optimization): {result_rep.fun:.6f}")
print(f"  TRUE value (noise-free): {true_val_rep:.6f}")

print(f"\nWith repeats + OCBA:")
print(f"  Found x: {result_ocba_comp.x}")
print(f"  Noisy value (from optimization): {result_ocba_comp.fun:.6f}")
print(f"  TRUE value (noise-free): {true_val_ocba:.6f}")

print(f"\n" + "="*70)
print("PERFORMANCE COMPARISON SUMMARY")
print("="*70)
```



```

print(f"\nComparison on noisy Rosenbrock ( $\sigma$ ={NOISE_STD}):")
print(f"{'Method':<30} {'Noisy Value':<15} {'TRUE Value':<15} {'Evaluations':<15}")
print(f"{'-'*75}")
print(f"{'Standard (no repeats)':<30} {result_std.fun:<15.6f} {true_val_std:<15.6f} {result_std.nfev:<15}")
print(f"{'With repeats (no OCBA)':<30} {result_rep.fun:<15.6f} {true_val_rep:<15.6f} {result_rep.nfev:<15}")
print(f"{'With repeats + OCBA':<30} {result_ocba_comp.fun:<15.6f} {true_val_ocba:<15.6f} {result_ocba_comp.nfev:<15}")
print(f"{'-'*75}")
print(f"{'True optimum':<30} {'0.0':<15} {'0.0':<15}")

print("\nKey observations:")
print(" • Noisy values (from optimization) are misleading due to noise")
print(" • TRUE values show actual solution quality")
print(" • Repeats help find better solutions despite noise")
print(" • OCBA focuses evaluations on distinguishing best solutions")
print(" • More evaluations generally lead to better TRUE performance")

```

PERFORMANCE COMPARISON

TRUE FUNCTION VALUE EVALUATION

Re-evaluating found optima with noise-free Rosenbrock function:

```

Standard (no repeats):
  Found x: [-0.83535584  0.47919121]
  Noisy value (from optimization): -11.628924
  TRUE value (noise-free): 8.148358

With repeats (no OCBA):
  Found x: [1.09501145  1.41292401]
  Noisy value (from optimization): -15.194050
  TRUE value (noise-free): 4.583233

With repeats + OCBA:
  Found x: [-0.83535584  0.47919121]
  Noisy value (from optimization): -11.628924
  TRUE value (noise-free): 8.148358

```

73. Performance Comparison on Noisy Functions

PERFORMANCE COMPARISON SUMMARY

=====

Comparison on noisy Rosenbrock ($\sigma=10.0$):

Method	Noisy Value	TRUE Value	Evaluations

Standard (no repeats)	-11.628924	8.148358	50
With repeats (no OCBA)	-15.194050	4.583233	50
With repeats + OCBA	-11.628924	8.148358	50

True optimum	0.0	0.0	

Key observations:

- Noisy values (from optimization) are misleading due to noise
- TRUE values show actual solution quality
- Repeats help find better solutions despite noise
- OCBA focuses evaluations on distinguishing best solutions
- More evaluations generally lead to better TRUE performance

74. Summary

74.1. Complete Workflow Diagram

```

XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
X                               SPOTOPTIM WORKFLOW                               X
X  optimize() → execute_optimization_run()                                     X
X    → optimize_sequential_run() → _run_sequential_loop()                     X
XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX

INITIALIZATION PHASE (optimize_sequential_run via _initialize_run)
X
X  XXX set_seed()
X    XXX Set random seed for reproducibility
X
X  XXX get_initial_design()
X    XXX Generate LHS or process user design
X
X  XXX curate_initial_design()
X    XXX Remove duplicates, add repeats
X
X  XXX evaluate_function()
X    XXX Evaluate objective function
X
X  XXX rm_initial_design_NA_values()
X    XXX Remove NaN/inf points
X
X  XXX check_size_initial_design()
X    XXX Validate sufficient points
X
X  XXX init_storage()
X    XXX Initialize X_, y_, n_iter_
X
X  XXX update_stats()
X    XXX Compute mean/variance (if noise)
X
X  XXX _init_tensorboard() [if enabled]
X    XXX Log initial design to TensorBoard

```

74. Summary

```

    fit_scheduler()
    get_best_xy_initial_design()
        Identify initial best

SEQUENTIAL OPTIMIZATION LOOP (_run_sequential_loop, until max_iter or max_time)
    fit_scheduler()
    transform_X() - Transform to internal scale
    fit_surrogate() - Fit GP to current data
    fit_selection_dispatcher() - Subset selection if needed
    apply_ocba() [if enabled]
    Allocate additional evaluations
    suggest_next_infill_point()
    _acquisition_function()
    _predict_with_uncertainty()
    Optimize to find next point
    update_repeats_infill_points()
    Repeat point if repeats_surrogate > 1
    evaluate_function()
    Evaluate at suggested point(s)
    _handle_NA_new_points()
    apply_penalty_NA()
    remove_nan()
    update_success_rate()
    Track evaluation success
    [Check for restart if success_rate == 0]
    Return "RESTART" if threshold reached
    update_storage()
    Append new evaluations (with scale conversion)
    update_stats()
    Update mean/variance
    _write_tensorboard_hparams() [if enabled]
    Log hyperparameters
```

```

    _write_tensorboard_scalars() [if enabled]
    _ Log scalar metrics
    _
    _update_best_main_loop()
    _ Update best if improved

```

FINALIZATION

```

    _
    _to_all_dim() [if dimensionality reduction used]
    _ Expand to full dimensions
    _
    _determine_termination()
    _ Set termination message
    _
    _close_tensorboard_writer() [if enabled]
    _ Close TensorBoard logging
    _
    _map_to_factor_values() [if factors used]
    _ Convert factors back to strings
    _
    _ Return OptimizeResult

```

74.2. Key Concepts

74.2.1. 1. Initial Design

- **Latin Hypercube Sampling:** Space-filling design for efficient exploration
- **Curation:** Remove duplicates from integer/factor rounding
- **Failure Handling:** Remove invalid points, no penalties

74.2.2. 2. Surrogate Model

- **Default:** Gaussian Process with Matérn kernel
- **Provides:** Mean $\mu(x)$ and uncertainty $\sigma(x)$ predictions
- **Purpose:** Learn function landscape with limited evaluations

74.2.3. 3. Acquisition Function

- **EI:** Expected Improvement (default) - best balance
- **PI:** Probability of Improvement - more exploitative

74. Summary

- **Mean:** Pure exploitation of surrogate predictions

74.2.4. 4. Noise Handling

- **Repeats:** Evaluate each point multiple times
- **Statistics:** Track mean and variance
- **OCBA:** Intelligently allocate additional evaluations

74.2.5. 5. Failure Handling

- **Initial Phase:** Remove invalid points
- **Sequential Phase:** Apply adaptive penalties with noise
- **Robustness:** Continue optimization despite failures

74.3. Best Practices

1. For deterministic functions:

- Use default settings (no repeats)
- Acquisition = 'ei'
- Focus on n_initial and max_iter

2. For noisy functions:

- Set repeats_initial ≥ 2
- Set repeats_surrogate ≥ 1
- Consider OCBA with ocba_delta ≥ 2

3. For unreliable functions:

- SpotOptim handles failures automatically
- No special configuration needed
- Penalties guide search away from bad regions

4. For expensive functions:

- Increase n_initial (better initial model)
- Use 'ei' acquisition (best sample efficiency)
- Consider max_time limit

74.4. Conclusion

SpotOptim provides a robust and flexible framework for surrogate-based optimization with:

- x Efficient space-filling initial designs (LHS)
- x Powerful Gaussian Process surrogate models
- x Smart acquisition functions (EI, PI, Mean)
- x Automatic noise handling with statistics
- x Intelligent budget allocation (OCBA)
- x Robust failure handling
- x Comprehensive progress tracking

The modular design allows easy customization while maintaining robust defaults for most use cases.

74.5. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

75. SpotOptim Execution Flow

76. SpotOptim Execution Flow

76.1. Example Without Restarts

This document outlines the sequence of methods called when `SpotOptim.optimize()` is executed, explaining the input, processing, and output at each step. As an example, we will use the following code:

```
from spotoptim import SpotOptim
from spotoptim.function import sphere

opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    n_initial=5,
    max_iter=20,
    verbose=True
)
result = opt.optimize()
```

76.1.1. SpotOptim.__init__

- **Input:** Configuration parameters (`fun`, `bounds`, `n_initial`, `max_iter`, `surrogate`, etc.).
- **Processing:**
 - Validates arguments (e.g., checks `max_iter >= n_initial`).
 - Initializes `SpotOptimConfig` (configuration) and `SpotOptimState` (mutable state).
 - Sets random seeds via `set_seed`.
 - Determines variable types and transformations.
 - Initializes the default Gaussian Process surrogate if none is provided.
 - Sets up the Latin Hypercube Sampler (`lhs_sampler`).
- **Output:** An initialized `SpotOptim` instance ready for optimization.

76.1.2. SpotOptim.optimize

- **Input:** Optional initial point `X0` (Natural Space).
- **Processing:**
 - Initializes the restart results list.
 - Enters the restart loop.
 - Calls `_optimize_single_run` to perform the optimization.
 - Manages restarts based on the return status of the single run.
- **Output:** `OptimizeResult` object containing the best parameters `x`, best value `fun`, and optimization history.

76.1.3. SpotOptim._optimize_single_run

- **Input:** `timeout_start`, `X0` (initial points), `y0_known` (for injected restarts).
- **Processing:**
 - Orchestrates the lifecycle of a single optimization run (initial design -> loop -> termination).
 - Calls sub-methods for generating design, evaluating points, fitting models, and selecting candidates.
- **Output:** Tuple (`status`, `OptimizeResult`), where `status` is “FINISHED” or “RESTART”.

76.1.4. SpotOptim.get_initial_design

- **Input:** `X0` (from `optimize` call).
- **Processing:**
 - Generates a space-filling design using `lhs_sampler.random()`.
 - Scales the design from `[0, 1]` to the problem bounds.
 - Combines the generated design with any user-provided `X0`.
- **Output:** Array `X` of initial design points in Natural Space.

76.1.5. SpotOptim.curate_initial_design

- **Input:** Raw initial design points `X`.
- **Processing:**
 - Removes duplicate points to avoid redundant evaluations.
 - Verification against existing points (relevant during restarts).
- **Output:** Unique array of initial design points `X` to be evaluated.

76.1.6. SpotOptim.evaluate_function (Initial Design)

- **Input:** Batch of design points `X`.
- **Processing:**
 - Iterates through points and calls the user-provided objective function `fun(x)`.
 - Handles `NaN/Inf` results via penalty mechanism (if enabled).
 - Converts multi-objective results to single-objective if needed (`fun_mo2so`).
- **Output:** Array of objective values `y`.

76.1.7. SpotOptim.init_storage

- **Input:** Initial `X` and `y`.
- **Processing:**
 - Populates `self.state.X_` and `self.state.y_` with the initial design data.
 - Initializes statistics tracking (mean/variance) for noisy problems.
- **Output:** Updated internal state with initial data.

76.1.8. Optimization Loop Start

The process enters a `while` loop that continues until `max_iter` evaluations are reached or time runs out.

76.1.9. SpotOptim._fit_scheduler

- **Input:** Current history `X_`, `y_`.
- **Processing:**
 - Determines if the surrogate model needs retraining (based on iteration intervals).
 - Calls `self.surrogate.fit(X, y)` using the current available data.
- **Output:** A trained surrogate model capable of predicting `y` for unseen `X`.

76.1.10. SpotOptim.suggest_next_infill_point

- **Input:** None (implicitly uses trained surrogate).
- **Processing:**
 - Dispatcher that determines how to find the next point.
 - Calls `_try_optimizer_candidates`.
- **Output:** New candidate point `x_next` (in Internal/Transformed Space).

76.1.11. `SpotOptim._try_optimizer_candidates`

- **Input:** None.
- **Processing:**
 - Calls `optimize_acquisition_func` to maximize the acquisition function (e.g., Expected Improvement).
 - Rounds/repairs the candidate based on variable types (e.g., integers).
 - Checks that `x_next` is not too close to existing points (`tolerance_x`).
- **Output:** Valid, unique candidate point(s).

76.1.12. `SpotOptim.optimize_acquisition_func`

- **Input:** Surrogate model, Acquisition function type (configured in `__init__`).
- **Processing:**
 - Uses a global optimizer (e.g., `differential_evolution`) to find inputs `x` that maximize the acquisition function derived from the surrogate's predictions.
- **Output:** The point(s) `x` that maximize potential improvement.

76.1.13. `SpotOptim.evaluate_function (Loop)`

- **Input:** Candidate point `x_next`.
- **Processing:**
 - Evaluates the user's objective function at the new candidate point.
- **Output:** New objective value `y_next`.

76.1.14. `SpotOptim.update_storage`

- **Input:** New `x_next`, `y_next`.
- **Processing:**
 - Appends the new evaluation to the history `X_` and `y_`.
- **Output:** Updated dataset used for the next surrogate fit.

76.1.15. SpotOptim._update_best_main_loop

- **Input:** New `y_next`.
- **Processing:**
 - Compares new value against `self.state.best_y_`.
 - Updates `best_x_` and `best_y_` if a better solution is found.
- **Output:** Updated best known solution.

76.1.16. Termination

- **Condition:** Validates if `len(y_) >= max_iter`.
- **Processing:**
 - Exits the loop.
 - Prepares the final `OptimizeResult`.
- **Output:** Final optimization result returned to the user.

76.2. Example With Restarts

Now we can look at an example with restarts. Since the argument `restart_after_n` is set to 10, the optimizer will restart if the success rate remains zero for 10 consecutive iterations.

```
opt = SpotOptim(
    fun=sphere,
    bounds=[(-5, 5), (-5, 5)],
    n_initial=5,
    max_iter=100,
    verbose=True,
    restart_after_n=10
)

result = opt.optimize()
```

76.2.1. SpotOptim.optimize (Global Loop)

- **Input:** Global configuration (`max_iter`, `restart_after_n`).
- **Processing:**
 - Initializes `restarts_results_` to store results from each run.

76. SpotOptim Execution Flow

- Enters a **while** loop that manages the global budget.
- Checks if enough budget remains for a new run (must allow at least `n_initial` evaluations).
- **Output:** Orchestrates the sequence of optimization runs.

76.2.2. SpotOptim._optimize_single_run (First Run)

- **Input:** Initial random state.
- **Processing:**
 - Performs standard optimization (Initial Design -> Loop).
 - **Restart Check:** During the loop, monitors the `success_rate`.
 - If `success_rate` stays at 0.0 for `restart_after_n` (10) consecutive iterations, the run is terminated early with `status="RESTART"`.
- **Output:** Partial OptimizeResult and status RESTART.

76.2.3. Restart Injection (Between Runs)

- **Input:** Results from previous runs (`restarts_results_`).
- **Processing:**
 - Identifies the best solution (`x_best`, `y_best`) found so far.
 - **Injection:** If `restart_inject_best=True` (default), sets `self.x0 = x_best`. This ensures the next run starts with the best known solution.
 - Prepares `y0_known = y_best` to skip re-evaluating this expensive point.

76.2.4. SpotOptim._optimize_single_run (Subsequent Runs)

- **Input:** `y0_known` (Value of injected best point).
- **Processing:**
 - **Fresh Initial Design:**
 - * Generates a *new* Latin Hypercube Design (continuing the sequence, providing fresh random points).
 - * Injects the previous best point (`self.x0`) into this design, replacing one of the random points.
 - **Evaluation:**
 - * Evaluates the new LHS points.
 - * Skips evaluation for the injected point, using `y0_known`.
 - **Optimization Loop:**
 - * Fits a new Surrogate Model on this fresh dataset.
 - * Optimizes to refine the best solution or explore new regions.
- **Output:** New OptimizeResult for this specific run.

76.2.5. Termination

- **Condition:** The process repeats until the total number of evaluations (summing all runs) reaches `max_iter` (100).
- **Output:** The `OptimizeResult` with the best objective value found across all runs is returned.

76.3. Description Based on Chunks

The source code is split into chunks of functionality.

76.3.1. Configuration & Helpers

Methods for initializing the optimizer, handling configuration, and processing variable types.

76.3.1.1. `set_seed`

Sets global random seeds for reproducibility across `random`, `numpy`, and `torch`.

```
spot = SpotOptim(fun=lambda x: x, bounds=[(0, 1)], seed=42)
spot.set_seed()
```

76.3.1.2. `detect_var_type`

Auto-detects variable types ('factor' or 'float') based on the provided bounds.

```
spot = SpotOptim(fun=lambda x: x, bounds=[('a', 'b'), (0, 1)])
types = spot.detect_var_type() # ['factor', 'float']
```

76.3.1.3. `modify_bounds_based_on_var_type`

Adjusts bounds representation based on variable types (e.g., converting to integer bounds).

```
spot = SpotOptim(fun=lambda x: x, bounds=[(0.5, 10.5)], var_type=['int'])
spot.modify_bounds_based_on_var_type()
print(spot.bounds) # [(1, 10)]
```

76. *SpotOptim Execution Flow*

76.3.1.4. `process_factor_bounds`

Processes factor variable bounds (tuples of strings) into integer indices and stores the mapping.

```
spot = SpotOptim(fun=lambda x: x, bounds=[('red', 'green')], var_name=['color'])
spot.process_factor_bounds()
# Bounds become [(0, 1)] and mapping is stored
```

76.3.1.5. `repair_non_numeric`

Rounds values for discrete variables (int, factor) to the nearest integer.

```
X = np.array([[1.2, 2.8]])
X_repaired = spot.repair_non_numeric(X, var_type=['int', 'int'])
# [[1.0, 3.0]]
```

76.3.2. **Dimension Reduction**

Methods for handling fixed dimensions (where lower bound equals upper bound).

76.3.2.1. `setup_dimension_reduction`

Identifies fixed dimensions and reduces the optimization search space to only active variables.

```
spot = SpotOptim(fun=lambda x: x, bounds=[(1, 1), (0, 10)])
spot.setup_dimension_reduction()
# Dimension 0 is fixed, optimization happens only on Dimension 1
```

76.3.2.2. `to_red_dim`

Transforms vectors from the full original space to the reduced optimization space.

```
X_full = np.array([[1, 5]]) # D=2
X_red = spot.to_red_dim(X_full) # [[5]] (D=1)
```

76.3.2.3. to_all_dim

Transforms vectors from the reduced optimization space back to the full original space by re-inserting fixed values.

```
X_red = np.array([[5]])
X_full = spot.to_all_dim(X_red) # [[1, 5]]
```

76.3.3. Variable Transformation

Methods for scaling and transforming variables for the surrogate model.

76.3.3.1. transform_value / inverse_transform_value

Applies or reverses transformations (e.g., log10, ln) on single values.

```
val = spot.transform_value(100, "log10") # 2.0
orig = spot.inverse_transform_value(2.0, "log10") # 100.0
```

76.3.3.2. transform_X / inverse_transform_X

Applies or reverses transformations on valid X arrays (batch processing).

```
X = np.array([[10, 100]])
# Assuming log10 transform for both
X_trans = spot.transform_X(X) # [[1.0, 2.0]]
```

76.3.4. Initial Design

Methods for generating the initial population of points.

76.3.4.1. get_initial_design

Generates or retrieves the initial set of points to evaluate. Supports optional user-provided starting points (X0).

```
X_init = spot.get_initial_design()
```

76. SpotOptim Execution Flow

76.3.4.2. generate_initial_design

Internal method to generate Latin Hypercube Sampling (LHS) design.

```
X_lhs = spot.generate_initial_design()
```

76.3.4.3. curate_initial_design

Combines generated design with user-provided x0 and ensures valid bounds.

```
X_curated = spot.curate_initial_design(X_lhs)
```

76.3.5. Surrogate & Acquisition

Methods for fitting the model and selecting new points.

76.3.5.1. _fit_surrogate

Fits the Gaussian Process surrogate model to the observed data.

```
spot._fit_surrogate(X_train, y_train)
```

76.3.5.2. _predict_with_uncertainty

Predicts mean and standard deviation for new points using the surrogate.

```
mean, std = spot._predict_with_uncertainty(X_new)
```

76.3.5.3. suggest_next_infill_point

Determines the next point to evaluate by optimizing the acquisition function.

```
X_next = spot.suggest_next_infill_point()
```

76.3.6. Optimization Loop

Core logic for the optimization process.

76.3.6.1. `optimize`

Main entry point. Runs the full optimization loop: initial design, surrogate fitting, and sequential infill.

```
result = spot.optimize()
```

76.3.6.2. `_optimize_single_run`

Internal method to run the optimization loop (called by `optimize`).

```
spot._optimize_single_run()
```

76.3.7. Results & Analysis

Methods for saving, loading, and analyzing results.

76.3.7.1. `save_result` / `load_result`

Saves or loads the complete optimizer state, including results.

```
spot.save_result(prefix="run1")
spot2 = SpotOptim.load_result("run1_res.pkl")
```

76.3.7.2. `save_experiment` / `load_experiment`

Saves or loads only the configuration (no results), useful for distributing experiments.

```
spot.save_experiment(prefix="exp1")
```

76.3.7.3. `get_results_table`

Displays relevant statistics about the tuned parameters, including importance if available.

```
spot.get_results_table()
```

76.3.8. Plotting

Visualizations for the optimization process.

76.3.8.1. `plot_progress`

Plots the optimization history (objective value vs. iteration).

```
spot.plot_progress()
```

76.3.8.2. `plot_surrogate`

Plots the surrogate landscape for two dimensions.

```
spot.plot_surrogate(i=0, j=1)
```

76.3.8.3. `plot_important_hyperparameter_contour`

Automatically plots landscapes for the most important parameters.

```
spot.plot_important_hyperparameter_contour()
```

76.3.9. TensorBoard Integration

Methods for logging training progress to TensorBoard.

76.3.9.1. `_init_tensorboard`

Sets up the TensorBoard logging directory and writer.

```
spot = SpotOptim(..., tensorboard_log=True)
spot._init_tensorboard()
```

76.3.9.2. `_write_tensorboard_scalars`

Logs scalar metrics (objective value, best so far) to TensorBoard at each iteration.

```
spot._write_tensorboard_scalars()
```

76.3.9.3. `_write_tensorboard_hparams`

Logs hyperparameters and their corresponding metrics for comparison.

```
spot._write_tensorboard_hparams(X, y)
```

76.4. Description of the Infill Point Acquisition Steps

Assume that the initial design has been generated and the surrogate model has been fitted. Now we consider in detail, step-by-step, how the next infill point is selected and describe the methods used by `SpotOptim` in detail.

The main driver for this process is `suggest_next_infill_point`, which acts as a dispatcher.

76.4.1. 1. Optimize Acquisition Function

The optimizer uses the specified `acquisition` function (e.g., Lower Confidence Bound, Expected Improvement) to find promising candidates in the transformed and mapped search space.

The specific optimization strategy depends on `acquisition_optimizer`:

- **TriCands** (`_optimize_acquisition_tricands`): Uses a geometric strategy based on Delaunay triangulation of existing points to propose candidates, evaluating the acquisition function on them.
- **Differential Evolution** (`_optimize_acquisition_de`): Uses `scipy.optimize.differential_evolution` to globally optimize the acquisition function. Can be cold-started or warm-started with the best solution found so far.
- **Scipy Minimize** (`_optimize_acquisition_scipy`): Uses standard gradient-based optimization (e.g., L-BFGS-B, Nelder-Mead).

76.4.2. 2. Candidate Validation and Selection

Once a candidate is proposed (`_try_optimizer_candidates`), it undergoes rigorous validation:

1. **Rounding** (`repair_non_numeric`): If variables are integer or categorical, the candidate values are rounded to the nearest valid level.

76. SpotOptim Execution Flow

2. **Uniqueness Check (`select_new`)**: The candidate is compared against all previously evaluated points (`self.X_`) using a distance tolerance (`self.tolerance_x`).
 - If the candidate is **unique** (distance > tolerance), it is accepted.
 - If the candidate is **too close** to an existing point, it is rejected to avoid redundant evaluations.

```
# Simplified logic flow
candidate = optimize_acquisition()
candidate = repair_types(candidate)
if is_unique(candidate, history):
    return candidate
else:
    reject_candidate()
```

76.4.3. 3. Failure Handling and Fallback Strategies

If the primary optimization strategy fails to find a unique, valid candidate (e.g., the optimizer converges to an already visited optimum), `SpotOptim` triggers a fallback mechanism (`_try_fallback_strategy` calling `_handle_acquisition_failure`).

- **Strategy**: By default (`acquisition_failure_strategy="random"`), it generates a random space-filling point using Latin Hypercube Sampling.
- **Retries**: This process repeats up to `max_attempts` (default 10) times until a unique point is found.
- **Last Resort**: If even the random fallback fails to find a unique point after all attempts, the last generated candidate is returned, even if it is a duplicate, to allow the loop to continue (though potentially inefficiently).

76.4.4. Key Arguments and Configuration

The behavior of the infill process is controlled by several key arguments passed to the `SpotOptim` constructor.

76.4.4.1. `tolerance_x`

- **Description**: Defines the minimum distance required between a proposed candidate and any existing point in the transformed search space. It determines when a point is considered “new” versus “duplicate”. The distance metric used is defined by `min_tol_metric`.
- **Default**: None. Internally sets to `np.sqrt(np.spacing(1))` (approx. $1.49\text{e-}8$).
- **Effect of Evaluation**:

76.4. Description of the Infill Point Acquisition Steps

- **Larger value:** Forces the optimizer to explore points further away from existing solutions. This promotes exploration but prevents fine-tuning (exploitation) in the very late stages.
- **Smaller value:** Allows the optimizer to sample points very close to existing ones. Useful for high-precision fine-tuning but increases the risk of numerical issues or redundant evaluations if the surface is flat.

76.4.4.2. `min_tol_metric`

- **Description:** Specifies the distance metric used when checking `tolerance_x`. This allows standardizing the concept of “closeness” based on the problem topology.
- **Default:** "chebyshev"
- **Options:** Supports metrics from `scipy.spatial.distance.cdist`:
 - "chebyshev": L-infinity distance. Checks if max coordinate difference is within tolerance (hypercube check). Default. Matches legacy behavior.
 - "euclidean": Standard L2 distance. Checks distances in a hypersphere.
 - "minkowski": Lp distance (default p=2).
 - Others: "cityblock" (Manhattan), "cosine", "canberra", etc.
- **Effect of Evaluation:**
 - “euclidean”: Ensures points are separated by a radial distance, appropriate for most continuous spaces.
 - “chebyshev”: Allows points closer on diagonals but strictly separated on axes, useful for box-constrained problems where axis-independence matters.

76.4.4.3. `acquisition_failure_strategy`

- **Description:** Defines the fallback method used when the acquisition optimizer fails to propose a unique point (usually because the model converges to a known optimum).
- **Default:** "random"
- **Effect of Evaluation:**
 - "random": Uses Latin Hypercube Sampling (LHS) to propose a space-filling point. This effectively “restarts” exploration in a new area when exploitation stalls.
 - *Custom:* Users could potentially implement other strategies (e.g., pure random uniform) by modifying the class, but "random" (LHS) is the standard robust choice provided.

76.4.4.4. `acquisition_optimizer`

- **Description:** Specifies the numerical optimization algorithm used to maximize/minimize the acquisition function.
- **Default:** "differential_evolution"
- **Effect of Evaluation:**
 - "differential_evolution": A global optimizer. More robust at finding the true global optimum of the complex, multi-modal acquisition landscape. generally preferred for better exploration.
 - "L-BFGS-B", "Nelder-Mead" (via scipy): Local optimizers. Much faster but liable to get stuck in local optima of the acquisition function, potentially missing the actual best area to sample next.

76.4.4.5. `acquisition_fun_return_size`

- **Description:** Specifies the number of top candidate points to return from the acquisition function optimization step.
- **Default:** 3
- **Effect of Evaluation:**
 - **Value > 1:** The optimizer returns the top N best candidates found by the acquisition optimizer (e.g., DE or TriCands). `SpotOptim` then iterates through these candidates in order of quality and selects the first one that is **unique** (i.e., satisfies `tolerance_x` with respect to existing points). This significantly improves robustness against converging to duplicates.
 - **Value = 1:** Only the single best candidate is returned. If this candidate is a duplicate of an existing point, the acquisition step fails immediately and triggers the fallback strategy (e.g., random sampling), potentially missing other valid high-quality candidates.

76.4.5. Example Usage

```
# Manually triggering the infill suggestion process
# 1. Fit the surrogate model
opt._fit_surrogate(opt.X_, opt.y_)

# 2. Ask for the next point
x_next = opt.suggest_next_infill_point()

# 3. Check if it handles duplicates
# Force a duplicate situation by asking again without updating training data
x_next_2 = opt.suggest_next_infill_point()
```

```

if np.allclose(x_next, x_next_2):
    print("Warning: Optimizer returned same point (model hasn't updated)")
else:
    print("Optimizer found a new point or used fallback")

```

76.5. Using the Surrogate Model

SpotOptim allows to use the surrogate model to predict the objective function values for new points. The default surrogate model is sklearn's Gaussian Process Regressor. Here we compare the default surrogate model with a simple linear model, which is provided as MLP_Surrogate in the `spotoptim.surrogate` module.

```

import numpy as np
import pandas as pd
from spotoptim.surrogate import Kriging, MLPSurrogate
from spotoptim.function.so import rosenbrock, ackley
from sklearn.metrics import mean_squared_error, r2_score

def compare_surrogates(func, dim, n_train, n_test=1000, bounds=(-5, 5)):
    # Generate data
    rng = np.random.default_rng(42)
    X_train = rng.uniform(bounds[0], bounds[1], (n_train, dim))
    y_train = func(X_train)

    X_test = rng.uniform(bounds[0], bounds[1], (n_test, dim))
    y_test = func(X_test)

    # Kriging (Default)
    kriging = Kriging(seed=42)
    kriging.fit(X_train, y_train)
    y_pred_kriging = kriging.predict(X_test)

    # MLP Surrogate
    # Note: MLPs typically need more data, but we'll test on the same set
    mlp = MLPSurrogate(seed=42, epochs=500, verbose=False)
    mlp.fit(X_train, y_train)
    y_pred_mlp = mlp.predict(X_test)

    # Metrics
    rmse_kriging = np.sqrt(mean_squared_error(y_test, y_pred_kriging))
    rmse_mlp = np.sqrt(mean_squared_error(y_test, y_pred_mlp))

```

```

    return {
        "Function": func.__name__,
        "Dim": dim,
        "N_Train": n_train,
        "RMSE (Kriging)": rmse_kriging,
        "RMSE (MLP)": rmse_mlp
    }

# Define Scenarios
scenarios = [
    (rosenbrock, 2, 50),
    (rosenbrock, 20, 200),
    (ackley, 2, 50),
    (ackley, 20, 200)
]

results = []
print(f"{'Function':<15} | {'Dim':<5} | {'N_Train':<8} | {'RMSE Kriging':<15} | {'RMSE MLP':<15}")
print("-" * 70)

for func, dim, n_train in scenarios:
    res = compare_surrogates(func, dim, n_train)
    results.append(res)
    print(f"{res['Function']:<15} | {res['Dim']:<5} | {res['N_Train']:<8} | {res['RMSE Kriging']:<15} | {res['RMSE MLP']:<15}")

```

Function	Dim	N_Train	RMSE Kriging	RMSE MLP
rosenbrock	2	50	171.8553	4868.6868
rosenbrock	20	200	42638.6846	80009.5260
ackley	2	50	0.8096	0.9328
ackley	20	200	0.2448	0.6952

76.6. Comparing the SpotOptim With the Default Surrogate Model and MLP Surrogate

Here we run SpotOptim using the default Kriging surrogate and the MLPSurrogate on standard test functions to compare their optimization performance.

```

from spotoptim import SpotOptim
from spotoptim.surrogate import MLPSurrogate
from spotoptim.function.so import rosenbrock, ackley, robot_arm_hard

```

76.6. Comparing the SpotOptim With the Default Surrogate Model and MLP Surrogate

```
import numpy as np

def run_spotoptim_comparison(func, dim, bounds, max_iter=50):
    print(f"\n--- optimizing {func.__name__} ({dim}D) ---")

    # Define bounds list
    bounds_list = [bounds] * dim

    # 1. Default Surrogate (Kriging)
    print("Running with Default Surrogate (Kriging)...")
    opt_kriging = SpotOptim(
        fun=func,
        bounds=bounds_list,
        max_iter=max_iter,
        n_initial=10,
        seed=42,
        verbose=False
    )
    res_kriging = opt_kriging.optimize()

    # 2. MLP Surrogate
    print("Running with MLPSurrogate...")
    # Use fewer epochs for speed in this demo, but enough to learn
    # Note: MLPSurrogate improved with standard defaults (128x3 relu)
    mlp = MLPSurrogate(epochs=200, verbose=False, seed=42)

    opt_mlp = SpotOptim(
        fun=func,
        bounds=bounds_list,
        surrogate=mlp,
        max_iter=max_iter,
        n_initial=10,
        seed=42,
        verbose=False
    )
    res_mlp = opt_mlp.optimize()

    return {
        "Function": func.__name__,
        "Dim": dim,
        "Best_Kriging": res_kriging.fun,
        "Best_MLP": res_mlp.fun
    }
```

76. SpotOptim Execution Flow

```
# Scenarios
scenarios_opt = [
    (rosenbrock, 2, (-5, 5)),
    (rosenbrock, 10, (-5, 5)),
    (ackley, 2, (-32.768, 32.768)),
    (ackley, 10, (-32.768, 32.768)),
    (robot_arm_hard, 10, (0.0, 1.0))
]

results_opt = []
print(f"\n{'Function':<15} | {'Dim':<5} | {'Best y (Default)':<20} | {'Best y (MLP)':<20}")
print("-" * 70)

for func, dim, bnds in scenarios_opt:
    res = run_spotoptim_comparison(func, dim, bnds, max_iter=50)
    results_opt.append(res)
    print(f"{res['Function']:<15} | {res['Dim']:<5} | {res['Best_Kriging']:<20.6e} | {res['Best_MLP']:<20.6e}")
```

Function	Dim	Best y (Default)	Best y (MLP)
----------	-----	------------------	--------------

--- optimizing rosenbrock (2D) ---

Running with Default Surrogate (Kriging)...

Running with MLPSurrogate...

rosenbrock	2	9.502985e-01	4.545263e-01
------------	---	--------------	--------------

--- optimizing rosenbrock (10D) ---

Running with Default Surrogate (Kriging)...

Running with MLPSurrogate...

rosenbrock	10	3.589675e+02	7.808923e+03
------------	----	--------------	--------------

--- optimizing ackley (2D) ---

Running with Default Surrogate (Kriging)...

Running with MLPSurrogate...

ackley	2	3.113849e-03	2.993562e+00
--------	---	--------------	--------------

--- optimizing ackley (10D) ---

Running with Default Surrogate (Kriging)...

Running with MLPSurrogate...

ackley	10	1.615921e+01	8.888862e+00
--------	----	--------------	--------------

--- optimizing robot_arm_hard (10D) ---

Running with Default Surrogate (Kriging)...

Running with MLPSurrogate...

robot_arm_hard | 10 | 2.226474e+01 | 2.226474e+01

76.7. Multi-Surrogate Switching in SpotOptim

SpotOptim also supports **probabilistic multi-surrogate optimization**. This allows the optimizer to switch between different surrogate models (e.g., Gaussian Process and MLP) during the optimization loop. This can be beneficial for escaping local optima or combining the strengths of different modeling approaches (e.g., local exploitation of GPs vs global approximation of Neural Networks).

To use this feature, pass a **list of surrogate objects** to the `surrogate` argument. You can optionally control the selection probability using `prob_surrogate`.

76.7.1. Comparison: Default (GP) vs Multi-Surrogate (GP + MLP) on Robot Arm

The following example demonstrates how to set up a hybrid optimization strategy combining: 1. **GaussianProcessRegressor (default)**: Good for valid uncertainty estimation and exploitation. 2. **MLPSurrogate**: Good for global trend modeling and scalability.

We compare the performance on the **10D Robot Arm** problem.

```
from spotoptim import SpotOptim
from spotoptim.surrogate import MLPSurrogate
from spotoptim.function.so import robot_arm_hard
from sklearn.gaussian_process import GaussianProcessRegressor
from sklearn.gaussian_process.kernels import Matern, ConstantKernel
import numpy as np

# 1. Setup the Benchmark
func = robot_arm_hard
dim = 10
bounds = [(0.0, 1.0)] * dim
max_iter = 1000
seed = 42

print(f"--- Optimizing {func.__name__} (10D) ---")

# --- Run 1: Multi-Surrogate (GP + MLP) ---
print("1. Running Multi-Surrogate SpotOptim (GP + MLP)...")
```

76. SpotOptim Execution Flow

```
# Define Model A: Gaussian Process (Same as default)
kernel = ConstantKernel(1.0, (1e-2, 1e12)) * Matern(
    length_scale=1.0, length_scale_bounds=(1e-4, 1e2), nu=2.5
)
gp_model = GaussianProcessRegressor(
    kernel=kernel,
    n_restarts_optimizer=10,
    normalize_y=True,
    random_state=seed
)

# Define Model B: MLP Surrogate
mlp_model = MLPSurrogate(
    epochs=1000,
    verbose=False,
    seed=seed,
    l1=128,
    num_hidden_layers=3,
    activation="relu"
)

# Run Optimization with both surrogates
opt_multi = SpotOptim(
    fun=func,
    bounds=bounds,
    surrogate=[gp_model, mlp_model],    # Pass list of models
    prob_surrogate=[0.5, 0.5],         # 50% chance for each at every step
    max_iter=max_iter,
    n_initial=20, # Increased for 10D to support Delaunay triangulation (needs > d+1 p
    seed=seed,
    verbose=True,
    acquisition_optimizer="de_tricands",
    prob_de_tricands=0.1,
    restart_after_n=10,
)
res_multi = opt_multi.optimize()
print(f"    Best y (Multi-Surrogate): {res_multi.fun:.6f}")

# --- Run 2: Default SpotOptim (Standard GP) ---
print("2. Running Default SpotOptim (GP only)...")
opt_default = SpotOptim(
    fun=func,
    bounds=bounds,
    max_iter=max_iter,
```



```

    n_initial=20,
    seed=seed,
    verbose=True,
    acquisition_optimizer="de_tricands",
    prob_de_tricands=0.1,
    restart_after_n=10,
)
res_default = opt_default.optimize()
print(f"    Best y (Default): {res_default.fun:.6f}")

# --- Summary ---
print("\n--- Results Summary ---")
print(f"{'Method':<25} | {'Best Objective':<15}")
print("-" * 40)
print(f"{'Default (GP Only)':<25} | {res_default.fun:<15.6f}")
print(f"{'Multi-Surrogate (GP+MLP)':<25} | {res_multi.fun:<15.6f}")

```

--- Optimizing robot_arm_hard (10D) ---

1. Running Multi-Surrogate SpotOptim (GP + MLP)...

TensorBoard logging disabled

Initial best: $f(x) = 23.358705$

Iter 1	Best: 23.358705	Curr: 3013.582310	Rate: 0.00	Evals: 2.1%
Iter 2	Best: 23.358705	Curr: 2196.052832	Rate: 0.00	Evals: 2.2%
Iter 3	Best: 23.358705	Curr: 163.284717	Rate: 0.00	Evals: 2.3%
Iter 4	Best: 23.358705	Curr: 3278.917129	Rate: 0.00	Evals: 2.4%
Iter 5	Best: 23.358705	Curr: 118.542966	Rate: 0.00	Evals: 2.5%
Iter 6	Best: 23.358705	Curr: 583.300464	Rate: 0.00	Evals: 2.6%
Iter 7	Best: 23.358705	Curr: 57.819223	Rate: 0.00	Evals: 2.7%
Iter 8	Best: 23.358705	Curr: 2424.948397	Rate: 0.00	Evals: 2.8%
Iter 9	Best: 23.358705	Curr: 88.433880	Rate: 0.00	Evals: 2.9%

Restarting optimization: success_rate 0 for 10 iterations.

Starting point x0 validated and processed successfully.

Original scale: [0.12897446 0.11253871 0.53861201 0.60847174 0.55760895 0.75440905
0.99729245 0.253624 0.50095326 0.4098481]

Internal scale: [0.12897446 0.11253871 0.53861201 0.60847174 0.55760895 0.75440905
0.99729245 0.253624 0.50095326 0.4098481]

Restart injection: Using best found point so far as starting point ($f(x)=23.358705$).

Including 1 starting points from x0 in initial design.

Skipping re-evaluation of injected best point.

Initial best: $f(x) = 23.358705$

Iter 1	Best: 23.358705	Curr: 1450.330578	Rate: 0.00	Evals: 5.0%
Iter 2	Best: 23.358705	Curr: 1613.658468	Rate: 0.00	Evals: 5.1%
Iter 3	Best: 23.358705	Curr: 3003.378806	Rate: 0.00	Evals: 5.2%

76. SpotOptim Execution Flow

```
Iter 4 | Best: 23.358705 | Curr: 3273.379367 | Rate: 0.00 | Evals: 5.3%
Iter 5 | Best: 23.358705 | Curr: 43.128173 | Rate: 0.00 | Evals: 5.4%
Iter 6 | Best: 23.358705 | Curr: 4009.929970 | Rate: 0.00 | Evals: 5.5%
Iter 7 | Best: 23.358705 | Curr: 178.984918 | Rate: 0.00 | Evals: 5.6%
Iter 8 | Best: 23.358705 | Curr: 208.038600 | Rate: 0.00 | Evals: 5.7%
Iter 9 | Best: 23.358705 | Curr: 104.536775 | Rate: 0.00 | Evals: 5.8%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.12897446 0.11253871 0.53861201 0.60847174 0.55760895 0.75440905
0.99729245 0.253624 0.50095326 0.4098481 ]
  Internal scale: [0.12897446 0.11253871 0.53861201 0.60847174 0.55760895 0.75440905
0.99729245 0.253624 0.50095326 0.4098481 ]
Restart injection: Using best found point so far as starting point (f(x)=23.358705).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 22.480133
Iter 1 | Best: 22.480133 | Curr: 1846.271433 | Rate: 0.00 | Evals: 7.9%
Iter 2 | Best: 22.480133 | Curr: 1082.780239 | Rate: 0.00 | Evals: 8.0%
Iter 3 | Best: 22.480133 | Curr: 2592.970683 | Rate: 0.00 | Evals: 8.1%
Iter 4 | Best: 22.480133 | Curr: 452.571413 | Rate: 0.00 | Evals: 8.2%
Iter 5 | Best: 22.480133 | Curr: 73.940254 | Rate: 0.00 | Evals: 8.3%
Iter 6 | Best: 22.480133 | Curr: 145.282247 | Rate: 0.00 | Evals: 8.4%
Iter 7 | Best: 22.480133 | Curr: 87.049739 | Rate: 0.00 | Evals: 8.5%
Iter 8 | Best: 22.480133 | Curr: 77.476187 | Rate: 0.00 | Evals: 8.6%
Iter 9 | Best: 22.480133 | Curr: 79.350699 | Rate: 0.00 | Evals: 8.7%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567
0.75120132 0.82059894 0.03724163 0.869633 ]
  Internal scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567
0.75120132 0.82059894 0.03724163 0.869633 ]
Restart injection: Using best found point so far as starting point (f(x)=22.480133).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 22.480133
Iter 1 | Best: 22.480133 | Curr: 83.761120 | Rate: 0.00 | Evals: 10.8%
Iter 2 | Best: 22.480133 | Curr: 662.854364 | Rate: 0.00 | Evals: 10.9%
Iter 3 | Best: 22.480133 | Curr: 41.042709 | Rate: 0.00 | Evals: 11.0%
Iter 4 | Best: 22.480133 | Curr: 3169.311680 | Rate: 0.00 | Evals: 11.1%
Iter 5 | Best: 22.480133 | Curr: 273.583299 | Rate: 0.00 | Evals: 11.2%
Iter 6 | Best: 22.480133 | Curr: 2911.398719 | Rate: 0.00 | Evals: 11.3%
Iter 7 | Best: 22.480133 | Curr: 153.554936 | Rate: 0.00 | Evals: 11.4%
Iter 8 | Best: 22.480133 | Curr: 47.280793 | Rate: 0.00 | Evals: 11.5%
Iter 9 | Best: 22.480133 | Curr: 181.229621 | Rate: 0.00 | Evals: 11.6%
Restarting optimization: success_rate 0 for 10 iterations.
```

76.7. Multi-Surrogate Switching in SpotOptim

Starting point x0 validated and processed successfully.

Original scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567 0.75120132 0.82059894 0.03724163 0.869633]

Internal scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567 0.75120132 0.82059894 0.03724163 0.869633]

Restart injection: Using best found point so far as starting point (f(x)=22.480133).

Including 1 starting points from x0 in initial design.

Skipping re-evaluation of injected best point.

Initial best: f(x) = 22.480133

Iter 1 | Best: 22.480133 | Curr: 5226.054629 | Rate: 0.00 | Evals: 13.7%

Iter 2 | Best: 22.480133 | Curr: 4080.652466 | Rate: 0.00 | Evals: 13.8%

Iter 3 | Best: 22.480133 | Curr: 374.331925 | Rate: 0.00 | Evals: 13.9%

Iter 4 | Best: 22.480133 | Curr: 123.000007 | Rate: 0.00 | Evals: 14.0%

Iter 5 | Best: 22.480133 | Curr: 122.579920 | Rate: 0.00 | Evals: 14.1%

Iter 6 | Best: 22.480133 | Curr: 87.489213 | Rate: 0.00 | Evals: 14.2%

Iter 7 | Best: 22.480133 | Curr: 969.811694 | Rate: 0.00 | Evals: 14.3%

Iter 8 | Best: 22.480133 | Curr: 151.983099 | Rate: 0.00 | Evals: 14.4%

Iter 9 | Best: 22.480133 | Curr: 458.801841 | Rate: 0.00 | Evals: 14.5%

Restarting optimization: success_rate 0 for 10 iterations.

Starting point x0 validated and processed successfully.

Original scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567 0.75120132 0.82059894 0.03724163 0.869633]

Internal scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567 0.75120132 0.82059894 0.03724163 0.869633]

Restart injection: Using best found point so far as starting point (f(x)=22.480133).

Including 1 starting points from x0 in initial design.

Skipping re-evaluation of injected best point.

Initial best: f(x) = 22.480133

Iter 1 | Best: 22.480133 | Curr: 696.009525 | Rate: 0.00 | Evals: 16.6%

Iter 2 | Best: 22.480133 | Curr: 422.066545 | Rate: 0.00 | Evals: 16.7%

Iter 3 | Best: 22.480133 | Curr: 94.069890 | Rate: 0.00 | Evals: 16.8%

Iter 4 | Best: 22.480133 | Curr: 412.702648 | Rate: 0.00 | Evals: 16.9%

Iter 5 | Best: 22.480133 | Curr: 2581.889598 | Rate: 0.00 | Evals: 17.0%

Iter 6 | Best: 22.480133 | Curr: 126.481272 | Rate: 0.00 | Evals: 17.1%

Iter 7 | Best: 22.480133 | Curr: 98.768601 | Rate: 0.00 | Evals: 17.2%

Iter 8 | Best: 22.480133 | Curr: 195.022091 | Rate: 0.00 | Evals: 17.3%

Iter 9 | Best: 22.480133 | Curr: 175.076969 | Rate: 0.00 | Evals: 17.4%

Restarting optimization: success_rate 0 for 10 iterations.

Starting point x0 validated and processed successfully.

Original scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567 0.75120132 0.82059894 0.03724163 0.869633]

Internal scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567 0.75120132 0.82059894 0.03724163 0.869633]

Restart injection: Using best found point so far as starting point (f(x)=22.480133).

Including 1 starting points from x0 in initial design.

76. SpotOptim Execution Flow

Skipping re-evaluation of injected best point.

Initial best: $f(x) = 22.480133$

Iter 1	Best: 22.480133	Curr: 4175.379221	Rate: 0.00	Evals: 19.5%
Iter 2	Best: 22.480133	Curr: 62.430499	Rate: 0.00	Evals: 19.6%
Iter 3	Best: 22.480133	Curr: 4664.342179	Rate: 0.00	Evals: 19.7%
Iter 4	Best: 22.480133	Curr: 573.590071	Rate: 0.00	Evals: 19.8%
Iter 5	Best: 22.480133	Curr: 39.752357	Rate: 0.00	Evals: 19.9%
Iter 6	Best: 22.480133	Curr: 42.773202	Rate: 0.00	Evals: 20.0%
Iter 7	Best: 22.480133	Curr: 31.829089	Rate: 0.00	Evals: 20.1%
Iter 8	Best: 21.578808	Rate: 0.12	Evals: 20.2%	
Iter 9	Best: 21.578808	Curr: 44.646063	Rate: 0.11	Evals: 20.3%
Iter 10	Best: 21.578808	Curr: 3103.358681	Rate: 0.10	Evals: 20.4%
Iter 11	Best: 21.578808	Curr: 43.589985	Rate: 0.10	Evals: 20.5%
Iter 12	Best: 21.578808	Curr: 533.764224	Rate: 0.10	Evals: 20.6%
Iter 13	Best: 21.578808	Curr: 83.694192	Rate: 0.10	Evals: 20.7%
Iter 14	Best: 21.578808	Curr: 86.099939	Rate: 0.10	Evals: 20.8%
Iter 15	Best: 21.578808	Curr: 78.222550	Rate: 0.10	Evals: 20.9%
Iter 16	Best: 21.578808	Curr: 108.432080	Rate: 0.10	Evals: 21.0%
Iter 17	Best: 21.578808	Curr: 33.734033	Rate: 0.10	Evals: 21.1%
Iter 18	Best: 21.578808	Curr: 99.782757	Rate: 0.00	Evals: 21.2%
Iter 19	Best: 21.578808	Curr: 59.657865	Rate: 0.00	Evals: 21.3%
Iter 20	Best: 19.108051	Rate: 0.10	Evals: 21.4%	
Iter 21	Best: 19.108051	Curr: 44.809311	Rate: 0.10	Evals: 21.5%
Iter 22	Best: 19.108051	Curr: 40.222936	Rate: 0.10	Evals: 21.6%
Iter 23	Best: 19.108051	Curr: 78.258148	Rate: 0.10	Evals: 21.7%
Iter 24	Best: 19.108051	Curr: 41.368402	Rate: 0.10	Evals: 21.8%
Iter 25	Best: 19.108051	Curr: 47.660263	Rate: 0.10	Evals: 21.9%
Iter 26	Best: 19.108051	Curr: 54.471952	Rate: 0.10	Evals: 22.0%
Iter 27	Best: 19.108051	Curr: 71.348053	Rate: 0.10	Evals: 22.1%
Iter 28	Best: 19.108051	Curr: 86.371536	Rate: 0.10	Evals: 22.2%
Iter 29	Best: 19.108051	Curr: 109.917884	Rate: 0.10	Evals: 22.3%
Iter 30	Best: 19.108051	Curr: 126.015952	Rate: 0.00	Evals: 22.4%
Iter 31	Best: 19.108051	Curr: 55.480825	Rate: 0.00	Evals: 22.5%
Iter 32	Best: 19.108051	Curr: 67.133710	Rate: 0.00	Evals: 22.6%
Iter 33	Best: 19.108051	Curr: 26.802925	Rate: 0.00	Evals: 22.7%
Iter 34	Best: 19.108051	Curr: 3055.508484	Rate: 0.00	Evals: 22.8%
Iter 35	Best: 19.108051	Curr: 480.069325	Rate: 0.00	Evals: 22.9%
Iter 36	Best: 19.108051	Curr: 1712.384713	Rate: 0.00	Evals: 23.0%
Iter 37	Best: 19.108051	Curr: 29.873473	Rate: 0.00	Evals: 23.1%
Iter 38	Best: 19.108051	Curr: 48.480064	Rate: 0.00	Evals: 23.2%

Restarting optimization: success_rate 0 for 10 iterations.

Starting point x0 validated and processed successfully.

Original scale: [0.50613526 0.23132681 0.40161636 0.34075494 0.41920951 0.25922529
0.60687194 0.50047366 0.63037481 0.51311913]

Internal scale: [0.50613526 0.23132681 0.40161636 0.34075494 0.41920951 0.25922529

76.7. Multi-Surrogate Switching in SpotOptim

```
0.60687194 0.50047366 0.63037481 0.51311913]
Restart injection: Using best found point so far as starting point (f(x)=19.108051).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 19.108051
Iter 1 | Best: 19.108051 | Curr: 335.833460 | Rate: 0.00 | Evals: 25.3%
Iter 2 | Best: 19.108051 | Curr: 210.941246 | Rate: 0.00 | Evals: 25.4%
Iter 3 | Best: 19.108051 | Curr: 197.293894 | Rate: 0.00 | Evals: 25.5%
Iter 4 | Best: 19.108051 | Curr: 129.776546 | Rate: 0.00 | Evals: 25.6%
Iter 5 | Best: 19.108051 | Curr: 2664.932954 | Rate: 0.00 | Evals: 25.7%
Iter 6 | Best: 19.108051 | Curr: 90.629203 | Rate: 0.00 | Evals: 25.8%
Iter 7 | Best: 19.108051 | Curr: 77.874227 | Rate: 0.00 | Evals: 25.9%
Iter 8 | Best: 19.108051 | Curr: 98.539003 | Rate: 0.00 | Evals: 26.0%
Iter 9 | Best: 19.108051 | Curr: 147.845929 | Rate: 0.00 | Evals: 26.1%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.50613526 0.23132681 0.40161636 0.34075494 0.41920951 0.25922529
0.60687194 0.50047366 0.63037481 0.51311913]
Internal scale: [0.50613526 0.23132681 0.40161636 0.34075494 0.41920951 0.25922529
0.60687194 0.50047366 0.63037481 0.51311913]
Restart injection: Using best found point so far as starting point (f(x)=19.108051).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 18.171750
Iter 1 | Best: 18.171750 | Curr: 2494.698672 | Rate: 0.00 | Evals: 28.2%
Iter 2 | Best: 18.171750 | Curr: 1412.024650 | Rate: 0.00 | Evals: 28.3%
Iter 3 | Best: 18.171750 | Curr: 228.972542 | Rate: 0.00 | Evals: 28.4%
Iter 4 | Best: 18.171750 | Curr: 215.478957 | Rate: 0.00 | Evals: 28.5%
Iter 5 | Best: 18.171750 | Curr: 226.282011 | Rate: 0.00 | Evals: 28.6%
Iter 6 | Best: 18.171750 | Curr: 145.058382 | Rate: 0.00 | Evals: 28.7%
Iter 7 | Best: 18.171750 | Curr: 191.670462 | Rate: 0.00 | Evals: 28.8%
Iter 8 | Best: 18.171750 | Curr: 240.576538 | Rate: 0.00 | Evals: 28.9%
Iter 9 | Best: 18.171750 | Curr: 57.244058 | Rate: 0.00 | Evals: 29.0%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.65238644 0.60930271 0.03433355 0.72954934 0.93074097 0.25509536
0.98995351 0.84219394 0.52650099 0.3756625 ]
Internal scale: [0.65238644 0.60930271 0.03433355 0.72954934 0.93074097 0.25509536
0.98995351 0.84219394 0.52650099 0.3756625 ]
Restart injection: Using best found point so far as starting point (f(x)=18.171750).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 18.171750
Iter 1 | Best: 18.171750 | Curr: 103.244355 | Rate: 0.00 | Evals: 31.1%
Iter 2 | Best: 18.171750 | Curr: 1830.807743 | Rate: 0.00 | Evals: 31.2%
```

76. SpotOptim Execution Flow

```
Iter 3 | Best: 18.171750 | Curr: 172.172955 | Rate: 0.00 | Evals: 31.3%
Iter 4 | Best: 18.171750 | Curr: 113.326626 | Rate: 0.00 | Evals: 31.4%
Iter 5 | Best: 18.171750 | Curr: 144.158206 | Rate: 0.00 | Evals: 31.5%
Iter 6 | Best: 18.171750 | Curr: 172.626343 | Rate: 0.00 | Evals: 31.6%
Iter 7 | Best: 18.171750 | Curr: 116.603354 | Rate: 0.00 | Evals: 31.7%
Iter 8 | Best: 18.171750 | Curr: 95.251654 | Rate: 0.00 | Evals: 31.8%
Iter 9 | Best: 18.171750 | Curr: 139.828612 | Rate: 0.00 | Evals: 31.9%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.65238644 0.60930271 0.03433355 0.72954934 0.93074097 0.25509536
0.98995351 0.84219394 0.52650099 0.3756625 ]
  Internal scale: [0.65238644 0.60930271 0.03433355 0.72954934 0.93074097 0.25509536
0.98995351 0.84219394 0.52650099 0.3756625 ]
Restart injection: Using best found point so far as starting point (f(x)=18.171750).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 18.171750
Iter 1 | Best: 18.171750 | Curr: 3964.939819 | Rate: 0.00 | Evals: 34.0%
Iter 2 | Best: 18.171750 | Curr: 102.861875 | Rate: 0.00 | Evals: 34.1%
Iter 3 | Best: 18.171750 | Curr: 138.510266 | Rate: 0.00 | Evals: 34.2%
Iter 4 | Best: 18.171750 | Curr: 88.629403 | Rate: 0.00 | Evals: 34.3%
Iter 5 | Best: 18.171750 | Curr: 4147.957863 | Rate: 0.00 | Evals: 34.4%
Iter 6 | Best: 18.171750 | Curr: 130.603489 | Rate: 0.00 | Evals: 34.5%
Iter 7 | Best: 18.171750 | Curr: 139.147290 | Rate: 0.00 | Evals: 34.6%
Iter 8 | Best: 18.171750 | Curr: 77.272890 | Rate: 0.00 | Evals: 34.7%
Optimizer candidate 1/3 was duplicate/invalid.
Iter 9 | Best: 17.524655 | Rate: 0.11 | Evals: 34.8%
Iter 10 | Best: 17.524655 | Curr: 87.532211 | Rate: 0.10 | Evals: 34.9%
Iter 11 | Best: 17.524655 | Curr: 50.103593 | Rate: 0.10 | Evals: 35.0%
Iter 12 | Best: 17.524655 | Curr: 153.384835 | Rate: 0.10 | Evals: 35.1%
Iter 13 | Best: 17.524655 | Curr: 131.893115 | Rate: 0.10 | Evals: 35.2%
Iter 14 | Best: 17.524655 | Curr: 2583.282311 | Rate: 0.10 | Evals: 35.3%
Iter 15 | Best: 17.524655 | Curr: 90.750887 | Rate: 0.10 | Evals: 35.4%
Iter 16 | Best: 17.524655 | Curr: 84.802566 | Rate: 0.10 | Evals: 35.5%
Iter 17 | Best: 17.524655 | Curr: 114.330915 | Rate: 0.10 | Evals: 35.6%
Iter 18 | Best: 17.524655 | Curr: 107.136596 | Rate: 0.10 | Evals: 35.7%
Iter 19 | Best: 17.524655 | Curr: 5664.479260 | Rate: 0.00 | Evals: 35.8%
Iter 20 | Best: 17.524655 | Curr: 74.487106 | Rate: 0.00 | Evals: 35.9%
Iter 21 | Best: 17.524655 | Curr: 115.439110 | Rate: 0.00 | Evals: 36.0%
Iter 22 | Best: 17.524655 | Curr: 2787.022338 | Rate: 0.00 | Evals: 36.1%
Iter 23 | Best: 17.524655 | Curr: 121.922003 | Rate: 0.00 | Evals: 36.2%
Iter 24 | Best: 17.524655 | Curr: 99.999049 | Rate: 0.00 | Evals: 36.3%
Iter 25 | Best: 17.524655 | Curr: 162.647820 | Rate: 0.00 | Evals: 36.4%
Iter 26 | Best: 17.524655 | Curr: 92.342658 | Rate: 0.00 | Evals: 36.5%
Iter 27 | Best: 17.524655 | Curr: 156.561025 | Rate: 0.00 | Evals: 36.6%
```

76.7. Multi-Surrogate Switching in SpotOptim

Restarting optimization: success_rate 0 for 10 iterations.

Starting point x0 validated and processed successfully.

Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]

Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]

Restart injection: Using best found point so far as starting point (f(x)=17.524655).

Including 1 starting points from x0 in initial design.

Skipping re-evaluation of injected best point.

Initial best: f(x) = 17.524655

Iter 1 | Best: 17.524655 | Curr: 8093.922072 | Rate: 0.00 | Evals: 38.7%

Iter 2 | Best: 17.524655 | Curr: 2562.858124 | Rate: 0.00 | Evals: 38.8%

Iter 3 | Best: 17.524655 | Curr: 3660.173963 | Rate: 0.00 | Evals: 38.9%

Iter 4 | Best: 17.524655 | Curr: 4990.529671 | Rate: 0.00 | Evals: 39.0%

Iter 5 | Best: 17.524655 | Curr: 5851.496931 | Rate: 0.00 | Evals: 39.1%

Iter 6 | Best: 17.524655 | Curr: 4628.946027 | Rate: 0.00 | Evals: 39.2%

Iter 7 | Best: 17.524655 | Curr: 400.409840 | Rate: 0.00 | Evals: 39.3%

Iter 8 | Best: 17.524655 | Curr: 124.984549 | Rate: 0.00 | Evals: 39.4%

Iter 9 | Best: 17.524655 | Curr: 1786.268960 | Rate: 0.00 | Evals: 39.5%

Restarting optimization: success_rate 0 for 10 iterations.

Starting point x0 validated and processed successfully.

Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]

Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]

Restart injection: Using best found point so far as starting point (f(x)=17.524655).

Including 1 starting points from x0 in initial design.

Skipping re-evaluation of injected best point.

Initial best: f(x) = 17.524655

Iter 1 | Best: 17.524655 | Curr: 25.253733 | Rate: 0.00 | Evals: 41.6%

Iter 2 | Best: 17.524655 | Curr: 3782.878397 | Rate: 0.00 | Evals: 41.7%

Iter 3 | Best: 17.524655 | Curr: 6403.842018 | Rate: 0.00 | Evals: 41.8%

Iter 4 | Best: 17.524655 | Curr: 109.116292 | Rate: 0.00 | Evals: 41.9%

Iter 5 | Best: 17.524655 | Curr: 671.044795 | Rate: 0.00 | Evals: 42.0%

Iter 6 | Best: 17.524655 | Curr: 86.730903 | Rate: 0.00 | Evals: 42.1%

Iter 7 | Best: 17.524655 | Curr: 212.632624 | Rate: 0.00 | Evals: 42.2%

Iter 8 | Best: 17.524655 | Curr: 181.312056 | Rate: 0.00 | Evals: 42.3%

Iter 9 | Best: 17.524655 | Curr: 177.461075 | Rate: 0.00 | Evals: 42.4%

Restarting optimization: success_rate 0 for 10 iterations.

Starting point x0 validated and processed successfully.

Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]

Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]

Restart injection: Using best found point so far as starting point (f(x)=17.524655).

76. SpotOptim Execution Flow

Including 1 starting points from x0 in initial design.

Skipping re-evaluation of injected best point.

Initial best: $f(x) = 17.524655$

Iter 1	Best: 17.524655	Curr: 55.563721	Rate: 0.00	Evals: 44.5%
Iter 2	Best: 17.524655	Curr: 41.682907	Rate: 0.00	Evals: 44.6%
Iter 3	Best: 17.524655	Curr: 38.626154	Rate: 0.00	Evals: 44.7%
Iter 4	Best: 17.524655	Curr: 22.477779	Rate: 0.00	Evals: 44.8%
Iter 5	Best: 17.524655	Curr: 32.863813	Rate: 0.00	Evals: 44.9%
Iter 6	Best: 17.524655	Curr: 26.178470	Rate: 0.00	Evals: 45.0%
Iter 7	Best: 17.524655	Curr: 22.028484	Rate: 0.00	Evals: 45.1%
Iter 8	Best: 17.524655	Curr: 34.769394	Rate: 0.00	Evals: 45.2%
Iter 9	Best: 17.524655	Curr: 43.298672	Rate: 0.00	Evals: 45.3%

Restarting optimization: success_rate 0 for 10 iterations.

Starting point x0 validated and processed successfully.

Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]

Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]

Restart injection: Using best found point so far as starting point ($f(x)=17.524655$).

Including 1 starting points from x0 in initial design.

Skipping re-evaluation of injected best point.

Initial best: $f(x) = 17.524655$

Iter 1	Best: 17.524655	Curr: 198.493439	Rate: 0.00	Evals: 47.4%
Iter 2	Best: 17.524655	Curr: 244.529544	Rate: 0.00	Evals: 47.5%
Iter 3	Best: 17.524655	Curr: 61.053913	Rate: 0.00	Evals: 47.6%
Iter 4	Best: 17.524655	Curr: 259.190424	Rate: 0.00	Evals: 47.7%
Iter 5	Best: 17.524655	Curr: 2217.373370	Rate: 0.00	Evals: 47.8%
Iter 6	Best: 17.524655	Curr: 2784.469544	Rate: 0.00	Evals: 47.9%
Iter 7	Best: 17.524655	Curr: 48.627565	Rate: 0.00	Evals: 48.0%
Iter 8	Best: 17.524655	Curr: 2943.675155	Rate: 0.00	Evals: 48.1%
Iter 9	Best: 17.524655	Curr: 146.018666	Rate: 0.00	Evals: 48.2%

Restarting optimization: success_rate 0 for 10 iterations.

Starting point x0 validated and processed successfully.

Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]

Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]

Restart injection: Using best found point so far as starting point ($f(x)=17.524655$).

Including 1 starting points from x0 in initial design.

Skipping re-evaluation of injected best point.

Initial best: $f(x) = 17.524655$

Iter 1	Best: 17.524655	Curr: 196.285313	Rate: 0.00	Evals: 50.3%
Iter 2	Best: 17.524655	Curr: 107.669733	Rate: 0.00	Evals: 50.4%
Iter 3	Best: 17.524655	Curr: 114.382385	Rate: 0.00	Evals: 50.5%
Iter 4	Best: 17.524655	Curr: 86.072078	Rate: 0.00	Evals: 50.6%

76.7. Multi-Surrogate Switching in SpotOptim

Iter 5 | Best: 17.524655 | Curr: 96.273889 | Rate: 0.00 | Evals: 50.7%
Iter 6 | Best: 17.524655 | Curr: 95.054240 | Rate: 0.00 | Evals: 50.8%
Iter 7 | Best: 17.524655 | Curr: 82.754604 | Rate: 0.00 | Evals: 50.9%
Iter 8 | Best: 17.524655 | Curr: 66.643461 | Rate: 0.00 | Evals: 51.0%
Iter 9 | Best: 17.524655 | Curr: 176.056821 | Rate: 0.00 | Evals: 51.1%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]
Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]
Restart injection: Using best found point so far as starting point (f(x)=17.524655).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.524655
Iter 1 | Best: 17.524655 | Curr: 3042.594320 | Rate: 0.00 | Evals: 53.2%
Iter 2 | Best: 17.524655 | Curr: 99.819259 | Rate: 0.00 | Evals: 53.3%
Iter 3 | Best: 17.524655 | Curr: 3354.825790 | Rate: 0.00 | Evals: 53.4%
Iter 4 | Best: 17.524655 | Curr: 99.786249 | Rate: 0.00 | Evals: 53.5%
Iter 5 | Best: 17.524655 | Curr: 91.984260 | Rate: 0.00 | Evals: 53.6%
Iter 6 | Best: 17.524655 | Curr: 80.578200 | Rate: 0.00 | Evals: 53.7%
Iter 7 | Best: 17.524655 | Curr: 78.054773 | Rate: 0.00 | Evals: 53.8%
Iter 8 | Best: 17.524655 | Curr: 130.193090 | Rate: 0.00 | Evals: 53.9%
Iter 9 | Best: 17.524655 | Curr: 132.365227 | Rate: 0.00 | Evals: 54.0%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]
Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]
Restart injection: Using best found point so far as starting point (f(x)=17.524655).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.524655
Iter 1 | Best: 17.524655 | Curr: 1415.070809 | Rate: 0.00 | Evals: 56.1%
Iter 2 | Best: 17.524655 | Curr: 100.244699 | Rate: 0.00 | Evals: 56.2%
Iter 3 | Best: 17.524655 | Curr: 3390.331584 | Rate: 0.00 | Evals: 56.3%
Iter 4 | Best: 17.524655 | Curr: 54.802857 | Rate: 0.00 | Evals: 56.4%
Iter 5 | Best: 17.524655 | Curr: 96.405676 | Rate: 0.00 | Evals: 56.5%
Iter 6 | Best: 17.524655 | Curr: 136.182956 | Rate: 0.00 | Evals: 56.6%
Iter 7 | Best: 17.524655 | Curr: 100.721728 | Rate: 0.00 | Evals: 56.7%
Iter 8 | Best: 17.524655 | Curr: 197.215796 | Rate: 0.00 | Evals: 56.8%
Iter 9 | Best: 17.524655 | Curr: 79.237452 | Rate: 0.00 | Evals: 56.9%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.

76. SpotOptim Execution Flow

Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]
Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]
Restart injection: Using best found point so far as starting point ($f(x)=17.524655$).
Including 1 starting points from x_0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: $f(x) = 17.524655$
Iter 1 | Best: 17.524655 | Curr: 4736.946183 | Rate: 0.00 | Evals: 59.0%
Iter 2 | Best: 17.524655 | Curr: 6021.453645 | Rate: 0.00 | Evals: 59.1%
Iter 3 | Best: 17.524655 | Curr: 28.087317 | Rate: 0.00 | Evals: 59.2%
Iter 4 | Best: 17.524655 | Curr: 282.574281 | Rate: 0.00 | Evals: 59.3%
Iter 5 | Best: 17.524655 | Curr: 32.789852 | Rate: 0.00 | Evals: 59.4%
Iter 6 | Best: 17.524655 | Curr: 982.292883 | Rate: 0.00 | Evals: 59.5%
Iter 7 | Best: 17.524655 | Curr: 2593.434131 | Rate: 0.00 | Evals: 59.6%
Iter 8 | Best: 17.524655 | Curr: 63.083826 | Rate: 0.00 | Evals: 59.7%
Iter 9 | Best: 17.524655 | Curr: 56.518143 | Rate: 0.00 | Evals: 59.8%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x_0 validated and processed successfully.
Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]
Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]
Restart injection: Using best found point so far as starting point ($f(x)=17.524655$).
Including 1 starting points from x_0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: $f(x) = 17.524655$
Iter 1 | Best: 17.524655 | Curr: 4138.063858 | Rate: 0.00 | Evals: 61.9%
Iter 2 | Best: 17.524655 | Curr: 528.644975 | Rate: 0.00 | Evals: 62.0%
Iter 3 | Best: 17.524655 | Curr: 2535.723907 | Rate: 0.00 | Evals: 62.1%
Iter 4 | Best: 17.524655 | Curr: 3302.996151 | Rate: 0.00 | Evals: 62.2%
Iter 5 | Best: 17.524655 | Curr: 51.054911 | Rate: 0.00 | Evals: 62.3%
Iter 6 | Best: 17.524655 | Curr: 176.402063 | Rate: 0.00 | Evals: 62.4%
Iter 7 | Best: 17.524655 | Curr: 1069.606619 | Rate: 0.00 | Evals: 62.5%
Iter 8 | Best: 17.524655 | Curr: 3324.808337 | Rate: 0.00 | Evals: 62.6%
Iter 9 | Best: 17.524655 | Curr: 1375.818939 | Rate: 0.00 | Evals: 62.7%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x_0 validated and processed successfully.
Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]
Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]
Restart injection: Using best found point so far as starting point ($f(x)=17.524655$).
Including 1 starting points from x_0 in initial design.
Skipping re-evaluation of injected best point.

76.7. Multi-Surrogate Switching in SpotOptim

Initial best: $f(x) = 17.524655$

Iter	Best	Curr	Rate	Evals
1	17.524655	976.770929	0.00	64.8%
2	17.524655	68.842198	0.00	64.9%
3	17.524655	64.882017	0.00	65.0%
4	17.524655	47.492065	0.00	65.1%
5	17.524655	44.813345	0.00	65.2%
6	17.524655	52.127410	0.00	65.3%
7	17.524655	3050.049864	0.00	65.4%
8	17.524655	93.110921	0.00	65.5%
9	17.524655	96.707900	0.00	65.6%

Restarting optimization: success_rate 0 for 10 iterations.

Starting point x0 validated and processed successfully.

Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]

Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]

Restart injection: Using best found point so far as starting point ($f(x)=17.524655$).

Including 1 starting points from x0 in initial design.

Skipping re-evaluation of injected best point.

Initial best: $f(x) = 17.524655$

Iter	Best	Curr	Rate	Evals
1	17.524655	3841.775107	0.00	67.7%
2	17.524655	129.623115	0.00	67.8%
3	17.524655	57.996385	0.00	67.9%
4	17.524655	143.629123	0.00	68.0%
5	17.524655	73.578630	0.00	68.1%
6	17.524655	35.371311	0.00	68.2%
7	17.524655	30.376581	0.00	68.3%
8	17.524655	60.832696	0.00	68.4%
9	17.524655	25.659987	0.00	68.5%

Restarting optimization: success_rate 0 for 10 iterations.

Starting point x0 validated and processed successfully.

Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]

Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]

Restart injection: Using best found point so far as starting point ($f(x)=17.524655$).

Including 1 starting points from x0 in initial design.

Skipping re-evaluation of injected best point.

Initial best: $f(x) = 17.524655$

Iter	Best	Curr	Rate	Evals
1	17.524655	3146.345858	0.00	70.6%
2	17.524655	124.257609	0.00	70.7%
3	17.524655	3837.618021	0.00	70.8%
4	17.524655	105.566264	0.00	70.9%
5	17.524655	60.294948	0.00	71.0%
6	17.524655	132.897096	0.00	71.1%

76. SpotOptim Execution Flow

Iter 7 | Best: 17.524655 | Curr: 118.155164 | Rate: 0.00 | Evals: 71.2%
Iter 8 | Best: 17.524655 | Curr: 61.112932 | Rate: 0.00 | Evals: 71.3%
Iter 9 | Best: 17.524655 | Curr: 122.176944 | Rate: 0.00 | Evals: 71.4%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]
Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]
Restart injection: Using best found point so far as starting point (f(x)=17.524655).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.524655
Iter 1 | Best: 17.524655 | Curr: 44.206794 | Rate: 0.00 | Evals: 73.5%
Iter 2 | Best: 17.524655 | Curr: 27.073664 | Rate: 0.00 | Evals: 73.6%
Iter 3 | Best: 17.524655 | Curr: 91.334237 | Rate: 0.00 | Evals: 73.7%
Iter 4 | Best: 17.524655 | Curr: 276.601766 | Rate: 0.00 | Evals: 73.8%
Iter 5 | Best: 17.524655 | Curr: 84.579087 | Rate: 0.00 | Evals: 73.9%
Iter 6 | Best: 17.524655 | Curr: 23.621980 | Rate: 0.00 | Evals: 74.0%
Iter 7 | Best: 17.524655 | Curr: 26.050063 | Rate: 0.00 | Evals: 74.1%
Iter 8 | Best: 17.524655 | Curr: 920.340728 | Rate: 0.00 | Evals: 74.2%
Iter 9 | Best: 17.524655 | Curr: 400.390212 | Rate: 0.00 | Evals: 74.3%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]
Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]
Restart injection: Using best found point so far as starting point (f(x)=17.524655).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.524655
Iter 1 | Best: 17.524655 | Curr: 3474.991777 | Rate: 0.00 | Evals: 76.4%
Iter 2 | Best: 17.524655 | Curr: 178.497924 | Rate: 0.00 | Evals: 76.5%
Iter 3 | Best: 17.524655 | Curr: 3679.449924 | Rate: 0.00 | Evals: 76.6%
Iter 4 | Best: 17.524655 | Curr: 203.037313 | Rate: 0.00 | Evals: 76.7%
Iter 5 | Best: 17.524655 | Curr: 544.872477 | Rate: 0.00 | Evals: 76.8%
Iter 6 | Best: 17.524655 | Curr: 125.796756 | Rate: 0.00 | Evals: 76.9%
Iter 7 | Best: 17.524655 | Curr: 327.990392 | Rate: 0.00 | Evals: 77.0%
Iter 8 | Best: 17.524655 | Curr: 209.665707 | Rate: 0.00 | Evals: 77.1%
Iter 9 | Best: 17.524655 | Curr: 212.906563 | Rate: 0.00 | Evals: 77.2%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]

76.7. Multi-Surrogate Switching in SpotOptim

Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]
Restart injection: Using best found point so far as starting point (f(x)=17.524655).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.524655
Iter 1 | Best: 17.524655 | Curr: 2460.199561 | Rate: 0.00 | Evals: 79.3%
Iter 2 | Best: 17.524655 | Curr: 2899.656101 | Rate: 0.00 | Evals: 79.4%
Iter 3 | Best: 17.524655 | Curr: 1256.963472 | Rate: 0.00 | Evals: 79.5%
Iter 4 | Best: 17.524655 | Curr: 180.412074 | Rate: 0.00 | Evals: 79.6%
Iter 5 | Best: 17.524655 | Curr: 50.689278 | Rate: 0.00 | Evals: 79.7%
Iter 6 | Best: 17.524655 | Curr: 204.903135 | Rate: 0.00 | Evals: 79.8%
Iter 7 | Best: 17.524655 | Curr: 203.829338 | Rate: 0.00 | Evals: 79.9%
Iter 8 | Best: 17.524655 | Curr: 191.211277 | Rate: 0.00 | Evals: 80.0%
Iter 9 | Best: 17.524655 | Curr: 2410.180948 | Rate: 0.00 | Evals: 80.1%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]
Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]
Restart injection: Using best found point so far as starting point (f(x)=17.524655).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.524655
Iter 1 | Best: 17.524655 | Curr: 1561.639393 | Rate: 0.00 | Evals: 82.2%
Iter 2 | Best: 17.524655 | Curr: 546.245929 | Rate: 0.00 | Evals: 82.3%
Iter 3 | Best: 17.524655 | Curr: 1324.077028 | Rate: 0.00 | Evals: 82.4%
Iter 4 | Best: 17.524655 | Curr: 934.605897 | Rate: 0.00 | Evals: 82.5%
Iter 5 | Best: 17.524655 | Curr: 2782.370675 | Rate: 0.00 | Evals: 82.6%
Iter 6 | Best: 17.524655 | Curr: 3688.808964 | Rate: 0.00 | Evals: 82.7%
Iter 7 | Best: 17.524655 | Curr: 2687.213726 | Rate: 0.00 | Evals: 82.8%
Iter 8 | Best: 17.524655 | Curr: 1477.588866 | Rate: 0.00 | Evals: 82.9%
Iter 9 | Best: 17.524655 | Curr: 864.129319 | Rate: 0.00 | Evals: 83.0%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]
Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]
Restart injection: Using best found point so far as starting point (f(x)=17.524655).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.524655
Iter 1 | Best: 17.524655 | Curr: 74.089718 | Rate: 0.00 | Evals: 85.1%

76. SpotOptim Execution Flow

```
Iter 2 | Best: 17.524655 | Curr: 158.601189 | Rate: 0.00 | Evals: 85.2%
Iter 3 | Best: 17.524655 | Curr: 124.048580 | Rate: 0.00 | Evals: 85.3%
Iter 4 | Best: 17.524655 | Curr: 141.378396 | Rate: 0.00 | Evals: 85.4%
Iter 5 | Best: 17.524655 | Curr: 102.691686 | Rate: 0.00 | Evals: 85.5%
Iter 6 | Best: 17.524655 | Curr: 80.839369 | Rate: 0.00 | Evals: 85.6%
Iter 7 | Best: 17.524655 | Curr: 3650.128986 | Rate: 0.00 | Evals: 85.7%
Iter 8 | Best: 17.524655 | Curr: 236.428605 | Rate: 0.00 | Evals: 85.8%
Iter 9 | Best: 17.524655 | Curr: 27.972527 | Rate: 0.00 | Evals: 85.9%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.6549495  0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]
  Internal scale: [0.6549495  0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]
Restart injection: Using best found point so far as starting point (f(x)=17.524655).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.524655
Iter 1 | Best: 17.524655 | Curr: 188.403229 | Rate: 0.00 | Evals: 88.0%
Iter 2 | Best: 17.524655 | Curr: 56.917797 | Rate: 0.00 | Evals: 88.1%
Iter 3 | Best: 17.524655 | Curr: 170.351541 | Rate: 0.00 | Evals: 88.2%
Iter 4 | Best: 17.524655 | Curr: 151.475543 | Rate: 0.00 | Evals: 88.3%
Iter 5 | Best: 17.524655 | Curr: 200.327281 | Rate: 0.00 | Evals: 88.4%
Iter 6 | Best: 17.524655 | Curr: 71.278926 | Rate: 0.00 | Evals: 88.5%
Iter 7 | Best: 17.524655 | Curr: 227.755474 | Rate: 0.00 | Evals: 88.6%
Iter 8 | Best: 17.524655 | Curr: 26.133245 | Rate: 0.00 | Evals: 88.7%
Iter 9 | Best: 17.524655 | Curr: 156.853258 | Rate: 0.00 | Evals: 88.8%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.6549495  0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]
  Internal scale: [0.6549495  0.36470321 0.18842231 0.63332954 0.91877645 0.41401838
0.96665621 0.93887881 0.45219443 0.45765655]
Restart injection: Using best found point so far as starting point (f(x)=17.524655).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.524655
Iter 1 | Best: 17.524655 | Curr: 1852.788107 | Rate: 0.00 | Evals: 90.9%
Iter 2 | Best: 17.524655 | Curr: 4630.560991 | Rate: 0.00 | Evals: 91.0%
Iter 3 | Best: 17.524655 | Curr: 170.311905 | Rate: 0.00 | Evals: 91.1%
Iter 4 | Best: 17.524655 | Curr: 2839.266633 | Rate: 0.00 | Evals: 91.2%
Iter 5 | Best: 17.524655 | Curr: 863.555998 | Rate: 0.00 | Evals: 91.3%
Iter 6 | Best: 17.524655 | Curr: 3174.588355 | Rate: 0.00 | Evals: 91.4%
Iter 7 | Best: 17.524655 | Curr: 2437.536587 | Rate: 0.00 | Evals: 91.5%
Iter 8 | Best: 17.524655 | Curr: 171.105445 | Rate: 0.00 | Evals: 91.6%
```

76.7. Multi-Surrogate Switching in SpotOptim

Iter 9 | Best: 17.524655 | Curr: 2062.660522 | Rate: 0.00 | Evals: 91.7%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]
Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]
Restart injection: Using best found point so far as starting point (f(x)=17.524655).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.524655
Iter 1 | Best: 17.524655 | Curr: 3136.926852 | Rate: 0.00 | Evals: 93.8%
Iter 2 | Best: 17.524655 | Curr: 58.622702 | Rate: 0.00 | Evals: 93.9%
Iter 3 | Best: 17.524655 | Curr: 2176.657312 | Rate: 0.00 | Evals: 94.0%
Iter 4 | Best: 17.524655 | Curr: 3582.777274 | Rate: 0.00 | Evals: 94.1%
Iter 5 | Best: 17.524655 | Curr: 1364.518690 | Rate: 0.00 | Evals: 94.2%
Iter 6 | Best: 17.524655 | Curr: 58.129595 | Rate: 0.00 | Evals: 94.3%
Iter 7 | Best: 17.524655 | Curr: 1359.391109 | Rate: 0.00 | Evals: 94.4%
Iter 8 | Best: 17.524655 | Curr: 2424.053411 | Rate: 0.00 | Evals: 94.5%
Iter 9 | Best: 17.524655 | Curr: 27.457708 | Rate: 0.00 | Evals: 94.6%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]
Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]
Restart injection: Using best found point so far as starting point (f(x)=17.524655).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.524655
Iter 1 | Best: 17.524655 | Curr: 2018.525553 | Rate: 0.00 | Evals: 96.7%
Iter 2 | Best: 17.524655 | Curr: 1389.088322 | Rate: 0.00 | Evals: 96.8%
Iter 3 | Best: 17.524655 | Curr: 3261.387331 | Rate: 0.00 | Evals: 96.9%
Iter 4 | Best: 17.524655 | Curr: 3619.053142 | Rate: 0.00 | Evals: 97.0%
Iter 5 | Best: 17.524655 | Curr: 1279.002731 | Rate: 0.00 | Evals: 97.1%
Iter 6 | Best: 17.524655 | Curr: 4838.186637 | Rate: 0.00 | Evals: 97.2%
Iter 7 | Best: 17.524655 | Curr: 9078.107585 | Rate: 0.00 | Evals: 97.3%
Iter 8 | Best: 17.524655 | Curr: 3560.629864 | Rate: 0.00 | Evals: 97.4%
Iter 9 | Best: 17.524655 | Curr: 5069.977193 | Rate: 0.00 | Evals: 97.5%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]
Internal scale: [0.6549495 0.36470321 0.18842231 0.63332954 0.91877645 0.41401838 0.96665621 0.93887881 0.45219443 0.45765655]

76. SpotOptim Execution Flow

Restart injection: Using best found point so far as starting point ($f(x)=17.524655$).
Including 1 starting points from x_0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: $f(x) = 17.524655$
Iter 1 | Best: 17.524655 | Curr: 628.319267 | Rate: 0.00 | Evals: 99.6%
Iter 2 | Best: 17.524655 | Curr: 2752.912696 | Rate: 0.00 | Evals: 99.7%
Iter 3 | Best: 17.524655 | Curr: 431.798088 | Rate: 0.00 | Evals: 99.8%
Iter 4 | Best: 17.524655 | Curr: 53.899066 | Rate: 0.00 | Evals: 99.9%
Iter 5 | Best: 17.524655 | Curr: 2834.826819 | Rate: 0.00 | Evals: 100.0%
Best y (Multi-Surrogate): 17.524655
2. Running Default SpotOptim (GP only)...
TensorBoard logging disabled
Initial best: $f(x) = 23.358705$
Iter 1 | Best: 23.358705 | Curr: 3013.582310 | Rate: 0.00 | Evals: 2.1%
Iter 2 | Best: 23.358705 | Curr: 3143.796348 | Rate: 0.00 | Evals: 2.2%
Iter 3 | Best: 23.358705 | Curr: 2678.952109 | Rate: 0.00 | Evals: 2.3%
Iter 4 | Best: 23.358705 | Curr: 1520.800700 | Rate: 0.00 | Evals: 2.4%
Iter 5 | Best: 23.358705 | Curr: 1695.568549 | Rate: 0.00 | Evals: 2.5%
Iter 6 | Best: 23.358705 | Curr: 3484.518880 | Rate: 0.00 | Evals: 2.6%
Iter 7 | Best: 23.358705 | Curr: 90.073676 | Rate: 0.00 | Evals: 2.7%
Iter 8 | Best: 23.358705 | Curr: 150.319954 | Rate: 0.00 | Evals: 2.8%
Iter 9 | Best: 23.358705 | Curr: 174.105445 | Rate: 0.00 | Evals: 2.9%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x_0 validated and processed successfully.
Original scale: [0.12897446 0.11253871 0.53861201 0.60847174 0.55760895 0.75440905
0.99729245 0.253624 0.50095326 0.4098481]
Internal scale: [0.12897446 0.11253871 0.53861201 0.60847174 0.55760895 0.75440905
0.99729245 0.253624 0.50095326 0.4098481]
Restart injection: Using best found point so far as starting point ($f(x)=23.358705$).
Including 1 starting points from x_0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: $f(x) = 23.358705$
Iter 1 | Best: 23.358705 | Curr: 1450.330578 | Rate: 0.00 | Evals: 5.0%
Iter 2 | Best: 23.358705 | Curr: 478.466260 | Rate: 0.00 | Evals: 5.1%
Iter 3 | Best: 23.358705 | Curr: 1789.325888 | Rate: 0.00 | Evals: 5.2%
Iter 4 | Best: 23.358705 | Curr: 79.030220 | Rate: 0.00 | Evals: 5.3%
Iter 5 | Best: 23.358705 | Curr: 159.629124 | Rate: 0.00 | Evals: 5.4%
Iter 6 | Best: 23.358705 | Curr: 1283.053855 | Rate: 0.00 | Evals: 5.5%
Iter 7 | Best: 23.358705 | Curr: 149.855884 | Rate: 0.00 | Evals: 5.6%
Iter 8 | Best: 23.358705 | Curr: 198.768529 | Rate: 0.00 | Evals: 5.7%
Iter 9 | Best: 23.358705 | Curr: 195.994700 | Rate: 0.00 | Evals: 5.8%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x_0 validated and processed successfully.
Original scale: [0.12897446 0.11253871 0.53861201 0.60847174 0.55760895 0.75440905
0.99729245 0.253624 0.50095326 0.4098481]

76.7. Multi-Surrogate Switching in SpotOptim

Internal scale: [0.12897446 0.11253871 0.53861201 0.60847174 0.55760895 0.75440905
0.99729245 0.253624 0.50095326 0.4098481]

Restart injection: Using best found point so far as starting point ($f(x)=23.358705$).
Including 1 starting points from x_0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: $f(x) = 22.480133$

Iter 1	Best: 22.480133	Curr: 62.880980	Rate: 0.00	Evals: 7.9%
Iter 2	Best: 22.480133	Curr: 73.567488	Rate: 0.00	Evals: 8.0%
Iter 3	Best: 22.480133	Curr: 95.129208	Rate: 0.00	Evals: 8.1%
Iter 4	Best: 22.480133	Curr: 92.465121	Rate: 0.00	Evals: 8.2%
Iter 5	Best: 22.480133	Curr: 31.052463	Rate: 0.00	Evals: 8.3%
Iter 6	Best: 22.480133	Curr: 65.173381	Rate: 0.00	Evals: 8.4%
Iter 7	Best: 22.480133	Curr: 101.746132	Rate: 0.00	Evals: 8.5%
Iter 8	Best: 22.480133	Curr: 3440.048645	Rate: 0.00	Evals: 8.6%
Iter 9	Best: 22.480133	Curr: 84.365207	Rate: 0.00	Evals: 8.7%

Restarting optimization: success_rate 0 for 10 iterations.
Starting point x_0 validated and processed successfully.

Original scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567
0.75120132 0.82059894 0.03724163 0.869633]

Internal scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567
0.75120132 0.82059894 0.03724163 0.869633]

Restart injection: Using best found point so far as starting point ($f(x)=22.480133$).
Including 1 starting points from x_0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: $f(x) = 22.480133$

Iter 1	Best: 22.480133	Curr: 83.761120	Rate: 0.00	Evals: 10.8%
Iter 2	Best: 22.480133	Curr: 4473.655028	Rate: 0.00	Evals: 10.9%
Iter 3	Best: 22.480133	Curr: 9354.144867	Rate: 0.00	Evals: 11.0%
Iter 4	Best: 22.480133	Curr: 322.290239	Rate: 0.00	Evals: 11.1%
Iter 5	Best: 22.480133	Curr: 1310.242902	Rate: 0.00	Evals: 11.2%
Iter 6	Best: 22.480133	Curr: 670.131948	Rate: 0.00	Evals: 11.3%
Iter 7	Best: 22.480133	Curr: 783.391586	Rate: 0.00	Evals: 11.4%
Iter 8	Best: 22.480133	Curr: 383.193813	Rate: 0.00	Evals: 11.5%
Iter 9	Best: 22.480133	Curr: 63.695172	Rate: 0.00	Evals: 11.6%

Restarting optimization: success_rate 0 for 10 iterations.
Starting point x_0 validated and processed successfully.

Original scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567
0.75120132 0.82059894 0.03724163 0.869633]

Internal scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567
0.75120132 0.82059894 0.03724163 0.869633]

Restart injection: Using best found point so far as starting point ($f(x)=22.480133$).
Including 1 starting points from x_0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: $f(x) = 22.480133$

Iter 1	Best: 22.480133	Curr: 5226.054629	Rate: 0.00	Evals: 13.7%
--------	-----------------	-------------------	------------	--------------

76. SpotOptim Execution Flow

```
Iter 2 | Best: 22.480133 | Curr: 267.199476 | Rate: 0.00 | Evals: 13.8%
Iter 3 | Best: 22.480133 | Curr: 68.489628 | Rate: 0.00 | Evals: 13.9%
Iter 4 | Best: 22.480133 | Curr: 598.663117 | Rate: 0.00 | Evals: 14.0%
Iter 5 | Best: 22.480133 | Curr: 864.010052 | Rate: 0.00 | Evals: 14.1%
Iter 6 | Best: 22.480133 | Curr: 247.702155 | Rate: 0.00 | Evals: 14.2%
Iter 7 | Best: 22.480133 | Curr: 1854.589202 | Rate: 0.00 | Evals: 14.3%
Iter 8 | Best: 22.480133 | Curr: 43.047980 | Rate: 0.00 | Evals: 14.4%
Iter 9 | Best: 22.480133 | Curr: 142.173696 | Rate: 0.00 | Evals: 14.5%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567
0.75120132 0.82059894 0.03724163 0.869633 ]
  Internal scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567
0.75120132 0.82059894 0.03724163 0.869633 ]
Restart injection: Using best found point so far as starting point (f(x)=22.480133).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 22.480133
Iter 1 | Best: 22.480133 | Curr: 696.009525 | Rate: 0.00 | Evals: 16.6%
Iter 2 | Best: 22.480133 | Curr: 1556.753437 | Rate: 0.00 | Evals: 16.7%
Iter 3 | Best: 22.480133 | Curr: 1947.423614 | Rate: 0.00 | Evals: 16.8%
Iter 4 | Best: 22.480133 | Curr: 2389.571978 | Rate: 0.00 | Evals: 16.9%
Iter 5 | Best: 22.480133 | Curr: 159.807210 | Rate: 0.00 | Evals: 17.0%
Iter 6 | Best: 22.480133 | Curr: 134.043080 | Rate: 0.00 | Evals: 17.1%
Iter 7 | Best: 22.480133 | Curr: 170.530168 | Rate: 0.00 | Evals: 17.2%
Iter 8 | Best: 22.480133 | Curr: 2880.020911 | Rate: 0.00 | Evals: 17.3%
Iter 9 | Best: 22.480133 | Curr: 210.862794 | Rate: 0.00 | Evals: 17.4%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567
0.75120132 0.82059894 0.03724163 0.869633 ]
  Internal scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567
0.75120132 0.82059894 0.03724163 0.869633 ]
Restart injection: Using best found point so far as starting point (f(x)=22.480133).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 22.480133
Iter 1 | Best: 22.480133 | Curr: 4175.379221 | Rate: 0.00 | Evals: 19.5%
Iter 2 | Best: 22.480133 | Curr: 2256.719140 | Rate: 0.00 | Evals: 19.6%
Iter 3 | Best: 22.480133 | Curr: 1469.835766 | Rate: 0.00 | Evals: 19.7%
Iter 4 | Best: 22.480133 | Curr: 73.857789 | Rate: 0.00 | Evals: 19.8%
Iter 5 | Best: 22.480133 | Curr: 56.743810 | Rate: 0.00 | Evals: 19.9%
Iter 6 | Best: 22.480133 | Curr: 67.987361 | Rate: 0.00 | Evals: 20.0%
Iter 7 | Best: 22.480133 | Curr: 110.259506 | Rate: 0.00 | Evals: 20.1%
Iter 8 | Best: 22.480133 | Curr: 220.182373 | Rate: 0.00 | Evals: 20.2%
```

76.7. Multi-Surrogate Switching in SpotOptim

Iter 9 | Best: 22.480133 | Curr: 216.667589 | Rate: 0.00 | Evals: 20.3%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567
0.75120132 0.82059894 0.03724163 0.869633]
Internal scale: [0.77583701 0.50025469 0.15778035 0.71883974 0.37284592 0.11704567
0.75120132 0.82059894 0.03724163 0.869633]
Restart injection: Using best found point so far as starting point (f(x)=22.480133).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 19.382131
Iter 1 | Best: 19.382131 | Curr: 2489.714100 | Rate: 0.00 | Evals: 22.4%
Iter 2 | Best: 19.382131 | Curr: 7337.744425 | Rate: 0.00 | Evals: 22.5%
Iter 3 | Best: 19.382131 | Curr: 394.603371 | Rate: 0.00 | Evals: 22.6%
Iter 4 | Best: 19.382131 | Curr: 2281.268220 | Rate: 0.00 | Evals: 22.7%
Iter 5 | Best: 19.382131 | Curr: 171.644430 | Rate: 0.00 | Evals: 22.8%
Iter 6 | Best: 19.382131 | Curr: 926.478530 | Rate: 0.00 | Evals: 22.9%
Iter 7 | Best: 19.382131 | Curr: 222.485683 | Rate: 0.00 | Evals: 23.0%
Iter 8 | Best: 19.382131 | Curr: 3869.142885 | Rate: 0.00 | Evals: 23.1%
Iter 9 | Best: 19.382131 | Curr: 408.350507 | Rate: 0.00 | Evals: 23.2%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.82762986 0.23097536 0.68598884 0.96544327 0.77713053 0.65822737
0.84033887 0.72758654 0.77639412 0.34639732]
Internal scale: [0.82762986 0.23097536 0.68598884 0.96544327 0.77713053 0.65822737
0.84033887 0.72758654 0.77639412 0.34639732]
Restart injection: Using best found point so far as starting point (f(x)=19.382131).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 18.171750
Iter 1 | Best: 18.171750 | Curr: 2209.164844 | Rate: 0.00 | Evals: 25.3%
Iter 2 | Best: 18.171750 | Curr: 2922.590435 | Rate: 0.00 | Evals: 25.4%
Iter 3 | Best: 18.171750 | Curr: 219.614553 | Rate: 0.00 | Evals: 25.5%
Iter 4 | Best: 18.171750 | Curr: 141.039762 | Rate: 0.00 | Evals: 25.6%
Iter 5 | Best: 18.171750 | Curr: 126.975941 | Rate: 0.00 | Evals: 25.7%
Iter 6 | Best: 18.171750 | Curr: 195.724835 | Rate: 0.00 | Evals: 25.8%
Iter 7 | Best: 18.171750 | Curr: 149.802841 | Rate: 0.00 | Evals: 25.9%
Iter 8 | Best: 18.171750 | Curr: 138.026558 | Rate: 0.00 | Evals: 26.0%
Iter 9 | Best: 18.171750 | Curr: 170.418401 | Rate: 0.00 | Evals: 26.1%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.65238644 0.60930271 0.03433355 0.72954934 0.93074097 0.25509536
0.98995351 0.84219394 0.52650099 0.3756625]
Internal scale: [0.65238644 0.60930271 0.03433355 0.72954934 0.93074097 0.25509536
0.98995351 0.84219394 0.52650099 0.3756625]

76. SpotOptim Execution Flow

Restart injection: Using best found point so far as starting point ($f(x)=18.171750$).
Including 1 starting points from x_0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: $f(x) = 18.171750$

Iter 1	Best: 18.171750	Curr: 195.878495	Rate: 0.00	Evals: 28.2%
Iter 2	Best: 18.171750	Curr: 207.357393	Rate: 0.00	Evals: 28.3%
Iter 3	Best: 18.171750	Curr: 2056.300039	Rate: 0.00	Evals: 28.4%
Iter 4	Best: 18.171750	Curr: 209.699030	Rate: 0.00	Evals: 28.5%
Iter 5	Best: 18.171750	Curr: 192.850643	Rate: 0.00	Evals: 28.6%
Iter 6	Best: 18.171750	Curr: 171.333960	Rate: 0.00	Evals: 28.7%
Iter 7	Best: 18.171750	Curr: 201.113992	Rate: 0.00	Evals: 28.8%
Iter 8	Best: 18.171750	Curr: 199.489816	Rate: 0.00	Evals: 28.9%
Iter 9	Best: 18.171750	Curr: 83.460894	Rate: 0.00	Evals: 29.0%

Restarting optimization: success_rate 0 for 10 iterations.
Starting point x_0 validated and processed successfully.
Original scale: [0.65238644 0.60930271 0.03433355 0.72954934 0.93074097 0.25509536 0.98995351 0.84219394 0.52650099 0.3756625]
Internal scale: [0.65238644 0.60930271 0.03433355 0.72954934 0.93074097 0.25509536 0.98995351 0.84219394 0.52650099 0.3756625]

Restart injection: Using best found point so far as starting point ($f(x)=18.171750$).
Including 1 starting points from x_0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: $f(x) = 18.171750$

Iter 1	Best: 18.171750	Curr: 1121.761518	Rate: 0.00	Evals: 31.1%
Iter 2	Best: 18.171750	Curr: 175.207222	Rate: 0.00	Evals: 31.2%
Iter 3	Best: 18.171750	Curr: 2482.001387	Rate: 0.00	Evals: 31.3%
Iter 4	Best: 18.171750	Curr: 2959.220827	Rate: 0.00	Evals: 31.4%
Iter 5	Best: 18.171750	Curr: 57.401547	Rate: 0.00	Evals: 31.5%
Iter 6	Best: 18.171750	Curr: 143.928700	Rate: 0.00	Evals: 31.6%
Iter 7	Best: 18.171750	Curr: 2230.975481	Rate: 0.00	Evals: 31.7%
Iter 8	Best: 18.171750	Curr: 5993.535411	Rate: 0.00	Evals: 31.8%
Iter 9	Best: 18.171750	Curr: 190.002577	Rate: 0.00	Evals: 31.9%

Restarting optimization: success_rate 0 for 10 iterations.
Starting point x_0 validated and processed successfully.
Original scale: [0.65238644 0.60930271 0.03433355 0.72954934 0.93074097 0.25509536 0.98995351 0.84219394 0.52650099 0.3756625]
Internal scale: [0.65238644 0.60930271 0.03433355 0.72954934 0.93074097 0.25509536 0.98995351 0.84219394 0.52650099 0.3756625]

Restart injection: Using best found point so far as starting point ($f(x)=18.171750$).
Including 1 starting points from x_0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: $f(x) = 17.988188$

Iter 1	Best: 17.988188	Curr: 852.147217	Rate: 0.00	Evals: 34.0%
Iter 2	Best: 17.988188	Curr: 564.802432	Rate: 0.00	Evals: 34.1%
Iter 3	Best: 17.988188	Curr: 877.100221	Rate: 0.00	Evals: 34.2%

76.7. Multi-Surrogate Switching in SpotOptim

Iter 4 | Best: 17.988188 | Curr: 243.357391 | Rate: 0.00 | Evals: 34.3%
Iter 5 | Best: 17.988188 | Curr: 2204.061616 | Rate: 0.00 | Evals: 34.4%
Iter 6 | Best: 17.988188 | Curr: 629.452237 | Rate: 0.00 | Evals: 34.5%
Iter 7 | Best: 17.988188 | Curr: 1490.338508 | Rate: 0.00 | Evals: 34.6%
Iter 8 | Best: 17.988188 | Curr: 3481.842329 | Rate: 0.00 | Evals: 34.7%
Iter 9 | Best: 17.988188 | Curr: 1933.138417 | Rate: 0.00 | Evals: 34.8%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
Iter 1 | Best: 17.988188 | Curr: 27.480679 | Rate: 0.00 | Evals: 36.9%
Iter 2 | Best: 17.988188 | Curr: 46.269769 | Rate: 0.00 | Evals: 37.0%
Iter 3 | Best: 17.988188 | Curr: 2787.079856 | Rate: 0.00 | Evals: 37.1%
Iter 4 | Best: 17.988188 | Curr: 2680.068053 | Rate: 0.00 | Evals: 37.2%
Iter 5 | Best: 17.988188 | Curr: 48.174325 | Rate: 0.00 | Evals: 37.3%
Iter 6 | Best: 17.988188 | Curr: 53.145760 | Rate: 0.00 | Evals: 37.4%
Iter 7 | Best: 17.988188 | Curr: 48.348795 | Rate: 0.00 | Evals: 37.5%
Iter 8 | Best: 17.988188 | Curr: 20.461050 | Rate: 0.00 | Evals: 37.6%
Iter 9 | Best: 17.988188 | Curr: 21.238648 | Rate: 0.00 | Evals: 37.7%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
Iter 1 | Best: 17.988188 | Curr: 1825.214975 | Rate: 0.00 | Evals: 39.8%
Iter 2 | Best: 17.988188 | Curr: 2549.641122 | Rate: 0.00 | Evals: 39.9%
Iter 3 | Best: 17.988188 | Curr: 72.789455 | Rate: 0.00 | Evals: 40.0%
Iter 4 | Best: 17.988188 | Curr: 2718.188283 | Rate: 0.00 | Evals: 40.1%
Iter 5 | Best: 17.988188 | Curr: 407.947214 | Rate: 0.00 | Evals: 40.2%
Iter 6 | Best: 17.988188 | Curr: 60.047906 | Rate: 0.00 | Evals: 40.3%
Iter 7 | Best: 17.988188 | Curr: 3727.478000 | Rate: 0.00 | Evals: 40.4%
Iter 8 | Best: 17.988188 | Curr: 1961.138592 | Rate: 0.00 | Evals: 40.5%
Iter 9 | Best: 17.988188 | Curr: 93.350112 | Rate: 0.00 | Evals: 40.6%
Restarting optimization: success_rate 0 for 10 iterations.

76. SpotOptim Execution Flow

Starting point x0 validated and processed successfully.

Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]

Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]

Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.

Skipping re-evaluation of injected best point.

Initial best: f(x) = 17.988188

Iter 1 | Best: 17.988188 | Curr: 2773.745889 | Rate: 0.00 | Evals: 42.7%

Iter 2 | Best: 17.988188 | Curr: 4463.572454 | Rate: 0.00 | Evals: 42.8%

Iter 3 | Best: 17.988188 | Curr: 1303.492669 | Rate: 0.00 | Evals: 42.9%

Iter 4 | Best: 17.988188 | Curr: 25.852038 | Rate: 0.00 | Evals: 43.0%

Iter 5 | Best: 17.988188 | Curr: 41.238297 | Rate: 0.00 | Evals: 43.1%

Iter 6 | Best: 17.988188 | Curr: 37.438776 | Rate: 0.00 | Evals: 43.2%

Iter 7 | Best: 17.988188 | Curr: 3979.434441 | Rate: 0.00 | Evals: 43.3%

Iter 8 | Best: 17.988188 | Curr: 76.199673 | Rate: 0.00 | Evals: 43.4%

Iter 9 | Best: 17.988188 | Curr: 132.283785 | Rate: 0.00 | Evals: 43.5%

Restarting optimization: success_rate 0 for 10 iterations.

Starting point x0 validated and processed successfully.

Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]

Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]

Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.

Skipping re-evaluation of injected best point.

Initial best: f(x) = 17.988188

Iter 1 | Best: 17.988188 | Curr: 146.373690 | Rate: 0.00 | Evals: 45.6%

Iter 2 | Best: 17.988188 | Curr: 146.796739 | Rate: 0.00 | Evals: 45.7%

Iter 3 | Best: 17.988188 | Curr: 62.838682 | Rate: 0.00 | Evals: 45.8%

Iter 4 | Best: 17.988188 | Curr: 198.356719 | Rate: 0.00 | Evals: 45.9%

Iter 5 | Best: 17.988188 | Curr: 158.002267 | Rate: 0.00 | Evals: 46.0%

Iter 6 | Best: 17.988188 | Curr: 173.117569 | Rate: 0.00 | Evals: 46.1%

Iter 7 | Best: 17.988188 | Curr: 186.882895 | Rate: 0.00 | Evals: 46.2%

Iter 8 | Best: 17.988188 | Curr: 178.809070 | Rate: 0.00 | Evals: 46.3%

Iter 9 | Best: 17.988188 | Curr: 175.931442 | Rate: 0.00 | Evals: 46.4%

Restarting optimization: success_rate 0 for 10 iterations.

Starting point x0 validated and processed successfully.

Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]

Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]

Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.

76.7. Multi-Surrogate Switching in SpotOptim

Skipping re-evaluation of injected best point.

Initial best: $f(x) = 17.988188$

Iter 1		Best: 17.988188		Curr: 2198.225906		Rate: 0.00		Evals: 48.5%
Iter 2		Best: 17.988188		Curr: 60.211946		Rate: 0.00		Evals: 48.6%
Iter 3		Best: 17.988188		Curr: 1241.808367		Rate: 0.00		Evals: 48.7%
Iter 4		Best: 17.988188		Curr: 52.516664		Rate: 0.00		Evals: 48.8%
Iter 5		Best: 17.988188		Curr: 58.303265		Rate: 0.00		Evals: 48.9%
Iter 6		Best: 17.988188		Curr: 2868.000550		Rate: 0.00		Evals: 49.0%
Iter 7		Best: 17.988188		Curr: 62.588314		Rate: 0.00		Evals: 49.1%
Iter 8		Best: 17.988188		Curr: 46.379883		Rate: 0.00		Evals: 49.2%
Iter 9		Best: 17.988188		Curr: 73.451372		Rate: 0.00		Evals: 49.3%

Restarting optimization: success_rate 0 for 10 iterations.

Starting point x0 validated and processed successfully.

Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856 0.34800416 0.3285074 0.40791185 0.64510118]

Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856 0.34800416 0.3285074 0.40791185 0.64510118]

Restart injection: Using best found point so far as starting point ($f(x)=17.988188$).

Including 1 starting points from x0 in initial design.

Skipping re-evaluation of injected best point.

Initial best: $f(x) = 17.988188$

Iter 1		Best: 17.988188		Curr: 2853.658131		Rate: 0.00		Evals: 51.4%
Iter 2		Best: 17.988188		Curr: 2242.829031		Rate: 0.00		Evals: 51.5%
Iter 3		Best: 17.988188		Curr: 168.051726		Rate: 0.00		Evals: 51.6%
Iter 4		Best: 17.988188		Curr: 191.604441		Rate: 0.00		Evals: 51.7%
Iter 5		Best: 17.988188		Curr: 220.533924		Rate: 0.00		Evals: 51.8%
Iter 6		Best: 17.988188		Curr: 4065.085485		Rate: 0.00		Evals: 51.9%
Iter 7		Best: 17.988188		Curr: 188.401813		Rate: 0.00		Evals: 52.0%
Iter 8		Best: 17.988188		Curr: 154.066332		Rate: 0.00		Evals: 52.1%
Iter 9		Best: 17.988188		Curr: 185.483402		Rate: 0.00		Evals: 52.2%

Restarting optimization: success_rate 0 for 10 iterations.

Starting point x0 validated and processed successfully.

Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856 0.34800416 0.3285074 0.40791185 0.64510118]

Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856 0.34800416 0.3285074 0.40791185 0.64510118]

Restart injection: Using best found point so far as starting point ($f(x)=17.988188$).

Including 1 starting points from x0 in initial design.

Skipping re-evaluation of injected best point.

Initial best: $f(x) = 17.988188$

Iter 1		Best: 17.988188		Curr: 5550.286749		Rate: 0.00		Evals: 54.3%
Iter 2		Best: 17.988188		Curr: 6021.453645		Rate: 0.00		Evals: 54.4%
Iter 3		Best: 17.988188		Curr: 688.338058		Rate: 0.00		Evals: 54.5%
Iter 4		Best: 17.988188		Curr: 3485.042969		Rate: 0.00		Evals: 54.6%
Iter 5		Best: 17.988188		Curr: 80.518743		Rate: 0.00		Evals: 54.7%

76. SpotOptim Execution Flow

```
Iter 6 | Best: 17.988188 | Curr: 20.626949 | Rate: 0.00 | Evals: 54.8%
Iter 7 | Best: 17.988188 | Curr: 100.579789 | Rate: 0.00 | Evals: 54.9%
Iter 8 | Best: 17.988188 | Curr: 497.440238 | Rate: 0.00 | Evals: 55.0%
Iter 9 | Best: 17.988188 | Curr: 1343.715335 | Rate: 0.00 | Evals: 55.1%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
  Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
Iter 1 | Best: 17.988188 | Curr: 591.491331 | Rate: 0.00 | Evals: 57.2%
Iter 2 | Best: 17.988188 | Curr: 211.173261 | Rate: 0.00 | Evals: 57.3%
Iter 3 | Best: 17.988188 | Curr: 1519.526479 | Rate: 0.00 | Evals: 57.4%
Iter 4 | Best: 17.988188 | Curr: 187.512161 | Rate: 0.00 | Evals: 57.5%
Iter 5 | Best: 17.988188 | Curr: 25.312220 | Rate: 0.00 | Evals: 57.6%
Iter 6 | Best: 17.988188 | Curr: 172.918242 | Rate: 0.00 | Evals: 57.7%
Iter 7 | Best: 17.988188 | Curr: 864.681360 | Rate: 0.00 | Evals: 57.8%
Iter 8 | Best: 17.988188 | Curr: 221.122501 | Rate: 0.00 | Evals: 57.9%
Iter 9 | Best: 17.988188 | Curr: 3451.540391 | Rate: 0.00 | Evals: 58.0%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
  Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
Iter 1 | Best: 17.988188 | Curr: 777.996129 | Rate: 0.00 | Evals: 60.1%
Iter 2 | Best: 17.988188 | Curr: 181.588998 | Rate: 0.00 | Evals: 60.2%
Iter 3 | Best: 17.988188 | Curr: 92.858567 | Rate: 0.00 | Evals: 60.3%
Iter 4 | Best: 17.988188 | Curr: 207.378723 | Rate: 0.00 | Evals: 60.4%
Iter 5 | Best: 17.988188 | Curr: 185.223873 | Rate: 0.00 | Evals: 60.5%
Iter 6 | Best: 17.988188 | Curr: 40.417783 | Rate: 0.00 | Evals: 60.6%
Iter 7 | Best: 17.988188 | Curr: 79.395958 | Rate: 0.00 | Evals: 60.7%
Iter 8 | Best: 17.988188 | Curr: 43.973869 | Rate: 0.00 | Evals: 60.8%
Iter 9 | Best: 17.988188 | Curr: 48.067173 | Rate: 0.00 | Evals: 60.9%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
```


76.7. Multi-Surrogate Switching in SpotOptim

```
0.34800416 0.3285074 0.40791185 0.64510118]
  Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
Iter 1 | Best: 17.988188 | Curr: 7141.230028 | Rate: 0.00 | Evals: 63.0%
Iter 2 | Best: 17.988188 | Curr: 195.924790 | Rate: 0.00 | Evals: 63.1%
Iter 3 | Best: 17.988188 | Curr: 1122.073996 | Rate: 0.00 | Evals: 63.2%
Iter 4 | Best: 17.988188 | Curr: 193.650688 | Rate: 0.00 | Evals: 63.3%
Iter 5 | Best: 17.988188 | Curr: 177.720539 | Rate: 0.00 | Evals: 63.4%
Iter 6 | Best: 17.988188 | Curr: 160.365697 | Rate: 0.00 | Evals: 63.5%
Iter 7 | Best: 17.988188 | Curr: 150.965501 | Rate: 0.00 | Evals: 63.6%
Iter 8 | Best: 17.988188 | Curr: 181.023097 | Rate: 0.00 | Evals: 63.7%
Iter 9 | Best: 17.988188 | Curr: 172.727634 | Rate: 0.00 | Evals: 63.8%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
  Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
Iter 1 | Best: 17.988188 | Curr: 7096.195754 | Rate: 0.00 | Evals: 65.9%
Iter 2 | Best: 17.988188 | Curr: 840.570735 | Rate: 0.00 | Evals: 66.0%
Iter 3 | Best: 17.988188 | Curr: 713.866983 | Rate: 0.00 | Evals: 66.1%
Iter 4 | Best: 17.988188 | Curr: 2596.884950 | Rate: 0.00 | Evals: 66.2%
Iter 5 | Best: 17.988188 | Curr: 151.463473 | Rate: 0.00 | Evals: 66.3%
Iter 6 | Best: 17.988188 | Curr: 52.226965 | Rate: 0.00 | Evals: 66.4%
Iter 7 | Best: 17.988188 | Curr: 134.728341 | Rate: 0.00 | Evals: 66.5%
Iter 8 | Best: 17.988188 | Curr: 2220.323045 | Rate: 0.00 | Evals: 66.6%
Iter 9 | Best: 17.988188 | Curr: 156.076630 | Rate: 0.00 | Evals: 66.7%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
  Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
```

76. SpotOptim Execution Flow

```
Iter 1 | Best: 17.988188 | Curr: 151.860156 | Rate: 0.00 | Evals: 68.8%
Iter 2 | Best: 17.988188 | Curr: 119.044755 | Rate: 0.00 | Evals: 68.9%
Iter 3 | Best: 17.988188 | Curr: 1611.117179 | Rate: 0.00 | Evals: 69.0%
Iter 4 | Best: 17.988188 | Curr: 155.451101 | Rate: 0.00 | Evals: 69.1%
Iter 5 | Best: 17.988188 | Curr: 123.567532 | Rate: 0.00 | Evals: 69.2%
Iter 6 | Best: 17.988188 | Curr: 166.691711 | Rate: 0.00 | Evals: 69.3%
Iter 7 | Best: 17.988188 | Curr: 101.713550 | Rate: 0.00 | Evals: 69.4%
Iter 8 | Best: 17.988188 | Curr: 59.441433 | Rate: 0.00 | Evals: 69.5%
Iter 9 | Best: 17.988188 | Curr: 94.200203 | Rate: 0.00 | Evals: 69.6%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
  Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
Iter 1 | Best: 17.988188 | Curr: 149.460119 | Rate: 0.00 | Evals: 71.7%
Iter 2 | Best: 17.988188 | Curr: 193.969893 | Rate: 0.00 | Evals: 71.8%
Iter 3 | Best: 17.988188 | Curr: 175.553259 | Rate: 0.00 | Evals: 71.9%
Iter 4 | Best: 17.988188 | Curr: 178.432348 | Rate: 0.00 | Evals: 72.0%
Iter 5 | Best: 17.988188 | Curr: 74.404462 | Rate: 0.00 | Evals: 72.1%
Iter 6 | Best: 17.988188 | Curr: 2169.369773 | Rate: 0.00 | Evals: 72.2%
Iter 7 | Best: 17.988188 | Curr: 57.079803 | Rate: 0.00 | Evals: 72.3%
Iter 8 | Best: 17.988188 | Curr: 139.383264 | Rate: 0.00 | Evals: 72.4%
Iter 9 | Best: 17.988188 | Curr: 79.518634 | Rate: 0.00 | Evals: 72.5%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
  Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
Iter 1 | Best: 17.988188 | Curr: 63.809655 | Rate: 0.00 | Evals: 74.6%
Iter 2 | Best: 17.988188 | Curr: 70.209822 | Rate: 0.00 | Evals: 74.7%
Iter 3 | Best: 17.988188 | Curr: 4512.662643 | Rate: 0.00 | Evals: 74.8%
Iter 4 | Best: 17.988188 | Curr: 2237.763357 | Rate: 0.00 | Evals: 74.9%
Iter 5 | Best: 17.988188 | Curr: 3518.229469 | Rate: 0.00 | Evals: 75.0%
Iter 6 | Best: 17.988188 | Curr: 118.595336 | Rate: 0.00 | Evals: 75.1%
Iter 7 | Best: 17.988188 | Curr: 120.463702 | Rate: 0.00 | Evals: 75.2%
```

76.7. Multi-Surrogate Switching in SpotOptim

Iter 8 | Best: 17.988188 | Curr: 67.179019 | Rate: 0.00 | Evals: 75.3%
Iter 9 | Best: 17.988188 | Curr: 51.977293 | Rate: 0.00 | Evals: 75.4%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
Iter 1 | Best: 17.988188 | Curr: 3034.176710 | Rate: 0.00 | Evals: 77.5%
Iter 2 | Best: 17.988188 | Curr: 4683.846226 | Rate: 0.00 | Evals: 77.6%
Iter 3 | Best: 17.988188 | Curr: 2949.664204 | Rate: 0.00 | Evals: 77.7%
Iter 4 | Best: 17.988188 | Curr: 529.625953 | Rate: 0.00 | Evals: 77.8%
Iter 5 | Best: 17.988188 | Curr: 4644.335129 | Rate: 0.00 | Evals: 77.9%
Iter 6 | Best: 17.988188 | Curr: 119.398623 | Rate: 0.00 | Evals: 78.0%
Iter 7 | Best: 17.988188 | Curr: 7614.879298 | Rate: 0.00 | Evals: 78.1%
Iter 8 | Best: 17.988188 | Curr: 3736.113700 | Rate: 0.00 | Evals: 78.2%
Iter 9 | Best: 17.988188 | Curr: 573.264799 | Rate: 0.00 | Evals: 78.3%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
Iter 1 | Best: 17.988188 | Curr: 3930.676660 | Rate: 0.00 | Evals: 80.4%
Iter 2 | Best: 17.988188 | Curr: 232.434667 | Rate: 0.00 | Evals: 80.5%
Iter 3 | Best: 17.988188 | Curr: 2711.224258 | Rate: 0.00 | Evals: 80.6%
Iter 4 | Best: 17.988188 | Curr: 3068.358370 | Rate: 0.00 | Evals: 80.7%
Iter 5 | Best: 17.988188 | Curr: 226.185037 | Rate: 0.00 | Evals: 80.8%
Iter 6 | Best: 17.988188 | Curr: 168.609846 | Rate: 0.00 | Evals: 80.9%
Iter 7 | Best: 17.988188 | Curr: 243.482707 | Rate: 0.00 | Evals: 81.0%
Iter 8 | Best: 17.988188 | Curr: 196.687525 | Rate: 0.00 | Evals: 81.1%
Iter 9 | Best: 17.988188 | Curr: 182.556546 | Rate: 0.00 | Evals: 81.2%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856

76. SpotOptim Execution Flow

```
0.34800416 0.3285074 0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
Iter 1 | Best: 17.988188 | Curr: 100.441252 | Rate: 0.00 | Evals: 83.3%
Iter 2 | Best: 17.988188 | Curr: 88.192975 | Rate: 0.00 | Evals: 83.4%
Iter 3 | Best: 17.988188 | Curr: 90.503345 | Rate: 0.00 | Evals: 83.5%
Iter 4 | Best: 17.988188 | Curr: 69.631755 | Rate: 0.00 | Evals: 83.6%
Iter 5 | Best: 17.988188 | Curr: 2782.577463 | Rate: 0.00 | Evals: 83.7%
Iter 6 | Best: 17.988188 | Curr: 3691.717233 | Rate: 0.00 | Evals: 83.8%
Iter 7 | Best: 17.988188 | Curr: 2657.043534 | Rate: 0.00 | Evals: 83.9%
Iter 8 | Best: 17.988188 | Curr: 113.031169 | Rate: 0.00 | Evals: 84.0%
Iter 9 | Best: 17.988188 | Curr: 96.965716 | Rate: 0.00 | Evals: 84.1%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
Iter 1 | Best: 17.988188 | Curr: 2657.071381 | Rate: 0.00 | Evals: 86.2%
Iter 2 | Best: 17.988188 | Curr: 42.525511 | Rate: 0.00 | Evals: 86.3%
Iter 3 | Best: 17.988188 | Curr: 191.567921 | Rate: 0.00 | Evals: 86.4%
Iter 4 | Best: 17.988188 | Curr: 44.911561 | Rate: 0.00 | Evals: 86.5%
Iter 5 | Best: 17.988188 | Curr: 175.733952 | Rate: 0.00 | Evals: 86.6%
Iter 6 | Best: 17.988188 | Curr: 124.787192 | Rate: 0.00 | Evals: 86.7%
Iter 7 | Best: 17.988188 | Curr: 146.714298 | Rate: 0.00 | Evals: 86.8%
Iter 8 | Best: 17.988188 | Curr: 205.011924 | Rate: 0.00 | Evals: 86.9%
Iter 9 | Best: 17.988188 | Curr: 694.079979 | Rate: 0.00 | Evals: 87.0%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
Iter 1 | Best: 17.988188 | Curr: 5040.684152 | Rate: 0.00 | Evals: 89.1%
Iter 2 | Best: 17.988188 | Curr: 1779.029711 | Rate: 0.00 | Evals: 89.2%
```

76.7. Multi-Surrogate Switching in SpotOptim

```
Iter 3 | Best: 17.988188 | Curr: 29.599194 | Rate: 0.00 | Evals: 89.3%
Iter 4 | Best: 17.988188 | Curr: 2310.952329 | Rate: 0.00 | Evals: 89.4%
Iter 5 | Best: 17.988188 | Curr: 5634.753301 | Rate: 0.00 | Evals: 89.5%
Iter 6 | Best: 17.988188 | Curr: 112.190367 | Rate: 0.00 | Evals: 89.6%
Iter 7 | Best: 17.988188 | Curr: 117.829421 | Rate: 0.00 | Evals: 89.7%
Iter 8 | Best: 17.988188 | Curr: 87.767625 | Rate: 0.00 | Evals: 89.8%
Iter 9 | Best: 17.988188 | Curr: 33.656731 | Rate: 0.00 | Evals: 89.9%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.84733651 0.01645646 0.0075814  0.06164297 0.86257478 0.74152856
0.34800416 0.3285074  0.40791185 0.64510118]
  Internal scale: [0.84733651 0.01645646 0.0075814  0.06164297 0.86257478 0.74152856
0.34800416 0.3285074  0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
Iter 1 | Best: 17.988188 | Curr: 198.474002 | Rate: 0.00 | Evals: 92.0%
Iter 2 | Best: 17.988188 | Curr: 65.248842 | Rate: 0.00 | Evals: 92.1%
Iter 3 | Best: 17.988188 | Curr: 144.627256 | Rate: 0.00 | Evals: 92.2%
Iter 4 | Best: 17.988188 | Curr: 514.118987 | Rate: 0.00 | Evals: 92.3%
Iter 5 | Best: 17.988188 | Curr: 162.030109 | Rate: 0.00 | Evals: 92.4%
Iter 6 | Best: 17.988188 | Curr: 38.813512 | Rate: 0.00 | Evals: 92.5%
Iter 7 | Best: 17.988188 | Curr: 118.144419 | Rate: 0.00 | Evals: 92.6%
Iter 8 | Best: 17.988188 | Curr: 115.527647 | Rate: 0.00 | Evals: 92.7%
Iter 9 | Best: 17.988188 | Curr: 122.607606 | Rate: 0.00 | Evals: 92.8%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.84733651 0.01645646 0.0075814  0.06164297 0.86257478 0.74152856
0.34800416 0.3285074  0.40791185 0.64510118]
  Internal scale: [0.84733651 0.01645646 0.0075814  0.06164297 0.86257478 0.74152856
0.34800416 0.3285074  0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
Iter 1 | Best: 17.988188 | Curr: 6218.958105 | Rate: 0.00 | Evals: 94.9%
Iter 2 | Best: 17.988188 | Curr: 1980.642477 | Rate: 0.00 | Evals: 95.0%
Iter 3 | Best: 17.988188 | Curr: 7320.199882 | Rate: 0.00 | Evals: 95.1%
Iter 4 | Best: 17.988188 | Curr: 3023.182600 | Rate: 0.00 | Evals: 95.2%
Iter 5 | Best: 17.988188 | Curr: 193.218741 | Rate: 0.00 | Evals: 95.3%
Iter 6 | Best: 17.988188 | Curr: 354.432689 | Rate: 0.00 | Evals: 95.4%
Iter 7 | Best: 17.988188 | Curr: 19.728206 | Rate: 0.00 | Evals: 95.5%
Iter 8 | Best: 17.988188 | Curr: 179.194727 | Rate: 0.00 | Evals: 95.6%
Iter 9 | Best: 17.988188 | Curr: 67.919629 | Rate: 0.00 | Evals: 95.7%
```

76. SpotOptim Execution Flow

```

Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
  Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Including 1 starting points from x0 in initial design.
Skipping re-evaluation of injected best point.
Initial best: f(x) = 17.988188
Iter 1 | Best: 17.988188 | Curr: 110.478749 | Rate: 0.00 | Evals: 97.8%
Iter 2 | Best: 17.988188 | Curr: 5546.119725 | Rate: 0.00 | Evals: 97.9%
Iter 3 | Best: 17.988188 | Curr: 2392.083263 | Rate: 0.00 | Evals: 98.0%
Iter 4 | Best: 17.988188 | Curr: 139.774509 | Rate: 0.00 | Evals: 98.1%
Iter 5 | Best: 17.988188 | Curr: 3757.731665 | Rate: 0.00 | Evals: 98.2%
Iter 6 | Best: 17.988188 | Curr: 102.201540 | Rate: 0.00 | Evals: 98.3%
Iter 7 | Best: 17.988188 | Curr: 102.068295 | Rate: 0.00 | Evals: 98.4%
Iter 8 | Best: 17.988188 | Curr: 828.574808 | Rate: 0.00 | Evals: 98.5%
Iter 9 | Best: 17.988188 | Curr: 58.374429 | Rate: 0.00 | Evals: 98.6%
Restarting optimization: success_rate 0 for 10 iterations.
Starting point x0 validated and processed successfully.
  Original scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
  Internal scale: [0.84733651 0.01645646 0.0075814 0.06164297 0.86257478 0.74152856
0.34800416 0.3285074 0.40791185 0.64510118]
Restart injection: Using best found point so far as starting point (f(x)=17.988188).
Global budget exhausted. Stopping restarts.
Best y (Default): 17.988188

```

--- Results Summary ---

Method	Best Objective
Default (GP Only)	17.988188
Multi-Surrogate (GP+MLP)	17.524655

Part XII.

Optimization

77. Aircraft Wing Weight Example

i Note

- This section is based on chapter 1.3 “A ten-variable weight function” in Forrester et al. (2008).
- The following Python packages are imported:

```
import math
import matplotlib.pyplot as plt
import numpy as np
from mpl_toolkits.axes_grid1 import ImageGrid
```

77.1. AWWWE Equation

- Example from Forrester et al. (2008)
- Understand the **weight** of an unpainted light aircraft wing as a function of nine design and operational parameters:

$$W = 0.036 S_W^{0.758} \times W_{fw}^{0.0035} \left(\frac{A}{\cos^2 \Lambda} \right)^{0.6} \times q^{0.006} \times \lambda^{0.04} \\ \times \left(\frac{100 R_{tc}}{\cos \Lambda} \right)^{-0.3} \times (N_z W_{dg})^{0.49}$$

77.2. AWWWE Parameters and Equations (Part 1)

Table 77.1.: Aircraft Wing Weight Parameters

Symbol	Parameter	Baseline	Mini- mum	Maxi- mum
S_W	Wing area (ft^2)	174	150	200
W_{fw}	Weight of fuel in wing (lb)	252	220	300

77. Aircraft Wing Weight Example

Symbol	Parameter	Baseline	Minimum	Maximum
A	Aspect ratio	7.52	6	10
Λ	Quarter-chord sweep (deg)	0	-10	10
q	Dynamic pressure at cruise (lb/ft^2)	34	16	45
λ	Taper ratio	0.672	0.5	1
R_{tc}	Aerofoil thickness to chord ratio	0.12	0.08	0.18
N_z	Ultimate load factor	3.8	2.5	6
W_{dg}	Flight design gross weight (lb)	2000	1700	2500
W_p	paint weight (lb/ft^2)	0.064	0.025	0.08

The study begins with a baseline Cessna C172 Skyhawk Aircraft as its reference point. It aims to investigate the impact of wing area and fuel weight on the overall weight of the aircraft. Two crucial parameters in this analysis are the aspect ratio (A), defined as the ratio of the wing's length to the average chord (thickness of the airfoil), and the taper ratio (λ), which represents the ratio of the maximum to the minimum thickness of the airfoil or the maximum to minimum chord.

It's important to note that the equation used in this context is not a computer simulation but will be treated as one for the purpose of illustration. This approach involves employing a true mathematical equation, even if it's considered unknown, as a useful tool for generating realistic settings to test the methodology. The functional form of this equation was derived by "calibrating" known physical relationships to curves obtained from existing aircraft data, as referenced in Raymer (2006). Essentially, it acts as a surrogate for actual measurements of aircraft weight.

Examining the mathematical properties of the AWWE (Aircraft Weight With Wing Area and Fuel Weight Equation), it is evident that the response is highly nonlinear concerning its inputs. While it's common to apply the logarithm to simplify equations with complex exponents, even when modeling the logarithm, which transforms powers into slope coefficients and products into sums, the response remains nonlinear due to the presence of trigonometric terms. Given the combination of nonlinearity and high input dimension, simple linear and quadratic response surface approximations are likely to be inadequate for this analysis.

77.3. Goals: Understanding and Optimization

The primary goals of this study revolve around understanding and optimization:

1. **Understanding:** One of the straightforward objectives is to gain a deep understanding of the input-output relationships in this context. Given the global perspective implied by this setting, it becomes evident that a more sophisticated

77.3. Goals: Understanding and Optimization

model is almost necessary. At this stage, let's focus on this specific scenario to establish a clear understanding.

2. **Optimization:** Another application of this analysis could be optimization. There may be an interest in minimizing the weight of the aircraft, but it's likely that there will be constraints in place. For example, the presence of wings with a nonzero area is essential for the aircraft to be capable of flying. In situations involving (constrained) optimization, a global perspective and, consequently, the use of flexible modeling are vital.

The provided Python code serves as a genuine computer implementation that “solves” a mathematical model. It accepts arguments encoded in the unit cube, with defaults used to represent baseline settings, as detailed in the table labeled as Table 77.1. To map values from the interval $[a, b]$ to the interval $[0, 1]$, the following formula can be employed:

$$y = f(x) = \frac{x - a}{b - a}. \quad (77.1)$$

To reverse this mapping and obtain the original values, the formula

$$g(y) = a + (b - a)y \quad (77.2)$$

can be used. The function `wingwt()` expects inputs from the unit cube, which are then transformed back to their original scales using Equation 77.2. The function is defined as follows:

```
def wingwt(Sw=0.48, Wfw=0.4, A=0.38, L=0.5, q=0.62, l=0.344, Rtc=0.4, Nz=0.37, Wdg=0.38):
    # put coded inputs back on natural scale
    Sw = Sw * (200 - 150) + 150
    Wfw = Wfw * (300 - 220) + 220
    A = A * (10 - 6) + 6
    L = (L * (10 - (-10)) - 10) * np.pi/180
    q = q * (45 - 16) + 16
    l = l * (1 - 0.5) + 0.5
    Rtc = Rtc * (0.18 - 0.08) + 0.08
    Nz = Nz * (6 - 2.5) + 2.5
    Wdg = Wdg*(2500 - 1700) + 1700
    # calculation on natural scale
    W = 0.036 * Sw**0.758 * Wfw**0.0035 * (A/np.cos(L)**2)**0.6 * q**0.006
    W = W * l**0.04 * (100*Rtc/np.cos(L))**(-0.3) * (Nz*Wdg)**(0.49)
    return(W)
```

77.4. Properties of the Python “Solver”

The compute time required by the “wingwt” solver is extremely short and can be considered trivial in terms of computational resources. The approximation error is exceptionally small, effectively approaching machine precision, which indicates the high accuracy of the solver’s results.

To simulate time-consuming evaluations, a deliberate delay is introduced by incorporating a `sleep(3600)` command, which effectively synthesizes a one-hour execution time for a particular evaluation.

Moving on to the AWWE visualization, plotting in two dimensions is considerably simpler than dealing with nine dimensions. To aid in creating visual representations, the code provided below establishes a grid within the unit square to facilitate the generation of sliced visuals. This involves generating a “meshgrid” as outlined in the code.

```
x = np.linspace(0, 1, 3)
y = np.linspace(0, 1, 3)
X, Y = np.meshgrid(x, y)
zp = zip(np.ravel(X), np.ravel(Y))
list(zp)
```

```
[(np.float64(0.0), np.float64(0.0)),
 (np.float64(0.5), np.float64(0.0)),
 (np.float64(1.0), np.float64(0.0)),
 (np.float64(0.0), np.float64(0.5)),
 (np.float64(0.5), np.float64(0.5)),
 (np.float64(1.0), np.float64(0.5)),
 (np.float64(0.0), np.float64(1.0)),
 (np.float64(0.5), np.float64(1.0)),
 (np.float64(1.0), np.float64(1.0))]
```


The coding used to transform inputs from natural units is largely a matter of taste, so long as it’s easy to undo for reporting back on original scales

```
%matplotlib inline
# plt.style.use('seaborn-white')
x = np.linspace(0, 1, 100)
y = np.linspace(0, 1, 100)
X, Y = np.meshgrid(x, y)
```

77.5. Plot 1: Load Factor (N_z) and Aspect Ratio (A)

We will vary N_z and A , with other inputs fixed at their baseline values.

```
z = wingwt(A = X, Nz = Y)
fig = plt.figure(figsize=(7., 5.))
plt.contourf(X, Y, z, 20, cmap='jet')
plt.xlabel("A")
plt.ylabel("Nz")
plt.title("Load factor (Nz) vs. Aspect Ratio (A)")
plt.colorbar()
plt.show()
```




Contour plots can be refined, e.g., by adding explicit contour lines as shown in the following figure.

```
contours = plt.contour(X, Y, z, 4, colors='black')
plt.clabel(contours, inline=True, fontsize=8)
```

77. Aircraft Wing Weight Example

```
plt.xlabel("A")
plt.ylabel("Nz")

plt.imshow(z, extent=[0, 1, 0, 1], origin='lower',
           cmap='jet', alpha=0.9)
plt.colorbar()
```



The interpretation of the AWWE plot can be summarized as follows:

- The figure displays the weight response as a function of two variables, N_z and A , using an image-contour plot.
- The slight curvature observed in the contours suggests an interaction between these two variables.
- Notably, the range of outputs depicted in the figure, spanning from approximately 160 to 320, nearly encompasses the entire range of outputs observed from various input settings within the full 9-dimensional input space.
- The plot indicates that aircraft wings tend to be heavier when the aspect ratios (A) are high.
- This observation aligns with the idea that wings are designed to withstand and accommodate high gravitational forces (g -forces, large N_z), and there may be a compounding effect where larger values of N_z contribute to increased wing weight.

77.6. Plot 2: Taper Ratio and Fuel Weight


- It's plausible that this phenomenon is related to the design considerations of fighter jets, which cannot have the efficient and lightweight glider-like wings typically found in other types of aircraft.

77.6. Plot 2: Taper Ratio and Fuel Weight

- The same experiment for two other inputs, e.g., taper ratio λ and fuel weight W_{fw}

```
z = wingwt(Wfw = X, Nz = Y)
contours = plt.contour(X, Y, z, 4, colors='black')
plt.clabel(contours, inline=True, fontsize=8)
plt.xlabel("WfW")
plt.ylabel("l")

plt.imshow(z, extent=[0, 1, 0, 1], origin='lower',
           cmap='jet', alpha=0.9)
plt.colorbar();
```



- Interpretation of Taper Ratio (l) and Fuel Weight (W_{fw})
 - Apparently, neither input has much effect on wing weight:

77. Aircraft Wing Weight Example

- * with λ having a marginally greater effect, covering less than 4 percent of the span of weights observed in the $A \times N_z$ plane
- There's no interaction evident in $\lambda \times W_{fw}$

77.7. The Big Picture: Combining all Variables

```
p1 = ["Sw", "Wfw", "A", "L", "q", "l", "Rtc", "Nz", "Wdg"]
```

```
Z = []
Zlab = []
l = len(p1)
# lc = math.comb(l,2)
for i in range(l):
    for j in range(i+1, l):
        # for j in range(l):
            # print(p1[i], p1[j])
            d = {p1[i]: X, p1[j]: Y}
            Z.append(wingwt(**d))
            Zlab.append([p1[i], p1[j]])
```

Now we can generate all 36 combinations, e.g., our first example is combination $p = 19$.


```
p = 19
Zlab[p]
```

```
['A', 'Nz']
```

To help interpret outputs from experiments such as this one—to level the playing field when comparing outputs from other pairs of inputs—code below sets up a color palette that can be re-used from one experiment to the next. We use the arguments `vmin=180` and `vmax=360` to implement comparability

```
plt.contourf(X, Y, Z[p], 20, cmap='jet', vmin=180, vmax=360)
plt.xlabel(Zlab[p][0])
plt.ylabel(Zlab[p][1])
plt.colorbar()
```


77.7. The Big Picture: Combining all Variables




- Let's plot the second example, taper ratio λ and fuel weight W_{fw}
- This is combination 11:

```
p = 11
Zlab[p]
```

```
['Wfw', '1']
```

```
plt.contourf(X, Y, Z[p], 20, cmap='jet', vmin=180, vmax=360)
plt.xlabel(Zlab[p][0])
plt.ylabel(Zlab[p][1])
plt.colorbar()
```

77. Aircraft Wing Weight Example



- Using a global colormap indicates that these variables have minor effects on the wing weight.
- Important factors can be detected by visual inspection
- Plotting the Big Picture: we can plot all 36 combinations in one figure.

```
fig = plt.figure(figsize=(20., 20.))
grid = ImageGrid(fig, 111, # similar to subplot(111)
                  nrows_ncols=(6,6), # creates 2x2 grid of axes
                  axes_pad=0.5, # pad between axes in inch.
                  share_all=True,
                  label_mode="all",
                  )

i = 0
for ax, im in zip(grid, Z):
    # Iterating over the grid returns the Axes.
    ax.set_xlabel(Zlab[i][0])
    ax.set_ylabel(Zlab[i][1])
    # ax.set_title(Zlab[i][1] + " vs. " + Zlab[i][0])
    ax.contourf(X, Y, im, 30, cmap = "jet", vmin = 180, vmax = 360)
    i = i + 1

plt.show()
```




Figure 77.1.: AWE Landscape

77.8. AWE Landscape

- Our Observations

1. The load factor N_z , which determines the magnitude of the maximum aerodynamic load on the wing, is very active and involved in interactions with other variables.
 - Classic example: the interaction of N_z with the aspect ratio A indicates a heavy wing for high aspect ratios and large g -forces

77. Aircraft Wing Weight Example

- This is the reason why highly manoeuvrable fighter jets cannot have very efficient, glider wings)
- 2. Aspect ratio A and airfoil thickness to chord ratio R_{tc} have nonlinear interactions.
- 3. Most important variables:
 - Ultimate load factor N_z , wing area S_w , and flight design gross weight W_{dg} .
- 4. Little impact: dynamic pressure q , taper ratio l , and quarter-chord sweep L .
- Expert Knowledge
 - Aircraft designers know that the overall weight of the aircraft and the wing area must be kept to a minimum
 - the latter usually dictated by constraints such as required stall speed, landing distance, turn rate, etc.

77.9. Summary of the First Experiments

- First, we considered two pairs of inputs, out of 36 total pairs
- Then, the “Big Picture”:
 - For each pair we evaluated `wingwt` 10,000 times
- Doing the same for all pairs would require 360K evaluations:
 - not a reasonable number with a real computer simulation that takes any non-trivial amount of time to evaluate
 - Only 1s per evaluation: > 100 hours
- Many solvers take minutes/hours/days to execute a single run
- And: three-way interactions?
- Consequence: a different strategy is needed

77.10. Exercise

77.10.1. Adding Paint Weight

- Paint weight is not considered.
- Add Paint Weight W_p to formula (the updated formula is shown below) and update the functions and plots in the notebook.

$$W = 0.036 S_W^{0.758} \times W_{fw}^{0.0035} \times \left(\frac{A}{\cos^2 \Lambda} \right)^{0.6} \times q^{0.006} \times \lambda^{0.04} \\ \times \left(\frac{100 R_{tc}}{\cos \Lambda} \right)^{-0.3} \times (N_z W_{dg})^{0.49} + S_w W_p$$

77.11. Jupyter Notebook

i Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

Part XIII.

Numerical Methods

78. Simulation and Surrogate Modeling

- We will consider the interplay between
 - mathematical models,
 - numerical approximation,
 - simulation,
 - computer experiments, and
 - field data
- Experimental design will play a key role in our developments, but not in the classical regression and response surface methodology sense
- Challenging real-data/real-simulation examples benefiting from modern surrogate modeling methodology
- We will consider the classical, response surface methodology (RSM) approach, and then move on to more modern approaches
- All approaches are based on surrogates

78.1. Surrogates

- Gathering data is **expensive**, and sometimes getting exactly the data you want is impossible or unethical
- **Surrogate**: substitute for the real thing
- In statistics, draws from predictive equations derived from a fitted model can act as a surrogate for the data-generating mechanism
- Benefits of the surrogate approach:
 - Surrogate could represent a cheaper way to explore relationships, and entertain “what ifs?”
 - Surrogates favor faithful yet pragmatic reproduction of dynamics:
 - * interpretation,
 - * establishing causality, or
 - * identification
 - Many numerical simulators are **deterministic**, whereas field observations are noisy or have measurement error

78.1.1. Costs of Simulation

- Computer simulations are generally cheaper (but not always!) than physical observation
- Some computer simulations can be just as expensive as field experimentation, but computer modeling is regarded as easier because:
 - the experimental apparatus is better understood
 - more aspects may be controlled.

78.1.2. Mathematical Models and Meta-Models

- Use of mathematical models leveraging numerical solvers has been commonplace for some time
- Mathematical models became more complex, requiring more resources to simulate/solve numerically
- Practitioners increasingly relied on **meta-models** built off of limited simulation campaigns

78.1.3. Surrogates = Trained Meta-models

- Data collected via expensive computer evaluations tuned flexible functional forms that could be used in lieu of further simulation to
 - save money or computational resources;
 - cope with an inability to perform future runs (expired licenses, off-line or over-impacted supercomputers)
- Trained meta-models became known as **surrogates**

78.1.4. Computer Experiments

- **Computer experiment**: design, running, and fitting meta-models.
 - Like an ordinary statistical experiment, except the data are generated by computer codes rather than physical or field observations, or surveys
- **Surrogate modeling** is statistical modeling of computer experiments

78.1.5. Limits of Mathematical Modeling

- Mathematical biologists, economists and others had reached the limit of equilibrium-based mathematical modeling with cute closed-form solutions
- **Stochastic simulations replace deterministic solvers** based on FEM, Navier–Stokes or Euler methods
- Agent-based simulation models are used to explore predator-prey (Lotka–Volterra) dynamics, spread of disease, management of inventory or patients in health insurance markets
- Consequence: the distinction between surrogate and statistical model is all but gone

78.1.6. Why Computer Simulations are Necessary

- You can't seed a real community with Ebola and watch what happens
- If there's (real) field data, say on a historical epidemic, further experimentation may be almost entirely limited to the mathematical and computer modeling side
- Classical statistical methods offer little guidance

78.1.7. Simulation Requirements

- Simulation should
 - enable rich **diagnostics** to help criticize that models
 - **understanding** its sensitivity to inputs and other configurations
 - providing the ability to **optimize** and
 - refine both **automatically** and with expert intervention
- And it has to do all that while remaining **computationally tractable**
- One perspective is so-called **response surface methods** (RSMs):
- a poster child from industrial statistics' heyday, well before information technology became a dominant industry

78.2. Applications of Surrogate Models

The four most common usages of surrogate models are:

- **Augmenting Expensive Simulations:** Surrogate models act as a 'curve fit' to approximate the results of expensive simulation codes, enabling predictions without rerunning the primary source. This provides significant speed improvements while maintaining useful accuracy.

78. *Simulation and Surrogate Modeling*

- **Calibration of Predictive Codes:** Surrogates bridge the gap between simpler, faster but less accurate models and more accurate, slower models. This multi-fidelity approach allows for improved accuracy without the full computational expense.
- **Handling Noisy or Missing Data:** Surrogates smooth out random or systematic errors in experimental or computational data, filling gaps and revealing overall trends while filtering out extraneous details.
- **Data Mining and Insight Generation:** Surrogates help identify functional relationships between variables and their impact on results. They enable engineers to focus on critical variables and visualize data trends effectively.

78.3. DACE and RSM

Mathematical models implemented in computer codes are used to circumvent the need for expensive field data collection. These models are particularly useful when dealing with highly nonlinear response surfaces, high signal-to-noise ratios (which often involve deterministic evaluations), and a global scope. As a result, a new approach is required in comparison to Response Surface Methodology (RSM), which is discussed in Section 81.1.

With the improvement in computing power and simulation fidelity, researchers gain higher confidence and a better understanding of the dynamics in physical, biological, and social systems. However, the expansion of configuration spaces and increasing input dimensions necessitates more extensive designs. High-performance computing (HPC) allows for thousands of runs, whereas previously only tens were possible. This shift towards larger models and training data presents new computational challenges.

Research questions for DACE (Design and Analysis of Computer Experiments) include how to design computer experiments that make efficient use of computation and how to meta-model computer codes to save on simulation effort. The choice of surrogate model for computer codes significantly impacts the optimal experiment design, and the preferred model-design pairs can vary depending on the specific goal.

The combination of computer simulation, design, and modeling with field data from similar real-world experiments introduces a new category of computer model tuning problems. The ultimate goal is to automate these processes to the greatest extent possible, allowing for the deployment of HPC with minimal human intervention.

One of the remaining differences between RSM and DACE lies in how they handle noise. DACE employs replication, a technique that would not be used in a deterministic setting, to separate signal from noise. Traditional RSM is best suited for situations where a substantial proportion of the variability in the data is due to noise, and where the acquisition of data values can be severely limited. Consequently, RSM is better suited for a different class of problems, aligning with its intended purposes.

Two very good texts on computer experiments and surrogate modeling are Santner et al. (2003) and Forrester et al. (2008). The former is the canonical reference in the statistics literature and the latter is perhaps more popular in engineering.

Example 78.1 (Example: DACE and RSM). Imagine you are a chemical engineer tasked with optimizing a chemical process to maximize yield. You can control temperature and pressure, but repeated experiments show variability in yield due to inconsistencies in raw materials.

- Using RSM: You would use RSM to design a series of experiments varying temperature and pressure. You would then fit a response surface (a mathematical model) to the data, helping you understand how changes in temperature and pressure affect yield. Using this model, you can identify optimal conditions for maximizing yield despite the noise.
- Using DACE: If instead you use a computational model to simulate the chemical process and want to account for numerical noise or uncertainty in model parameters, you might use DACE. You would run simulations at different conditions, possibly repeating them to assess variability and build a surrogate model that accurately predicts yields, which can be optimized to find the best conditions.

78.3.1. Noise Handling in RSM and DACE

Noise in RSM: In experimental settings, noise often arises due to variability in experimental conditions, measurement errors, or other uncontrollable factors. This noise can significantly affect the response variable, Y . Replication is a standard procedure for handling noise in RSM. In the context of computer experiments, noise might not be present in the traditional sense since simulations can be deterministic. However, variability can arise from uncertainty in input parameters or model inaccuracies. DACE predominantly utilizes advanced interpolation to construct accurate models of deterministic data, sometimes considering statistical noise modeling if needed.

78.4. Updating a Surrogate Model

A surrogate model is updated by incorporating new data points, known as infill points, into the model to improve its accuracy and predictive capabilities. This process is iterative and involves the following steps:

- Identify Regions of Interest: The surrogate model is analyzed to determine areas where it is inaccurate or where further exploration is needed. This could be regions with high uncertainty or areas where the model predicts promising results (e.g., potential optima).

78. Simulation and Surrogate Modeling

- **Select Infill Points:** Infill points are new data points chosen based on specific criteria, such as:
- **Exploitation:** Sampling near predicted optima to refine the solution. **Exploration:** Sampling in regions of high uncertainty to improve the model globally. **Balanced Approach:** Combining exploitation and exploration to ensure both local and global improvements.
- **Evaluate the True Function:** The true function (e.g., a simulation or experiment) is evaluated at the selected infill points to obtain their corresponding outputs.
- **Update the Surrogate Model:** The surrogate model is retrained or updated using the new data, including the infill points, to improve its accuracy.
- **Repeat:** The process is repeated until the model meets predefined accuracy criteria or the computational budget is exhausted.

Definition 78.1 (Infill Points). Infill points are strategically chosen new data points added to the surrogate model. They are selected to:

- Reduce uncertainty in the model.
- Improve predictions in regions of interest.
- Enhance the model's ability to identify optima or trends.

The selection of infill points is often guided by infill criteria, such as:

- **Expected Improvement (EI):** Maximizing the expected improvement over the current best solution.
- **Uncertainty Reduction:** Sampling where the model's predictions have high variance.
- **Probability of Improvement (PI):** Sampling where the probability of improving the current best solution is highest.

The iterative infill-points updating process ensures that the surrogate model becomes increasingly accurate and useful for optimization or decision-making tasks.

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

79. Sampling Plans

Note

- This section is based on chapter 1 in Forrester et al. (2008).
- The following Python packages are imported:

```
import pandas as pd
import numpy as np
from typing import Tuple, Optional
import matplotlib.pyplot as plt

from spotoptim.sampling.lhs import rlh
from spotoptim.sampling.effects import screening_plot, screeningplan
from spotoptim.function.so import wingwt

from spotoptim.sampling.mm import (bestlh,
                                   jd, mm, mmphi, mmsort, perturb, mmlhs, phisort, mmphi_intensive)

from spotoptim.sampling.design import generate_clustered_design, generate_sobol_design, generate_c
```

79.1. Ideas and Concepts

Definition 79.1 (Sampling Plan). In the context of computer experiments, the term sampling plan refers to the set of input values, say X , at which the computer code is evaluated.

The goal of a sampling plan is to efficiently explore the input space to understand the behavior of the computer code and build a surrogate model that accurately represents the code's behavior. Traditionally, Response Surface Methodology (RSM) has been used to design sampling plans for computer experiments. These sampling plans are based on procedures that generate points by means of a rectangular grid or a factorial design.

However, more recently, Design and Analysis of Computer Experiments (DACE) has emerged as a more flexible and powerful approach for designing sampling plans.

79. Sampling Plans

Engineering design often requires the construction of a surrogate model $\hat{f}$ to approximate the expensive response of a black-box function f . The function $f(x)$ represents a continuous metric (e.g., quality, cost, or performance) defined over a design space $D \subset \mathbb{R}^k$, where x is a k -dimensional vector of design variables. Since evaluating f is costly, only a sparse set of samples is used to construct $\hat{f}$, which can then provide inexpensive predictions for any $x \in D$.

The process involves:

- Sampling discrete observations:
- Using these samples to construct an approximation $\hat{f}$.
- Ensuring the surrogate model is well-posed, meaning it is mathematically valid and can generalize predictions effectively.

A sampling plan

$$X = \{x^{(i)} \in D | i = 1, \dots, n\}$$

determines the spatial arrangement of observations. While some models require a minimum number of data points n , once this threshold is met, a surrogate model can be constructed to approximate f efficiently.

A well-posed model does not always perform well because its ability to generalize depends heavily on the sampling plan used to collect data. If the sampling plan is poorly designed, the model may fail to capture critical behaviors in the design space. For example:

- **Extreme Sampling:** Measuring performance only at the extreme values of parameters may miss important behaviors in the center of the design space, leading to incomplete understanding.
- **Uneven Sampling:** Concentrating samples in certain regions while neglecting others forces the model to extrapolate over unsampled areas, potentially resulting in inaccurate or misleading predictions. Additionally, in some cases, the data may come from external sources or be limited in scope, leaving little control over the sampling plan. This can further restrict the model's ability to generalize effectively.

79.1.1. The ‘Curse of Dimensionality’ and How to Avoid It

The “curse of dimensionality” refers to the exponential increase in computational complexity and data requirements as the number of dimensions (variables) in a problem grows. For a one-dimensional space, sampling n locations may suffice for accurate predictions. In high-dimensional spaces, the amount of data needed to maintain the same level of accuracy or coverage increases dramatically. For example, if a one-dimensional space requires n samples for a certain accuracy, a k -dimensional space

would require n^k samples. This makes tasks like optimization, sampling, and modeling computationally expensive and often impractical in high-dimensional settings.

Example 79.1 (Example: Curse of Dimensionality). Consider a simple example where we want to model the cost of a car tire based on its wheel diameter. If we have one variable (wheel diameter), we might need 10 simulations to get a good estimate of the cost. Now, if we add 8 more variables (e.g., tread pattern, rubber type, etc.), the number of simulations required increases to 10^8 (10 million). This is because the number of combinations of design variables grows exponentially with the number of dimensions. This means that the computational budget required to evaluate all combinations of design variables becomes infeasible. In this case, it would take 11,416 years to complete the simulations, making it impractical to explore the design space fully.

79.1.2. Physical versus Computational Experiments

Physical experiments are prone to experimental errors from three main sources:

- Human error: Mistakes made by the experimenter.
- Random error: Measurement inaccuracies that vary unpredictably.
- Systematic error: Consistent bias due to flaws in the experimental setup.

The key distinction is repeatability: systematic errors remain constant across repetitions, while random errors vary.

Computational experiments, on the other hand, are deterministic and free from random errors. However, they are still affected by:

- Human error: Bugs in code or incorrect boundary conditions.
- Systematic error: Biases from model simplifications (e.g., inviscid flow approximations) or finite resolution (e.g., insufficient mesh resolution).

The term “noise” is used differently in physical and computational contexts. In physical experiments, it refers to random errors, while in computational experiments, it often refers to systematic errors.

Understanding these differences is crucial for designing experiments and applying techniques like Gaussian process-based approximations. For physical experiments, replication mitigates random errors, but this is unnecessary for deterministic computational experiments.

79.1.3. Designing Preliminary Experiments (Screening)

Minimizing the number of design variables $x_1, x_2, \dots, x_k$ is crucial before modeling the objective function f . This process, called screening, aims to reduce dimensionality without compromising the analysis. If f is at least once differentiable over the design domain D , the partial derivative $\frac{\partial f}{\partial x_i}$ can be used to classify variables:

- Negligible Variables: If $\frac{\partial f}{\partial x_i} = 0, \forall x \in D$, the variable x_i can be safely neglected.
- Linear Additive Variables: If $\frac{\partial f}{\partial x_i} = \text{constant} \neq 0, \forall x \in D$, the effect of x_i is linear and additive.
- Nonlinear Variables: If $\frac{\partial f}{\partial x_i} = g(x_i), \forall x \in D$, where $g(x_i)$ is a non-constant function, f is nonlinear in x_i .
- Interactive Nonlinear Variables: If $\frac{\partial f}{\partial x_i} = g(x_i, x_j, \dots), \forall x \in D$, where $g(x_i, x_j, \dots)$ is a function involving interactions with other variables, f is nonlinear in x_i and interacts with x_j .

Measuring $\frac{\partial f}{\partial x_i}$ across the entire design space is often infeasible due to limited budgets. The percentage of time allocated to screening depends on the problem: If many variables are expected to be inactive, thorough screening can significantly improve model accuracy by reducing dimensionality. If most variables are believed to impact the objective, focus should shift to modeling instead. Screening is a trade-off between computational cost and model accuracy, and its effectiveness depends on the specific problem context.

79.1.3.1. Estimating the Distribution of Elementary Effects

In order to simplify the presentation of what follows, we make, without loss of generality, the assumption that the design space $D = [0, 1]^k$; that is, we normalize all variables into the unit cube. We shall adhere to this convention for the rest of the book and strongly urge the reader to do likewise when implementing any algorithms described here, as this step not only yields clearer mathematics in some cases but also safeguards against scaling issues.

Before proceeding with the description of the Morris algorithm, we need to define an important statistical concept. Let us restrict our design space D to a k -dimensional, p -level full factorial grid, that is,

$$x_i \in \{0, \frac{1}{p-1}, \frac{2}{p-1}, \dots, 1\}, \quad \text{for } i = 1, \dots, k.$$

Definition 79.2 (Elementary Effect). For a given baseline value $x \in D$, let $d_i(x)$ denote the elementary effect of x_i , where:

$$d_i(x) = \frac{f(x_1, \dots, x_i + \Delta, \dots, x_k) - f(x_1, \dots, x_i - \Delta, \dots, x_k)}{2\Delta}, \quad i = 1, \dots, k, \quad (79.1)$$

where Δ is the step size, which is defined as the distance between two adjacent levels in the grid. In other words, we have:

with

$$\Delta = \frac{\xi}{p-1}, \quad \xi \in \mathbb{N}^*, \quad \text{and} \quad x \in D, \quad \text{such that its components } x_i \leq 1 - \Delta.$$

Δ is the step size. The elementary effect $d_i(x)$ measures the sensitivity of the function f to changes in the variable x_i at the point x .

Morris's method aims to estimate the parameters of the distribution of elementary effects associated with each variable. A large measure of central tendency indicates that a variable has a significant influence on the objective function across the design space, while a large measure of spread suggests that the variable is involved in interactions or contributes to the nonlinearity of f . In practice, the sample mean and standard deviation of a set of $d_i(x)$ values, calculated in different parts of the design space, are used for this estimation.

To ensure efficiency, the preliminary sampling plan X should be designed so that each evaluation of the objective function f contributes to the calculation of two elementary effects, rather than just one (as would occur with a naive random spread of baseline x values and adding Δ to one variable). Additionally, the sampling plan should provide a specified number (e.g., r) of elementary effects for each variable, independently drawn with replacement. For a detailed discussion on constructing such a sampling plan, readers are encouraged to consult Morris's original paper (Morris, 1991). Here, we focus on describing the process itself.

The random orientation of the sampling plan B can be constructed as follows:

- Let B be a $(k+1) \times k$ matrix of 0s and 1s, where for each column i , two rows differ only in their i -th entries.
- Compute a random orientation of B , denoted B^* :

$$B^* = (1_{k+1,k} x^* + (\Delta/2) [(2B - 1_{k+1,k}) D^* + 1_{k+1,k}]) P^*,$$

where:

- D^* is a k -dimensional diagonal matrix with diagonal elements ± 1 (equal probability),
- $\mathbf{1}$ is a matrix of 1s,
- x^* is a randomly chosen point in the p -level design space (limited by Δ),

79. Sampling Plans

- P^* is a $k \times k$ random permutation matrix with one 1 per column and row.

spotpython provides a Python implementation to compute B^* , see <https://github.com/sequential-parameter-optimization/spotPython/blob/main/src/spotpython/utils/effects.py>.

Here is the corresponding code:

```
def randorient(k, p, xi, seed=None):
    # Initialize random number generator with the provided seed
    if seed is not None:
        rng = np.random.default_rng(seed)
    else:
        rng = np.random.default_rng()

    # Step length
    Delta = xi / (p - 1)

    m = k + 1

    # A truncated p-level grid in one dimension
    xs = np.arange(0, 1 - Delta, 1 / (p - 1))
    xsl = len(xs)
    if xsl < 1:
        print(f"xi = {xi}.")
        print(f"p = {p}.")
        print(f"Delta = {Delta}.")
        print(f"p - 1 = {p - 1}.")
        raise ValueError(f"The number of levels xsl is {xsl}, but it must be greater t

    # Basic sampling matrix
    B = np.vstack((np.zeros((1, k)), np.tril(np.ones((k, k)))))

    # Randomization

    # Matrix with +1s and -1s on the diagonal with equal probability
    Dstar = np.diag(2 * rng.integers(0, 2, size=k) - 1)

    # Random base value
    xstar = xs[rng.integers(0, xsl, size=k)]

    # Permutation matrix
    Pstar = np.zeros((k, k))
    rp = rng.permutation(k)
    for i in range(k):
```

```

Pstar[i, rp[i]] = 1

# A random orientation of the sampling matrix
Bstar = (np.ones((m, 1)) @ xstar.reshape(1, -1) +
         (Delta / 2) * ((2 * B - np.ones((m, k))) @ Dstar +
                        np.ones((m, k)))) @ Pstar

return Bstar

```

The code following snippet generates a random orientation of a sampling matrix **Bstar** using the `randorient()` function. The input parameters are:

- $k = 3$: The number of design variables (dimensions).
- $p = 3$: The number of levels in the grid for each variable.
- $\xi = 1$: A parameter used to calculate the step size Δ .

Step-size calculation is performed as follows: $\Delta = \xi / (p - 1) = 1 / (3 - 1) = 0.5$, which determines the spacing between levels in the grid.

Next, random sampling matrix construction is computed:

- A truncated grid is created with levels $[0, 0.5]$ (based on Δ).
- A basic sampling matrix **B** is constructed, which is a lower triangular matrix with 0s and 1s.

Then, randomization is applied:

- **Dstar**: A diagonal matrix with random entries of +1 or -1.
- **xstar**: A random starting point from the grid.
- **Pstar**: A random permutation matrix.

Random orientation is applied to the basic sampling matrix **B** to create **Bstar**. This involves scaling, shifting, and permuting the rows and columns of **B**.

The final output is the matrix **Bstar**, which represents a random orientation of the sampling plan. Each row corresponds to a sampled point in the design space, and each column corresponds to a design variable.

Example 79.2 (Random Orientation of the Sampling Matrix in 2-D).

```

k = 2
p = 3
xi = 1
Bstar = randorient(k, p, xi, seed=123)
print(f"Random orientation of the sampling matrix:\n{Bstar}")

```

79. Sampling Plans

Random orientation of the sampling matrix:

```
[[0.5 0. ]  
 [0.  0. ]  
 [0.  0.5]]
```

We can visualize the random orientation of the sampling matrix in 2-D as shown in Figure 79.1.

```
plt.figure(figsize=(6, 6))  
plt.scatter(Bstar[:, 0], Bstar[:, 1], color='blue', s=50, label='Hypercube Points')  
for i in range(Bstar.shape[0]):  
    plt.text(Bstar[i, 0] + 0.01, Bstar[i, 1] + 0.01, str(i), fontsize=9)  
plt.xlim(-0.1, 1.1)  
plt.ylim(-0.1, 1.1)  
plt.xlabel('x1')  
plt.ylabel('x2')  
plt.grid()
```




Figure 79.1.: Random orientation of the sampling matrix in 2-D. The labels indicate the row index of the points.

Example 79.3 (Random Orientation of the Sampling Matrix).

```
k = 3
p = 3
xi = 1
Bstar = randorient(k, p, xi)
print(f"Random orientation of the sampling matrix:\n{Bstar}")
```

Random orientation of the sampling matrix:

79. Sampling Plans

```
[[0.5 0.  0. ]
 [0.5 0.  0.5]
 [0.5 0.5 0.5]
 [0.  0.5 0.5]]
```

To obtain r elementary effects for each variable, the screening plan is built from r random orientations:

$$X = \begin{pmatrix} B_1^* \\ B_2^* \\ \vdots \\ B_r^* \end{pmatrix}$$

The function `screeningplan()` generates a screening plan by calling the `randorient()` function r times. It creates a list of random orientations and then concatenates them into a single array, which represents the screening plan. The screening plan implementation in Python is as follows (see <https://github.com/sequential-parameter-optimization/spotPython/blob/main/src/spotpython/utils/effects.py>):

```
def screeningplan(k, p, xi, r):
    # Empty list to accumulate screening plan rows
    X = []
    for i in range(r):
        X.append(randorient(k, p, xi))
    # Concatenate list of arrays into a single array
    X = np.vstack(X)
    return X
```

It works like follows:

- The value of the objective function f is computed for each row of the screening plan matrix X . These values are stored in a column vector t of size $(r * (k + 1)) \times 1$, where:
 - r is the number of random orientations.
 - k is the number of design variables.

The elementary effects are calculated using the following formula:

- For each random orientation, adjacent rows of the screening plan matrix X and their corresponding function values from t are used.
- These values are inserted into Equation 79.1 to compute elementary effects for each variable. An elementary effect measures the sensitivity of the objective function to changes in a specific variable.

Results can be used for a statistical analysis. After collecting a sample of r elementary effects for each variable:

- The sample mean (central tendency) is computed to indicate the overall influence of the variable.
- The sample standard deviation (spread) is computed to capture variability, which may indicate interactions or nonlinearity.

The results (sample means and standard deviations) are plotted on a chart for comparison. This helps identify which variables have the most significant impact on the objective function and whether their effects are linear or involve interactions. This is implemented in the function `screening_plot()` in Python, which uses the helper function `_screening()` to calculate the elementary effects and their statistics.

```
def _screening(X, fun, xi, p, labels, bounds=None) -> tuple:
    """Helper function to calculate elementary effects for a screening design.

    Args:
        X (np.ndarray): The screening plan matrix, typically structured
            within a [0,1]^k box.
        fun (object): The objective function to evaluate at each
            design point in the screening plan.
        xi (float): The elementary effect step length factor.
        p (int): Number of discrete levels along each dimension.
        labels (list of str): A list of variable names corresponding to
            the design variables.
        bounds (np.ndarray): A 2xk matrix where the first row contains
            lower bounds and the second row contains upper bounds for
            each variable.

    Returns:
        tuple: A tuple containing two arrays:
            - sm: The mean of the elementary effects for each variable.
            - ssd: The standard deviation of the elementary effects for
                each variable.
    """
    k = X.shape[1]
    r = X.shape[0] // (k + 1)

    # Scale each design point
    t = np.zeros(X.shape[0])
    for i in range(X.shape[0]):
        if bounds is not None:
            X[i, :] = bounds[0, :] + X[i, :] * (bounds[1, :] - bounds[0, :])
        t[i] = np.asarray(fun(X[i, :])).item()
```

79. Sampling Plans

```

# Elementary effects
F = np.zeros((k, r))
for i in range(r):
    for j in range(i * (k + 1), i * (k + 1) + k):
        idx = np.where(X[j, :] - X[j + 1, :] != 0)[0][0]
        F[idx, i] = (t[j + 1] - t[j]) / (xi / (p - 1))

# Statistical measures (divide by n)
ssd = np.std(F, axis=1, ddof=0)
sm = np.mean(F, axis=1)
return sm, ssd

def screening_plot(X, fun, xi, p, labels, bounds=None, show=True) -> None:
    """Generates a plot with elementary effect screening metrics.

    This function calculates the mean and standard deviation of the
    elementary effects for a given set of design variables and plots
    the results.

    Args:
        X (np.ndarray):
            The screening plan matrix, typically structured within a  $[0,1]^k$  box.
        fun (object):
            The objective function to evaluate at each design point in the screening p
        xi (float):
            The elementary effect step length factor.
        p (int):
            Number of discrete levels along each dimension.
        labels (list of str):
            A list of variable names corresponding to the design variables.
        bounds (np.ndarray):
            A 2xk matrix where the first row contains lower bounds and
            the second row contains upper bounds for each variable.
        show (bool):
            If True, the plot is displayed. Defaults to True.

    Returns:
        None: The function generates a plot of the results.
    """
    k = X.shape[1]
    sm, ssd = _screening(X=X, fun=fun, xi=xi, p=p, labels=labels, bounds=bounds)
    plt.figure()
    for i in range(k):

```

```
plt.text(sm[i], ssd[i], labels[i], fontsize=10)
plt.axis([min(sm), 1.1 * max(sm), min(ssd), 1.1 * max(ssd)])
plt.xlabel("Sample means")
plt.ylabel("Sample standard deviations")
plt.gca().tick_params(labelsize=10)
plt.grid(True)
if show:
    plt.show()
```

79.1.4. Special Considerations When Deploying Screening Algorithms

When implementing the screening algorithm described above, two specific scenarios require special attention:

- **Duplicate Design Points:** If the dimensionality k of the space is relatively low and you can afford a large number of elementary effects r , we should be aware of the increased probability of duplicate design points appearing in the sampling plan X . *Since the responses at sample points are deterministic, there's no value in evaluating the same point multiple times. Fortunately, this issue is relatively uncommon in practice, as screening high-dimensional spaces typically requires large numbers of elementary effects, which naturally reduces the likelihood of duplicates.
- **Failed Simulations:** Numerical simulation codes occasionally fail to return valid results due to meshing errors, non-convergence of partial differential equation solvers, numerical instabilities, or parameter combinations outside the stable operating range.

From a screening perspective, this is particularly problematic because an entire random orientation B^* becomes compromised if even a single point within it fails to evaluate properly. Implementing error handling strategies or fallback methods to manage such cases should be considered.

For robust screening studies, monitoring simulation success rates and having contingency plans for failed evaluations are important aspects of the experimental design process.

79.2. Analyzing Variable Importance in Aircraft Wing Weight

Let us consider the following analytical expression used as a conceptual level estimate of the weight of a light aircraft wing as discussed in Chapter 77.

79. Sampling Plans

```
k = 10
p = 10
xi = 1
r = 25
X = screeningplan(k=k, p=p, xi=xi, r=r) # shape (r x (k+1), k)
value_range = np.array([
    [150, 220, 6, -10, 16, 0.5, 0.08, 2.5, 1700, 0.025],
    [200, 300, 10, 10, 45, 1.0, 0.18, 6.0, 2500, 0.08 ],
])
labels = [
    "S_W", "W_fw", "A", "Lambda",
    "q", "lambda", "tc", "N_z",
    "W_dg", "W_p"
]
screening_plot(
    X=X,
    fun=wingwt,
    bounds=np.array([np.zeros(10), np.ones(10)]),
    xi=xi,
    p=p,
    labels=labels,
)
```




Figure 79.2.: Estimated means and standard deviations of the elementary effects for the 10 design variables of the wing weight function. Example based on Forrester et al. (2008).

i Nondeterministic Results

The code will generate a slightly different screening plan each time, as it uses random orientations of the sampling matrix B .

Figure 79.2 provides valuable insights into variable activity without requiring domain expertise. The screening study with $r = 25$ elementary effects reveals distinct patterns in how variables affect wing weight:

- Variables with Minimal Impact: A clearly defined group of variables clusters around the origin - indicating their minimal impact on wing weight:
 - Paint weight (W_p) - as expected, contributes little to overall wing weight
 - Dynamic pressure (q) - within our chosen range, this has limited effect (essentially representing different cruise altitudes at the same speed)
 - Taper ratio (λ) and quarter-chord sweep (Λ) - these geometric parameters have minor influence within the narrow range (-10° to 10°) typical of light aircraft
- Variables with Linear Effects:

79. Sampling Plans

- While still close to the origin, fuel weight (W_{fw}) shows a slightly larger central tendency with very low standard deviation. This indicates moderate importance but minimal involvement in interactions with other variables.
- Variables with Nonlinear/Interactive Effects:
 - Aspect ratio (A) and airfoil thickness ratio (R_{tc}) show similar importance levels, but their high standard deviations suggest significant nonlinear behavior and interactions with other variables.
- Dominant Variables: The most significant impacts come from:
 - Flight design gross weight (W_{dg})
 - Wing area (S_W)
 - Ultimate load factor (N_z)

These variables show both large central tendency values and high standard deviations, indicating strong direct effects and complex interactions. The interaction between aspect ratio and load factor is particularly important - high values of both create extremely heavy wings, explaining why highly maneuverable fighter jets cannot use glider-like wing designs.

What makes this screening approach valuable is its ability to identify critical variables without requiring engineering knowledge or expensive modeling. In real-world applications, we rarely have the luxury of creating comprehensive parameter space visualizations, which is precisely why surrogate modeling is needed. After identifying the active variables through screening, we can design a focused sampling plan for these key variables. This forms the foundation for building an accurate surrogate model of the objective function.

When the objective function is particularly expensive to evaluate, we might recycle the runs performed during screening for the actual model fitting step. This is most effective when some variables prove to have no impact at all. However, since completely inactive variables are rare in practice, engineers must carefully balance the trade-off between reusing expensive simulation runs and introducing potential noise into the model.

79.3. Designing a Sampling Plan

79.3.1. Stratification

A feature shared by all of the approximation models discussed in Forrester et al. (2008) is that they are more accurate in the vicinity of the points where we have evaluated the objective function. In later chapters we will delve into the laws that quantify our decaying trust in the model as we move away from a known, sampled point, but for the purposes of the present discussion we shall merely draw the intuitive conclusion

79.3. Designing a Sampling Plan

that a uniform level of model accuracy throughout the design space requires a uniform spread of points. A sampling plan possessing this feature is said to be space-filling.

The most straightforward way of sampling a design space in a uniform fashion is by means of a rectangular grid of points. This is the full factorial sampling technique.

Here is the simplified version of a **Python** function that will sample the unit hypercube at all levels in all dimensions, with the k -vector q containing the number of points required along each dimension, see <https://github.com/sequential-parameter-optimization/spotPython/blob/main/src/spotpython/utils/sampling.py>.

The variable **Edges** specifies whether we want the points to be equally spaced from edge to edge (**Edges=1**) or we want them to be in the centres of $n = q_1 \times q_2 \times \dots \times q_k$ bins filling the unit hypercube (for any other value of **Edges**).

```
def fullfactorial(q_param, Edges=1) -> np.ndarray:
    """Generates a full factorial sampling plan in the unit cube.

    Args:
        q (list or np.ndarray):
            A list or array containing the number of points along each dimension (k-vector).
        Edges (int, optional):
            Determines spacing of points. If `Edges=1`, points are equally spaced from edge to edge.
            Otherwise, points will be in the centers of  $n = q[0] \times q[1] \times \dots \times q[k-1]$  bins filling the unit cube.

    Returns:
        (np.ndarray): Full factorial sampling plan as an array of shape (n, k), where n is the total number of points.

    Raises:
        ValueError: If any dimension in `q` is less than 2.
    """
    q_levels = np.array(q_param) # Use a distinct variable for original levels
    if np.min(q_levels) < 2:
        raise ValueError("You must have at least two points per dimension.")

    n = np.prod(q_levels)
    k = len(q_levels)
    X = np.zeros((n, k))

    # q_for_prod_calc is used for calculating repetitions, includes the phantom element.
    # This matches the logic of the user-provided snippet where 'q' was modified.
    q_for_prod_calc = np.append(q_levels, 1)

    for j in range(k): # k is the original number of dimensions
        # current_dim_levels is the number of levels for the current dimension j
        # In the user's snippet, q[j] correctly refers to the original level count
```

79. Sampling Plans

```

# as j ranges from 0 to k-1, and q_for_prod_calc[j] = q_levels[j] for this ran
current_dim_levels = q_for_prod_calc[j]

if Edges == 1:
    one_d_slice = np.linspace(0, 1, int(current_dim_levels))
else:
    # Corrected calculation for bin centers
    if current_dim_levels == 1: # Should not be hit if np.min(q_levels) >= 2
        one_d_slice = np.array([0.5])
    else:
        one_d_slice = np.linspace(1 / (2 * current_dim_levels),
                                   1 - 1 / (2 * current_dim_levels),
                                   int(current_dim_levels))

column = np.array([])
# The product q_for_prod_calc[j + 1 : k] correctly calculates
# the product of remaining original dimensions' levels.
num_consecutive_repeats = np.prod(q_for_prod_calc[j + 1 : k])

# This loop structure replicates the logic from the user's snippet
while len(column) < n:
    for ll_idx in range(int(current_dim_levels)): # Iterate through levels of
        val_to_repeat = one_d_slice[ll_idx]
        column = np.append(column, np.ones(int(num_consecutive_repeats)) * va
    X[:, j] = column
return X

```

```

q = [3, 2]
X = fullfactorial(q, Edges=0)
print(X)

```

```

[[0.16666667 0.25      ]
 [0.16666667 0.75      ]
 [0.5         0.25      ]
 [0.5         0.75      ]
 [0.83333333 0.25      ]
 [0.83333333 0.75      ]]

```

Figure 79.3 shows the points in the unit hypercube for the case of 3x2 points.

```

X = fullfactorial(q, Edges=1)
print(X)

```





Figure 79.3.: 2D Full Factorial Sampling (3x2 Points). Edges = 0

79. Sampling Plans

```
[[0.  0. ]  
 [0.  1. ]  
 [0.5 0. ]  
 [0.5 1. ]  
 [1.  0. ]  
 [1.  1. ]]
```

Figure 79.4 shows the points in the unit hypercube for the case of 3x2 points with edges.




Figure 79.4.: 2D Full Factorial Sampling (3x2 Points). Edges = 1

The full factorial sampling plan method generates a uniform sampling design by creating a grid of points across all dimensions. For example, calling `fullfactorial([3, 4, 5], 1)` produces a three-dimensional sampling plan with 3, 4, and 5 levels along each dimension, respectively. While this approach satisfies the uniformity criterion, it has two significant limitations:

- **Restricted Design Sizes:** The method only works for designs where the total number of points n can be expressed as the product of the number of levels in each dimension, i.e., $n = q_1 \times q_2 \times \dots \times q_k$.
- **Overlapping Projections:** When the sampling points are projected onto individual axes, sets of points may overlap, reducing the effectiveness of the sampling plan. This can lead to non-uniform coverage in the projections, which may not fully represent the design space.

79.3.2. Latin Squares and Random Latin Hypercubes

To improve the uniformity of projections for any individual variable, the range of that variable can be divided into a large number of equal-sized bins, and random subsamples of equal size can be generated within these bins. This method is called stratified random sampling. Extending this idea to all dimensions results in a stratified sampling plan, commonly implemented using Latin hypercube sampling.

Definition 79.3 (Latin Squares and Hypercubes). In the context of statistical sampling, a square grid containing sample positions is a Latin square if (and only if) there is only one sample in each row and each column. A Latin hypercube is the generalisation of this concept to an arbitrary number of dimensions, whereby each sample is the only one in each axis-aligned hyperplane containing it

For two-dimensional discrete variables, a Latin square ensures uniform projections. An $(n \times n)$ Latin square is constructed by filling each row and column with a permutation of $\{1, 2, \dots, n\}$, ensuring each number appears only once per row and column.

Example 79.4 (Latin Square). For $n = 4$, a Latin square might look like this:

2	1	3	4
3	2	4	1
1	4	2	3
4	3	1	2

Latin Hypercubes are the multidimensional extension of Latin squares. The design space is divided into equal-sized hypercubes (bins), and one point is placed in each bin. The placement ensures that moving along any axis from an occupied bin does not encounter another occupied bin. This guarantees uniform projections across all dimensions. To construct a Latin hypercube, the following steps are taken:

79. Sampling Plans

- Represent the sampling plan as an $n \times k$ matrix X , where n is the number of points and k is the number of dimensions.
- Fill each column of X with random permutations of $\{1, 2, \dots, n\}$.
- Normalize the plan into the unit hypercube $[0, 1]^k$.

This approach ensures multidimensional stratification and uniformity in projections. Here is the code:

```
def rlh(n: int, k: int, edges: int = 0) -> np.ndarray:
    # Initialize array
    X = np.zeros((n, k), dtype=float)

    # Fill with random permutations
    for i in range(k):
        X[:, i] = np.random.permutation(n)

    # Adjust normalization based on the edges flag
    if edges == 1:
        # [X=0..n-1] -> [0..1]
        X = X / (n - 1)
    else:
        # Points at true midpoints
        # [X=0..n-1] -> [0.5/n..(n-0.5)/n]
        X = (X + 0.5) / n

    return X
```

Example 79.5 (Random Latin Hypercube). The following code can be used to generate a 2D Latin hypercube with 5 points and edges=0:

```
X = rlh(n=5, k=2, edges=0)
print(X)
```

```
[[0.1 0.3]
 [0.5 0.9]
 [0.7 0.7]
 [0.3 0.5]
 [0.9 0.1]]
```

Figure 79.5 shows the points in the unit hypercube for the case of 5 points with edges=0.

Example 79.6 (Random Latin Hypercube with Edges). The following code can be used to generate a 2D Latin hypercube with 5 points and edges=1:




Figure 79.5.: 2D Latin Hypercube Sampling (5 Points, Edges=0)

```
X = r1h(n=5, k=2, edges=1)
print(X)
```

```
[[1.    1.   ]
 [0.75 0.5  ]
 [0.5  0.   ]
 [0.    0.75]
 [0.25 0.25]]
```

Figure 79.6 shows the points in the unit hypercube for the case of 5 points with `edges=1`.

79.3.3. Space-filling Designs: Maximin Plans

A widely adopted measure for assessing the uniformity, or ‘space-fillingness’, of a sampling plan is the maximin metric, initially proposed by Johnson et al. (1990). This criterion can be formally defined as follows.

Consider a sampling plan X . Let $d_1, d_2, \dots, d_m$ represent the unique distances between all possible pairs of points within X , arranged in ascending order. Furthermore, let $J_1, J_2, \dots, J_m$ be defined such that J_j denotes the count of point pairs in X separated by the distance d_j .

Definition 79.4 (Maximin plan). A sampling plan X is considered a maximin plan if, among all candidate plans, it maximizes the smallest inter-point distance d_1 . Among plans that satisfy this condition, it further minimizes J_1 , the number of pairs separated by this minimum distance.

While this definition is broadly applicable to any collection of sampling plans, our focus is narrowed to Latin hypercube designs to preserve their desirable stratification properties. However, even within this restricted class, Definition 79.4 may identify multiple equivalent maximin designs. To address this, a more comprehensive ‘tie-breaker’ definition, as proposed by Morris and Mitchell (1995), is employed:

Definition 79.5 (Maximin plan with tie-breaker). A sampling plan X is designated as the maximin plan if it sequentially optimizes the following conditions: it maximizes d_1 ; among those, it minimizes J_1 ; among those, it maximizes d_2 ; among those, it minimizes J_2 ; and so forth, concluding with minimizing J_m .

Johnson et al. (1990) established that the maximin criterion (Definition 79.4) is equivalent to the D-optimality criterion used in linear regression. However, the extended maximin criterion incorporating a tie-breaker (Definition 79.5) is often preferred due to




Figure 79.6.: 2D Latin Hypercube Sampling (5 Points, Edges=1)

79. Sampling Plans

its intuitive nature and practical utility. Given that the sampling plans under consideration make no assumptions about model structure, the latter criterion (Definition 79.5) will be employed.

To proceed, a precise definition of ‘distance’ within these contexts is necessary. The p-norm is the most widely adopted metric for this purpose:

Definition 79.6 (p-norm). The p-norm of a vector $\vec{x} = (x_1, x_2, \dots, x_k)$ is defined as:

$$d_p(\vec{x}^{(i_1)}, \vec{x}^{(i_2)}) = \left(\sum_{j=1}^k |x_j^{(i_1)} - x_j^{(i_2)}|^p \right)^{1/p}. \quad (79.2)$$

When $p = 1$, Equation 79.2 defines the rectangular distance, occasionally referred to as the Manhattan norm (an allusion to a grid-like city layout). Setting $p = 2$ yields the Euclidean norm. The existing literature offers limited evidence to suggest the superiority of one norm over the other for evaluating sampling plans when no model structure assumptions are made. It is important to note, however, that the rectangular distance is considerably less computationally demanding. This advantage can be quite significant, particularly when evaluating large sampling plans.

For the computational implementation of Definition 79.5, the initial step involves constructing the vectors $d_1, d_2, \dots, d_m$ and $J_1, J_2, \dots, J_m$. The `jd` function facilitates this task.

79.3.3.1. The Function `jd`

The function `jd` computes the distinct p-norm distances between all pairs of points in a given set and counts their occurrences. It returns two arrays: one for the distinct distances and another for their multiplicities.

```
def jd(X: np.ndarray, p: float = 1.0) -> Tuple[np.ndarray, np.ndarray]:
    """
    Args:
        X (np.ndarray):
            A 2D array of shape (n, d) representing n points
            in d-dimensional space.
        p (float, optional):
            The distance norm to use.
            p=1 uses the Manhattan (L1) norm, while p=2 uses the
            Euclidean (L2) norm. Defaults to 1.0 (Manhattan norm).

    Returns:
        (np.ndarray, np.ndarray):
```



```

    A tuple (J, distinct_d), where:
    - distinct_d is a 1D float array of unique,
      sorted distances between points.
    - J is a 1D integer array that provides
      the multiplicity (occurrence count)
      of each distance in distinct_d.
"""
n = X.shape[0]

# Allocate enough space for all pairwise distances
# (n*(n-1))/2 pairs for an n-point set
pair_count = n * (n - 1) // 2
d = np.zeros(pair_count, dtype=float)

# Fill the distance array
idx = 0
for i in range(n - 1):
    for j in range(i + 1, n):
        # Compute the p-norm distance
        d[idx] = np.linalg.norm(X[i] - X[j], ord=p)
        idx += 1

# Find unique distances and their multiplicities
distinct_d = np.unique(d)
J = np.zeros_like(distinct_d, dtype=int)
for i, val in enumerate(distinct_d):
    J[i] = np.sum(d == val)
return J, distinct_d

```

Example 79.7 (The Function `jd`). Consider a small 3-point set in 2D space, with points located at (0,0), (1,1), and (2,2) as shown in Figure 79.7. The distinct distances and their occurrences can be computed using the `jd` function, as shown in the following code:

```

J, distinct_d = jd(X, p=2.0)
print("Distinct distances (d_i):", distinct_d)
print("Occurrences (J_i):", J)

```

```

Distinct distances (d_i): [1.41421356 2.82842712]
Occurrences (J_i): [2 1]

```




Figure 79.7.: 3-Point Set in 2D Space

79.3.4. Memory Management

A computationally intensive part of the calculation performed with the `jd`-function is the creation of the vector $\vec{d}$ containing all pairwise distances. This is particularly true for large sampling plans; for instance, a 1000-point plan requires nearly half a million distance calculations.

Definition 79.7 (Pre-allocation of Memory). Pre-allocation of memory is a programming technique where a fixed amount of memory is reserved for a data structure (like an array or vector) before it is actually filled with data. This is done to avoid the computational overhead associated with dynamic memory allocation, which involves repeatedly requesting and resizing memory as new elements are added.

Consequently, pre-allocating memory for the distance vector $\vec{d}$ is essential. This necessitates a slightly less direct method for computing the indices of $\vec{d}$, rather than appending each new element, which would involve costly dynamic memory allocation.

The implementation of Definition 79.5 is now required. Finding the most space-filling design involves pairwise comparisons. This problem can be approached using a ‘divide and conquer’ strategy, simplifying it to the task of selecting the better of two sampling plans. The function `mm(X1,X2,p)` is designed for this purpose. It returns an index indicating which of the two designs is more space-filling, or 0 if they are equally space-filling, based on the p -norm for distance computation.

79.3.4.1. The Function `mm`

The function `mm` compares two sampling plans based on the Morris-Mitchell criterion. It uses the `jd` function to compute the distances and multiplicities, constructs vectors for comparison, and determines which plan is more space-filling.

```
def mm(X1: np.ndarray, X2: np.ndarray, p: Optional[float] = 1.0) -> int:
    """
    Args:
        X1 (np.ndarray): A 2D array representing the first sampling plan.
        X2 (np.ndarray): A 2D array representing the second sampling plan.
        p (float, optional): The distance metric. p=1 uses Manhattan (L1) distance,
            while p=2 uses Euclidean (L2). Defaults to 1.0.

    Returns:
        int:
            - 0 if both plans are identical or equally space-filling
            - 1 if X1 is more space-filling
            - 2 if X2 is more space-filling
    """
```

79. Sampling Plans

```
X1_sorted = X1[np.lexsort(np.rot90(X1))]
X2_sorted = X2[np.lexsort(np.rot90(X2))]
if np.array_equal(X1_sorted, X2_sorted):
    return 0 # Identical sampling plans

# Compute distance multiplicities for each plan
J1, d1 = jd(X1, p)
J2, d2 = jd(X2, p)
m1, m2 = len(d1), len(d2)

# Construct V1 and V2: alternate distance and negative multiplicity
V1 = np.zeros(2 * m1)
V1[0::2] = d1
V1[1::2] = -J1

V2 = np.zeros(2 * m2)
V2[0::2] = d2
V2[1::2] = -J2

# Trim the longer vector to match the size of the shorter
m = min(m1, m2)
V1 = V1[:m]
V2 = V2[:m]

# Compare element-by-element:
# c[i] = 1 if V1[i] > V2[i], 2 if V1[i] < V2[i], 0 otherwise.
c = (V1 > V2).astype(int) + 2 * (V1 < V2).astype(int)

if np.sum(c) == 0:
    # Equally space-filling
    return 0
else:
    # The first non-zero entry indicates which plan is better
    idx = np.argmax(c != 0)
    return c[idx]
```

Example 79.8 (The Function `mm`). We can use the `mm` function to compare two sampling plans. The following code creates two 3-point sampling plans in 2D (shown in Figure 79.8) and compares them using the Morris-Mitchell criterion:

```
X1 = np.array([[0.0, 0.0], [0.5, 0.5], [0.0, 1.0], [1.0, 1.0]])
X2 = np.array([[0.1, 0.1], [0.4, 0.6], [0.1, 0.9], [0.9, 0.9]])
```

We can compare which plan has better space-filling (Morris-Mitchell). The output is




Figure 79.8.: Comparison of Two Sampling Plans

79. Sampling Plans

either 0, 1, or 2 depending on which plan is more space-filling.

```
better = mm(X1, X2, p=2.0)
print(f"Plan {better} is more space-filling.")
```

Plan 1 is more space-filling.

79.3.4.2. The Function `mmphi`

Searching across a space of potential sampling plans can be accomplished by pairwise comparisons. An optimization algorithm could, in theory, be written with `mm` as the comparative objective. However, experimental evidence (Morris and Mitchell 1995) suggests that the resulting optimization landscape can be quite deceptive, making it difficult to search reliably. This difficulty arises because the comparison process terminates upon finding the first non-zero element in the comparison array `c`. Consequently, the remaining values in the distance $(d_1, d_2, \dots, d_m)$ and multiplicity $(J_1, J_2, \dots, J_m)$ arrays are disregarded. These disregarded values, however, might contain potentially useful ‘slope’ information about the global landscape for the optimization process.

To address this, Morris and Mitchell (1995) defined the following scalar-valued criterion function, which is used to rank competing sampling plans. This function, while based on the logic of Definition 79.5, incorporates the complete vectors $d_1, d_2, \dots, d_m$ and $J_1, J_2, \dots, J_m$.

Definition 79.8 (Morris-Mitchell Criterion). The Morris-Mitchell criterion is defined as:

$$\Phi_q(X) = \left(\sum_{j=1}^m J_j d_j^{-q} \right)^{1/q}, \quad (79.3)$$

where X is the sampling plan, d_j is the distance between points, J_j is the multiplicity of that distance, and q is a user-defined exponent. The parameter q can be adjusted to control the influence of smaller distances on the overall metric.

The smaller the value of Φ_q , the better the space-filling properties of X will be.

The function `mmphi` computes the Morris-Mitchell sampling plan quality criterion for a given sampling plan. It takes a 2D array of points and calculates the space-fillingness metric based on the distances between points. This can be implemented in Python as follows:

```

def mmphi(X: np.ndarray,
          q: Optional[float] = 2.0,
          p: Optional[float] = 1.0) -> float:
    """
    Args:
        X (np.ndarray):
            A 2D array representing the sampling plan,
            where each row is a point in
            d-dimensional space (shape: (n, d)).
        q (float, optional):
            Exponent used in the computation of the metric.
            Defaults to 2.0.
        p (float, optional):
            The distance norm to use.
            For example, p=1 is Manhattan (L1),
            p=2 is Euclidean (L2). Defaults to 1.0.

    Returns:
        float:
            The space-fillingness metric  $\Phi_q$ . Larger values typically indicate a more
            space-filling plan according to the Morris-Mitchell criterion.
    """
    # Compute the distance multiplicities: J, and unique distances: d
    J, d = jd(X, p)
    # Summation of  $J[i] * d[i]^{-q}$ , then raised to  $1/q$ 
    # This follows the Morris-Mitchell definition.
     $\Phi_q$  = np.sum(J * (d ** (-q))) ** (1.0 / q)
    return  $\Phi_q$ 

```

Example 79.9 (The Function `mmphi`). We can use the `mmphi` function to evaluate the space-filling quality of the two sampling plans from Example 79.8. The following code uses these two 3-point sampling plans in 2D and computes their quality using the Morris-Mitchell criterion:

```

# Two simple sampling plans from above
quality1 = mmphi(X1, q=2, p=2)
quality2 = mmphi(X2, q=2, p=2)
print(f"Quality of sampling plan X1: {quality1}")
print(f"Quality of sampling plan X2: {quality2}")

```

```

Quality of sampling plan X1: 2.91547594742265
Quality of sampling plan X2: 3.917162046269215

```

79. Sampling Plans

This equation provides a more compact representation of the maximin criterion, but the selection of the q value is an important consideration. Larger values of q ensure that terms in the sum corresponding to smaller inter-point distances (the d_j values, which are sorted in ascending order) have a dominant influence. As a result, Φ_q will rank sampling plans in a way that closely emulates the original maximin definition (Definition 79.5). This implies that the optimization landscape might retain the challenging characteristics that the Φ_q metric, especially with smaller q values, is intended to alleviate. Conversely, smaller q values tend to produce a Φ_q landscape that, while not perfectly aligning with the original definition, is generally more conducive to optimization.

To illustrate the relationship between Equation 79.3 and the maximin criterion of Definition 79.5, sets of 50 random Latin hypercubes of varying sizes and dimensionalities were considered by Forrester et al. (2008). The correlation plots from this analysis suggest that as the sampling plan size increases, a smaller q value is needed for the Φ_q -based ranking to closely match the ranking derived from Definition 79.5.

Rankings based on both the direct maximin comparison (`mm`) and the Φ_q metric (`mmphi`), determined using a simple bubble sort algorithm, are implemented in the Python function `mmsort`.

79.3.4.3. The Function `mmsort`

The function `mmsort` is designed to rank multiple sampling plans based on their space-filling properties using the Morris-Mitchell criterion. It takes a 3D array of sampling plans and returns the indices of the plans sorted in ascending order of their space-filling quality.

```
def mmsort(X3D: np.ndarray, p: Optional[float] = 1.0) -> np.ndarray:
    """
    Args:
        X3D (np.ndarray):
            A 3D NumPy array of shape (n, d, m), where m is the number of
            sampling plans, and each plan is an (n, d) matrix of points.
        p (float, optional):
            The distance metric to use. p=1 for Manhattan (L1), p=2 for
            Euclidean (L2). Defaults to 1.0.

    Returns:
        np.ndarray:
            A 1D integer array of length m that holds the plan indices in
            ascending order of space-filling quality. The first index in the
            returned array corresponds to the most space-filling plan.
    """
    # Number of plans (m)
    m = X3D.shape[2]
```



```

# Create index array (1-based to match original MATLAB convention)
Index = np.arange(1, m + 1)

swap_flag = True
while swap_flag:
    swap_flag = False
    i = 0
    while i < m - 1:
        # Compare plan at Index[i] vs. Index[i+1] using mm()
        # Note: subtract 1 from each index to convert to 0-based array indexing
        if mm(X3D[:, :, Index[i] - 1], X3D[:, :, Index[i + 1] - 1], p) == 2:
            # Swap indices if the second plan is more space-filling
            Index[i], Index[i + 1] = Index[i + 1], Index[i]
            swap_flag = True
        i += 1

return Index

```

Example 79.10 (The Function `mmsort`). The `mmsort` function can be used to rank multiple sampling plans based on their space-filling properties. The following code demonstrates how to use `mmsort` to compare two 3-point sampling plans in 3D space:

Suppose we have two 3-point sampling plans X_1 and X_2 from above. They are sorted using the Morris-Mitchell criterion with $p = 2.0$. For example, the output `[1, 2]` indicates that X_1 is more space-filling than X_2 :

```

X3D = np.stack([X1, X2], axis=2)
ranking = mmsort(X3D, p=2.0)
print(ranking)

```

```
[1 2]
```

To determine the optimal Latin hypercube for a specific application, a recommended approach by Morris and Mitchell (1995) involves minimizing Φ_q for a set of q values (1, 2, 5, 10, 20, 50, and 100). Subsequently, the best plan from these results is selected based on the actual maximin definition. The `mmsort` function can be utilized for this purpose: a 3D matrix, X_{3D} , can be constructed where each 2D slice represents the best sampling plan found for each Φ_q . Applying `mmsort(X3D, 1)` then ranks these plans according to Definition 79.5, using the rectangular distance metric. The subsequent discussion will address the methods for finding these optimized Φ_q designs.

79.3.4.4. The Function phisort

`phisort` only differs from `mmsort` in having q as an additional argument, as well as the comparison line being:

```
if mmphi(X3D[:, :, Index[i] - 1], q=q, p=p) >
    mmphi(X3D[:, :, Index[i + 1] - 1], q=q, p=p):
```

```
def phisort(X3D: np.ndarray,
            q: Optional[float] = 2.0,
            p: Optional[float] = 1.0) -> np.ndarray:
    """
    Args:
        X3D (np.ndarray):
            A 3D array of shape (n, d, m),
            where m is the number of sampling plans.
        q (float, optional):
            Exponent for the mmphi metric. Defaults to 2.0.
        p (float, optional):
            Distance norm for mmphi.
            p=1 is Manhattan; p=2 is Euclidean.
            Defaults to 1.0.

    Returns:
        np.ndarray:
            A 1D integer array of length m, giving the plan indices in ascending
            order of mmphi. The first index in the returned array corresponds
            to the numerically lowest mmphi value.
    """
    # Number of 2D sampling plans
    m = X3D.shape[2]
    # Create a 1-based index array
    Index = np.arange(1, m + 1)
    # Bubble-sort: plan with lower mmphi() climbs toward the front
    swap_flag = True
    while swap_flag:
        swap_flag = False
        for i in range(m - 1):
            # Retrieve mmphi values for consecutive plans
            val_i = mmphi(X3D[:, :, Index[i] - 1], q=q, p=p)
            val_j = mmphi(X3D[:, :, Index[i + 1] - 1], q=q, p=p)

            # Swap if the left plan's mmphi is larger (i.e. 'worse')
            if val_i > val_j:
```

```

        Index[i], Index[i + 1] = Index[i + 1], Index[i]
        swap_flag = True
    return Index

```

Example 79.11 (The Function `phisort`). The `phisort` function can be used to rank multiple sampling plans based on the Morris-Mitchell criterion. The following code demonstrates how to use `phisort` to compare two 3-point sampling plans in 3D space:

```

X1 = bestlh(n=5, k=2, population=5, iterations=10)
X2 = bestlh(n=5, k=2, population=15, iterations=20)
X3 = bestlh(n=5, k=2, population=25, iterations=30)
# Map X1 and X2 so that X3D has the two sampling plans
# in X3D[:, :, 0] and X3D[:, :, 1]
X3D = np.array([X1, X2])
print(phisort(X3D))
X3D = np.array([X3, X2])
print(phisort(X3D))

```

```

[2 1]
[1 2]

```

79.3.5. Optimizing the Morris-Mitchell Criterion Φ_q

Once a criterion for assessing the quality of a Latin hypercube sampling plan has been established, a systematic method for optimizing this metric across the space of Latin hypercubes is required. This task is non-trivial; as the reader may recall from the earlier discussion on Latin squares, this search space is vast. In fact, its vastness means that for many practical applications, locating the globally optimal solution is often infeasible. Therefore, the objective becomes finding the best possible sampling plan achievable within a specific computational time budget.

This budget is influenced by the computational cost associated with obtaining each objective function value. Determining the optimal allocation of total computational effort—between generating the sampling plan and actually evaluating the objective function at the selected points—remains an open research question. However, it is typical for no more than approximately 5% of the total available time to be allocated to the task of generating the sampling plan itself.

Forrester et al. (2008) draw an analogy to the process of devising a revision timetable before an exam. While a well-structured timetable enhances the effectiveness of revision, an excessive amount of the revision time itself should not be consumed by the planning phase.

79. Sampling Plans

A significant challenge in devising a sampling plan optimizer is ensuring that the search process remains confined to the space of valid Latin hypercubes. As previously discussed, the defining characteristic of a Latin hypercube X is that each of its columns represents a permutation of the possible levels for the corresponding variable. Consequently, the smallest modification that can be applied to a Latin hypercube—without compromising its crucial multidimensional stratification property—involves swapping two elements within any single column of X . A Python implementation for ‘mutating’ a Latin hypercube through such an operation, generalized to accommodate random changes applied to multiple sites, is provided below:

79.3.5.1. The Function `perturb()`

The function `perturb` randomly swaps elements in a Latin hypercube sampling plan. It takes a 2D array representing the sampling plan and performs a specified number of random element swaps, ensuring that the result remains a valid Latin hypercube.

```
def perturb(X: np.ndarray,
            PertNum: Optional[int] = 1) -> np.ndarray:
    """
    Args:
        X (np.ndarray):
            A 2D array (sampling plan) of shape (n, k),
            where each row is a point
            and each column is a dimension.
        PertNum (int, optional):
            The number of element swaps (perturbations)
            to perform. Defaults to 1.

    Returns:
        np.ndarray:
            The perturbed sampling plan,
            identical in shape to the input, with
            one or more random column swaps executed.
    """
    # Get dimensions of the plan
    n, k = X.shape
    if n < 2 or k < 2:
        raise ValueError("Latin hypercubes require at least 2 points and 2 dimensions")
    for _ in range(PertNum):
        # Pick a random column
        col = int(np.floor(np.random.rand() * k))
        # Pick two distinct row indices
        el1, el2 = 0, 0
        while el1 == el2:
```

```

    e11 = int(np.floor(np.random.rand() * n))
    e12 = int(np.floor(np.random.rand() * n))
    # Swap the two selected elements in the chosen column
    X[e11, col], X[e12, col] = X[e12, col], X[e11, col]
return X

```

Example 79.12 (The Function `perturb()`). The `perturb` function can be used to randomly swap elements in a Latin hypercube sampling plan. The following code demonstrates how to use `perturb` to create a perturbed version of a 4x2 sampling plan:

```

X_original = np.array([[1, 3],[2, 4],[3, 1],[4, 2]])
print("Original Sampling Plan:")
print(X_original)
print("Perturbed Sampling Plan:")
X_perturbed = perturb(X_original, PertNum=1)
print(X_perturbed)

```

Original Sampling Plan:

```

[[1 3]
 [2 4]
 [3 1]
 [4 2]]

```

Perturbed Sampling Plan:

```

[[1 1]
 [2 4]
 [3 3]
 [4 2]]

```

Forrester et al. (2008) uses the term ‘mutation’, because this problem lends itself to nature-inspired computation. Morris and Mitchell (1995) use a simulated annealing algorithm, the detailed pseudocode of which can be found in their paper. As an alternative, a method based on evolutionary operation (EVOP) is offered by Forrester et al. (2008).

79.3.6. Evolutionary Operation

As introduced by Box (1957), evolutionary operation was designed to optimize chemical processes. The current parameters of the reaction would be recorded in a box at the centre of a board, with a series of ‘offspring’ boxes along the edges containing values of the parameters slightly altered with respect to the central, ‘parent’ values. Once the reaction was completed for all of these sets of variable values and the corresponding yields recorded, the contents of the central box would be replaced with that of the

79. Sampling Plans

setup with the highest yield and this would then become the parent of a new set of peripheral boxes.

This is generally viewed as a local search procedure, though this depends on the mutation step sizes, that is on the differences between the parent box and its offspring. The longer these steps, the more global is the scope of the search.

For the purposes of the Latin hypercube search, a variable scope strategy is applied. The process starts with a long step length (that is a relatively large number of swaps within the columns) and, as the search progresses, the current best basin of attraction is gradually approached by reducing the step length to a single change.

In each generation the parent is mutated (randomly, using the `perturb` function) a pertnum number of times. The sampling plan that yields the smallest Φ_q value (as per the Morris-Mitchell criterion, calculated using `mmphi`) among all offspring and the parent is then selected; in evolutionary computation parlance this selection philosophy is referred to as elitism.

The EVOP based search for space-filling Latin hypercubes is thus a truly evolutionary process: the optimized sampling plan results from the nonrandom survival of random variations.

79.3.7. Putting it all Together

All the pieces of the optimum Latin hypercube sampling process puzzle are now in place: the random hypercube generator as a starting point for the optimization process, the ‘spacefillingness’ metric that needs to be optimized, the optimization engine that performs this task and the comparison function that selects the best of the optima found for the various q ’s. These pieces just need to be put into a sequence. Here is the Python embodiment of the completed puzzle. It results in a function `bestlh` that uses the function `mmlhs` to find the best Latin hypercube sampling plan for a given set of parameters.

79.3.7.1. The Function `mmlhs`

Performs an evolutionary search (using perturbations) to find a Morris-Mitchell optimal Latin hypercube, starting from an initial plan `X_start`.

This function does the following:

1. Initializes a “best” Latin hypercube (`X_best`) from the provided `X_start`.
2. Iteratively perturbs `X_best` to create offspring.
3. Evaluates the space-fillingness of each offspring via the Morris-Mitchell metric (using `mmphi`).
4. Updates the best plan whenever a better offspring is found.

```

def mmlhs(X_start: np.ndarray,
          population: int,
          iterations: int,
          q: Optional[float] = 2.0,
          plot=False) -> np.ndarray:
    """
    Args:
        X_start (np.ndarray):
            A 2D array of shape (n, k) providing the initial Latin hypercube
            (n points in k dimensions).
        population (int):
            Number of offspring to create in each generation.
        iterations (int):
            Total number of generations to run the evolutionary search.
        q (float, optional):
            The exponent used by the Morris-Mitchell space-filling criterion.
            Defaults to 2.0.
        plot (bool, optional):
            If True, a simple scatter plot of the first two dimensions will be
            displayed at each iteration. Only if k >= 2. Defaults to False.

    Returns:
        np.ndarray:
            A 2D array representing the most space-filling Latin hypercube found
            after all iterations, of the same shape as X_start.
    """
    n = X_start.shape[0]
    if n < 2:
        raise ValueError("Latin hypercubes require at least 2 points")
    k = X_start.shape[1]
    if k < 2:
        raise ValueError("Latin hypercubes are not defined for dim k < 2")
    # Initialize best plan and its metric
    X_best = X_start.copy()
    Phi_best = mmphi(X_best, q=q)
    # After 85% of iterations, reduce the mutation rate to 1
    leveloff = int(np.floor(0.85 * iterations))
    for it in range(1, iterations + 1):
        # Decrease number of mutations over time
        if it < leveloff:
            mutations = int(round(1 + (0.5 * n - 1) * (leveloff - it) / (leveloff - 1)))
        else:
            mutations = 1
        X_improved = X_best.copy()

```

```

Phi_improved = Phi_best
# Create offspring, evaluate, and keep the best
for _ in range(population):
    X_try = perturb(X_best.copy(), mutations)
    Phi_try = mmphi(X_try, q=q)

    if Phi_try < Phi_improved:
        X_improved = X_try
        Phi_improved = Phi_try
# Update the global best if we found a better plan
if Phi_improved < Phi_best:
    X_best = X_improved
    Phi_best = Phi_improved
# Simple visualization of the first two dimensions
if plot and (X_best.shape[1] >= 2):
    plt.clf()
    plt.scatter(X_best[:, 0], X_best[:, 1], marker="o")
    plt.grid(True)
    plt.title(f"Iteration {it} - Current Best Plan")
    plt.pause(0.01)
return X_best

```

Example 79.13 (The Function `mmlhs`). The `mmlhs` function can be used to optimize a Latin hypercube sampling plan. The following code demonstrates how to use `mmlhs` to optimize a 4x2 Latin hypercube starting from an initial plan:

```

# Suppose we have an initial 4x2 plan
X_start = np.array([[0.1, 0.3], [.1, .4], [.2, .9], [.9, .2]])
print("Initial plan:")
print(X_start)
# Search for a more space-filling plan
X_opt = mmlhs(X_start, population=10, iterations=100, q=2)
print("Optimized plan:")
print(X_opt)

```

```

Initial plan:
[[0.1 0.3]
 [0.1 0.4]
 [0.2 0.9]
 [0.9 0.2]]
Optimized plan:
[[0.9 0.3]
 [0.1 0.2]]

```



```
[0.2 0.4]
[0.1 0.9]]
```

Figure 79.9 shows the initial and optimized plans in 2D. The blue points represent the initial plan, while the red points represent the optimized plan.




Figure 79.9.: Comparison of the initial and optimized plans in 2D.

79.3.7.2. The Function `bestlh`

Generates an optimized Latin hypercube by evolving the Morris-Mitchell criterion across multiple exponents (q values) and selecting the best plan.

```
def bestlh(n: int,
           k: int,
           population: int,
           iterations: int,
           p=1,
```

```

        plot=False,
        verbosity=0,
        edges=0,
        q_list=[1, 2, 5, 10, 20, 50, 100]) -> np.ndarray:
"""
Args:
    n (int):
        Number of points required in the Latin hypercube.
    k (int):
        Number of design variables (dimensions).
    population (int):
        Number of offspring in each generation of the evolutionary search.
    iterations (int):
        Number of generations for the evolutionary search.
    p (int, optional):
        The distance norm to use. p=1 for Manhattan (L1), p=2 for Euclidean (L2).
        Defaults to 1 (faster than 2).
    plot (bool, optional):
        If True, a scatter plot of the optimized plan in the first two dimensions
        will be displayed. Only if k>=2. Defaults to False.
    verbosity (int, optional):
        Verbosity level. 0 is silent, 1 prints the best q value found. Defaults to 0.
    edges (int, optional):
        If 1, places centers of the extreme bins at the domain edges ([0,1]).
        Otherwise, bins are fully contained within the domain, i.e. midpoints.
        Defaults to 0.
    q_list (list, optional):
        A list of q values to optimize. Defaults to [1, 2, 5, 10, 20, 50, 100].
        These values are used to evaluate the space-fillingness of the Latin
        hypercube. The best plan is selected based on the lowest mmphi value.

Returns:
    np.ndarray:
        A 2D array of shape (n, k) representing an optimized Latin hypercube.
"""
if n < 2:
    raise ValueError("Latin hypercubes require at least 2 points")
if k < 2:
    raise ValueError("Latin hypercubes are not defined for dim k < 2")

# A list of exponents (q) to optimize

# Start with a random Latin hypercube
X_start = rlh(n, k, edges=edges)

```

```

# Allocate a 3D array to store the results for each q
# (shape: (n, k, number_of_q_values))
X3D = np.zeros((n, k, len(q_list)))

# Evolve the plan for each q in q_list
for i, q_val in enumerate(q_list):
    if verbosity > 0:
        print(f"Now optimizing for q={q_val}...")
    X3D[:, :, i] = mmlhs(X_start, population, iterations, q_val)

# Sort the set of evolved plans according to the Morris-Mitchell criterion
index_order = mmsort(X3D, p=p)

# index_order is a 1-based array of plan indices; the first element is the best
best_idx = index_order[0] - 1
if verbosity > 0:
    print(f"Best lh found using q={q_list[best_idx]}...")

# The best plan in 3D array order
X = X3D[:, :, best_idx]

# Plot the first two dimensions
if plot and (k >= 2):
    plt.scatter(X[:, 0], X[:, 1], c="r", marker="o")
    plt.title(f"Morris-Mitchell optimum plan found using q={q_list[best_idx]}")
    plt.xlabel("x_1")
    plt.ylabel("x_2")
    plt.grid(True)
    plt.show()

return X

```

Example 79.14 (The Function `bestlh`). The `bestlh` function can be used to generate an optimized Latin hypercube sampling plan. The following code demonstrates how to use `bestlh` to create a 5x2 Latin hypercube with a population of 5 and 10 iterations:

```
Xbestlh= bestlh(n=5, k=2, population=5, iterations=10)
```

Figure 79.10 shows the best Latin hypercube sampling in 2D. The red points represent the optimized plan.

Sorting all candidate plans in ascending order is not strictly necessary - after all, only the best one is truly of interest. Nonetheless, the added computational complexity is




Figure 79.10.: Best Latin Hypercube Sampling

minimal (the vector will only ever contain as many elements as there are candidate q values, and only an index array is sorted, not the actual repository of plans). This sorting gives the reader the opportunity to compare, if desired, how different choices of q influence the resulting plans.

79.4. Experimental Analysis of the Morris-Mitchell Criterion

Morris-Mitchell Criterion Experimental Analysis

- Number of points: 16, Dimensions: 2
- mmphi parameters: q (exponent) = 2.0, p (distance norm) = 2.0 (1=Manhattan, 2=Euclidean)

```
N_POINTS = 16
N_DIM = 2
RANDOM_SEED = 42
q = 2.0
p = 2.0
```

79.4.1. Evaluation of Sampling Designs

We generate various sampling designs and evaluate their space-filling properties using the Morris-Mitchell criterion.

```
designs = {}
bounds = [(0, 1)] * N_DIM

if int(np.sqrt(N_POINTS))**2 == N_POINTS:
    designs["Grid (4x4)"] = generate_grid_design(bounds, n_design=N_POINTS)
else:
    print(f"Skipping grid design as N_POINTS={N_POINTS} is not a perfect square for a simple 2D grid")

designs["LHS"] = generate_qmc_lhs_design(bounds, n_design=N_POINTS, seed=42)

designs["Sobol"] = generate_sobol_design(bounds, n_design=N_POINTS, seed=42)

designs["Random"] = generate_uniform_design(bounds, n_design=N_POINTS, seed=42)

designs["Collinear"] = generate_collinear_design(bounds, n_design=N_POINTS, seed=42)
```

79. Sampling Plans

```
designs["Clustered (3 clusters)"] = generate_clustered_design(bounds, n_design=N_POINTS)

results = {}

print("Calculating Morris-Mitchell metric (smaller is better):")
for name, X_design in designs.items():
    metric_val = mmpm(X_design, q=q, p=p)
    results[name] = metric_val
    print(f" {name}: {metric_val:.4f}")
```

```
Calculating Morris-Mitchell metric (smaller is better):
Grid (4x4): 20.2617
LHS: 33.0365
Sobol: 30.2998
Random: 36.1032
Collinear: 84.8903
Clustered (3 clusters): 90.3702
```

```
if N_DIM == 2:
    num_designs = len(designs)
    cols = 2
    rows = int(np.ceil(num_designs / cols))
    fig, axes = plt.subplots(rows, cols, figsize=(5 * cols, 5 * rows))
    axes = axes.ravel() # Flatten axes array for easy iteration


    for i, (name, X_design) in enumerate(designs.items()):
        ax = axes[i]
        ax.scatter(X_design[:, 0], X_design[:, 1], s=50, edgecolors='k', alpha=0.7)
        ax.set_title(f"{name} \nmmpm = {results[name]:.3f}", fontsize=10)
        ax.set_xlabel("X1")
        ax.set_ylabel("X2")
        ax.set_xlim(-0.05, 1.05)
        ax.set_ylim(-0.05, 1.05)
        ax.set_aspect('equal', adjustable='box')
        ax.grid(True, linestyle='--', alpha=0.6)

    # Hide any unused subplots
    for j in range(i + 1, len(axes)):
        fig.delaxes(axes[j])

    plt.tight_layout()
    plt.suptitle(f"Comparison of 2D Sampling Designs ({N_POINTS} points each)", fontsize=12)
    plt.show()
```

79.4. Experimental Analysis of the Morris-Mitchell Criterion

Comparison of 2D Sampling Designs (16 points each)



79.4.2. Demonstrate the Impact of mmphi Parameters

Demonstrating Impact of mmphi Parameters on 'LHS' Design

```
X_lhs = designs["LHS"]

# 1. Default parameters (already calculated)
print(f"  LHS (q={q}, p={p} Euclidean): {results['LHS']:.4f}")

# 2. Change q (main exponent, literature's p or k)
q_high = 15.0
metric_lhs_q_high = mmphi(X_lhs, q=q_high, p=p)
print(f"  LHS (q={q_high}, p={p} Euclidean): {metric_lhs_q_high:.4f} (Higher q penalizes small distances more)")

# 3. Change p (distance norm, literature's q or m)
p_manhattan = 1.0
metric_lhs_p_manhattan = mmphi(X_lhs, q=q, p=p_manhattan)
print(f"  LHS (q={q}, p={p_manhattan} Manhattan): {metric_lhs_p_manhattan:.4f} (Using L1 distance)")
```

```
LHS (q=2.0, p=2.0 Euclidean): 33.0365
LHS (q=15.0, p=2.0 Euclidean): 11.7143 (Higher q penalizes small distances more)
LHS (q=2.0, p=1.0 Manhattan): 25.5294 (Using L1 distance)
```

79.4.3. Morris-Mitchell Criterion: Impact of Adding Points

Impact of adding a point to a 2x2 grid design

```
# Initial 2x2 Grid Design
X_initial = np.array([[0.0, 0.0], [1.0, 0.0], [0.0, 1.0], [1.0, 1.0]])
mmphi_initial = mmphi(X_initial, q=q, p=p)

print(f"Parameters: q (exponent) = {q}, p (distance) = {p} (Euclidean)\n")
print(f"Initial 2x2 Grid Design (4 points):")
print(f"  Points:\n{X_initial}")
print(f"  Morris-Mitchell Criterion (Phi_q): {mmphi_initial:.4f}\n")
```

```
Parameters: q (exponent) = 2.0, p (distance) = 2.0 (Euclidean)
```

```
Initial 2x2 Grid Design (4 points):
```

```
  Points:
[[0. 0.]
 [1. 0.]
```


79.4. Experimental Analysis of the Morris-Mitchell Criterion

```
[0. 1.]
[1. 1.]]
Morris-Mitchell Criterion (Phi_q): 2.2361
```

Scenarios for adding a 5th point:

```
scenarios = {
    "Scenario 1: Add to Center": {
        "new_point": np.array([[0.5, 0.5]]),
        "description": "Adding a point in the center of the grid."
    },
    "Scenario 2: Add Close to Existing (Cluster)": {
        "new_point": np.array([[0.1, 0.1]]),
        "description": "Adding a point very close to an existing point (0,0).",
    },
    "Scenario 3: Add on Edge": {
        "new_point": np.array([[0.5, 0.0]]),
        "description": "Adding a point on an edge between (0,0) and (1,0).",
    }
}

results_summary = []
augmented_designs_for_plotting = {"Initial Design": X_initial}

for name, scenario_details in scenarios.items():
    new_point = scenario_details["new_point"]
    X_augmented = np.vstack((X_initial, new_point))
    augmented_designs_for_plotting[name] = X_augmented

    mmphi_augmented = mmphi(X_augmented, q=q, p=p)
    change = mmphi_augmented - mmphi_initial

    print(f"{name}:")
    print(f"  Description: {scenario_details['description']}")
    print(f"  New Point Added: {new_point}")
    # print(f"  Augmented Design (5 points):\n{X_augmented}") # Optional: print full matrix
    print(f"  Morris-Mitchell Criterion (Phi_q): {mmphi_augmented:.4f}")
    print(f"  Change from Initial Phi_q: {change:+.4f}\n")

    results_summary.append({
        "Scenario": name,
        "Initial Phi_q": mmphi_initial,
        "Augmented Phi_q": mmphi_augmented,
        "Change": change
    })
```

Scenario 1: Add to Center:

Description: Adding a point in the center of the grid.

New Point Added: $[[0.5 \ 0.5]]$ Morris-Mitchell Criterion (Φ_q): 3.6056Change from Initial Φ_q : +1.3695

Scenario 2: Add Close to Existing (Cluster):

Description: Adding a point very close to an existing point (0,0).

New Point Added: $[[0.1 \ 0.1]]$ Morris-Mitchell Criterion (Φ_q): 7.6195Change from Initial Φ_q : +5.3834

Scenario 3: Add on Edge:

Description: Adding a point on an edge between (0,0) and (1,0).

New Point Added: $[[0.5 \ 0. \]]$ Morris-Mitchell Criterion (Φ_q): 3.8210Change from Initial Φ_q : +1.5849

```

num_designs = len(augmented_designs_for_plotting)
cols = 2
rows = int(np.ceil(num_designs / cols))

fig, axes = plt.subplots(rows, cols, figsize=(6 * cols, 5 * rows))
axes = axes.ravel()

plot_idx = 0
# Plot initial design first
ax = axes[plot_idx]
ax.scatter(X_initial[:, 0], X_initial[:, 1], s=100, edgecolors='k', alpha=0.7, label=
ax.set_title(f"Initial Design\nPhi_q = {mmphi_initial:.3f}", fontsize=10)
ax.set_xlabel("X1")
ax.set_ylabel("X2")
ax.set_xlim(-0.1, 1.1)
ax.set_ylim(-0.1, 1.1)
ax.set_aspect('equal', adjustable='box')
ax.grid(True, linestyle='--', alpha=0.6)
ax.legend(fontsize='small')
plot_idx += 1

# Plot augmented designs
for name, X_design in augmented_designs_for_plotting.items():
    if name == "Initial Design":

```

```

        continue # Already plotted

    ax = axes[plot_idx]
    # Highlight original vs new point
    original_points = X_design[:-1, :]
    new_point = X_design[-1, :].reshape(1,2)


    ax.scatter(original_points[:, 0], original_points[:, 1], s=100, edgecolors='k', alpha=0.7, label='original')
    ax.scatter(new_point[:, 0], new_point[:, 1], s=150, color='red', edgecolors='k', marker='X', label='new point')

    current_phi_q = next(item['Augmented Phi_q'] for item in results_summary if item["Scenario"] == "q")
    ax.set_title(f"{name}\nPhi_q = {current_phi_q:.3f}", fontsize=10)
    ax.set_xlabel("X1")
    ax.set_ylabel("X2")
    ax.set_xlim(-0.1, 1.1)
    ax.set_ylim(-0.1, 1.1)
    ax.set_aspect('equal', adjustable='box')
    ax.grid(True, linestyle='--', alpha=0.6)
    ax.legend(fontsize='small')
    plot_idx += 1

# Hide any unused subplots
for j in range(plot_idx, len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout(rect=[0, 0, 1, 0.96]) # Adjust layout to make space for supitle
plt.suptitle(f"Impact of Adding a Point to a 2x2 Grid Design (q={q}, p={p})", fontsize=14)
plt.show()

```



Summary Table (Conceptual):

Scenario	Initial Φ_q	Augmented Φ_q	Change
Baseline (2x2 Grid)	2.236	—	—
Scenario 1: Add to Center	2.236	3.606	+1.369
Scenario 2: Add Close to Existing (Cluster)	2.236	7.619	+5.383
Scenario 3: Add on Edge	2.236	3.821	+1.585

79.5. A Sample-Size Invariant Version of the Morris-Mitchell Criterion

79.5.1. Comparison of `mmphi()` and `mmphi_intensive()`

The Morris-Mitchell criterion is a widely used metric for evaluating the space-filling properties of Latin hypercube sampling designs. However, it is sensitive to the number of points in the design, which can lead to misleading comparisons between designs with different sample sizes. To address this issue, a sample-size invariant version of the Morris-Mitchell criterion has been proposed. It is available in the `spotpython` package as `mmphi_intensive()`, see [SOURCE].

The functions `mmphi()` and `mmphi_intensive()` both calculate a Morris-Mitchell criterion, but they differ in their normalization, which makes `mmphi_intensive()` invariant to the sample size.

Let X be a sampling plan with n points $\{x_1, x_2, \dots, x_n\}$ in a k -dimensional space. Let $d_{ij} = \|x_i - x_j\|_p$ be the p -norm distance between points x_i and x_j . Let J_l be the multiplicity of the l -th unique distance d_l among all pairs of points in X . Let m be the total number of unique distances.

1. `mmphi()` (Morris-Mitchell Criterion Φ_q)

The `mmphi()` function, as defined in the context and implemented in `sampling.py`, calculates the Morris-Mitchell criterion Φ_q as:

$$\Phi_q(X) = \left(\sum_{l=1}^m J_l d_l^{-q} \right)^{1/q},$$

where:

- J_l is the number of pairs of points separated by the unique distance d_l .
- d_l are the unique pairwise distances.
- q is a user-defined exponent (typically $q > 0$).

This formulation is directly based on the sum of inverse powers of distances. The value of Φ_q is generally dependent on the number of points n in the design X , as the sum $\sum J_l d_l^{-q}$ will typically increase with more points (and thus more pairs).

2. `mmphi_intensive()` (Intensive Morris-Mitchell Criterion)

The `mmphi_intensive()` function, as implemented in `sampling.py` calculates a sample-size invariant version of the Morris-Mitchell criterion, which will be referred to as Φ_q^I . The formula is:

$$\Phi_q^I(X) = \left(\frac{1}{M} \sum_{l=1}^m J_l d_l^{-q} \right)^{1/q}$$

where:

- $M = \binom{n}{2} = \frac{n(n-1)}{2}$ is the total number of unique pairs of points in the design X .
- The other terms J_l , d_l , q are the same as in `mmphi()`.

The key mathematical difference is the normalization factor $\frac{1}{M}$ inside the parentheses before the outer exponent $1/q$ is applied.

- `mmphi()`: Calculates $(\text{SumTerm})^{1/q}$, where $\text{SumTerm} = \sum J_l d_l^{-q}$.
- `mmphi_intensive()`: Calculates $\left(\frac{\text{SumTerm}}{M} \right)^{1/q}$.

By dividing the sum $\sum J_l d_l^{-q}$ by M (the total number of pairs), `mmphi_intensive()` effectively calculates an *average* contribution per pair to the $-q$ -th power of distance, before taking the q -th root. This normalization makes the criterion less dependent on the absolute number of points n and allows for more meaningful comparisons of space-fillingness between designs of different sizes. A smaller value indicates a better (more space-filling) design for both criteria.

79.5.2. Plotting the Two Morris-Mitchell Criteria for Different Sample Sizes

Figure 79.11 shows the comparison of the two Morris-Mitchell criteria for different sample sizes using the `plot_mmphi_vs_n_lhs` function. The red line represents the standard Morris-Mitchell criterion, while the blue line represents the sample-size invariant version. Note the difference in the y-axis scales, which highlights how the sample-size invariant version remains consistent across varying sample sizes.

```
def plot_mmphi_vs_n_lhs(k_dim: int,
                        seed: int,
                        n_min: int = 10,
                        n_max: int = 100,
                        n_step: int = 5,
                        q_phi: float = 2.0,
                        p_phi: float = 2.0):
    """
    Generates LHS designs for varying n, calculates mmphi and mmphi_intensive,
    and plots them against the number of samples (n).

    Args:
        k_dim (int): Number of dimensions for the LHS design.
```

79.5. A Sample-Size Invariant Version of the Morris-Mitchell Criterion

```

seed (int): Random seed for reproducibility.
n_min (int): Minimum number of samples.
n_max (int): Maximum number of samples.
n_step (int): Step size for increasing n.
q_phi (float): Exponent q for the Morris-Mitchell criteria.
p_phi (float): Distance norm p for the Morris-Mitchell criteria.
"""
n_values = list(range(n_min, n_max + 1, n_step))
if not n_values:
    print("Warning: n_values list is empty. Check n_min, n_max, and n_step.")
    return
mmphi_results = []
mmphi_intensive_results = []
# lhs_generator removed

print(f"Calculating for n from {n_min} to {n_max} with step {n_step}...")
for n_points in n_values:
    if n_points < 2 : # mmphi requires at least 2 points to calculate distances
        print(f"Skipping n={n_points} as it's less than 2.")
        mmphi_results.append(np.nan)
        mmphi_intensive_results.append(np.nan)
        continue
    try:
        X_design = generate_qmc_lhs_design([(0, 1)] * k_dim, n_design=n_points, seed=seed)
        phi = mmphi(X_design, q=q_phi, p=p_phi)
        phi_intensive, _, _ = mmphi_intensive(X_design, q=q_phi, p=p_phi)
        mmphi_results.append(phi)
        mmphi_intensive_results.append(phi_intensive)
    except Exception as e:
        print(f"Error calculating for n={n_points}: {e}")
        mmphi_results.append(np.nan)
        mmphi_intensive_results.append(np.nan)

fig, ax1 = plt.subplots(figsize=(9, 6))

color = 'tab:red'
ax1.set_xlabel('Number of Samples (n)')
ax1.set_ylabel('mmphi (Phiq)', color=color)
ax1.plot(n_values, mmphi_results, color=color, marker='o', linestyle='-', label='mmphi (Phiq)')
ax1.tick_params(axis='y', labelcolor=color)
ax1.grid(True, linestyle='--', alpha=0.7)

ax2 = ax1.twinx() # instantiate a second axes that shares the same x-axis
color = 'tab:blue'

```

79. Sampling Plans

```
ax2.set_ylabel('mmphi_intensive (PhiqI)', color=color) # we already handled the x
ax2.plot(n_values, mmphi_intensive_results, color=color, marker='x', linestyle='--')
ax2.tick_params(axis='y', labelcolor=color)

fig.tight_layout() # otherwise the right y-label is slightly clipped
plt.title(f'Morris-Mitchell Criteria vs. Number of Samples (n)\nLHS (k={k_dim}, q={q_dim})')
# Add legends
lines, labels = ax1.get_legend_handles_labels()
lines2, labels2 = ax2.get_legend_handles_labels()
ax2.legend(lines + lines2, labels + labels2, loc='best')
plt.show()
```

```
N_DIM = 2
RANDOM_SEED = 42
plot_mmphi_vs_n_lhs(k_dim=N_DIM, seed=RANDOM_SEED, n_min=10, n_max=100, n_step=5)
```

Calculating for n from 10 to 100 with step 5...




Figure 79.11.: Comparison of the two Morris-Mitchell Criteria for Different Sample Sizes

79.6. Jupyter Notebook

i Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

80. Constructing a Surrogate

Note

This section is based on chapter 2 in Forrester et al. (2008).

Definition 80.1 (Black Box Problem). We are trying to learn a mapping that converts the vector $\vec{x}$ into a scalar output y , i.e., we are trying to learn a function

$$y = f(x).$$

If function is hidden (“lives in a black box”), so that the physics of the problem is not known, the problem is called a black box problem.

This black box could take the form of either a physical or computer experiment, for example, a finite element code, which calculates the maximum stress (σ) for given product dimensions ($\vec{x}$).

Definition 80.2 (Generic Solution). The generic solution method is to collect the output values $y^{(1)}, y^{(2)}, \dots, y^{(n)}$ that result from a set of inputs $\vec{x}^{(1)}, \vec{x}^{(2)}, \dots, \vec{x}^{(n)}$ and find a best guess $\hat{f}(\vec{x})$ for the black box mapping f , based on these known observations.

80.1. Stage One: Preparing the Data and Choosing a Modelling Approach

The first step is the identification, through a small number of observations, of the inputs that have a significant impact on f ; that is the determination of the shortest design variable vector $\vec{x} = \{x_1, x_2, \dots, x_k\}^T$ that, by sweeping the ranges of all of its variables, can still elicit most of the behavior the black box is capable of. The ranges of the various design variables also have to be established at this stage.

The second step is to recruit n of these k -vectors into a list

$$X = \{\vec{x}^{(1)}, \vec{x}^{(2)}, \dots, \vec{x}^{(n)}\}^T,$$

where each $\vec{x}^{(i)}$ is a k -vector. The corresponding responses are collected in a vector such that this represents the design space as thoroughly as possible.

80. Constructing a Surrogate

In the surrogate modeling process, the number of samples n is often limited, as it is constrained by the computational cost (money and/or time) associated with obtaining each observation.

It is advisable to scale $\vec{x}$ at this stage into the unit cube $[0, 1]^k$, a step that can simplify the subsequent mathematics and prevent multidimensional scaling issues.

We now focus on the attempt to learn f through data pairs

$$\{(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(n)}, y^{(n)})\}.$$

This supervised learning process essentially involves searching across the space of possible functions $\hat{f}$ that would replicate observations of f . This space of functions is infinite. Any number of hypersurfaces could be drawn to pass through or near the known observations, accounting for experimental error. However, most of these would generalize poorly; they would be practically useless at predicting responses at new sites, which is the ultimate goal.

Example 80.1 (The Needle(s) in the Haystack Function). An extreme example is the ‘needle(s) in the haystack’ function:

$$f(x) = \begin{cases} y^{(1)}, & \text{if } x = \vec{x}^{(1)} \\ y^{(2)}, & \text{if } x = \vec{x}^{(2)} \\ \vdots & \\ y^{(n)}, & \text{if } x = \vec{x}^{(n)} \\ 0, & \text{otherwise.} \end{cases}$$

While this predictor reproduces all training data, it seems counter-intuitive and unsettling to predict 0 everywhere else for most engineering functions. Although there is a small chance that the function genuinely resembles the equation above and we sampled exactly where the needles are, it is highly unlikely.

There are countless other configurations, perhaps less contrived, that still generalize poorly. This suggests a need for systematic means to filter out nonsensical predictors. In our approach, we embed the structure of f into the model selection algorithm and search over its parameters to fine-tune the approximation to observations. For instance, consider one of the simplest models,

$$f(x, \vec{w}) = \vec{w}^T \vec{x} + v. \quad (80.1)$$

Learning f with this model implies that its structure—a hyperplane—is predetermined, and the fitting process involves finding the $k + 1$ parameters (the slope vector $\vec{w}$ and the intercept v) that best fit the data. This will be accomplished in Stage Two.

Complicating this further is the noise present in observed responses (we assume design vectors $\vec{x}$ are not corrupted). Here, we focus on learning from such data, which sometimes risks overfitting.

Definition 80.3 (Overfitting). Overfitting occurs when the model becomes too flexible and captures not only the underlying trend but also the noise in the data.

In the surrogate modeling process, the second stage as described in Section 80.2, addresses this issue of complexity control by estimating the parameters of the fixed structure model. However, foresight is necessary even at the model type selection stage.

Model selection often involves physics-based considerations, where the modeling technique is chosen based on expected underlying responses.

Example 80.2 (Model Selection). Modeling stress in an elastically deformed solid due to small strains may justify using a simple linear approximation. Without insights into the physics, and if one fails to account for the simplicity of the data, a more complex and excessively flexible model may be incorrectly chosen. Although parameter estimation might still adjust the approximation to become linear, an opportunity to develop a simpler and robust model may be lost.

- Simple linear (or polynomial) models, despite their lack of flexibility, have advantages like applicability in further symbolic computations.
- Conversely, if we incorrectly assume a quadratic process when multiple peaks and troughs exist, the parameter estimation stage will not compensate for an unsuitable model choice. A quadratic model is too rigid to fit a multimodal function, regardless of parameter adjustments.

80.2. Stage Two: Parameter Estimation and Training

Assuming that Stage One helped identify the k critical design variables, acquire the learning data set, and select a generic model structure $f(\vec{x}, \vec{w})$, the task now is to estimate parameters $\vec{w}$ to ensure the model fits the data optimally. Among several estimation criteria, we will discuss two methods here.

Definition 80.4 (Maximum Likelihood Estimation). Given a set of parameters $\vec{w}$, the model $f(\vec{x}, \vec{w})$ allows computation of the probability of the data set

$$\{(\vec{x}^{(1)}, y^{(1)} \pm \epsilon), (\vec{x}^{(2)}, y^{(2)} \pm \epsilon), \dots, (\vec{x}^{(n)}, y^{(n)} \pm \epsilon)\}$$

resulting from f (where ϵ is a small error margin around each data point).

i Maximum Likelihood Estimation

Section 88.18 presents a more detailed discussion of the maximum likelihood estimation (MLE) method.

80. Constructing a Surrogate

Taking Equation 88.36 and assuming errors ϵ are independently and normally distributed with standard deviation σ , the probability of the data set is given by:

$$P = \frac{1}{(2\pi\sigma^2)^{n/2}} \exp \left[-\frac{1}{2\sigma^2} \sum_{i=1}^n (y^{(i)} - f(\vec{x}^{(i)}, \vec{w}))^2 \epsilon \right].$$

Intuitively, this is equivalent to the likelihood of the parameters given the data. Accepting this intuitive relationship as a mathematical one aids in model parameter estimation. This is achieved by maximizing the likelihood or, more conveniently, minimizing the negative of its natural logarithm:

$$\min_{\vec{w}} \sum_{i=1}^n \frac{[y^{(i)} - f(\vec{x}^{(i)}, \vec{w})]^2}{2\sigma^2} + \frac{n}{2} \ln \epsilon. \quad (80.2)$$

If we assume σ and ϵ are constants, Equation 80.2 simplifies to the well-known least squares criterion:

$$\min_{\vec{w}} \sum_{i=1}^n [y^{(i)} - f(\vec{x}^{(i)}, \vec{w})]^2.$$

Cross-validation is another method used to estimate model performance.

Definition 80.5 (Cross-Validation). Cross-validation splits the data randomly into q roughly equal subsets, and then cyclically removing each subset and fitting the model to the remaining $q - 1$ subsets. A loss function L is then computed to measure the error between the predictor and the withheld subset for each iteration, with contributions summed over all q iterations. More formally, if a mapping $\theta : \{1, \dots, n\} \rightarrow \{1, \dots, q\}$ describes the allocation of the n training points to one of the q subsets and $f^{(-\theta(i))}(\vec{x})$ is the predicted value by removing the subset $\theta(i)$ (i.e., the subset where observation i belongs), the cross-validation measure, used as an estimate of prediction error, is:

$$CV = \frac{1}{n} \sum_{i=1}^n L(y^{(i)}, f^{(-\theta(i))}(\vec{x}^{(i)})). \quad (80.3)$$

Introducing the squared error as the loss function and considering our generic model f still dependent on undetermined parameters, we write Equation 80.3 as:

$$CV = \frac{1}{n} \sum_{i=1}^n [y^{(i)} - f^{(-\theta(i))}(\vec{x}^{(i)})]^2. \quad (80.4)$$

The extent to which Equation 80.4 is an unbiased estimator of true risk depends on q . It is shown that if $q = n$, the leave-one-out cross-validation (LOOCV) measure is

almost unbiased. However, LOOCV can have high variance because subsets are very similar. Hastie et al. (2017) suggest using compromise values like $q = 5$ or $q = 10$. Using fewer subsets also reduces the computational cost of the cross-validation process, see also Arlot et al. (2010) and Kohavi (1995).

80.3. Stage Three: Model Testing

If there is a sufficient amount of observational data, a random subset should be set aside initially for model testing. Hastie et al. (2017) recommend setting aside approximately $0.25n$ of $\vec{x} \rightarrow y$ pairs for testing purposes. These observations must remain untouched during Stages One and Two, as their sole purpose is to evaluate the testing error—the difference between true and approximated function values at the test sites—once the model has been built. Interestingly, if the main goal is to construct an initial surrogate for seeding a global refinement criterion-based strategy (as discussed in Section 3.2 in Forrester et al. (2008)), the model testing phase might be skipped.

It is noted that, ideally, parameter estimation (Stage Two) should also rely on a separate subset. However, observational data is rarely abundant enough to afford this luxury (if the function is cheap to evaluate and evaluation sites are selectable, a surrogate model might not be necessary).

When data are available for model testing and the primary objective is a globally accurate model, using either a root mean square error (RMSE) metric or the correlation coefficient (r^2) is recommended. To test the model, a test data set of size n_t is used alongside predictions at the corresponding locations to calculate these metrics.

The RMSE is defined as follows:

Definition 80.6 (Root Mean Square Error (RMSE)).

$$\text{RMSE} = \sqrt{\frac{1}{n_t} \sum_{i=1}^{n_t} (y^{(i)} - \hat{y}^{(i)})^2},$$

Ideally, the RMSE should be minimized, acknowledging its limitation by errors in the objective function f calculation. If the error level is known, like a standard deviation, the aim might be to achieve an RMSE within this value. Often, the target is an RMSE within a specific percentage of the observed data's objective value range.

The squared correlation coefficient r , see Equation 88.13, between the observed y and predicted $\hat{y}$ values can be computed as:

$$r^2 = \left(\frac{\text{cov}(y, \hat{y})}{\sqrt{\text{var}(y)\text{var}(\hat{y})}} \right)^2, \quad (80.5)$$

80. Constructing a Surrogate

Equation 80.5 and can be expanded as:

$$r^2 = \left(\frac{n_t \sum_{i=1}^{n_t} y^{(i)} \hat{y}^{(i)} - \sum_{i=1}^{n_t} y^{(i)} \sum_{i=1}^{n_t} \hat{y}^{(i)}}{\sqrt{\left(n_t \sum_{i=1}^{n_t} (y^{(i)})^2 - \left(\sum_{i=1}^{n_t} y^{(i)} \right)^2 \right) \left(n_t \sum_{i=1}^{n_t} (\hat{y}^{(i)})^2 - \left(\sum_{i=1}^{n_t} \hat{y}^{(i)} \right)^2 \right)}} \right)^2.$$

The correlation coefficient r^2 does not require scaling the data sets and only compares landscape shapes, not values. An $r^2 > 0.8$ typically indicates a surrogate with good predictive capability.

The methods outlined provide quantitative assessments of model accuracy, yet visual evaluations can also be insightful. In general, the RMSE will not reach zero but will stabilize around a low value. At this point, the surrogate model is saturated with data, and further additions do not enhance the model globally (though local improvements can occur at newly added points if using an interpolating model).

80.3.0.1. The Tea and Sugar Analogy

Forrester et al. (2008) illustrates this saturation point using a comparison with a cup of tea and sugar. The tea represents the surrogate model, and sugar represents data. Initially, the tea is unsweetened, and adding sugar increases its sweetness. Eventually, a saturation point is reached where no more sugar dissolves, and the tea cannot get any sweeter. Similarly, a more flexible model, like one with additional parameters or employing interpolation rather than regression, can increase the saturation point—akin to making a hotter cup of tea for dissolving more sugar.

80.4. Jupyter Notebook

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

81. Response Surface Methods

This part deals with numerical implementations of optimization methods. The goal is to understand the implementation of optimization methods and to solve real-world problems numerically and efficiently. We will focus on the implementation of surrogate models, because they are the most efficient way to solve real-world problems.

Starting point is the well-established response surface methodology (RSM). It will be extended to the design and analysis of computer experiments (DACE). The DACE methodology is a modern extension of the response surface methodology. It is based on the use of surrogate models, which are used to replace the real-world problem with a simpler problem. The simpler problem is then solved numerically. The solution of the simpler problem is then used to solve the real-world problem.

! Numerical methods: Goals

- Understand implementation of optimization methods
- Solve real-world problems numerically and efficiently

81.1. What is RSM?

Response Surface Methods (RSM) refer to a collection of statistical and mathematical tools that are valuable for developing, improving, and optimizing processes. The overarching theme of RSM involves studying how input variables that control a product or process can potentially influence a response that measures performance or quality characteristics.

The advantages of RSM include a rich literature, well-established methods often used in manufacturing, the importance of careful experimental design combined with a well-understood model, and the potential to add significant value to scientific inquiry, process refinement, optimization, and more. However, there are also drawbacks to RSM, such as the use of simple and crude surrogates, the hands-on nature of the methods, and the limitation of local methods.

RSM is related to various fields, including Design of Experiments (DoE), quality management, reliability, and productivity. Its applications are widespread in industry and manufacturing, focusing on designing, developing, and formulating new products and improving existing ones, as well as from laboratory research. RSM is commonly

81. Response Surface Methods

applied in domains such as materials science, manufacturing, applied chemistry, climate science, and many others.

An example of RSM involves studying the relationship between a response variable, such as yield (y) in a chemical process, and two process variables: reaction time (ξ_1) and reaction temperature (ξ_2). The provided code illustrates this scenario, following a variation of the so-called “banana function.”

In the context of visualization, RSM offers the choice between 3D plots and contour plots. In a 3D plot, the independent variables ξ_1 and ξ_2 are represented, with y as the dependent variable.

```
import numpy as np
import matplotlib.pyplot as plt


def fun_rosen(x1, x2):
    b = 10
    return (x1-1)**2 + b*(x2-x1**2)**2

fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')
x = np.arange(-2.0, 2.0, 0.05)
y = np.arange(-1.0, 3.0, 0.05)
X, Y = np.meshgrid(x, y)
zs = np.array(fun_rosen(np.ravel(X), np.ravel(Y)))
Z = zs.reshape(X.shape)

ax.plot_surface(X, Y, Z)

ax.set_xlabel('X1')
ax.set_ylabel('X2')
ax.set_zlabel('Y')

plt.show()
```




- contour plot example:

- x_1 and x_2 are the independent variables
- y is the dependent variable

```
import numpy as np
import matplotlib.cm as cm
import matplotlib.pyplot as plt

delta = 0.025
x1 = np.arange(-2.0, 2.0, delta)
x2 = np.arange(-1.0, 3.0, delta)
X1, X2 = np.meshgrid(x1, x2)
Y = fun_rosen(X1, X2)
fig, ax = plt.subplots()
CS = ax.contour(X1, X2, Y, 50)
ax.clabel(CS, inline=True, fontsize=10)
ax.set_title("Rosenbrock's Banana Function")
```

```
Text(0.5, 1.0, "Rosenbrock's Banana Function")
```



- Visual inspection: yield is optimized near (ξ_1, ξ_2)

81.1.1. Visualization: Problems in Practice

- True response surface is unknown in practice
- When yield evaluation is not as simple as a toy banana function, but a process requiring care to monitor, reconfigure and run, it's far too expensive to observe over a dense grid
- And, measuring yield may be a noisy/inexact process
- That's where stats (RSM) comes in

81.1.2. RSM: Strategies

- RSMs consist of experimental strategies for
- **exploring** the space of the process (i.e., independent/input) variables (above ξ_1 and ξ_2)
- empirical statistical **modeling** targeted toward development of an appropriate approximating relationship between the response (yield) and process variables local to a study region of interest

- **optimization** methods for sequential refinement in search of the levels or values of process variables that produce desirable responses (e.g., that maximize yield or explain variation)
- RSM used for fitting an Empirical Model
- True response surface driven by an unknown physical mechanism
- Observations corrupted by noise
- Helpful: fit an empirical model to output collected under different process configurations
- Consider response Y that depends on controllable input variables $\xi_1, \xi_2, \dots, \xi_m$
- RSM: Equations of the Empirical Model
 - $Y = f(\xi_1, \xi_2, \dots, \xi_m) + \epsilon$
 - $\mathbb{E}\{Y\} = \eta = f(\xi_1, \xi_2, \dots, \xi_m)$
 - ϵ is treated as zero mean idiosyncratic noise possibly representing
 - * inherent variation, or
 - * the effect of other systems or
 - * variables not under our purview at this time

81.1.3. RSM: Noise in the Empirical Model

- Typical simplifying assumption: $\epsilon \sim N(0, \sigma^2)$
- We seek estimates for f and σ^2 from noisy observations Y at inputs ξ

81.1.4. RSM: Natural and Coded Variables

- Inputs $\xi_1, \xi_2, \dots, \xi_m$ called **natural variables**:
 - expressed in natural units of measurement, e.g., degrees Celsius, pounds per square inch (psi), etc.
- Transformed to **coded variables** $x_1, x_2, \dots, x_m$:
 - to mitigate hassles and confusion that can arise when working with a multitude of scales of measurement
- Typical **Transformations** offering dimensionless inputs $x_1, x_2, \dots, x_m$
 - in the unit cube, or
 - scaled to have a mean of zero and standard deviation of one, are common choices.
- Empirical model becomes $\eta = f(x_1, x_2, \dots, x_m)$

81.1.5. RSM Low-order Polynomials

- Low-order polynomial make the following simplifying Assumptions
 - Learning about f is lots easier if we make some simplifying approximations
 - Appealing to **Taylor's theorem**, a low-order polynomial in a small, localized region of the input (x) space is one way forward
 - Classical RSM:
 - * disciplined application of **local analysis** and
 - * **sequential refinement** of locality through conservative extrapolation
 - Inherently a **hands-on process**

81.2. First-Order Models (Main Effects Model)

- **First-order model** (sometimes called main effects model) useful in parts of the input space where it's believed that there's little curvature in f :

$$\eta = \beta_0 + \beta_1 x_1 + \beta_2 x_2$$

- For example:

$$\eta = 50 + 8x_1 + 3x_2$$

- In practice, such a surface would be obtained by fitting a model to the outcome of a designed experiment
- First-Order Model in python Evaluated on a Grid
- Evaluate model on a grid in a double-unit square centered at the origin
- Coded units are chosen arbitrarily, although one can imagine deploying this approximating function nearby $x^{(0)} = (0, 0)$


```
def fun_1(x1,x2):
    return 50 + 8*x1 + 3*x2
```

```
import numpy as np
import matplotlib.cm as cm
import matplotlib.pyplot as plt

delta = 0.025
x1 = np.arange(-1.0, 1.0, delta)
x2 = np.arange(-1.0, 1.0, delta)
X1, X2 = np.meshgrid(x1, x2)
Y = fun_1(X1,X2)
fig, ax = plt.subplots()
CS = ax.contour(X1, X2, Y)
ax.clabel(CS, inline=True, fontsize=10)
ax.set_title('First Order Model: $50 + 8x_1 + 3x_2$')
```

81.2. First-Order Models (Main Effects Model)

```
Text(0.5, 1.0, 'First Order Model: $50 + 8x_1 + 3x_2$')
```



81.2.1. First-Order Model Properties

- First-order model in 2d traces out a **plane** in $y \times (x_1, x_2)$ space
- Only be appropriate for the most trivial of response surfaces, even when applied in a highly localized part of the input space
- Adding **curvature** is key to most applications:
 - First-order model with **interactions** induces limited degree of curvature via different rates of change of y as x_1 is varied for fixed x_2 , and vice versa:

$$\eta = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_{12} x_{12}$$

- For example $\eta = 50 + 8x_1 + 3x_2 - 4x_1x_2$

81.2.2. First-order Model with Interactions in python


- Code below facilitates evaluations for pairs (x_1, x_2)
- Responses may be observed over a mesh in the same double-unit square

81. Response Surface Methods

```
def fun_11(x1,x2):  
    return 50 + 8 * x1 + 3 * x2 - 4 * x1 * x2
```

```
import numpy as np  
import matplotlib.cm as cm  
import matplotlib.pyplot as plt  
  
delta = 0.025  
x1 = np.arange(-2.0, 2.0, delta)  
x2 = np.arange(-2.0, 2.0, delta)  
X1, X2 = np.meshgrid(x1, x2)  
Y = fun_11(X1,X2)  
fig, ax = plt.subplots()  
CS = ax.contour(X1, X2, Y, 20)  
ax.clabel(CS, inline=True, fontsize=10)  
ax.set_title('First Order Model with Interactions')
```

```
Text(0.5, 1.0, 'First Order Model with Interactions')
```



81.2.3. Observations: First-Order Model with Interactions

- Mean response η is increasing marginally in both x_1 and x_2 , or conditional on a fixed value of the other until x_1 is 0.75
- Rate of increase slows as both coordinates grow simultaneously since the coefficient in front of the interaction term x_1x_2 is negative
- Compared to the first-order model (without interactions): surface is far more useful locally
- Least squares regressions often flag up significant interactions when fit to data collected on a design far from local optima

81.3. Second-Order Models

- Second-order model may be appropriate near local optima where f would have substantial curvature:

$$\eta = \beta_0 + \beta_1x_1 + \beta_2x_2 + \beta_{11}x_1^2 + \beta_{22}x_2^2 + \beta_{12}x_1x_2$$

- For example

$$\eta = 50 + 8x_1 + 3x_2 - 7x_1^2 - 3x_2^2 - 4x_1x_2$$

- Implementation of the Second-Order Model as `fun_2()`.


```
def fun_2(x1,x2):
    return 50 + 8 * x1 + 3 * x2 - 7 * x1**2 - 3*x2**2 - 4 * x1 * x2
```

```
import numpy as np
import matplotlib.cm as cm
import matplotlib.pyplot as plt

delta = 0.025
x1 = np.arange(-2.0, 2.0, delta)
x2 = np.arange(-2.0, 2.0, delta)
X1, X2 = np.meshgrid(x1, x2)
Y = fun_2(X1,X2)
fig, ax = plt.subplots()
CS = ax.contour(X1, X2, Y, 20)
ax.clabel(CS, inline=True, fontsize=10)
ax.set_title('Second Order Model with Interactions. Maximum near about $(0.6,0.2)$')
```

```
Text(0.5, 1.0, 'Second Order Model with Interactions. Maximum near about $(0.6,0.2)$')
```

Second Order Model with Interactions. Maximum near about (0.6, 0.2)



81.3.1. Second-Order Models: Properties

- Not all second-order models would have a single stationary point (in RSM jargon called “a simple maximum”)
- In “yield maximizing” setting we’re presuming response surface is **concave** down from a global viewpoint
 - even though local dynamics may be more nuanced
- Exact criteria depend upon the eigenvalues of a certain matrix built from those coefficients
- Box and Draper (2007) provide a diagram categorizing all of the kinds of second-order surfaces in RSM analysis, where finding local maxima is the goal

81.3.2. Example: Stationary Ridge

- Example set of coefficients describing what’s called a **stationary ridge** is provided by the code below

```
def fun_ridge(x1, x2):
    return 80 + 4*x1 + 8*x2 - 3*x1**2 - 12*x2**2 - 12*x1*x2
```


```

import numpy as np
import matplotlib.cm as cm
import matplotlib.pyplot as plt

delta = 0.025
x1 = np.arange(-2.0, 2.0, delta)
x2 = np.arange(-2.0, 2.0, delta)
X1, X2 = np.meshgrid(x1, x2)
Y = fun_ridge(X1, X2)
fig, ax = plt.subplots()
CS = ax.contour(X1, X2, Y, 20)
ax.clabel(CS, inline=True, fontsize=10)
ax.set_title('Example of a stationary ridge')

```

```
Text(0.5, 1.0, 'Example of a stationary ridge')
```



81.3.3. Observations: Second-Order Model (Ridge)

- **Ridge:** a whole line of stationary points corresponding to maxima
- Situation means that the practitioner has some flexibility when it comes to optimizing:

81. Response Surface Methods

- can choose the precise setting of (x_1, x_2) either arbitrarily or (more commonly) by consulting some tertiary criteria


81.3.4. Example: Rising Ridge

- An example of a rising ridge is implemented by the code below.

```
def fun_ridge_rise(x1, x2):  
    return 80 - 4*x1 + 12*x2 - 3*x1**2 - 12*x2**2 - 12*x1*x2
```

```
import numpy as np  
import matplotlib.cm as cm  
import matplotlib.pyplot as plt  
  
delta = 0.025  
x1 = np.arange(-2.0, 2.0, delta)  
x2 = np.arange(-2.0, 2.0, delta)  
X1, X2 = np.meshgrid(x1, x2)  
Y = fun_ridge_rise(X1, X2)  
fig, ax = plt.subplots()  
CS = ax.contour(X1, X2, Y, 20)  
ax.clabel(CS, inline=True, fontsize=10)  
ax.set_title('Rising ridge:  $\eta = 80 + 4x_1 + 8x_2 - 3x_1^2 - 12x_2^2 - 12x_1x_2$ ')
```

```
Text(0.5, 1.0, 'Rising ridge:  $\eta = 80 + 4x_1 + 8x_2 - 3x_1^2 - 12x_2^2 - 12x_1x_2$ ')
```



81.3.5. Summary: Rising Ridge

- The stationary point is remote to the study region
- Continuum of (local) stationary points along any line going through the 2d space, excepting one that lies directly on the ridge
- Although estimated response will increase while moving along the axis of symmetry toward its stationary point, this situation indicates
 - either a poor fit by the approximating second-order function, or
 - that the study region is not yet precisely in the vicinity of a local optima—often both.

81.3.6. Falling Ridge

- Inversion of a rising ridge is a falling ridge
- Similarly indicating one is far from local optima, except that the response decreases as you move toward the stationary point
- Finding a falling ridge system can be a back-to-the-drawing-board affair.

81.3.7. Saddle Point


- Finally, we can get what's called a saddle or minimax system.

```
def fun_saddle(x1, x2):
    return 80 + 4*x1 + 8*x2 - 2*x2**2 - 12*x1*x2

import numpy as np
import matplotlib.cm as cm
import matplotlib.pyplot as plt

delta = 0.025
x1 = np.arange(-2.0, 2.0, delta)
x2 = np.arange(-2.0, 2.0, delta)
X1, X2 = np.meshgrid(x1, x2)
Y = fun_saddle(X1, X2)
fig, ax = plt.subplots()
CS = ax.contour(X1, X2, Y, 20)
ax.clabel(CS, inline=True, fontsize=10)
ax.set_title('Saddle Point:  $\eta = 80 + 4x_1 + 8x_2 - 2x_2^2 - 12x_1x_2$ ')
```

```
Text(0.5, 1.0, 'Saddle Point:  $\eta = 80 + 4x_1 + 8x_2 - 2x_2^2 - 12x_1x_2$ ')
```



81.3.8. Interpretation: Saddle Points

- Likely further data collection, and/or outside expertise, is needed before determining a course of action in this situation

81.3.9. Summary: Ridge Analysis

- Finding a simple maximum, or stationary ridge, represents ideals in the spectrum of second-order approximating functions
- But getting there can be a bit of a slog
- Using models fitted from data means uncertainty due to noise, and therefore uncertainty in the type of fitted second-order model
- A ridge analysis attempts to offer a principled approach to navigating uncertainties when one is seeking local maxima
- The two-dimensional setting exemplified above is convenient for visualization, but rare in practice
- Complications compound when studying the effect of more than two process variables

81.4. General RSM Models

- General **first-order model** on m process variables $x_1, x_2, \dots, x_m$ is

$$\eta = \beta_0 + \beta_1 x_1 + \dots + \beta_m x_m$$

- General **second-order model** on m process variables

$$\eta = \beta_0 + \sum_{j=1}^m \beta_j x_j + \sum_{j=1}^m \beta_{jj} x_j^2 + \sum_{j=2}^m \sum_{k=1}^{j-1} \beta_{jk} x_k x_j.$$

81.4.1. Ordinary Least Squares

- Inference from data is carried out by **ordinary least squares** (OLS)
- For an excellent review including R examples, see Sheather (2009)
- OLS and maximum likelihood estimators (MLEs) are in the typical Gaussian linear modeling setup basically equivalent

81.5. General Linear Regression

We are considering a model, which can be written in the form

$$Y = X\beta + \epsilon,$$

where Y is an $(n \times 1)$ vector of observations (responses), X is an $(n \times p)$ matrix of known form, β is a $(1 \times p)$ vector of unknown parameters, and ϵ is an $(n \times 1)$ vector of errors. Furthermore, $E(\epsilon) = 0$, $Var(\epsilon) = \sigma^2 I$ and the ϵ_i are uncorrelated.

Using the normal equations

$$(X'X)b = X'Y,$$

the solution is given by

$$b = (X'X)^{-1}X'Y.$$

Example 81.1 (Linear Regression).

```
import numpy as np
n = 8
X = np.linspace(0, 2*np.pi, n, endpoint=False).reshape(-1,1)
print(np.round(X, 2))
y = np.sin(X)
print(np.round(y, 2))
# fit an OLS model to the data, predict the response based on the 100 x values
m = 100
x = np.linspace(0, 2*np.pi, m, endpoint=False).reshape(-1,1)
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X, y)
y_pred = model.predict(x)
# visualize the data and the fitted model
import matplotlib.pyplot as plt
plt.scatter(X, y, color='black')
plt.plot(x, y_pred, color='blue', linewidth=1)
# add the ground truth (sine function) in orange
plt.plot(x, np.sin(x), color='orange', linewidth=1)
plt.show()
```


```
[[0.  ]
 [0.79]
 [1.57]
 [2.36]
```



```

[3.14]
[3.93]
[4.71]
[5.5 ]]
[[ 0. ]
 [ 0.71]
 [ 1.  ]
 [ 0.71]
 [ 0.  ]
 [-0.71]
 [-1.  ]
 [-0.71]]

```



81.6. Designs

- Important: Organize the data collection phase of a response surface study carefully
- **Design:** choice of x 's where we plan to observe y 's, for the purpose of approximating f
- Analyses and designs need to be carefully matched
- When using a first-order model, some designs are preferred over others
- When using a second-order model to capture curvature, a different sort of design is appropriate

81. Response Surface Methods

- Design choices often contain features enabling modeling assumptions to be challenged
 - e.g., to check if initial impressions are supported by the data ultimately collected

81.6.1. Different Designs

- **Screening designs:** determine which variables matter so that subsequent experiments may be smaller and/or more focused
- Then there are designs tailored to the form of model (first- or second-order, say) in the screened variables
- And then there are more designs still

81.7. RSM Experimentation

81.7.1. First Step

- RSM-based experimentation begins with a **first-order model**, possibly with interactions
- Presumption: current process operating **far from optimal** conditions
- Collect data and apply **method of steepest ascent** (gradient) on fitted surfaces to move to the optimum

81.7.2. Second Step

- Eventually, if all goes well after several such carefully iterated refinements, **second-order models** are used on appropriate designs in order to zero-in on ideal operating conditions
- Careful analysis of the fitted surface:
 - Ridge analysis with further refinement using gradients of, and
 - standard errors associated with, the fitted surfaces, and so on

81.7.3. Third Step

- Once the practitioner is satisfied with the full arc of
 - design(s),
 - fit(s), and
 - decision(s):

- A small experiment called **confirmation test** may be performed to check if the predicted optimal settings are realizable in practice

81.8. RSM: Review and General Considerations

- First Glimpse, RSM seems sensible, and pretty straightforward as quantitative statistics-based analysis goes
- But: RSM can get complicated, especially when input dimensions are not very low
- Design considerations are particularly nuanced, since the goal is to obtain reliable estimates of main effects, interaction, and curvature while minimizing sampling effort/expense
- RSM Downside: Inefficiency
 - Despite intuitive appeal, several RSM downsides become apparent upon reflection
 - Problems in practice
 - Stepwise nature of sequential decision making is inefficient:
 - * Not obvious how to re-use or update analysis from earlier phases, or couple with data from other sources/related experiments
- RSM Downside: Locality
 - In addition to being local in experiment-time (stepwise approach), it's local in experiment-space
 - Balance between
 - * exploration (maybe we're barking up the wrong tree) and
 - * exploitation (let's make things a little better) is modest at best
- RSM Downside: Expert Knowledge
 - Interjection of expert knowledge is limited to hunches about relevant variables (i.e., the screening phase), where to initialize search, how to design the experiments
 - Yet at the same time classical RSMs rely heavily on constant examination throughout stages of modeling and design and on the instincts of seasoned practitioners
- RSM Downside: Replicability
 - Parallel analyses, conducted according to the same best intentions, rarely lead to the same designs, model fits and so on
 - Sometimes that means they lead to different conclusions, which can be cause for concern

81.8.1. Historical Considerations about RSM

- In spite of those criticisms, however, there was historically little impetus to revise the status quo
- Classical RSM was comfortable in its skin, consistently led to improvements or compelling evidence that none can reasonably be expected
- But then in the late 20th century came an explosive expansion in computational capability, and with it a means of addressing many of those downsides

81.8.2. Status Quo

- Nowadays, field experiments and statistical models, designs and optimizations are coupled with mathematical models
- Simple equations are not regarded as sufficient to describe real-world systems anymore
- Physicists figured that out fifty years ago; industrial engineers followed, biologists, social scientists, climate scientists and weather forecasters, etc.
- Systems of equations are required, solved over meshes (e.g., finite elements), or stochastically interacting agents
- Goals for those simulation experiments are as diverse as their underlying dynamics
- Optimization of systems is common, e.g., to identify worst-case scenarios

81.8.3. The Role of Statistics

- Solving systems of equations, or interacting agents, requires computing
- Statistics involved at various stages:
 - choosing the mathematical model
 - solving by stochastic simulation (Monte Carlo)
 - designing the computer experiment
 - smoothing over idiosyncrasies or noise
 - finding optimal conditions, or
 - calibrating mathematical/computer models to data from field experiments

81.8.4. New RSM is needed: DACE

- Classical RSMs are not well-suited to any of those tasks, because
 - they lack the fidelity required to model these data
 - their intended application is too local
 - they're also too hands-on.

- Once computers are involved, a natural inclination is to automate—to remove humans from the loop and set the computer running on the analysis in order to maximize computing throughput, or minimize idle time
- **Design and Analysis of Computer Experiments** as a modern extension of RSM
- Experimentation is changing due to advances in machine learning
- **Gaussian process** (GP) regression is the canonical surrogate model
- Origins in geostatistics (gold mining)
- Wide applicability in contexts where prediction is king
- Machine learners exposed GPs as powerful predictors for all sorts of tasks:
 - from regression to classification,
 - active learning/sequential design,
 - reinforcement learning and optimization,
 - latent variable modeling, and so on

81.9. Exercises

1. Generate 3d Plots for the Contour Plots in this notebook.
2. Write a `plot_3d` function, that takes the objective function `fun` as an argument.
 - It should provide the following interface: `plot_3d(fun)`.
3. Write a `plot_contour` function, that takes the objective function `fun` as an argument:
 - It should provide the following interface: `plot_contour(fun)`.
4. Consider further arguments that might be useful for both function, e.g., ranges, size, etc.

81.10. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

82. Polynomial Models

i Note

- This section is based on chapter 2.2 in Forrester et al. (2008).
- The following Python packages are imported:

```
import numpy as np
import matplotlib.pyplot as plt
```

82.1. Fitting a Polynomial

We will consider one-variable cases, i.e., $k = 1$, first.

Let us consider the scalar-valued function $f : \mathbb{R} \rightarrow \mathbb{R}$ observed according to the sampling plan $X = \{x^{(1)}, x^{(2)} \dots, x^{(n)}\}^T$, yielding the responses $\vec{y} = \{y^{(1)}, y^{(2)}, \dots, y^{(n)}\}^T$.

A polynomial approximation of f of order m can be written as:

$$\hat{f}(x, m, \vec{w}) = \sum_{i=0}^m w_i x^i.$$

In the spirit of the earlier discussion of maximum likelihood parameter estimation, we seek to estimate $w = w_0, w_1, \dots, w_m^T$ through a least squares solution of:

$$\Phi \vec{w} = \vec{y}$$

where Φ is the Vandermonde matrix:

$$\Phi = \begin{bmatrix} 1 & x_1 & x_1^2 & \dots & x_1^m \\ 1 & x_2 & x_2^2 & \dots & x_2^m \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_n & x_n^2 & \dots & x_n^m \end{bmatrix}.$$

The maximum likelihood estimate of w is given by:

82. Polynomial Models

$$\vec{w} = (\Phi^T \Phi)^{-1} \Phi^T y,$$

where $\Phi^+ = (\Phi^T \Phi)^{-1} \Phi^T$ is the Moore-Penrose pseudo-inverse of Φ (see Section 85.3).

The polynomial approximation of order m is essentially a truncated Taylor series expansion. While higher values of m yield more accurate approximations, they risk overfitting the noise in the data.

To prevent this, we estimate m using cross-validation. This involves minimizing the cross-validation error over a discrete set of possible orders m (e.g., $m \in 1, 2, \dots, 15$).

For each m , the data is split into q subsets. The model is trained on $q - 1$ subsets, and the error is computed on the left-out subset. This process is repeated for all subsets, and the cross-validation error is summed. The order m with the smallest cross-validation error is chosen.

82.2. Polynomial Fitting in Python

82.2.1. Fitting the Polynomial

```
from sklearn.model_selection import KFold
def polynomial_fit(X, Y, max_order=15, q=5):
    """
    Fits a one-variable polynomial to one-dimensional data using cross-validation.

    Args:
        X (array-like): Training data vector (independent variable).
        Y (array-like): Training data vector (dependent variable).
        max_order (int): Maximum polynomial order to consider. Default is 15.
        q (int): Number of cross-validation folds. Default is 5.

    Returns:
        best_order (int): The optimal polynomial order.
        coeff (array): Coefficients of the best-fit polynomial.
        mnstd (tuple): Normalization parameters (mean, std) for X.
    """
    X = np.array(X)
    Y = np.array(Y)
    n = len(X)
    # Normalize X
    mnstd = (np.mean(X), np.std(X))
    X_norm = (X - mnstd[0]) / mnstd[1]
```



```

# Cross-validation setup
kf = KFold(n_splits=q, shuffle=True, random_state=42)
cross_val_errors = np.zeros(max_order)
for order in range(1, max_order + 1):
    fold_errors = []
    for train_idx, val_idx in kf.split(X_norm):
        X_train, X_val = X_norm[train_idx], X_norm[val_idx]
        Y_train, Y_val = Y[train_idx], Y[val_idx]
        # Fit polynomial
        coeff = np.polyfit(X_train, Y_train, order)
        # Predict on validation set
        Y_pred = np.polyval(coeff, X_val)
        # Compute mean squared error
        mse = np.mean((Y_val - Y_pred) ** 2)
        fold_errors.append(mse)
    cross_val_errors[order - 1] = np.mean(fold_errors)
# Find the best order
best_order = np.argmin(cross_val_errors) + 1
# Fit the best polynomial on the entire dataset
best_coeff = np.polyfit(X_norm, Y, best_order)
return best_order, best_coeff, mnstd

```

82.2.2. Explaining the k -fold Cross-Validation

The line

```

kf = KFold(n_splits=q, shuffle=True, random_state=42)

```

initializes a k -Fold cross-validator object from the `sklearn.model_selection` library. The `n_splits` parameter specifies the number of folds. The data will be divided into q parts. In each iteration of the cross-validation, one part will be used as the validation set, and the remaining $q-1$ parts will be used as the training set.

The `kf.split` method takes the dataset `X_norm` as input and yields pairs of index arrays for each fold: * `train_idx`: In each iteration, `train_idx` is an array containing the indices of the data points that belong to the training set for that specific fold. * `val_idx`: Similarly, `val_idx` is an array containing the indices of the data points that belong to the validation (or test) set for that specific fold.

The loop will run q times (the number of splits). In each iteration, a different fold serves as the validation set, while the other $q-1$ folds form the training set.

Here's a Python example to demonstrate the values of `train_idx` and `val_idx`:

```

# Sample data (e.g., X_norm)
X_norm = np.array([0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0])
print(f"Original data indices: {np.arange(len(X_norm))}\n")
# Number of splits (folds)
q = 3 # Let's use 3 folds for this example
# Initialize KFold
kf = KFold(n_splits=q, shuffle=True, random_state=42)
# Iterate through the splits and print the indices
fold_number = 1
for train_idx, val_idx in kf.split(X_norm):
    print(f"--- Fold {fold_number} ---")
    print(f"Train indices: {train_idx}")
    print(f"Validation indices: {val_idx}")
    print(f"Training data for this fold: {X_norm[train_idx]}")
    print(f"Validation data for this fold: {X_norm[val_idx]}\n")
    fold_number += 1

```

Original data indices: [0 1 2 3 4 5 6 7 8 9]

--- Fold 1 ---

Train indices: [2 3 4 6 7 9]

Validation indices: [0 1 5 8]

Training data for this fold: [0.3 0.4 0.5 0.7 0.8 1.]

Validation data for this fold: [0.1 0.2 0.6 0.9]

--- Fold 2 ---

Train indices: [0 1 3 4 5 6 8]

Validation indices: [2 7 9]

Training data for this fold: [0.1 0.2 0.4 0.5 0.6 0.7 0.9]

Validation data for this fold: [0.3 0.8 1.]

--- Fold 3 ---

Train indices: [0 1 2 5 7 8 9]

Validation indices: [3 4 6]

Training data for this fold: [0.1 0.2 0.3 0.6 0.8 0.9 1.]

Validation data for this fold: [0.4 0.5 0.7]

82.2.3. Making Predictions

To make predictions, we can use the coefficients. The data is standardized around its mean in the polynomial function, which is why the vector `mnstd` is required. The coefficient vector is computed based on the normalized data, and this must be taken into account if further analytical calculations are performed on the fitted model.

The polynomial approximation of C_D is:

$$C_D(x) = w_8x^8 + w_7x^7 + \dots + w_1x + w_0,$$

where x is normalized as:

$$\bar{x} = \frac{x - \mu(X)}{\sigma(X)}$$

```
def predict_polynomial_fit(X, coeff, mnstd):
    """
    Generates predictions for the polynomial fit.

    Args:
        X (array-like): Original independent variable data.
        coeff (array): Coefficients of the best-fit polynomial.
        mnstd (tuple): Normalization parameters (mean, std) for X.

    Returns:
        tuple: De-normalized predicted X values and corresponding Y predictions.
    """
    # Normalize X
    X_norm = (X - mnstd[0]) / mnstd[1]

    # Generate predictions
    X_pred = np.linspace(min(X_norm), max(X_norm), 100)
    Y_pred = np.polyval(coeff, X_pred)

    # De-normalize X for plotting
    X_pred_original = X_pred * mnstd[1] + mnstd[0]

    return X_pred_original, Y_pred
```

82.2.4. Plotting the Results

```
def plot_polynomial_fit(X, Y, X_pred_original, Y_pred, best_order, y_true=None):
    """
    Visualizes the polynomial fit.

    Args:
        X (array-like): Original independent variable data.
```

```

        Y (array-like): Original dependent variable data.
        X_pred_original (array): De-normalized predicted X values.
        Y_pred (array): Predicted Y values.
        y_true (array): True Y values.
        best_order (int): The optimal polynomial order.
    """
    plt.scatter(X, Y, label="Training Data", color="grey", marker="o")
    plt.plot(X_pred_original, Y_pred, label=f"Order {best_order} Polynomial", color="blue")
    if y_true is not None:
        plt.plot(X, y_true, label="True Function", color="blue", linestyle="--")
    plt.title(f"Polynomial Fit (Order {best_order})")
    plt.xlabel("X")
    plt.ylabel("Y")
    plt.legend()
    plt.show()

```

82.3. Example One: Aerofoil Drag

The circles in Figure 82.1 represent 101 drag coefficient values obtained through a numerical simulation by iterating each member of a family of aerofoils towards a target lift value (see the Appendix, Section A.3 in Forrester et al. (2008)). The members of the family have different shapes, as determined by the sampling plan:

$$X = x_1, x_2, \dots, x_{101}$$

The responses are:

$$C_D = \{C_D^{(1)}, C_D^{(2)}, \dots, C_D^{(101)}\}$$

These responses are corrupted by “noise,” which are deviations of the systematic variety caused by small changes in the computational mesh from one design to the next.

The original data is measured in natural units, i.e., from -0.3 unit to 0.1 unit. The data is normalized to the range of 0 to 1 for the computation with the `aerofoilcd` function. The data is then fitted with a polynomial of order m . To obtain the best polynomial through this data, the following Python code can be used:

```

from spotoptim.function import aerofoilcd
import numpy as np
import matplotlib.pyplot as plt
X = np.linspace(-0.3, 0.1, 101)
# normalize the data so that it will be in the range of 0 to 1

```

```

a = np.min(X)
b = np.max(X)
X_cod = (X - a) / (b - a)
y = aerofoilcd(X_cod)
best_order, best_coeff, mnstd = polynomial_fit(X, y)
X_pred_original, Y_pred = predict_polynomial_fit(X, best_coeff, mnstd)

plot_polynomial_fit(X, y, X_pred_original, Y_pred, best_order)

```




Figure 82.1.: Aerofoil drag data

Figure 82.1 shows an eighth-order polynomial fitted through the aerofoil drag data. The order was selected via cross-validation, and the coefficients were determined through likelihood maximization. Results, i.e., the best polynomial order and coefficients, are printed in the console. The coefficients are stored in the vector `best_coeff`, which contains the coefficients of the polynomial in descending order. The first element is the coefficient of x^8 , and the last element is the constant term. The vector `mnstd`, containing the mean and standard deviation of X , is:

82. Polynomial Models

```
print(f"Best polynomial order: {best_order}\n")
print(f"Coefficients (starting with w0):\n {best_coeff}\n")
print(f"Normalization parameters (mean, std):\n {mnstd}\n")
```

Best polynomial order: 8

Coefficients (starting with w0):

```
[-0.00022964 -0.00014636  0.00116742  0.00052988 -0.0016912  -0.00047398
  0.00244373  0.00270342  0.03041508]
```

Normalization parameters (mean, std):

```
(np.float64(-0.09999999999999999), np.float64(0.11661903789690602))
```

82.4. Example Two: A Multimodal Test Case

Let us consider the one-variable test function:

$$f(x) = (6x - 2)^2 \sin(12x - 4).$$

```
import numpy as np
from spotoptim.function import onevar
X = np.linspace(0, 1, 51)
y_true = onevar(X)
# initialize random seed
np.random.seed(42)
y = y_true + np.random.normal(0, 1, len(X))*1.1
best_order, best_coeff, mnstd = polynomial_fit(X, y)
X_pred_original, Y_pred = predict_polynomial_fit(X, best_coeff, mnstd)
```

```
plot_polynomial_fit(X, y, X_pred_original, Y_pred, best_order, y_true=y_true)
```




Figure 82.2.: Onevar function

This function, depicted by the dotted line in Figure 82.2, has local minima of different depths, which can be deceptive to some surrogate-based optimization procedures. Here, we use it as an example of a multimodal function for polynomial fitting.

We generate the training data (depicted by circles in Figure 82.2) by adding normally distributed noise to the function. Figure 82.2 shows a seventh-order polynomial fitted through the noisy data. This polynomial was selected as it minimizes the cross-validation metric.

82.5. Extending to Multivariable Polynomial Models

While the examples above focus on the one-variable case, real-world engineering problems typically involve multiple input variables. For k -variable problems, polynomial approximation becomes significantly more complex but follows the same fundamental principles. For a k -dimensional input space, the polynomial approximation can be expressed as:

82. Polynomial Models

$$\hat{f}(\vec{x}) = \sum_{i=1}^N w_i \phi_i(\vec{x}),$$

where $\phi_i(\vec{x})$ represents multivariate basis functions, and N is the total number of terms in the polynomial. Unlike the univariate case, these basis functions include all possible combinations of variables up to the selected polynomial order m , which might result in a “basis function explosion” as the number of variables increases.

For a third-order polynomial ($m = 3$) with three variables ($k = 3$), the complete set of basis functions would include 20 terms:

$$\text{Constant term: } 1 \quad (82.1)$$

$$\text{First-order terms: } x_1, x_2, x_3 \quad (82.2)$$

$$\text{Second-order terms: } x_1^2, x_2^2, x_3^2, x_1x_2, x_1x_3, x_2x_3 \quad (82.3)$$

$$\text{Third-order terms: } x_1^3, x_2^3, x_3^3, x_1^2x_2, x_1^2x_3, x_2^2x_1, x_2^2x_3, x_3^2x_1, x_3^2x_2, x_1x_2x_3 \quad (82.4)$$

The total number of terms grows combinatorially as $N = \binom{k+m}{m}$, which quickly becomes prohibitive as dimensionality increases. For example, a 10-variable cubic polynomial requires $\binom{13}{3} = 286$ coefficients! This exponential growth creates three interrelated challenges:

- **Model Selection:** Determining the appropriate polynomial order m that balances complexity with generalization ability
- **Coefficient Estimation:** Computing the potentially large number of weights $\vec{w}$ while avoiding numerical instability
- **Term Selection:** Identifying which specific basis functions should be included, as many may be irrelevant to the response

Several techniques have been developed to address these challenges:

- Regularization methods (LASSO, ridge regression) that penalize model complexity
- Stepwise regression algorithms that incrementally add or remove terms
- Dimension reduction techniques that project the input space to lower dimensions
- Orthogonal polynomials that improve numerical stability for higher-order models

These limitations of polynomial models in higher dimensions motivate the exploration of more flexible surrogate modeling approaches like Radial Basis Functions and Kriging, which we'll examine in subsequent sections.

82.6. Jupyter Notebook

i Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

83. Radial Basis Function Models

Note

- This section is based on chapter 2.3 in Forrester et al. (2008).
- The following Python packages are imported:

```
import numpy as np
import matplotlib.pyplot as plt
```

83.1. Radial Basis Function Models

Scientists and engineers frequently tackle complex functions by decomposing them into a “vocabulary” of simpler, well-understood basic functions. These fundamental building blocks possess properties that make them easier to analyze mathematically and implement computationally. We explored this concept earlier with multivariable polynomials, where complex behaviors were modeled using combinations of polynomial terms such as 1 , x_1 , x_2 , x_1^2 , and x_1x_2 . This approach is not limited to polynomials; it extends to various function classes, including trigonometric functions, exponential functions, and even more complex structures.

While Fourier analysis—perhaps the most widely recognized example of this approach—excels at representing periodic phenomena through sine and cosine functions, the focus in Forrester et al. (2008) is broader. They aim to approximate arbitrary smooth, continuous functions using strategically positioned basis functions. Specifically, radial basis function (RBF) models employ symmetrical basis functions centered at selected points distributed throughout the design space. These basis functions have the unique property that their output depends only on the distance from their center point.

First, we give a definition of the Euclidean distance, which is the most common distance measure used in RBF models.

Definition 83.1 (Euclidean Distance). The Euclidean distance between two points in a k -dimensional space is defined as:

83. Radial Basis Function Models

$$\|\vec{x} - \vec{c}\| = \sqrt{\sum_{i=1}^k (x_i - c_i)^2}, \quad (83.1)$$

where:

- $\vec{x} = (x_1, x_2, \dots, x_d)$ is the first point,
- $\vec{c} = (c_1, c_2, \dots, c_d)$ is the second point, and
- k is the number of dimensions.

The Euclidean distance measure represents the straight-line distance between two points in Euclidean space.

Using the Euclidean distance, we can define the radial basis function (RBF) model.

Definition 83.2 (Radial Basis Function (RBF)). Mathematically, a radial basis function ψ can be expressed as:

$$\psi(\vec{x}) = \psi(\|\vec{x} - \vec{c}\|), \quad (83.2)$$

where $\vec{x}$ is the input vector, $\vec{c}$ is the center of the function, and $\|\vec{x} - \vec{c}\|$ denotes the Euclidean distance between $\vec{x}$ and $\vec{c}$.

In the context of RBFs, the Euclidean distance calculation determines how much influence a particular center point $\vec{c}$ has on the prediction at point $\vec{x}$.

We will first examine interpolating RBF models, which assume noise-free data and pass exactly through all training points. This approach provides an elegant mathematical foundation before we consider more practical scenarios where data contains measurement or process noise.

83.1.1. Fitting Noise-Free Data

Let us consider the scalar valued function f observed without error, according to the sampling plan $X = \{\vec{x}^{(1)}, \vec{x}^{(2)}, \dots, \vec{x}^{(n)}\}^T$, yielding the responses $\vec{y} = \{y^{(1)}, y^{(2)}, \dots, y^{(n)}\}^T$.

For a given set of n_c centers $\vec{c}^{(i)}$, we would like to express the RBF model as a linear combination of the basis functions centered at these points. The goal is to find the weights $\vec{w}$ that minimize the error between the predicted and observed values. Thus, we seek a radial basis function approximation to f of the fixed form:

$$\hat{f}(\vec{x}) = \sum_{i=1}^{n_c} w_i \psi(\|\vec{x} - \vec{c}^{(i)}\|), \quad (83.3)$$

where

- w_i are the weights of the n_c basis functions,
- $\vec{c}^{(i)}$ are the n_c centres of the basis functions, and
- ψ is a radial basis function.

The notation $\|\cdot\|$ denotes the Euclidean distance between two points in the design space as defined in Equation 83.1.

83.1.1.1. Selecting Basis Functions: From Fixed to Parametric Forms

When implementing a radial basis function model, we initially have one undetermined parameter per basis function: the weight applied to each function's output. This simple parameterization remains true when we select from several standard fixed-form basis functions, such as:

- Linear ($\psi(r) = r$): The simplest form, providing a response proportional to distance
- Cubic ($\psi(r) = r^3$): Offers stronger emphasis on points farther from the center
- Thin plate spline ($\psi(r) = r^2 \ln r$): Models the physical bending of a thin sheet, providing excellent smoothness properties

While these fixed basis functions are computationally efficient, they offer limited flexibility in how they generalize across the design space. For more adaptive modeling power, we can employ parametric basis functions that introduce additional tunable parameters:

- Gaussian ($\psi(r) = e^{-r^2/(2\sigma^2)}$): Produces bell-shaped curves with σ controlling the width of influence
- Multiquadric ($\psi(r) = (r^2 + \sigma^2)^{1/2}$): Provides broader coverage with less localized effects
- Inverse multiquadric ($\psi(r) = (r^2 + \sigma^2)^{-1/2}$): Offers sharp peaks near centers with asymptotic behavior

The parameter σ in these functions serves as a shape parameter that controls how rapidly the function's influence decays with distance. This added flexibility enables significantly better generalization, particularly when modeling complex responses, though at the cost of a more involved parameter estimation process requiring optimization of both weights and shape parameters.

Example 83.1 (Gaussian RBF). Using the general definition of a radial basis function (Equation 83.2), we can express the Gaussian RBF as:

$$\psi(\vec{x}) = \exp\left(-\frac{\|\vec{x} - \vec{c}\|^2}{2\sigma^2}\right) = \exp\left(-\frac{\sum_{j=1}^k (x_j - c_j)^2}{2\sigma^2}\right) \quad (83.4)$$

where:

83. Radial Basis Function Models

- $\vec{x}$ is the input vector,
- $\vec{c}$ is the center vector,
- $\|\vec{x} - \vec{c}\|$ is the Euclidean distance between the input and center, and
- σ is the width parameter that controls how quickly the function's response diminishes with distance from the center.

The Gaussian RBF produces a bell-shaped response that reaches its maximum value of 1 when $\vec{x} = \vec{c}$ and asymptotically approaches zero as the distance increases. The parameter σ determines how “localized” the response is—smaller values create a narrower peak with faster decay, while larger values produce a broader, more gradual response across the input space. Figure 83.1 shows the Gaussian RBF for different values of σ in an one-dimensional space. The center of the RBF is set at 0, and the width parameter σ varies to illustrate how it affects the shape of the function.

```
def gaussian_rbf(x, center, sigma):  
    """  
    Compute the Gaussian Radial Basis Function.  
  
    Args:  
        x (ndarray): Input points  
        center (float): Center of the RBF  
        sigma (float): Width parameter  
  
    Returns:  
        ndarray: RBF values  
    """  
    return np.exp(-((x - center)**2) / (2 * sigma**2))
```

The sum of Gaussian RBFs can be visualized by summing the individual Gaussian RBFs centered at different points as shown in Figure 83.2. The following code snippet demonstrates how to create this plot showing the sum of three Gaussian RBFs with different centers and a common width parameter σ .

83.1.1.2. The Interpolation Condition: Elegant Solutions Through Linear Systems

A remarkable property of radial basis function models is that regardless of which basis functions we choose—parametric or fixed—determining the weights $\vec{w}$ remains straightforward through interpolation. The core principle is elegantly simple: we require our model to exactly reproduce the observed data points:

$$\hat{f}(\vec{x}^{(i)}) = y^{(i)}, \quad i = 1, 2, \dots, n. \quad (83.5)$$

This constraint produces one of the most powerful aspects of RBF modeling: while the system in Equation 83.5 is linear with respect to the weights $\vec{w}$, the resulting predictor




Figure 83.1.: Gaussian RBF

$\hat{f}$ can capture highly nonlinear relationships in the data. The RBF approach transforms a complex nonlinear modeling problem into a solvable linear algebra problem.

For a unique solution to exist, we require that the number of basis functions equals the number of data points ($n_c = n$). The standard practice, which greatly simplifies implementation, is to center each basis function at a training data point, setting $\vec{c}^{(i)} = \vec{x}^{(i)}$ for all $i = 1, 2, \dots, n$. This choice allows us to express the interpolation condition as a compact matrix equation:

$$\Psi \vec{w} = \vec{y}.$$

Here, Ψ represents the Gram matrix (also called the design matrix or kernel matrix), whose elements measure the similarity between data points:

$$\Psi_{i,j} = \psi(\|\vec{x}^{(i)} - \vec{x}^{(j)}\|), \quad i, j = 1, 2, \dots, n.$$

The solution for the weight vector becomes:

$$\vec{w} = \Psi^{-1} \vec{y}.$$




Figure 83.2.: Sum of Gaussian RBFs

This matrix inversion step is the computational core of the RBF model fitting process, and the numerical properties of this operation depend critically on the chosen basis function. Different basis functions produce Gram matrices with distinct conditioning properties, directly affecting both computational stability and the model's generalization capabilities.

83.1.2. Numerical Stability Through Positive Definite Matrices

A significant advantage of Gaussian and inverse multiquadric basis functions lies in their mathematical guarantees. Vapnik (1998) demonstrated that these functions always produce symmetric positive definite Gram matrices when using strictly positive definite kernels (see Section 85.4), which is a critical property for numerical reliability. Unlike other basis functions that may lead to ill-conditioned systems, these functions ensure the existence of unique, stable solutions.

This positive definiteness enables the use of Cholesky factorization, which offers substantial computational advantages over standard matrix inversion techniques. The Cholesky approach reduces the computational cost (reducing from $O(n^3)$ to roughly $O(n^3/3)$) while significantly improving numerical stability when handling the inevitable rounding errors in floating-point arithmetic. This robustness to numerical issues explains why

Gaussian and inverse multiquadric basis functions remain the preferred choice in many practical RBF implementations.

Furthermore, the positive definiteness guarantee provides theoretical assurances about the model's interpolation properties—ensuring that the RBF interpolant exists and is unique for any distinct set of centers. This mathematical foundation gives practitioners confidence in the method's reliability, particularly for complex engineering applications where model stability is paramount.

The computational advantage stems from how a symmetric positive definite matrix Ψ can be efficiently decomposed into the product of an upper triangular matrix U and its transpose:

$$\Psi = U^T U.$$

This decomposition transforms the system

$$\Psi \vec{w} = \vec{y}$$

into

$$U^T U \vec{w} = \vec{y},$$

which can be solved through two simpler triangular systems:

- First solve $U^T \vec{v} = \vec{y}$ for the intermediate vector $\vec{v}$
- Then solve $U \vec{w} = \vec{v}$ for the desired weights $\vec{w}$

In Python implementations, this process is elegantly handled using NumPy's or SciPy's Cholesky decomposition functions, followed by specialized solvers that exploit the triangular structure:

```
from scipy.linalg import cholesky, cho_solve
# Compute the Cholesky factorization
L = cholesky(Psi, lower=True) # L is the lower triangular factor
weights = cho_solve((L, True), y) # Efficient solver for (L L^T)w = y
```

83.1.3. Ill-Conditioning

An important numerical consideration in RBF modeling is that points positioned extremely close to each other in the input space X can lead to severe ill-conditioning of the Gram matrix (Micchelli 1986). This ill-conditioning manifests as nearly linearly dependent rows and columns in Ψ , potentially causing the Cholesky factorization to fail.

While this problem rarely arises with initial space-filling experimental designs (such as Latin Hypercube or quasi-random sequences), it frequently emerges during sequential

83. Radial Basis Function Models

optimization processes that adaptively add infill points in promising regions. As these clusters of points concentrate in areas of high interest, the condition number of the Gram matrix deteriorates, jeopardizing numerical stability.

Several mitigation strategies exist: regularization through ridge-like penalties (modifying the standard RBF interpolation problem by adding a penalty term to the diagonal of the Gram matrix. This creates a literal “ridge” along the diagonal of the matrix), removing nearly coincident points, clustering, or applying more sophisticated approaches. One theoretically elegant solution involves augmenting non-conditionally positive definite basis functions with polynomial terms (Keane and Nair 2005). This technique not only improves conditioning but also ensures polynomial reproduction properties, enhancing the approximation quality for certain function classes while maintaining numerical stability.

Beyond determining $\vec{w}$, there is, of course, the additional task of estimating any other parameters introduced via the basis functions. A typical example is the σ of the Gaussian basis function, usually taken to be the same for all basis functions, though a different one can be selected for each centre, as is customary in the case of the Kriging basis function, to be discussed shortly (once again, we trade additional parameter estimation complexity versus increased flexibility and, hopefully, better generalization).

83.1.4. Parameter Optimization: A Two-Level Approach

When building RBF models, we face two distinct parameter estimation challenges:

- Determining the weights ($\vec{w}$): These parameters ensure our model precisely reproduces the training data. For any fixed basis function configuration, we can calculate these weights directly through linear algebra as shown earlier.
- Optimizing shape parameters (like σ in Gaussian RBF): These parameters control how the model generalizes to new, unseen data. Unlike weights, there’s no direct formula to find their optimal values.

To address this dual challenge, we employ a nested optimization strategy (inner and outer levels):

83.1.4.1. Inner Level ($\vec{w}$)

For each candidate value of shape parameters (e.g., σ), we determine the corresponding optimal weights $\vec{w}$ by solving the linear system. The `estim_weights()` method implements the inner level optimization by calculating the optimal weights $\vec{w}$ for a given shape parameter (σ):

```

def estim_weights(self):
    # [...]

    # Construct the Phi (Psi) matrix
    self.Phi = np.zeros((n, n))
    for i in range(n):
        for j in range(i+1):
            self.Phi[i, j] = self.basis(d[i, j], self.sigma)
            self.Phi[j, i] = self.Phi[i, j]

    # Calculate weights using appropriate method
    if self.code == 4 or self.code == 6:
        # Use Cholesky factorization for Gaussian or inverse multiquadric
        try:
            L = cholesky(self.Phi, lower=True)
            self.weights = cho_solve((L, True), self.y)
            self.success = True
        except np.linalg.LinAlgError:
            # Error handling...
    else:
        # Use direct solve for other basis functions
        try:
            self.weights = np.linalg.solve(self.Phi, self.y)
            self.success = True
        except np.linalg.LinAlgError:
            # Error handling...

    return self

```

This method:

- Creates the Gram matrix (Phi) based on distances between points
- Solves the linear system $\Psi \vec{w} = \vec{y}$ for weights
- Uses appropriate numerical methods based on the basis function type (Cholesky factorization or direct solve)

83.1.4.2. Outer Level (σ)

We use cross-validation to evaluate how well the model generalizes with different shape parameter values. The outer level optimization is implemented within the `fit()` method, where cross-validation is used to evaluate different σ values:

```

def fit(self):
    if self.code < 4:
        # Fixed basis function, only w needs estimating
        self.estim_weights()
    else:
        # Basis function requires a sigma, estimate first using cross-validation
        # [...]

        # Generate candidate sigma values
        sigmas = np.logspace(-2, 2, 30)

        # Setup cross-validation (determine number of folds)
        # [...]

        cross_val = np.zeros(len(sigmas))

        # For each candidate sigma value
        for sig_index, sigma in enumerate(sigmas):
            print(f"Computing cross-validation metric for Sigma={sigma:.4f}...")

            # Perform k-fold cross-validation
            for j in range(len(from_idx)):
                # Create and fit model on training subset
                temp_model = Rbf(
                    X=X_orig[xs_temp],
                    y=y_orig[ys_temp],
                    code=self.code
                )
                temp_model.sigma = sigma

                # Call inner level optimization
                temp_model.estim_weights()

                # Evaluate on held-out data
                # [...]

            # Select best sigma based on cross-validation performance
            min_cv_index = np.argmin(cross_val)
            best_sig = sigmas[min_cv_index]

        # Use the best sigma for final model
        self.sigma = best_sig
        self.estim_weights() # Call inner level again with optimal sigma

```

The outer level:

- Generates a range of candidate σ values
- For each σ , performs k-fold cross-validation:
 - Creates models on subsets of the data
 - Calls the inner level method (`estim_weights()`) to determine weights
 - Evaluates prediction quality on held-out data
- Selects the σ that minimizes cross-validation error
- Performs a final call to the inner level method with the optimal σ

This two-level approach is particularly critical for parametric basis functions (Gaussian, multiquadric, etc.), where the wrong choice of shape parameter could lead to either overfitting (too much flexibility) or underfitting (too rigid). Cross-validation provides an unbiased estimate of how well different parameter choices will perform on new data, helping us balance the trade-off between fitting the training data perfectly and generalizing well.

83.2. Python Implementation of the RBF Model

Section 83.2 shows a Python implementation of this parameter estimation process (based on a cross-validation routine), which will represent the surrogate, once its parameters have been estimated. The model building process is very simple.

Instead of using a dictionary for bookkeeping, we implement a Python class `Rbf` that encapsulates all the necessary data and functionality. The class stores the sampling plan X as the `X` attribute and the corresponding n -vector of responses y as the `y` attribute. The `code` attribute specifies the type of basis function to be used. After fitting the model, the class will also contain the estimated parameter values $\vec{w}$ and, if a parametric basis function is used, σ . These are stored in the `weights` and `sigma` attributes respectively.

Finally, a note on prediction error estimation. We have already indicated that the guarantee of a positive definite Ψ is one of the advantages of Gaussian radial basis functions. They also possess another desirable feature: it is relatively easy to estimate their prediction error at any $\vec{x}$ in the design space. Additionally, the expectation function of the improvement in minimum (or maximum) function value with respect to the minimum (or maximum) known so far can also be calculated quite easily, both of these features being very useful when the optimization of f is the goal of the surrogate modelling process.

83.2.1. The Rbf Class

The Rbf class implements the Radial Basis Function model. It encapsulates all the data and methods needed for fitting the model and making predictions.

```
import numpy as np
from scipy.linalg import cholesky, cho_solve
import numpy.random as rnd

class Rbf:
    """Radial Basis Function model implementation.

    Attributes:
        X (ndarray): The sampling plan (input points).
        y (ndarray): The response vector.
        code (int): Type of basis function to use.
        weights (ndarray, optional): The weights vector (set after fitting).
        sigma (float, optional): Parameter for parametric basis functions.
        Phi (ndarray, optional): The Gram matrix (set during fitting).
        success (bool, optional): Flag indicating successful fitting.
    """

    def __init__(self, X=None, y=None, code=3):
        """Initialize the RBF model.

        Args:
            X (ndarray, optional):
                The sampling plan.
            y (ndarray, optional):
                The response vector.
            code (int, optional):
                Type of basis function.
                Default is 3 (thin plate spline).
        """
        self.X = X
        self.y = y
        self.code = code
        self.weights = None
        self.sigma = None
        self.Phi = None
        self.success = None

    def basis(self, r, sigma=None):
        """Compute the value of the basis function.
```

```

Args:
    r (float): Radius (distance)
    sigma (float, optional): Parameter for parametric basis functions

Returns:
    float: Value of the basis function
"""
# Use instance sigma if not provided
if sigma is None and hasattr(self, 'sigma'):
    sigma = self.sigma

if self.code == 1:
    # Linear function
    return r
elif self.code == 2:
    # Cubic
    return r**3
elif self.code == 3:
    # Thin plate spline
    if r < 1e-200:
        return 0
    else:
        return r**2 * np.log(r)
elif self.code == 4:
    # Gaussian
    return np.exp(-(r**2)/(2*sigma**2))
elif self.code == 5:
    # Multi-quadric
    return (r**2 + sigma**2)**0.5
elif self.code == 6:
    # Inverse Multi-Quadric
    return (r**2 + sigma**2)**(-0.5)
else:
    raise ValueError("Invalid basis function code")

def estim_weights(self):
    """Estimates the basis function weights if sigma is known or not required.

Returns:
    self: The updated model instance
"""
    # Check if sigma is required but not provided
    if self.code > 3 and self.sigma is None:
        raise ValueError("The basis function requires a sigma parameter")

```

83. Radial Basis Function Models

```
# Number of points
n = len(self.y)

# Build distance matrix
d = np.zeros((n, n))
for i in range(n):
    for j in range(i+1):
        d[i, j] = np.linalg.norm(self.X[i] - self.X[j])
        d[j, i] = d[i, j]

# Construct the Phi (Psi) matrix
self.Phi = np.zeros((n, n))
for i in range(n):
    for j in range(i+1):
        self.Phi[i, j] = self.basis(d[i, j], self.sigma)
        self.Phi[j, i] = self.Phi[i, j]

# Calculate weights using appropriate method
if self.code == 4 or self.code == 6:
    # Use Cholesky factorization for Gaussian or inverse multiquadric
    try:
        L = cholesky(self.Phi, lower=True)
        self.weights = cho_solve((L, True), self.y)
        self.success = True
    except np.linalg.LinAlgError:
        print("Cholesky factorization failed.")
        print("Two points may be too close together.")
        self.weights = None
        self.success = False
else:
    # Use direct solve for other basis functions
    try:
        self.weights = np.linalg.solve(self.Phi, self.y)
        self.success = True
    except np.linalg.LinAlgError:
        self.weights = None
        self.success = False

return self

def fit(self):
    """Estimates the parameters of the Radial Basis Function model.

    Returns:
```



```

        self: The updated model instance
    """
    if self.code < 4:
        # Fixed basis function, only w needs estimating
        self.estim_weights()
    else:
        # Basis function also requires a sigma, estimate first
        # Save original model data
        X_orig = self.X.copy()
        y_orig = self.y.copy()

        # Direct search between 10^-2 and 10^2
        sigmas = np.logspace(-2, 2, 30)

        # Number of cross-validation subsets
        if len(self.X) < 6:
            q = 2
        elif len(self.X) < 15:
            q = 3
        elif len(self.X) < 50:
            q = 5
        else:
            q = 10

        # Number of sample points
        n = len(self.X)

        # X split into q randomly selected subsets
        xs = rnd.permutation(n)
        full_xs = xs.copy()

        # The beginnings of the subsets...
        from_idx = np.arange(0, n, n//q)
        if from_idx[-1] >= n:
            from_idx = from_idx[:-1]

        # ...and their ends
        to_idx = np.zeros_like(from_idx)
        for i in range(len(from_idx) - 1):
            to_idx[i] = from_idx[i+1] - 1
        to_idx[-1] = n - 1

        cross_val = np.zeros(len(sigmas))

```

83. Radial Basis Function Models

```
# Cycling through the possible values of Sigma
for sig_index, sigma in enumerate(sigmas):
    print(f"Computing cross-validation metric for Sigma={sigma:.4f}...")

    cross_val[sig_index] = 0

    # Model fitting to subsets of the data
    for j in range(len(from_idx)):
        removed = xs[from_idx[j]:to_idx[j]+1]
        xs_temp = np.delete(xs, np.arange(from_idx[j], to_idx[j]+1))

        # Create a temporary model for CV
        temp_model = Rbf(
            X=X_orig[xs_temp],
            y=y_orig[xs_temp],
            code=self.code
        )
        temp_model.sigma = sigma

        # Sigma and subset chosen, now estimate w
        temp_model.estim_weights()

        if temp_model.weights is None:
            cross_val[sig_index] = 1e20
            xs = full_xs.copy()
            break

        # Compute vector of predictions at the removed sites
        pr = np.zeros(len(removed))
        for jj, idx in enumerate(removed):
            pr[jj] = temp_model.predict(X_orig[idx])

        # Calculate cross-validation error
        cross_val[sig_index] += np.sum((y_orig[removed] - pr)**2) / len(r

    xs = full_xs.copy()

    # Now attempt Cholesky on the full set, in case the subsets could
    # be fitted correctly, but the complete X could not
    temp_model = Rbf(
        X=X_orig,
        y=y_orig,
        code=self.code
    )
```

```

        temp_model.sigma = sigma
        temp_model.estim_weights()

        if temp_model.weights is None:
            cross_val[sig_index] = 1e20
            print("Failed to fit complete sample data.")

        # Find the best sigma
        min_cv_index = np.argmin(cross_val)
        best_sig = sigmas[min_cv_index]

        # Set the best sigma and recompute weights
        print(f"Selected sigma={best_sig:.4f}")
        self.sigma = best_sig
        self.estim_weights()

    return self

def predict(self, x):
    """Calculates the value of the Radial Basis Function surrogate model at x.

    Args:
        x (ndarray): Point at which to make prediction

    Returns:
        float: Predicted value
    """
    # Calculate distances to all sample points
    d = np.zeros(len(self.X))
    for k in range(len(self.X)):
        d[k] = np.linalg.norm(x - self.X[k])

    # Calculate basis function values
    phi = np.zeros(len(self.X))
    for k in range(len(self.X)):
        phi[k] = self.basis(d[k], self.sigma)

    # Calculate prediction
    y = np.dot(phi, self.weights)
    return y

```

83.3. RBF Example: The One-Dimensional sin Function

```
import numpy as np
import matplotlib.pyplot as plt
from scipy.linalg import cholesky, cho_solve

# Define the data points for fitting
x_centers = np.array([np.pi/2, np.pi, 3*np.pi/2]).reshape(-1, 1) # Centers for RBFs
y_values = np.sin(x_centers.flatten()) # Sine values at these points

# Create and fit the RBF model
rbf_model = Rbf(X=x_centers, y=y_values, code=4) # Code 4 is Gaussian RBF
rbf_model.sigma = 1.0 # Set sigma parameter directly
rbf_model.estimate_weights() # Calculate optimal weights

# Print the weights
print("RBF model weights:", rbf_model.weights)

# Create a grid for visualization
x_grid = np.linspace(0, 2*np.pi, 1000).reshape(-1, 1)
y_true = np.sin(x_grid.flatten()) # True sine function

# Generate predictions using the RBF model
y_pred = np.zeros(len(x_grid))
for i in range(len(x_grid)):
    y_pred[i] = rbf_model.predict(x_grid[i])

# Calculate individual basis functions for visualization
basis_funcs = np.zeros((len(x_grid), len(x_centers)))
for i in range(len(x_grid)):
    for j in range(len(x_centers)):
        # Calculate distance
        distance = np.linalg.norm(x_grid[i] - x_centers[j])
        # Compute basis function value scaled by its weight
        basis_funcs[i, j] = rbf_model.basis(distance, rbf_model.sigma) * rbf_model.weights[j]

# Plot the results
plt.figure(figsize=(6, 4))

# Plot the true sine function
plt.plot(x_grid, y_true, 'k-', label='True sine function', linewidth=2)

# Plot individual basis functions
```

83.3. RBF Example: The One-Dimensional $\sin$ Function

```
for i in range(len(x_centers)):
    plt.plot(x_grid, basis_funcs[:, i], '--',
             label=f'Basis function at x={x_centers[i][0]:.2f}')


# Plot the RBF fit (sum of basis functions)
plt.plot(x_grid, y_pred, 'r-', label='RBF fit', linewidth=2)

# Plot the sample points
plt.scatter(x_centers, y_values, color='blue', s=100, label='Sample points')

# Add horizontal line at y=0
plt.axhline(y=0, color='gray', linestyle='--', alpha=0.3)

plt.title('RBF Approximation of Sine Function')
plt.xlabel('x')
plt.ylabel('y')
plt.grid(True)
plt.legend()
plt.tight_layout()
plt.show()
```

RBF model weights: [1.00724398e+00 2.32104414e-16 -1.00724398e+00]



83.4. RBF Example: The Two-Dimensional dome Function

The `dome` function is an example of a test function that can be used to evaluate the performance of the Radial Basis Function model. It is a simple mathematical function defined over a two-dimensional space.

```
def dome(x) -> float:
    """
    Dome test function.

    Args:
        x (ndarray): Input vector (1D array of length 2)

    Returns:
        float: Function value

    Examples:
        dome(np.array([0.5, 0.5]))
    """
    return np.sum(1 - (2*x - 1)**2) / len(x)
```

The following code demonstrates how to use the Radial Basis Function model to approximate a function. It generates a Latin Hypercube sample, computes the objective function values, estimates the model parameters, and plots the results.

```
def generate_rbf_data(n_samples=10, grid_points=41):
    """
    Generates data for RBF visualization.

    Args:
        n_samples (int): Number of samples for the RBF model
        grid_points (int): Number of grid points for prediction

    Returns:
        tuple: (rbf_model, X, Y, Z, Z_0) - Model and grid data for plotting
    """
    from spotoptim.sampling.mm import bestlh as best_lh
    # Generate sampling plan
    X_samples = best_lh(n_samples, 2, population=10, iterations=100)
    # Compute objective function values
    y_samples = np.zeros(len(X_samples))
    for i in range(len(X_samples)):
```

83.4. RBF Example: The Two-Dimensional dome Function

```
    y_samples[i] = dome(X_samples[i])
# Create and fit RBF model
rbf_model = Rbf(X=X_samples, y=y_samples, code=3) # Thin plate spline
rbf_model.fit()
# Generate grid for prediction
x = np.linspace(0, 1, grid_points)
y = np.linspace(0, 1, grid_points)
X, Y = np.meshgrid(x, y)
Z_0 = np.zeros_like(X)
Z = np.zeros_like(X)

# Evaluate model at grid points
for i in range(len(x)):
    for j in range(len(y)):
        Z_0[j, i] = dome(np.array([x[i], y[j]]))
        Z[j, i] = rbf_model.predict(np.array([x[i], y[j]]))

return rbf_model, X, Y, Z, Z_0

def plot_rbf_results(rbf_model, X, Y, Z, Z_0=None, n_contours=10):
    """
    Plots RBF approximation results.

    Args:
        rbf_model (Rbf): Fitted RBF model
        X (ndarray): Grid X-coordinates
        Y (ndarray): Grid Y-coordinates
        Z (ndarray): RBF model predictions
        Z_0 (ndarray, optional): True function values for comparison
        n_contours (int): Number of contour levels to plot
    """
    import matplotlib.pyplot as plt

    plt.figure(figsize=(10, 8))

    # Plot the contour
    contour = plt.contour(X, Y, Z, n_contours)

    if Z_0 is not None:
        contour_0 = plt.contour(X, Y, Z_0, n_contours, colors='k', linestyle='dashed')

    # Plot the sample points
    plt.scatter(rbf_model.X[:, 0], rbf_model.X[:, 1],
               c='r', marker='o', s=50)
```

```
plt.title('RBF Approximation (Thin Plate Spline)')
plt.xlabel('x1')
plt.ylabel('x2')
plt.colorbar(label='f(x1, x2)')
plt.show()
```

Figure 83.3 shows the contour plots of the underlying function $f(x_1, x_2) = 0.5[-(2x_1 - 1)^2 - (2x_2 - 1)^2]$ and its thin plate spline radial basis function approximation, along with the 10 points of a Morris-Mitchell optimal Latin hypercube sampling plan (obtained via `best_lh()`).

```
rbf_model, X, Y, Z, Z_0 = generate_rbf_data(n_samples=10, grid_points=41)
plot_rbf_results(rbf_model, X, Y, Z, Z_0)
```




Figure 83.3.: RBF Approximation.

83.4.1. The Connection Between RBF Models and Neural Networks

Radial basis function models share a profound architectural similarity with artificial neural networks, specifically with what's known as RBF networks. This connection provides valuable intuition about how RBF models function. A radial basis function model can be viewed as a specialized neural network with the following structure:

- Input Layer: Receives the feature vector $\vec{x}$
- Hidden Layer: Contains neurons (basis functions) that compute radial distances
- Output Layer: Produces a weighted sum of the hidden unit activations

Unlike traditional neural networks that use dot products followed by nonlinear activation functions, RBF networks measure the distance between inputs and learned center points. This distance is then transformed by the radial basis function.

Mathematically, the equivalence between RBF models and RBF networks can be expressed as follows:

The RBF model equation:

$$\hat{f}(\vec{x}) = \sum_{i=1}^{n_c} w_i \psi(\|\vec{x} - \vec{c}^{(i)}\|)$$

directly maps to the following neural network components:

- $\vec{x}$: Input vector
- $\vec{c}^{(i)}$: Center vectors for each hidden neuron
- $\psi(\cdot)$: Activation function (Gaussian, inverse multiquadric, etc.)
- w_i : Output weights
- $\hat{f}(\vec{x})$: Network output

Example 83.2 (Comparison of RBF Networks and Traditional Neural Networks). Consider approximating a simple 1D function $f(x) = \sin(2\pi x)$ over the interval $[0, 1]$:

The neural network approach would use multiple layers with neurons computing $\sigma(w \cdot x + b)$. It would require a large number of neurons and layers to capture the sine wave's complexity. The network would learn both weights and biases, making it less interpretable.

The RBF network approach, on the other hand, places basis functions at strategic points (e.g., 5 evenly spaced centers). Each neuron computes $\psi(\|x - c_i\|)$ (e.g., using Gaussian RBF). The output layer combines these values with learned weights. If we place Gaussian RBFs with $\sigma = 0.15$ at 0.1, 0.3, 0.5, 0.7, 0.9, each neuron responds strongly when the input is close to its center and weakly otherwise. The network can then learn weights that, when multiplied by these response patterns and summed, closely approximate the sine function.

83. Radial Basis Function Models

This locality property gives RBF networks a notable advantage: they offer more interpretable internal representations and often require fewer neurons for certain types of function approximation compared to traditional multilayer perceptrons.

The key insight is that while standard neural networks create complex decision boundaries through compositions of hyperplanes, RBF networks directly model functions using a set of overlapping “bumps” positioned strategically in the input space.

83.5. Radial Basis Function Models for Noisy Data

When the responses $\vec{y} = y^{(1)}, y^{(2)}, \dots, y^{(n)T}$ contain measurement or simulation noise, the standard RBF interpolation approach can lead to overfitting—the model captures both the underlying function and the random noise. This compromises generalization performance on new data points. Two principal strategies address this challenge:

83.5.1. Ridge Regularization Approach

The most straightforward solution involves introducing regularization through the parameter λ (Poggio and Girosi 1990). This is implemented by adding λ to the diagonal elements of the Gram matrix, creating a “ridge” that improves numerical stability. Mathematically, the weights are determined by:

$$\vec{w} = (\Psi + \lambda I)^{-1} \vec{y},$$

where I is an $n \times n$ identity matrix. This regularized solution balances two competing objectives:

- fitting the training data accurately versus
- keeping the magnitude of weights controlled to prevent overfitting.

Theoretically, optimal performance is achieved when λ equals the variance of the noise in the response data $\vec{y}$ (Keane and Nair 2005). Since this information is rarely available in practice, λ is typically estimated through cross-validation alongside other model parameters.

83.5.2. Reduced Basis Approach

An alternative strategy involves reducing m , the number of basis functions. This might result in a non-square Ψ matrix. With a non-square Ψ matrix, the weights are found through least squares minimization:

$$\vec{w} = (\Psi^T \Psi)^{-1} \Psi^T \vec{y}$$

This approach introduces an important design decision: which subset of points should serve as basis function centers? Several selection strategies exist:

- Clustering methods that identify representative points
- Greedy algorithms that sequentially select influential centers
- Support vector regression techniques (discussed elsewhere in the literature)

Additional parameters such as the width parameter σ in Gaussian bases can be optimized through cross-validation to minimize generalization error estimates.

The ridge regularization and reduced basis approaches can be combined, allowing for a flexible modeling framework, though at the cost of a more complex parameter estimation process. This hybrid approach often yields superior results for highly noisy datasets or when the underlying function has varying complexity across the input space.

The broader challenge of building accurate models from noisy observations is examined comprehensively in the context of Kriging models, which provide a statistical framework for explicitly modeling both the underlying function and the noise process.

83.6. Jupyter Notebook

i Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

84. Kriging (Gaussian Process Regression)

i Note

- This section is based on chapter 2.4 in Forrester et al. (2008).
- The following Python packages are imported:

```
import matplotlib.pyplot as plt
import numpy as np
from numpy import (array, zeros, power, ones, exp, multiply,
                  eye, linspace, spacing, sqrt, arange,
                  append, ravel)
from numpy.linalg import cholesky, solve
from scipy.spatial.distance import squareform, pdist, cdist
```

84.1. From Gaussian RBF to Kriging Basis Functions

Kriging can be explained using the concept of radial basis functions (RBFs), which were introduced in Chapter 83. An RBF is a real-valued function whose value depends only on the distance from a certain point, called the center, usually in a multidimensional space. The basis function is a function of the distance between the point $\vec{x}$ and the center $\vec{x}^{(i)}$. Other names for basis functions are *kernel* or *covariance* functions.

Definition 84.1 (The Kriging Basis Functions). Kriging (also known as Gaussian Process Regression) uses k -dimensional basis functions of the form

$$\psi^{(i)}(\vec{x}) = \psi(\vec{x}^{(i)}, \vec{x}) = \exp \left(- \sum_{j=1}^k \theta_j |x_j^{(i)} - x_j|^{p_j} \right), \quad (84.1)$$

where $\vec{x}$ and $\vec{x}^{(i)}$ denote the k -dim vector $\vec{x} = (x_1, \dots, x_k)^T$ and $\vec{x}^{(i)} = (x_1^{(i)}, \dots, x_k^{(i)})^T$, respectively.

□

84. Kriging (Gaussian Process Regression)

Kriging uses a specialized basis function that offers greater flexibility than standard RBFs. Examining Equation 84.1, we can observe how Kriging builds upon and extends the Gaussian basis concept. The key enhancements of Kriging over Gaussian RBF can be summarized as follows:

- Dimension-specific width parameters: While a Gaussian RBF uses a single width parameter $1/\sigma^2$, Kriging employs a vector $\tilde{\theta} = (\theta_1, \theta_2, \dots, \theta_k)^T$. This allows the model to automatically adjust its sensitivity to each input dimension, effectively performing automatic feature relevance determination.
- Flexible smoothness control: The Gaussian RBF fixes the exponent at 2, producing uniformly smooth functions. In contrast, Kriging's dimension-specific exponents $\tilde{p} = (p_1, p_2, \dots, p_k)^T$ (typically with $p_j \in [1, 2]$) enable precise control over smoothness properties in each dimension.
- Unifying framework: When all exponents are set to $p_j = 2$ and all width parameters are equal ($\theta_j = 1/\sigma^2$ for all j), the Kriging basis function reduces exactly to the Gaussian RBF. This makes Gaussian RBF a special case within the more general Kriging framework.

These enhancements make Kriging particularly well-suited for engineering problems where variables may operate at different scales or exhibit varying degrees of smoothness across dimensions. For now, we will only consider Kriging interpolation. We will cover Kriging regression later.

84.2. Building the Kriging Model

Consider sample data X and $\vec{y}$ from n locations that are available in matrix form: X is a $(n \times k)$ matrix, where k denotes the problem dimension and $\vec{y}$ is a $(n \times 1)$ vector. We want to find an expression for a predicted values at a new point $\vec{x}$, denoted as $\hat{y}$.

We start with an abstract, not really intuitive concept: The observed responses $\vec{y}$ are considered as if they are from a stochastic process, which will be denoted as

$$\begin{pmatrix} Y(\vec{x}^{(1)}) \\ \vdots \\ Y(\vec{x}^{(n)}) \end{pmatrix}. \quad (84.2)$$

The set of random vectors from Equation 84.2 (also referred to as a *random field*) has a mean of $\vec{1}\mu$, which is a $(n \times 1)$ vector. The random vectors are correlated with each other using the basis function expression from Equation 84.1:

$$\text{cor}(Y(\vec{x}^{(i)}), Y(\vec{x}^{(l)})) = \exp\left(-\sum_{j=1}^k \theta_j |x_j^{(i)} - x_j^{(l)}|^{p_j}\right). \quad (84.3)$$

Using Equation 84.3, we can compute the $(n \times n)$ correlation matrix Ψ of the observed sample data as shown in Equation 84.4,

$$\Psi = \begin{pmatrix} \text{cor}(Y(\vec{x}^{(1)}), Y(\vec{x}^{(1)})) & \dots & \text{cor}(Y(\vec{x}^{(1)}), Y(\vec{x}^{(n)})) \\ \vdots & \ddots & \vdots \\ \text{cor}(Y(\vec{x}^{(n)}), Y(\vec{x}^{(1)})) & \dots & \text{cor}(Y(\vec{x}^{(n)}), Y(\vec{x}^{(n)})) \end{pmatrix}, \quad (84.4)$$

and a covariance matrix as shown in Equation 84.5,

$$\text{Cov}(Y, Y) = \sigma^2 \Psi. \quad (84.5)$$

This assumed correlation between the sample data reflects our expectation that an engineering function will behave in a certain way and it will be smoothly and continuous.

Remark 84.1 (Note on Stochastic Processes). See Section 94.4 for a more detailed discussion on realizations of stochastic processes.

□

We now have a set of n random variables ($\mathbf{Y}$) that are correlated with each other as described in the $(n \times n)$ correlation matrix Ψ , see Equation 84.4. The correlations depend on the absolute distances in dimension j between the i -th and the l -th sample point $|x_j^{(i)} - x_j^{(l)}|$ and the corresponding parameters p_j and θ_j for dimension j . The correlation is intuitive, because when

- two points move close together, then $|x_j^{(i)} - x_j| \rightarrow 0$ and $\exp(-|x_j^{(i)} - x_j|^{p_j}) \rightarrow 1$ (these points show very close correlation and $Y(x_j^{(i)}) = Y(x_j)$).
- two points move far apart, then $|x_j^{(i)} - x_j| \rightarrow \infty$ and $\exp(-|x_j^{(i)} - x_j|^{p_j}) \rightarrow 0$ (these points show very low correlation).

Example 84.1 (Correlations for different p_j). Three different correlations are shown in Figure 84.1: $p_j = 0.1, 1, 2$. The smoothness parameter p_j affects the correlation:

- With $p_j = 0.1$, there is basically no immediate correlation between the points and there is a near discontinuity between the points $Y(\vec{x}_j^{(i)})$ and $Y(\vec{x}_j)$.
- With $p_j = 2$, the correlation is more smooth and we have a continuous gradient through $x_j^{(i)} - x_j$.

Reducing p_j increases the rate at which the correlation initially drops with distance. This is shown in Figure 84.1.

□

Example 84.2 (Correlations for different θ). Figure 84.2 visualizes the correlation between two points $Y(\vec{x}_j^{(i)})$ and $Y(\vec{x}_j)$ for different values of θ . The parameter θ can be seen as a *width* parameter:

84. Kriging (Gaussian Process Regression)




Figure 84.1.: Correlations with varying p . θ set to 1.

- low θ_j means that all points will have a high correlation, with $Y(x_j)$ being similar across the sample.
- high θ_j means that there is a significant difference between the $Y(x_j)$'s.
- θ_j is a measure of how *active* the function we are approximating is.
- High θ_j indicate important parameters, see Figure 84.2.


□

Considering the activity parameter θ is useful in high-dimensional problems where it is difficult to visualize the design landscape and the effect of the variable is unknown. By examining the elements of the vector θ , we can identify the most important variables and focus on them. This is a crucial step in the optimization process, as it allows us to reduce the dimensionality of the problem and focus on the most important variables.

Example 84.3 (The Correlation Matrix (Detailed Computation)). Let $n = 4$ and $k = 3$. The sample plan is represented by the following matrix X :

$$X = \begin{pmatrix} x_{11} & x_{12} & x_{13} \\ x_{21} & x_{22} & x_{23} \\ x_{31} & x_{32} & x_{33} \\ x_{41} & x_{42} & x_{43} \end{pmatrix}$$

To compute the elements of the matrix Ψ , the following k (one for each of the k dimensions) (n, n) -matrices have to be computed:

Figure 84.2.: Correlations with varying θ . p set to 2.

- For $k = 1$, i.e., the first column of X :

$$D_1 = \begin{pmatrix} x_{11} - x_{11} & x_{11} - x_{21} & x_{11} - x_{31} & x_{11} - x_{41} \\ x_{21} - x_{11} & x_{21} - x_{21} & x_{21} - x_{31} & x_{21} - x_{41} \\ x_{31} - x_{11} & x_{31} - x_{21} & x_{31} - x_{31} & x_{31} - x_{41} \\ x_{41} - x_{11} & x_{41} - x_{21} & x_{41} - x_{31} & x_{41} - x_{41} \end{pmatrix}$$

- For $k = 2$, i.e., the second column of X :

$$D_2 = \begin{pmatrix} x_{12} - x_{12} & x_{12} - x_{22} & x_{12} - x_{32} & x_{12} - x_{42} \\ x_{22} - x_{12} & x_{22} - x_{22} & x_{22} - x_{32} & x_{22} - x_{42} \\ x_{32} - x_{12} & x_{32} - x_{22} & x_{32} - x_{32} & x_{32} - x_{42} \\ x_{42} - x_{12} & x_{42} - x_{22} & x_{42} - x_{32} & x_{42} - x_{42} \end{pmatrix}$$

- For $k = 3$, i.e., the third column of X :

$$D_3 = \begin{pmatrix} x_{13} - x_{13} & x_{13} - x_{23} & x_{13} - x_{33} & x_{13} - x_{43} \\ x_{23} - x_{13} & x_{23} - x_{23} & x_{23} - x_{33} & x_{23} - x_{43} \\ x_{33} - x_{13} & x_{33} - x_{23} & x_{33} - x_{33} & x_{33} - x_{43} \\ x_{43} - x_{13} & x_{43} - x_{23} & x_{43} - x_{33} & x_{43} - x_{43} \end{pmatrix}$$

Since the matrices are symmetric and the main diagonals are zero, it is sufficient to

84. Kriging (Gaussian Process Regression)

compute the following matrices:

$$D_1 = \begin{pmatrix} 0 & x_{11} - x_{21} & x_{11} - x_{31} & x_{11} - x_{41} \\ 0 & 0 & x_{21} - x_{31} & x_{21} - x_{41} \\ 0 & 0 & 0 & x_{31} - x_{41} \\ 0 & 0 & 0 & 0 \end{pmatrix}$$

$$D_2 = \begin{pmatrix} 0 & x_{12} - x_{22} & x_{12} - x_{32} & x_{12} - x_{42} \\ 0 & 0 & x_{22} - x_{32} & x_{22} - x_{42} \\ 0 & 0 & 0 & x_{32} - x_{42} \\ 0 & 0 & 0 & 0 \end{pmatrix}$$

$$D_3 = \begin{pmatrix} 0 & x_{13} - x_{23} & x_{13} - x_{33} & x_{13} - x_{43} \\ 0 & 0 & x_{23} - x_{33} & x_{23} - x_{43} \\ 0 & 0 & 0 & x_{33} - x_{43} \\ 0 & 0 & 0 & 0 \end{pmatrix}$$

We will consider $p_l = 2$. The differences will be squared and multiplied by θ_i , i.e.:

$$D_1 = \theta_1 \begin{pmatrix} 0 & (x_{11} - x_{21})^2 & (x_{11} - x_{31})^2 & (x_{11} - x_{41})^2 \\ 0 & 0 & (x_{21} - x_{31})^2 & (x_{21} - x_{41})^2 \\ 0 & 0 & 0 & (x_{31} - x_{41})^2 \\ 0 & 0 & 0 & 0 \end{pmatrix}$$

$$D_2 = \theta_2 \begin{pmatrix} 0 & (x_{12} - x_{22})^2 & (x_{12} - x_{32})^2 & (x_{12} - x_{42})^2 \\ 0 & 0 & (x_{22} - x_{32})^2 & (x_{22} - x_{42})^2 \\ 0 & 0 & 0 & (x_{32} - x_{42})^2 \\ 0 & 0 & 0 & 0 \end{pmatrix}$$

$$D_3 = \theta_3 \begin{pmatrix} 0 & (x_{13} - x_{23})^2 & (x_{13} - x_{33})^2 & (x_{13} - x_{43})^2 \\ 0 & 0 & (x_{23} - x_{33})^2 & (x_{23} - x_{43})^2 \\ 0 & 0 & 0 & (x_{33} - x_{43})^2 \\ 0 & 0 & 0 & 0 \end{pmatrix}$$

The sum of the three matrices $D = D_1 + D_2 + D_3$ will be calculated next:

$$\begin{pmatrix} 0 & \theta_1(x_{11} - x_{21})^2 + \theta_2(x_{12} - x_{22})^2 + \theta_3(x_{13} - x_{23})^2 & \theta_1(x_{11} - x_{31})^2 + \theta_2(x_{12} - x_{32})^2 + \theta_3(x_{13} - x_{33})^2 & \theta_1(x_{11} - x_{41})^2 + \theta_2(x_{12} - x_{42})^2 + \theta_3(x_{13} - x_{43})^2 \\ 0 & 0 & \theta_1(x_{21} - x_{31})^2 + \theta_2(x_{22} - x_{32})^2 + \theta_3(x_{23} - x_{33})^2 & \theta_1(x_{21} - x_{41})^2 + \theta_2(x_{22} - x_{42})^2 + \theta_3(x_{23} - x_{43})^2 \\ 0 & 0 & 0 & \theta_1(x_{31} - x_{41})^2 + \theta_2(x_{32} - x_{42})^2 + \theta_3(x_{33} - x_{43})^2 \\ 0 & 0 & 0 & 0 \end{pmatrix}$$

Finally,

$$\Psi = \exp(-D)$$

is computed.

Next, we will demonstrate how this computation can be implemented in Python. We will consider four points in three dimensions and compute the correlation matrix Ψ using the basis function from Equation 84.1. These points are placed at the origin, at the unit vectors, and at the points (100,100,100) and (101,100,100). So, they form two clusters: one at the origin and one at (100,100,100).

```
theta = np.array([1,2,3])
X = np.array([ [1,0,0], [0,1,0], [100, 100, 100], [101, 100, 100]])
X
```

```
array([[ 1,  0,  0],
       [ 0,  1,  0],
       [100, 100, 100],
       [101, 100, 100]])
```

```
def build_Psi(X, theta):
    n = X.shape[0]
    k = X.shape[1]
    D = zeros((k, n, n))
    for l in range(k):
        for i in range(n):
            for j in range(i, n):
                D[l, i, j] = theta[l]*(X[i,l] - X[j,l])**2
    D = sum(D)
    D = D + D.T
    return exp(-D)
```

```
Psi = build_Psi(X, theta)
Psi
```

```
array([[1.          , 0.04978707, 0.          , 0.          ],
       [0.04978707, 1.          , 0.          , 0.          ],
       [0.          , 0.          , 1.          , 0.36787944],
       [0.          , 0.          , 0.36787944, 1.          ]])
```

□

Example 84.4 (Example: The Correlation Matrix (Using Existing Functions)). The same result as computed in Example 84.3 can be obtained with existing python functions, e.g., from the package `scipy`.

84. Kriging (Gaussian Process Regression)




Figure 84.3.: Correlation matrix Ψ .

```
def build_Psi(X, theta, eps=sqrt(spacing(1))):
    return exp(- squareform(pdist(X,
        metric='sqeuclidean',
        out=None,
        w=theta))) + multiply(eye(X.shape[0]),
        eps)

Psi = build_Psi(X, theta, eps=.0)
Psi
```

```
array([[1.          , 0.04978707, 0.          , 0.          ],
       [0.04978707, 1.          , 0.          , 0.          ],
       [0.          , 0.          , 1.          , 0.36787944],
       [0.          , 0.          , 0.36787944, 1.          ]])
```

The condition number of the correlation matrix Ψ is a measure of how well the matrix can be inverted. A high condition number indicates that the matrix is close to singular, which can lead to numerical instability in computations involving the inverse of the matrix, see Section 85.2.

```
np.linalg.cond(Psi)
```

```
np.float64(2.163953413738652)
```

□

84.3. MLE to estimate θ and p

84.3.1. The Log-Likelihood

Until now, the observed data $\vec{y}$ was not used. We know what the correlations mean, but how do we estimate the values of θ_j and where does our observed data y come in? To estimate the values of $\vec{\theta}$ and $\vec{p}$, they are chosen to maximize the likelihood of $\vec{y}$,

$$L = L(Y(\vec{x}^{(1)}), \dots, Y(\vec{x}^{(n)}) | \mu, \sigma) = \frac{1}{(2\pi\sigma^2)^{n/2}} \exp \left[-\frac{\sum_{i=1}^n (Y(\vec{x}^{(i)}) - \mu)^2}{2\sigma^2} \right], \quad (84.6)$$

where μ is the mean of the observed data $\vec{y}$ and σ is the standard deviation of the errors ϵ , which can be expressed in terms of the sample data

$$L = \frac{1}{(2\pi\sigma^2)^{n/2} |\vec{\Psi}|^{1/2}} \exp \left[-\frac{(\vec{y} - \vec{1}\mu)^T \vec{\Psi}^{-1} (\vec{y} - \vec{1}\mu)}{2\sigma^2} \right]. \quad (84.7)$$

Remark 84.2. The transition from Equation 84.6 to Equation 84.7 reflects a shift from assuming independent errors in the observed data to explicitly modeling the *correlation structure* between the observed responses, which is a key aspect of the stochastic process framework used in methods like Kriging. It can be explained as follows:

1. **Initial Likelihood Expression (assuming independent errors):** Equation 84.6 is an expression for the likelihood of the data set, which is based on the assumption that the errors ϵ are *independently randomly distributed according to a normal distribution with standard deviation σ* . This form is characteristic of the likelihood of n independent observations $Y(\vec{x}^{(i)})$, each following a normal distribution with mean μ and variance σ^2 .
2. **Using Vector Notation.** The sum in the exponent, i.e.,

$$\sum_{i=1}^n (Y(\vec{x}^{(i)}) - \mu)^2$$

is equivalent to

$$(\vec{y} - \vec{1}\mu)^T (\vec{y} - \vec{1}\mu),$$

assuming $Y(\vec{x}^{(i)}) = y^{(i)}$ and using vector notation for $\vec{y}$ and $\vec{1}\mu$.

3. **Assuming Independent Observations:** Equation 84.6 assumes that the observations are independent, which means that the covariance between any two observations $Y(\vec{x}^{(i)})$ and $Y(\vec{x}^{(l)})$ is zero for $i \neq l$. In this case, the covariance matrix of the observations would be a diagonal matrix with σ^2 along the diagonal (i.e., $\sigma^2 I$), where I is the identity matrix.
4. **Stochastic Process and Correlation:** In the context of Kriging, the observed responses $\vec{y}$ are considered as if they are from a *stochastic process* or *random field*. This means the random variables $Y(\vec{x}^{(i)})$ at different locations $\vec{x}^{(i)}$ are not independent, but they correlated with each other. This correlation is described by an $(n \times n)$ **correlation matrix** Ψ , which is used instead of $\sigma^2 I$. The strength of the correlation between two points $Y(\vec{x}^{(i)})$ and $Y(\vec{x}^{(l)})$ depends on the distance between them and model parameters θ_j and p_j .
5. **Multivariate Normal Distribution:** When random variables are correlated, their joint probability distribution is generally described by a *multivariate distribution*. Assuming the stochastic process follows a Gaussian process, the joint distribution of the observed responses $\vec{y}$ is a **multivariate normal distribution**. A multivariate normal distribution for a vector $\vec{Y}$ with mean vector $\vec{\mu}$ and covariance matrix Σ has a probability density function given by:

$$p(\vec{y}) = \frac{1}{\sqrt{(2\pi)^n |\Sigma|}} \exp \left[-\frac{1}{2} (\vec{y} - \vec{\mu})^T \Sigma^{-1} (\vec{y} - \vec{\mu}) \right].$$

6. **Connecting the Expressions:** In the stochastic process framework, the following holds:

- The mean vector of the observed data $\vec{y}$ is $\vec{1}\mu$.
- The covariance matrix Σ is constructed by considering both the variance σ^2 and the correlations Ψ .
- The covariance between $Y(\vec{x}^{(i)})$ and $Y(\vec{x}^{(l)})$ is $\sigma^2 \text{cor}(Y(\vec{x}^{(i)}), Y(\vec{x}^{(l)}))$.
- Therefore, the covariance matrix is $\Sigma = \sigma^2 \vec{\Psi}$.
- Substituting $\vec{\mu} = \vec{1}\mu$ and $\Sigma = \sigma^2 \vec{\Psi}$ into the multivariate normal PDF formula, we get:

$$\Sigma^{-1} = (\sigma^2 \vec{\Psi})^{-1} = \frac{1}{\sigma^2} \vec{\Psi}^{-1}$$

and

$$|\Sigma| = |\sigma^2 \vec{\Psi}| = (\sigma^2)^n |\vec{\Psi}|.$$

The PDF becomes:

$$p(\vec{y}) = \frac{1}{\sqrt{(2\pi)^n (\sigma^2)^n |\vec{\Psi}|}} \exp \left[-\frac{1}{2} (\vec{y} - \vec{1}\mu)^T \left(\frac{1}{\sigma^2} \vec{\Psi}^{-1} \right) (\vec{y} - \vec{1}\mu) \right]$$

and simplifies to:

$$p(\vec{y}) = \frac{1}{(2\pi\sigma^2)^{n/2} |\vec{\Psi}|^{1/2}} \exp \left[-\frac{1}{2\sigma^2} (\vec{y} - \vec{1}\mu)^T \vec{\Psi}^{-1} (\vec{y} - \vec{1}\mu) \right].$$

This is the **likelihood of the sample data** $\vec{y}$ given the parameters μ , σ , and the correlation structure defined by the parameters within $\vec{\Psi}$ (i.e., $\vec{\theta}$ and $\vec{p}$).

In summary, the Equation 84.6 represents the likelihood under a simplified assumption of independent errors, whereas Equation 84.7 is the likelihood derived from the assumption that the observed data comes from a **multivariate normal distribution** where observations are correlated according to the matrix $\vec{\Psi}$. Equation 84.7, using the sample data vector $\vec{y}$ and the correlation matrix $\vec{\Psi}$, properly accounts for the dependencies between data points inherent in the stochastic process model. Maximizing this likelihood is how the correlation parameters $\vec{\theta}$ and $\vec{p}$ are estimated in Kriging.

□

Equation 84.7 can be formulated as the log-likelihood:

$$\ln(L) = -\frac{n}{2} \ln(2\pi\sigma) - \frac{1}{2} \ln |\vec{\Psi}| - \frac{(\vec{y} - \vec{1}\mu)^T \vec{\Psi}^{-1} (\vec{y} - \vec{1}\mu)}{2\sigma^2}. \quad (84.8)$$

84.3.2. Differentiation with Respect to μ

Looking at the log-likelihood function, only the last term depends on μ :

$$\frac{1}{2\sigma^2} (\vec{y} - \vec{1}\mu)^T \vec{\Psi}^{-1} (\vec{y} - \vec{1}\mu)$$

To differentiate this with respect to the scalar μ , we can use matrix calculus rules.

Let $\mathbf{v} = \vec{y} - \vec{1}\mu$. $\vec{y}$ is a constant vector with respect to μ , and $\vec{1}\mu$ is a vector whose derivative with respect to the scalar μ is $\vec{1}$. So, $\frac{\partial \mathbf{v}}{\partial \mu} = -\vec{1}$.

The term is in the form $\mathbf{v}^T \mathbf{A} \mathbf{v}$, where $\mathbf{A} = \vec{\Psi}^{-1}$ is a symmetric matrix. The derivative of $\mathbf{v}^T \mathbf{A} \mathbf{v}$ with respect to $\mathbf{v}$ is $2\mathbf{A} \mathbf{v}$ as explained in Remark 84.3.

Remark 84.3 (Derivative of a Quadratic Form). Consider the derivative of $\mathbf{v}^T \mathbf{A} \mathbf{v}$ with respect to $\mathbf{v}$:

- The derivative of a scalar function $f(\mathbf{v})$ with respect to a vector $\mathbf{v}$ is a vector (the gradient).
- For a quadratic form $\mathbf{v}^T \mathbf{A} \mathbf{v}$, where $\mathbf{A}$ is a matrix and $\mathbf{v}$ is a vector, the general formula for the derivative with respect to $\mathbf{v}$ is $\frac{\partial}{\partial \mathbf{v}} (\mathbf{v}^T \mathbf{A} \mathbf{v}) = \mathbf{A} \mathbf{v} + \mathbf{A}^T \mathbf{v}$. (This is a standard result in matrix calculus and explained in Equation 85.1).
- Since $\mathbf{A} = \vec{\Psi}^{-1}$ is a symmetric matrix, its transpose $\mathbf{A}^T$ is equal to $\mathbf{A}$.
- Substituting $\mathbf{A}^T = \mathbf{A}$ into the general derivative formula, we get $\mathbf{A} \mathbf{v} + \mathbf{A} \mathbf{v} = 2\mathbf{A} \mathbf{v}$.

□

84. Kriging (Gaussian Process Regression)

Using the chain rule for differentiation with respect to the scalar μ :

$$\frac{\partial}{\partial \mu}(\mathbf{v}^T \mathbf{A} \mathbf{v}) = 2 \left(\frac{\partial \mathbf{v}}{\partial \mu} \right)^T \mathbf{A} \mathbf{v}$$

Substituting $\frac{\partial \mathbf{v}}{\partial \mu} = -\vec{1}$ and $\mathbf{v} = \vec{y} - \vec{1}\mu$:

$$\frac{\partial}{\partial \mu}(\vec{y} - \vec{1}\mu)^T \vec{\Psi}^{-1}(\vec{y} - \vec{1}\mu) = 2(-\vec{1})^T \vec{\Psi}^{-1}(\vec{y} - \vec{1}\mu) = -2\vec{1}^T \vec{\Psi}^{-1}(\vec{y} - \vec{1}\mu)$$

Now, differentiate the full log-likelihood term depending on μ :

$$\frac{\partial}{\partial \mu} \left(-\frac{1}{2\sigma^2}(\vec{y} - \vec{1}\mu)^T \vec{\Psi}^{-1}(\vec{y} - \vec{1}\mu) \right) = -\frac{1}{2\sigma^2} (-2\vec{1}^T \vec{\Psi}^{-1}(\vec{y} - \vec{1}\mu)) = \frac{1}{\sigma^2} \vec{1}^T \vec{\Psi}^{-1}(\vec{y} - \vec{1}\mu)$$

Setting this to zero for maximization gives:

$$\frac{1}{\sigma^2} \vec{1}^T \vec{\Psi}^{-1}(\vec{y} - \vec{1}\mu) = 0.$$

Rearranging gives:

$$\vec{1}^T \vec{\Psi}^{-1}(\vec{y} - \vec{1}\mu) = 0.$$

Solving for μ gives:

$$\vec{1}^T \vec{\Psi}^{-1} \vec{y} = \mu \vec{1}^T \vec{\Psi}^{-1} \vec{1}.$$

84.3.3. Differentiation with Respect to σ

Let $\nu = \sigma^2$ for simpler differentiation notation. The log-likelihood becomes:

$$\ln(L) = C_1 - \frac{n}{2} \ln(\nu) - \frac{(\vec{y} - \vec{1}\mu)^T \vec{\Psi}^{-1}(\vec{y} - \vec{1}\mu)}{2\nu},$$

where $C_1 = -\frac{n}{2} \ln(2\pi) - \frac{1}{2} \ln |\vec{\Psi}|$ is a constant with respect to $\nu = \sigma^2$.

We differentiate with respect to ν :

$$\frac{\partial \ln(L)}{\partial \nu} = \frac{\partial}{\partial \nu} \left(-\frac{n}{2} \ln(\nu) \right) + \frac{\partial}{\partial \nu} \left(-\frac{(\vec{y} - \vec{1}\mu)^T \vec{\Psi}^{-1}(\vec{y} - \vec{1}\mu)}{2\nu} \right).$$

The first term's derivative is straightforward:

$$\frac{\partial}{\partial \nu} \left(-\frac{n}{2} \ln(\nu) \right) = -\frac{n}{2} \cdot \frac{1}{\nu} = -\frac{n}{2\sigma^2}.$$

For the second term, let $C_2 = (\vec{y} - \vec{1}\mu)^T \vec{\Psi}^{-1}(\vec{y} - \vec{1}\mu)$. This term is constant with respect to σ^2 . The derivative is:

$$\frac{\partial}{\partial \nu} \left(-\frac{C_2}{2\nu} \right) = -\frac{C_2}{2} \frac{\partial}{\partial \nu} (\nu^{-1}) = -\frac{C_2}{2} (-\nu^{-2}) = \frac{C_2}{2\nu^2} = \frac{(\vec{y} - \vec{1}\mu)^T \vec{\Psi}^{-1}(\vec{y} - \vec{1}\mu)}{2(\sigma^2)^2}.$$

Combining the derivatives, the gradient of the log-likelihood with respect to σ^2 is:

$$\frac{\partial \ln(L)}{\partial \sigma^2} = -\frac{n}{2\sigma^2} + \frac{(\vec{y} - \vec{1}\mu)^T \vec{\Psi}^{-1}(\vec{y} - \vec{1}\mu)}{2(\sigma^2)^2}.$$

Setting this to zero for maximization gives:

$$-\frac{n}{2\sigma^2} + \frac{(\vec{y} - \vec{1}\mu)^T \vec{\Psi}^{-1}(\vec{y} - \vec{1}\mu)}{2(\sigma^2)^2} = 0.$$

84.3.4. Results of the Optimizations

Optimization of the log-likelihood by taking derivatives with respect to μ and σ results in

$$\hat{\mu} = \frac{\vec{1}^T \vec{\Psi}^{-1} \vec{y}^T}{\vec{1}^T \vec{\Psi}^{-1} \vec{1}^T} \quad (84.9)$$

and

$$\hat{\sigma}^2 = \frac{(\vec{y} - \vec{1}\mu)^T \vec{\Psi}^{-1}(\vec{y} - \vec{1}\mu)}{n}. \quad (84.10)$$

84.3.5. The Concentrated Log-Likelihood Function

Combining the equations, i.e., substituting Equation 84.9 and Equation 84.10 into Equation 84.8 leads to the concentrated log-likelihood function:

$$\ln(L) \approx -\frac{n}{2} \ln(\hat{\sigma}) - \frac{1}{2} \ln |\vec{\Psi}|. \quad (84.11)$$

Remark 84.4 (The Concentrated Log-Likelihood).

- The first term in Equation 84.11 requires information about the measured point (observations) y_i .
- To maximize $\ln(L)$, optimal values of $\vec{\theta}$ and $\vec{p}$ are determined numerically, because the function (Equation 84.11) is not differentiable.

□

84.3.6. Optimizing the Parameters $\vec{\theta}$ and $\vec{p}$

The concentrated log-likelihood function is very quick to compute. We do not need a statistical model, because we are only interested in the maximum likelihood estimate (MLE) of θ and p . Optimizers such as Nelder-Mead, Conjugate Gradient, or Simulated Annealing can be used to determine optimal values for θ and p . After the optimization, the correlation matrix Ψ is built with the optimized θ and p values. This is best (most likely) Kriging model for the given data y .

Observing Figure 84.2, there's significant change between $\theta = 0.1$ and $\theta = 1$, just as there is between $\theta = 1$ and $\theta = 10$. Hence, it is sensible to search for θ on a logarithmic scale. Suitable search bounds typically range from 10^{-3} to 10^2 , although this is not a stringent requirement. Importantly, the scaling of the observed data does not affect the values of $\hat{\theta}$, but the scaling of the design space does. Therefore, it is advisable to consistently scale variable ranges between zero and one to ensure consistency in the degree of activity $\hat{\theta}_j$ represents across different problems.

84.3.7. Correlation and Covariance Matrices Revisited

The covariance matrix Σ is constructed by considering both the variance σ^2 and the correlation matrix Ψ . They are related as follows:

1. **Covariance vs. Correlation:** Covariance is a measure of the joint variability of two random variables, while correlation is a standardized measure of this relationship, ranging from -1 to 1. The relationship between covariance and correlation for two random variables X and Y is given by $\text{cor}(X, Y) = \text{cov}(X, Y) / (\sigma_X \sigma_Y)$, where σ_X and σ_Y are their standard deviations.
2. **The Covariance Matrix Σ :** The **covariance matrix** Σ (or $\text{Cov}(Y, Y)$ for the vector $\vec{Y}$) captures the **pairwise covariances** between all elements of the vector of observed responses.
3. **Connecting σ^2 and Ψ to Σ :** In the Kriging framework described, the variance of each observation is often assumed to be constant, σ^2 . The covariance between any two observations $Y(\vec{x}^{(i)})$ and $Y(\vec{x}^{(l)})$ is given by σ^2 multiplied by their correlation. That is,

$$\text{cov}(Y(\vec{x}^{(i)}), Y(\vec{x}^{(l)})) = \sigma^2 \text{cor}(Y(\vec{x}^{(i)}), Y(\vec{x}^{(l)})).$$

This relationship holds for *all* pairs of points. When expressed in matrix form, the covariance matrix Σ is the product of the variance σ^2 (a scalar) and the correlation matrix Ψ :

$$\Sigma = \sigma^2 \Psi.$$

In essence, the correlation matrix Ψ defines the *structure* or *shape* of the dependencies between the data points based on their locations. The parameter σ^2 acts as a **scaling factor** that converts these unitless correlation values (which are between -1 and 1) into

actual covariance values with units of variance, setting the overall level of variability in the system.

So, σ^2 tells us about the general spread or variability of the underlying process, while Ψ tells you *how* that variability is distributed and how strongly points are related to each other based on their positions. Together, they completely define the covariance structure of your observed data in the multivariate normal distribution used in Kriging.

84.4. Implementing an MLE of the Model Parameters

The matrix algebra necessary for calculating the likelihood is the most computationally intensive aspect of the Kriging process. It is crucial to ensure that the code implementation is as efficient as possible.

Given that Ψ (our correlation matrix) is symmetric, only half of the matrix needs to be computed before adding it to its transpose. When calculating the log-likelihood, several matrix inversions are required. The fastest approach is to conduct one Cholesky factorization and then apply backward and forward substitution for each inverse.

The Cholesky factorization is applicable only to positive-definite matrices, which Ψ generally is. However, if Ψ becomes nearly singular, such as when the $\vec{x}^{(i)}$'s are densely packed, the Cholesky factorization might fail. In these cases, one could employ an LU-decomposition, though the result might be unreliable. When Ψ is near singular, the best course of action is to either use regression techniques or, as we do here, assign a poor likelihood value to parameters generating the near singular matrix, thus diverting the MLE search towards better-conditioned Ψ matrices.

When working with correlation matrices, increasing the values on the main diagonal of a matrix will increase the absolute value of its determinant. A critical numerical consideration in calculating the concentrated log-likelihood is that for poorly conditioned matrices, $\det(\Psi)$ approaches zero, leading to potential numerical instability. To address this issue, it is advisable to calculate $\ln(|\Psi|)$ in Equation 84.11 using twice the sum of the logarithms of the diagonal elements of the Cholesky factorization. This approach provides a more numerically stable method for computing the log-determinant, as the Cholesky decomposition $\Psi = LL^T$ allows us to express $\ln(|\Psi|) = 2 \sum_{i=1}^n \ln(L_{ii})$, avoiding the direct computation of potentially very small determinant values.

84.5. Kriging Prediction

We will use the Kriging correlation Ψ to predict new values based on the observed data. The presentation follows the approach described in Forrester et al. (2008) and Bartz et al. (2022).

84. Kriging (Gaussian Process Regression)

Main idea for prediction is that the new $Y(\vec{x})$ should be consistent with the old sample data X . For a new prediction $\hat{y}$ at $\vec{x}$, the value of $\hat{y}$ is chosen so that it maximizes the likelihood of the sample data X and the prediction, given the (optimized) correlation parameter $\vec{\theta}$ and $\vec{p}$ from above. The observed data $\vec{y}$ is augmented with the new prediction $\hat{y}$ which results in the augmented vector $\vec{\tilde{y}} = (\vec{y}^T, \hat{y})^T$. A vector of correlations between the observed data and the new prediction is defined as

$$\vec{\psi} = \begin{pmatrix} \text{cor}(Y(\vec{x}^{(1)}), Y(\vec{x})) \\ \vdots \\ \text{cor}(Y(\vec{x}^{(n)}), Y(\vec{x})) \end{pmatrix} = \begin{pmatrix} \psi^{(1)} \\ \vdots \\ \psi^{(n)} \end{pmatrix}.$$

Definition 84.2 (The Augmented Correlation Matrix). The augmented correlation matrix is constructed as

$$\tilde{\Psi} = \begin{pmatrix} \vec{\Psi} & \vec{\psi} \\ \vec{\psi}^T & 1 \end{pmatrix}.$$

□

The log-likelihood of the augmented data is

$$\ln(L) = -\frac{n}{2} \ln(2\pi) - \frac{n}{2} \ln(\hat{\sigma}^2) - \frac{1}{2} \ln |\tilde{\Psi}| - \frac{(\vec{\tilde{y}} - \vec{1}\hat{\mu})^T \tilde{\Psi}^{-1} (\vec{\tilde{y}} - \vec{1}\hat{\mu})}{2\hat{\sigma}^2}, \quad (84.12)$$

where $\vec{1}$ is a vector of ones and $\hat{\mu}$ and $\hat{\sigma}^2$ are the MLEs from Equation 84.9 and Equation 84.10. Only the last term in Equation 84.12 depends on $\hat{y}$, so we need only consider this term in the maximization. Details can be found in Forrester et al. (2008). Finally, the MLE for $\hat{y}$ can be calculated as

$$\hat{y}(\vec{x}) = \hat{\mu} + \vec{\psi}^T \tilde{\Psi}^{-1} (\vec{\tilde{y}} - \vec{1}\hat{\mu}). \quad (84.13)$$

Equation 84.13 reveals two important properties of the Kriging predictor:

- **Basis functions:** The basis function impacts the vector $\vec{\psi}$, which contains the n correlations between the new point $\vec{x}$ and the observed locations. Values from the n basis functions are added to a mean base term μ with weightings

$$\vec{w} = \tilde{\Psi}^{(-1)} (\vec{\tilde{y}} - \vec{1}\hat{\mu}).$$

- **Interpolation:** The predictions interpolate the sample data. When calculating the prediction at the i th sample point, $\vec{x}^{(i)}$, the i th column of $\tilde{\Psi}^{-1}$ is $\vec{\psi}$, and $\vec{\psi} \tilde{\Psi}^{-1}$ is the i th unit vector. Hence,

$$\hat{y}(\vec{x}^{(i)}) = y^{(i)}.$$

84.6. Kriging Example: Sinusoid Function

Toy example in 1d where the response is a simple sinusoid measured at eight equally spaced x -locations in the span of a single period of oscillation.

84.6.1. Calculating the Correlation Matrix Ψ

The correlation matrix Ψ is based on the pairwise squared distances between the input locations. Here we will use $n = 8$ sample locations and θ is set to 1.0.

```
n = 8
X = np.linspace(0, 2*np.pi, n, endpoint=False).reshape(-1,1)
print(np.round(X, 2))
```

```
[[0.  ]
 [0.79]
 [1.57]
 [2.36]
 [3.14]
 [3.93]
 [4.71]
 [5.5 ]]
```

Evaluate at sample points

```
y = np.sin(X)
print(np.round(y, 2))
```

```
[[ 0.  ]
 [ 0.71]
 [ 1.  ]
 [ 0.71]
 [ 0.  ]
 [-0.71]
 [-1.  ]
 [-0.71]]
```

We have the data points shown in Table 84.1.

84. Kriging (Gaussian Process Regression)

Table 84.1.: Data points for the sinusoid function

x	y
0.0	0.0
0.79	0.71
1.57	1.0
2.36	0.71
3.14	0.0
3.93	-0.71
4.71	-1.0
5.5	-0.71

The data points are visualized in Figure 84.4.

```
plt.plot(X, y, "bo")
plt.title(f"Sin(x) evaluated at {n} points")
plt.grid()
plt.show()
```




Figure 84.4.: Sin(x) evaluated at 8 points.

84.6.2. Computing the Ψ Matrix

We will use the `build_Psi` function from Example 84.4 to compute the correlation matrix Ψ . θ should be an array of one value, because we are only working in one dimension ($k = 1$).

```
theta = np.array([1.0])
Psi = build_Psi(X, theta)
print(np.round(Psi, 2))
```

```
[[1.  0.54 0.08 0.  0.  0.  0.  0. ]
 [0.54 1.  0.54 0.08 0.  0.  0.  0. ]
 [0.08 0.54 1.  0.54 0.08 0.  0.  0. ]
 [0.  0.08 0.54 1.  0.54 0.08 0.  0. ]
 [0.  0.  0.08 0.54 1.  0.54 0.08 0. ]
 [0.  0.  0.  0.08 0.54 1.  0.54 0.08]
 [0.  0.  0.  0.  0.08 0.54 1.  0.54]
 [0.  0.  0.  0.  0.  0.08 0.54 1. ]]
```

Figure 84.5 visualizes the $(8, 8)$ correlation matrix Ψ .




Figure 84.5.: Correlation matrix Ψ for the sinusoid function.

84.6.3. Selecting the New Locations

We would like to predict at $m = 100$ new locations (or testign locations) in the interval $[0, 2\pi]$. The new locations are stored in the variable $\mathbf{x}$.

```
m = 100
x = np.linspace(0, 2*np.pi, m, endpoint=False).reshape(-1,1)
```

84.6.4. Computing the ψ Vector

Distances between testing locations x and training data locations X .

```
def build_psi(X, x, theta, eps=sqrt(spacing(1))):
    n = X.shape[0]
    k = X.shape[1]
    m = x.shape[0]
    psi = zeros((n, m))
    theta = theta * ones(k)
    D = zeros((n, m))
    D = cdist(x.reshape(-1, k),
              X.reshape(-1, k),
              metric='sqeuclidean',
              out=None,
              w=theta)
    psi = exp(-D)
    # return psi transpose to be consistent with the literature
    print(f"Dimensions of psi: {psi.T.shape}")
    return(psi.T)

psi = build_psi(X, x, theta)
```

Dimensions of psi: (8, 100)

Figure 84.6 visualizes the (8,100) prediction matrix ψ .




Figure 84.6.: Visualization of the predition matrix ψ

84.6.5. Predicting at New Locations

Computation of the predictive equations.


```
U = cholesky(Psi).T
one = np.ones(n).reshape(-1,1)
mu = (one.T.dot(solve(U, solve(U.T, y)))) / one.T.dot(solve(U, solve(U.T, one)))
f = mu * ones(m).reshape(-1,1) + psi.T.dot(solve(U, solve(U.T, y - one * mu)))
print(f"Dimensions of f: {f.shape}")
```

Dimensions of f: (100, 1)

To compute f , Equation 84.13 is used.

84.6.6. Visualization

```
plt.plot(x, f, color = "orange", label="Fitted")
plt.plot(x, np.sin(x), color = "grey", label="Original")
plt.plot(X, y, "bo", label="Measurements")
plt.title("Kriging prediction of sin(x) with {} points.\n theta: {}".format(n, theta[0]))
plt.legend(loc='upper right')
plt.show()
```



84.6.7. The Complete Python Code for the Example

Here is the self-contained Python code for direct use in a notebook:

```
import numpy as np
import matplotlib.pyplot as plt
from numpy import (array, zeros, power, ones, exp, multiply, eye, linspace, spacing, s
from numpy.linalg import cholesky, solve
from scipy.spatial.distance import squareform, pdist, cdist

# --- 1. Kriging Basis Functions (Defining the Correlation) ---
# The core of Kriging uses a specialized basis function for correlation:
#  $\psi(x^{(i)}, x) = \exp(-\sum_{j=1}^k \theta_{\theta_j} |x_j^{(i)} - x_j|^{p_j})$ 
# For this 1D example ( $k=1$ ), and with  $p_j=2$  (squared Euclidean distance implicit from
# and  $\theta_{\theta_j} = \theta$  (a single value), it simplifies.

def build_Psi(X, theta, eps=sqrt(spacing(1))):
    """
    Computes the correlation matrix Psi based on pairwise squared Euclidean distances
    between input locations, scaled by theta.
    Adds a small epsilon to the diagonal for numerical stability (nugget effect).
```

84.6. Kriging Example: Sinusoid Function

```
"""
# Calculate pairwise squared Euclidean distances (D) between points in X
D = squareform(pdist(X, metric='sqeuclidean', out=None, w=theta))
# Compute Psi = exp(-D)
Psi = exp(-D)
# Add a small value to the diagonal for numerical stability (nugget)
# This is often done in Kriging implementations, though a regression method
# with a 'nugget' parameter (Lambda) is explicitly mentioned for noisy data later.
# The source code snippet for build_Psi explicitly includes `multiply(eye(X.shape), eps)`.
# FIX: Use X.shape to get the number of rows for the identity matrix
Psi += multiply(eye(X.shape[0]), eps) # Corrected line
return Psi

def build_psi(X_train, x_predict, theta):
    """
    Computes the correlation vector (or matrix) psi between new prediction locations
    and training data locations.
    """
    # Calculate pairwise squared Euclidean distances (D) between prediction points (x_predict)
    # and training points (X_train).
    # `cdist` computes distances between each pair of the two collections of inputs.
    D = cdist(x_predict, X_train, metric='sqeuclidean', out=None, w=theta)
    # Compute psi = exp(-D)
    psi = exp(-D)
    return psi.T # Return transpose to be consistent with literature (n x m or n x 1)

# --- 2. Data Points for the Sinusoid Function Example ---
# The example uses a 1D sinusoid measured at eight equally spaced x-locations [153, Table 9.1].
n = 8 # Number of sample locations
X_train = np.linspace(0, 2 * np.pi, n, endpoint=False).reshape(-1, 1) # Generate x-locations
y_train = np.sin(X_train) # Corresponding y-values (sine of x)

print("--- Training Data (X_train, y_train) ---")
print("x values:\n", np.round(X_train, 2))
print("y values:\n", np.round(y_train, 2))
print("-" * 40)

# Visualize the data points
plt.figure(figsize=(8, 5))
plt.plot(X_train, y_train, "bo", label=f"Measurements ({n} points)")
plt.title(f"Sin(x) evaluated at {n} points")
plt.xlabel("x")
plt.ylabel("sin(x)")
plt.grid(True)
```

84. Kriging (Gaussian Process Regression)

```
plt.legend()
plt.show()

# --- 3. Calculating the Correlation Matrix (Psi) ---
# Psi is based on pairwise squared distances between input locations.
# theta is set to 1.0 for this 1D example.
theta = np.array([1.0])
Psi = build_Psi(X_train, theta)

print("\n--- Computed Correlation Matrix (Psi) ---")
print("Dimensions of Psi:", Psi.shape) # Should be (8, 8)
print("First 5x5 block of Psi:\n", np.round(Psi[:5,:5], 2))
print("-" * 40)

# --- 4. Selecting New Locations (for Prediction) ---
# We want to predict at m = 100 new locations in the interval [0, 2*pi].
m = 100 # Number of new locations
x_predict = np.linspace(0, 2 * np.pi, m, endpoint=True).reshape(-1, 1)

print("\n--- New Locations for Prediction (x_predict) ---")
print(f"Number of prediction points: {m}")
print("First 5 prediction points:\n", np.round(x_predict[:5], 2).flatten())
print("-" * 40)

# --- 5. Computing the psi Vector ---
# This vector contains correlations between each of the n observed data points
# and each of the m new prediction locations.
psi = build_psi(X_train, x_predict, theta)

print("\n--- Computed Prediction Correlation Matrix (psi) ---")
print("Dimensions of psi:", psi.shape) # Should be (8, 100)
print("First 5x5 block of psi:\n", np.round(psi[:5,:5], 2))
print("-" * 40)

# --- 6. Predicting at New Locations (Kriging Prediction) ---
# The Maximum Likelihood Estimate (MLE) for y_hat is calculated using the formula:
# y_hat(x) = mu_hat + psi.T @ Psi_inv @ (y - 1 * mu_hat) [p. 2 of previous response, a
# Matrix inversion is efficiently performed using Cholesky factorization.

# Step 6a: Cholesky decomposition of Psi
U = cholesky(Psi).T # Note: `cholesky` in numpy returns lower triangular L, we need U

# Step 6b: Calculate mu_hat (estimated mean)
# mu_hat = (one.T @ Psi_inv @ y) / (one.T @ Psi_inv @ one) [p. 2 of previous response]
```

84.6. Kriging Example: Sinusoid Function

```
one = np.ones(n).reshape(-1, 1) # Vector of ones
mu_hat = (one.T @ solve(U, solve(U.T, y_train))) / (one.T @ solve(U, solve(U.T, one)))
mu_hat = mu_hat.item() # Extract scalar value

print("\n--- Kriging Prediction Calculation ---")
print(f"Estimated mean (mu_hat): {np.round(mu_hat, 4)}")

# Step 6c: Calculate predictions f (y_hat) at new locations
# f = mu_hat * ones(m) + psi.T @ Psi_inv @ (y - one * mu_hat)
f_predict = mu_hat * np.ones(m).reshape(-1, 1) + psi.T @ solve(U, solve(U.T, y_train - one * mu_hat))

print(f"Dimensions of predicted values (f_predict): {f_predict.shape}") # Should be (100, 1)
print("First 5 predicted f values:\n", np.round(f_predict[:5], 2).flatten())
print("-" * 40)


# --- 7. Visualization ---
# Plot the original sinusoid function, the measured points, and the Kriging predictions.

plt.figure(figsize=(10, 6))
plt.plot(x_predict, f_predict, color="orange", label="Kriging Prediction")
plt.plot(x_predict, np.sin(x_predict), color="grey", linestyle='--', label="True Sinusoid Function")
plt.plot(X_train, y_train, "bo", markersize=8, label="Measurements")
plt.title(f"Kriging prediction of sin(x) with {n} points. (theta: {theta})")
plt.xlabel("x")
plt.ylabel("y")
plt.legend(loc='upper right')
plt.grid(True)
plt.show()

--- Training Data (X_train, y_train) ---
x values:
[[0. ]
 [0.79]
 [1.57]
 [2.36]
 [3.14]
 [3.93]
 [4.71]
 [5.5 ]]
y values:
[[ 0. ]
 [ 0.71]
 [ 1. ]
 [ 0.71]
```

84. Kriging (Gaussian Process Regression)

```
[ 0. ]
[-0.71]
[-1. ]
[-0.71]]
```



--- Computed Correlation Matrix (Psi) ---

Dimensions of Psi: (8, 8)

First 5x5 block of Psi:

```
[[1.  0.54 0.08 0.  0. ]
 [0.54 1.  0.54 0.08 0. ]
 [0.08 0.54 1.  0.54 0.08]
 [0.  0.08 0.54 1.  0.54]
 [0.  0.  0.08 0.54 1.  ]]
```

--- New Locations for Prediction (x_predict) ---

Number of prediction points: 100

First 5 prediction points:

```
[0.  0.06 0.13 0.19 0.25]
```

84.6. Kriging Example: Sinusoid Function

--- Computed Prediction Correlation Matrix (psi) ---

Dimensions of psi: (8, 100)

First 5x5 block of psi:

```
[[1.   1.   0.98 0.96 0.94]
 [0.54 0.59 0.65 0.7  0.75]
 [0.08 0.1  0.12 0.15 0.18]
 [0.   0.01 0.01 0.01 0.01]
 [0.   0.   0.   0.   0.  ]]
```


--- Kriging Prediction Calculation ---

Estimated mean (μ_{hat}): -0.0499

Dimensions of predicted values (f_{predict}): (100, 1)

First 5 predicted f values:

```
[0.   0.05 0.1  0.15 0.21]
```



84.7. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

85. Matrices

85.1. Derivatives of Quadratic Forms

We present a step-by-step derivation of the general formula

$$\frac{\partial}{\partial \mathbf{v}}(\mathbf{v}^T \mathbf{A} \mathbf{v}) = \mathbf{A} \mathbf{v} + \mathbf{A}^T \mathbf{v}. \quad (85.1)$$

1. Define the components. Let $\mathbf{v}$ be a vector of size $n \times 1$, and let $\mathbf{A}$ be a matrix of size $n \times n$.
2. Write out the quadratic form in summation notation. The product $\mathbf{v}^T \mathbf{A} \mathbf{v}$ is a scalar. It can be expanded and be rewritten as a double summation:

$$\mathbf{v}^T \mathbf{A} \mathbf{v} = \sum_{i=1}^n \sum_{j=1}^n v_i a_{ij} v_j.$$

3. Calculate the partial derivative with respect to a component v_k : The derivative of the scalar $\mathbf{v}^T \mathbf{A} \mathbf{v}$ with respect to the vector $\mathbf{v}$ is the gradient vector, whose k -th component is $\frac{\partial}{\partial v_k}(\mathbf{v}^T \mathbf{A} \mathbf{v})$. We need to find $\frac{\partial}{\partial v_k} \left(\sum_{i=1}^n \sum_{j=1}^n v_i a_{ij} v_j \right)$. Consider the terms in the summation that involve v_k . A term $v_i a_{ij} v_j$ involves v_k if $i = k$ or $j = k$ (or both).

- Terms where $i = k$: $v_k a_{kj} v_j$. The derivative with respect to v_k is $a_{kj} v_j$.
- Terms where $j = k$: $v_i a_{ik} v_k$. The derivative with respect to v_k is $v_i a_{ik}$.
- The term where $i = k$ and $j = k$: $v_k a_{kk} v_k = a_{kk} v_k^2$. Its derivative with respect to v_k is $2a_{kk} v_k$. Notice this term is included in both cases above when $i = k$ and $j = k$. When $i = k$, the term is $v_k a_{kk} v_k$, derivative is $a_{kk} v_k$. When $j = k$, the term is $v_k a_{kk} v_k$, derivative is $v_k a_{kk}$. Summing these two gives $2a_{kk} v_k$.

4. Let's differentiate the sum $\sum_{i=1}^n \sum_{j=1}^n v_i a_{ij} v_j$ with respect to v_k :

$$\frac{\partial}{\partial v_k} \left(\sum_{i=1}^n \sum_{j=1}^n v_i a_{ij} v_j \right) = \sum_{i=1}^n \sum_{j=1}^n \frac{\partial}{\partial v_k} (v_i a_{ij} v_j).$$

5. The partial derivative $\frac{\partial}{\partial v_k} (v_i a_{ij} v_j)$ is non-zero only if $i = k$ or $j = k$.

- If $i = k$ and $j \neq k$: $\frac{\partial}{\partial v_k}(v_k a_{kj} v_j) = a_{kj} v_j$.
 - If $i \neq k$ and $j = k$: $\frac{\partial}{\partial v_k}(v_i a_{ik} v_k) = v_i a_{ik}$.
 - If $i = k$ and $j = k$: $\frac{\partial}{\partial v_k}(v_k a_{kk} v_k) = \frac{\partial}{\partial v_k}(a_{kk} v_k^2) = 2a_{kk} v_k$.
6. So, the partial derivative is the sum of derivatives of all terms involving v_k :
 $\frac{\partial}{\partial v_k}(\mathbf{v}^T \mathbf{A} \mathbf{v}) = \sum_{j \neq k} (a_{kj} v_j) + \sum_{i \neq k} (v_i a_{ik}) + (2a_{kk} v_k)$.
7. We can rewrite this by including the $i = k, j = k$ term back into the summations:
 $\sum_{j \neq k} (a_{kj} v_j) + a_{kk} v_k + \sum_{i \neq k} (v_i a_{ik}) + v_k a_{kk}$ (since $v_k a_{kk} = a_{kk} v_k$) $= \sum_{j=1}^n a_{kj} v_j + \sum_{i=1}^n v_i a_{ik}$.
8. Convert back to matrix/vector notation: The first summation $\sum_{j=1}^n a_{kj} v_j$ is the k -th component of the matrix-vector product $\mathbf{A} \mathbf{v}$. The second summation $\sum_{i=1}^n v_i a_{ik}$ can be written as $\sum_{i=1}^n a_{ik} v_i$. Recall that the element in the k -th row and i -th column of the transpose matrix $\mathbf{A}^T$ is $(\mathbf{A}^T)_{ki} = a_{ik}$. So, $\sum_{i=1}^n a_{ik} v_i = \sum_{i=1}^n (\mathbf{A}^T)_{ki} v_i$, which is the k -th component of the matrix-vector product $\mathbf{A}^T \mathbf{v}$.
9. Assemble the gradient vector: The k -th component of the gradient $\frac{\partial}{\partial \mathbf{v}}(\mathbf{v}^T \mathbf{A} \mathbf{v})$ is $(\mathbf{A} \mathbf{v})_k + (\mathbf{A}^T \mathbf{v})_k$. Since this holds for all $k = 1, \dots, n$, the gradient vector is the sum of the two vectors $\mathbf{A} \mathbf{v}$ and $\mathbf{A}^T \mathbf{v}$. Therefore, the general formula for the derivative is $\frac{\partial}{\partial \mathbf{v}}(\mathbf{v}^T \mathbf{A} \mathbf{v}) = \mathbf{A} \mathbf{v} + \mathbf{A}^T \mathbf{v}$.

85.2. The Condition Number

A small value, `eps`, can be passed to the function `build_Psi` to improve the condition number. For example, `eps=sqrt(spacing(1))` can be used. The numpy function `spacing()` returns the distance between a number and its nearest adjacent number.

The condition number of a matrix is a measure of its sensitivity to small changes in its elements. It is used to estimate how much the output of a function will change if the input is slightly altered.

A matrix with a low condition number is well-conditioned, which means its behavior is relatively stable, while a matrix with a high condition number is ill-conditioned, meaning its behavior is unstable with respect to numerical precision.

```
import numpy as np

# Define a well-conditioned matrix (low condition number)
A = np.array([[1, 0.1], [0.1, 1]])
print("Condition number of A: ", np.linalg.cond(A))

# Define an ill-conditioned matrix (high condition number)
```

```
B = np.array([[1, 0.99999999], [0.99999999, 1]])
print("Condition number of B: ", np.linalg.cond(B))
```

```
Condition number of A: 1.2222222222222225
Condition number of B: 200000000.57495335
```

85.3. The Moore-Penrose Pseudoinverse

85.3.1. Definitions

The Moore-Penrose pseudoinverse is a generalization of the inverse matrix for non-square or singular matrices. It is computed as

$$A^+ = (A^*A)^{-1}A^*,$$

where A^* is the conjugate transpose of A .

It satisfies the following properties:

1. $AA^+A = A$
2. $A^+AA^+ = A^+$
3. $(AA^+)^* = AA^+$.
4. $(A^+A)^* = A^+A$
5. $A^+ = (A^*)^+$
6. $A^+ = A^T$ if A is a square matrix and A is invertible.

The pseudoinverse can be computed using Singular Value Decomposition (SVD).

85.3.2. Implementation in Python

```
import numpy as np
from numpy.linalg import pinv
A = np.array([[1, 2], [3, 4], [5, 6]])
print(f"Matrix A:\n {A}")
A_pseudo_inv = pinv(A)
print(f"Moore-Penrose Pseudoinverse:\n {A_pseudo_inv}")
```

Matrix A:

```
[[1 2]
 [3 4]
 [5 6]]
```

Moore-Penrose Pseudoinverse:

```
[[-1.33333333 -0.33333333  0.66666667]
 [ 1.08333333  0.33333333 -0.41666667]]
```

85.4. Strictly Positive Definite Kernels

85.4.1. Definition

Definition 85.1 (Strictly Positive Definite Kernel). A kernel function $k(x, y)$ is called strictly positive definite if for any finite collection of distinct points $x_1, x_2, \dots, x_n$ in the input space and any non-zero vector of coefficients $\alpha = (\alpha_1, \alpha_2, \dots, \alpha_n)$, the following inequality holds:

$$\sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j k(x_i, x_j) > 0. \quad (85.2)$$

In contrast, a kernel function $k(x, y)$ is called positive definite (but not strictly) if the “>” sign is replaced by “ $\geq$ ” in the above inequality.

85.4.2. Connection to Positive Definite Matrices

The connection between strictly positive definite kernels and positive definite matrices lies in the Gram matrix construction:

- When we evaluate a kernel function $k(x, y)$ at all pairs of data points in our sample, we construct the Gram matrix K where $K_{ij} = k(x_i, x_j)$.
- If the kernel function k is strictly positive definite, then for any set of distinct points, the resulting Gram matrix will be symmetric positive definite.

A symmetric matrix is positive definite if and only if for any non-zero vector α , the quadratic form $\alpha^T K \alpha > 0$, which directly corresponds to the kernel definition above.

85.4.3. Connection to RBF Models

For RBF models, the kernel function is the radial basis function itself:

$$k(x, y) = \psi(\|x - y\|).$$

The Gaussian RBF kernel $\psi(r) = e^{-r^2/(2\sigma^2)}$ is strictly positive definite in $\mathbb{R}^n$ for any dimension n . The inverse multiquadric kernel $\psi(r) = (r^2 + \sigma^2)^{-1/2}$ is also strictly positive definite in any dimension.

This mathematical property guarantees that the interpolation problem has a unique solution (the weight vector $\vec{w}$ is uniquely determined). The linear system $\Psi\vec{w} = \vec{y}$ can be solved reliably using Cholesky decomposition. The RBF interpolant exists and is unique for any distinct set of centers.

85.5. Cholesky Decomposition and Positive Definite Matrices

We consider the definiteness of a matrix, before discussing the Cholesky decomposition.

Definition 85.2 (Positive Definite Matrix). A symmetric matrix A is positive definite if all its eigenvalues are positive.

Example 85.1 (Positive Definite Matrix). Given a symmetric matrix $A = \begin{pmatrix} 9 & 4 \\ 4 & 9 \end{pmatrix}$, the eigenvalues of A are $\lambda_1 = 13$ and $\lambda_2 = 5$. Since both eigenvalues are positive, the matrix A is positive definite.

Definition 85.3 (Negative Definite, Positive Semidefinite, and Negative Semidefinite Matrices). Similarly, a symmetric matrix A is negative definite if all its eigenvalues are negative. It is positive semidefinite if all its eigenvalues are non-negative, and negative semidefinite if all its eigenvalues are non-positive.

The covariance matrix must be positive definite for a multivariate normal distribution for a couple of reasons:

- Semidefinite vs Definite: A covariance matrix is always symmetric and positive semidefinite. However, for a multivariate normal distribution, it must be positive definite, not just semidefinite. This is because a positive semidefinite matrix can have zero eigenvalues, which would imply that some dimensions in the distribution have zero variance, collapsing the distribution in those dimensions. A positive definite matrix has all positive eigenvalues, ensuring that the distribution has positive variance in all dimensions.

85. Matrices

- Invertibility: The multivariate normal distribution's probability density function involves the inverse of the covariance matrix. If the covariance matrix is not positive definite, it may not be invertible, and the density function would be undefined.

In summary, the covariance matrix being positive definite ensures that the multivariate normal distribution is well-defined and has positive variance in all dimensions.

The definiteness of a matrix can be checked by examining the eigenvalues of the matrix. If all eigenvalues are positive, the matrix is positive definite.

```
import numpy as np

def is_positive_definite(matrix):
    return np.all(np.linalg.eigvals(matrix) > 0)

matrix = np.array([[9, 4], [4, 9]])
print(is_positive_definite(matrix)) # Outputs: True
```

True

However, a more efficient way to check the definiteness of a matrix is through the Cholesky decomposition.

Definition 85.4 (Cholesky Decomposition). For a given symmetric positive-definite matrix $A \in \mathbb{R}^{n \times n}$, there exists a unique lower triangular matrix $L \in \mathbb{R}^{n \times n}$ with positive diagonal elements such that:

$$A = LL^T.$$

Here, L^T denotes the transpose of L .

Example 85.2 (Cholesky decomposition using `numpy`). `linalg.cholesky` computes the Cholesky decomposition of a matrix, i.e., it computes a lower triangular matrix L such that $LL^T = A$. If the matrix is not positive definite, an error (`LinAlgError`) is raised.

```
import numpy as np

# Define a Hermitian, positive-definite matrix
A = np.array([[9, 4], [4, 9]])

# Compute the Cholesky decomposition
L = np.linalg.cholesky(A)
```

```
print("L = \n", L)
print("L*LT = \n", np.dot(L, L.T))
```

```
L =
[[3.          0.          ]
 [1.33333333  2.68741925]]
L*LT =
[[9.  4.]
 [4.  9.]]
```

Example 85.3 (Cholesky Decomposition). Given a symmetric positive-definite matrix $A = \begin{pmatrix} 9 & 4 \\ 4 & 9 \end{pmatrix}$, the Cholesky decomposition computes the lower triangular matrix L such that $A = LL^T$. The matrix L is computed as:

$$L = \begin{pmatrix} 3 & 0 \\ 4/3 & 2 \end{pmatrix},$$

so that

$$LL^T = \begin{pmatrix} 3 & 0 \\ 4/3 & \sqrt{65}/3 \end{pmatrix} \begin{pmatrix} 3 & 4/3 \\ 0 & \sqrt{65}/3 \end{pmatrix} = \begin{pmatrix} 9 & 4 \\ 4 & 9 \end{pmatrix} = A.$$

An efficient implementation of the definiteness-check based on Cholesky is already available in the `numpy` library. It provides the `np.linalg.cholesky` function to compute the Cholesky decomposition of a matrix. This more efficient `numpy`-approach can be used as follows:

```
import numpy as np

def is_pd(K):
    try:
        np.linalg.cholesky(K)
        return True
    except np.linalg.linalg.LinAlgError as err:
        if 'Matrix is not positive definite' in err.message:
            return False
        else:
            raise
matrix = np.array([[9, 4], [4, 9]])
print(is_pd(matrix)) # Outputs: True
```

True

85.5.1. Example of Cholesky Decomposition

We consider dimension $k = 1$ and $n = 2$ sample points. The sample points are located at $x_1 = 1$ and $x_2 = 5$. The response values are $y_1 = 2$ and $y_2 = 10$. The correlation parameter is $\theta = 1$ and p is set to 1. Using Equation 84.1, we can compute the correlation matrix Ψ :

$$\Psi = \begin{pmatrix} 1 & e^{-1} \\ e^{-1} & 1 \end{pmatrix}.$$

To determine MLE as in Equation 84.13, we need to compute Ψ^{-1} :

$$\Psi^{-1} = \frac{e}{e^2 - 1} \begin{pmatrix} e & -1 \\ -1 & e \end{pmatrix}.$$

Cholesky-decomposition of Ψ is recommended to compute Ψ^{-1} . Cholesky decomposition is a decomposition of a positive definite symmetric matrix into the product of a lower triangular matrix L , a diagonal matrix D and the transpose of L , which is denoted as L^T . Consider the following example:

$$\begin{aligned} LDL^T &= \begin{pmatrix} 1 & 0 \\ l_{21} & 1 \end{pmatrix} \begin{pmatrix} d_{11} & 0 \\ 0 & d_{22} \end{pmatrix} \begin{pmatrix} 1 & l_{21} \\ 0 & 1 \end{pmatrix} = \\ &= \begin{pmatrix} d_{11} & 0 \\ d_{11}l_{21} & d_{22} \end{pmatrix} \begin{pmatrix} 1 & l_{21} \\ 0 & 1 \end{pmatrix} = \begin{pmatrix} d_{11} & d_{11}l_{21} \\ d_{11}l_{21} & d_{11}l_{21}^2 + d_{22} \end{pmatrix}. \end{aligned} \quad (85.3)$$

Using Equation 85.3, we can compute the Cholesky decomposition of Ψ :

1. $d_{11} = 1$,
2. $l_{21}d_{11} = e^{-1} \Rightarrow l_{21} = e^{-1}$, and
3. $d_{11}l_{21}^2 + d_{22} = 1 \Rightarrow d_{22} = 1 - e^{-2}$.

The Cholesky decomposition of Ψ is

$$\Psi = \begin{pmatrix} 1 & 0 \\ e^{-1} & 1 \end{pmatrix} \begin{pmatrix} 1 & 0 \\ 0 & 1 - e^{-2} \end{pmatrix} \begin{pmatrix} 1 & e^{-1} \\ 0 & 1 \end{pmatrix} = LDL^T$$

Some programs use U instead of L . The Cholesky decomposition of Ψ is

$$\Psi = LDL^T = U^T DU.$$

Using

$$\sqrt{D} = \begin{pmatrix} 1 & 0 \\ 0 & \sqrt{1 - e^{-2}} \end{pmatrix},$$

we can write the Cholesky decomposition of Ψ without a diagonal matrix D as

$$\Psi = \begin{pmatrix} 1 & 0 \\ e^{-1} & \sqrt{1-e^{-2}} \end{pmatrix} \begin{pmatrix} 1 & e^{-1} \\ 0 & \sqrt{1-e^{-2}} \end{pmatrix} = U^T U.$$

85.5.2. Inverse Matrix Using Cholesky Decomposition

To compute the inverse of a matrix using the Cholesky decomposition, you can follow these steps:

1. Decompose the matrix A into L and L^T , where L is a lower triangular matrix and L^T is the transpose of L .
2. Compute L^{-1} , the inverse of L .
3. The inverse of A is then $(L^{-1})^T L^{-1}$.

Please note that this method only applies to symmetric, positive-definite matrices.

The inverse of the matrix Ψ from above is:

$$\Psi^{-1} = \frac{e}{e^2 - 1} \begin{pmatrix} e & -1 \\ -1 & e \end{pmatrix}.$$

Here's an example of how to compute the inverse of a matrix using Cholesky decomposition in Python:

```
import numpy as np
from scipy.linalg import cholesky, inv
E = np.exp(1)

# Psi is a symmetric, positive-definite matrix
Psi = np.array([[1, 1/E], [1/E, 1]])
L = cholesky(Psi, lower=True)
L_inv = inv(L)
# The inverse of A is (L^-1)^T * L^-1
Psi_inv = np.dot(L_inv.T, L_inv)

print("Psi:\n", Psi)
print("Psi Inverse:\n", Psi_inv)
```

```
Psi:
[[1.          0.36787944]
 [0.36787944  1.          ]]
Psi Inverse:
[[ 1.15651764 -0.42545906]
 [-0.42545906  1.15651764]]
```

85.6. Nyström Approximation

85.6.1. What's the Big Idea?

Imagine you have a huge, detailed map of a country. Working with the full, high-resolution map is slow and takes up a lot of computer memory. The Nyström method is like creating a smaller-scale summary map by only looking at a few key, representative locations.

In machine learning, we often work with a **kernel matrix** (or Gram matrix), which tells us how similar every pair of data points is to each other. For very large datasets, this matrix can become massive, making it computationally expensive to store and process.

The Nyström method provides an efficient way to create a **low-rank approximation** of this large kernel matrix. In simple terms, it finds a “simpler” version of the matrix that captures its most important properties without needing to compute or store the whole thing.

85.6.2. How Does It Work?

The core idea is to select a small, random subset of the columns of the full kernel matrix and use them to reconstruct the entire matrix. Let's say our full kernel matrix is K .

1. **Sample:** Randomly select l columns from the n total columns of K . Let C be the $n \times l$ matrix of these sampled columns.
2. **Intersect:** Take the rows of C corresponding to the sampled column indices to form the $l \times l$ matrix W .
3. **Approximate:** Using C and W , calculate the Nyström approximation $\tilde{K}$ of K :

$$\tilde{K} \approx CW^+C^T$$

where W^+ is the pseudoinverse of W .

85.6.3. Example

Suppose we have 4 data points and the full kernel matrix K is:

$$K = \begin{pmatrix} 9 & 6 & 3 & 1 \\ 6 & 4 & 2 & 0.5 \\ 3 & 2 & 1 & 0.25 \\ 1 & 0.5 & 0.25 & 0.1 \end{pmatrix}$$

Let's approximate it by sampling 2 columns ($l = 2$):

1. **Sample:** Pick the 1st and 3rd columns:

$$C = \begin{pmatrix} 9 & 3 \\ 6 & 2 \\ 3 & 1 \\ 1 & 0.25 \end{pmatrix}$$

2. **Intersect:** Take the 1st and 3rd rows from C to form W :

$$W = \begin{pmatrix} 9 & 3 \\ 3 & 1 \end{pmatrix}$$

3. **Approximate:** Suppose the pseudoinverse of W is:

$$W^+ = \begin{pmatrix} 0.09 & -0.27 \\ -0.27 & 0.81 \end{pmatrix}$$

Then,

$$\tilde{K} = CW^+C^T = \begin{pmatrix} 9 & 6 & 3 & 0.675 \\ 6 & 4 & 2 & 0.45 \\ 3 & 2 & 1 & 0.225 \\ 0.675 & 0.45 & 0.225 & 0.05 \end{pmatrix}$$

$\tilde{K}$ is a good approximation of the original K , especially in the top-left portion.

85.6.4. Why Is This Useful?

- **Speed:** The Nyström method is much faster than computing the full kernel matrix. The complexity is roughly $O(l^2n)$ instead of $O(n^2d)$ (where d is the number of features).
- **Scalability:** It allows kernel methods (like SVM or Kernel PCA) to be used on much larger datasets.

- **Feature Mapping:** The method can be used to project new data points into the same feature space for prediction tasks.

The quality of the approximation depends on the columns you sample. Uniform random sampling is common and often effective, but more advanced techniques exist to select more informative columns.

85.6.5. Applying the Nyström Approximation: How Nyström Approximation Helps Kriging

Kriging can significantly benefit from the Nyström approximation, especially when dealing with large datasets. Kriging is a spatial interpolation method used to estimate values at unmeasured locations based on observed points. It relies on a **covariance matrix** (often denoted as $\mathbf{K}$) that describes the spatial correlation between all observed data points.

The Problem with Standard Kriging:

The main computational challenge in Kriging is solving for the weights needed for prediction, which requires **inverting the covariance matrix $\mathbf{K}$** . For n data points, $\mathbf{K}$ is an $n \times n$ matrix, and inverting it has computational complexity $O(n^3)$. This becomes impractical for large datasets.

The Nyström Solution:

Since the covariance matrix in Kriging is a type of kernel matrix, we can use the Nyström method to create a low-rank approximation, $\tilde{\mathbf{K}}$. Instead of inverting the full matrix, we use the **Woodbury matrix identity** on the Nyström approximation, allowing us to efficiently compute $\tilde{\mathbf{K}}^{-1}$ without forming the full matrix. This reduces computational complexity to roughly $O(l^2n)$, where l is the number of sampled columns.

In summary, Nyström makes Kriging feasible for large-scale problems by replacing expensive matrix inversion with a faster, memory-efficient approximation.

85.6.6. Example: Predicting Temperature with Nyström-Kriging

Suppose we have temperature readings from 100 weather stations ($n=100$) and want to predict the temperature at a new location.

Data:

- Observed Locations ($\mathbf{X}$): 100 coordinate pairs
- Observed Temperatures ($\mathbf{y}$): 100 values
- Prediction Location ($\mathbf{x}^*$): Coordinates of the new location

85.6.6.1. Step 1: Nyström Approximation of the Covariance Matrix

1. **Sample Representative Points:** Randomly select 1=10 stations as landmarks.
2. **Compute C and W:**
 - **C:** Covariance between all 100 stations and the 10 landmarks (100x10 matrix)
 - **W:** Covariance among the 10 landmarks (10x10 matrix)

Nyström approximation: $\tilde{K} = CW^+C^T$

85.6.6.2. Step 2: Modeling and Prediction

Standard Kriging prediction:

$$y(x^*) = \mathbf{k}^{*T} \mathbf{K}^{-1} \mathbf{y}$$

where $\mathbf{k}^{*T}$ is the covariance vector between the prediction location and all observed locations.

Nyström-Kriging prediction:

$$y(x^*) \approx \mathbf{k}^{*T} (\text{fast_approx_inverse}(\mathbf{C}, \mathbf{W})) \mathbf{y}$$

Prediction Steps:

1. Calculate $\mathbf{k}^{*T}$: Covariance between new location and all stations.
2. Approximate the inverse term using the Woodbury identity with **C** and **W**.
3. Make the prediction: Take the dot product of $\mathbf{k}^{*T}$ and the weights vector.

This yields an accurate prediction efficiently, enabling rapid mapping for large regions.

85.6.7. Details: Woodbury Matrix Identity for Avoiding the Big Inversion

First, what is the **Woodbury matrix identity**? It's a mathematical rule that tells you how to find the inverse of a matrix that's been modified slightly. Its most useful form is for a "low-rank update":

$$(A + UCV)^{-1} = A^{-1} - A^{-1}U(C^{-1} + VA^{-1}U)^{-1}VA^{-1}$$

This looks complicated, but the core idea is simple:

- If you have a matrix A that is **easy to invert** (like a diagonal matrix).
- And you add a low-rank matrix to it (the UCV part, where C is small).

- You can find the new inverse without directly inverting the big $(A + UCV)$ matrix. Instead, you only need to invert the much smaller matrix in the middle of the formula: $(C^{-1} + VA^{-1}U)$.

How does this apply to the Nyström approximation?

In many machine learning and Kriging applications, we don't just need the kernel matrix $\tilde{K}$, but a “regularized” version, $(\lambda I + \tilde{K})$, where λI is a diagonal matrix that helps prevent overfitting. We need to find the inverse of this:

$$(\lambda I + \tilde{K})^{-1}$$

Substituting the Nyström formula $\tilde{K} = CW^+C^T$, we get:

$$(\lambda I + CW^+C^T)^{-1}$$

This expression fits the Woodbury identity perfectly!

- $A = \lambda I$ (very easy to invert: $A^{-1} = \frac{1}{\lambda}I$)
- $U = C$ (our $n \times l$ matrix)
- C (middle matrix) $= W^+$ (our small $l \times l$ matrix)
- $V = C^T$ (our $l \times n$ matrix)

By plugging these into the Woodbury formula, we get an expression for the inverse that only requires inverting a small $l \times l$ matrix. This means we never have to build the full $n \times n$ matrix $\tilde{K}$ or invert it directly. This is the source of the massive speed-up.

85.6.8. The Example: Step-by-Step

Let's reuse our 4-point example and show both the slow way and the fast Woodbury way.

Recall our matrices:

- $C = \begin{pmatrix} 9 & 3 \\ 6 & 2 \\ 3 & 1 \\ 1 & 0.25 \end{pmatrix}$
- $W^+ = \begin{pmatrix} 0.09 & -0.27 \\ -0.27 & 0.81 \end{pmatrix}$
- Let's use a regularization value $\lambda = 0.1$.

85.6.8.1. Method 1: The Slow Way (Forming the full matrix)

1. **Construct $\tilde{K}$:** First, we explicitly calculate the full 4×4 Nyström approximation $\tilde{K} = CW^+C^T$.

$$\tilde{K} = \begin{pmatrix} 9 & 6 & 3 & 0.675 \\ 6 & 4 & 2 & 0.45 \\ 3 & 2 & 1 & 0.225 \\ 0.675 & 0.45 & 0.225 & 0.05 \end{pmatrix}$$

2. **Add the regularization:** Now we compute $(\lambda I + \tilde{K})$.

$$(\lambda I + \tilde{K}) = \begin{pmatrix} 9.1 & 6 & 3 & 0.675 \\ 6 & 4.1 & 2 & 0.45 \\ 3 & 2 & 1.1 & 0.225 \\ 0.675 & 0.45 & 0.225 & 0.15 \end{pmatrix}$$

3. **Invert the 4×4 matrix:** This is the expensive step. The result is:

$$(\lambda I + \tilde{K})^{-1} \approx \begin{pmatrix} 9.85 & -14.78 & -0.07 & 0.27 \\ -14.78 & 22.22 & 0.09 & -0.41 \\ -0.07 & 0.09 & 0.91 & -0.03 \\ 0.27 & -0.41 & -0.03 & 6.67 \end{pmatrix}$$

This works for our tiny 4×4 example, but it would be computationally infeasible if n was 10,000.

85.6.8.2. Method 2: The Fast Way (Using Woodbury Identity)

We use the Woodbury formula to get the same result without ever creating a 4×4 matrix. The formula simplifies to:

$$(\lambda I + \tilde{K})^{-1} = \frac{1}{\lambda} I - \frac{1}{\lambda^2} C \left(W + \frac{1}{\lambda} C^T C \right)^{-1} C^T$$

1. **Compute the small 2×2 pieces:**

- $C^T C = \begin{pmatrix} 127 & 42.25 \\ 42.25 & 14.0625 \end{pmatrix}$
- $W = \begin{pmatrix} 9 & 3 \\ 3 & 1 \end{pmatrix}$

- The matrix to invert is $W + \frac{1}{0.1} C^T C = W + 10 \cdot (C^T C)$, which is:

$$\begin{pmatrix} 9 & 3 \\ 3 & 1 \end{pmatrix} + \begin{pmatrix} 1270 & 422.5 \\ 422.5 & 140.625 \end{pmatrix} = \begin{pmatrix} 1279 & 425.5 \\ 425.5 & 141.625 \end{pmatrix}$$

2. **Invert the small 2×2 matrix:** This is the only inversion we need, and it's extremely fast.

$$(W + \frac{1}{\lambda} C^T C)^{-1} \approx \begin{pmatrix} 0.22 & -0.66 \\ -0.66 & 1.99 \end{pmatrix}$$

3. **Combine the results:** Now we plug this small inverse back into the full formula. The rest is just matrix multiplication, no more inversions.

- First, calculate the middle term: $M = \frac{1}{\lambda^2} C(\dots)^{-1} C^T$. This will result in a 4×4 matrix.
- Then, calculate the final result: $\frac{1}{\lambda} I - M$.

After performing these multiplications, you will get the **exact same 4×4 inverse matrix** as in the slow method.

The crucial difference is that the most expensive operation—the matrix inversion—was performed on a tiny 2×2 matrix instead of a 4×4 one. For a large-scale problem, this is the difference between a calculation that takes seconds and one that could take hours or even be impossible.

85.7. Extending spotpython's Kriging Surrogate with Nyström Approximation for Enhanced Scalability

85.7.1. Introduction: Overcoming the Scalability Challenge in Kriging for Sequential Optimization

The Sequential Parameter Optimization Toolbox (spotpython) is a framework for hyperparameter tuning and black-box optimization based on Sequential Model-Based Optimization (SMBO). At the core of SMBO lies a surrogate model that approximates the true, expensive objective. Kriging (Gaussian Process regression) is a premier choice because it provides both predictions and a principled measure of uncertainty. This uncertainty enables a balance between exploration and exploitation. In each SMBO iteration, the Kriging model is updated with new evaluations, refining its approximation and proposing the next points.

Standard Kriging requires constructing and inverting an $n \times n$ covariance matrix, where n is the number of data points. Matrix inversion scales as $O(n^3)$. During SMBO, n can reach hundreds or thousands; refitting the surrogate each iteration becomes

prohibitively expensive. This cubic scaling is the key obstacle to applying Kriging at larger scales.

We integrate the Nyström method into the `spotpython` Kriging class. The Nyström method yields a low-rank approximation of a symmetric positive semidefinite (SPSD) kernel matrix by selecting $l \ll n$ “landmark” points. It approximates the full $n \times n$ covariance while requiring inversion of only an $l \times l$ matrix, reducing fitting cost from $O(n^3)$ to $O(nl^2)$. This makes Kriging viable even when the number of function evaluations is large.

85.7.2. Report Objectives and Structure

- Review theoretical foundations of Kriging and Nyström approximation
- Present documented Python code updates for Kriging (as in `kriging.py`)
- Explain changes to `__init__`, `fit`, and `predict`
- Show how mixed variable types are preserved via `build_Psi` and `build_psi_vec`
- Provide practical usage guidance and a formal complexity analysis

85.8. Theoretical Foundations: The Nyström–Kriging Framework

85.8.1. A Primer on Kriging (Gaussian Process Regression)

Kriging models $f(x)$ as a Gaussian Process with mean function $m(\cdot)$ and covariance (kernel) $k(\cdot, \cdot)$. For training inputs $X = \{x_1, \dots, x_n\}$ and observations $y = \{y_1, \dots, y_n\}$:

$$y \sim \mathcal{N}(m(X), K(X, X) + \sigma_n^2 I)$$

For a new point x_* :

$$\mu(x_*) = k(x_*, X) [K(X, X) + \sigma_n^2 I]^{-1} y$$

$$\sigma^2(x_*) = k(x_*, x_*) - k(x_*, X) [K(X, X) + \sigma_n^2 I]^{-1} k(X, x_*)$$

The challenge is inverting the $n \times n$ matrix $K(X, X) + \sigma_n^2 I$.

85.8.2. The Nyström Method for Low-Rank Kernel Approximation

Select l landmark points $X_m \subset X$. Let: - $C = K_{nm} = K(X, X_m) \in \mathbb{R}^{n \times l}$ - $W = K_{mm} = K(X_m, X_m) \in \mathbb{R}^{l \times l}$ Then the Nyström approximation is:

$$\tilde{K}_{nn} = C W^+ C^\top = K_{nm} K_{mm}^+ K_{mn}$$

where W^+ is the pseudoinverse of W . The approximation has rank $\leq l$.

85.8.3. Justification for Landmark Selection

Uniform sampling without replacement is an effective and inexpensive strategy for selecting landmarks across varied datasets and kernels.

85.9. Implementation: A Scalable Kriging Class for spotpython

85.9.1. Updated kriging.py with Nyström Approximation (excerpt)

```
"""
Kriging surrogate with optional Nyström approximation.
"""
import numpy as np
from scipy.spatial.distance import cdist
from scipy.linalg import cholesky, cho_solve, solve_triangular

class Kriging:
    def __init__(self, fun_control, n_theta=None, theta=None, p=2.0,
                 corr="squared_exponential", isotropic=False,
                 approximation="None", n_landmarks=100):
        self.fun_control = fun_control
        self.dim = self.fun_control["lower"].shape
        self.p = p
        self.corr = corr
        self.isotropic = isotropic
        self.approximation = approximation
        self.n_landmarks = n_landmarks
        self.factor_mask = self.fun_control["var_type"] == "factor"
```

```

self.ordered_mask = ~self.factor_mask
self.n_theta = 1 if isotropic else (n_theta or self.dim)
self.theta = np.full(self.n_theta, 0.1) if theta is None else theta
self.X_, self.y_, self.L_, self.alpha_ = None, None, None, None
self.landmarks_, self.W_cho_, self.nystrom_alpha_ = None, None, None

def fit(self, X, y):
    self.X_, self.y_ = X, y
    n_samples = X.shape[0]
    if self.approximation.lower() == "nystroem" and n_samples > self.n_landmarks:
        return self._fit_nystrom(X, y)
    return self._fit_standard(X, y)

def _fit_standard(self, X, y):
    Psi = self.build_Psi(X, X)
    Psi[np.diag_indices_from(Psi)] += 1e-8
    try:
        self.L_ = cholesky(Psi, lower=True)
        self.alpha_ = cho_solve((self.L_, True), y)
    except np.linalg.LinAlgError:
        self.L_ = None
        self.alpha_ = np.linalg.pinv(Psi) @ y

def _fit_nystrom(self, X, y):
    n_samples = X.shape[0]
    idx = np.random.choice(n_samples, self.n_landmarks, replace=False)
    self.landmarks_ = X[idx, :]
    W = self.build_Psi(self.landmarks_, self.landmarks_) + 1e-8 * np.eye(self.n_landmarks)
    C = self.build_Psi(X, self.landmarks_)
    try:
        self.W_cho_ = cholesky(W, lower=True)
        self.nystrom_alpha_ = cho_solve((self.W_cho_, True), C.T @ y)
    except np.linalg.LinAlgError:
        self.W_cho_ = None
        self._fit_standard(X, y)

def predict(self, X_star):
    if self.approximation.lower() == "nystroem" and self.landmarks_ is not None:
        return self._predict_nystrom(X_star)
    return self._predict_standard(X_star)

def _predict_standard(self, X_star):
    psi = self.build_Psi(X_star, self.X_)
    y_pred = psi @ self.alpha_

```

```

    if self.L_ is not None:
        v = solve_triangular(self.L_, psi.T, lower=True)
        y_mse = 1.0 - np.sum(v**2, axis=0)
    else:
        Psi = self.build_Psi(self.X_, self.X_) + 1e-8 * np.eye(self.X_.shape[0])
        pi_Psi = np.linalg.pinv(Psi)
        y_mse = 1.0 - np.sum((psi @ pi_Psi) * psi, axis=1)
    y_mse[y_mse < 0] = 0
    return y_pred, y_mse.reshape(-1, 1)

def _predict_nystrom(self, X_star):
    psi_star_m = self.build_Psi(X_star, self.landmarks_)
    y_pred = psi_star_m @ self.nystrom_alpha_
    if self.W_cho_ is not None:
        v = cho_solve((self.W_cho_, True), psi_star_m.T)
        quad = np.sum(psi_star_m * v.T, axis=1)
        y_mse = 1.0 - quad
    else:
        y_mse = np.ones(X_star.shape[0])
    y_mse[y_mse < 0] = 0
    return y_pred, y_mse.reshape(-1, 1)

def build_Psi(self, X1, X2):
    n1 = X1.shape[0]
    Psi = np.zeros((n1, X2.shape[0]))
    for i in range(n1):
        Psi[i, :] = self.build_psi_vec(X1[i, :], X2)
    return Psi

def build_psi_vec(self, x, X_):
    theta10 = np.full(self.dim, 10**self.theta) if self.isotropic else 10**self.t
    D = np.zeros(X_.shape[0])
    if self.ordered_mask.any():
        Xo = X_[ :, self.ordered_mask]
        xo = x[self.ordered_mask]
        D += cdist(xo.reshape(1, -1), Xo, metric="sqeuclidean",
                    w=theta10[self.ordered_mask]).ravel()
    if self.factor_mask.any():
        Xf = X_[ :, self.factor_mask]
        xf = x[self.factor_mask]
        D += cdist(xf.reshape(1, -1), Xf, metric="hamming",
                    w=theta10[self.factor_mask]).ravel() * self.factor_mask.sum()
    return np.exp(-D) if self.corr == "squared_exponential" else np.exp(-(D**self

```

85.10. Implementation Details

85.10.1. Architectural Enhancements to `init`

- New argument `approximation="None"` for backward-compatible selection between exact Kriging and Nyström
- New argument `n_landmarks` (default 100) controls the number of inducing points when using Nyström
- State attributes for both exact and Nyström paths are maintained separately

85.10.2. The `fit()` Method: A Dual-Pathway Approach

- Dispatcher selecting exact or Nyström pathway
- The Nyström fit Pathway (`_fit_nystrom`):
 - Landmark selection via uniform sampling without replacement
 - Core matrices:
 - * $W = K_{mm}$ (landmark-landmark)
 - * $C = K_{nm}$ (data-landmark)
 - Cholesky factorization of W (with jitter) for stability
 - Pre-computation: $\alpha_{nys} = W^{-1}C^T y$ via `cho_solve`
- The Standard fit Pathway (`_fit_standard`):
 - Full Ψ construction, Cholesky decomposition, and solve for α
 - Fallback to pseudoinverse if Cholesky fails

85.10.3. The `predict()` Method: Conditional Prediction Logic

- Routes to Nyström or standard prediction path based on fitted model state
- The Nyström predict Pathway (`_predict_nystrom`):
 - Cross-covariance ψ between test points and landmarks
 - Mean: $\psi \cdot \alpha_{nys}$
 - Variance: uses `cho_solve` with W Cholesky; non-negative clipping
- The Standard predict Pathway (`_predict_standard`):
 - Cross-covariance with all training points
 - Mean from α ; variance via triangular solves or pseudoinverse fallback

85.10.4. Critical Detail: Preserving Mixed Variable Type Functionality

The Significance of `build_psi_vec`:

- Mixed spaces: continuous (ordered) and categorical (factor) variables
- Distances:
 - Weighted squared Euclidean for ordered variables
 - Weighted Hamming for factors
- Anisotropic kernel via per-dimension length-scales θ
- Nyström path reuses `build_Psi` $\rightarrow$ `build_psi_vec`, preserving mixed-type handling

85.10.5. Seamless Integration into the Nyström Workflow

All covariance computations (W , C , predictive cross-covariance) use `build_Psi`, ensuring identical handling for mixed variable types in both standard and Nyström modes.

85.11. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

86. Infill Criteria

In the context of computer experiments and surrogate modeling, a sampling plan refers to the set of input values, often denoted as X , at which a computer code is evaluated. The primary objective of a sampling plan is to efficiently explore the input space to understand the behavior of the computer code and to construct a surrogate model that accurately represents that behavior. Historically, Response Surface Methodology (RSM) provided methods for designing such plans, often based on rectangular grids or factorial designs. More recently, Design and Analysis of Computer Experiments (DACE) has emerged as a more flexible and powerful approach for this purpose.

A surrogate model, or $\hat{f}$, is built to approximate the expensive response of a black-box function $f(x)$. Since evaluating f is costly, only a sparse set of samples is used to construct $\hat{f}$, which can then provide inexpensive predictions for any point in the design space. However, as a surrogate model is inherently an approximation of the true function, its accuracy and predictive capabilities can be significantly improved by incorporating new data points, known as infill points. Infill points are strategically chosen to either reduce uncertainty, improve predictions in specific regions of interest, or enhance the model's ability to identify optima or trends.

The process of updating a surrogate model with infill points is iterative. It typically involves:

- **Identifying Regions of Interest:** Analyzing the current surrogate model to determine areas where it is inaccurate, has high uncertainty, or predicts promising results (e.g., potential optima).
- **Selecting Infill Points:** Choosing new data points based on specific criteria that balance different objectives.
- **Evaluating the True Function:** Running the actual simulation or experiment at the selected infill points to obtain their corresponding outputs.
- **Updating the Surrogate Model:** Retraining or updating the surrogate model using the new, augmented dataset.
- **Repeating:** Iterating this process until the model meets predefined accuracy criteria or the computational budget is exhausted.

86.1. Balancing Exploitation and Exploration

A crucial aspect of selecting infill points is navigating the inherent trade-off between exploitation and exploration.

Definition 86.1 (Exploitation). Exploitation refers to sampling near predicted optima to refine the solution. This strategy aims to rapidly converge on a good solution by focusing computational effort where the surrogate model suggests the best values might lie.

Definition 86.2 (Exploration). Exploration involves sampling in regions of high uncertainty to improve the global accuracy of the model. This approach ensures that the model is well-informed across the entire design space, preventing it from getting stuck in local optima.

Forrester et al. (2008) emphasizes that effective infill criteria are designed to combine both exploitation and exploration.

86.2. Expected Improvement (EI)

Expected Improvement (EI) is one of the most influential and widely-used infill criteria. Formalized by Jones et al. (1998) and building upon the work of Moćkus (1974), EI provides a mathematically elegant framework that naturally balances exploitation and exploration. Rather than simply picking the point with the best predicted value (pure exploitation) or the point with the highest uncertainty (pure exploration), EI asks a more nuanced question: “How much improvement over the current best solution can we *expect* to gain by evaluating the true function at a new point x ?”.

The Expected Improvement, $EI(x)$, can be calculated using the following formula:

$$EI(x) = \sigma(x) [Z\Phi(Z) + \phi(Z)]$$

where:

- $\mu(x)$ (or $\hat{y}(x)$) is the Kriging prediction (mean of the stochastic process) at a new, unobserved point x .
- $\sigma(x)$ (or $\hat{s}(x)$) is the estimated standard deviation (square root of the variance $\hat{s}^2(x)$) of the prediction at point x .
- f_{best} (or y_{min}) is the best (minimum, for minimization problems) observed function value found so far.
- $Z = \frac{f_{best} - \mu(x)}{\sigma(x)}$ is the standardized improvement.
- $\Phi(Z)$ is the cumulative distribution function (CDF) of the standard normal distribution.

86.3. Expected Improvement in the Hyperparameter Tuning Cookbook (Python Implementation)

- $\phi(Z)$ is the probability density function (PDF) of the standard normal distribution.

If $\sigma(x) = 0$ (meaning there is no uncertainty at point x , typically because it's an already sampled point), then $EI(x) = 0$, reflecting the intuition that no further improvement can be expected at a known point. A maximization of Expected Improvement as an infill criterion will eventually lead to the global optimum.

The elegance of the EI formula lies in its combination of two distinct terms:

- **Exploitation Term:** $(f_{best} - \mu(x))\Phi(Z)$. This part of the formula contributes more when the predicted value $\mu(x)$ is significantly lower (better) than the current best observed value f_{best} . It is weighted by the probability $\Phi(Z)$ that the true function value at x will indeed be an improvement over f_{best} .
- **Exploration Term:** $\sigma(x)\phi(Z)$. This term becomes larger when there is high uncertainty ($\sigma(x)$ is large) in the model's prediction at x . It accounts for the potential of discovering unexpectedly good values in areas that have not been thoroughly explored, even if the current mean prediction there is not the absolute best.

Expected Improvement offers several significant practical benefits:

- **Automatic Balance:** It inherently balances exploitation and exploration without requiring any manual adjustment of weights or parameters.
- **Scale Invariance:** EI is relatively insensitive to the scaling of the objective function, making it robust across various problem types.
- **Theoretical Foundation:** It is underpinned by a strong theoretical basis derived from decision theory and information theory.
- **Efficient Optimization:** The smooth and differentiable nature of the EI function allows for efficient optimization using gradient-based algorithms to find the next infill point.
- **Proven Performance:** EI has demonstrated consistent and strong performance in numerous real-world applications across various domains.

86.3. Expected Improvement in the Hyperparameter Tuning Cookbook (Python Implementation)

Within the context of the Hyperparameter Tuning Cookbook, Expected Improvement serves a critical role in Sequential Model-Based Optimization. It systematically guides the selection of which hyperparameter configurations to evaluate next, facilitating the efficient utilization of computational resources. By intelligently balancing the need to exploit promising regions and explore uncertain areas, EI helps identify optimal hyperparameters with a reduced number of expensive model training runs. This

provides a principled and automated method for navigating complex hyperparameter spaces without extensive manual intervention.

While the foundational concepts in Forrester et al. (2008) are often illustrated with MATLAB code, the Hyperparameter Tuning Cookbook emphasizes and provides implementations in Python. The `spotpython` package, consistent with the Cookbook's approach, provides a Python implementation of Expected Improvement within its Kriging class. For minimization problems, `spotpython` typically calculates and returns the negative Expected Improvement, aligning with standard optimization algorithm conventions. Furthermore, to enhance numerical stability and mitigate issues when EI values are very small, `spotpython` often works with a logarithmic transformation of EI and incorporates a small epsilon value.

86.4. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

87. Remote Optimization Challenge: A Benchmarking Study

This report performs an experimental comparison of two optimization algorithms on a remote objective function. The goal is to evaluate the efficiency and effectiveness of `SpotOptim` (a surrogate-based optimizer) compared to a standard gradient-based approach (Scipy L-BFGS-B) in a limited budget scenario.

87.1. Benchmarking Best Practices

Benchmarking is a critical aspect of validating optimization algorithms. According to *Bartz-Beielstein et al. (2020)* in “Benchmarking in Optimization: Best Practice and Open Issues” (arXiv:2007.03488), a sound benchmarking study should adhere to several key principles:

1. Clearly Stated Goals: The objective of this study is to compare the convergence speed and final solution quality of two distinct optimization strategies given a constrained budget of function evaluations.
2. Well-Specified Problems: We employ the “Wing Weight” problem, hosted remotely, which represents a realistic, expensive black-box engineering problem. It is a 10-dimensional problem defined on the unit hypercube $[0, 1]^{10}$.
3. Adequate Performance Measures: We report the *best-so-far* objective value as a function of the number of evaluations (convergence) and the distribution of the final solutions found.
4. Repetitions: Since the remote function (and potentially the algorithms) may exhibit stochastic behavior (noise in the function or random initialization in the algorithms), we perform multiple independent repetitions to ensure statistical reliability.

87.2. Experimental Setup

87.2.1. The Remote Objective Function

The objective function is accessed via `robot_arm_hard`. It simulates the weight of a light aircraft wing.

87. Remote Optimization Challenge: A Benchmarking Study

- Input: $X \in [0, 1]^{10}$
- Output: $y \in \mathbb{R}$ (to be minimized)
- Characteristics: Expensive to evaluate, potentially noisy.

87.2.2. Algorithms

1. SpotOptim: A Sequential Parameter Optimization tool that uses surrogate models (e.g., Kriging/Gaussian Processes) to guide the search. It is designed to be sample-efficient.
2. Scipy L-BFGS-B: A standard quasi-Newton method. While efficient for smooth local optimization, it typically requires many function evaluations to estimate gradients and may get stuck in local optima. We use random restarts (via the repetitions) to explore different basins.

87.2.3. Budget and Settings

- Maximum Evaluations: 50 per run.
- Repetitions: 5 independent runs (kept low for demonstration speed, usually 10-30 are recommended).
- SpotOptim Settings: Initial design $n = 10$, surrogate-based optimization for the remaining budget.
- Scipy Settings: L-BFGS-B with numeric differentiation.

87.3. Comparison of the Difficulty

To estimate the landscape difficulty of the remote function, we perform random linear cuts through the input space and plot the objective values along these cuts. We compare the following four functions:

1. Wing Weight (10D)
2. Lennard-Jones Potential (39D)
3. Remote Objective (10D), SPOTSeven Optimization Challenge

```
import matplotlib.pyplot as plt
import time
import numpy as np
# from spotoptim.function.remote import objective_remote
from spotoptim.function.so import wingwt, lennard_jones, robot_arm_hard

def get_landscape_cut(func, dim, steps=200):
    np.random.seed(42)
```

```

# Define two random points in  $[0, 1]^{\text{dim}}$ 
start_point = np.random.rand(dim)
end_point = np.random.rand(dim)

# Interpolate
alphas = np.linspace(0, 1, steps)
# Shape (steps, dim)
trajectory = np.outer(1 - alphas, start_point) + np.outer(alphas, end_point)

# Time execution
t0 = time.time()
values = func(trajectory)
dt = time.time() - t0

return alphas, values, dt * 1000 # ms

# Run Analysis
steps = 300

# 1. Wing Weight (10 dim)
a1, v1, t1 = get_landscape_cut(wingwt, 10, steps)

# 2. Lennard-Jones (39 dim)
a2, v2, t2 = get_landscape_cut(lennard_jones, 39, steps)

# 3. Challenge Remote (10 dim)
a3, v3, t3 = get_landscape_cut(robot_arm_hard, 10, steps)

fig, axes = plt.subplots(3, 1, figsize=(10, 9))

# Plot 1: Wing Weight
axes[0].plot(a1, v1, color='blue', linewidth=2)
axes[0].set_title(f"Wing Weight (10D)\nSmooth & Unimodal")
axes[0].set_xlabel("Interpolation  $\alpha$ ")
axes[0].set_ylabel("Weight")
axes[0].grid(True, alpha=0.3)

# Plot 2: Lennard-Jones
axes[1].plot(a2, v2, color='green', linewidth=2)
axes[1].set_title(f"Lennard-Jones (39D)\nRugged & Multimodal")
axes[1].set_xlabel("Interpolation  $\alpha$ ")
axes[1].set_ylabel("Potential Energy")
axes[1].grid(True, alpha=0.3)

```

87. Remote Optimization Challenge: A Benchmarking Study

```
# Plot 3: Remote Challenge
axes[2].plot(a3, v3, color='red', linewidth=2)
axes[2].set_title(f"Remote Challenge (10D)\nNoisy & Complex")
axes[2].set_xlabel("Interpolation  $\alpha$ ")
axes[2].set_ylabel("Objective Value")
axes[2].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```




Figure 87.1.: 1D Slice through High-Dimensional Landscapes

87.4. Statistical Power Analysis

To compare the algorithms fairly in the presence of noise, we cannot simply take the final objective value of a single run. The objective function $f(x)$ is noisy, meaning $y = f(x) + \epsilon$. We want to compare the true performance $\mu = E[f(x)]$ of the best solutions found.

We first estimate the noise level (σ) of the objective function and then determine the number of repetitions (n) required to distinguish between two solutions with sufficient statistical power.

87.4.1. 1. Estimating Noise (σ)

We evaluate a random design point 30 times to estimate the standard deviation of the noise.

```
# import numpy as np
# import statistics
# from spotoptim.function.so import robot_arm_hard

# # 1. Estimate Noise
# print("Estimating noise...")
# x_sample = np.array([[0.5] * 10])
# noise_samples = []
# for _ in range(30):
#     noise_samples.append(robot_arm_hard(x_sample)[0])
# # remove None values
# noise_samples = [x for x in noise_samples if x is not None]
# # heck if there are enough values for stdev
# if len(noise_samples) < 2:
#     raise ValueError("Not enough samples for stdev estimation")
# else:
#     sigma_est = statistics.stdev(noise_samples)
#     print(f"Estimated Noise (sigma): {sigma_est:.4f}")
```

87.4.2. 2. Determining Sample Size (n)

We perform a power analysis to calculate the number of evaluations (n) needed to validate a solution. We aim to detect a “large” effect size (Cohen’s $d = 1.0$, which corresponds to a difference of $1.0 \times \sigma$) with a power of 0.8 and significance level $\alpha = 0.05$.

87. Remote Optimization Challenge: A Benchmarking Study

```
# from spotoptim.utils.stats import get_sample_size

# delta = 1.0 * sigma_est # Detect difference of 1 sigma
# alpha = 0.05             # Significance level
# beta = 0.2               # Power = 0.8 implies beta = 0.2

# n_required = get_sample_size(alpha, beta, sigma_est, delta)
# n_validate = int(np.ceil(n_required))
# print(f"Required samples (n) for validation (delta={delta:.4f}, sigma={sigma_est:.4f})")
```

Hack: set `n_validate` to 10 for demonstration purposes:

```
n_validate = 10
```

87.5. Experimental Comparison

We perform m independent optimization runs for each algorithm. For each run, we take the *best design vector* found (x^*) and evaluate it n times to obtain a robust estimate of its true quality.

- Number of Optimization Runs (m): 5
- Budget per Run: 50 evaluations
- Validation Size (n): `n_validate` (calculated above)

```
import matplotlib.pyplot as plt
from scipy.optimize import minimize
from spotoptim import SpotOptim

# --- Configuration ---
M_RUNS = 2          # Number of independent optimization runs, here only 2 for spot
MAX_EVALS = 50      # Optimization budget
DIM = 10
BOUNDS = [(0.0, 1.0) for _ in range(DIM)]

# Storage for VALIDATED means
final_means_spot = []
final_means_scipy = []

# Storage for NOISY history (for convergence plots)
results_spot_hist = []
results_scipy_hist = []
```



```

def get_validated_mean(x_best, n_evals):
    """Evaluate x_best n_evals times and return the mean."""
    vals = []
    # Reshape for remote call if needed (SpotOptim returns 1D or 2D depending on version, handle both)
    X = np.atleast_2d(x_best)

    # Batch evaluation if objective supports it, or loop
    # robot_arm_hard supports batch, let's use it to save overhead
    # Create batch of n_evals identical rows
    X_batch = np.repeat(X, n_evals, axis=0)
    y_batch = robot_arm_hard(X_batch)
    return np.mean(y_batch)

print(f"Starting Benchmark: {M_RUNS} runs, opt budget {MAX_EVALS}, validation n={n_validate}...")

# --- 1. SpotOptim Loop ---
print("\n--- SpotOptim ---")
for i in range(M_RUNS):
    print(f"Run {i+1}/{M_RUNS}")

    optimizer = SpotOptim(
        fun=robot_arm_hard,
        bounds=BOUNDS,
        max_iter=MAX_EVALS,
        n_initial=10,
        acquisition='y',
        seed=i,
        verbose=False
    )

    try:
        res = optimizer.optimize()
        x_best = res.x

        # Collect History
        y_hist = res.y.flatten().tolist()
        results_spot_hist.append(y_hist)

        # Validation Step
        val_mean = get_validated_mean(x_best, n_validate)
        final_means_spot.append(val_mean)
        print(f" Best found: {res.fun:.4f} -> Validated Mean: {val_mean:.4f}")

    except Exception as e:

```

```

        print(f"    Error in SpotOptim run {i}: {e}")

# --- 2. Scipy L-BFGS-B Loop ---
print("\n--- Scipy L-BFGS-B ---")

# Wrapper for Scipy
class HistoryWrapper:
    def __init__(self, func):
        self.func = func
        self.best_y = np.inf
        self.best_x = None
        self.history = []
        self.count = 0

    def __call__(self, x):
        if self.count >= MAX_EVALS:
            return self.best_y # Stop

        X = x.reshape(1, -1)
        y = self.func(X)[0]

        self.history.append(y)

        if y < self.best_y:
            self.best_y = y
            self.best_x = x.copy()

        self.count += 1
        return y

for i in range(M_RUNS):
    print(f"Run {i+1}/{M_RUNS}")

    np.random.seed(i)
    x0 = np.random.rand(DIM)

    wrapper = HistoryWrapper(robot_arm_hard)

    try:
        minimize(
            wrapper,
            x0,
            method='L-BFGS-B',
            bounds=BOUNDS,

```

```

        options={'maxfun': MAX_EVALS}
    )

    # Collect History
    results_scipy_hist.append(wrapper.history)

    # Validation Step
    if wrapper.best_x is not None:
        val_mean = get_validated_mean(wrapper.best_x, n_validate)
        final_means_scipy.append(val_mean)
        print(f"  Best found: {wrapper.best_y:.4f} -> Validated Mean: {val_mean:.4f}")
    else:
        print("  No solutions evaluated.")

except Exception as e:
    print(f"  Error in Scipy run {i}: {e}")
    # Append worst case if failed? Or skip.
    # final_means_scipy.append(np.inf)

print("\nExperiments completed.")

```

Starting Benchmark: 2 runs, opt budget 50, validation n=10...

--- SpotOptim ---

Run 1/2

Best found: 17.5429 -> Validated Mean: 17.5429

Run 2/2

Best found: 26.9641 -> Validated Mean: 26.9641

--- Scipy L-BFGS-B ---

Run 1/2

Best found: 26.8058 -> Validated Mean: 26.8058

Run 2/2

Best found: 50.5134 -> Validated Mean: 50.5134

Experiments completed.

87.6. Convergence Analysis (Noisy)

We can visualize the progress of the optimizers by plotting the *best-so-far* objective value against the number of function evaluations. Note that because the objective

87. Remote Optimization Challenge: A Benchmarking Study

function is noisy, these curves may fluctuate or appear “better” than the true value due to lucky noise samples.

```
def get_best_so_far(history):
    """Convert a list of evaluations to a best-so-far trajectory."""
    bsf = []
    current_min = np.inf
    for val in history:
        if val < current_min:
            current_min = val
        bsf.append(current_min)
    return bsf

def process_histories(results_hist, max_evals):
    matrix = np.zeros((len(results_hist), max_evals))
    for i, hist in enumerate(results_hist):
        if not hist:
            bsf = [np.nan] * max_evals
        else:
            bsf = get_best_so_far(hist)
            if len(bsf) < max_evals:
                bsf += [bsf[-1]] * (max_evals - len(bsf))
            matrix[i, :] = bsf[:max_evals]
    return matrix

# Process data
spot_bsf_matrix = process_histories(results_spot_hist, MAX_EVALS)
scipy_bsf_matrix = process_histories(results_scipy_hist, MAX_EVALS)

# Calculate stats omitting NaNs
spot_mean = np.nanmean(spot_bsf_matrix, axis=0)
spot_std = np.nanstd(spot_bsf_matrix, axis=0)
scipy_mean = np.nanmean(scipy_bsf_matrix, axis=0)
scipy_std = np.nanstd(scipy_bsf_matrix, axis=0)

# Plot
plt.figure(figsize=(10, 6))
x_axis = np.arange(1, MAX_EVALS + 1)

# SpotOptim
plt.plot(x_axis, spot_mean, label='SpotOptim', color='blue', linewidth=2)
plt.fill_between(x_axis, spot_mean - spot_std, spot_mean + spot_std, color='blue', alpha=0.2)

# Scipy
plt.plot(x_axis, scipy_mean, label='Scipy L-BFGS-B', color='red', linewidth=2)
```

```
plt.fill_between(x_axis, scipy_mean - scipy_std, scipy_mean + scipy_std, color='red', alpha=0.2)

plt.xlabel('Function Evaluations')
plt.ylabel('Best Objective Value (Noisy)')
plt.title('Optimization Progress Comparison (Noisy)')
plt.legend()
plt.grid(True, linestyle='--', alpha=0.7)
plt.show()
```




Figure 87.2.: Convergence Plot: Mean Best-So-Far Objective Value (Noisy) vs. Evaluations

87.7. Comparative Analysis

We now statistically compare the validated performance of the two algorithms.

```
import scipy.stats as stats
plt.figure(figsize=(8, 5))
plt.boxplot([final_means_spot, final_means_scipy], tick_labels=['SpotOptim', 'Scipy L-BFGS-B'], patch_artist=True)
plt.ylabel('True Objective Value (Estimate)')
plt.title(f'Comparison over {M_RUNS} Runs (Validation n={n_validate})')
```

87. Remote Optimization Challenge: A Benchmarking Study

```
plt.grid(True, axis='y', linestyle='--', alpha=0.7)
plt.show()
```




Figure 87.3.: Comparison of Validated True Means (Lower is Better)

t-test:

```
# 2. T-Test
t_stat, p_val = stats.ttest_ind(final_means_spot, final_means_scipy, equal_var=False)

print("Statistical Comparison (Welch's t-test):")
print(f" SpotOptim Mean: {np.mean(final_means_spot):.4f} (std: {np.std(final_means_spot):.4f})")
print(f" Scipy Mean:      {np.mean(final_means_scipy):.4f} (std: {np.std(final_means_scipy):.4f})")
print(f" T-statistic: {t_stat:.4f}")
print(f" P-value: {p_val:.4f}")

if p_val < 0.05:
    print(" Result: Significant difference detected (p < 0.05).")
else:
    print(" Result: No significant difference detected (p >= 0.05).")
```

Statistical Comparison (Welch's t-test):

```

SpotOptim Mean: 22.2535 (std: 4.7106)
Scipy Mean:      38.6596 (std: 11.8538)
T-statistic: -1.2862
P-value: 0.3803
Result: No significant difference detected (p >= 0.05).

```

87.8. Conclusion

This study followed a rigorous benchmarking protocol: 1. Noise Estimation: Objective function noise was characterized ($\sigma \approx \text{sigma_est}$). 2. Power Analysis: Sample size n was determined to ensure statistical power. 3. Robust Evaluation: Algorithms were compared based on the *validated mean* performance of their final solutions, not single noisy evaluations.

The results indicate that [SpotOptim/Scipy] provides superior performance on this remote engineering problem given the fixed budget. This methodology ensures that the conclusions are not artifacts of random noise.

87.9. Landscape Analysis of robot_arm_hard

To gain further insight into the difficulty of the `robot_arm_hard` problem, we visualize its landscape using 2D slices. Similar to the AWWWE example, we vary pairs of variables while keeping the others fixed at the center of the domain (0.5). This reveals the complex constraint boundaries and “valleys” created by the obstacles.

```

import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits.axes_grid1 import ImageGrid
from spotoptim.function.so import robot_arm_hard

# 1. Setup Grid
N_RES = 40
x = np.linspace(0, 1, N_RES)
X_mesh, Y_mesh = np.meshgrid(x, x)

# 2. Compute Slices for all 45 pairs
n_dim = 10
default_val = 0.5
Z_list = []
Z_labels = []

# Loop over pairs (i, j)

```

```

for i in range(n_dim):
    for j in range(i + 1, n_dim):
        # Create batch input
        inputs = np.full((N_RES * N_RES, n_dim), default_val)
        inputs[:, i] = X_mesh.ravel()
        inputs[:, j] = Y_mesh.ravel()

        # Evaluate
        z = robot_arm_hard(inputs).reshape(X_mesh.shape)
        Z_list.append(z)
        Z_labels.append((f"x{i+1}", f"x{j+1}"))

# Only keep the last 6 pairs
Z_list = Z_list[-6:]
Z_labels = Z_labels[-6:]

# 3. Plotting
fig = plt.figure(figsize=(12., 8.))
grid = ImageGrid(fig, 111,
                  nrows_ncols=(2, 3), # Last 6 pairs
                  axes_pad=0.3,
                  share_all=True,
                  cbar_location="right",
                  cbar_mode="single",
                  cbar_size="5%",
                  cbar_pad=0.1,
                  label_mode="all"
                  )

# Set color limits to highlight the 'valleys' vs 'walls'
# Unconstrained values are ~0-200. Obstacle penalties are >1000.
# We saturate at 500 to show the structure of the feasible space.
vmin, vmax = 225, 280

counter = 0
for ax in grid:
    if counter < len(Z_list):
        ax.set_xticks([])
        ax.set_yticks([])

        im = ax.contourf(X_mesh, Y_mesh, Z_list[counter], 150,
                         cmap="jet", vmin=vmin, vmax=vmax)
    else:
        ax.set_visible(False)

```



```

    counter += 1

# Add Colorbar
cax = grid.cbar_axes[0]
cax.colorbar(im)

plt.suptitle("Robot Arm Hard: Selected Pairwise Landscape Slices", y=0.95, fontsize=16)
plt.show()

```

Robot Arm Hard: Selected Pairwise Landscape Slices




Figure 87.4.: Robot Arm Hard Landscape (2D Slices). Blue regions are low cost (feasible/close to target), red regions are high cost (obstacles or far from target).

87.10. References

- Bartz-Beielstein, T., et al. (2020). Benchmarking in Optimization: Best Practice and Open Issues. arXiv:2007.03488.

87.11. Jupyter Notebook

Note

- The Jupyter-Notebook of this chapter is available on GitHub in the Sequential Parameter Optimization Cookbook Repository

Part XIV.

Data-Driven Modeling and Optimization

88. Basic Statistics and Data Analysis

This chapter covers basic statistical concepts, namely descriptive statistics, probability distributions, and hypothesis testing. These concepts are fundamental to understanding data and making informed decisions based on data analysis. The chapter also introduces the concept of exploratory data analysis (EDA), data preprocessing (Principal Component Analysis), and data visualization techniques.

88.1. Exploratory Data Analysis

88.1.1. Histograms

Creating a histogram and calculating the probabilities from a dataset can be approached with scientific precision

1. **Data Collection:** Obtain the dataset you wish to analyze. This dataset could represent any quantitative measure, such to examine its distribution.
2. **Decide on the Number of Bins:** The number of bins influences the histogram's granularity. There are several statistical rules to determine an optimal number of bins:
 - **Square-root rule:** suggests using the square root of the number of data points as the number of bins.
 - **Sturges' formula:** $k = 1 + 3.322 \log_{10}(n)$, where n is the number of data points and k is the suggested number of bins.
 - **Freedman-Diaconis rule:** uses the interquartile range (IQR) and the cube root of the number of data points n to calculate bin width as $2 \frac{IQR}{n^{1/3}}$.
3. **Determine Range and Bin Width:** Calculate the range of data by subtracting the minimum data point value from the maximum. Divide this range by the number of bins to determine the width of each bin.
4. **Allocate Data Points to Bins:** Iterate through the data, sorting each data point into the appropriate bin based on its value.
5. **Draw the Histogram:** Use a histogram to visualize the frequency or relative frequency (probability) of data points within each bin.

6. Calculate Probabilities: The relative frequency of data within each bin represents the probability of a randomly selected data point falling within that bin's range.

Below is a Python script that demonstrates how to generate a histogram and compute probabilities using the `matplotlib` library for visualization and `numpy` for data manipulation.

```
import numpy as np
import matplotlib.pyplot as plt

# Sample data: Randomly generated for demonstration
data = np.random.normal(0, 1, 1000) # 1000 data points with a normal distribution

# Step 2: Decide on the number of bins
num_bins = int(np.ceil(1 + 3.322 * np.log10(len(data)))) # Sturges' formula

# Step 3: Determine range and bin width -- handled internally by matplotlib


# Steps 4 & 5: Sort data into bins and draw the histogram
fig, ax = plt.subplots()
n, bins, patches = ax.hist(data, bins=num_bins, density=True, alpha=0.75, edgecolor='k')

# Calculate probabilities (relative frequencies) manually, if needed
bin_width = np.diff(bins) # np.diff finds the difference between adjacent bin boundaries
probabilities = n * bin_width # n is already normalized to form a probability density

# Adding labels and title for clarity
ax.set_xlabel('Data Value')
ax.set_ylabel('Probability Density')
ax.set_title('Histogram with Probability Density')
```

```
Text(0.5, 1.0, 'Histogram with Probability Density')
```

(a) Histogram with Probability Density



(b)

Figure 88.1.

```
for i, prob in enumerate(probabilities):
    print(f"Bin {i+1} Probability: {prob:.4f}")

# Ensure probabilities sum to 1 (or very close, due to floating-point arithmetic)
print(f"Sum of probabilities: {np.sum(probabilities)}")
```

```
Bin 1 Probability: 0.0030
Bin 2 Probability: 0.0210
Bin 3 Probability: 0.0420
Bin 4 Probability: 0.1130
Bin 5 Probability: 0.1830
Bin 6 Probability: 0.2280
Bin 7 Probability: 0.1780
Bin 8 Probability: 0.1400
Bin 9 Probability: 0.0600
```

88. *Basic Statistics and Data Analysis*

```
Bin 10 Probability: 0.0260
Bin 11 Probability: 0.0060
Sum of probabilities: 1.0000000000000002
```

This code segment goes through the necessary steps to generate a histogram and calculate probabilities for a synthetic dataset. It demonstrates important scientific and computational practices including binning, visualization, and probability calculation in Python.

Key Points:

- The histogram represents the distribution of data, with the histogram's bins outlining the data's spread and density.
- The option `density=True` in `ax.hist()` normalizes the histogram so that the total area under the histogram sums to 1, thereby converting frequencies to probability densities.
- The choice of bin number and width has a significant influence on the histogram's shape and the insights that can be drawn from it, highlighting the importance of selecting appropriate binning strategies based on the dataset's characteristics and the analysis objectives.
- Video: [Histograms, Clearly Explained](#)

88.1.2. Boxplots

- Video: [Boxplots are Awesome](#)

88.2. Probability Distributions

What happens when we use smaller bins in a histogram? The histogram becomes more detailed, revealing the distribution of data points with greater precision. However, as the bin size decreases, the number of data points within each bin may decrease, leading to sparse or empty bins. This sparsity can make it challenging to estimate probabilities accurately, especially for data points that fall within these empty bins.

Advantages, when using a probability distribution, include:

- Blanks can be filled
- Probabilities can be calculated
- Parameters are sufficient to describe the distribution, e.g., mean and variance for the normal distribution

Probability distributions offer a powerful solution to the challenges posed by limited data in estimating probabilities. When data is scarce, constructing a histogram to determine the probability of certain outcomes can lead to inaccurate or unreliable results due to the lack of detail in the dataset. However, collecting vast amounts of data to populate a histogram for more precise estimates can often be impractical, time-consuming, and expensive.

A probability distribution is a mathematical function that provides the probabilities of occurrence of different possible outcomes for an experiment. It is a more efficient approach to understanding the likelihood of various outcomes than relying solely on extensive data collection. For continuous data, this is often represented graphically by a smooth curve.

- Video: The Main Ideas behind Probability Distributions

88.2.1. Sampling from a Distribution

- Video: Sampling from a Distribution, Clearly Explained!!!

88.3. Discrete Distributions

Discrete probability distributions are essential tools in statistics, providing a mathematical foundation to model and analyze situations with discrete outcomes. Histograms, which can be seen as discrete distributions with data organized into bins, offer a way to visualize and estimate probabilities based on the collected data. However, they come with limitations, especially when data is scarce or when we encounter gaps in the data (blank spaces in histograms). These gaps can make it challenging to accurately estimate probabilities.

A more efficient approach, especially for discrete data, is to use mathematical equations—particularly those defining discrete probability distributions—to calculate probabilities directly, thus bypassing the intricacies of data collection and histogram interpretation.

88.3.1. Bernoulli Distribution

The Bernoulli distribution, named after Swiss scientist Jacob Bernoulli, is a discrete probability distribution, which takes value 1 with success probability p and value 0 with failure probability $q = 1 - p$. So if X is a random variable with this distribution, we have:

$$P(X = 1) = 1 - P(X = 0) = p = 1 - q.$$

88.3.2. Binomial Distribution

The Binomial Distribution is a prime example of a discrete probability distribution that is particularly useful for binary outcomes (e.g., success/failure, yes/no, pumpkin pie/blueberry pie). It leverages simple mathematical principles to calculate the probability of observing a specific number of successes (preferred outcomes) in a fixed number of trials, given the probability of success in each trial.

Example 88.1 (Pie Preference). Consider a scenario from “StatLand” where 70% of people prefer pumpkin pie over blueberry pie. The question is: What is the probability that, out of three people asked, the first two prefer pumpkin pie and the third prefers blueberry pie?

Using the concept of the Binomial Distribution, the probability of such an outcome can be calculated without the need to layout every possible combination by hand. This process not only simplifies calculations but also provides a clear and precise method to determine probabilities in scenarios involving discrete choices. We will use Python to calculate the probability of observing exactly two out of three people prefer pumpkin pie, given the 70% preference rate:

```
from scipy.stats import binom
n = 3 # Number of trials (people asked)
p = 0.7 # Probability of success (preferring pumpkin pie)
x = 2 # Number of successes (people preferring pumpkin pie)
# Probability calculation using Binomial Distribution
prob = binom.pmf(x, n, p)
print(f"The probability that exactly 2 out of 3 people prefer pumpkin pie is: {prob:.3f}")
```

The probability that exactly 2 out of 3 people prefer pumpkin pie is: 0.441

This code uses the `binom.pmf()` function from `scipy.stats` to calculate the probability mass function (PMF) of observing exactly `x` successes in `n` trials, where each trial has a success probability of `p`.

A Binomial random variable is the sum of n independent, identically distributed Bernoulli random variables, each with probability p of success. We may indicate a random variable X with Bernoulli distribution using the notation $X \sim \text{Bi}(1, \theta)$. Then, the notation for the Binomial is $X \sim \text{Bi}(n, \theta)$. Its probability and distribution functions are, respectively,

$$p_X(x) = \binom{n}{x} \theta^x (1 - \theta)^{n-x}, \quad F_X(x) = \Pr\{X \leq x\} = \sum_{i=0}^x \binom{n}{i} \theta^i (1 - \theta)^{n-i}.$$

The mean of the binomial distribution is $E[X] = n\theta$. The variance of the distribution is $\text{Var}[X] = n\theta(1 - \theta)$ (see next section).

A process consists of a sequence of n independent trials, i.e., the outcome of each trial does not depend on the outcome of previous trials. The outcome of each trial is either a success or a failure. The probability of success is denoted as p , and p is constant for each trial. Coin tossing is a classical example for this setting.

The binomial distribution is a statistical distribution giving the probability of obtaining a specified number of successes in a binomial experiment; written $\text{Binomial}(n, p)$, where n is the number of trials, and p the probability of success in each.

Definition 88.1 (Binomial Distribution). The binomial distribution with parameters n and p , where n is the number of trials, and p the probability of success in each, is

$$p(x) = \binom{n}{k} p^x (1 - p)^{n-x} \quad x = 0, 1, \dots, n. \quad (88.1)$$

The mean μ and the variance σ^2 of the binomial distribution are

$$\mu = np \quad (88.2)$$

and

$$\sigma^2 = np(1 - p). \quad (88.3)$$

Note, the Bernoulli distribution is simply $\text{Binomial}(1, p)$.

88.4. Continuous Distributions

Our considerations regarding probability distributions, expectations, and standard deviations will be extended from discrete distributions to continuous distributions. One simple example of a continuous distribution is the uniform distribution. Continuous distributions are defined by probability density functions.

88.4.1. Distribution functions: PDFs and CDFs

The density for a continuous distribution is a measure of the relative probability of “getting a value close to x .” Probability density functions f and cumulative distribution function F are related as follows.

$$f(x) = \frac{d}{dx} F(x) \quad (88.4)$$

88.5. Background: Expectation, Mean, Standard Deviation

The distribution of a random vector is characterized by some indexes. These are the expectation, the mean, and the standard deviation. The expectation is a measure of the central tendency of a random variable, while the standard deviation quantifies the spread of the distribution. These indexes are essential for understanding the behavior of random variables and making predictions based on them.

Definition 88.2 (Random Variable). A random variable X is a mapping from the sample space of a random experiment to the real numbers. It assigns a numerical value to each outcome of the experiment. Random variables can be either:

- Discrete: If X takes on a countable number of distinct values.
- Continuous: If X takes on an uncountable number of values.

Mathematically, a random variable is a function $X : \Omega \rightarrow \mathbb{R}$, where Ω is the sample space.

Definition 88.3 (Probability Distribution). A probability distribution describes how the values of a random variable are distributed. It is characterized for a discrete random variable X by the probability mass function (PMF) $p_X(x)$ and for a continuous random variable X by the probability density function (PDF) $f_X(x)$.

Definition 88.4 (Probability Mass Function (PMF)). $p_X(x) = P(X = x)$ gives the probability that X takes the value x .

Definition 88.5 (Probability Density Function (PDF)). $f_X(x)$ is a function such that for any interval $[a, b]$, the probability that X falls within this interval is given by the integral $\int_a^b f_X(x)dx$.

The distribution function must satisfy:

$$\sum_{x \in D_X} p_X(x) = 1$$

for discrete random variables, where D_X is the domain of X and

$$\int_{-\infty}^{\infty} f_X(x)dx = 1$$

for continuous random variables.

With these definitions in place, we can now introduce the definition of the expectation, which is a fundamental measure of the central tendency of a random variable.

88.5. Background: Expectation, Mean, Standard Deviation

Definition 88.6 (Expectation). The expectation or expected value of a random variable X , denoted $E[X]$, is defined as follows:

For a discrete random variable X :

$$E[X] = \sum_{x \in D_X} xp_X(x) \quad \text{if } X \text{ is discrete.}$$

For a continuous random variable X :

$$E[X] = \int_{x \in D_X} xf_X(x)dx \quad \text{if } X \text{ is continuous.}$$

The mean, μ , of a probability distribution is a measure of its central tendency or location. That is, $E(X)$ is defined as the average of all possible values of X , weighted by their probabilities.

Example 88.2 (Expectation). Let X denote the number produced by rolling a fair die. Then

$$E(X) = 1 \times 1/6 + 2 \times 1/6 + 3 \times 1/6 + 4 \times 1/6 + 5 \times 1/6 + 6 \times 1/6 = 3.5.$$

Definition 88.7 (Sample Mean). The sample mean is an important estimate of the population mean. The sample mean of a sample $\{x_i\}$ ($i = 1, 2, \dots, n$) is defined as

$$\bar{x} = \frac{1}{n} \sum_i x_i.$$

While both the expectation of a random variable and the sample mean provide measures of central tendency, they differ in their context, calculation, and interpretation.

- The expectation is a theoretical measure that characterizes the average value of a random variable over an infinite number of repetitions of an experiment. The expectation is calculated using a probability distribution and provides a parameter of the entire population or distribution. It reflects the long-term average or central value of the outcomes generated by the random process.
- The sample mean is a statistic. It provides an estimate of the population mean based on a finite sample of data. It is computed directly from the data sample, and its value can vary between different samples from the same population. It serves as an approximation or estimate of the population mean. It is used in statistical inference to make conclusions about the population mean based on sample data.

If we are trying to predict the value of a random variable X by its mean $\mu = E(X)$, the error will be $X - \mu$. In many situations it is useful to have an idea how large this deviation or error is. Since $E(X - \mu) = E(X) - \mu = 0$, it is necessary to use the absolute value or the square of $(X - \mu)$. The squared error is the first choice, because the derivatives are easier to calculate. These considerations motivate the definition of the variance:

Definition 88.8 (Variance). The variance of a random variable X is the mean squared deviation of X from its expected value $\mu = E(X)$.

$$\text{Var}(X) = E[(X - \mu)^2]. \quad (88.5)$$

The variance is a measure of the spread of a distribution. It quantifies how much the values of a random variable differ from the mean. A high variance indicates that the values are spread out over a wide range, while a low variance indicates that the values are clustered closely around the mean.

Definition 88.9 (Standard Deviation). Taking the square root of the variance to get back to the same scale of units as X gives the standard deviation. The standard deviation of X is the square root of the variance of X .

$$\text{sd}(X) = \sqrt{\text{Var}(X)}. \quad (88.6)$$

88.6. Distributions and Random Numbers in Python

Results from computers are deterministic, so it sounds like a contradiction in terms to generate random numbers on a computer. Standard computers generate pseudorandom numbers, i.e., numbers that behave as if they were drawn randomly.

Deterministic Random Numbers

- Idea: Generate deterministically numbers that **look** (behave) as if they were drawn randomly.

88.6.1. Calculation of the Standard Deviation with Python

The function `numpy.std` returns the standard deviation, a measure of the spread of a distribution, of the array elements. The argument `ddof` specifies the Delta Degrees of Freedom. The divisor used in calculations is $N - \text{ddof}$, where N represents the number of elements. By default `ddof` is zero, i.e., `std` uses the formula

$$\sqrt{\frac{1}{N} \sum_i (x_i - \bar{x})^2} \quad \text{with} \quad \bar{x} = \sum_{i=1}^N x_i / N.$$

Example 88.3 (Standard Deviation with Python). Consider the array $[1, 2, 3]$: Since $\bar{x} = 2$, the following value is computed:

$$\sqrt{1/3 \times ((1-2)^2 + (2-2)^2 + (3-2)^2)} = \sqrt{2/3}.$$

```
import numpy as np
a = np.array([[1, 2, 3]])
np.std(a)
```

```
np.float64(0.816496580927726)
```

The empirical standard deviation (which uses $N-1$), $\sqrt{1/2 \times ((1-2)^2 + (2-2)^2 + (3-2)^2)} = \sqrt{2/2}$, can be calculated in Python as follows:

```
np.std(a, ddof=1)
```

```
np.float64(1.0)
```

88.6.2. The Argument “axis”

When you compute `np.std` with `axis=0`, it calculates the standard deviation along the vertical axis, meaning it computes the standard deviation for each column of the array. On the other hand, when you compute `np.std` with `axis=1`, it calculates the standard deviation along the horizontal axis, meaning it computes the standard deviation for each row of the array. If the axis parameter is not specified, `np.std` computes the standard deviation of the flattened array, i.e., it calculates the standard deviation of all the elements in the array.

Example 88.4 (Axes along which the standard deviation is computed).

```
A = np.array([[1, 2], [3, 4]])
A
```

```
array([[1, 2],
       [3, 4]])
```

First, we calculate the standard deviation of all elements in the array:

```
np.std(A)
```

```
np.float64(1.118033988749895)
```

88. Basic Statistics and Data Analysis

Setting `axis=0` calculates the standard deviation along the vertical axis (column-wise):

```
np.std(A, axis=0)
```

```
array([1., 1.])
```

Finally, setting `axis=1` calculates the standard deviation along the horizontal axis (row-wise):

```
np.std(A, axis=1)
```

```
array([0.5, 0.5])
```

88.7. Expectation (Continuous)

Definition 88.10 (Expectation (Continuous)).

$$E(X) = \int_{-\infty}^{\infty} x f(x) dx \quad (88.7)$$

88.8. Variance and Standard Deviation (Continuous)

Definition 88.11 (Variance (Continuous)). Variance can be calculated with $E(X)$ and

$$E(X^2) = \int_{-\infty}^{\infty} x^2 f(x) dx \quad (88.8)$$

as

$$\text{Var}(X) = E(X^2) - [E(X)]^2.$$

□

Definition 88.12 (Standard Deviation (Continuous)). Standard deviation can be calculated as

$$\text{sd}(X) = \sqrt{\text{Var}(X)}.$$

□

- Video: Population and Estimated Parameters, Clearly Explained
- Video: Calculating the Mean, Variance, and Standard Deviation

88.8.1. The Uniform Distribution

Definition 88.13 (The Uniform Distribution). The probability density function of the uniform distribution is defined as:

$$f_X(x) = \frac{1}{b-a} \quad \text{for } x \in [a, b].$$

Generate 10 random numbers from a uniform distribution between $a = 0$ and $b = 1$:

```
import numpy as np
# Initialize the random number generator
rng = np.random.default_rng(seed=123456789)
n = 10
x = rng.uniform(low=0.0, high=1.0, size=n)
x
```

```
array([0.02771274, 0.90670006, 0.88139355, 0.62489728, 0.79071481,
       0.82590801, 0.84170584, 0.47172795, 0.95722878, 0.94659153])
```


Generate 10,000 random numbers from a uniform distribution between 0 and 10 and plot a histogram of the numbers:

```
import numpy as np
import matplotlib.pyplot as plt

# Initialize the random number generator
rng = np.random.default_rng(seed=123456789)

# Generate random numbers from a uniform distribution
x = rng.uniform(low=0, high=10, size=10000)

# Plot a histogram of the numbers
plt.hist(x, bins=50, density=True, edgecolor='black')
plt.title('Uniform Distribution [0,10]')
plt.xlabel('Value')
plt.ylabel('Frequency')
plt.show()
```



88.9. The Uniform Distribution

This variable is defined in the interval $[a, b]$. We write it as $X \sim U[a, b]$. Its density and cumulative distribution functions are, respectively,

$$f_X(x) = \frac{I_{[a,b]}(x)}{b-a}, \quad F_X(x) = \frac{1}{b-a} \int_{-\infty}^x I_{[a,b]}(t) dt = \frac{x-a}{b-a},$$

where $I_{[a,b]}(\cdot)$ is the indicator function of the interval $[a, b]$. Note that, if we set $a = 0$ and $b = 1$, we obtain $F_X(x) = x$, $x \in [0, 1]$.

A typical example is the following: the cdf of a continuous r.v. is uniformly distributed in $[0, 1]$. The proof of this statement is as follows: For $u \in [0, 1]$, we have

$$\begin{aligned} \Pr\{F_X(X) \leq u\} &= \Pr\{F_X^{-1}(F_X(X)) \leq F_X^{-1}(u)\} = \Pr\{X \leq F_X^{-1}(u)\} \\ &= F_X(F_X^{-1}(u)) = u. \end{aligned}$$

This means that, when X is continuous, there is a one-to-one relationship (given by the cdf) between $x \in D_X$ and $u \in [0, 1]$.

The *uniform distribution* has a constant density over a specified interval, say $[a, b]$. The uniform $U(a, b)$ distribution has density

$$f(x) = \begin{cases} 1/(b-a) & \text{if } a < x < b, \\ 0 & \text{otherwise} \end{cases} \quad (88.9)$$

88.9.1. The Normal Distribution

A normally distributed random variable is a random variable whose associated probability distribution is the normal (or Gaussian) distribution. The normal distribution is a continuous probability distribution characterized by a symmetric bell-shaped curve.

The distribution is defined by two parameters: the mean μ and the standard deviation σ . The mean indicates the center of the distribution, while the standard deviation measures the spread or dispersion of the distribution.

This distribution is widely used in statistics and the natural and social sciences as a simple model for random variables with unknown distributions.

Definition 88.14 (The Normal Distribution). The probability density function of the normal distribution is defined as:

$$f_X(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2\right), \quad (88.10)$$

where: μ is the mean; σ is the standard deviation.

To generate ten random numbers from a normal distribution, the following command can be used.

```
import numpy as np
rng = np.random.default_rng()
n = 10
mu, sigma = 2, 0.1
x = rng.normal(mu, sigma, n)
x
```

```
array([1.76202311, 1.95481236, 1.97375374, 2.0532046 , 2.01146411,
       2.09286743, 1.86278569, 1.95777461, 1.85435487, 1.83424859])
```

Verify the mean:

```
abs(mu - np.mean(x))
```


```
np.float64(0.06427108811699078)
```

Note: To verify the standard deviation, we use `ddof = 1` (empirical standard deviation):

```
abs(sigma - np.std(x, ddof=1))
```

```
np.float64(0.004848758841445949)
```

```
plot_normal_distribution(mu=0, sigma=1, num_samples=10000)
```




88.9.2. Visualization of the Standard Deviation

The standard deviation of normal distributed can be visualized in terms of the histogram of X :

- about 68% of the values will lie in the interval within one standard deviation of the mean
- 95% lie within two standard deviation of the mean

- and 99.9% lie within 3 standard deviations of the mean.



88.9.3. Realizations of a Normal Distribution

Realizations of a normal distribution refers to the actual values that you get when you draw samples from a normal distribution. Each sample drawn from the distribution is a realization of that distribution.

Example 88.5 (Realizations of a Normal Distribution). If you have a normal distribution with a mean of 0 and a standard deviation of 1, each number you draw from that distribution is a realization. Here is a Python example that generates 10 realizations of a normal distribution with a mean of 0 and a standard deviation of 1:

```
import numpy as np
mu = 0
sigma = 1
realizations = np.random.normal(mu, sigma, 10)
print(realizations)
```

```
[ 0.48951662  0.23879586 -0.44811181 -0.610795  -2.02994507  0.60794659
 -0.35410888  0.15258149  0.50127485 -0.78640277]
```

In this code, `np.random.normal` generates ten realizations of a normal distribution with a mean of 0 and a standard deviation of 1. The `realizations` array contains the actual values drawn from the distribution.

88.10. The Normal Distribution

A commonly encountered probability distribution is the normal distribution, known for its characteristic bell-shaped curve. This curve represents how the values of a variable are distributed: most of the observations cluster around the mean (or center) of the distribution, with frequencies gradually decreasing as values move away from the mean.

The normal distribution is particularly useful because of its defined mathematical properties. It is determined entirely by its mean (μ) and its standard deviation (σ). The area under the curve represents probability, making it possible to calculate the likelihood of a random variable falling within a specific range.

- Video: The Normal Distribution, Clearly Explained!!!

Example 88.6 (Normal Distribution: Estimating Probabilities). Consider we are interested in the heights of adults in a population. Instead of measuring the height of every adult (which would be impractical), we can use the normal distribution to estimate the probability of adults' heights falling within certain intervals, assuming we know the mean and standard deviation of the heights.

```
import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import norm
mu = 170 # e.g., mu height of adults in cm
sd = 10 # e.g., standard deviation of heights in cm
heights = np.linspace(mu - 3*sd, mu + 3*sd, 1000)
# Calculate the probability density function for the normal distribution
pdf = norm.pdf(heights, mu, sd)
# Plot the normal distribution curve
plt.plot(heights, pdf, color='blue', linewidth=2)
plt.fill_between(heights, pdf, where=(heights >= mu - 2 * sd) & (heights <= mu + 2*sd))
plt.xlabel('Height (cm)')
plt.ylabel('Probability Density')
plt.show()
```




Figure 88.2.: Normal Distribution Curve with Highlighted Probability Area. 95 percent of the data falls within two standard deviations of the mean.

This Python code snippet generates a plot of the normal distribution for adult heights, with a mean of 170 cm and a standard deviation of 10 cm. It visually approximates a histogram with a blue bell-shaped curve, and highlights (in grey) the area under the curve between $\mu \pm 2 \times \sigma$. This area corresponds to the probability of randomly selecting an individual whose height falls within this range.

By using the area under the curve, we can efficiently estimate probabilities without needing to collect and analyze a vast amount of data. This method not only saves time and resources but also provides a clear and intuitive way to understand and communicate statistical probabilities.

Definition 88.15 (Normal Distribution). This variable is defined on the support $D_X = \mathbb{R}$ and its density function is given by

$$f_X(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{1}{2\sigma^2} (x - \mu)^2 \right\}.$$

The density function is identified by the pair of parameters (μ, σ^2) , where $\mu \in \mathbb{R}$ is the mean (or location parameter) and $\sigma^2 > 0$ is the variance (or dispersion parameter) of X . $\square$

The density function is symmetric around μ . The normal distribution belongs to the location-scale family distributions. This means that, if $Z \sim N(0, 1)$ (read, Z has a standard normal distribution; i.e., with $\mu = 0$ and $\sigma^2 = 1$), and we consider the linear transformation $X = \mu + \sigma Z$, then $X \sim N(\mu, \sigma^2)$ (read, X has a normal distribution with mean μ and variance σ^2). This means that one can obtain the probability of any interval $(-\infty, x]$, $x \in \mathbb{R}$ for any normal distribution (i.e., for any pair of the parameters μ and σ) once the quantiles of the standard normal distribution are known. Indeed

$$\begin{aligned} F_X(x) &= \Pr\{X \leq x\} = \Pr\left\{\frac{X - \mu}{\sigma} \leq \frac{x - \mu}{\sigma}\right\} \\ &= \Pr\left\{Z \leq \frac{x - \mu}{\sigma}\right\} = F_Z\left(\frac{x - \mu}{\sigma}\right) \quad x \in \mathbb{R}. \end{aligned}$$

The quantiles of the standard normal distribution are available in any statistical program. The density and cumulative distribution function of the standard normal r.v.~at point x are usually denoted by the symbols $\phi(x)$ and $\Phi(x)$.

The standard normal distribution is based on the *standard normal density function*

$$\varphi(z) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{z^2}{2}\right). \quad (88.11)$$

An important application of the standardization introduced in Equation 88.11 reads as follows. In case the distribution of X is approximately normal, the distribution of $X^* = \frac{X - \mu}{\sigma}$ is approximately standard normal. That is

$$P(X \leq b) = P\left(\frac{X - \mu}{\sigma} \leq \frac{b - \mu}{\sigma}\right) = P(X^* \leq \frac{b - \mu}{\sigma})$$

The probability $P(X \leq b)$ can be approximated by $\Phi(\frac{b - \mu}{\sigma})$, where Φ is the standard normal cumulative distribution function.

If X is a normal random variable with mean μ and variance σ^2 , i.e., $X \sim N(\mu, \sigma^2)$, then

$$X = \mu + \sigma Z \text{ where } Z \sim N(0, 1). \quad (88.12)$$

If $Z \sim N(0, 1)$ and $X \sim N(\mu, \sigma^2)$, then

$$X = \mu + \sigma Z.$$

The probability of getting a value in a particular interval is the area under the corresponding part of the curve. Consider the density function of the normal distribution. It can be plotted using the following commands. The result is shown in Figure 88.3.


```

import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import norm
x = np.arange(-4, 4, 0.1)
# Calculating the normal distribution's density function values for each point in x
y = norm.pdf(x, 0, 1)
plt.plot(x, y, linestyle='-', linewidth=2)
plt.title('Normal Distribution')
plt.xlabel('X')
plt.ylabel('Density')
plt.grid(True)
plt.show()

```




Figure 88.3.: Normal Distribution Density Function

The *cumulative distribution function* (CDF) describes the probability of “hitting” x or less in a given distribution. We consider the CDF function of the normal distribution. It can be plotted using the following commands. The result is shown in Figure 88.4.

```

import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import norm

# Generating a sequence of numbers from -4 to 4 with 0.1 intervals
x = np.arange(-4, 4, 0.1)

# Calculating the cumulative distribution function value of the normal distribution f
y = norm.cdf(x, 0, 1) # mean=0, stddev=1

# Plotting the results. The equivalent of 'type="l"' in R (line plot) becomes the defa
plt.plot(x, y, linestyle='-', linewidth=2)
plt.title('Normal Distribution CDF')
plt.xlabel('X')
plt.ylabel('Cumulative Probability')
plt.grid(True)
plt.show()

```




Figure 88.4.: Normal Distribution Cumulative Distribution Function

88.11. The Exponential Distribution

The exponential distribution is a continuous probability distribution that describes the time between events in a Poisson process, where events occur continuously and independently at a constant average rate. It is characterized by a single parameter, the rate parameter λ , which represents the average number of events per unit time.

88.11.1. Standardization of Random Variables

To compare statistical properties of random variables which use different units, it is a common practice to transform these random variables into standardized variables.

Definition 88.16 (Standard Units). If a random variable X has expectation $E(X) = \mu$ and standard deviation $sd(X) = \sigma > 0$, the random variable

$$X^* = (X - \mu)/\sigma$$

is called X in standard units. It has $E(X^*) = 0$ and $sd(X^*) = 1$.

88.12. Covariance and Correlation

88.12.1. The Multivariate Normal Distribution

The multivariate normal, multinormal, or Gaussian distribution serves as a generalization of the one-dimensional normal distribution to higher dimensions. We will consider n -dimensional random vectors $X = (X_1, X_2, \dots, X_n)$. When drawing samples from this distribution, it results in a set of values represented as $\{x_1, x_2, \dots, x_n\}$. To fully define this distribution, it is necessary to specify its mean μ and covariance matrix Σ . These parameters are analogous to the mean, which represents the central location, and the variance (squared standard deviation) of the one-dimensional normal distribution introduced in Equation 88.10.

Definition 88.17 (The Multivariate Normal Distribution). The probability density function (PDF) of the multivariate normal distribution is defined as:

$$f_X(x) = \frac{1}{\sqrt{(2\pi)^n \det(\Sigma)}} \exp\left(-\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu)\right),$$

where: μ is the $n \times 1$ mean vector; Σ is the $n \times n$ covariance matrix. The covariance matrix Σ is assumed to be positive definite, so that its determinant is strictly positive.

In the context of the multivariate normal distribution, the mean takes the form of a coordinate within an k -dimensional space. This coordinate represents the location where samples are most likely to be generated, akin to the peak of the bell curve in a one-dimensional or univariate normal distribution.

Definition 88.18 (Covariance of two random variables). For two random variables X and Y , the covariance is defined as the expected value (or mean) of the product of their deviations from their individual expected values:

$$\text{cov}(X, Y) = E[(X - E[X])(Y - E[Y])]$$

For discrete random variables, covariance can be written as:

$$\text{cov}(X, Y) = \frac{1}{n} \sum_{i=1}^n (x_i - E(X))(y_i - E(Y)).$$

The covariance within the multivariate normal distribution denotes the extent to which two variables vary together. The elements of the covariance matrix, such as Σ_{ij} , represent the covariances between the variables x_i and x_j . These covariances describe how the different variables in the distribution are related to each other in terms of their variability.

Example 88.7 (The Bivariate Normal Distribution with Positive Covariances).

Figure 88.5 shows draws from a bivariate normal distribution with $\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$ and $\Sigma = \begin{pmatrix} 9 & 4 \\ 4 & 9 \end{pmatrix}$.

The covariance matrix of a bivariate normal distribution determines the shape, orientation, and spread of the distribution in the two-dimensional space.

The diagonal elements of the covariance matrix (σ_1^2, σ_2^2) are the variances of the individual variables. They determine the spread of the distribution along each axis. A larger variance corresponds to a greater spread along that axis.

The off-diagonal elements of the covariance matrix (σ_{12}, σ_{21}) are the covariances between the variables. They determine the orientation and shape of the distribution. If the covariance is positive, the distribution is stretched along the line $y = x$, indicating that the variables tend to increase together. If the covariance is negative, the distribution is stretched along the line $y = -x$, indicating that one variable tends to decrease as the other increases. If the covariance is zero, the variables are uncorrelated and the distribution is axis-aligned.

In Figure 88.5, the variances are identical and the variables are correlated (covariance is 4), so the distribution is stretched along the line $y = x$.

Bivariate Normal. Mean zero and positive covariance: $[[9, 4], [4, 9]]$




Figure 88.5.: Bivariate Normal. Mean zero and covariance $\Sigma = \begin{pmatrix} 9 & 4 \\ 4 & 9 \end{pmatrix}$

Bivariate Normal Distribution




Figure 88.6.: Bivariate Normal. Mean zero and covariance $\Sigma = \begin{pmatrix} 9 & 4 \\ 4 & 9 \end{pmatrix}$.

Example 88.8 (The Bivariate Normal Distribution with Mean Zero and Zero Covariances). The Bivariate Normal Distribution with Mean Zero and Zero Covariances $\sigma_{12} = \sigma_{21} = 0$.

$$\Sigma = \begin{pmatrix} 9 & 0 \\ 0 & 9 \end{pmatrix}$$

Bivariate Normal. Mean zero and covariance: $[[9, 0], [0, 9]]$




Figure 88.7.: Bivariate Normal. Mean zero and covariance $\Sigma = \begin{pmatrix} 9 & 0 \\ 0 & 9 \end{pmatrix}$

Example 88.9 (The Bivariate Normal Distribution with Mean Zero and Negative Covariances). The Bivariate Normal Distribution with Mean Zero and Negative Covariances $\sigma_{12} = \sigma_{21} = -4$.

$$\Sigma = \begin{pmatrix} 9 & -4 \\ -4 & 9 \end{pmatrix}$$

88.13. Covariance

In statistics, understanding the relationship between random variables is crucial for making inferences and predictions. Two common measures of such relationships are covariance and correlation. Covariance is a measure of how much two random variables change together. If the variables tend to show similar behavior (i.e., when one increases,

Bivariate Normal. Mean zero and covariance: $[[9, -4], [-4, 9]]$




Figure 88.8.: Bivariate Normal. Mean zero and covariance $\Sigma = \begin{pmatrix} 9 & -4 \\ -4 & 9 \end{pmatrix}$

the other tends to increase), the covariance is positive. Conversely, if they tend to move in opposite directions, the covariance is negative.

Definition 88.19 (Covariance). Covariance is calculated as:

$$\text{Cov}(X, Y) = E[(X - E[X])(Y - E[Y])] = E[XY] - E[X]E[Y]$$

Here, $E[X]$ and $E[Y]$ are the expected values (means) of X and Y , respectively. Covariance has units that are the product of the units of X and Y .

For a vector of random variables $\mathbf{Y} = (Y^{(1)}, \dots, Y^{(n)})^T$, the covariance matrix Σ encapsulates the covariances between each pair of variables:

$$\Sigma = \text{Cov}(\mathbf{Y}, \mathbf{Y}) = \begin{pmatrix} \text{Var}(Y^{(1)}) & \text{Cov}(Y^{(1)}, Y^{(2)}) & \dots \\ \text{Cov}(Y^{(2)}, Y^{(1)}) & \text{Var}(Y^{(2)}) & \dots \\ \vdots & \vdots & \ddots \end{pmatrix}$$

The diagonal elements represent the variances, while the off-diagonal elements are the covariances.

88.14. Correlation

88.14.1. Definitions

Definition 88.20 ((Pearson) Correlation Coefficient). The Pearson correlation coefficient, often denoted by ρ for the population or r for a sample, is calculated by dividing the covariance of two variables by the product of their standard deviations.

$$\rho_{XY} = \frac{\text{Cov}(X, Y)}{\sigma_X \sigma_Y}, \quad (88.13)$$

where $\text{Cov}(X, Y)$ is the covariance between variables X and Y , and σ_X and σ_Y are the standard deviations of X and Y , respectively.

Correlation, specifically the correlation coefficient, is a normalized measure of the linear relationship between two variables. It provides a value ranging from -1 to 1 , which is scale-free, making it easier to interpret:

- -1 : Perfect negative correlation, indicating that as one variable increases, the other decreases.
- 0 : No correlation, indicating no linear relationship between the variables.
- 1 : Perfect positive correlation, indicating that both variables increase together.

The correlation matrix Ψ provides a way to quantify the linear relationship between multiple variables, extending the notion of the correlation coefficient beyond just pairs of variables. It is derived from the covariance matrix Σ by normalizing each element with respect to the variances of the relevant variables.

Definition 88.21 (The Correlation Matrix Ψ). Given a set of random variables $X_1, X_2, \dots, X_n$, the covariance matrix Σ is:

$$\Sigma = \begin{pmatrix} \sigma_{11} & \sigma_{12} & \cdots & \sigma_{1n} \\ \sigma_{21} & \sigma_{22} & \cdots & \sigma_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ \sigma_{n1} & \sigma_{n2} & \cdots & \sigma_{nn} \end{pmatrix},$$

where $\sigma_{ij} = \text{cov}(X_i, X_j)$ is the covariance between the i^{th} and j^{th} variables. The correlation matrix Ψ is then defined as:

$$\Psi = (\rho_{ij}) = \left(\frac{\sigma_{ij}}{\sqrt{\sigma_{ii}\sigma_{jj}}} \right),$$

where:

- ρ_{ij} is the correlation coefficient between the i^{th} and j^{th} variables.
- σ_{ii} is the variance of the i^{th} variable, i.e., σ_i^2 .
- $\sqrt{\sigma_{ii}}$ is the standard deviation of the i^{th} variable, denoted as σ_i .

Thus, Ψ can also be expressed as:

$$\Psi = \begin{pmatrix} 1 & \frac{\sigma_{12}}{\sigma_1\sigma_2} & \cdots & \frac{\sigma_{1n}}{\sigma_1\sigma_n} \\ \frac{\sigma_{21}}{\sigma_2\sigma_1} & 1 & \cdots & \frac{\sigma_{2n}}{\sigma_2\sigma_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\sigma_{n1}}{\sigma_n\sigma_1} & \frac{\sigma_{n2}}{\sigma_n\sigma_2} & \cdots & 1 \end{pmatrix}$$

The correlation matrix Ψ has the following properties:

- The matrix Ψ is symmetric, meaning $\rho_{ij} = \rho_{ji}$.
- The diagonal elements are all 1, as $\rho_{ii} = \frac{\sigma_{ii}}{\sigma_i\sigma_i} = 1$.
- Each off-diagonal element is constrained between -1 and 1, indicating the strength and direction of the linear relationship between pairs of variables.

88.14.2. Computations

Example 88.10 (Computing a Correlation Matrix). Suppose you have a dataset consisting of three variables: X , Y , and Z . You can compute the correlation matrix as follows:

1. Calculate the covariance matrix Σ , which contains covariances between all pairs of variables.
2. Extract the standard deviations for each variable from the diagonal elements of Σ .
3. Use the standard deviations to compute the correlation matrix Ψ .

Suppose we have two sets of data points:

- $X = [1, 2, 3]$
- $Y = [4, 5, 6]$

We want to compute the correlation matrix Ψ for these variables. First, calculate the mean of each variable.

$$\bar{X} = \frac{1 + 2 + 3}{3} = 2$$

$$\bar{Y} = \frac{4 + 5 + 6}{3} = 5$$

Second, compute the covariance between X and Y . The covariance is calculated as:

$$\text{Cov}(X, Y) = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})$$

For our data:

$$\begin{aligned} \text{Cov}(X, Y) &= \frac{1}{3-1} [(1-2)(4-5) + (2-2)(5-5) + (3-2)(6-5)] \\ &= \frac{1}{2} [1 + 0 + 1] = 1 \end{aligned}$$

Third, calculate the variances of X and Y . Variance is calculated as:

$$\text{Var}(X) = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$$

$$= \frac{1}{2} [(1-2)^2 + (2-2)^2 + (3-2)^2] = \frac{1}{2}(1+0+1) = 1$$

Similarly,

$$\text{Var}(Y) = \frac{1}{2} [(4-5)^2 + (5-5)^2 + (6-5)^2] = \frac{1}{2}(1+0+1) = 1$$

Then, compute the correlation coefficient. The correlation coefficient ρ_{XY} is:

$$\begin{aligned} \rho_{XY} &= \frac{\text{Cov}(X, Y)}{\sqrt{\text{Var}(X)} \cdot \sqrt{\text{Var}(Y)}} \\ &= \frac{1}{\sqrt{1} \cdot \sqrt{1}} = 1 \end{aligned}$$

Finally, construct the correlation matrix. The correlation matrix Ψ is given as:

$$\Psi = \begin{pmatrix} 1 & \rho_{XY} \\ \rho_{XY} & 1 \end{pmatrix} = \begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix}$$

Thus, for these two variables, the correlation matrix indicates a perfect positive linear relationship (correlation coefficient of 1) between X and Y .

88.14.3. The Outer-product and the `np.outer` Function

The function `np.outer` from the NumPy library computes the outer product of two vectors. The outer product of two vectors results in a matrix, where each element is the product of an element from the first vector and an element from the second vector.

Definition 88.22 (Outer Product). For two vectors $\mathbf{a}$ and $\mathbf{b}$, the outer product is defined in terms of their elements as:

$$\text{outer}(\mathbf{a}, \mathbf{b}) = \begin{pmatrix} a_1 \cdot b_1 & a_1 \cdot b_2 & \cdots & a_1 \cdot b_n \\ a_2 \cdot b_1 & a_2 \cdot b_2 & \cdots & a_2 \cdot b_n \\ \vdots & \vdots & \ddots & \vdots \\ a_m \cdot b_1 & a_m \cdot b_2 & \cdots & a_m \cdot b_n \end{pmatrix},$$

where $\mathbf{a}$ is a vector of length m and $\mathbf{b}$ is a vector of length n .

Example 88.11 (Computing the Outer Product). We will consider two vectors, **a** and **b**:

```
import numpy as np

a = np.array([1, 2, 3])
b = np.array([4, 5])
outer_product = np.outer(a, b)
print("Vector a:", a)
print("Vector b:", b)
print("Outer Product:\n", outer_product)
```

```
Vector a: [1 2 3]
Vector b: [4 5]
Outer Product:
[[ 4  5]
 [ 8 10]
 [12 15]]
```

For the vectors defined:

$$\mathbf{a} = [1, 2, 3], \quad \mathbf{b} = [4, 5]$$

The outer product will be:

$$\begin{pmatrix} 1 \cdot 4 & 1 \cdot 5 \\ 2 \cdot 4 & 2 \cdot 5 \\ 3 \cdot 4 & 3 \cdot 5 \end{pmatrix} = \begin{pmatrix} 4 & 5 \\ 8 & 10 \\ 12 & 15 \end{pmatrix}.$$

Thus, `np.outer` creates a matrix with dimensions $m \times n$, where m is the length of the first vector and n is the length of the second vector. The function is particularly useful in various mathematical and scientific computations where matrix representations of vector relationships are needed.

Example 88.12 (Computing the Covariance and the Correlation Matrix). The following Python code computes the covariance and correlation matrices using the NumPy library.

```
import numpy as np

def calculate_cov_corr_matrices(data, rowvar=False)->(np.array, np.array):
    """
    Calculate the covariance and correlation matrices of the input data.
```

```

Args:
    data (np.array):
        Input data array.
    rowvar (bool):
        Whether the data is row-wise or column-wise.
        Default is False (column-wise).

Returns:
    np.array: Covariance matrix.
    np.array: Correlation matrix.

Examples:
    >>> data = np.array([[1, 2, 3],
                        [4, 5, 6],
                        [7, 8, 9]])
    >>> calculate_cov_corr_matrices(data)
    """
    cov_matrix = np.cov(data, rowvar=rowvar)
    std_devs = np.sqrt(np.diag(cov_matrix))
    # check whether the standard deviations are zero
    # and throw an error if they are
    if np.any(std_devs == 0):
        raise ValueError("Correlation matrix cannot be computed,"+
                        " because one or more variables have zero variance.")
    corr_matrix = cov_matrix / np.outer(std_devs, std_devs)
    return cov_matrix, corr_matrix

```

```

A = np.array([[0,1],
             [1,0]])
print(f"Input matrix:\n {A}")
Sigma, Psi = calculate_cov_corr_matrices(A)
print(f"Covariance matrix:\n {Sigma}")
print(f"Correlation matrix:\n {Psi}")

```

```

Input matrix:
[[0 1]
 [1 0]]
Covariance matrix:
[[ 0.5 -0.5]
 [-0.5  0.5]]
Correlation matrix:
[[ 1. -1.]
 [-1.  1.]]

```

88.14.4. Correlation and Independence

Definition 88.23 (Statistical Independence (Independence of Random Vectors)). Two random vectors are statistically independent if the joint distribution of the vectors is equal to the product of their marginal distributions.

This means that knowing the realization of one vector gives you no information about the realization of the other vector. This independence is a probabilistic concept used in statistics and probability theory to denote that two sets of random variables do not affect each other. Independence implies that all pairwise covariances between the components of the two vectors are zero, but zero covariance does not imply independence unless certain conditions are met (e.g., normality). Statistical independence is a stronger condition than zero covariance. Statistical independence is not related to the linear independence of vectors in linear algebra.

Example 88.13 (Covariance of Independent Variables). Consider a covariance matrix where variables are independent:

```
A = np.array([[1,-1],
[2,0],
[3,1],
[4,0],
[5,-1]])
print(f"Input matrix:\n {A}")
Sigma, Psi = calculate_cov_corr_matrices(A)
print(f"Covariance matrix:\n {Sigma}")
print(f"Correlation matrix:\n {Psi}")
```

```
Input matrix:
[[ 1 -1]
 [ 2  0]
 [ 3  1]
 [ 4  0]
 [ 5 -1]]
Covariance matrix:
[[2.5 0. ]
 [0.  0.7]]
Correlation matrix:
[[1. 0.]
 [0. 1.]]
```

Here, since the off-diagonal elements are 0, the variables are uncorrelated. X increases linearly, while Y alternates in a simple pattern with no trend that is linearly related to X .

Example 88.14 (Strong Correlation). For a covariance matrix with strong positive correlation:

```
A = np.array([[10,-1],
[20,0],
[30,1],
[40,2],
[50,3]])
print(f"Input matrix:\n {A}")
Sigma, Psi = calculate_cov_corr_matrices(A)
print(f"Covariance matrix:\n {Sigma}")
print(f"Correlation matrix:\n {Psi}")
```

Input matrix:

```
[[10 -1]
[20  0]
[30  1]
[40  2]
[50  3]]
```

Covariance matrix:

```
[[250.  25. ]
[ 25.   2.5]]
```

Correlation matrix:

```
[[1. 1.]
[1. 1.]]
```

A value close to 1 suggests a strong positive relationship between the variables.

Example 88.15 (Strong Negative Correlation).

```
A = np.array([[10,1],
[20,0],
[30,-1],
[40,-2],
[50,-3]])
print(f"Input matrix:\n {A}")
Sigma, Psi = calculate_cov_corr_matrices(A)
print(f"Covariance matrix:\n {Sigma}")
print(f"Correlation matrix:\n {Psi}")
```

Input matrix:

```
[[10  1]
[20  0]]
```

```
[30 -1]
[40 -2]
[50 -3]]
Covariance matrix:
[[250. -25. ]
 [-25.   2.5]]
Correlation matrix:
[[ 1. -1.]
 [-1.  1.]]
```

This matrix indicates a perfect negative correlation where one variable increases as the other decreases.

88.14.5. Pearson's Correlation

- Video: [Pearson's Correlation, Clearly Explained]

88.14.6. Interpreting the Correlation: Correlation Squared

Rummel (1976) describes how to interpret correlations as follows:

Seldom, indeed, will a correlation be zero or perfect. Usually, the covariation between things will be something like .56 or $-.16$. Clearly .56 is positive, indicating positive covariation; $-.16$ is negative, indicating some negative covariation. Moreover, we can say that the positive correlation is greater than the negative. But, we require more than. If we have a correlation of .56 between two variables, for example, what precisely can we say other than the correlation is positive and .56? The squared correlation describes the proportion of variance in common between the two variables. If we multiply this by 100 we then get the percent of variance in common between two variables. That is:

$$r_{XY}^2 \times 100 = \text{percent of variance in common between } X \text{ and } Y.$$

For example, we found that the correlation between a nation's power and its defense budget was .66. This correlation squared is .45, which means that across the fourteen nations constituting the sample 45 percent of their variance on the two variables is in common (or 55 percent is not in common). In thus squaring correlations and transforming covariance to percentage terms we have an easy to understand meaning of correlation. And we are then in a position to evaluate a particular correlation. As a matter of routine it is the squared correlations that should be interpreted. This is because the correlation coefficient is misleading in suggesting the existence of more covariation than exists, and this problem gets worse as the correlation approaches zero.

Example 88.16 (The relationship between study time and test scores). Imagine we are examining the relationship between the number of hours students study for a subject (Variable A) and their scores on a test (Variable B). After analyzing the data, we calculate a correlation of 0.8 between study time and test scores. When we square this correlation coefficient ($0.8^2 = 0.64$), we get 0.64 or 64%. This means that 64% of the variability in test scores can be accounted for by the variability in study hours. This indicates that a substantial part of why students score differently on the test can be explained by how much they studied. However, there remains 36% of the variability in test scores that needs to be explained by other factors, such as individual abilities, the difficulty of the test, or other external influences.

88.14.7. Partial Correlation

Often, a correlation between two variables X and Y can be found only because both variables are correlated with a third variable Z . The correlation between X and Y is then a spurious correlation. Therefore, it is often of interest to determine the correlation between X and Y while partializing a variable Z , i.e., the correlation between X and Y that exists without the influence of Z . Such a correlation $\rho_{(X,Y)/Z}$ is called the partial correlation of X and Y while holding Z constant. It is given by

$$\rho_{(X,Y)/Z} = \frac{\rho_{XY} - \rho_{XZ}\rho_{YZ}}{\sqrt{(1 - \rho_{XZ}^2)(1 - \rho_{YZ}^2)}}, \quad (88.14)$$

where ρ_{XY} is the correlation between X and Y , ρ_{XZ} is the correlation between X and Z , and ρ_{YZ} is the correlation between Y and Z (Hartung et al. 1995).

If the variables X , Y and Z are jointly normally distributed in the population of interest, one can estimate $\rho_{(X,Y)/Z}$ based on n realizations $x_1, \dots, x_n$, $y_1, \dots, y_n$ and $z_1, \dots, z_n$ of the random variables X , Y and Z by replacing the simple correlations ρ_{XY} , ρ_{XZ} and ρ_{YZ} with the empirical correlations $\hat{\rho}_{XY}$, $\hat{\rho}_{XZ}$ and $\hat{\rho}_{YZ}$. The partial correlation coefficient $\hat{\rho}_{(X,Y)/Z}$ is then estimated using

$$r_{(X,Y)/Z} = \frac{r_{XY} - r_{XZ}r_{YZ}}{\sqrt{(1 - r_{XZ}^2)(1 - r_{YZ}^2)}}. \quad (88.15)$$

Based on this estimated value for the partial correlation, a test at the α level for partial uncorrelatedness or independence of X and Y under Z can also be carried out. The hypothesis

$$H_0 : \rho_{(X,Y)/Z} = 0 \quad (88.16)$$

is tested against the alternative

$$H_1 : \rho_{(X,Y)/Z} \neq 0 \quad (88.17)$$

at the level α is discarded if

$$\left| \frac{r_{(X,Y)/Z} \sqrt{n-3}}{\sqrt{1-r_{(X,Y)/Z}^2}} \right| > t_{n-3, 1-\alpha/2}$$

applies. Here $t_{n-3, 1-\alpha/2}$ is the $(1-\alpha/2)$ -quantile of the t -distribution with $n-3$ degrees of freedom.

Example 88.17. For example, given economic data on the consumption X , income Y , and wealth Z of various individuals, consider the relationship between consumption and income. Failing to control for wealth when computing a correlation coefficient between consumption and income would give a misleading result, since income might be numerically related to wealth which in turn might be numerically related to consumption; a measured correlation between consumption and income might actually be contaminated by these other correlations. The use of a partial correlation avoids this problem (Wikipedia contributors 2024).

Example 88.18 (Partial Correlation. Numerical Example). Given the following data, calculate the partial correlation between A and B , controlling for C .

$$A = \begin{pmatrix} 2 \\ 4 \\ 15 \\ 20 \end{pmatrix}, \quad B = \begin{pmatrix} 1 \\ 2 \\ 3 \\ 4 \end{pmatrix}, \quad C = \begin{pmatrix} 0 \\ 0 \\ 1 \\ 1 \end{pmatrix}$$

```
from spotoptim.utils.stats import partial_correlation
import numpy as np
import pandas as pd
data = pd.DataFrame({
    'A': [2, 4, 15, 20],
    'B': [1, 2, 3, 4],
    'C': [0, 0, 1, 1]
})
print(f"Correlation between A and B: {data['A'].corr(data['B'])}")
pc = partial_correlation(data, method='pearson')
print(f"Partial Correlation between A and B: {pc[\"estimate\"][0, 1]}")
```

Correlation between A and B: 0.9695015519208121

Partial Correlation between A and B: 0.9191450300180576

Instead of considering only one variable Z , multiple variables Z_i can be considered. The formal definition of partial correlation reads as follows:

Definition 88.24 (Partial Correlation). Formally, the partial correlation between X and Y given a set of n controlling variables $\mathbf{Z} = \{Z_1, Z_2, \dots, Z_n\}$, written $\rho_{XY \cdot \mathbf{Z}}$, is the correlation between the residuals e_X and e_Y resulting from the linear regression of X with $\mathbf{Z}$ and of Y with $\mathbf{Z}$, respectively. The first-order partial correlation (i.e., when $n = 1$) is the difference between a correlation and the product of the removable correlations divided by the product of the coefficients of alienation of the removable correlations (Wikipedia contributors 2024).

Like the correlation coefficient, the partial correlation coefficient takes on a value in the range from -1 to 1. The value -1 conveys a perfect negative correlation controlling for some variables (that is, an exact linear relationship in which higher values of one variable are associated with lower values of the other); the value 1 conveys a perfect positive linear relationship, and the value 0 conveys that there is no linear relationship (Wikipedia contributors 2024).

Lemma 88.1 (Matrix Representation of the Partial Correlation). *The partial correlation can also be written in terms of the joint precision matrix (Wikipedia contributors 2024). Consider a set of random variables, $\mathbf{V} = \{X_1, \dots, X_n\}$ of cardinality n . We want the partial correlation between two variables X_i and X_j given all others, i.e., $\mathbf{V} \setminus \{X_i, X_j\}$. Suppose the (joint/full) covariance matrix $\Sigma = (\sigma_{ij})$ is positive definite and therefore invertible. If the precision matrix is defined as $\Omega = (p_{ij}) = \Sigma^{-1}$, then*

$$\rho_{X_i X_j \cdot \mathbf{V} \setminus \{X_i, X_j\}} = -\frac{p_{ij}}{\sqrt{p_{ii} p_{jj}}} \quad (88.18)$$

The semipartial correlation statistic is similar to the partial correlation statistic; both compare variations of two variables after certain factors are controlled for. However, to calculate the semipartial correlation, one holds the third variable constant for either X or Y but not both; whereas for the partial correlation, one holds the third variable constant for both (Wikipedia contributors 2024).

88.15. Hypothesis Testing and the Null-Hypothesis

88.15.1. Alternative Hypotheses, Main Ideas

88.15.2. p-values: What they are and how to interpret them

88.15.2.1. How to calculate p-values

88.15.2.2. p-hacking: What it is and how to avoid it

88.16. Statistical Power

- Video: Statistical Power, Clearly Explained

88.16.0.1. Power Analysis

- Video: Power Analysis, Clearly Explained!!!

88.17. The Central Limit Theorem

- Video: The Central Limit Theorem, Clearly Explained!!!

88.18. Maximum Likelihood

Maximum likelihood estimation is a method used to estimate the parameters of a statistical model. It is based on the principle of choosing the parameter values that maximize the likelihood of the observed data. The likelihood function represents the probability of observing the data given the model parameters. By maximizing this likelihood, we can find the parameter values that best explain the observed data.

Example 88.19 (Maximum Likelihood Estimation: Bernoulli Experiment). Bernoulli experiment for the event A , repeated n times, with the probability of success p . Result given as n tuple with entries A and $\bar{A}$. A appears k times. The probability of this event is given by

$$L(p) = p^k(1 - p)^{n-k} \quad (88.19)$$

Applying maximum likelihood estimation, we find the maximum of the likelihood function $L(p)$, i.e., we are trying to find the value of p that maximizes the probability of observing the data. This value will be denoted as $\hat{p}$.

Differentiating the likelihood function with respect to p and setting the derivative to zero, we find the maximum likelihood estimate $\hat{p}$. We get

$$\frac{d}{dp}L(p) = kp^{k-1}(1-p)^{n-k} - p^k(n-k)(1-p)^{n-k-1} \quad (88.20)$$

$$= p^{k-1}(1-p)^{n-k-1}(k(1-p) - p(n-k)) = 0 \quad (88.21)$$

Because

$$p \neq 0 \text{ and } (1-p)p \neq 0,$$

we can divide by $p^{k-1}(1-p)^{n-k-1}$ and get

$$k(1-p) - p(n-k) = 0. \quad (88.22)$$

Solving for p gives

$$\hat{p} = \frac{k}{n} \quad (88.23)$$

Therefore, the maximum likelihood estimate for the probability of success in a Bernoulli experiment is the ratio of the number of successes to the total number of trials.

□

Example 88.20 (Maximum Likelihood Estimation: Normal Distribution). Random variable $X \sim \mathcal{N}(\mu, \sigma^2)$ with n observations $x_1, x_2, \dots, x_n$. The likelihood function is given by

$$L(x_1, x_2, \dots, x_n, \mu, \sigma^2) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x_i - \mu)^2}{2\sigma^2}\right) \quad (88.24)$$

Taking the logarithm of the likelihood function, we get

$$\log L(x_1, x_2, \dots, x_n, \mu, \sigma^2) = -\frac{n}{2} \log(2\pi) - \frac{n}{2} \log(\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 \quad (88.25)$$

Partial derivative with respect to μ is

$$\frac{\partial}{\partial \mu} \log L(x_1, x_2, \dots, x_n, \mu, \sigma^2) = \frac{1}{\sigma^2} \sum_{i=1}^n (x_i - \mu) = 0 \quad (88.26)$$

We obtain the maximum likelihood estimate for μ as

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i, \quad (88.27)$$

which is the sample mean $\bar{x}$.

The partial derivative with respect to σ^2 is

$$\frac{\partial}{\partial \sigma^2} \log L(x_1, x_2, \dots, x_n, \mu, \sigma^2) = -\frac{n}{2\sigma^2} + \frac{1}{2(\sigma^2)^2} \sum_{i=1}^n (x_i - \mu)^2 = 0 \quad (88.28)$$

This can be simplified to

$$-n + \frac{1}{\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 = 0 \quad (88.29)$$

$$\Rightarrow n\sigma^2 = \sum_{i=1}^n (x_i - \mu)^2 \quad (88.30)$$

Using the maximum likelihood estimate for μ from Equation 88.27, we get

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu})^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2 \quad (88.31)$$

$$= \frac{n-1}{n} \frac{\sum_{i=1}^n (x_i - \hat{\mu})^2}{n-1} = \frac{n-1}{n} s^2, \quad (88.32)$$

where

$$s = \sqrt{\frac{\sum_{i=1}^n (x_i - \bar{x})^2}{n-1}} \quad (88.33)$$

is the sample standard deviation. We obtain the maximum likelihood estimate for σ^2 as

$$\hat{\sigma}^2 = \frac{n-1}{n} s^2 \quad (88.34)$$

- Video: Maximum Likelihood, clearly explained!!!
- Video: Probability is not Likelihood. Find out why!!!

88.19. Maximum Likelihood Estimation: Multivariate Normal Distribution

88.19.1. The Joint Probability Density Function of the Multivariate Normal Distribution

Consider the first n terms of an identically and independently distributed (i.i.d.) sequence $X^{(j)}$ of k -dimensional multivariate normal random vectors, i.e.,

$$X^{(j)} \sim N(\mu, \Sigma), j = 1, 2, \dots. \quad (88.35)$$

The joint probability density function of the j -th term of the sequence is

$$f_X(x_j) = \frac{1}{\sqrt{(2\pi)^k \det(\Sigma)}} \exp\left(-\frac{1}{2}(x_j - \mu)^T \Sigma^{-1}(x_j - \mu)\right),$$

where: μ is the $k \times 1$ mean vector; Σ is the $k \times k$ covariance matrix. The covariance matrix Σ is assumed to be positive definite, so that its determinant is strictly positive. We use $x_1, \dots, x_n$, i.e., the realizations of the first n random vectors in the sequence, to estimate the two unknown parameters μ and Σ .

88.19.2. The Log-Likelihood Function

Definition 88.25 (Likelihood Function). The likelihood function is defined as the joint probability density function of the observed data, viewed as a function of the unknown parameters.

Since the terms in the sequence Equation 88.35 are independent, their joint density is equal to the product of their marginal densities. As a consequence, the likelihood function can be written as the product of the individual densities:

$$\begin{aligned} L(\mu, \Sigma) &= \prod_{j=1}^n f_X(x_j) = \prod_{j=1}^n \frac{1}{\sqrt{(2\pi)^k \det(\Sigma)}} \exp\left(-\frac{1}{2}(x_j - \mu)^T \Sigma^{-1}(x_j - \mu)\right) \\ &= \frac{1}{(2\pi)^{nk/2} \det(\Sigma)^{n/2}} \exp\left(-\frac{1}{2} \sum_{j=1}^n (x_j - \mu)^T \Sigma^{-1}(x_j - \mu)\right). \end{aligned} \quad (88.36)$$

Taking the natural logarithm of the likelihood function, we obtain the log-likelihood function:

Example 88.21 (Log-Likelihood Function of the Multivariate Normal Distribution). The log-likelihood function of the multivariate normal distribution is given by

$$\ell(\mu, \Sigma) = -\frac{nk}{2} \ln(2\pi) - \frac{n}{2} \ln(\det(\Sigma)) - \frac{1}{2} \sum_{j=1}^n (x_j - \mu)^T \Sigma^{-1}(x_j - \mu).$$

The likelihood function is well-defined only if $\det(\Sigma) > 0$.

88.20. Cross-Validation

- Video: Machine Learning Fundamentals: Cross Validation

88.20.0.1. Bias and Variance

- Video: Machine Learning Fundamentals: Bias and Variance

88.21. Mutual Information

R^2 works only for numerical data. Mutual information can also be used, when the dependent variable is boolean or categorical. Mutual information provides a way to quantify the relationship of a mixture of continuous and discrete variables to a target variable. Mutual information explains how closely related two variables are. It is a measure of the amount of information that one variable provides about another variable.

Definition 88.26 (Mutual Information).

$$\text{MI} = \sum_{x \in X} \sum_{y \in Y} p(x, y) \log \left(\frac{p(x, y)}{p(x)p(y)} \right) \quad (88.37)$$

The terms in the nominator and denominator in Equation 88.37 are the joint probability, $p(x, y)$, the marginal probability of X , $p(x)$ and the marginal probability of Y , $p(y)$ respectively. Joint probabilities are the probabilities of two events happening together. Marginal probabilities are the probabilities of individual events happening. If $p(x, y) = p(x)p(y)$, then X and Y are independent. In this case, the mutual information is one.

The mutual information of two variables X and Y is

$$\text{MI}(X, Y) = p(X, Y) \log \left(\frac{p(X, Y)}{p(X)p(Y)} \right) + p(X, \bar{Y}) \log \left(\frac{p(X, \bar{Y})}{p(X)p(\bar{Y})} \right) + \quad (88.38)$$

$$p(\bar{X}, Y) \log \left(\frac{p(\bar{X}, Y)}{p(\bar{X})p(Y)} \right) + p(\bar{X}, \bar{Y}) \log \left(\frac{p(\bar{X}, \bar{Y})}{p(\bar{X})p(\bar{Y})} \right) \quad (88.39)$$

In general, when at least one of the two features has no variance (i.e., it is constant), the mutual information is zero, because something that never changes cannot tell us about something that does. When two features change, but change in exact the same way, then MI is 1/2. When two factors change, but in exact the opposite way, then MI is 1/2. When both features change, it does not matter if they change in the exact same or exact opposite ways; both result in the same MI value

**Mutual Information**

- Video: Mutual Information, Clearly Explained

88.22. **t-SNE**

- Video: *t*-SNE, Clearly Explained

89. Hypothesis Testing

89.1. Paired-Sample Analysis

The paired-sample experiment is characterized by a structure where each observation in one sample is uniquely and non-randomly coupled with an observation in another sample. This coupling is not a matter of convenience but a deliberate methodological choice intended to control for extraneous variability. Common applications include “before-and-after” studies, where the same subjects are measured prior to and following an intervention; studies involving naturally matched pairs, such as twins or spouses; or scenarios where two different measurement techniques are applied to the same set of specimens.

The analytical elegance of the paired design lies in its transformation of a two-sample comparison into a more straightforward one-sample problem. Rather than analyzing the distributions of the two raw measurements, x_1 and x_2 , the procedure begins by calculating a single vector of difference scores for each pair, $d_i = x_{1i} - x_{2i}$.

The entire statistical inference is then conducted on this sample of differences, focusing on a single population parameter: the mean of the paired differences, denoted as μ_d .

The research question is thus reframed from “Is the mean of x_1 different from the mean of x_2 ?” to “Is the mean of the differences, μ_d , different from a specified value?” In the vast majority of cases, this specified value is zero, representing the null hypothesis of no difference between the paired measurements.

The primary statistical advantage of this transformation is a significant enhancement in statistical power, which is the ability of a test to detect a true effect. This increase in power is a direct mathematical consequence of the design’s structure, because the variance of the difference between two random variables, x_1 and x_2 , is defined by the formula:

$$\text{Var}(x_1 - x_2) = \text{Var}(x_1) + \text{Var}(x_2) - 2 \cdot \text{Cov}(x_1, x_2).$$

In a paired-sample design, the measurements x_1 and x_2 are derived from the same or closely matched subjects, which typically induces a positive correlation between them. This positive correlation results in a positive covariance term, $\text{Cov}(x_1, x_2)$. As this term is subtracted in the variance formula, the resulting variance of the difference scores, $\text{Var}(d)$, is smaller than the sum of the individual variances that would be relevant in an independent-samples design. A smaller variance leads to a smaller standard error of the mean difference, which in turn produces a larger test statistic for a given observed

difference. This amplified “signal-to-noise” ratio makes it more likely that a true effect will be deemed statistically significant, thereby increasing the power of the test.

The paired design should therefore be understood not merely as a data structure but as an active experimental control strategy. Much of the “noise” or unexplained variance in an experiment stems from stable, pre-existing differences between individual subjects (e.g., genetic predispositions, baseline health, cognitive ability). By using each subject as its own control, the paired design effectively isolates the effect of the intervention from this inter-subject variability.

Paired designs are more powerful, they often require a smaller sample size to detect an effect of a given magnitude compared to an independent-groups design.

89.1.1. Directional Hypotheses for a Paired Experiment

The formulation of statistical hypotheses is the formal process of translating a research question into a testable claim about a population parameter. This process is governed by a strict logical framework involving two competing statements: the null hypothesis (H_0) and the alternative hypothesis (H_1). The null hypothesis represents the claim of “no effect” or the status quo. It is the default position that is assumed to be true unless sufficient evidence is presented to refute it. In many experimental contexts, H_0 is the proposition that the researcher aims to reject. The alternative hypothesis, conversely, embodies the research hypothesis—it is the claim that an effect exists, for which the researcher is seeking evidence.

A critical requirement of this framework is that H_0 and H_1 must be mutually exclusive and exhaustive. This means that they cannot both be true simultaneously, and together they must cover all possible values of the population parameter under investigation.

89.1.1.1. Defining the Parameter of Interest and the Alternative Hypothesis

For a paired-sample experiment, the statistical inference does not focus on the individual population means of the two measurements (μ_{x_1} and μ_{x_2}). Instead, the central parameter of interest is the population mean of the paired differences, $\mu_d = \mu_{x_1} - \mu_{x_2}$. The user’s specific research question—that x_1 is better than x_2 , where “better” is defined as a smaller value—must be translated into a precise mathematical statement about this parameter. The claim x_1 is smaller than x_2 corresponds to the population-level statement $\mu_{x_1} < \mu_{x_2}$. By rearranging this inequality, we can express it in terms of the parameter μ_d : $\mu_{x_1} - \mu_{x_2} < 0$, which is equivalent to $\mu_d < 0$.

This statement, $\mu_d < 0$, becomes the alternative hypothesis (H_1). It is the specific, directional claim that the experiment is designed to support. The formulation is as follows:

Alternative Hypothesis,

$$H_1 : \mu_d < 0$$

This hypothesis dictates that the statistical test will be a one-tailed test, specifically a lower-tailed or left-tailed test. The term “one-tailed” signifies that the researcher is interested only in evidence indicating a difference in a particular direction (in this case, a negative difference). The entire region of rejection for the null hypothesis will be located in the extreme left tail of the test statistic’s sampling distribution.

89.1.1.2. Establishing the Null Hypothesis

Once the alternative hypothesis is defined, the null hypothesis is constructed as its logical complement.

The composite null hypothesis includes the boundary case of zero difference as well as any possible difference in the direction opposite to the alternative hypothesis:

Composite Null: $H_0 : \mu_d \geq 0$

This statement posits that the mean difference is either zero or positive (i.e., x_1 is not smaller than x_2).

89.1.1.3. The Critical Justification for a One-Tailed Test

The choice to conduct a one-tailed test is a significant methodological decision that must be justified a priori—that is, before the data are collected or analyzed. A one-tailed test offers a distinct advantage in statistical power. By concentrating the entire significance level (alpha, α) in one tail of the distribution, the test is more sensitive to detecting an effect in the specified direction. However, this increased power comes at a substantial cost: the test is rendered completely incapable of detecting a statistically significant effect in the opposite direction, no matter how large that effect might be.

Therefore, a one-tailed test is only appropriate when there is a strong theoretical foundation or overwhelming prior evidence suggesting that an effect in the opposite direction is either impossible or of no substantive interest.

89.1.1.4. One-tailed Test Justification Example

If a new manufacturing process is being tested that can only possibly improve efficiency or have no effect, but cannot logically make it worse, a one-tailed test might be justified. However, consider a scenario where a new drug is developed with the belief that it will be an improvement over an existing drug. Opting for a one-tailed test to maximize the chances of detecting this improvement would mean failing to test for the possibility that the new drug is actually less effective or even harmful. Such an oversight could have serious ethical and practical consequences.

89.1.2. The Paired Samples t-Test: A Technical Examination

The paired samples t-test is the statistical procedure used to evaluate the hypotheses formulated for a paired experimental design. It is a parametric test, meaning its validity relies on certain assumptions about the underlying data. Crucially, these assumptions apply not to the raw scores of the two measurements, but to the calculated difference scores (d) between the pairs.

89.1.2.1. 3.1 Core Assumptions of the Paired t-Test

The validity of the conclusions drawn from a paired t-test is contingent upon the satisfaction of four primary assumptions.

- **Continuous Data:** The dependent variable, and therefore the calculated differences, must be measured on a continuous scale, which includes interval and ratio levels of measurement. Examples include variables like temperature, weight, blood pressure, or test scores.
- **Independence of Pairs:** The observations between pairs must be independent of one another. For instance, the outcome for one subject in a before-after study should not influence the outcome for any other subject. This assumption is typically met through the use of random sampling from the population of interest.
- **Absence of Significant Outliers:** The set of difference scores should not contain extreme outliers. Outliers are data points that deviate markedly from the overall pattern of the data. They can exert a disproportionate influence on the sample mean and standard deviation, potentially distorting the test statistic and leading to invalid results. The presence of outliers can be effectively diagnosed by examining a boxplot of the difference scores.
- **Normality of Differences:** The population of difference scores from which the sample is drawn is assumed to be approximately normally distributed. The t-test is known to be fairly “robust” to violations of this assumption, especially as the sample size increases. This means that even if the data deviate moderately from a perfect normal distribution, the test can still provide valid results. This assumption can be checked visually using tools like histograms or Q-Q plots of the differences, or more formally with statistical tests such as the Shapiro-Wilk test.

The assumptions of the paired t-test are not arbitrary rules but are direct consequences of the test’s mathematical derivation. The test statistic,

$$t = \bar{d} / (s_d / \sqrt{n}),$$

is constructed under the premise that the sampling distribution of the mean difference ($\bar{d}$) is normally distributed. The Central Limit Theorem ensures this will be approximately true for large samples, but for smaller samples, the normality of the underlying population of differences is a more critical condition. Furthermore, the formula relies on

the sample standard deviation of the differences (s_d) as an estimate of the population standard deviation (σ_d). The presence of outliers can drastically inflate s_d , which in turn increases the denominator of the t-statistic. This artificially reduces the magnitude of the test statistic, thereby decreasing the power of the test to detect a true effect. Therefore, satisfying these assumptions is essential for ensuring that the calculated t-statistic is a reliable measure of the evidence against the null hypothesis.

89.1.2.2. A Nuanced Approach to Assumption Checking

While formal statistical tests for assumptions like normality exist, a sophisticated and practical approach is warranted. Over-reliance on tests like the Shapiro-Wilk test can be misleading. The central issue is that no real-world dataset is perfectly normal. Consequently, with a large sample size, a normality test will have high power and will likely detect even trivial and inconsequential deviations from normality, leading to a rejection of the null hypothesis of normality. Conversely, with a small sample size, the test will have low power and may fail to detect substantial and problematic deviations from normality that could invalidate the t-test.

A more informative approach focuses on assessing the impact of any non-normality. Visual inspection of the distribution of differences using histograms and Q-Q plots is often more valuable than a p-value from a formal test. The primary concerns are severe skewness or the presence of very heavy tails (i.e., a high propensity for outliers), as these can most seriously affect the validity of the t-test. The robustness of the t-test, particularly for sample sizes greater than 30, means that minor to moderate deviations from normality are generally not a cause for alarm.

89.1.2.3. Calculating the Test Statistic

The calculation of the paired t-test statistic is mathematically equivalent to performing a one-sample t-test on the vector of difference scores, testing whether their population mean is equal to zero. The procedure involves four steps:

- Calculate the mean of the differences ($\bar{d}$): This is the sample average of all the calculated difference scores.

$$\bar{d} = \frac{\sum_{i=1}^n d_i}{n}$$

where d_i is the difference for the i-th pair and n is the number of pairs.

- Calculate the standard deviation of the differences (s_d): This measures the variability of the difference scores around their mean.

$$s_d = \sqrt{\frac{\sum_{i=1}^n (d_i - \bar{d})^2}{n - 1}}$$

89. Hypothesis Testing

- Calculate the standard error of the mean difference ($SE_{\bar{d}}$): This is an estimate of the standard deviation of the sampling distribution of the mean difference.

$$SE_{\bar{d}} = \frac{s_d}{\sqrt{n}}$$

- Calculate the t-statistic (t): This is the ratio of the observed mean difference to its standard error. It quantifies how many standard errors the observed mean difference is away from the null-hypothesized value of zero.

$$t = \frac{\bar{d} - \mu_0}{SE_{\bar{d}}} = \frac{\bar{d}}{s_d/\sqrt{n}}$$

where μ_0 is the value from the null hypothesis, which is 0 in this case.

89.1.2.4. The Decision Framework

The calculated t-statistic is then compared to the theoretical Student's t-distribution with $n - 1$ degrees of freedom (df) to determine the statistical significance of the result. This comparison can be made using two complementary approaches.

- **The Critical Value Approach:** A critical value is determined from a t-distribution table or calculator based on the chosen significance level (α) and the degrees of freedom (df). For the lower-tailed test ($H_1 : \mu_d < 0$), this will be a negative value in the left tail of the distribution (e.g., t_{critical} for $\alpha = 0.05$ and $\text{df}=29$ is -1.699). The decision rule is: if the calculated t-statistic is less than or equal to the negative critical value ($t_{\text{calculated}} \leq t_{\text{critical}}$), the null hypothesis is rejected.
- **The p-value Approach:** The p-value represents the probability of observing a sample mean difference as extreme as, or more extreme than, the one obtained in the study, assuming that the null hypothesis is true. For a lower-tailed test, the p-value is the area under the t-distribution curve to the left of the calculated t-statistic: $p = P(T_{df} \leq t_{\text{calculated}})$. The decision rule is: if the p-value is less than or equal to the significance level ($p \leq \alpha$), the null hypothesis is rejected. This approach is generally preferred in modern statistical practice as it provides a more nuanced measure of the strength of evidence against the null hypothesis.

89.1.3. A Priori Sample Size Determination for Desired Power

One of the most critical steps in the design of a rigorous quantitative study is the a priori determination of the required sample size. Conducting a study with too few subjects runs a high risk of a Type II error—failing to detect a real effect that exists in the population. Conversely, a study with too many subjects wastes resources

and may needlessly expose participants to experimental conditions. A formal power analysis allows researchers to calculate the minimum sample size needed to have a high probability of detecting an effect of a specified magnitude, thereby balancing statistical rigor with practical and ethical constraints.

89.1.3.1. The Four Pillars of Power Analysis

The calculation of sample size for a paired t-test is a function of four interrelated parameters. A researcher must specify values for each of these to proceed.

Definition 89.1 (Significance Level (α)). The significant level is the probability of committing a Type I error, which is the error of rejecting the null hypothesis when it is actually true (a “false positive”). This value represents the threshold for statistical significance and is conventionally set at 0.05, corresponding to a 5% risk of a false positive conclusion.

Definition 89.2 (Statistical Power ($1 - \beta$)). Power is the probability of correctly rejecting the null hypothesis when it is false. It is the probability that the test will detect a true effect of a specified magnitude. The desired power is typically set at 0.80 or 0.90, which means the researcher desires an 80% or 90% chance of finding a statistically significant result if the hypothesized effect truly exists in the population. The term β represents the probability of a Type II error (a “false negative”).

89.1.4. Effect Size (Δ)

The effect size (Δ) is the magnitude of the mean difference (μ_d) that the study is designed to detect. It should represent the smallest difference that is considered scientifically, clinically, or practically meaningful within the research domain. This is often the most challenging parameter to specify, as it requires substantive expertise rather than purely statistical knowledge.

Definition 89.3 (Standard Deviation of Differences (σ_d)). The Standard Deviation of Differences (σ_d) is an estimate of the population standard deviation of the paired differences. It quantifies the expected variability, or “noise,” in the difference scores.

Example 89.1 (Sample Size Calculation). For instructional purposes, the sample size formula can be presented using an approximation based on the standard normal (Z) distribution. This formula is generally accurate for larger sample sizes (e.g., $n > 30$). For a one-tailed test, the formula is:

$$n = \frac{(Z_\alpha + Z_\beta)^2 \sigma_d^2}{\Delta^2}$$

Where:

89. Hypothesis Testing

- n is the required number of pairs.
- Z_α is the critical Z-value corresponding to the one-tailed significance level α . For $\alpha = 0.05$, Z_α is 1.645.
- Z_β is the critical Z-value corresponding to the desired power. For a power of 0.80 ($\beta = 0.20$), Z_β is 0.84.
- σ_d is the estimated standard deviation of the differences.
- Δ is the target effect size or the mean difference to be detected.

While this formula provides a good estimate, more precise calculations, particularly for smaller sample sizes, are performed using the non-central t-distribution.

89.1.4.1. Practical Challenges: Estimating σ_d and Δ

The primary practical hurdles in any sample size calculation are obtaining reasonable estimates for the effect size (Δ) and the variability (σ_d).

- **Estimating Δ :** The choice of Δ should be driven by subject-matter expertise. The question to ask is not “What effect do I expect to find?” but rather “What is the smallest effect that would be considered important and meaningful in my field?” This ensures the study is powered to detect a difference that matters.
- **Estimating σ_d :** Several strategies can be employed to estimate the standard deviation of the differences:
 - **Conduct a Pilot Study:** A small-scale preliminary study can provide a direct sample estimate (s_d) of the variability.
 - **Use Data from Previous Research:** Similar studies published in the literature may report the standard deviation of the differences, which can be used as an estimate.
 - **Estimate from Individual Standard Deviations and Correlation:** This is a particularly powerful technique. If previous research provides estimates for the standard deviations of the two individual measurements (σ_1 and σ_2) and the correlation between them (ρ), the standard deviation of the differences can be calculated using the formula:

$$\sigma_d = \sqrt{\sigma_1^2 + \sigma_2^2 - 2\rho\sigma_1\sigma_2}$$

This last formula reveals a crucial design principle. The most effective way to design a more efficient experiment (i.e., one that requires a smaller sample size) is to reduce σ_d . The formula shows that σ_d decreases as the correlation ρ increases. This provides a direct, actionable insight for experimental design: researchers should strive to maximize the correlation between the paired measurements. In a before-after study, this might mean minimizing the time between measurements to reduce the influence of other factors. In a matched-pairs study, it means matching subjects on variables that are known to be strongly related to the outcome measure. This transforms the abstract

89.1. Paired-Sample Analysis

statistical concept of correlation into a practical tool for creating more powerful and efficient research.

The following table serves as a practical checklist for researchers preparing to conduct a power analysis, ensuring all necessary components are considered.

Parameter	Symbol	Description	Common Values / Source
Significance Level	α	Probability of a Type I error (false positive).	0.05, 0.01
Statistical Power	$1 - \beta$	Probability of detecting a true effect.	0.80, 0.90
Directionality	Tails	Number of tails in the hypothesis test.	One-tailed (for directional H_1)
Effect Size	Δ	Minimum meaningful mean difference (μ_d) to detect.	Based on domain knowledge, clinical relevance.
Variability	σ_d	Standard deviation of the paired differences.	Estimated from pilot data, literature, or correlation.

Example 89.2 (Testing a New Fuel Additive). A team of automotive engineers has developed a new fuel additive (c_1) that is hypothesized to reduce fuel consumption in passenger vehicles compared to the standard fuel (c_2). In this context, a “better” outcome is a smaller value for fuel consumption, typically measured in liters per 100 kilometers (L/100km). To test this hypothesis, the team plans a paired-sample experiment. They will use a fleet of cars, and each car will be driven over a standardized test course twice: once with the standard fuel and once with the new additive. This design ensures that factors specific to each car (e.g., engine wear, tire type) are controlled for, as each car serves as its own control.

89.1.4.2. Formulating the Hypotheses

- **Research Question:** Does the new fuel additive reduce fuel consumption?
- **Parameter of Interest:** The population mean of the paired differences in fuel consumption, $\mu_d = \mu_{\text{additive}} - \mu_{\text{standard}}$.
- **Hypotheses:** The research question implies a directional prediction. A reduction in consumption means the value for the additive (μ_{additive}) should be less than the value for the standard fuel (μ_{standard}), resulting in a negative mean difference.

89. Hypothesis Testing

- **Null Hypothesis (H_0):** $\mu_d \geq 0$. (The new additive does not reduce fuel consumption or increases it.)
- **Alternative Hypothesis (H_1):** $\mu_d < 0$. (The new additive reduces fuel consumption.)
- **Type of Test:** This setup requires a one-tailed (lower-tailed) paired t-test.

89.1.4.3. A Priori Sample Size Calculation

Before commencing the experiment, the engineers perform a power analysis to determine the number of cars required for the study.

- **Specification of Parameters:**
 - **Significance Level (α):** They choose a conventional α of 0.05. For a one-tailed test, the corresponding Z_α value is 1.645.
 - **Statistical Power ($1 - \beta$):** They desire a power of 0.80, which is a standard choice in many fields. This corresponds to $\beta = 0.20$, and the Z_β value is 0.84.
 - **Effect Size (Δ):** After deliberation, the team decides that a reduction of 0.5 L/100km is the minimum effect size that would be commercially viable and practically meaningful. Thus, $\Delta = -0.5$.
 - **Standard Deviation of Differences (σ_d):** Based on a small pilot study with a few vehicles, they estimate the standard deviation of the differences to be $\sigma_d = 1.2$ L/100km.
- **Calculation (using the Z-approximation):**

$$n = \frac{(Z_\alpha + Z_\beta)^2 \sigma_d^2}{\Delta^2} = \frac{(1.645 + 0.84)^2 (1.2)^2}{(-0.5)^2} = \frac{(2.485)^2 (1.44)}{0.25} = \frac{(6.175)(1.44)}{0.25} \approx 35.58$$

- **Conclusion:** Since the number of subjects must be an integer, the result is rounded up. The team concludes that a sample size of 36 cars is required to achieve 80% power to detect a mean reduction of 0.5 L/100km at a 0.05 significance level.

89.1.4.4. Data Collection and Presentation (Hypothetical)

The experiment is conducted with 36 cars. For each car, fuel consumption is measured with the standard fuel (c_2) and the new additive (c_1). The difference ($d = c_1 - c_2$) is calculated for each car.

89.1.4.5. Checking the Assumptions

The engineers analyze the 36 difference scores. A boxplot reveals no extreme outliers. A histogram and a Q-Q plot of the differences show a distribution that is roughly symmetric and unimodal, suggesting that the assumption of normality is reasonably met, especially given the sample size of 36, which is large enough for the Central Limit Theorem to lend robustness to the t-test.

89.1.4.6. Calculating the Test Statistic

The analysis of the hypothetical data from the 36 cars yields the following summary statistics for the difference scores:

- Sample Mean of the Differences ($\bar{d}$): -0.65 L/100km
- Sample Standard Deviation of the Differences (s_d): 1.15 L/100km

Using these values, the test statistic is calculated:

- **Standard Error of the Mean Difference ($SE_{\bar{d}}$):**

$$SE_{\bar{d}} = \frac{s_d}{\sqrt{n}} = \frac{1.15}{\sqrt{36}} = \frac{1.15}{6} \approx 0.1917$$

- **t-statistic (t):**

$$t = \frac{\bar{d}}{SE_{\bar{d}}} = \frac{-0.65}{0.1917} \approx -3.39$$

- **Degrees of Freedom (df):** $df = n - 1 = 36 - 1 = 35$.

89.1.4.7. Determining the p-value

The calculated t-statistic ($t = -3.39$) with 35 degrees of freedom is used to find the p-value. Using statistical software or a t-distribution calculator for this lower-tailed test, the p-value is found to be approximately 0.0008.

89.1.4.8. Drawing a Conclusion

This final step involves making a statistical decision and interpreting it in the context of the original research question.

- **Statistical Decision:** The calculated p-value (0.0008) is substantially smaller than the pre-specified significance level ($\alpha = 0.05$). Therefore, the decision is to reject the null hypothesis (H_0).

- **Interpretation in Context:** There is strong, statistically significant evidence to conclude that the new fuel additive is effective in reducing fuel consumption. The observed average reduction in the sample was 0.65 L/100km.
- **Confidence Interval for the Mean Difference:** To provide a more complete picture, a confidence interval for the true mean difference is calculated. Because a one-tailed test with $\alpha = 0.05$ was performed, a corresponding 90% two-sided confidence interval is appropriate (as it leaves 5% in each tail). The formula is $\bar{d} \pm t_{\text{critical}}^* \cdot SE_{\bar{d}}$, where t_{critical}^* for a 90% CI with 35 df is approximately 1.69. Interval = $-0.65 \pm 1.69 \cdot 0.1917$ Interval = -0.65 ± 0.324 90% Confidence Interval: [-0.974, -0.326]
- **Final Report Conclusion:** The study provides statistically significant evidence ($t(35) = -3.39, p = 0.0008$) that the new fuel additive reduces fuel consumption. The average reduction observed in the sample of 36 cars was 0.65 L/100km. We are 90% confident that the true mean reduction in fuel consumption for the entire population of similar vehicles lies between 0.326 and 0.974 L/100km. Since this interval does not contain zero and consists entirely of negative values (indicating a reduction), the result is consistent with the rejection of the null hypothesis and provides a range of plausible magnitudes for the additive's true effect.

89.2. Non-inferiority Test

- **The Goal:** To show that the performance of a model `c1` is not substantially worse than the performance of a model `c2`. “Worse” in this context means that the loss `c1` being larger than the loss `c2`. A specific, acceptable margin for this difference, which is `delta`, is defined. The goal is to prove that the true mean difference ($\mu_d = \mu_{c1} - \mu_{c2}$) is less than or equal to this margin.
- **Alternative Hypothesis (H_1):** This is the research claim. It can be stated as:
 - $H_1 : \mu_d \leq \text{delta}$
- **Null Hypothesis (H_0):** This is the opposite of the claim, representing the scenario to find evidence against. It can be stated as:
 - $H_0 : \mu_d > \text{delta}$
- **Type of Test:** This is a **one-tailed (left-tailed) test**. The null will be rejected only if the sample data provides strong evidence that the mean difference is significantly *less than* `delta`.

The key difference from a standard superiority test (where we test against zero) is that the comparison point is a non-zero value, `delta`. This is handled by modifying the t-statistic formula. The analysis still uses a paired t-test framework, but it compares the sample mean difference ($\bar{d}$) to `delta` instead of 0. The formula for the test statistic becomes:

$$t = \frac{\bar{d} - \text{delta}}{s_d / \sqrt{n}},$$

where:

- $\bar{d}$ is the sample mean of the differences ($c_1 - c_2$).
- **delta** is the specified margin of non-inferiority.
- s_d is the sample standard deviation of the differences.
- n is the number of pairs.

This calculated t-statistic is compared to a critical value from the t-distribution with $n - 1$ degrees of freedom.

Example 89.3 (Non-Inferiority of a New Manufacturing Process). A factory wants to switch to a new, cheaper manufacturing process (**c1**) to produce a specific component. The old process is **c2**. A key quality metric is the number of micro-fractures per component; a lower number is better.

The new process is only acceptable if it's not substantially worse than the old one. Management defines "not substantially worse" as producing, on average, no more than 5 extra micro-fractures per component.

1. State the Goal and Hypotheses

- **Goal:** Show that the new process is non-inferior to the old one, within a margin of **delta** = 5 micro-fractures.
- **Analysis Variable:** $d = c_1 - c_2$ (Difference in micro-fractures).
- **Hypotheses:**
 - $H_0 : \mu_d > 5$ (The new process is unacceptably worse; it adds more than 5 micro-fractures on average).
 - $H_1 : \mu_d \leq 5$ (The new process is acceptable; it is non-inferior).
- **Significance Level:** We'll use $\alpha = 0.05$.

2. Collect Sample Data (Hypothetical)

The factory runs a test on $n = 30$ components, using both processes on each component and counting the micro-fractures.

- **Sample Mean Difference ($\bar{d}$):** 3.8 micro-fractures. (On average, the new process produced 3.8 more fractures than the old one).
- **Sample Standard Deviation of Differences (s_d):** 4.5 micro-fractures.

3. Calculate the Test Statistic

Using the modified formula with **delta** = 5:

$$t = \frac{\bar{d} - \text{delta}}{s_d/\sqrt{n}}$$

$$t = \frac{3.8 - 5}{4.5/\sqrt{30}} \approx -1.46$$

4. Determine the Critical Value and Make a Decision

- **Degrees of Freedom (df):** $n - 1 = 30 - 1 = 29$.
- **Critical Value:** We need the critical t-value for a **one-tailed (left-tailed)** test with $\alpha = 0.05$ and $df = 29$. Looking this up in a t-distribution table, we find:
 $-t_{critical} \approx -1.699$
- **Rejection Rule:** We reject H_0 if our calculated t-statistic is *less than* (more negative than) the critical value.
- **Comparison:** Is $-1.46 < -1.699$? No, it is not.

5. State the Conclusion

Since the calculated t-statistic (-1.46) is not less than the critical value (-1.699), it does not fall into the rejection region. Therefore, we fail to reject the null hypothesis (H_0).

In the context of the problem, this means: **There is not enough statistical evidence to conclude that the new manufacturing process is non-inferior.** Even though the sample average of 3.8 was within the acceptable margin of 5, the result was not statistically significant enough to rule out the possibility that the true average difference could be greater than 5. The factory should not adopt the new process based on these results.

90. Addressing Multicollinearity: Principle Component Analysis (PCA) and Factor Analysis (FA)

90.1. Introduction

The concepts of Principal Component Analysis (PCA) and Factor Analysis (FA) are both dimensionality reduction techniques. They operate on different assumptions and serve distinct purposes. PCA aims to transform correlated variables into a smaller set of uncorrelated principal components that capture maximum variance, whereas Factor Analysis seeks to explain the correlations between observed variables in terms of a smaller number of unobserved, underlying factors.

After loading and preprocessing the data in Section 90.2, we will explore these methods to reduce dimensions and address multicollinearity. In Section 90.3 we will conduct linear regression on the extracted components or factors. Section 90.4 provides diagnostics for multicollinearity, including the coefficient table, eigenvalues, condition indices, and the KMO measure. Section 90.5 explains how PCA is applied to the data, while Section 90.6 discusses Factor Analysis. Both methods are used to mitigate multicollinearity issues in regression models. Section 90.8 shows how the reduced dimensions can be used in other machine learning models, such as Random Forests.

The following packages are used in this chapter:

```
import pandas as pd
import numpy as np
import statsmodels.api as sm
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from spoptim.factor_analyzer import FactorAnalyzer
from spoptim.factor_analyzer import calculate_kmo
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
import copy
```

```
from spotoptim.utils.stats import condition_index
from spotoptim.utils.pca import (get_pca, plot_pca_scee, plot_pca1vs2, get_pca_topk,
```

90.2. The Car-Sales Data Set

First, the data is preprocessed to ensure that it does not contain any NaN or infinite values. We load the data set, which contains information about car sales, including various features such as price, engine size, horsepower, and more. The initial shape of the DataFrame is (157, 27).

```
df = pd.read_csv("data/car_sales.csv", encoding="utf-8", index_col=None)
print(df.shape)
df.head()
```

(157, 27)

	Unnamed: 0	manufact	model	sales	resale	type	price	engine_s	horsepow	wheel
0	1	Acura	Integra	16.919	16.360	0	21.50	1.8	140.0	101.2
1	2	Acura	TL	39.384	19.875	0	28.40	3.2	225.0	108.1
2	3	Acura	CL	14.114	18.225	0	NaN	3.2	225.0	106.9
3	4	Acura	RL	8.588	29.725	0	42.00	3.5	210.0	114.6
4	5	Audi	A4	20.397	22.255	0	23.99	1.8	150.0	102.6

The first column is removed as it's an index or non-informative column.

```
df = df.drop(df.columns[0], axis=1)
df.head()
```

	manufact	model	sales	resale	type	price	engine_s	horsepow	wheelbas	width
0	Acura	Integra	16.919	16.360	0	21.50	1.8	140.0	101.2	67.3
1	Acura	TL	39.384	19.875	0	28.40	3.2	225.0	108.1	70.3
2	Acura	CL	14.114	18.225	0	NaN	3.2	225.0	106.9	70.6
3	Acura	RL	8.588	29.725	0	42.00	3.5	210.0	114.6	71.4
4	Audi	A4	20.397	22.255	0	23.99	1.8	150.0	102.6	68.2

90.2.1. The Target Variable

The `sales` variable, which is our target, is transformed to a log scale. Missing or zero values are handled by replacing them with the median.

```
df['ln_sales'] = np.log(df['sales'].replace(0, np.nan))
if df['ln_sales'].isnull().any() or np.isinf(df['ln_sales']).any():
    df['ln_sales'] = df['ln_sales'].fillna(df['ln_sales'].median()) # Or any other strategy
y = df['ln_sales']
```

90.2.2. The Features

90.2.2.1. Numerical Features

The following steps are performed during data preprocessing for numerical features:

1. Check for NaN or infinite values in X.
2. Replace NaN and infinite values with the median of the respective column.
3. Remove constant or nearly constant columns (not explicitly shown in code but stated in preprocessing steps).
4. Standardize the numerical predictors in X using `StandardScaler`.
5. Verify that `X_scaled` does not contain any NaN or infinite values.

```
# Use columns from 'price' to 'mpg' as predictors
independent_var_columns = ['price', 'engine_s', 'horsepow', 'wheelbas',
                           'width', 'length', 'curb_wgt', 'fuel_cap', 'mpg']
# Select those columns, ensuring they are numeric
X = df[independent_var_columns].apply(pd.to_numeric, errors='coerce')
# Handle missing/nans in features by using an appropriate imputation strategy
X = X.fillna(X.median()) # Impute with median or any other appropriate strategy
# Display the first few rows of the features
X.head()
```

	price	engine_s	horsepow	wheelbas	width	length	curb_wgt	fuel_cap	mpg
0	21.500	1.8	140.0	101.2	67.3	172.4	2.639	13.2	28.0
1	28.400	3.2	225.0	108.1	70.3	192.9	3.517	17.2	25.0
2	22.799	3.2	225.0	106.9	70.6	192.0	3.470	17.2	26.0
3	42.000	3.5	210.0	114.6	71.4	196.6	3.850	18.0	22.0
4	23.990	1.8	150.0	102.6	68.2	178.0	2.998	16.4	27.0

```

if X.isnull().any().any():
    print("NaNs detected in X. Filling with column medians.")
    X = X.fillna(X.median())

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
if np.isnan(X_scaled).any() or np.isinf(X_scaled).any():
    raise ValueError("X_scaled contains NaN or infinite values after preprocessing.")
# Convert the scaled data back to a DataFrame
X_scaled = pd.DataFrame(X_scaled, columns=X.columns)
# Display the first few rows of the scaled features
X_scaled.head()

```

	price	engine_s	horsepow	wheelbas	width	length	curb_wgt	fuel_cap	n
0	-0.410053	-1.214376	-0.814577	-0.827661	-1.121287	-1.119971	-1.182726	-1.228700	0
1	0.075070	0.134385	0.694066	0.081122	-0.246689	0.416070	0.223285	-0.193381	0
2	-0.318723	0.134385	0.694066	-0.076927	-0.159229	0.348634	0.148020	-0.193381	0
3	1.031255	0.423406	0.427835	0.937221	0.073997	0.693306	0.756545	0.013683	-
4	-0.234987	-1.214376	-0.637089	-0.643270	-0.858907	-0.700370	-0.607830	-0.400444	0

90.2.2.2. Categorical Features

Categorical features (like ‘type’) are one-hot encoded and then combined with the scaled numerical features.

```

categorical_cols = ['type'] # Replace if more categorical variables exist
encoder = OneHotEncoder(drop='first', sparse_output=False)
X_categorical_encoded = encoder.fit_transform(df[categorical_cols])
# Convert encoded data into a DataFrame
X_categorical_encoded_df = pd.DataFrame(X_categorical_encoded,
                                       columns=encoder.get_feature_names_out(categorical_cols))
X_categorical_encoded_df.describe(include='all')

```

	type_1
count	157.000000
mean	0.261146
std	0.440665
min	0.000000
25%	0.000000
50%	0.000000

90.3. Fit the Linear Regression Model

	type_1
75%	1.000000
max	1.000000

90.2.3. Combining Non-categorical and Categorical (encoded) Data

The final feature set `X_encoded` is created by concatenating the scaled numerical features and the one-hot encoded categorical features. This combined DataFrame will be used for regression analysis.

```
X_encoded = pd.concat([X_scaled, X_categorical_encoded_df], axis=1)
print(f"Dimension: {X_encoded.shape}")
print(list(X_encoded.columns))
```

```
Dimension: (157, 10)
['price', 'engine_s', 'horsepow', 'wheelbas', 'width', 'length', 'curb_wgt', 'fuel_cap', 'mpg', 'type_1']
```

```
X_encoded.describe(include='all')
```

	price	engine_s	horsepow	wheelbas	width	length	curb_w
count	1.570000e+02	1.570000e+02	1.570000e+02	1.570000e+02	1.570000e+02	1.570000e+02	1.570000e+02
mean	-5.091469e-16	1.018294e-16	1.584012e-16	3.892145e-15	-2.489162e-16	-1.584012e-16	-2.94173e-16
std	1.003200e+00	1.003200e+00	1.003200e+00	1.003200e+00	1.003200e+00	1.003200e+00	1.003200e+00
min	-1.272376e+00	-1.985097e+00	-2.323220e+00	-1.960346e+00	-2.491490e+00	-2.843334e+00	-2.37415e+00
25%	-6.459349e-01	-7.326758e-01	-6.370894e-01	-5.905870e-01	-8.006008e-01	-7.303412e-01	-6.44662e-01
50%	-3.187230e-01	-5.829498e-02	-1.489990e-01	-6.375655e-02	-1.738055e-01	4.142561e-02	-5.69560e-01
75%	3.232563e-01	4.234056e-01	5.165788e-01	6.211231e-01	6.570627e-01	6.558419e-01	6.41245e-01
max	4.089638e+00	4.758711e+00	4.687533e+00	4.111375e+00	2.552025e+00	2.783820e+00	3.51411e+00

90.3. Fit the Linear Regression Model

An Ordinary Least Squares (OLS) regression model is fitted using the preprocessed and combined features (`X_encoded`).

```
X_encoded_with_const = sm.add_constant(X_encoded) # Adds a constant term (intercept) to the model
model = sm.OLS(df['ln_sales'], X_encoded_with_const).fit()
```

90.3.1. Model Summary and Interpretation

90.3.1.1. Model Summary (ANOVA Table)

The ANOVA table shows a significant F-value (Prob (F-statistic) close to zero), indicating that the model is statistically significant and better than simply estimating the mean. The Adj. R-squared value, close to 0.40, suggests that nearly 40% of the variation in `ln_sales` is explained by the model.

```
print(model.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  ln_sales    R-squared:                0.485
Model:                            OLS      Adj. R-squared:           0.449
Method:                 Least Squares    F-statistic:               13.73
Date:                Wed, 01 Jul 2026    Prob (F-statistic):       7.69e-
17
Time:                  21:26:56    Log-Likelihood:            -
213.62
No. Observations:                157      AIC:                      449.2
Df Residuals:                    146      BIC:                      482.9
Df Model:                        10
Covariance Type:                nonrobust
=====
                    coef    std err          t      P>|t|      [0.025      0.975]
-----
const                3.0678      0.114     26.962      0.000        2.843        3.293
price               -0.6451      0.177     -3.655      0.000       -0.994         -
0.296
engine_s             0.3557      0.192      1.854      0.066       -0.023        0.735
horsepow            -0.1364      0.229     -0.596      0.552       -0.589        0.316
wheelbas             0.3166      0.174      1.816      0.071       -0.028        0.661
width               -0.0763      0.140     -0.547      0.586       -0.352        0.200
length              0.2029      0.185      1.099      0.273       -0.162        0.568
curb_wgt             0.0842      0.211      0.399      0.691       -0.333        0.501
fuel_cap            -0.2284      0.179     -1.276      0.204       -0.582        0.125
mpg                  0.3232      0.167      1.941      0.054       -0.006        0.652
type_1              0.8735      0.317      2.756      0.007        0.247        1.500
=====
Omnibus:                 41.296    Durbin-Watson:           1.423
Prob(Omnibus):            0.000    Jarque-Bera (JB):       107.145
Skew:                    -1.064    Prob(JB):               5.42e-
24

```

Kurtosis: 6.442 Cond. No. 11.4
 =====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Despite the positive model fit, many predictors show non-significant coefficients ($P > |t|$ much larger than 0.05), suggesting they contribute little to the model.

90.4. Collinearity Diagnostics

90.4.1. The Coefficient Table

The coefficient table provides further evidence of multicollinearity. The function `compute_coefficients_table()` from the `spotoptim` package is used here for comprehensive diagnostics.

```
from spotoptim.utils.stats import compute_coefficients_table
coeffs_table = compute_coefficients_table(
    model=model, X_encoded=X_encoded_with_const, y=y, vif_table=None
)
print("\nCoefficients Table:")
print(coeffs_table)
```

Coefficients Table:

	Variable	Zero-Order r	Partial r	Semipartial r	Tolerance	VIF
0	price	-0.551325	-0.289521	-0.217155	0.195662	5.110865
1	engine_s	-0.139066	0.151682	0.110172	0.165615	6.038084
2	horsepow	-0.386896	-0.049229	-0.035386	0.116244	8.602562
3	wheelbas	0.292461	0.148618	0.107895	0.200566	4.985881
4	width	0.040572	-0.045185	-0.032473	0.312675	3.198207
5	length	0.216882	0.090597	0.065310	0.178835	5.591740
6	curb_wgt	-0.040042	0.032981	0.023691	0.136742	7.313045
7	fuel_cap	-0.017278	-0.105041	-0.075831	0.190355	5.253349
8	mpg	0.119998	0.158587	0.115313	0.219810	4.549388
9	type_1	0.273500	0.222382	0.163754	0.314477	3.179880

For most predictors, the partial correlations (Partial r) decrease significantly compared to the zero-order correlations (Zero-Order r), which suggests multicollinearity. Tolerance values (1 minus the proportion of variance explained by other predictors) are low, indicating that approximately 70%-90% of a given predictor's variance can be explained

by other predictors. Tolerances close to 0 signify high multicollinearity. A Variance Inflation Factor (VIF) greater than 2 is typically considered problematic, and in this table, the smallest VIF is already greater than 2, confirming serious multicollinearity.

90.4.2. Eigenvalues and Condition Indices

Eigenvalues indicate how many factors or components can be meaningfully extracted. An eigenvalue greater than 1 suggests that the factor/component explains more variance than a single variable.

90.4.2.1. Eigenvalues

We use the `FactorAnalyzer` function from the `factor_analyzer` package to compute eigenvalues.

```
fa_temp = FactorAnalyzer(n_factors=X_encoded.shape[1], method="principal", rotation=None)
try:
    fa_temp.fit(X_encoded)
    ev, _ = fa_temp.get_eigenvalues()
    ev = np.sort(ev) # The source prints in ascending order
    print("Eigenvalues for each component:\n", ev)
except Exception as e:
    print(f"Error during factor analysis fitting: {e}")
    print("Consider reducing multicollinearity or removing problematic features.")
```

Eigenvalues for each component:

```
[0.06453844 0.09238346 0.13143422 0.15658714 0.20129034 0.25457714
 0.33712112 1.14556836 1.64880929 5.96769049]
```

The eigenvalue-based diagnostics confirm severe multicollinearity. Several eigenvalues are close to 0, indicating strong correlations among predictors.

90.4.2.2. Condition Indices

From `spotoptim.utils.stats`, we can compute the condition index, which is a measure of multicollinearity. A condition index greater than 15 suggests potential multicollinearity issues, and values above 30 indicate severe problems.

Condition indices, calculated as the square roots of the ratios of the largest eigenvalue to each subsequent eigenvalue, also highlight the issue.

Definition 90.1 (Condition Index). The Condition Index (CI_i) for the i -th eigenvalue is defined as:

$$CI_i = \sqrt{\frac{\lambda_{\max}}{\lambda_i}},$$

where $\lambda_{\max}$ is the largest eigenvalue of the scaled predictor correlation matrix, and λ_i is the i -th eigenvalue of the same matrix.

CI_i -values greater than 15 suggest a potential problem, and values over 30 indicate a severe problem.

```
X_cond = copy.deepcopy(X_encoded)
condition_index_df = condition_index(X_cond)
print("\nCondition Index:")
print(condition_index_df)
```

Condition Index:

	Index	Eigenvalue	Condition Index
0	0	0.047116	11.150268
1	1	0.067199	9.336595
2	2	0.121066	6.955955
3	3	0.146634	6.320499
4	4	0.157663	6.095428
5	5	0.248119	4.858905
6	6	0.338187	4.161884
7	7	0.736900	2.819449
8	8	1.531162	1.955951
9	9	5.857833	1.000000

90.4.3. Kayser-Meyer-Olkin (KMO) Measure

The KMO (Kaiser-Meyer-Olkin) measure is a metric for assessing the suitability of data for Factor Analysis. A KMO value of 0.6 or higher is generally considered acceptable, while a value below 0.5 indicates that the data is not suitable for Factor Analysis.

The KMO measure is based on the correlation and partial correlation between variables. It is calculated as the ratio of the squared sums of correlations to the squared sums of correlations plus the squared sums of partial correlations. KMO values range between 0 and 1, where values close to 1 suggest strong correlations and suitability for Factor Analysis, and values close to 0 indicate weak correlations and unsuitability.

```
kmo_all, kmo_model = calculate_kmo(X_encoded)
print(f"\nKMO measure: {kmo_model:.3f} (0.6+ is often considered acceptable)")
```

KMO measure: 0.835 (0.6+ is often considered acceptable)

A KMO measure of 0.835 indicates that the data is well-suited for Factor Analysis.

90.5. Addressing Multicollinearity with Principal Component Analysis (PCA)

Definition 90.2 (Multicollinearity and Multicorrelation). Multicorrelation is a general term that describes correlation between multiple variables. Multicollinearity is a specific problem in regression models caused by strong correlations between independent variables, making model interpretation difficult.

90.5.1. Introduction to PCA

Principal Component Analysis (PCA) is a popular unsupervised dimensionality reduction technique. It transforms a set of possibly correlated variables into a set of linearly uncorrelated variables called principal components. The first principal component accounts for as much of the variability in the data as possible, and each succeeding component accounts for as much of the remaining variability as possible. PCA is primarily used for data compression and simplifying complex datasets.

90.5.2. Application of PCA in Regression Problems:

- **Dimensionality Reduction:** PCA reduces the number of explanatory variables by transforming original variables into a smaller set of uncorrelated principal components, making regression algorithms less prone to overfitting, especially with many features.
- **Reducing Multicollinearity:** PCA effectively eliminates multicollinearity in linear regression models because the resulting principal components are orthogonal (uncorrelated) to each other, leading to more stable coefficient estimates.
- **Handling High-Dimensional Data:** It can reduce the dimensions of datasets with many variables to a manageable level before regression.
- **Reduced Overfitting Tendencies:** By removing redundant and highly correlated variables, PCA helps reduce the risk of overfitting by focusing the model on the most influential features.

90.5. Addressing Multicollinearity with Principal Component Analysis (PCA)

- **Improved Model Performance:** Performing regression on the most important principal components often leads to better generalization and improved model performance on new data.
- **Interpretation of Feature Importance:** PCA provides insights into the importance of original features through the variance explained by each principal component, which can identify combinations of variables best representing the data.

90.5.3. Scree Plot

Definition 90.3 (Scree Plot). A scree plot is a graphical representation of the eigenvalues of a covariance or correlation matrix in descending order. It is used to determine the number of significant components or factors in dimensionality reduction techniques.

Mathematically, the eigenvalues $\lambda_1, \lambda_2, \dots, \lambda_p$ are plotted against their corresponding component or factor indices $i = 1, 2, \dots, p$, where p is the total number of components or factors.

The eigenvalues are defined as:

$$\lambda_i = \text{Var}(\mathbf{z}_i),$$

where $\mathbf{z}_i$ is the i -th principal component or factor, and $\text{Var}(\mathbf{z}_i)$ is its variance.

The scree plot is constructed by plotting the points (i, λ_i) for $i = 1, 2, \dots, p$. The “elbow” in the plot, where the eigenvalues start to level off, indicates the optimal number of components or factors to retain.

90.5.4. Loading Scores (for PCA)

Loading scores in the context of Principal Component Analysis (PCA) represent the correlation or relationship between the original variables and the principal components.

Definition 90.4 (Loading Scores). The loading score for the j -th variable on the i -th principal component is defined as:

$$L_{ij} = \mathbf{a}_i^\top \mathbf{x}_j,$$

where:

$\mathbf{a}_i$ is the eigenvector corresponding to the i -th principal component, $\mathbf{x}_j$ is the standardized value of the j -th variable.

In PCA, the loading scores indicate how much each original variable contributes to a given principal component. High absolute values of L_{ij} suggest that the j -th variable strongly influences the i -th principal component. In PCA, loading scores can be viewed as directional vectors in the feature space. The magnitude of the score indicates how dominant the variable is in a component, while the sign represents the direction of the relationship. A high positive loading means a positive influence and correlation with the component, and a high negative loading indicates a negative correlation. Loading score values also show how much each original variable contributes to the explained variance in its respective principal component.

i Summary of Loading Scores

Loading scores are used in Principal Component Analysis (PCA).

- **Definition:** Loading scores represent the correlation or relationship between the original variables and the principal components.
- **Purpose:** They indicate how much each original variable contributes to a given principal component.
- **Mathematical Representation:** In PCA, the loading scores are the elements of the eigenvectors of the covariance (or correlation) matrix, scaled by the square root of the corresponding eigenvalues.
- **Interpretation:** High absolute values of loading scores suggest that the variable strongly influences the corresponding principal component.

Section 90.7.3 explains the difference between loading scores in PCA and factor loadings in FA.

90.5.5. PCA for Car Sales Example

90.5.5.1. Computing the Principal Components

The Principal Component Analysis (PCA) is applied only to the features (`X_encoded`), not to the target variable. We will use functions from `spotoptim.utils.pca`, which are based on `sklearn.decomposition.PCA` to perform PCA.

Step 1: Perform PCA and scale the data

```
pca, scaled_data, feature_names, sample_names, df_pca_components = get_pca(df=X_encoded)
```

Step 2: Plot the scree plot

```
plot_pca_scree(pca, df_name="Car Sales Data", max_scree=10)
```

90.5. Addressing Multicollinearity with Principal Component Analysis (PCA)




Figure 90.1.: Scree plot for PCA showing the explained variance ratio for each principal component.

Step 3: Plot the first two principal components

```
plot_pca1vs2(pca, df_pca_components, df_name="Car Sales Data")
```

90. Addressing Multicollinearity: Principle Component Analysis (PCA) and Factor Analysis (FA)




Figure 90.2.: Scatter plot of the first two principal components (PC1 vs PC2) for the Car Sales Data.

Step 4: Get the top k features influencing PC1 and PC2

```
top_k_features_pc1, top_k_features_pc2 = get_pca_topk(pca, feature_names, k=10)
print("Top 10 features influencing PC1:", top_k_features_pc1)
print("Top 10 features influencing PC2:", top_k_features_pc2)
```

Top 10 features influencing PC1: ['curb_wgt', 'engine_s', 'fuel_cap', 'mpg', 'width',
Top 10 features influencing PC2: ['price', 'wheelbas', 'horsepow', 'length', 'engine_s']

90.5.5.2. Loading Scores for PCA (10 Components)

```
# Get and print loading scores
loading_scores_df = get_loading_scores(pca, X_encoded.columns)
print("PCA Loading Scores (10 Components):\n", loading_scores_df)
```

PCA Loading Scores (10 Components):

	PC1	PC2	PC3	PC4	PC5	PC6	\
price	0.251214	0.568904	0.145341	-0.484049	0.335697	-0.164197	
engine_s	0.364662	0.204801	0.120816	0.392355	-0.419180	0.432765	
horsepow	0.317527	0.438619	0.305988	0.022026	-0.199571	-0.072824	

90.5. Addressing Multicollinearity with Principal Component Analysis (PCA)

wheelbas	0.300637	-0.470803	0.154529	-0.380962	-0.080738	-0.283304
width	0.343110	-0.186317	0.239367	0.582904	0.634863	-0.193732
length	0.302587	-0.398562	0.440896	-0.197738	-0.274463	0.077374
curb_wgt	0.382220	-0.036672	-0.252691	-0.118033	0.162548	0.314687
fuel_cap	0.358821	-0.102589	-0.424633	-0.203499	0.220563	0.384133
mpg	-0.351143	-0.076695	0.468213	-0.176015	0.326376	0.640074
type_1	0.072646	-0.095107	-0.362901	0.030564	-0.072634	0.019721

	PC7	PC8	PC9	PC10
price	-0.009121	-0.098316	-0.395916	0.227288
engine_s	0.279081	-0.161115	-0.433737	-0.039611
horsepow	0.132544	0.184063	0.718133	-0.019567
wheelbas	0.568127	-0.236665	-0.001192	-0.231065
width	0.030011	0.022386	-0.009774	0.091688
length	-0.517766	0.214795	-0.090943	0.335102
curb_wgt	-0.410187	-0.591957	0.262427	-0.248584
fuel_cap	0.163440	0.640656	-0.004986	-0.074928
mpg	0.243304	-0.105455	0.148714	0.109367
type_1	0.241967	-0.236377	0.200080	0.832423

Figure 90.3 shows the loading scores heatmap for the first 10 principal components. The heatmap visualizes how much each original feature contributes to each principal component, with darker colors indicating stronger contributions.

```
plot_loading_scores(loading_scores_df)
```

90. Addressing Multicollinearity: Principle Component Analysis (PCA) and Factor Analysis (FA)




Figure 90.3.: PCA Loading Scores Heatmap showing the influence of original features on the principal components.

90.5.6. Creating the Regression Model with Principal Components

Now, a linear regression model is fitted using the principal components derived from PCA. These components are uncorrelated, which should eliminate multicollinearity issues.

```
X_pca_model_with_const = sm.add_constant(df_pca_components)
model_pca = sm.OLS(y, X_pca_model_with_const).fit()
print("\nRegression on PCA Components:")
print(model_pca.summary())
```

Regression on PCA Components:

OLS Regression Results

```
=====
Dep. Variable:          ln_sales    R-squared:                0.485
Model:                  OLS        Adj. R-squared:            0.449
```


90.5. Addressing Multicollinearity with Principal Component Analysis (PCA)

Method:	Least Squares		F-statistic:		13.73	
Date:	Wed, 01 Jul 2026		Prob (F-statistic):		7.69e-	
17						
Time:	21:26:56		Log-Likelihood:		-	
213.62						
No. Observations:	157		AIC:		449.2	
Df Residuals:	146		BIC:		482.9	
Df Model:	10					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	3.2959	0.078	42.215	0.000	3.142	3.450
PC1	-0.0450	0.032	-1.392	0.166	-0.109	0.019
PC2	-0.6572	0.063	-10.383	0.000	-0.782	-
0.532						
PC3	-0.0624	0.091	-0.683	0.495	-0.243	0.118
PC4	0.2500	0.135	1.856	0.065	-0.016	0.516
PC5	-0.4628	0.157	-2.943	0.004	-0.774	-
0.152						
PC6	0.3734	0.197	1.893	0.060	-0.016	0.763
PC7	0.3777	0.205	1.847	0.067	-0.027	0.782
PC8	-0.4887	0.225	-2.171	0.032	-0.934	-
0.044						
PC9	0.2311	0.302	0.765	0.445	-0.366	0.828
PC10	0.5885	0.361	1.631	0.105	-0.125	1.302
=====						
Omnibus:	41.296		Durbin-Watson:		1.423	
Prob(Omnibus):	0.000		Jarque-Bera (JB):		107.145	
Skew:	-1.064		Prob(JB):		5.42e-	
24						
Kurtosis:	6.442		Cond. No.		11.2	
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

When all principal components are retained, the PCA regression model performs identically to the original OLS model in terms of R-squared, Adjusted R-squared, MSE, and RMSE. This is because PCA merely rotates the data, preserving all variance if all components are used. Its benefit lies in handling multicollinearity and enabling dimensionality reduction if fewer components are chosen without significant loss of information.

90.5.7. Collinearity Diagnostics for PCA Regression Model

Consider the eigenvalues of the PCA components to verify that they are uncorrelated. The eigenvalues should be close to 1, indicating that the components are orthogonal and do not exhibit multicollinearity.

```
fa_temp = FactorAnalyzer(n_factors=df_pca_components.shape[1], method="principal", rotation="varimax")
try:
    fa_temp.fit(df_pca_components)
    ev, _ = fa_temp.get_eigenvalues()
    ev = np.sort(ev) # The source prints in ascending order
    print("Eigenvalues for each component:\n", ev)
except Exception as e:
    print(f"Error during factor analysis fitting: {e}")
    print("Consider reducing multicollinearity or removing problematic features.")
```

Eigenvalues for each component:
[1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]

Next, we compute the condition indices for the PCA components to confirm that they are uncorrelated.

```
coeffs_table = compute_coefficients_table(
    model=model_pca, X_encoded=X_pca_model_with_const, y=y, vif_table=None
)
print("\nCoefficients Table:")
print(coeffs_table)
```

Coefficients Table:

	Variable	Zero-Order r	Partial r	Semipartial r	Tolerance	VIF
0	PC1	-0.082694	-0.114428	-0.082694	1.0	1.0
1	PC2	-0.616895	-0.651723	-0.616895	1.0	1.0
2	PC3	-0.040608	-0.056473	-0.040608	1.0	1.0
3	PC4	0.110272	0.151818	0.110272	1.0	1.0
4	PC5	-0.174882	-0.236673	-0.174882	1.0	1.0
5	PC6	0.112489	0.154797	0.112489	1.0	1.0
6	PC7	0.109722	0.151078	0.109722	1.0	1.0
7	PC8	-0.129005	-0.176859	-0.129005	1.0	1.0
8	PC9	0.045456	0.063189	0.045456	1.0	1.0
9	PC10	0.096903	0.133763	0.096903	1.0	1.0

90.5. Addressing Multicollinearity with Principal Component Analysis (PCA)

As expected, results indicate that there is no multicollinearity among the principal components. This confirms that PCA successfully addresses the multicollinearity problem. The R-squared and Adjusted R-squared values remain the same as the original OLS model since PCA preserves the total variance when all components are retained.

90.5.8. PCA: Creating the Regression Model with three Principle Components only

```
# Create a regression model using only the first three principal components
df_pc_reduced = df_pca_components.iloc[:, :3] # select the first three factors
X_model_pc_reduced = sm.add_constant(df_pc_reduced)
model_pc_reduced = sm.OLS(y, X_model_pc_reduced).fit()
print("\nRegression on PCs (three PCs only):")
print(model_pc_reduced.summary())

# Verify collinearity statistics for reduced PCs scores
coeffs_table_pc_reduced = compute_coefficients_table(
    model=model_pc_reduced, X_encoded=X_model_pc_reduced, y=y, vif_table=None
)
print("\nCoefficients Table (Reduced PCs Analysis Model):")
print(coeffs_table_pc_reduced)

# Verify condition indices for reduced FA scores
X_cond_pc_reduced = copy.deepcopy(df_pc_reduced)
condition_index_df_pc_reduced = condition_index(X_cond_pc_reduced)
print("\nCondition Index (Reduced PC Analysis Model):")
print(condition_index_df_pc_reduced)
```

Regression on PCs (three PCs only):

OLS Regression Results			
=====			
Dep. Variable:	ln_sales	R-squared:	0.389
Model:	OLS	Adj. R-squared:	0.377
Method:	Least Squares	F-statistic:	32.48
Date:	Wed, 01 Jul 2026	Prob (F-statistic):	2.66e-
16			
Time:	21:26:56	Log-Likelihood:	-
226.97			
No. Observations:	157	AIC:	461.9
Df Residuals:	153	BIC:	474.2

90. Addressing Multicollinearity: Principle Component Analysis (PCA) and Factor Analysis (FA)

Df Model:	3					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	3.2959	0.083	39.693	0.000	3.132	3.460
PC1	-0.0450	0.034	-1.309	0.193	-0.113	0.023
PC2	-0.6572	0.067	-9.762	0.000	-0.790	-
0.524						
PC3	-0.0624	0.097	-0.643	0.521	-0.254	0.129
Omnibus:		43.520	Durbin-Watson:			1.413
Prob(Omnibus):		0.000	Jarque-Bera (JB):			125.210
Skew:		-1.081	Prob(JB):			6.47e-
28						
Kurtosis:		6.804	Cond. No.			2.82

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Coefficients Table (Reduced PCs Analysis Model):

	Variable	Zero-Order r	Partial r	Semipartial r	Tolerance	VIF
0	PC1	-0.082694	-0.105209	-0.082694	1.0	1.0
1	PC2	-0.616895	-0.619530	-0.616895	1.0	1.0
2	PC3	-0.040608	-0.051883	-0.040608	1.0	1.0

Condition Index (Reduced PC Analysis Model):

	Index	Eigenvalue	Condition Index
0	0	0.736900	2.819449
1	1	1.531162	1.955951
2	2	5.857833	1.000000

90.6. Addressing Multicollinearity and Latent Structure with Factor Analysis (FA)

90.6.1. Introduction to Factor Analysis

Factor Analysis (FA) is a statistical method used to describe variability among observed, correlated variables in terms of a potentially lower number of unobserved variables called factors or latent variables. Unlike PCA, which is primarily a data reduction technique focused on maximizing variance explained, FA assumes that the observed

90.6. Addressing Multicollinearity and Latent Structure with Factor Analysis (FA)

variables are linear combinations of these underlying factors plus an error term. FA's main goal is to uncover the underlying structure that explains the correlations among observed variables.

90.6.2. Determining the Number of Factors for Factor Analysis

For Factor Analysis, the number of factors to extract is a crucial decision. A common approach, consistent with the KMO measure, is to consider factors with eigenvalues greater than 1 (Kaiser's criterion). Factor analysis is then performed, often with a rotation method like Varimax to improve factor interpretability.

```
anz_fak = 10 # Number of factors to extract, similar to the components in PCA
n_factors = min(anz_fak, X_encoded.shape[1])
fa = FactorAnalyzer(n_factors=n_factors, method="principal", rotation="varimax")
fa.fit(X_encoded) # Fit the Factor Analyzer
actual_factors = fa.loadings_.shape[1] # Number of factors actually extracted
print(f"actual_factors: {actual_factors}")
if actual_factors < n_factors:
    print(
        f"\nWarning: Only {actual_factors} factors could be extracted "
        f"(requested {n_factors})."
    )
factor_columns = [f"Factor{i+1}" for i in range(actual_factors)]
```

```
actual_factors: 10
```

90.6.3. Scree Plot for Factor Analysis

Figure 90.4 shows the eigenvalues for each factor extracted from Factor Analysis. The scree plot helps in determining the number of factors to retain by identifying the “elbow” point where the eigenvalues start to level off, indicating diminishing returns in explained variance.

```
plt.figure(figsize=(10, 6))
ev_fa, _ = fa.get_eigenvalues()
plt.plot(range(1, len(ev_fa) + 1), ev_fa, marker='o', linestyle='--')
plt.title('Scree Plot for Factor Analysis')
plt.xlabel('Number of Factors')
plt.ylabel('Eigenvalue')
plt.grid(True)
plt.xticks(range(1, len(ev_fa) + 1))
plt.show()
```




Figure 90.4.: Scree plot for Factor Analysis showing the eigenvalues for each factor.

90.6.4. Factor Loadings

Factor Loadings indicate how strongly each original variable is correlated with the extracted factors. High absolute values suggest that the variable has a significant influence on, or is strongly associated with, that factor. Loadings help in interpreting the meaning of each underlying factor.

i Summary of Factor Loadings

Factor loadings are used in Factor Analysis (FA). * Definition: Factor loadings represent the correlation or relationship between the observed variables and the latent factors. * Purpose: They indicate how much each observed variable is explained by a given factor. * Mathematical Representation: In FA, factor loadings are derived from the factor model, where observed variables are expressed as linear combinations of latent factors plus error terms. * Interpretation: High absolute values of factor loadings suggest that the variable is strongly associated with the corresponding factor.

Section 90.7.3 explains the difference between loading scores in PCA and factor loadings in FA.

90.6. Addressing Multicollinearity and Latent Structure with Factor Analysis (FA)


```
# Print factor loadings with 2 decimals
factor_loadings = fa.loadings_
print("Factor Loadings (rounded to 2 decimals):\n", np.round(factor_loadings, 2))

# Create a DataFrame for the factor loadings for better visualization
factor_loadings_df = pd.DataFrame(
    factor_loadings, index=X_encoded.columns, # Original feature names
    columns=factor_columns # Factor names
)

# Plot the heatmap for factor loadings
plt.figure(figsize=(10, 8))
sns.heatmap(
    factor_loadings_df, annot=True, # Annotate with values
    fmt=".2f", # Format values to 2 decimals
    cmap="coolwarm", # Color map
    cbar=True # Show color bar
)
plt.title("Factor Loadings Heatmap")
plt.xlabel("Factors")
plt.ylabel("Features")
plt.tight_layout()
plt.show()
```

```
Factor Loadings (rounded to 2 decimals):
[[ 0.12  0.97 -0.01  0.06 -0.01  0.08  0.01  0.09  0.01 -0.15]
 [ 0.12  0.6   0.18  0.28  0.61  0.2   0.31  0.09 -0.01  0.02]
 [ 0.06  0.87 -0.01  0.19  0.26  0.16  0.19  0.01 -0.02  0.26]
 [ 0.13  0.07  0.26  0.18  0.04  0.09  0.9   0.03  0.26 -0.01]
 [ 0.14  0.25  0.12  0.77  0.15  0.14  0.51  0.08  0.   0.01]
 [ 0.09  0.12  0.   0.2   0.14  0.11  0.92  0.09 -0.23  0.03]
 [ 0.28  0.43  0.41  0.26  0.15  0.25  0.44  0.48 -0.01 -0. ]
 [ 0.63  0.32  0.46  0.22  0.11  0.24  0.4   0.12  0.01  0. ]
 [-0.2  -0.41 -0.46 -0.2  -0.17 -0.66 -0.25 -0.1  -0.   -0.01]
 [ 0.1  -0.06  0.98  0.06  0.05  0.12  0.1   0.05  0.01 -0.  ]]
```

90. Addressing Multicollinearity: Principle Component Analysis (PCA) and Factor Analysis (FA)



90.6.5. Factor Scores

The factor scores are the transformed values of the original variables based on the extracted factors. These scores represent the values of the latent factors for each observation and can be used as new features in regression models, similar to principal components in PCA.

Definition 90.5 (Factor Scores). A **factor score** represents the value of a latent factor for a given observation, calculated as a linear combination of the observed variables weighted by the factor score coefficients.

Mathematically, the factor score for the i -th factor and the j -th observation is defined as:

$$F_{ji} = w_{i1}x_{j1} + w_{i2}x_{j2} + \cdots + w_{ip}x_{jp} = \sum_{k=1}^p w_{ik}x_{jk},$$

90.6. Addressing Multicollinearity and Latent Structure with Factor Analysis (FA)

where

- F_{ji} is the factor score for factor i and observation j ,
- w_{ik} is the factor score coefficient for variable k on factor i ,
- x_{jk} is the standardized value of variable k for observation j , and
- p is the number of observed variables.

```
# Factor scores for each row (shape: [n_samples, actual_factors])
X_factor_scores = fa.transform(X_encoded)
print(f"X_factor_scores shape: {X_factor_scores.shape}")

# Adapt the factor column names to the actual factor count
df_factors = pd.DataFrame(X_factor_scores, columns=factor_columns)
print(f"df_factors shape: {df_factors.shape}")
print(f"df_factors head:\n{df_factors.head()}")
```

X_factor_scores shape: (157, 10)

df_factors shape: (157, 10)

df_factors head:

	Factor1	Factor2	Factor3	Factor4	Factor5	Factor6	Factor7	\
0	-0.647996	-0.310986	-0.395620	-0.514476	-0.753763	-0.171572	-0.691765	
1	-0.171241	0.352069	-0.579629	-0.677204	0.113380	-0.329903	0.434305	
2	0.077192	0.050156	-0.595317	-0.396626	0.412052	-0.688322	0.246025	
3	-0.683708	0.820534	-0.676114	-0.796906	-0.241928	0.602161	1.058645	
4	0.615152	-0.262258	-0.541357	-0.489288	-1.207964	-0.186946	-0.485740	

	Factor8	Factor9	Factor10
0	-0.233725	0.567292	-0.139248
1	0.852994	-0.099874	1.690789
2	0.941176	-0.209195	2.468886
3	1.063771	1.022527	-1.245557
4	0.259073	0.073952	0.308099

90.6.6. Creating the Regression Model with Extracted Factors (from FA)

A linear regression model is built using all ten extracted factors from Factor Analysis. The expectation is that these factors are uncorrelated, addressing multicollinearity.

90. Addressing Multicollinearity: Principle Component Analysis (PCA) and Factor Analysis (FA)

```
X_model_fa = sm.add_constant(df_factors)
model_factors = sm.OLS(y, X_model_fa).fit()
print("\nRegression on Factor Scores (all 10 factors):")
print(model_factors.summary())

# Verify collinearity statistics for Factor Analysis scores (VIF and Tolerance)
coeffs_table_fa = compute_coefficients_table(
    model=model_factors, X_encoded=X_model_fa, y=y, vif_table=None
)
print("\nCoefficients Table (Factor Analysis Model):")
print(coeffs_table_fa)

# Verify condition indices
X_cond_fa = copy.deepcopy(df_factors)
condition_index_df_fa = condition_index(X_cond_fa)
print("\nCondition Index (Factor Analysis Model):")
print(condition_index_df_fa)
```

Regression on Factor Scores (all 10 factors):

OLS Regression Results

=====						
Dep. Variable:	ln_sales	R-squared:	0.485			
Model:	OLS	Adj. R-squared:	0.449			
Method:	Least Squares	F-statistic:	13.73			
Date:	Wed, 01 Jul 2026	Prob (F-statistic):	7.69e-			
17						
Time:	21:26:57	Log-Likelihood:	-			
213.62						
No. Observations:	157	AIC:	449.2			
Df Residuals:	146	BIC:	482.9			
Df Model:	10					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	3.2959	0.078	42.215	0.000	3.142	3.450
Factor1	-0.1366	0.078	-1.749	0.082	-0.291	0.018
Factor2	-0.7022	0.078	-8.994	0.000	-0.856	-
0.548						
Factor3	0.3035	0.078	3.888	0.000	0.149	0.458
Factor4	9.177e-06	0.078	0.000	1.000	-0.154	0.154
Factor5	0.1719	0.078	2.201	0.029	0.018	0.326

90.6. Addressing Multicollinearity and Latent Structure with Factor Analysis (FA)

Factor6	-0.1653	0.078	-2.117	0.036	-0.320	-
0.011						
Factor7	0.4130	0.078	5.290	0.000	0.259	0.567
Factor8	-0.0072	0.078	-0.092	0.927	-0.161	0.147
Factor9	0.0317	0.078	0.407	0.685	-0.123	0.186
Factor10	0.0665	0.078	0.852	0.396	-0.088	0.221
=====						
Omnibus:		41.296	Durbin-Watson:			1.423
Prob(Omnibus):		0.000	Jarque-Bera (JB):			107.145
Skew:		-1.064	Prob(JB):			5.42e-
24						
Kurtosis:		6.442	Cond. No.			1.00
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Coefficients Table (Factor Analysis Model):

	Variable	Zero-Order r	Partial r	Semipartial r	Tolerance	VIF
0	Factor1	-0.103920	-0.143257	-0.103920	1.0	1.0
1	Factor2	-0.534367	-0.597080	-0.534367	1.0	1.0
2	Factor3	0.231004	0.306300	0.231004	1.0	1.0
3	Factor4	0.000007	0.000010	0.000007	1.0	1.0
4	Factor5	0.130790	0.179228	0.130790	1.0	1.0
5	Factor6	-0.125772	-0.172560	-0.125772	1.0	1.0
6	Factor7	0.314284	0.401023	0.314284	1.0	1.0
7	Factor8	-0.005478	-0.007630	-0.005478	1.0	1.0
8	Factor9	0.024158	0.033630	0.024158	1.0	1.0
9	Factor10	0.050594	0.070298	0.050594	1.0	1.0

Condition Index (Factor Analysis Model):

	Index	Eigenvalue	Condition Index
0	0	1.00641	1.0
1	1	1.00641	1.0
2	2	1.00641	1.0
3	3	1.00641	1.0
4	4	1.00641	1.0
5	5	1.00641	1.0
6	6	1.00641	1.0
7	7	1.00641	1.0
8	8	1.00641	1.0
9	9	1.00641	1.0

As expected, the collinearity statistics (VIF and Tolerance) for the factor values show that they are uncorrelated (VIF=1, Tolerance=1). The condition indices are also all

90. Addressing Multicollinearity: Principle Component Analysis (PCA) and Factor Analysis (FA)

close to 1, confirming that Factor Analysis successfully mitigates multicollinearity. The coefficient estimates are larger relative to their standard errors compared to the original model, which can lead to more factors being identified as statistically significant.

If the R-squared and Adjusted R-squared values for `model_factors` are close to those of the original `model`, it indicates that the regression model based on Factor Analysis performs similarly well, while successfully reducing multicollinearity. When all factors are used, the predictive performance metrics are identical to the original OLS model.

90.6.7. Factor Analysis: Creating the Regression Model with three Extracted Factors only

90.6.7.1. Setting Up the Regression Model with Reduced Factors

To demonstrate the effect of dimensionality reduction, a regression model is created using only the first three extracted factors from Factor Analysis.

```
# Create a regression model using only the first three factors
df_factors_reduced = df_factors.iloc[:, :3] # select the first three factors
X_model_fa_reduced = sm.add_constant(df_factors_reduced)
model_factors_reduced = sm.OLS(y, X_model_fa_reduced).fit()
print("\nRegression on Factor Scores (three factors only):")
print(model_factors_reduced.summary())

# Verify collinearity statistics for reduced FA scores
coeffs_table_fa_reduced = compute_coefficients_table(
    model=model_factors_reduced, X_encoded=X_model_fa_reduced, y=y, vif_table=None
)
print("\nCoefficients Table (Reduced Factor Analysis Model):")
print(coeffs_table_fa_reduced)

# Verify condition indices for reduced FA scores
X_cond_fa_reduced = copy.deepcopy(df_factors_reduced)
condition_index_df_fa_reduced = condition_index(X_cond_fa_reduced)
print("\nCondition Index (Reduced Factor Analysis Model):")
print(condition_index_df_fa_reduced)
```

Regression on Factor Scores (three factors only):

OLS Regression Results

```
=====
Dep. Variable:          ln_sales    R-squared:                0.350
Model:                  OLS        Adj. R-squared:            0.337
```

90.6. Addressing Multicollinearity and Latent Structure with Factor Analysis (FA)

```

Method:                Least Squares    F-statistic:                27.43
Date:                  Wed, 01 Jul 2026  Prob (F-statistic):        2.99e-
14
Time:                  21:26:57    Log-Likelihood:            -
231.87
No. Observations:      157    AIC:                          471.7
Df Residuals:          153    BIC:                          484.0
Df Model:               3
Covariance Type:       nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----+-----
const          3.2959      0.086     38.474     0.000      3.127      3.465
Factor1       -0.1366      0.086     -1.594     0.113     -0.306      0.033
Factor2       -0.7022      0.086     -8.197     0.000     -0.871      -
0.533
Factor3         0.3035      0.086      3.543     0.001      0.134      0.473
=====
Omnibus:                43.992    Durbin-Watson:              1.418
Prob(Omnibus):           0.000    Jarque-Bera (JB):          134.618
Skew:                   -1.068    Prob(JB):                  5.86e-
30
Kurtosis:                7.002    Cond. No.                  1.00
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Coefficients Table (Reduced Factor Analysis Model):

	Variable	Zero-Order r	Partial r	Semipartial r	Tolerance	VIF
0	Factor1	-0.103920	-0.127811	-0.103920	1.0	1.0
1	Factor2	-0.534367	-0.552381	-0.534367	1.0	1.0
2	Factor3	0.231004	0.275385	0.231004	1.0	1.0

Condition Index (Reduced Factor Analysis Model):

	Index	Eigenvalue	Condition Index
0	0	1.00641	1.0
1	1	1.00641	1.0
2	2	1.00641	1.0

The collinearity statistics for the reduced factor set continue to show that they are uncorrelated, with VIFs of 1.0 and condition indices close to 1.

90.6.7.2. Comparison of Model Performance of the Reduced FA Model and the Full OLS Model

When reducing the number of factors from 10 to 3, the R-squared and Adjusted R-squared values for the Factor Analysis model decrease significantly (from ~0.48 to ~0.35). This indicates a trade-off: while reducing dimensionality successfully addresses multicollinearity, retaining too few factors can lead to information loss and reduced predictive accuracy. Lower MSE and RMSE values still suggest better predictive performance for the full OLS model in this specific comparison, as it retains more information.

90.7. Summary: Comparing OLS, PCA, and Factor Analysis Models

Multicollinearity is a common issue in regression models that can lead to unstable and difficult-to-interpret coefficients. Both Principal Component Analysis (PCA) and Factor Analysis (FA) are powerful techniques for addressing multicollinearity and reducing dimensionality.

- **PCA** is a standard method for addressing multicollinearity by transforming correlated variables into uncorrelated principal components. These components can be effectively used in linear regression and other models like Random Forest. While PCA components are not always easy to interpret directly in terms of original variables, they excel at data compression and reducing model complexity.
- **Factor Analysis** provides a way to simplify data by identifying underlying latent structures (factors) that explain correlations among variables. It also results in uncorrelated factors, making it suitable for regression problems affected by multicollinearity. Interpretation of factors relies on factor loadings.

The choice between PCA and Factor Analysis depends on the specific goals: PCA for dimensionality reduction and variance explanation, FA for discovering latent constructs. Both are valuable tools in the data scientist's toolkit for handling complex, highly correlated datasets.

90.7.1. Interpretation of the Regression Models

- **OLS Model (`model`):** This model uses the original variables directly. Coefficients indicate the direct relationship between each original variable and the target variable.

- **PCA Regression Model (`model_pca`):** This model uses principal components, which are linear combinations of the original variables, as predictors. The coefficients show the relationship between the target variable and these abstract components.
- **Factor Analysis Model (`model_factors`):** This model uses extracted factors, which are also linear combinations of original variables, designed to capture underlying latent structures. Coefficients indicate the relationship between the target variable and these latent factors.

90.7.2. Differences Compared to the Standard OLS Model

Feature	OLS Model (Standard)	PCA Regression Model	Factor Analysis Model
Input Variables	Uses original variables (e.g., <code>X_encoded</code>) as predictors.	Uses principal components (e.g., <code>df_pca_components</code>) as predictors.	Uses extracted factors (e.g., <code>df_factors</code>) as predictors.
Multicollinearity	Can suffer from multicollinearity if predictors are highly correlated, leading to unstable coefficients and inflated standard errors.	Reduces multicollinearity because principal components are orthogonal (uncorrelated).	Reduces multicollinearity by using uncorrelated factors as predictors.
Interpretability	Coefficients correspond directly to original variables, making interpretation straightforward.	Coefficients relate to abstract principal components, making direct interpretation of original variable influence more challenging.	Coefficients relate to abstract factors, making interpretation more challenging. Factor loadings must be analyzed for meaning.
Dimensionality	Uses all original variables, potentially including redundant or irrelevant features.	Reduces the number of predictors by combining original variables into fewer principal components.	Reduces the number of predictors by combining original variables into fewer factors.
Purpose	Direct relationship modeling, inference.	Dimensionality reduction, variance maximization, multicollinearity mitigation.	Discovering latent structures, explaining correlations.
Assumptions	None on underlying structure beyond linear relationship.	Does not assume an underlying causal model.	Assumes observed variables are caused by underlying factors.

Fea- ture	OLS Model (Standard)	PCA Regression Model	Factor Analysis Model
Er- ror Vari- ance	Does not explicitly separate unique variance.	Does not separate unique variance from common variance.	Explicitly models unique variance for each variable.

90.7.3. Key Differences Between Loading Scores (PCA) and Factor Loadings (FA)

Aspect	Loading Scores (PCA)	Factor Loadings (FA)
Context	Principal Component Analysis (PCA)	Factor Analysis (FA)
Purpose	Describe the contribution of variables to principal components.	Describe the relationship between variables and latent factors.
Underlying Model	No assumption of latent structure; purely variance-based.	Assumes a latent structure explaining observed variables.
Error Term	PCA does not explicitly model error variance.	FA explicitly models unique (error) variance for each variable.
Inter- pretability	Components are orthogonal (uncorrelated).	Factors may not be orthogonal, depending on rotation.

While both loading scores and factor loadings describe relationships between variables and derived components or factors, **loading scores** are specific to PCA and focus on maximizing variance, while **factor loadings** are specific to FA and aim to uncover latent structures.

90.7.4. Advantages of Using PCA and FA

Principal Component Analysis (PCA):

- **Reduced Multicollinearity:** By using uncorrelated principal components, the model avoids instability caused by multicollinearity.
- **Dimensionality Reduction:** The model uses fewer predictors if desired, improving computational efficiency and potentially generalization by removing noise.
- **Variance Maximization:** Components are constructed to capture the maximum possible variance from the original data.

Factor Analysis (FA):

- **Reduced Multicollinearity:** Similar to PCA, using uncorrelated factors prevents instability from multicollinearity.
- **Dimensionality Reduction:** Reduces the number of predictors, improving computational efficiency and generalization.
- **Focus on Underlying Structure:** Factor analysis aims to capture the latent structure of the data, potentially providing better insights into the fundamental relationships between variables.

90.7.5. Disadvantages of Using PCA and FA

Principal Component Analysis (PCA):

- **Loss of Interpretability:** Principal components are abstract combinations of the original variables, making it harder to directly interpret the coefficients. Understanding individual variable influence requires examining loading scores.
- **Potential Information Loss:** If too few components are retained, information from the original variables may be lost, potentially reducing predictive accuracy.

Factor Analysis (FA):

- **Loss of Interpretability:** Factors are abstract combinations of the original variables, making it harder to directly interpret the coefficients. Factor loadings must be analyzed to understand the influence of individual variables.
- **Potential Information Loss:** If too few factors are retained, information from the original variables may be lost, reducing predictive accuracy.
- **Complexity:** The process of extracting factors and interpreting their meaning adds complexity to the modeling process.
- **Dependence on Factor Selection:** The number of factors to retain is subjective and can affect model performance. Too few factors may oversimplify, while too many may reintroduce multicollinearity.
- **Assumption of Latent Structure:** Relies on the assumption that an underlying latent structure exists, which may not always be true for all datasets.

90.7.6. When to Use Which Method

- Use PCA when the primary goal is **dimensionality reduction**, **data compression**, and **multicollinearity resolution**, especially when interpretability of the new components is secondary to predictive performance.
- Use Factor Analysis when the goal is to **uncover underlying latent constructs or factors** that explain the relationships among variables, and when you seek to understand the conceptual meaning of these latent variables, even if it adds complexity.

- The original OLS model is preferable when **interpretability of original variables is crucial** and multicollinearity is not a significant issue.

90.8. Using Principal Components / Factors in Other Models

The principal components from PCA or factors from Factor Analysis can also be effectively used as predictors in other machine learning models, not just linear regression.

90.8.1. Random Forest Regressor with the Full Dataset

First, a Random Forest Regressor is trained using the original, full dataset (`X_encoded`).

```
# 1. Prepare Data #
# Use the original input features (X_encoded) as predictors
X_original = X_encoded
# Split the data into training and testing sets
X_train_orig, X_test_orig, y_train_orig, y_test_orig = train_test_split(X_original, y

# 2. Fit Random Forest Model
from sklearn.ensemble import RandomForestRegressor
rf_model_orig = RandomForestRegressor(n_estimators=100, random_state=42)
# Train the model on the training data
rf_model_orig.fit(X_train_orig, y_train_orig)

# 3. Evaluate the Model
# Make predictions on the test set
y_pred_orig = rf_model_orig.predict(X_test_orig)
# Calculate evaluation metrics
r2_rf_orig = r2_score(y_test_orig, y_pred_orig)
mse_rf_orig = mean_squared_error(y_test_orig, y_pred_orig)
rmse_rf_orig = np.sqrt(mse_rf_orig)

# Print the results
print("\nRandom Forest Model (using original data):")
print(f"R-squared: {r2_rf_orig:.4f}")
print(f"MSE: {mse_rf_orig:.4f}")
print(f"RMSE: {rmse_rf_orig:.4f}")
```

Random Forest Model (using original data):
 R-squared: 0.3987
 MSE: 1.3218
 RMSE: 1.1497

90.8.2. Random Forest Regressor with PCA Components

Next, a Random Forest Regressor is trained using the principal components derived from PCA. This tests if the dimensionality reduction and multicollinearity resolution of PCA benefit non-linear models.

```
# 1. Prepare Data
# Use the extracted PCA components as predictors (using the 10 components)
X_pca_rf = df_pca_components

# Split the data into training and testing sets
X_train_pca_rf, X_test_pca_rf, y_train_pca_rf, y_test_pca_rf = train_test_split(X_pca_rf, y, test_si

# 2. Fit Random Forest Model
# Initialize the Random Forest Regressor
rf_model_pca = RandomForestRegressor(n_estimators=100, random_state=42)
# Train the model on the training data
rf_model_pca.fit(X_train_pca_rf, y_train_pca_rf)

# 3. Evaluate the Model
# Make predictions on the test set
y_pred_pca_rf = rf_model_pca.predict(X_test_pca_rf)
# Calculate evaluation metrics
r2_rf_pca = r2_score(y_test_pca_rf, y_pred_pca_rf)
mse_rf_pca = mean_squared_error(y_test_pca_rf, y_pred_pca_rf)
rmse_rf_pca = np.sqrt(mse_rf_pca)

# Print the results
print("\nRandom Forest Model (using PCA components):")
print(f"R-squared: {r2_rf_pca:.4f}")
print(f"MSE: {mse_rf_pca:.4f}")
print(f"RMSE: {rmse_rf_pca:.4f}")
```

Random Forest Model (using PCA components):
 R-squared: 0.2871
 MSE: 1.5670
 RMSE: 1.2518

90.8.3. Random Forest Regressor with Extracted Factors (from FA)

Finally, a Random Forest Regressor is trained using the extracted factors from Factor Analysis (using the 3 factors from the reduced model for this example to illustrate potential impact of reduction).

```
# 1. Prepare Data
# Use the extracted factors as predictors (using the 3 factors from the reduced FA model)
X_factors_rf = df_factors_reduced

# Split the data into training and testing sets
X_train_fa_rf, X_test_fa_rf, y_train_fa_rf, y_test_fa_rf = train_test_split(X_factors_rf, y,
                                                                              test_size=0.3,
                                                                              random_state=42)

# 2. Fit Random Forest Model
# Initialize the Random Forest Regressor
rf_model_fa = RandomForestRegressor(n_estimators=100, random_state=42)
# Train the model on the training data
rf_model_fa.fit(X_train_fa_rf, y_train_fa_rf)

# 3. Evaluate the Model
# Make predictions on the test set
y_pred_fa_rf = rf_model_fa.predict(X_test_fa_rf)
# Calculate evaluation metrics
r2_rf_fa = r2_score(y_test_fa_rf, y_pred_fa_rf)
mse_rf_fa = mean_squared_error(y_test_fa_rf, y_pred_fa_rf)
rmse_rf_fa = np.sqrt(mse_rf_fa)

# Print the results
print("\nRandom Forest Model (using extracted factors):")
print(f"R-squared: {r2_rf_fa:.4f}")
print(f"MSE: {mse_rf_fa:.4f}")
print(f"RMSE: {rmse_rf_fa:.4f}")
```

```
Random Forest Model (using extracted factors):
R-squared: 0.2901
MSE: 1.5605
RMSE: 1.2492
```

90.8.4. Comparison of the Random Forest Models

```
# Print comparison of Random Forest models
print("\nComparison of Random Forest Models:")
print("\nUsing Original Data:")
print(f"R-squared: {r2_rf_orig:.4f}")
print(f"MSE: {mse_rf_orig:.4f}")
print(f"RMSE: {rmse_rf_orig:.4f}")

print("\nUsing PCA Components:")
print(f"R-squared: {r2_rf_pca:.4f}")
print(f"MSE: {mse_rf_pca:.4f}")
print(f"RMSE: {rmse_rf_pca:.4f}")

print("\nUsing Extracted Factors (from FA):")
print(f"R-squared: {r2_rf_fa:.4f}")
print(f"MSE: {mse_rf_fa:.4f}")
print(f"RMSE: {rmse_rf_fa:.4f}")
```

Comparison of Random Forest Models:

Using Original Data:

R-squared: 0.3987

MSE: 1.3218

RMSE: 1.1497

Using PCA Components:

R-squared: 0.2871

MSE: 1.5670

RMSE: 1.2518

Using Extracted Factors (from FA):

R-squared: 0.2901

MSE: 1.5605

RMSE: 1.2492

In this example, for Random Forest, using the reduced set of 3 factors from PCA and FA led to a decrease in R-squared and an increase in MSE/RMSE compared to using the original variables. This highlights that while dimensionality reduction can be beneficial, choosing too few components or factors can lead to information loss, negatively impacting predictive performance.

90.9. Videos: Principal Component Analysis (PCA)

- Video: Principal Component Analysis (PCA), Step-by-Step
- Video: PCA - Practical Tips
- Video: PCA in Python

90.10. Jupyter Notebook

Note

- The Jupyter-Notebook of this lecture is available on GitHub in the Hyperparameter-Tuning-Cookbook Repository

91. Regression

91.1. Supervised and Unsupervised Learning

Two important types: supervised and unsupervised learning. There is even more, e.g., semi-supervised learning.

91.1.1. Starting point

- Outcome measurement Y (dependent variable, response, target).
- Vector of p predictor measurements X (inputs, regressors, covariates, features, independent variables).
- Training data $(x_1, y_1), \dots, (x_N, y_N)$. These are observations (examples, instances) of these measurements.

In the *regression* problem, Y is quantitative (e.g., price, blood pressure). In the *classification* problem, Y takes values in a finite, unordered set (e.g., survived/died, digit 0-9, cancer class of tissue sample).

91.1.2. Philosophy

It is important to understand the ideas behind the various techniques, in order to know how and when to use them. One has to understand the simpler methods first, in order to grasp the more sophisticated ones. It is important to accurately assess the performance of a method, to know how well or how badly it is working (simpler methods often perform as well as fancier ones!) This is an exciting research area, having important applications in science, industry and finance. Statistical learning is a fundamental ingredient in the training of a modern data scientist.

91.1.3. Supervised Learning

Objectives of supervised learning: On the basis of the training data we would like to:

- Accurately predict unseen test cases.
- Understand which inputs affect the outcome, and how.
- Assess the quality of our predictions and inferences.

91. Regression

Note: Supervised means Y is known.

Exercise 91.1.

- Do children learn supervised?
- When do you learn supervised?
- Can learning be unsupervised?

91.1.4. Unsupervised Learning

No outcome variable, just a set of predictors (features) measured on a set of samples. The objective is more fuzzy—find groups of samples that behave similarly, find features that behave similarly, find linear combinations of features with the most variation. It is difficult to know how well you are doing. Unsupervised learning different from supervised learning, but can be useful as a pre-processing step for supervised learning. Clustering and principle component analysis are important techniques.

Unsupervised: Y is unknown, there is no Y , no trainer, no teacher, but: distances between the inputs values (features). A distance (or similarity) measure is necessary.

91.1.4.0.1. Statistical Learning

We consider supervised learning first. Sales figures from a marketing campaign. Trend shown using regression. First seems to be stronger than the third.

Can we predict $Y = \text{Sales}$ using these three? Perhaps we can do better using a model

$$Y = \text{Sales} \approx f(X_1 = \text{TV}, X_2 = \text{Radio}, X_3 = \text{Newspaper})$$

modeling the joint relationship.

Here Sales is a response or target that we wish to predict. We generically refer to the response as Y . TV is a feature, or input, or predictor; we name it X_1 . Likewise name Radio as X_2 , and so on. We can refer to the input vector collectively as

$$X = \begin{pmatrix} X_1 \\ X_2 \\ X_3 \end{pmatrix}$$

Now we write our model as

$$Y = f(X) + \epsilon$$

where ϵ captures measurement errors and other discrepancies.

What is f good for? With a good f we can make predictions of Y at new points $X = x$. We can understand which components of $X = (X_1, X_2, \dots, X_p)$ are important in explaining Y , and which are irrelevant.

For example, Seniority and Years of Education have a big impact on Income, but Marital Status typically does not. Depending on the complexity of f , we may be able to understand how each component X_j of X affects Y .

91.2. Linear Regression

91.2.1. The main ideas of fitting a line to data (The main ideas of least squares and linear regression.)

- Video: The main ideas of fitting a line to data (The main ideas of least squares and linear regression.)

91.2.1.1. Linear Regression

- Video: Linear Regression, Clearly Explained

91.2.1.2. Multiple Regression

- Video: Multiple Regression, Clearly Explained

91.2.1.3. A Gentle Introduction to Machine Learning

- Video: A Gentle Introduction to Machine Learning

91.2.1.4. Regression Function

Is there an ideal $f(X)$? In particular, what is a good value for $f(X)$ at any selected value of X , say $X = 4$? There can be many Y values at $X = 4$. A good value is

$$f(4) = E(Y|X = 4).$$

$E(Y|X = 4)$ means **expected value** (average) of Y given $X = 4$.

The ideal $f(x) = E(Y|X = x)$ is called the **regression function**. Read: The regression function gives the conditional expectation of Y given X .

The regression function $f(x)$ is also defined for the vector X ; e.g., $f(x) = f(x_1, x_2, x_3) = E(Y|X_1 = x_1, X_2 = x_2, X_3 = x_3)$.

91.2.2. Optimal Predictor

The regression function is the **ideal** or **optimal predictor** of Y with regard to mean-squared prediction error: It means that $f(x) = E(Y|X = x)$ is the function that minimizes

$$E[(Y - g(X))^2|X = x]$$

over all functions g at all points $X = x$.

91.2.2.1. Residuals, Reducible and Irreducible Error

At each point X we make mistakes:

$$\epsilon = Y - f(x)$$

is the **residual**. Even if we knew $f(x)$, we would still make errors in prediction, since at each $X = x$ there is typically a distribution of possible Y values. For any estimate $\hat{f}(x)$ of $f(x)$, we have

$$E[(Y - \hat{f}(X))^2|X = x] = [f(x) - \hat{f}(x)]^2 + \text{var}(\epsilon),$$

and $[f(x) - \hat{f}(x)]^2$ is the **reducible** error, because it depends on the model (changing the model f might reduce this error), and $\text{var}(\epsilon)$ is the **irreducible** error.

91.2.2.2. Local Regression (Smoothing)

Typically we have few if any data points with $X = 4$ exactly. So we cannot compute $E(Y|X = x)$! Idea: Relax the definition and let

$$\hat{f}(x) = \text{Ave}(Y|X \in N(x)),$$

where $N(x)$ is some neighborhood of x .

Nearest neighbor averaging can be pretty good for small p , i.e., $p \leq 4$ and large-ish N . We will discuss smoother versions, such as kernel and spline smoothing later in the course.

91.2.3. Curse of Dimensionality and Parametric Models

Local, e.g., nearest neighbor, methods can be lousy when p is large. Reason: **the curse of dimensionality**, i.e., nearest neighbors tend to be far away in high dimensions. We need to get a reasonable fraction of the N values of y_i to average to bring the variance down—e.g., 10%. A 10% neighborhood in high dimensions need no longer be

local, so we lose the spirit of estimating $E(Y|X = x)$ by local averaging. If the curse of dimensionality does not exist, nearest neighbor models would be perfect prediction models.

We will use structured (parametric) models to deal with the curse of dimensionality. The linear model is an important example of a parametric model:

$$f_L(X) = \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p.$$

A linear model is specified in terms of $p + 1$ parameters $\beta_1, \beta_2, \dots, \beta_p$. We estimate the parameters by fitting the model to *training data*. Although it is almost never correct, a linear model often serves as a good and interpretable approximation to the unknown true function $f(X)$.

The linear model is avoiding the curse of dimensionality, because it is not relying on any local properties. Linear models belong to the class of *model-based* approaches: they replace the problem of estimating f with estimating a fixed set of coefficients β_i , with $i = 1, 2, \dots, p$.

A linear model

$$\hat{f}_L(X) = \hat{\beta}_0 + \hat{\beta}_1 X$$

gives a reasonable fit. A quadratic model

$$\hat{f}_Q(X) = \hat{\beta}_0 + \hat{\beta}_1 X + \hat{\beta}_2 X^2$$

gives a slightly improved fit. The linear regression model

$$\hat{f}(\text{education}, \text{seniority}) = \hat{\beta}_0 + \hat{\beta}_1 \times \text{education} + \hat{\beta}_2 \times \text{seniority}$$

captures the important information. But it does not capture everything. More flexible regression model

$$\hat{f}_S(\text{education}, \text{seniority})$$

fit to the simulated data. Here we use a technique called a **thin-plate spline** to fit a flexible surface. Even more flexible spline regression model

$$\hat{f}_S(\text{education}, \text{seniority})$$

fit to the simulated data. Here the fitted model makes no errors on the training data! Also known as **overfitting**.

91.2.3.1. Trade-offs

- Prediction accuracy versus interpretability: Linear models are easy to interpret; thin-plate splines are not.
- Good fit versus over-fit or under-fit: How do we know when the fit is just right?
- Parsimony (Occam's razor) versus black-box: We often prefer a simpler model involving fewer variables over a black-box predictor involving them all.

91.2.4. Assessing Model Accuracy and Bias-Variance Trade-off

Suppose we fit a model $f(x)$ to some training data $Tr = \{x_i, y_i\}_1^N$, and we wish to see how well it performs. We could compute the average squared prediction error over Tr :

$$MSE_{Tr} = Ave_{i \in Tr} [y_i - \hat{f}(x_i)]^2.$$

This may be biased toward more overfit models. Instead we should, if possible, compute it using fresh **test data** $Te = \{x_i, y_i\}_1^N$:

$$MSE_{Te} = Ave_{i \in Te} [y_i - \hat{f}(x_i)]^2.$$

The red curve, which illustrated the test error, can be estimated by holding out some data to get the test-data set.

91.2.4.1. Bias-Variance Trade-off

Suppose we have fit a model $f(x)$ to some training data Tr , and let (x_0, y_0) be a test observation drawn from the population. If the true model is

$$Y = f(X) + \epsilon \quad \text{with } f(x) = E(Y|X = x),$$

then

$$E(y_0 - \hat{f}(x_0))^2 = \text{var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{var}(\epsilon). \quad (91.1)$$

Here, $\text{var}(\epsilon)$ is the irreducible error. The reducible error consists of two components:

- $\text{var}(\hat{f}(x_0))$ is the variance that comes from different training sets. Different training sets result in different functions $\hat{f}$.
- $\text{Bias}(\hat{f}(x_0)) = E[\hat{f}(x_0)] - f(x_0)$.

The expectation averages over the variability of y_0 as well as the variability in Tr . Note that

$$\text{Bias}(\hat{f}(x_0)) = E[\hat{f}(x_0)] - f(x_0).$$

Typically as the flexibility of $\hat{f}$ increases, its variance increases (because the fits differ from training set to training set), and its bias decreases. So choosing the flexibility based on average test error amounts to a bias-variance trade-off.

If we add the two components (reducible and irreducible error), we get the MSE in Equation 91.1.

- Video: Machine Learning Fundamentals: Bias and Variance

91.3. Multiple Regression

91.4. R-squared

91.4.1. R-Squared in Simple Linear Regression

In simple linear regression, the relationship between the independent variable X and the dependent variable Y is modeled using the equation:

$$Y = \beta_0 + \beta_1 X + \epsilon$$

Here, β_0 is the intercept, β_1 is the slope or regression coefficient, and ϵ is the error term.

Definition 91.1 (R-Squared (R^2)). R^2 is a measure of how well the regression model explains the variance in the dependent variable. It is calculated as the square of the correlation coefficient (r) between the actual values Y and the predicted values $\hat{Y}$ from the regression model. It ranges from 0 to 1, where:

- 1 indicates that the regression predictions perfectly fit the data.
- 0 indicates that the model does not explain any of the variability in the target data around its mean.

In simple linear regression, where there is one independent variable X and one dependent variable Y , the R-squared (R^2) is the square of the Pearson correlation coefficient (r) between the observed values of the dependent variable and the values predicted by the regression model. That is, in simple linear regression, we have

$$R^2 = r^2.$$

This equivalence holds specifically for simple linear regression due to the direct relationship between the linear fit and the correlation of two variables. In multiple linear regression, while R^2 still represents the proportion of variance explained by the model, it is not simply the square of a single correlation coefficient as it involves multiple predictors.

- Video: R-squared, Clearly Explained

91.5. Assessing Confounding Effects in Multiple Regression

Confounding is a bias introduced by the imbalanced distribution of extraneous risk factors among comparison groups (Wang 2007). `spotoptim` provides tools for assessing confounding effects in multiple regression models.

Example 91.1 (Assessing Confounding Effects in Multiple Regression with `spotoptim`). Consider the following data generation function `generate_data` and the `fit_ols_model` function to fit an ordinary least squares (OLS) regression model.

```
import numpy as np
import pandas as pd
import statsmodels.formula.api as smf

def generate_data(n_samples=100, b0=0, b1=-1, b2=0, b3=10, b12=0, b13=0, b23=0, b123=0,
                noise_std=1):
    """
    Generate data for the linear formula  $y \sim b_0 + b_1x_1 + b_2x_2 + b_3x_3 + b_{12}x_1x_2 + b_{13}x_1x_3 + b_{23}x_2x_3 + b_{123}x_1x_2x_3$ .

    Args:
        n_samples (int): Number of samples to generate.
        b0 (float): Coefficient for the intercept.
        b1 (float): Coefficient for  $x_1$ .
        b2 (float): Coefficient for  $x_2$ .
        b3 (float): Coefficient for  $x_3$ .
        b12 (float): Coefficient for the interaction term  $x_1x_2$ .
        b13 (float): Coefficient for the interaction term  $x_1x_3$ .
        b23 (float): Coefficient for the interaction term  $x_2x_3$ .
        b123 (float): Coefficient for the interaction term  $x_1x_2x_3$ .
        noise_std (float): Standard deviation of the Gaussian noise added to  $y$ .

    Returns:
        pd.DataFrame: A DataFrame containing the generated data with columns ['x1', 'x2', 'x3', 'y'].
    """
    np.random.seed(42) # For reproducibility
    x1 = np.random.uniform(0, 1, n_samples)
    x2 = np.random.uniform(0, 1, n_samples)
    x3 = np.random.uniform(0, 1, n_samples)

    y = (b0 + b1*x1 + b2*x2 + b3*x3 + b12*x1*x2 + b13*x1*x3 + b23*x2*x3 + b123*x1*x2*x3 +
         np.random.normal(0, noise_std, n_samples))

    data = pd.DataFrame({'y': y, 'x1': x1, 'x2': x2, 'x3': x3})
    return data
```

```
def fit_ols_model(formula, data) -> dict:
    """
    Fit an OLS model using the given formula and data, and print the results.

    Args:
        formula (str): The formula for the OLS model.
        data (pd.DataFrame): The data frame containing the variables.

    Returns:
        dict: A dictionary containing the p-values, estimates, confidence intervals, and AIC value.
    """
    mod_0 = smf.ols(formula=formula, data=data).fit()
    p = mod_0.pvalues.iloc[1]
    estimate = mod_0.params.iloc[1]
    conf_int = mod_0.conf_int().iloc[1]
    aic_value = mod_0.aic

    print(f"p-values: {p}")
    print(f"estimate: {estimate}")
    print(f"conf_int: {conf_int}")
    print(f"aic: {aic_value}")
```

These functions can be used to generate data and fit an OLS model. Here we use the model

$$y = f(x_1, x_2, x_3) + \epsilon = x_1 + 10x_3 + \epsilon.$$

We set up the basic model $y_0 = f_0(x_1)$ and analyze how the model fit changes when adding x_2 and x_3 to the model. If the p -values are decreasing by adding a variable, this indicates that the variable is relevant for the model. Similarly, if the p -values are increasing by removing a variable, this indicates that the variable is not relevant for the model.

```
data = generate_data(b0=0, b1=1, b2=0, b3=10, b12=0, b13=0, b23=0, b123=0, noise_std=1)
fit_ols_model("y ~ x1", data)
fit_ols_model("y ~ x1 + x2", data)
fit_ols_model("y ~ x1 + x3", data)
fit_ols_model("y ~ x1 + x2 + x3", data)
```

```
p-values: 0.34343741859526244
estimate: 1.025306391110114
conf_int: 0    -1.111963
1      3.162575
Name: x1, dtype: float64
```

91. Regression

```
aic: 517.6397392012537
p-values: 0.36375118507784565
estimate: 0.9810502049698092
conf_int: 0    -1.152698
1      3.114798
Name: x1, dtype: float64
aic: 518.1426513151566
p-values: 4.946760674421827e-05
estimate: 1.407792346942117
conf_int: 0      0.750106
1      2.065479
Name: x1, dtype: float64
aic: 282.73524524532
p-values: 4.849840959643549e-05
estimate: 1.4159292625696243
conf_int: 0      0.755494
1      2.076364
Name: x1, dtype: float64
aic: 284.34665447613634
```

The function `fit_all_lm()` simplifies this procedure. It can be used to fit all possible linear models with the given data and print the results in a systematic way for various combinations of variables.

```
from spotoptim.utils.stats import fit_all_lm, plot_coeff_vs_pvals, plot_coeff_vs_pvals
res = fit_all_lm("y ~ x1", ["x2", "x3"], data)
print(res["estimate"])
```

The basic model is: `y ~ x1`
The following features will be used for fitting the basic model: `Index(['y', 'x3', 'x2'])`
p-values: 0.34343741859526244
estimate: 1.025306391110114
conf_int: 0 -1.111963
1 3.162575
Name: x1, dtype: float64
aic: 517.6397392012537
Combinations: `[('x2',), ('x3',), ('x2', 'x3')]`

	variables	estimate	conf_low	conf_high	p	aic	n
0	basic	1.025306	-1.111963	3.162575	0.343437	517.639739	100
1	x2	0.981050	-1.152698	3.114798	0.363751	518.142651	100
2	x3	1.407792	0.750106	2.065479	0.000049	282.735245	100
3	x2, x3	1.415929	0.755494	2.076364	0.000048	284.346654	100

Interpreting the results, we can see that the p -values decrease when adding x_3 (as well as both x_2 and x_3) to the model, indicating that x_3 is relevant for the model. Adding only x_2 does not significantly improve the model fit.

In addition to the textual output, the function `plot_coeff_vs_pvals_by_included()` can be used to visualize the coefficients and p -values of the fitted models.

```
plot_coeff_vs_pvals_by_included(res)
```




Figure 91.1.: Coefficients vs. p -values for different models. The right plot indicates that x_3 should be included in the model, whereas the left plot shows that x_2 is not relevant.

Figure 91.1 shows the coefficients and p -values for different models. Because y depends on x_1 and x_3 , the p -value much smaller if x_3 is included in the model as can be seen in the right plot in Figure 91.1. The left plot shows that including x_2 in the model does not significantly improve the model fit.

91.6. Cross-Validation

- Video: Machine Learning Fundamentals: Cross Validation

92. Classification

In classification we have a qualitative response variable.

Here the response variable Y is qualitative, e.g., email is one of $C = (\text{spam}, \text{ham})$, where ham is good email, digit class is one of $C = \{0, 1, \dots, 9\}$. Our goals are to:

- Build a classifier $C(X)$ that assigns a class label from C to a future unlabeled observation X .
- Assess the uncertainty in each classification
- Understand the roles of the different predictors among $X = (X_1, X_2, \dots, X_p)$.

Simulation example depicted in @fig-0218a. Y takes two values, zero and one, and X has only one value. Big sample: each single vertical bar indicates an occurrence of a zero (orange) or one (blue) as a function of the X s. Black curve generated the data: it is the probability of generating a one. For high values of X , the probability of ones is increasing. What is an ideal classifier $C(X)$?

Suppose the K elements in C are numbered $1, 2, \dots, K$. Let

$$p_k(x) = \Pr(Y = k | X = x), k = 1, 2, \dots, K.$$

These are the **conditional class probabilities** at x ; e.g. see little barplot at $x = 5$. Then the **Bayes optimal classifier** at x is

$$C(x) = j \quad \text{if } p_j(x) = \max\{p_1(x), p_2(x), \dots, p_K(x)\}.$$

At $x = 5$ there is an 80% probability of one, and an 20% probability of a zero. So, we classify this point to the class with the highest probability, the majority class.

Nearest-neighbor averaging can be used as before. This is illustrated in Fig.~??. Here, we consider 100 points only. Nearest-neighbor averaging also breaks down as dimension grows. However, the impact on $\hat{C}(x)$ is less than on $\hat{p}_k(x)$, $k = 1, \dots, K$.

92.1. Classification: Some Details

Average number of errors made to measure the performance. Typically we measure the performance of $\hat{C}(x)$ using the **misclassification error rate**:

$$Err_{Te} = Ave_{i \in Te} I[y_i \neq \hat{C}(x_i)].$$

The Bayes classifier (using the true $p_k(x)$) has smallest error (in the population).

92.2. k-Nearest Neighbor Classification

Consider k-nearest neighbors in two dimensions. Orange and blue dots label the true class memberships of the underlying points in the 2-dim plane. Dotted line is the decision boundary, that is the contour with equal probability for both classes.

Nearest-neighbor averaging in 2-dim. At any given point we want to classify, we spread out a little neighborhood, say $K = 10$ points from the neighborhood and calculated the percentage of blue and orange. We assign the color with the highest probability to this point. If this is done for every point in the plane, we obtain the solid black curve as the estimated decision boundary.

We can use $K = 1$. This is the **nearest-neighbor classifier**. The decision boundary is piecewise linear. Islands occur. Approximation is rather noisy.

$K = 100$ leads to a smooth decision boundary. But gets uninteresting.

K large means higher bias, so $1/K$ is chosen, because we go from low to high complexity on the x -error. Horizontal dotted line is the base error.

92.2.1. Minkowski Distance

The Minkowski distance of order p (where p is an integer) between two points $X = (x_1, x_2, \dots, x_n)$ and $Y = (y_1, y_2, \dots, y_n) \in \mathbb{R}^n$ is defined as:

$$D(X, Y) = \left(\sum_{i=1}^n |x_i - y_i|^p \right)^{\frac{1}{p}}.$$

- Video: StatQuest: K-nearest neighbors, Clearly Explained

92.3. Decision and Classification Trees

92.3.1. Decision Trees

92.3.1.1. Decision and Classification Trees, Clearly Explained

92.3.1.2. StatQuest: Decision Trees, Part 2 - Feature Selection and Missing Data

92.3.2. Regression Trees

92.3.2.1. Regression Trees, Clearly Explained!!!

92.3.2.2. How to Prune Regression Trees, Clearly Explained!!!

92.4. The Confusion Matrix

- Video: Machine Learning Fundamentals: The Confusion Matrix

92.4.1. Sensitivity and Specificity

- Video: Machine Learning Fundamentals: Sensitivity and Specificity

92.5. Naive Bayes

- Video: Naive Bayes, Clearly Explained!!!

92.6. Gaussian Naive Bayes

- Video: Gaussian Naive Bayes, Clearly Explained!!!

93. Clustering

93.1. DBSCAN

- Video: Clustering with DBSCAN, Clearly Explained!!!

93.2. k-Means Clustering

The k -means algorithm is an unsupervised learning algorithm that has a loose relationship to the k -nearest neighbor classifier. The k -means algorithm works as follows:

- Step 1: Randomly choose k centers. Assign points to cluster.
- Step 2: Determine the distances of each data point to the centroids and re-assign each point to the closest cluster centroid based upon minimum distance
- Step 3: Calculate cluster centroids again
- Step 4: Repeat steps 2 and 3 until we reach global optima where no improvements are possible and no switching of data points from one cluster to other.
- Video: K-means clustering

93.3. DDMO-Additional Videos

- Odds and Log(Odds), Clearly Explained!!!
- One-Hot, Label, Target and K-Fold Target Encoding, Clearly Explained!!!
- Maximum Likelihood for the Exponential Distribution, Clearly Explained!!!
- ROC and AUC, Clearly Explained!
- Entropy (for data science) Clearly Explained!!!
- Classification Trees in Python from Start to Finish: Long live video!

93.4. DDMO-Exercises

Exercise 93.1 (Smaller Bins). What happens when we use smaller bins in a histogram?

Exercise 93.2 (Density Curve). Why plot a curve to approximate a histogram?

Exercise 93.3 (TwoSDQuestion). How many samples are plus/minus two SD around the mean?

Exercise 93.4 (OneSDQuestion). How many samples are plus/minus one SD around the mean?

Exercise 93.5 (ThreeSDQuestion). How many samples are plus/minus three SD around the mean?

Exercise 93.6 (DataRangeQuestion). You have a mean at 100 and a SD of 10. Where are 95% of the data?

Exercise 93.7 (PeakHeightQuestion). If the peak is very high, is the SD low or high?

Exercise 93.8 (ProbabilityQuestion). If we have a certain curve and want to calculate the probability of values equal to 20 if the mean is 20.

Exercise 93.9 (MeanDifferenceQuestion). The difference between μ and $\bar{x}$?

Exercise 93.10 (EstimateMeanQuestion). How do you calculate the sample mean?

Exercise 93.11 (SigmaSquaredQuestion). What is sigma squared?

Exercise 93.12 (EstimatedSDQuestion). What is the formula for the estimated standard deviation?

Exercise 93.13 (VarianceDifferenceQuestion). Difference between the variance and the estimated variance?

Exercise 93.14 (ModelBenefitsQuestion). What are the benefits of using models?

Exercise 93.15 (SampleDefinitionQuestion). What is a sample in statistics?

Exercise 93.16 (RejectHypothesisQuestion). What does it mean to reject a hypothesis?

Exercise 93.17 (NullHypothesisQuestion). What is a null hypothesis?

Exercise 93.18 (BetterDrugQuestion). How can you show that you have found a better drug?

Exercise 93.19 (PValueIntroductionQuestion). What is the reason for introducing the p-value?

Exercise 93.20 (PValueRangeQuestion). Is there any range for p-values? Can it be negative?

Exercise 93.21 (PValueRangeQuestion). Is there any range for p-values? Can it be negative?

Exercise 93.22 (TypicalPValueQuestion). What are typical values of the p-value and what does it mean? 5%?

Exercise 93.23 (FalsePositiveQuestion). What is a false-positive?

Exercise 93.24 (CalculatePValueQuestion). How to calculate p-value?

Exercise 93.25 (SDCalculationQuestion). What is the SD if the mean is 155 and in the range from 142 - 169 there are 95% of the data?

Exercise 93.26 (SidedPValueQuestion). When do we need the two-sided p-value and when the one-sided?

Exercise 93.27 (CoinTestQuestion). Test a coin with Tail-Head-Head. What is the p-value?

Exercise 93.28 (BorderPValueQuestion). If you get exactly the 0.05 border value, can you reject?

Exercise 93.29 (OneSidedPValueCautionQuestion). Why should you be careful with a one-sided p-test?

Exercise 93.30 (BinomialDistributionQuestion). What is the binomial distribution?

Exercise 93.31 (PHackingWaysQuestion). Name two typical ways of p-hacking.

Exercise 93.32 (AvoidPHackingQuestion). How can p-hacking be avoided?

Exercise 93.33 (MultipleTestingProblemQuestion). What is the multiple testing problem?

93.4.0.1. Covariance

Exercise 93.34 (CovarianceDefinitionQuestion). What is covariance?

Exercise 93.35 (CovarianceMeaningQuestion). What is the meaning of covariance?

Exercise 93.36 (CovarianceVarianceRelationshipQuestion). What is the relationship between covariance and variance?

Exercise 93.37 (HighCovarianceQuestion). If covariance is high, is there a strong relationship?

Exercise 93.38 (ZeroCovarianceQuestion). What if the covariance is zero?

Exercise 93.39 (NegativeCovarianceQuestion). Can covariance be negative?

Exercise 93.40 (NegativeVarianceQuestion). Can variance be negative?

Exercise 93.41 (CorrelationValueQuestion). What do you do if the correlation value is 10?

Exercise 93.42 (CorrelationRangeQuestion). What is the possible range of correlation values?

Exercise 93.43 (CorrelationFormulaQuestion). What is the formula for correlation?

Exercise 93.44 (UnderstandingStatisticalPower). What is the definition of power in a statistical test?

Exercise 93.45 (DistributionEffectOnPower). What is the implication for power analysis if the samples come from the same distribution?

Exercise 93.46 (IncreasingPower). How can you increase the power if the distributions are very similar?

Exercise 93.47 (PreventingPHacking). What should be done to avoid p-hacking when the distributions are close to each other?

Exercise 93.48 (SampleSizeAndPower). If there is overlap and the sample size is small, will the power be high or low?

Exercise 93.49 (FactorsAffectingPower). Which are the two main factors that affect power?

Exercise 93.50 (PurposeOfPowerAnalysis). What does power analysis tell us?

Exercise 93.51 (ExperimentRisks). What are the two risks faced when performing an experiment?

Exercise 93.52 (PerformingPowerAnalysis). How do you perform a power analysis?

Exercise 93.53 (CentralLimitTheoremExplanation). What does the Central Limit Theorem state?

Exercise 93.54 (MedianInBoxplot). What is represented by the middle line in a boxplot?

Exercise 93.55 (BoxContentInBoxplot). What does the box in a boxplot represent?

Exercise 93.56 (RSquaredDefinition). What is R-squared? Show the formula.

Exercise 93.57 (NegativeRSquared). Can the R-squared value be negative?

Exercise 93.58 (RSquaredCalculation). Perform a calculation involving R-squared.

Exercise 93.59 (LeastSquaresMeaning). What is the meaning of the least squares method?

Exercise 93.60 (RegressionVsClassification). What is the difference between regression and classification?

Exercise 93.61 (LikelihoodConcept). What is the idea of likelihood?

Exercise 93.62 (ProbabilityVsLikelihood). What is the difference between probability and likelihood?

Exercise 93.63 (TrainVsTestData). What is the difference between training and testing data?

Exercise 93.64 (SingleValidationIssue). What is the problem if you validate the model only once?

Exercise 93.65 (FoldDefinition). What is a fold in cross-validation?

Exercise 93.66 (LeaveOneOutValidation). What is leave-one-out cross-validation?

Exercise 93.67 (DrawingConfusionMatrix). Draw the confusion matrix.

93. Clustering

Exercise 93.68 (SensitivitySpecificityCalculation1). Calculate the sensitivity and specificity for a given confusion matrix.

Exercise 93.69 (SensitivitySpecificityCalculation2). Calculate the sensitivity and specificity for a given confusion matrix.

Exercise 93.70 (BiasAndVariance). What are bias and variance?

Exercise 93.71 (MutualInformationExample). Provide an example and calculate if mutual information is high or low.

Exercise 93.72 (WhatIsPCA). What is PCA?

Exercise 93.73 (ScreePlotExplanation). What is a scree plot?

Exercise 93.74 (LeastSquaresInPCA). Does PCA use least squares?

Exercise 93.75 (PCASteps). Which steps are performed by PCA?

Exercise 93.76 (EigenvaluePC1). What is the eigenvalue of the first principal component?

Exercise 93.77 (DifferencesBetweenPoints). Are the differences between red and yellow the same as the differences between red and blue points?

Exercise 93.78 (ScalingInPCA). How to scale data in PCA?

Exercise 93.79 (DetermineNumberOfComponents). How to determine the number of principal components?

Exercise 93.80 (LimitingNumberOfComponents). How is the number of principal components limited?

Exercise 93.81 (WhyUseTSNE). Why use t-SNE?

Exercise 93.82 (MainIdeaOfTSNE). What is the main idea of t-SNE?

Exercise 93.83 (BasicConceptOfTSNE). What is the basic concept of t-SNE?

Exercise 93.84 (TSNESteps). What are the steps in t-SNE?

Exercise 93.85 (HowKMeansWorks). How does K-means clustering work?

Exercise 93.86 (QualityOfClusters). How can the quality of the resulting clusters be calculated?

Exercise 93.87 (IncreasingK). Why is it not a good idea to increase k too much?

Exercise 93.88 (CorePointInDBSCAN). What is a core point in DBSCAN?

Exercise 93.89 (AddingVsExtending). What is the difference between adding and extending in DBSCAN?

Exercise 93.90 (OutliersInDBSCAN). What are outliers in DBSCAN?

Exercise 93.91 (AdvantagesAndDisadvantagesOfK). What are the advantages and disadvantages of $k = 1$ and $k = 100$ in K -nearest neighbors?

Exercise 93.92 (NaiveBayesFormula). What is the formula for Naive Bayes?

Exercise 93.93 (CalculateProbabilities). Calculate the probabilities for a given example using Naive Bayes.

Exercise 93.94 (UnderflowProblem). Why is underflow a problem in Gaussian Naive Bayes?

Exercise 93.95 (Tree Usage). For what can we use trees?

Exercise 93.96 (Tree Usage). Based on a shown tree graph:

- How can you use this tree?
- What is the root node?
- What are branches and internal nodes?
- What are the leafs?
- Are the leafs pure or impure?
- Which of the leafs is more impure?

Exercise 93.97 (Tree Feature Importance). Is the most or least important feature on top?

Exercise 93.98 (Tree Feature Imputation). How can you fill a gap/missing data?

Solution 93.1 (Tree Feature Imputation).

- Mean
- Median
- Comparing to column with high correlation

93. *Clustering*

Exercise 93.99 (Regression Tree Limitations). What are limitations?

Exercise 93.100 (Regression Tree Score). How is the tree score calculated?

Exercise 93.101 (Regression Tree Alpha Value Small). What can we say about the tree if the alpha value is small?

Exercise 93.102 (Regression Tree Increase Alpha Value). What happens if you increase alpha?

Exercise 93.103 (Regression Tree Pruning). What is the meaning of pruning?

Part XV.

Appendix

94. Gaussian Processes—Some Background Information

The concept of GP (Gaussian Process) regression can be understood as a simple extension of linear modeling. It is worth noting that this approach goes by various names and acronyms, including “kriging,” a term derived from geostatistics, as introduced by Matheron in 1963. Additionally, it is referred to as Gaussian spatial modeling or a Gaussian stochastic process, and machine learning (ML) researchers often use the term Gaussian process regression (GPR). In all of these instances, the central focus is on regression. This involves training on both inputs and outputs, with the ultimate objective of making predictions and quantifying uncertainty (referred to as uncertainty quantification or UQ).

However, it’s important to emphasize that GPs are not a universal solution for every problem. Specialized tools may outperform GPs in specific, non-generic contexts, and GPs have their own set of limitations that need to be considered.

94.1. Gaussian Process Prior

In the context of GP, any finite collection of realizations, which is represented by n observations, is modeled as having a multivariate normal (MVN) distribution. The characteristics of these realizations can be fully described by two key parameters:

1. Their mean, denoted as an n -vector μ .
2. The covariance matrix, denoted as an $n \times n$ matrix Σ . This covariance matrix encapsulates the relationships and variability between the individual realizations within the collection.

94.2. Covariance Function

The covariance function is defined by inverse exponentiated squared Euclidean distance:

$$\Sigma(\vec{x}, \vec{x}') = \exp\{-\|\vec{x} - \vec{x}'\|^2\},$$

94. Gaussian Processes—Some Background Information

where $\vec{x}$ and $\vec{x}'$ are two points in the k -dimensional input space and $\|\cdot\|$ denotes the Euclidean distance, i.e.,

$$\|\vec{x} - \vec{x}'\|^2 = \sum_{i=1}^k (x_i - x'_i)^2.$$

An 1-d example is shown in Figure 94.1.




Figure 94.1.: One-dim inverse exponentiated squared Euclidean distance

The covariance function is also referred to as the kernel function. The *Gaussian* kernel uses an additional parameter, σ^2 , to control the rate of decay. This parameter is referred to as the length scale or the characteristic length scale. The covariance function is then defined as

$$\Sigma(\vec{x}, \vec{x}') = \exp\{-\|\vec{x} - \vec{x}'\|^2 / (2\sigma^2)\}. \quad (94.1)$$

The covariance decays exponentially fast as $\vec{x}$ and $\vec{x}'$ become farther apart. Observe that

$$\Sigma(\vec{x}, \vec{x}) = 1$$

and

$$\Sigma(\vec{x}, \vec{x}') < 1$$

94.3. Construction of the Covariance Matrix

for $\vec{x} \neq \vec{x}'$. The function $\Sigma(\vec{x}, \vec{x}')$ must be positive definite.

Remark 94.1 (Kriging and Gaussian Basis Functions). The Kriging basis function (Equation 84.1) is related to the 1-dim Gaussian basis function (Equation 94.1), which is defined as

$$\Sigma(\vec{x}^{(i)}, \vec{x}^{(j)}) = \exp\{-\|\vec{x}^{(i)} - \vec{x}^{(j)}\|^2/(2\sigma^2)\}. \quad (94.2)$$

There are some differences between Gaussian basis functions and Kriging basis functions:

- Where the Gaussian basis function has $1/(2\sigma^2)$, the Kriging basis has a vector $\theta = [\theta_1, \theta_2, \dots, \theta_k]^T$.
- The θ vector allows the width of the basis function to vary from dimension to dimension.
- In the Gaussian basis function, the exponent is fixed at 2, Kriging allows this exponent p_i to vary (typically from 1 to 2).

94.2.1. Positive Definiteness

Positive definiteness in the context of the covariance matrix Σ_n is a fundamental requirement. It is determined by evaluating $\Sigma(x_i, x_j)$ at pairs of n $\vec{x}$ -values, denoted as $\vec{x}_1, \vec{x}_2, \dots, \vec{x}_n$. The condition for positive definiteness is that for all $\vec{x}$ vectors that are not equal to zero, the expression $\vec{x}^T \Sigma_n \vec{x}$ must be greater than zero. This property is essential when intending to use Σ_n as a covariance matrix in multivariate normal (MVN) analysis. It is analogous to the requirement in univariate Gaussian distributions where the variance parameter, σ^2 , must be positive.

Gaussian Processes (GPs) can be effectively utilized to generate random data that follows a smooth functional relationship. The process involves the following steps:

1. Select a set of $\vec{x}$ -values, denoted as $\vec{x}_1, \vec{x}_2, \dots, \vec{x}_n$.
2. Define the covariance matrix Σ_n by evaluating $\Sigma_n^{ij} = \Sigma(\vec{x}_i, \vec{x}_j)$ for $i, j = 1, 2, \dots, n$.
3. Generate an n -variate realization Y that follows a multivariate normal distribution with a mean of zero and a covariance matrix Σ_n , expressed as $Y \sim \mathcal{N}_n(0, \Sigma_n)$.
4. Visualize the result by plotting it in the x - y plane.

94.3. Construction of the Covariance Matrix

Here is an one-dimensional example. The process begins by creating an input grid using $\vec{x}$ -values. This grid consists of 100 elements, providing the basis for further analysis and visualization.

```
import numpy as np
n = 100
X = np.linspace(0, 10, n, endpoint=False).reshape(-1,1)
```

94.3. Construction of the Covariance Matrix

```


0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    ,
0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    ,
0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    ,
0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    ,
0.    ],
[0.995, 1.    , 0.995, 0.98 , 0.956, 0.923, 0.882, 0.835, 0.783,
0.726, 0.667, 0.607, 0.546, 0.487, 0.43 , 0.375, 0.325, 0.278,
0.236, 0.198, 0.164, 0.135, 0.11 , 0.089, 0.071, 0.056, 0.044,
0.034, 0.026, 0.02 , 0.015, 0.011, 0.008, 0.006, 0.004, 0.003,
0.002, 0.002, 0.001, 0.001, 0.    , 0.    , 0.    , 0.    , 0.    ,
0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    ,
0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    ,
0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    ,
0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    ,
0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    ,
0.    ],
[0.98 , 0.995, 1.    , 0.995, 0.98 , 0.956, 0.923, 0.882, 0.835,
0.783, 0.726, 0.667, 0.607, 0.546, 0.487, 0.43 , 0.375, 0.325,
0.278, 0.236, 0.198, 0.164, 0.135, 0.11 , 0.089, 0.071, 0.056,
0.044, 0.034, 0.026, 0.02 , 0.015, 0.011, 0.008, 0.006, 0.004,
0.003, 0.002, 0.002, 0.001, 0.001, 0.    , 0.    , 0.    , 0.    ,
0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    ,
0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    ,
0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    ,
0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    , 0.    ,
0.    ],
0.    ]])

```

```

import matplotlib.pyplot as plt
plt.imshow(Sigma, cmap='hot', interpolation='nearest')
plt.colorbar()
plt.show()

```



94.4. Generation of Random Samples and Plotting the Realizations of the Random Function

In the context of the multivariate normal distribution, the next step is to utilize the previously constructed covariance matrix denoted as **Sigma**. It is used as an essential component in generating random samples from the multivariate normal distribution.

The function `multivariate_normal` is employed for this purpose. It serves as a random number generator specifically designed for the multivariate normal distribution. In this case, the mean of the distribution is set equal to `mean`, and the covariance matrix is provided as `Psi`. The argument `size` specifies the number of realizations, which, in this specific scenario, is set to one.

By default, the mean vector is initialized to zero. To match the number of samples, which is equivalent to the number of rows in the `X` and `Sigma` matrices, the argument `zeros(n)` is used, where `n` represents the number of samples (here taken from the size of the matrix, e.g.,: `Sigma.shape[0]`).

```
rng = np.random.default_rng(seed=12345)
Y = rng.multivariate_normal(zeros(Sigma.shape[0]), Sigma, size = 1, check_valid="raise")
Y.shape
```

```
(100, 1)
```

94.4. Generation of Random Samples and Plotting the Realizations of the Random Function

Now we can plot the results, i.e., a finite realization of the random function $Y()$ under a GP prior with a particular covariance structure. We will plot those X and Y pairs as connected points on an x - y plane.

```
import matplotlib.pyplot as plt
plt.plot(X, Y)
plt.title("Realization of Random Functions under a GP prior.\n sigma2: {}".format(sigma2[0]))
plt.show()
```




Figure 94.2.: Realization of one random function under a GP prior. sigma2: 1.0

```
rng = np.random.default_rng(seed=12345)
Y = rng.multivariate_normal(zeros(Sigma.shape[0]), Sigma, size = 3, check_valid="raise")
plt.plot(X, Y.T)
plt.title("Realization of Three Random Functions under a GP prior.\n sigma2: {}".format(sigma2[0]))
plt.show()
```

Realization of Three Random Functions under a GP prior.
sigma2: 1.0




Figure 94.3.: Realization of three random functions under a GP prior. sigma2: 1.0

94.5. Properties of the 1d Example

94.5.1. Several Bumps:

In this analysis, we observe several bumps in the x -range of $[0, 10]$. These bumps in the function occur because shorter distances exhibit high correlation, while longer distances tend to be essentially uncorrelated. This leads to variations in the function's behavior:

- When x and x' are one σ unit apart, the correlation is $\exp(-\sigma^2/(2\sigma^2)) = \exp(-1/2) \approx 0.61$, i.e., a relative high correlation.
- 2σ apart means correlation $\exp(-2^2/2) \approx 0.14$, i.e., only small correlation.
- 4σ apart means correlation $\exp(-4^2/2) \approx 0.0003$, i.e., nearly no correlation—variables are considered independent for almost all practical application.

94.5.2. Smoothness:

The function plotted in Figure 94.2 represents only a finite realization, which means that we have data for a limited number of pairs, specifically 100 points. These points appear smooth in a tactile sense because they are closely spaced, and the plot function connects the dots with lines to create the appearance of smoothness. The complete surface, which can be conceptually extended to an infinite realization over a compact domain, is exceptionally smooth in a calculus sense due to the covariance function's property of being infinitely differentiable.

94.5.3. Scale of Two:

Regarding the scale of the Y values, they have a range of approximately $[-2, 2]$, with a 95% probability of falling within this range. In standard statistical terms, 95% of the data points typically fall within two standard deviations of the mean, which is a common measure of the spread or range of data.

```
import numpy as np
from numpy import array, zeros, power, ones, exp, multiply, eye, linspace, spacing, sqrt, arange, ap
from numpy.random import multivariate_normal

def build_Sigma(X, sigma2):
    n = X.shape[0]
    k = X.shape[1]
    D = zeros((k, n, n))
    for l in range(k):
        for i in range(n):
            for j in range(i, n):
                D[l, i, j] = 1/(2*sigma2[l])*(X[i,l] - X[j,l])**2
    D = sum(D)
    D = D + D.T
    return exp(-D)

def plot_mvn( a=0, b=10, sigma2=1.0, size=1, n=100, show=True):
    X = np.linspace(a, b, n, endpoint=False).reshape(-1,1)
    sigma2 = np.array([sigma2])
    Sigma = build_Sigma(X, sigma2)
    rng = np.random.default_rng(seed=12345)
    Y = rng.multivariate_normal(zeros(Sigma.shape[0]), Sigma, size = size, check_valid="raise")
    plt.plot(X, Y.T)
    plt.title("Realization of Random Functions under a GP prior.\n sigma2: {}".format(sigma2[0]))
    if show:
        plt.show()
```

```
plot_mvn(a=0, b=10, sigma2=10.0, size=3, n=250)
```




Figure 94.4.: Realization of Random Functions under a GP prior. sigma2: 10

```
plot_mvn(a=0, b=10, sigma2=0.1, size=3, n=250)
```




Figure 94.5.: Realization of Random Functions under a GP prior. sigma2: 0.1

94.6. Jupyter Notebook

i Note

- The Jupyter-Notebook of this lecture is available on GitHub in the Hyperparameter-Tuning-Cookbook Repository

References

- Arlot, Sylvain, Alain Celisse, et al. 2010. “A Survey of Cross-Validation Procedures for Model Selection.” *Statistics Surveys* 4: 40–79.
- Bartz, Eva, Thomas Bartz-Beielstein, Martin Zaefferer, and Olaf Mersmann, eds. 2022. *Hyperparameter Tuning for Machine and Deep Learning with R - A Practical Guide*. Springer.
- Box, G E P. 1957. “Evolutionary operation: A method for increasing industrial productivity.” *Applied Statistics* 6: 81–101.
- Forrester, Alexander, András Sóbester, and Andy Keane. 2008. *Engineering Design via Surrogate Modelling*. Wiley.
- Hartung, Joachim, Bärbel Elpert, and Karl-Heinz Klösener. 1995. *Statistik*. Oldenbourg.
- Hastie, Trevor, Robert Tibshirani, and Jerome Friedman. 2017. *The Elements of Statistical Learning*. Second. Springer.
- Johnson, M. E., L. M. Moore, and D. Ylvisaker. 1990. “Minimax and Maximin Distance Designs.” *Journal of Statistical Planning and Inference* 26 (2): 131–48.
- Jones, Donald R., Matthias Schonlau, and William J. Welch. 1998. “Efficient Global Optimization of Expensive Black-Box Functions.” *Journal of Global Optimization* 13 (4): 455–92.
- Keane, Andrew J, and Prasanth B Nair. 2005. *Computational Approaches for Aerospace Design: The Pursuit of Excellence*. Wiley.
- Kohavi, Ron. 1995. “A Study of Cross-Validation and Bootstrap for Accuracy Estimation and Model Selection.” *Proceedings of the 14th International Joint Conference on Artificial Intelligence - Volume 2* (San Francisco, CA, USA), IJCAI’95, 1137–43.
- Micchelli, Charles A. 1986. “Interpolation of Scattered Data: Distance Matrices and Conditionally Positive Definite Functions.” *Constructive Approximation* 2 (1): 11–22. <https://doi.org/10.1007/BF01893414>.

References

- Močkus, J. 1974. “On Bayesian Methods for Seeking the Extremum.” *Optimization Techniques IFIP Technical Conference*, 400–404.
- Morris, Max D., and Toby J. Mitchell. 1995. “Exploratory Designs for Computational Experiments.” *Journal of Statistical Planning and Inference* 43 (3): 381–402. [https://doi.org/https://doi.org/10.1016/0378-3758\(94\)00035-T](https://doi.org/https://doi.org/10.1016/0378-3758(94)00035-T).
- Poggio, T, and F Girosi. 1990. “Regularization Algorithms for Learning That Are Equivalent to Multilayer Networks.” *Science* 247 (4945): 978–82. <https://doi.org/10.1126/science.247.4945.978>.
- Raymer, Daniel P. 2006. *Aircraft Design: A Conceptual Approach*. AIAA.
- Rummel, R. J. 1976. *Understanding Correlation*. <https://www.hawaii.edu/powerkills/UC.HTM>.
- Santner, T J, B J Williams, and W I Notz. 2003. *The Design and Analysis of Computer Experiments*. Springer.
- Vapnik, V N. 1998. *Statistical learning theory*. Wiley; Wiley.
- Wang, Zhiqiang. 2007. “Two Postestimation Commands for Assessing Confounding Effects in Epidemiological Studies.” *The Stata Journal* 7 (2): 183–96.
- Wikipedia contributors. 2024. *Partial Correlation — Wikipedia, the Free Encyclopedia*. https://en.wikipedia.org/w/index.php?title=Partial_correlation&oldid=1253637419.